

53-1003628-01
31 March 2015

FastIron Ethernet Switch Platform and Layer 2 Switching

Configuration Guide

Supporting FastIron Software Release 08.0.30

BROCADE®

© 2015, Brocade Communications Systems, Inc. All Rights Reserved.

ADX, Brocade, Brocade Assurance, the B-wing symbol, DCX, Fabric OS, HyperEdge, ICX, MLX, MyBrocade, OpenScript, The Effortless Network, VCS, VDX, Vplane, and Vyatta are registered trademarks, and Fabric Vision and vADX are trademarks of Brocade Communications Systems, Inc., in the United States and/or in other countries. Other brands, products, or service names mentioned may be trademarks of others.

Notice: This document is for informational purposes only and does not set forth any warranty, expressed or implied, concerning any equipment, equipment feature, or service offered or to be offered by Brocade. Brocade reserves the right to make changes to this document at any time, without notice, and assumes no responsibility for its use. This informational document describes features that may not be currently available. Contact a Brocade sales office for information on feature and product availability. Export of technical data contained in this document may require an export license from the United States government.

The authors and Brocade Communications Systems, Inc. assume no liability or responsibility to any person or entity with respect to the accuracy of this document or any loss, cost, liability, or damages arising from the information contained herein or the computer programs that accompany it.

The product described by this document may contain open source software covered by the GNU General Public License or other open source license agreements. To find out which open source software is included in Brocade products, view the licensing terms applicable to the open source software, and obtain a copy of the programming source code, please visit <http://www.brocade.com/support/oscd>.

Contents

Preface.....	11
Document conventions.....	11
Text formatting conventions.....	11
Command syntax conventions.....	11
Notes, cautions, and warnings.....	12
Brocade resources.....	13
Contacting Brocade Technical Support.....	13
Document feedback.....	14
 About This Document.....	 15
What's new in this document.....	15
Supported Hardware.....	15
How command information is presented in this guide.....	16
 Basic Layer 2 Features	 17
About port regions.....	17
FastIron X Series device port regions.....	17
FCX device port regions.....	18
ICX 6610 device port regions.....	18
ICX 6430 device port regions.....	18
ICX 6450 device port regions.....	18
ICX 6650 device port regions.....	19
ICX 7450 device port regions.....	19
ICX 7750 device port regions.....	19
ICX 7250 device port regions.....	19
Enabling or disabling the Spanning Tree Protocol (STP).....	19
Modifying STP bridge and port parameters.....	19
Management MAC address for stackable devices.....	20
MAC learning rate control.....	20
Changing the MAC age time and disabling MAC address learning.....	20
Disabling the automatic learning of MAC addresses.....	21
Displaying the MAC address table.....	21
Static MAC entry configuration.....	22
Multi-port static MAC address.....	22
VLAN-based static MAC entries configuration.....	23
Configuring a VLAN to drop static MAC entries.....	24
Clearing MAC address entries.....	24
Flow-based MAC address learning.....	24
Flow-based learning overview.....	24
Flow-based learning configuration considerations.....	26
Configuring flow-based MAC address learning.....	26
Displaying information about flow-based MACs.....	27
Clearing flow-based MAC address entries.....	28
Enabling port-based VLANs.....	28
Assigning IEEE 802.1Q tagging to a port.....	29
Defining MAC address filters.....	29
Monitoring MAC address movement.....	30
Configuring the MAC address movement threshold rate.....	30

Viewing the MAC address movement threshold rate configuration.....	31
Configuring an interval for collecting MAC address move notifications.....	32
Viewing MAC address movement statistics for the interval history.....	32
SNMP MAC-notification trap support.....	33
Requirements and limitations for MAC-notification trap support.....	34
Configuring SNMP traps for MAC-notification	34
MAC-notification events.....	35
Displaying and modifying system parameter default settings.....	36
System default settings configuration considerations.....	37
Displaying system parameter default values.....	37
Modifying system parameter default values.....	42
Cut-through switching.....	43
Dynamic buffer allocation for QoS priorities for FastIron X Series devices.....	44
Default queue depth limits for FastIron X Series devices.....	44
Configuring the total transmit queue depth limit for FastIron X Series devices.....	45
Configuring the transmit queue depth limit for a given traffic class on FastIron X Series devices.....	45
Removing buffer allocation limits on FastIron X Series devices.....	46
Buffer profile configuration.....	46
Default queue depth limits for FastIron X Series devices.....	48
Dynamic buffer allocation for FCX and ICX devices.....	48
Configuring buffer profiles.....	48
Configuring buffer sharing on FCX and ICX devices.....	57
Removing buffer allocation limits on FCX and ICX.....	60
Buffer profiles for VoIP on FastIron stackable devices.....	61
Buffer allocation for ICX 6650.....	61
Buffer and descriptor maximum and default allocation values for ICX 6650.....	62
ICX 6650 buffer sharing levels.....	62
Displaying buffer sharing information on the ICX 6650.....	63
Ethernet loopback.....	63
Ethernet loopback operational modes.....	63
Ethernet loopback configuration considerations.....	64
Configuring Ethernet loopback in VLAN-unaware mode.....	65
Configuring Ethernet loopback in VLAN-aware mode.....	66
Ethernet loopback syslog messages.....	67
Remote Fault Notification on 1Gbps fiber connections.....	67
Enabling and disabling remote fault notification.....	67
Link Fault Signaling for 10Gbps Ethernet devices.....	68
Enabling Link Fault Signaling.....	68
Viewing the status of LFS-enabled links.....	68
Jumbo frame support.....	69
Packet InError Detection.....	69
Configuring Packet InError Detection.....	70
Syslog message for error-disabled port due to inError packets.....	70
Metro Features.....	71
Topology groups.....	71
Master VLAN and member VLANs.....	71
Control ports and free ports.....	72
Topology group configuration considerations.....	72
Configuring a topology group.....	72
Displaying topology group information.....	73

Metro Ring Protocol	74
Metro Ring Protocol configuration notes.....	76
MRP rings without shared interfaces (MRP Phase 1).....	76
MRP rings with shared interfaces (MRP Phase 2).....	77
Ring initialization.....	79
How ring breaks are detected and healed.....	83
Master VLANs and customer VLANs.....	85
Metro Ring Protocol configuration.....	87
Metro Ring Protocol diagnostics.....	89
Displaying MRP information.....	90
MRP CLI example.....	93
VSRP.....	95
VSRP configuration notes and feature limitations.....	97
Layer 2 and Layer 3 redundancy.....	97
Master election and failover.....	97
VSRP-aware security features.....	102
VSRP parameters.....	102
Configuring basic VSRP parameters.....	106
Configuring optional VSRP parameters.....	106
Displaying VSRP information.....	115
VSRP fast start.....	118
VSRP and MRP signaling.....	119
UDLD and Protected Link Groups.....	123
UDLD overview.....	123
UDLD for tagged ports.....	123
Configuration notes and feature limitations for UDLD.....	124
Enabling UDLD.....	124
Enabling UDLD for tagged ports.....	125
Changing the Keepalive interval.....	125
Changing the Keepalive retries.....	125
Displaying UDLD information.....	125
Clearing UDLD statistics.....	128
Protected link groups.....	128
About active ports.....	128
Using UDLD with protected link groups.....	128
UDLD with protected link groups configuration notes.....	129
Creating a protected link group and assigning an active port.....	130
Link Aggregation.....	133
Overview of link aggregation.....	133
LAG formation rules.....	134
Configuration notes for FastIron devices in a traditional stack.....	135
Maximum number of LAGs.....	137
Migrating from a previous release to 08.0.00a LAG or LACP configuration.	138
Downgrade considerations.....	139
LAG Load Sharing.....	139
LAG hashing on stacking products	140
Removing Layer 2 information from LAG hash output.....	140
Configuring a LAG.....	141
Creating a Link Aggregation Group (LAG).....	141
Creating a Link Aggregation Group (LAG) using the LAG ID option.	142
Deploying a LAG.....	146
Commands available under LAG once it is deployed.....	147
Disabling ports within a LAG.....	147
Enabling ports within a LAG.....	147

Adding a Port to Currently Deployed LAG.....	148
Deleting a Port from a Currently Deployed LAG.....	148
Monitoring an individual LAG port.....	148
Assigning a name to a port within a LAG.....	149
Enabling sFlow forwarding on a port in a LAG.....	150
Setting the sFlow sampling rate for a port in a LAG.....	150
IP assignment within a LAG.....	151
Renaming an existing LAG.....	151
Displaying LAG information.....	151
Displaying information about LAG interface.....	156
Enabling LAG hardware failover	156
Preboot eXecution Environment boot support.....	157
Enabling PXE boot support on a port.....	157
Multi-Chassis Trunking.....	159
Multi-Chassis Trunking Overview.....	159
How MCT works.....	159
MCT terminology.....	160
MCT data flow.....	161
MCT and VLANs.....	164
Cluster client automatic configuration.....	164
MCT feature interaction.....	165
Basic MCT configuration.....	166
Setting up cluster client automatic configuration	170
MCT failover scenarios.....	171
Layer 2 behavior with MCT.....	174
MAC operations.....	174
Port loop detection.....	178
MCT Layer 2 protocols.....	178
Protocol-based VLANs.....	179
Uplink switch.....	180
Layer 2 multicast snooping over MCT.....	180
Forcing a port up in a basic MCT configuration.....	183
Displaying MCT information.....	184
Displaying peer and client states.....	184
Displaying state machine information.....	185
Displaying cluster, peer, and client states.....	186
Displaying information about Ethernet interfaces.....	186
Displaying STP information.....	187
Displaying information for multicast snooping.....	187
MCT configuration examples	189
Single-level MCT example.....	189
Two-level MCT example.....	192
Multicast snooping configuration example.....	197
MCT configuration examples using STP	199
GVRP.....	205
GVRP overview.....	205
GVRP application examples.....	205
Dynamic core and fixed edge.....	206
Dynamic core and dynamic edge.....	207
Fixed core and dynamic edge.....	207
Fixed core and fixed edge.....	207
VLAN names created by GVRP.....	208
Configuration notes for GVRP.....	208
GVRP configuration.....	209

Changing the GVRP base VLAN ID.....	209
Increasing the maximum configurable value of the Leaveall timer... ..	210
Enabling GVRP.....	210
Disabling VLAN advertising.....	211
Disabling VLAN learning.....	211
Changing the GVRP timers.....	211
Converting a VLAN created by GVRP into a statically-configured VLAN.....	213
Displaying GVRP information.....	214
Displaying GVRP configuration information.....	214
Displaying GVRP VLAN information.....	217
Displaying GVRP statistics.....	219
Clearing GVRP statistics.....	220
GVRP CLI examples.....	221
Dynamic core and fixed edge.....	221
Dynamic core and dynamic edge.....	222
Fixed core and dynamic edge.....	222
Fixed core and fixed edge.....	223
Port mirroring and Monitoring.....	225
Port mirroring and monitoring overview.....	225
Port mirroring and monitoring configuration.....	225
Configuration notes for port mirroring and monitoring.....	226
Command syntax for port mirroring and monitoring.....	227
Mirroring configuration on a traditional stack.....	228
Configuration notes for traditional stack mirroring.....	228
ACL-based inbound mirroring.....	229
Creating an ACL-based inbound mirror clause for FCX and ICX devices.....	229
ACL-based inbound mirror clauses for FastIron X Series devices....	230
Destination mirror port	230
MAC address filter-based mirroring.....	233
Configuring MAC address filter-based mirroring.....	233
VLAN-based mirroring.....	234
Configuring VLAN-based mirroring.....	235
Displaying VLAN-based mirroring status.....	235
Configuration notes for VLAN-based mirroring.....	235
VLAN-based mirroring.....	236
Spanning Tree Protocol.....	241
STP overview.....	241
Standard STP parameter configuration.....	241
STP parameters and defaults.....	242
Enabling or disabling the Spanning Tree Protocol (STP).....	243
Changing STP bridge and port parameters.....	245
STP protection enhancement.....	246
Displaying STP information.....	248
STP feature configuration.....	255
Fast port span.....	255
Fast Uplink Span.....	257
802.1W Rapid Spanning Tree (RSTP).....	260
802.1W Draft 3.....	296
Single Spanning Tree (SSTP).....	300
STP per VLAN group.....	302
PVST/PVST+ compatibility.....	306
Overview of PVST and PVST+.....	307
VLAN tags and dual mode.....	308

Configuring PVST+ support.....	308
Displaying PVST+ support information.....	309
PVST+ configuration examples.....	310
PVRST compatibility.....	312
BPDU guard.....	312
Enabling BPDU protection by port.....	313
Re-enabling ports disabled by BPDU guard.....	313
Displaying the BPDU guard status.....	313
BPDU guard status example console messages	314
Root guard.....	315
Enabling STP root guard.....	315
Displaying the STP root guard.....	316
Displaying the root guard by VLAN.....	316
Designated Protection.....	316
Enabling Designated Protection on a port.....	317
Syslog message for a port in designated inconsistent state.....	317
Error disable recovery.....	317
Enabling an error-disabled port automatically.....	317
Enabling an error-disabled port manually.....	318
Setting the recovery interval.....	318
Displaying the error disable recovery state by interface	318
Displaying the recovery state for all conditions.....	319
Displaying the recovery state by port number and cause.....	319
Errdisable Syslog messages.....	319
802.1s Multiple Spanning Tree Protocol.....	319
Multiple spanning-tree regions	320
Configuration notes.....	321
Configuring MSTP mode and scope.....	321
Reduced occurrences of MSTP reconvergence.....	322
Configuring additional MSTP parameters.....	324
VLANs.....	335
VLAN overview.....	335
Types of VLANs.....	335
Modifying a port-based VLAN.....	341
Default VLAN.....	350
802.1Q tagging.....	351
Spanning Tree Protocol (STP).....	353
Virtual routing interfaces.....	354
VLAN and virtual routing interface groups.....	355
Dynamic, static, and excluded port membership.....	356
Super aggregated VLANs.....	359
Trunk group ports and VLAN membership.....	359
Summary of VLAN configuration rules.....	359
Routing between VLANs.....	360
Virtual routing interfaces (Layer 2 Switches only).....	360
Routing between VLANs using virtual routing interfaces (Layer 3 Switches only).....	361
Dynamic port assignment (Layer 2 Switches and Layer 3 Switches).....	362
Assigning a different VLAN ID to the default VLAN.....	362
Assigning different VLAN IDs to reserved VLANs 4091 and 4092.	362
Assigning trunk group ports.....	363
Enable spanning tree on a VLAN.....	363
Configuring IP subnet, IPX network and protocol-based VLANs.....	365
IP subnet, IPX network, and protocol-based VLAN configuration example.....	365

IP subnet, IPX network, and protocol-based VLANs within port-based VLANs.....	366
Configuring Layer 3 VLANs on FSX-A.....	367
Configuring Layer 3 VLANs on FSX-B.....	368
Configuring Layer 3 VLANs on FSX-C.....	369
IPv6 protocol VLAN configuration.....	369
Routing between VLANs using virtual routing interfaces (Layer 3 Switches only).....	370
Configuring Layer 3 VLANs and virtual routing interfaces on the FSX-A.....	371
Configuring Layer 3 VLANs and virtual routing interfaces for FSX-B.....	373
Configuring Layer 3 VLANs and virtual routing interfaces for FSX-C.....	374
Configuring protocol VLANs with dynamic ports.....	375
Aging of dynamic ports.....	375
Configuration guidelines for membership aging of dynamic VLAN ports.....	376
Configuring an IP, IPX, or AppleTalk Protocol VLAN with Dynamic Ports.....	377
Configuring an IP subnet VLAN with dynamic ports.....	377
Configuring an IPX network VLAN with dynamic ports.....	378
Configuring uplink ports within a port-based VLAN.....	379
Configuration considerations for uplink ports within a port-based VLAN.....	379
Configuration syntax for uplink ports within a port-based VLAN.....	379
IP subnet address on multiple port-based VLAN configuration.....	379
VLAN groups and virtual routing interface group	383
Configuring a VLAN group.....	384
Configuring a virtual routing interface group.....	385
Displaying the VLAN group and virtual routing interface group information.....	386
Allocating memory for more VLANs or virtual routing interfaces.....	387
Super-aggregated VLAN configuration.....	388
Configuration notes for aggregated VLANs.....	391
Configuring aggregated VLANs.....	391
Verifying the aggregated VLAN configuration.....	392
Complete CLI examples for aggregated VLANs.....	393
802.1ad tagging configuration.....	395
Configuration rules for 802.1ad tagging.....	396
Enabling 802.1ad tagging.....	397
Example 802.1ad configuration.....	397
Configuring 802.1ad tag profiles.....	398
Private VLAN configuration.....	399
Configuration notes for PVLANS and standard VLANs.....	403
Enabling broadcast or unknown unicast traffic to the PVLAN on FSX devices.....	406
CLI example for a general PVLAN network.....	407
CLI example for a PVLAN network with switch-switch link ports.....	407
Multiple promiscuous ports support in private VLANs	408
Dual-mode VLAN ports.....	409
Displaying VLAN information.....	412
Displaying VLANs in alphanumeric order.....	413
Displaying system-wide VLAN information.....	413
Displaying global VLAN information.....	414
Displaying VLAN information for specific ports.....	414
Displaying a port VLAN membership.....	414
Displaying a port dual-mode VLAN membership.....	415

Displaying port default VLAN IDs (PVIDs).....	415
Displaying PVLAN information.....	416

Preface

• Document conventions.....	11
• Brocade resources.....	13
• Contacting Brocade Technical Support.....	13
• Document feedback.....	14

Document conventions

The document conventions describe text formatting conventions, command syntax conventions, and important notice formats used in Brocade technical documentation.

Text formatting conventions

Text formatting conventions such as boldface, italic, or Courier font may be used in the flow of the text to highlight specific words or phrases.

Format	Description
bold text	Identifies command names
	Identifies keywords and operands
	Identifies the names of user-manipulated GUI elements
	Identifies text to enter at the GUI
<i>italic</i> text	Identifies emphasis
	Identifies variables
	Identifies document titles
Courier font	Identifies CLI output
	Identifies command syntax examples

Command syntax conventions

Bold and italic text identify command syntax components. Delimiters and operators define groupings of parameters and their logical relationships.

Convention	Description
bold text	Identifies command names, keywords, and command options.
<i>italic</i> text	Identifies a variable.
value	In Fibre Channel products, a fixed value provided as input to a command option is printed in plain text, for example, --show WWN.

Convention	Description
[]	Syntax components displayed within square brackets are optional. Default responses to system prompts are enclosed in square brackets.
{ x y z }	A choice of required parameters is enclosed in curly brackets separated by vertical bars. You must select one of the options. In Fibre Channel products, square brackets may be used instead for this purpose.
x y	A vertical bar separates mutually exclusive elements.
< >	Nonprinting characters, for example, passwords, are enclosed in angle brackets.
...	Repeat the previous element, for example, member[member...].
\	Indicates a “soft” line break in command examples. If a backslash separates two lines of a command input, enter the entire command at the prompt without the backslash.

Notes, cautions, and warnings

Notes, cautions, and warning statements may be used in this document. They are listed in the order of increasing severity of potential hazards.

NOTE

A Note provides a tip, guidance, or advice, emphasizes important information, or provides a reference to related information.

ATTENTION

An Attention statement indicates a stronger note, for example, to alert you when traffic might be interrupted or the device might reboot.



CAUTION

A Caution statement alerts you to situations that can be potentially hazardous to you or cause damage to hardware, firmware, software, or data.



DANGER

A Danger statement indicates conditions or situations that can be potentially lethal or extremely hazardous to you. Safety labels are also attached directly to products to warn of these conditions or situations.

Brocade resources

Visit the Brocade website to locate related documentation for your product and additional Brocade resources.

You can download additional publications supporting your product at www.brocade.com. Select the Brocade Products tab to locate your product, then click the Brocade product name or image to open the individual product page. The user manuals are available in the resources module at the bottom of the page under the Documentation category.

To get up-to-the-minute information on Brocade products and resources, go to [MyBrocade](#). You can register at no cost to obtain a user ID and password.

Release notes are available on [MyBrocade](#) under Product Downloads.

White papers, online demonstrations, and data sheets are available through the [Brocade website](#).

Contacting Brocade Technical Support

As a Brocade customer, you can contact Brocade Technical Support 24x7 online, by telephone, or by e-mail. Brocade OEM customers contact their OEM/Solutions provider.

Brocade customers

For product support information and the latest information on contacting the Technical Assistance Center, go to <http://www.brocade.com/services-support/index.html>.

If you have purchased Brocade product support directly from Brocade, use one of the following methods to contact the Brocade Technical Assistance Center 24x7.

Online	Telephone	E-mail
<p>Preferred method of contact for non-urgent issues:</p> <ul style="list-style-type: none"> • My Cases through MyBrocade • Software downloads and licensing tools • Knowledge Base 	<p>Required for Sev 1-Critical and Sev 2-High issues:</p> <ul style="list-style-type: none"> • Continental US: 1-800-752-8061 • Europe, Middle East, Africa, and Asia Pacific: +800-AT FIBREE (+800 28 34 27 33) • For areas unable to access toll free number: +1-408-333-6061 • Toll-free numbers are available in many countries. 	<p>support@brocade.com</p> <p>Please include:</p> <ul style="list-style-type: none"> • Problem summary • Serial number • Installation details • Environment description

Brocade OEM customers

If you have purchased Brocade product support from a Brocade OEM/Solution Provider, contact your OEM/Solution Provider for all of your product support needs.

- OEM/Solution Providers are trained and certified by Brocade to support Brocade® products.
- Brocade provides backline support for issues that cannot be resolved by the OEM/Solution Provider.

- Brocade Supplemental Support augments your existing OEM support contract, providing direct access to Brocade expertise. For more information, contact Brocade or your OEM.
- For questions regarding service levels and response times, contact your OEM/Solution Provider.

Document feedback

To send feedback and report errors in the documentation you can use the feedback form posted with the document or you can e-mail the documentation team.

Quality is our first concern at Brocade and we have made every effort to ensure the accuracy and completeness of this document. However, if you find an error or an omission, or you think that a topic needs further development, we want to hear from you. You can provide feedback in two ways:

- Through the online feedback form in the HTML documents posted on www.brocade.com.
- By sending your feedback to documentation@brocade.com.

Provide the publication title, part number, and as much detail as possible, including the topic heading and page number if applicable, as well as your suggestions for improvement.

About This Document

• What's new in this document	15
• Supported Hardware	15
• How command information is presented in this guide	16

What's new in this document

The following table includes descriptions of the new information added to this guide for the FastIron 08.0.30 release.

TABLE 1 Summary of Enhancements in FastIron release 08.0.30

Feature	Description	Location
Ethernet loopback	Ethernet loopback provides support to loopback the incoming traffic to the source device.	Refer to Ethernet loopback on page 63
Renaming an existing LAG	The name of an existing LAG can be changed without causing any impact on the functionality of the LAG.	Refer to Renaming an existing LAG on page 151
Displaying LAG interface details	Details of the LAG interface including counters for each port of the LAG are displayed. The details of a LAG can be viewed by specifying the LAG name or LAG ID.	Refer to Displaying information about LAG interface on page 156
Maximum number of ports in a LAG	On Brocade ICX 7750, Brocade ICX 7450, and Brocade ICX 7250 devices, the maximum number of ports supported in a LAG is increased to 16.	Refer to the chapter <i>Link Aggregation</i> .

Supported Hardware

This guide supports the following product families from Brocade:

- FCX Series
- FastIron X Series (FSX 800 and FSX 1600)
- ICX 6610 Series
- ICX 6430 Series (ICX 6430, ICX 6430-C12)
- ICX 6450 Series (ICX 6450, ICX 6450-C12-PD)
- ICX 6650 Series
- ICX 7750 Series
- ICX 7450 Series
- ICX 7250 Series

NOTE

The Brocade ICX 6430-C switch supports the same feature set as the Brocade ICX 6430 switch unless otherwise noted.

NOTE

The Brocade ICX 6450-C12-PD switch supports the same feature set as the Brocade ICX 6450 switch unless otherwise noted.

For information about the specific models and modules supported in a product family, refer to the hardware installation guide for that product family.

How command information is presented in this guide

For all new content supported in FastIron Release 08.0.20 and later, command information is documented in a standalone command reference guide.

In an effort to provide consistent command line interface (CLI) documentation for all products, Brocade is in the process of completing a standalone command reference for the FastIron platforms. This process involves separating command syntax and parameter descriptions from configuration tasks. Until this process is completed, command information is presented in two ways:

- For all new content supported in FastIron Release 08.0.20 and later, the CLI is documented in separate command pages included in the *FastIron Command Reference*. Command pages are compiled in alphabetical order and follow a standard format to present syntax, parameters, usage guidelines, examples, and command history.

NOTE

Many commands from previous FastIron releases are also included in the command reference.

- Legacy content in configuration guides continues to include command syntax and parameter descriptions in the chapters where the features are documented.

If you do not find command syntax information embedded in a configuration task, refer to the *FastIron Command Reference*.

Basic Layer 2 Features

• About port regions.....	17
• Enabling or disabling the Spanning Tree Protocol (STP).....	19
• Management MAC address for stackable devices.....	20
• MAC learning rate control.....	20
• Changing the MAC age time and disabling MAC address learning.....	20
• Static MAC entry configuration.....	22
• VLAN-based static MAC entries configuration.....	23
• Clearing MAC address entries.....	24
• Flow-based MAC address learning.....	24
• Enabling port-based VLANs.....	28
• Defining MAC address filters.....	29
• Monitoring MAC address movement.....	30
• SNMP MAC-notification trap support.....	33
• Displaying and modifying system parameter default settings.....	36
• Cut-through switching.....	43
• Dynamic buffer allocation for QoS priorities for FastIron X Series devices.....	44
• Dynamic buffer allocation for FCX and ICX devices.....	48
• Buffer allocation for ICX 6650.....	61
• Ethernet loopback.....	63
• Remote Fault Notification on 1Gbps fiber connections.....	67
• Link Fault Signaling for 10Gbps Ethernet devices.....	68
• Jumbo frame support.....	69
• Packet InError Detection.....	69

About port regions

This section describes port regions on FastIron switches.

FastIron X Series device port regions

Ports on the FastIron X Series devices are grouped into regions. For a few features, you will need to know the region to which a port belongs. However, for most features, a port region does not affect configuration or operation of the feature.

NOTE

Port regions do not apply to trunk group configurations on the FastIron X Series devices. However, port regions do apply to port monitoring and unknown unicast configurations on FastIron X Series devices.

- SX-FI48GPP: ports 1 - 48 belong to port region 0
- SX-FI-2XG: ports 1 - 2 belong to port region 0
- SX-FI-8XG: ports 1 - 8 belong to port region 0

- SX-FI-24HF: ports 1 - 24 belong to port region 0
- SX-FI-24GPP: ports 1 - 24 belong to port region 0

FCX device port regions

The port region rules for FCX devices are as follows:

- For all platforms, a 24-port Gbps module has one port region. In addition, any 10 Gbps ports on the device also belong to this single port region.
- For all platforms, the 48-port Gbps module has two port regions:
 - Ports 1 - 24 belong to port region 0
 - Ports 25 - 48 belong to port region 1
- For FCX648 devices with two 10 Gbps XFP ports, and a two 10 Gbps CX4 stacking ports:
 - The two 10 Gbps XFP ports belong to port region 0 (along with ports 1 -24)
 - The two 10 Gbps CX4 stacking ports belong to port region 1 (along with ports 25 - 48)
- For FCX648 devices with four 10 Gbps SFP+ ports:
 - 10 Gbps SFP+ ports 3 and 4 belong to port region 0 (along with ports 1 -24)
 - 10 Gbps SFP+ ports 1 and 2 ports belong to port region 1 (along with ports 25 - 48)

ICX 6610 device port regions

- A 24-port Gbps has the following port regions:
 - Ports 1-24 belong to region 0
 - 10 Gbps SFP+ ports belong to region 0
- A 48-port Gbps has the following port regions:
 - Ports 1-24 belong to region 0
 - Ports 25-48 belong to region 1
 - 10 Gbps SFP+ ports belong to region 0

ICX 6430 device port regions

- A 24-port Gbps module has one port region. The four SFP ports on the device also belong to this single port region.
- A 48-port Gbps module has two port regions:
 - Ports 1-24 and SFP ports 1& 2 belong to port region 0
 - Ports 25-48 and SFP ports 3 & 4 belong to port region 1

ICX 6450 device port regions

- A 24-port Gbps module has one port region. The four SFP+ ports on the device also belong to this single port region.
- A 48-port Gbps module has two port regions:
 - Ports 1-24 and SFP+ ports 1& 2 belong to port region 0
 - Ports 25-48 and SFP+ ports 3 & 4 belong to port region 1

ICX 6650 device port regions

ICX 6650 device has only one port region.

Brocade ICX 6650 has only one port region. All ports belong to region 0.

ICX 7450 device port regions

Brocade ICX 7450 24 port has only one port region.

Brocade ICX 7450 48 port has two port regions.

ICX 7750 device port regions

ICX 7750 device has only one port region.

Brocade ICX 7750 has only one port region. All ports belong to region 0.

ICX 7250 device port regions

ICX 7250 device has only one port region. All ports belong to region 0.

Enabling or disabling the Spanning Tree Protocol (STP)

STP (IEEE 802.1D bridge protocol) is supported on all Brocade devices. STP detects and eliminates logical loops in the network. STP also ensures that the least cost path is taken when multiple paths exist between ports or VLANs. If the selected path fails, STP searches for and then establishes an alternate path to prevent or limit retransmission of data.

NOTE

This section provides instructions for enabling and disabling STP. For configuration procedures and more information about STP, refer to [Spanning Tree Protocol](#) on page 241 in this guide.

STP must be enabled at the system level to allow assignment of this capability on the VLAN level. On devices running Layer 2 code, STP is enabled by default. On devices running Layer 3 code, STP is disabled by default.

To enable STP for all ports on a Brocade device, enter the following command.

```
device(config)#spanning tree
```

Syntax: [no] spanning-tree

You can also enable and disable spanning tree on a port-based VLAN and on an individual port basis, and enable advanced STP features. Refer to [Spanning Tree Protocol](#) on page 241.

Modifying STP bridge and port parameters

You can modify the following STP Parameters:

- Bridge parameters - forward delay, maximum age, hello time, and priority
- Port parameters - priority and path cost

For configuration details, refer to [Changing STP bridge and port parameters](#) on page 245.

Management MAC address for stackable devices

In a traditional stack, the management MAC address of the Active Controller is always used as the STP bridge ID. The Active Controller management MAC address is always used for control protocols for the following reasons:

- Unlike standalone devices, each stack member has a different range of MAC addresses.
- In a stack, the management MAC address is software generated, and is always the MAC address of the Active Controller first port. This ensures consistency across the stack during resets, assuming that the Active Controller is always the same unit.
- This helps avoid the disruption of frequent topology changes in the stack.

For more information about stacking and Brocade stackable devices, refer to the *FastIron Ethernet Switch Stacking Configuration Guide*.

MAC learning rate control

NOTE

The MAC learning rate control feature is not supported on ICX 7750, ICX 7450 and ICX 7250 devices.

You can set a rate limit to control CPU address updating. The range for this rate limit is 200 to 50,000 per second. The MAC learning rate limit applies to each packet processor, which means that for a system with two packet processors, each processor can send address messages to the CPU at the established rate limit.

Syntax: [no] **cpu-limit addr-msgs msgsRateLimit**

NOTE

Actual rates in hardware may have a variance of +200 or -100.

Changing the MAC age time and disabling MAC address learning

To change the MAC address age timer, enter a command such as the following.

```
device(config) #mac-age-time 60
```

Syntax: [no] **mac-age-time secs**

secs specifies the number of seconds. Possible values differ depending on the version of software running on your device, as follows:

- On Brocade FCX Series devices, you can configure the MAC address age timer to 0 or a value from 10-1000 (seconds). If you set the MAC age time to 0, aging is disabled.
- On FastIron X Series devices, you can configure 0 or a value from 10 - 86,400 (seconds), in 10-second intervals. If you set the MAC age time to 0, aging is disabled.
- On ICX Series devices, you can configure the MAC address age timer to 0 or a value from 10-570 (seconds). If you set the MAC age time to 0, aging is disabled.
- If the total MAC addresses in the system is more than 16000, Brocade recommends a MAC age timer greater than 60 seconds. If the total MAC addresses in the system is more than 64000, Brocade recommends a MAC age timer greater than 120 seconds.

On ICX 6650, you can configure a value from 0 or a value from 10-600. The default is 300. If you set the MAC age time to 0, aging is disabled. On ICX 7450 and ICX 7750 you can configure a value of 60-86400.

NOTE

Usually, the actual MAC age time is from one to two times the configured value. For example, if you set the MAC age timer to 60 seconds, learned MAC entries age out after remaining unused for between 60 - 120 seconds. However, if all of the following conditions are met, then the MAC entries age out after a longer than expected duration:

- The MAC age timer is greater than 630 seconds.
- The number of MAC entries is over 6000.
- All MAC entries are learned from the same packet processor.
- All MAC entries age out at the same time.

Disabling the automatic learning of MAC addresses

By default, when a packet with an unknown Source MAC address is received on a port, the Brocade device learns this MAC address on the port.

You can prevent a physical port from learning MAC addresses by entering the following command.

```
device(config)#interface ethernet 3/1
device(config-if-e1000-3/1)#mac-learn-disable
```

Syntax: [no] mac-learn disable

Use the no form of the command to allow a physical port to learn MAC addresses.

MAC address learning configuration notes and feature limitations

- This command is not available on virtual routing interfaces. Also, if this command is configured on the primary port of a trunk, MAC address learning will be disabled on all the ports in the trunk.
- Entering the **mac-learn-disable** command on tagged ports disables MAC learning for that port in all VLANs to which that port is a member. For example, if tagged port 3/1 is a member of VLAN 10, 20, and 30 and you issue the **mac-learn-disable** command on port 3/1, port 3/1 will not learn MAC addresses, even if it is a member of VLAN 10, 20, and 30.

Displaying the MAC address table

To display the MAC table, enter the show mac-address command.

```
device#show mac-address
```

```
Total active entries from all ports = 3
Total static entries from all ports = 1
  MAC-Address   Port   Type   VLAN
0000.0034.1234    15   Static    1
0000.0038.2f24    14   Dynamic   1
0000.0038.2f00    13   Dynamic   1
0000.0086.b159    10   Dynamic   1
```

In the output of the **show mac-address** command, the *Type* column indicates whether the MAC entry is static or dynamic. A static entry is one you create using the **static-mac-address** command. A dynamic entry is one that is learned by the software from network traffic.

The output of the **show mac-address** command on FSX devices include an *Index* column which indicates the index where the entry exists in the hardware MAC table.

NOTE

The **show mac-address** command output does not include MAC addresses for management ports, since these ports do not support typical MAC learning and MAC-based forwarding.

Static MAC entry configuration

Static MAC addresses can be assigned to Brocade devices.

NOTE

Brocade devices running Layer 3 code also support the assignment of static IP Routes, static ARP, and static RARP entries. For details on configuring these types of static entries, refer to "Adding a static IP route" and "Adding a static ARP entry" sections in the *FastIron Ethernet Switch Layer 3 Routing Configuration Guide*.

You can manually input the MAC address of a device to prevent it from being aged out of the system address table.

This option can be used to prevent traffic for a specific device, such as a server, from flooding the network with traffic when it is down. Additionally, the static MAC address entry is used to assign higher priorities to specific MAC addresses.

You can specify traffic priority (QoS) and VLAN membership (VLAN ID) for the MAC Address as well as specify the device type of either router or host.

The default and maximum configurable MAC table sizes can differ depending on the device. To determine the default and maximum MAC table sizes for your device, display the system parameter values. Refer to [Displaying and modifying system parameter default settings](#) on page 36.

Multi-port static MAC address

Many applications, such as Microsoft NLB, Juniper IPS, and Netscreen Firewall, use the same MAC address to announce load-balancing services. As a result, a switch must be able to learn the same MAC address on several ports. Multi-port static MAC allows you to statically configure a MAC address on multiple ports using a single command.

Multi-port static MAC address configuration notes

- This feature is applicable for Layer 2 traffic.
- This feature can be used to configure unicast as well as IPv4 and IPv6 multicast MAC addresses on one or more ports. However, when a multicast MAC address is configured, the corresponding MAC address entry cannot be used for IGMP snooping. For IPv4 multicast addresses (range 0100.5e00.0000 to 0100.5e7f.ffff) and IPv6 multicast addresses (range 3333.0000.0000 to 3333.ffff.ffff), use IGMP/MLD snooping. Other multicast addresses can also be configured on the ports using this feature.
- FastIron devices support a maximum of 15 multi-port static MAC addresses.
- Hosts or physical interfaces normally join multicast groups dynamically, but you can also statically configure a host or an interface to join a multicast group.

Configuring a multi-port static MAC address

For example, to add a static entry for a server with a MAC address of 0000.0063.67ff and a priority of 7, enter the following command. If the system has only default VLAN, the command has to be issued from the global configuration mode.

```
device(config)#static-mac-address 0000.0063.67ff ethernet 4/2 ethernet 4/3 ethernet  
4/4 priority 7
```

If the system has multiple VLANs, the command has to be issued from the VLAN configuration mode.

```
device(config-vlan-30)#static-mac-address 0000.0063.67ff ethernet 1/1/1
```

To specify a range of ports, enter the following command.

```
device(config)#static-mac-address 0000.0063.67ff ethernet 4/2 to 4/6 priority 7
```

Syntax: [no] static-mac-address mac-addr ethernet [slotnum/]portnum ethernet [slotnum/]portnum
ethernet [slotnum/]portnum....[priority num]

or

Syntax: [no] static-mac-address mac-addr ethernet [slotnum/]portnum to ethernet
[slotnum/]portnum [prioritynum]

The **slotnum** parameter is required on chassis devices.

The **portnum** parameter is a valid port number.

The **priority num** is optional and can be a value from 0 - 7 (0 is lowest priority and 7 is highest priority). The default priority is 0.

VLAN-based static MAC entries configuration

You can configure a VLAN to drop packets that have a particular source or destination MAC address.

You can configure a maximum of 2048 static MAC address drop entries on a Brocade device.

Use the CLI command **show running-config** to view the static MAC address drop entries currently configured on the device.

Configuring a VLAN to drop static MAC entries

To configure a VLAN to drop packets with a source or destination MAC address of 0000.0063.67FF, enter the following commands.

```
device(config)#vlan 2  
device(config-vlan-2)#static-mac-address 0000.0063.67FF drop
```

Syntax: [no] static-mac-address *mac-addr* drop

Use the **no** form of the command to remove the static MAC address drop configuration.

Clearing MAC address entries

You can remove learned MAC address entries from the MAC address table. The types of MAC address that can be removed are as follows:

- All MAC address entries
- All MAC address entries for a specified Ethernet port
- All MAC address entries for a specified VLAN
- All specified MAC address entry in all VLANs

For example, to remove entries for the MAC address 0000.0080.00d0 in all VLANs, enter the following command at the Privilege EXEC level of the CLI.

```
device#clear mac-address 0000.0080.00d0
```

Syntax: clear mac-address { *mac-address* | ethernet *port-num* | vlan *vlan-num* }

If you enter **clear mac-address** without any parameter, the software removes all MAC address entries.

Use the *mac-address* parameter to remove a specific MAC address from all VLANs. Specify the MAC address in the following format: HHHH.HHHH.HHHH.

Use the **ethernet** *port-num* parameter to remove all MAC addresses for a specific Ethernet port.

Use the **vlan** *vlan-num* parameter to remove all MAC addresses for a specific VLAN.

Flow-based MAC address learning

NOTE

Flow-based MAC address learning is only supported on FastIron X Series devices.

This section describes flow-based MAC address learning and how to enable it on a FastIron X Series switch.

Flow-based learning overview

With regular MAC address learning, when a new MAC address is learned, it is programmed in the same location (hardware index) in all packet processors in a FastIron Layer 2 or Layer 3 switch. There are multiple packet processors (one per port region) in a compact switch, and in each module in a

chassis-based switch. With regular MAC address learning, MAC addresses are *global*, meaning the hardware MAC table is identical across all packet processors.

With the introduction of **flow-based MAC address learning**, when a new source MAC address is learned, it is programmed only in the source packet processor (the processor that received the packet). The destination MAC address gets added to other packet processors on demand, whenever a traffic flow that needs it is detected. With flow-based MAC address learning, the MAC address is programmed in different hardware locations and the hardware MAC table is different across all packet processors.

The benefits of flow-based learning

With global MAC address learning, all MAC addresses are programmed in all packet processors, even though they may not be required and are never used by all packet processors. Global MAC address learning wastes some space in the hardware MAC table and limits the number of supported MAC addresses to 16K.

With flow-based MAC address learning, MAC addresses are learned and programmed selectively, only in the packet processors that need them. Since the MAC addresses are distributed across several packet processors, flow-based learning frees up space in the hardware MAC table and increases the number of supported MAC addresses from 16K to 32K.

How flow-based learning works

When a packet processor, for example, *PP 1*, receives an incoming packet with source *MAC address X*, it sends a *new address* message to the CPU. The system learns *MAC address X* by adding it to the software MAC table in the CPU, then programming it in the hardware MAC table in the source packet processor, in this case *PP 1*. If the MAC address is learned on a trunk port, the MAC address is also programmed on all of the packet processors that have ports in the same trunk group.

When another packet processor, let call it *PP 2*, receives an incoming packet and the packet destination MAC address matches *source MAC address X*, it floods the packet in hardware as an unknown unicast packet and copies the packet to the CPU. The system locates the MAC address in the software MAC table, then programs the MAC address in the hardware MAC table in *PP 2*. If the MAC address is learned on a trunk port, the MAC address is also programmed on all of the packet processors that have ports in the same trunk group. Once the MAC address is programmed in hardware, subsequent packets with this destination MAC are forwarded as known unicast packets and are not copied to the CPU.

Flow-based MAC addresses are aged out by the source packet processor according to the MAC age time learned on the local port. Furthermore, when a flow-based MAC address is aged out from the source packet processor, it is also aged out from all other packet processors on which the address is programmed. In the above example, when MAC address *X* is aged out from *PP 1*, it is also aged out from *PP 2*.

NOTE

Even when flow-based MAC address learning is enabled, some MAC addresses, including but not limited to control MACs, static MACs, multicast MACs, and MAC addresses resolved through ARP, will continue to be global MAC addresses. These MAC addresses are always programmed in all packet processors in a Layer 2 or Layer 3 switch.

NOTE

Global MAC addresses have priority over dynamic flow-based MAC addresses. To ensure that global MAC addresses are in sync across all packet processors, flow-based MAC addresses may be overwritten in one or more packet processors. The MAC addresses will be relearned and reprogrammed using the flow-based method as needed by incoming traffic flows.

Flow-based learning configuration considerations

When configuring flow-based MAC learning, consider the rules and limitations in this section.

- Flow-based MAC learning is not supported with the following features:
 - Disabling the automatic learning of MAC addresses (CLI command **mac-learn-disable**).
 - Globally disabling Layer 2 switching (CLI command route-only)
- Configuring flow-based mac address learning and Multi-Chassis Trunking on the same device is not supported.
- When flow-based MAC learning is enabled, unknown unicast packets are copied to the CPU. Therefore, flow-based MAC learning should not be enabled if a continuous high rate of unknown unicast packet flooding is expected, as this will cause high CPU utilization.
- Unknown unicast flooding can occur for a known destination MAC address, if the system fails to program that destination MAC address because the hardware MAC table or hash bucket is full. This condition can also lead to high CPU utilization.
- A source MAC address is learned only on the ingress (source) packet processor. The MAC address is added to other packet processors as needed by their incoming traffic flows. During a brief period until the destination MAC address is successfully added to the hardware MAC table, unknown unicast flooding is expected on the VLAN.
- When a flow-based MAC address moves, it is deleted from all of the packet processors, then relearned on each packet processor individually, as needed by incoming traffic flows.
- The software MAC address table in the CPU uses a hashing algorithm. Because hash collisions can occur and may consume software resources, the FastIron may not be able to support up to 32K MAC addresses.
- The system can scale up to 32K MAC addresses, however, each packet processor is limited to a maximum of 16K MAC addresses. This limit still applies, as this is a hardware limitation.
- On SXL with authentication protocols and high traffic, the clear mac and mac flush operations log a lot of new address messages resulting in high CPU utilization for few minutes.

Configuring flow-based MAC address learning

To configure flow-based MAC address learning, simply enable it globally. If necessary, increase the capacity of the MAC address table as well.

Enabling flow-based MAC address learning

To enable flow-based MAC address learning, enter the following command at the Global CONFIG level of the CLI.

```
device (config) #mac-learning-flow-based
```

This command enables flow-based MAC address learning. All dynamically-learned MAC addresses are flushed from the hardware and software MAC tables and are subsequently learned using flow-based MAC address learning.

Syntax: **[no] mac-learning-flow-based**

Use the **no** form of the command to disable flow-based MAC address learning. When disabled, all dynamically-learned MAC addresses are flushed from the hardware and software MAC tables and are subsequently learned using global MAC address learning.

Increasing the capacity of the MAC address table (optional)

After enabling support for flow-based MACs, you can increase the capacity of the MAC address table of up to 32K MAC addresses. By default, up to 32K MAC addresses are supported.

NOTE

FCX devices do not support flow-based MACs and the capacity of the MAC address table cannot be altered on FCX devices.

NOTE

On ICX 6650, you can increase the capacity of the MAC address table of up to 64K MAC addresses. By default, up to 64K MAC addresses are supported.

To increase the capacity of the MAC table, enter commands such as the following.

```
device(config)#system-max mac 32768
device(config)#write memory
device(config)#exit
device#reload
```

NOTE

You must save the configuration and reload the software to place the **system-max mac** change into effect.

Syntax: system-max mac max-flow-MACs

The *max-flow-MACs* parameter specifies the maximum number of MAC addresses in the MAC table. For flow-based MACs, the minimum value is 16K and the default value is 32K (16k for FSX).

Use the command **show default values** to display the default, maximum, and currently configured values for the MAC address table.

Displaying information about flow-based MACs

The **show mac-address** command includes information related to flow-based MAC address learning. The following shows an example show mac output.

```
device# show mac
Total active entries from all ports = 15
MAC-Address      Port          Type      Index
0000.0000.0001   1/1          Dynamic   NA
0000.0000.0002   1/1          Dynamic   NA
```

In the above example, since both MAC address entries are flow-based and are located on different packet processors (hardware index), the **Index** field displays **NA** (not applicable).

Syntax: show mac

To display all of the packet processors that have a particular flow-based MAC address, use the **show mac-address vlan** command.

```
device#show mac-address vlan 1 0000.0000.0001
Total active entries from all ports = 16
MAC-Address      Port          Type      Index
0000.0000.0001   1/1          Dynamic   NA
Present in following devices (at hw index) :-
0 (8196 )        4 (8196 )
```

In the above example, the MAC address 0000.0000.0001 is programmed in packet processors 0 and 4, and the hardware index is 8196.

Syntax: **show mac-address vlan** *vlan-num mac-address*

Clearing flow-based MAC address entries

To remove dynamically-learned MAC addresses from the MAC table, use the CLI command **clear mac**. This command clears all dynamically-learned MACs from the hardware and software MAC tables.

Enabling port-based VLANs

When using the CLI, port and protocol-based VLANs are created by entering one of the following commands at the global CONFIG level of the CLI.

To create a port-based VLAN, enter commands such as the following.

```
device(config)#vlan 222 by port  
device(config)#vlan 222 name Mktg
```

Syntax: **vlan** *num* **by port**

Syntax: **vlan** *num* **name** *string*

The num parameter specifies the VLAN ID. The valid range for VLAN IDs starts at 1 on all systems but the upper limit of the range differs depending on the device. In addition, you can change the upper limit on some devices using the **system max-vlans...** command.

The string parameter is the VLAN name and can be a string up to 32 characters. You can use blank spaces in the name if you enclose the name in double quotes (for example, "Product Marketing").

Depending on device support, you can configure up to 4000 port-based VLANs. Each port-based VLAN can contain either tagged or untagged ports. A port cannot be a member of more than one port-based VLAN unless the port is tagged. On both device types, valid VLAN IDs are 1 - 4095. You can configure up to the maximum number of VLANs within that ID range.

NOTE

VLAN IDs 4087, 4090, and 4093 are reserved for Brocade internal use only. VLAN 4094 is reserved for use by Single STP. Also, if you are running an earlier release, VLAN IDs 4091 and 4092 may be reserved for Brocade internal use only. If you want to use VLANs 4091 and 4092 as configurable VLANs, you can assign them to different VLAN IDs. For more information, refer to [Assigning different VLAN IDs to reserved VLANs 4091 and 4092](#) on page 362

NOTE

The second command is optional and also creates the VLAN if the VLAN does not already exist. You can enter the first command after you enter the second command if you first exit to the global CONFIG level of the CLI.

Assigning IEEE 802.1Q tagging to a port

When a port is tagged, it allows communication among the different VLANs to which it is assigned. A common use for this might be to place an email server that multiple groups may need access to on a tagged port, which in turn, is resident in all VLANs that need access to the server.

NOTE

Tagging does not apply to the default VLAN.

When using the CLI, ports are defined as either tagged or untagged at the VLAN level.

Command syntax for assigning 802.1Q tagging to a port

Suppose you want to make port 5 a member of port-based VLAN 4, a tagged port. To do so, enter the following.

```
device(config)#vlan 4
device(config-vlan-4)#tagged e 5
```

Syntax: `tagged ethernet [stack[slot]port] portnum [to [stack[port] portnum [ethernet [slotnum/] portnum...]]`

The *slotnum* parameter is required on chassis devices.

Defining MAC address filters

MAC layer filtering enables you to build access lists based on MAC layer headers in the Ethernet/IEEE 802.3 frame. You can filter on the source and destination MAC addresses. The filters apply to incoming traffic only.

You configure MAC address filters globally, then apply them to individual interfaces. To apply MAC address filters to an interface, you add the filters to that interface MAC address filter group.

The device takes the action associated with the first matching filter. If the packet does not match any of the filters in the access list, the default action is to drop the packet. If you want the system to permit traffic by default, you must specifically indicate this by making the last entry in the access list a permit filter. An example is given below.

Syntax: `mac filter last-index-number permit any any`

For devices running Layer 3 code, the MAC address filter is applied to all inbound Ethernet packets, including routed traffic. This includes those port associated with a virtual routing interface. However, the filter is not applied to the virtual routing interface. It is applied to the physical port.

When you create a MAC address filter, it takes effect immediately. You do not need to reset the system. However, you do need to save the configuration to flash memory to retain the filters across system resets.

Monitoring MAC address movement

MAC address movement notification allows you to monitor the movement of MAC addresses that migrate from port to port. It enables you to distinguish between legitimate movement and malicious movement by allowing you to define malicious use as a threshold number of times a MAC address moves within a specific interval.

Malicious use typically involves many MAC address moves, while legitimate use usually involves a single move. Malicious movement is often the result of MAC address spoofing, in which a malicious user masquerades as a legitimate user by changing his own MAC address to that of a legitimate user. As a result, the MAC address moves back and forth between the ports where the legitimate and malicious users are connected. A legitimate use might be to spoof the MAC address of a failed device in order to continue access using a different device.

You can monitor MAC address movements in the following ways:

- Threshold-rate notifications allow you to configure the maximum number of movements over a specified interval for each MAC address before a notification is sent. For example you could define the malicious move rate as three moves every 30 seconds.
- Interval-history notifications are best suited for a statistical analysis of the number of MAC address movements for a configured time interval. For example, you may want to find out how many MAC addresses have moved in the system over a given interval or how many times a specific MAC address has moved during that interval. However, it is not possible to get this information for every MAC address if there are a lot of MAC addresses that moved during the interval. Consequently, the number of MAC addresses that can have a recorded history is limited.

NOTE

MAC address move notification does not detect MAC movements across an MCT cluster between MCT peers. It only detects MAC movements locally within a cluster MCT peer.

Configuring the MAC address movement threshold rate

To enable notification of MAC address moves, enter the **mac-movement notification threshold-rate** command at the global configuration level. This command enables a corresponding SNMP trap. Notification is triggered when a threshold number of MAC address moves occurs within a specified period for the same MAC address. This command sets the threshold level and the sampling interval.

Avoid threshold rates and sampling intervals that are too small. If you choose a small threshold and a sampling interval that is also small, an unnecessarily high number of traps could occur.

The following example enables notification of MAC address moves and sends an SNMP trap when any MAC address moves to a different port five times in a 10-second interval.

```
device(config)# mac-movement notification threshold-rate 5 sampling-interval 10
```

To disable notification of MAC address moves and disable the SNMP trap, use the **no** form of the command, as shown in the following example.

```
device(config)# no mac-movement notification threshold-rate 5 sampling-interval 10
```

Syntax: [no] mac-movement notification threshold-rate *move-count* sampling-interval *interval*

The *move-count* variable indicates the number of times a MAC address can move within the specified period until an SNMP trap is sent. It has no default value.

The *interval* variable specifies the sampling period in seconds. It has no default value.

Viewing the MAC address movement threshold rate configuration

To display the configuration of the MAC address movement threshold rate, enter the **show notification mac-movement threshold-rate** command at the privileged EXEC level. This command also displays ongoing statistics for the current sampling interval.

```
device# show notification mac-movement threshold-rate
Threshold-Rate Mac Movement Notification is ENABLED
Configured Threshold-Rate : 5 moves
Configured Sampling-Interval : 30 seconds
Number of entries in the notification table : 100
MAC-Address      from-Port     to-Port      Last Move-Time   Vlan-id
-----  -----  -----  -----  -----
0000.0000.0022    7/1        7/2          10               Apr
29 18:29:35
0000.0000.0021    7/1        7/2          10               Apr
29 18:29:35
0000.0000.0020    7/1        7/2          10               Apr
29 18:29:35
0000.0000.001f    7/1        7/2          10               Apr
29 18:29:35
0000.0000.0024    7/1        7/2          10               Apr
29 18:29:35
0000.0000.001e    7/1        7/2          10               Apr
29 18:29:35
0000.0000.0023    7/1        7/2          10               Apr
29 18:29:35
0000.0000.001d    7/1        7/2          10               Apr
29 18:29:35
0000.0000.001c    7/1        7/2          10               Apr
29 18:29:35
(output truncated)
```

Syntax: show notification mac-movement threshold-rate

The following table defines the fields in the output of the **show notification mac-movement threshold-rate** command.

TABLE 2 Field definitions for the show notification mac-movement threshold-rate command

Field	Description
Threshold-Rate Mac Movement Notification is	Specifies whether the MAC movement notification threshold rate is enabled.
Configured Threshold-Rate	The rate in MAC address moves per sampling interval after which a notification is issued. The range is from 1 through 50000.
Configured Sampling-Interval	The sampling interval in seconds over which the number of MAC address moves is measured. The range is from 1 through 86400, which is the number of seconds in a day.
Number of entries in the notification table	One entry for each time a MAC address notification threshold was reached.
MAC-Address	The MAC address that has moved to a different port.
from-Port	The port from which the MAC address moved.
to-Port	The port to which the MAC address moved.

TABLE 2 Field definitions for the show notification mac-movement threshold-rate command
(Continued)

Field	Description
Last Move-Time	The time of the last move occurred. It uses the system up time if there is no time server configured.
Vlan-id	The VLAN for the port where the MAC address movement was detected.

Configuring an interval for collecting MAC address move notifications

To configure an interval for collecting statistical data about MAC address moves, enter the **mac-movement notification interval-history** command at the privileged EXEC level. This command enables a corresponding SNMP trap. This history includes statistical information such as the number of MAC addresses that move over the specified period, the total number of MAC address moves, which MAC addresses have moved, and how many times a MAC address has moved.

The software places an upper limit on the number of MAC addresses for which MAC address-specific data is reported. This limit is necessary to do this because it is not possible to report on all MAC addresses when many move.

The following example configures a history interval of 10 seconds.

```
device(config)# mac-movement notification interval-history 10
```

To disable the feature and the corresponding SNMP trap, enter the **no** version of the command, as shown in the following example.

```
device(config)# no mac-movement notification interval-history 10
```

Syntax: [no] mac-movement notification interval-history *interval*

The *interval* variable represents the amount of time in seconds during which the MAC address movement notification data is collected. It has no default value.

Viewing MAC address movement statistics for the interval history

To display the collected history of MAC address movement notification, enter the **show notification mac-movement interval-history** command at the privileged EXEC level. This command displays how the history interval is configured in addition to the MAC address move data itself.

```
device# show notification mac-movement interval-history
Interval-History Mac Movement Notification is ENABLED
Configured Interval : 30 seconds
Number of macs that moved in the interval : 100
Total number of moves in the interval : 98654
MAC-Address      from-Port      to-Port      Interval Move-Count      Last Move-Time
Vlan-id
-----
0000.0000.0052      7/1          7/2
1000
15
01:13:20
10
0000.0000.0051      7/1          7/2
1002
15
01:13:20
10
```

```

0000.0000.0050    7/1      7/2
1012
15
01:13:20
0000.0000.004f    7/1      7/2      10
1018
15
01:13:20
0000.0000.004e    7/1      7/2      10
1012
15
01:13:20
10
(output truncated)

```

The following table defines the fields in the output of the **show notification mac-movement interval-history** command.

TABLE 3 Field definitions for the show notification mac-movement interval-history command

Field	Description
Interval-History Mac Movement Notification is	Specifies whether the interval-history data collection is enabled.
Configured Interval	The interval over which the MAC address movement statistics were collected.
Number of macs that moved in the interval	The number of MAC addresses that moved during the configured interval, regardless of how many times each address moved.
Total number of moves in the interval	The total number of MAC address moves over the configured interval.
MAC-Address	The MAC address that has moved to a different port.
from-Port	The port from which the MAC address moved.
to-Port	The port to which the MAC address moved.
Interval Move-Count	The number of times the MAC address has moved within the interval.
Last Move-Time	The time of the last MAC move occurred. It uses the system up time if there is no time server configured.
Vlan-id	The VLAN for the port where the MAC address movement was detected.

SNMP MAC-notification trap support

The SNMP MAC-notification trap functionality allows an SNMPv3 trap to be sent to the SNMP manager when MAC addresses are added or deleted in the device. The SNMP manager or management software can then use these traps to define a security policy based on the requirement of the enterprise where the device is installed. With this functionality, management software can easily monitor Brocade devices and build a security policy for enterprise networks.

Access ports can be manually configured to enable the MAC-notification feature. While enabling MAC-notification on a particular port, you can configure the interval at which the trap messages will be sent to management software, and the buffer size which maintains maximum trap events that can be maintained in the system. Ports enabled for MAC-notification will send SNMP traps to management software for various MAC address events such as addition, deletion, and MAC address movement.

The access devices in an enterprise network typically connect to the end host, and MAC-notification can be deployed on such devices on the access port only. An access port by definition is a port that connects to an end host and typically does not result in a network loop.

Requirements and limitations for MAC-notification trap support

The following requirements and limitations apply to MAC-notification trap support:

- MAC-notification is only supported on access ports.
- The network administrator must ensure that there are no loops in the ports enabled for MAC-notification, because high volume and frequent MAC address movement is not expected on the access port.
- The expected MAC scaling with the MAC-notification functionality is 800 MAC addresses per system, on the access ports where it is enabled. An extra buffer queue size is reserved to absorb any burst.
- The MAC-notification could be bursty in nature. This could be due to a set of hosts that could join at a specific time or a security policy change that could move a set of MAC addresses from one VLAN to another. Such bursty events need to be queued, resulting in delayed notifications to the management software.
- The number of events that can be queued is finite.
- All queued events are notified during the notification interval. The notification interval should be tuned based on the requirements of the enterprise. However, a very aggressive timer coupled with bursty traffic could load the system and result in a loss of MAC-notification events.
- Static and control MAC events are not considered for MAC-notification event generation.
- MAC-notification is supported at an interface level on a device. When enabled, each MAC address addition or deletion is logged as an event in a buffer-queue.
- MAC-notification is currently not supported on MCT (Multi Chassis Trunking).

Configuring SNMP traps for MAC-notification

1. Use the **mac-notification interval** command with the specified interval value to enable MAC-notification.
2. Use the **interface ethernet** command with the specified Ethernet interface to enable MAC-notification on the individual interface.
3. Use the **snmp-server enable traps mac-notification** command to enable MAC-notification on the specified interface.
4. Use the **system-max mac-notification-buffer** command to change the value of the MAC-notification buffer size.

The following example shows enabling SNMP traps for MAC-notification on Ethernet interface 1/1/5:

```
device(config)# mac-notification interval 30
device(config)# interface ethernet 1/1/5
device(config-if-e1000-1/1/5)# snmp-server enable traps mac-notification
device(config-if-e1000-1/1/5)# exit
device (config)# system-max mac-notification-buffer 4000
```

Use the **show interfaces ethernet** command to check whether a MAC-notification SNMP trap is enabled or disabled on an interface. You can also use the **show mac-notification** command to view other statistics such as the configured interval, the number of traps sent, and the number of events sent.

MAC-notification events

NOTE

MAC-notifications for LAG should be enabled on the primary port.

When enabled, each MAC address addition or deletion is logged as an event in a buffer-queue. Each event is 11 bytes long and contains information about the following:

Value	Description
MAC address	The MAC address added or deleted on the device.
VLAN	The VLAN to which the MAC address is associated. The valid range is 1 to 4094.
Interface	The interface on which the MAC address is added or deleted.
Action	The event that occurred.

The following table lists the various events that can occur, along with the VLAN interface values and their interpretation for each event:

TABLE 4 MAC address notification events and values

Event	Action Value	Description	Expected action by management software	VLAN and port values
ADD-MAC	1	This event is generated when a new MAC address is learnt.	Management software should add the MAC address to its forwarding table.	(VLAN, Port)
REMOVE-MAC	2	This event is generated when the MAC address ages out.	Management software should delete the MAC address from its forwarding table.	(VLAN, Port)
REMOVE-ALL-MAC-ON-SYSTEM	3	This event is generated when all the MAC addresses on the system are flushed, for example, by using the clear mac command.	Management software should clear all the MAC addresses from its forwarding table.	(0, 0)

TABLE 4 MAC address notification events and values (Continued)

Event	Action Value	Description	Expected action by management software	VLAN and port values
REMOVE-ALL-MAC-ON-PORT	4	This event is generated when all the MAC addresses on a particular port are flushed, for example, when the link goes down.	Management software should clear all the MAC addresses learnt on this particular port from its forwarding table.	(0, Port)
REMOVE-ALL-MAC-ON-VLAN	5	This event is generated when the MAC addresses learnt on all ports, in a particular VLAN are flushed, for example, by using the no vlan command.	Management software should clear all the MAC addresses learnt on this particular VLAN from its forwarding table.	(VLAN, 0)
REMOVE-ALL-MAC-ON-VLAN-PORT	6	This event is generated when the MAC addresses, are flushed for a particular port in a particular VLAN, for example by a protocol flush event.	Management software should clear all the MAC addresses learnt on this particular VLAN and port from its forwarding table.	(VLAN, Port)
MAC-MOVE	7	This event is generated when the MAC address moves from an old port to a new port in the same VLAN.	Management software should move the MAC address from the old port to the specified new port learnt in its forwarding table.	(VLAN, new port)

Working with MAC-notification events

- Each event stored in the buffer queue is in the order in which the event occurred in the system.
- The number of events that can be stored in the buffer queue is by default 4000. This value is configurable up to 16000 through the command line interface.
- An out-of-band buffer full event trap is sent to the management software in the event of a buffer full. The system then flushes the existing buffer queue.
- You can configure a periodic interval at which point a MAC-notification trap should be sent to the management software. The interval can range from 1 to 3600 seconds. The default is 3 seconds.
- Each trap message sent on the notification interval can have one or more MAC-notification events taken from the buffer queue in the first-in first-out order.
- One or more SNMP trap messages can be sent on the expiry of a MAC-notification interval. However, the maximum number of trap messages that can be sent is limited to 5.

Displaying and modifying system parameter default settings

Brocade devices have default table sizes for the system parameters shown in the following display outputs. The table sizes determine the maximum number of entries the tables can hold. You can adjust individual table sizes to accommodate your configuration needs.

The tables you can configure, as well as the default values and valid ranges for each table, differ depending on the Brocade device you are configuring. To display the adjustable tables on your Brocade device, use the **show default values** command. The following shows example outputs.

System default settings configuration considerations

- Changing the table size for a parameter reconfigures the device memory. Whenever you reconfigure the memory on a Brocade device, you must save the change to the startup-config file, then reload the software to place the change into effect.
- Configurable tables and their defaults and maximum values differ on Brocade IPv4 devices versus IPv6-capable devices.
- For more information about Layer 3 system parameter limits, refer to "Displaying Layer 3 system parameter limits" section in the *FastIron Ethernet Switch Layer 3 Routing Configuration Guide*.

Displaying system parameter default values

To display the configurable tables and their defaults and maximum values, enter the **show default values** command at any level of the CLI.

The following shows an example output of the **show default values** command on a FastIron Layer 2 device.

```
device#show default values
sys log buffers:50          mac age time:300 sec      telnet sessions:5
System Parameters   Default    Maximum   Current   Configured
igmp-max-group-addr 4096      8192      1024
ip-filter-sys        2048      4096      4096
13-vlan              32         1024      1024
mac                  32768     32768     32768
vlan                64         4095      4095
spanning-tree        32         255       255
mac-filter-port      32         256       256
mac-filter-sys       64         512       512
view                10         65535     65535
rmon-entries         1024      32768     32768
mld-max-group-addr  8192      32768     32768
igmp-snoop-mcache   512       8192      8192
mld-snoop-mcache    512       8192      8192
```

The following shows an example output of the **show default values** command on a FastIron Layer 2 ICX 6650 device.

```
device#show default values
sys log buffers:50          mac age time:300 sec      telnet sessions:5
System Parameters   Default    Maximum   Current
igmp-max-group-addr 4096      8192      4096
ip-filter-port       2045      2045      2045
ip-filter-sys        2048      8192      2048
13-vlan              32         1024      32
mac                  65536     65536     65536
vlan                64         4095      64
spanning-tree        32         254       32
mac-filter-port      32         256       32
mac-filter-sys       64         512       64
view                10         65535     10
rmon-entries         1024      32768     1024
mld-max-group-addr  8192      32768     8192
igmp-snoop-mcache   512       8192      512
mld-snoop-mcache    512       8192      512
```

The following shows an example output on a FastIron IPV4 device running Layer 3 software.

```
device#show default values
sys log buffers:50          mac age time:300 sec      telnet sessions:5
ip arp age:10 min           bootp relay max hops:4      ip ttl:64 hops
```

Basic Layer 2 Features

```

ip addr per intf:24
when multicast enabled :
igmp group memb.:260 sec    igmp query:125 sec      hardware drop: enabled
when ospf enabled :
ospf dead:40 sec            ospf hello:10 sec       ospf retrans:5 sec
ospf transit delay:1 sec
when bgp enabled :
bgp local pref.:100          bgp keep alive:60 sec   bgp hold:180 sec
bgp metric:10                bgp local as:1           bgp cluster id:0
bgp ext. distance:20          bgp int. distance:200  bgp local distance:200
System Parameters  Default  Maximum  Current
ip-arp            6000     64000    6000
ip-static-arp     512       6000     512
multicast-route   64        8192     64
dvmrp-route      2048     32000    2048
dvmrp-mcache     512       4096     512
pim-mcache       1024     4096     1024
igmp-max-group-addr 4096     8192     4096
ip-cache          10000    32768    10000
ip-filter-port    1015     1015     1015
ip-filter-sys     2048     8192     2048
13-vlan           32        1024     32
ip-qos-session    1024     16000    1024
mac               16384    32768    16384
ip-route          80000    262144   80000
ip-static-route   64        2048     64
vlan              64        4095     64
spanning-tree     32        255      32
mac-filter-port   16        256      16
mac-filter-sys    32        512      32
ip-subnet-port   24        128      24
session-limit     65536    160000   65536
view              10        65535    10
virtual-interface 255       512      255
hw-ip-next-hop   2048     6144     2048
hw-logical-interface 4096    4096     4096
hw-ip-mcast-mll  1024     4096     1024
hw-traffic-condition 50      1024     50
rmon-entries      2048     32768    2048
mld-max-group-addr 8192    32768    8192
igmp-snoop-mcache 512      8192     512
mld-snoop-mcache  512      8192     512
msdp-sa-cache     4096    8192     4096

```

The following shows an example output on a FastIron IPV4 ICX 6650 device running Layer 3 software.

```

device#show default values
sys log buffers:50          mac age time:300 sec      telnet sessions:5
                           bootp relay max hops:4    ip ttl:64 hops
ip arp age:10 min          ip addr per intf:24
                           when multicast enabled :
                           igmp group memb.:260 sec    igmp query:125 sec      hardware drop: enabled
                           when ospf enabled :
                           ospf dead:40 sec            ospf hello:10 sec       ospf retrans:5 sec
                           ospf transit delay:1 sec
                           when bgp enabled :
                           bgp local pref.:100          bgp keep alive:60 sec   bgp hold:180 sec
                           bgp metric:10                bgp local as:1           bgp cluster id:0
                           bgp ext. distance:20          bgp int. distance:200  bgp local distance:200
System Parameters  Default  Maximum  Current
ip-arp            4000     64000    64000
ip-static-arp     512       6000     6000
multicast-route   64        8192     8192
pim-mcache       1024     4096     4096
igmp-max-group-addr 4096     8192     8192
ip-cache          10000    32768    32768
ip-filter-port    2045     2045     2045
ip-filter-sys     2048     8192     8192
13-vlan           32        1024     1024
ip-qos-session    1024     16000    16000
mac               65536    65536    65536
ip-route          5120     7168     6500

```

ip-static-route	64	2048	2048
vlan	64	4095	4095
spanning-tree	32	254	254
mac-filter-port	16	256	256
mac-filter-sys	32	512	512
ip-subnet-port	24	128	128
session-limit	8192	16384	16384
view	10	65535	65535
virtual-interface	255	512	512
hw-traffic-condition	896	896	896
rmon-entries	1024	32768	32768
mld-max-group-addr	8192	32768	32768
igmp-snoop-mcache	512	8192	8192
mld-snoop-mcache	512	8192	8192
ip6-route	580	1348	187
ip6-static-route	37	269	37
ip6-cache	93	674	93
gre-tunnels	16	64	64
hw-ip-route-tcam	8192	8192	8192

The following shows an example output on a FCX serving as a management host in an IPv6 network and running the Layer 3 software image.

```
device#show default values
sys log buffers:50          mac age time:300 sec      telnet sessions:5
ip arp age:10 min           bootp relay max hops:4   ip ttl:64 hops
ip addr per intf:24
when multicast enabled :
igmp group memb.:260 sec    igmp query:125 sec
when ospf enabled :
ospf dead:40 sec            ospf hello:10 sec
ospf transit delay:1 sec
when bgp enabled :
bgp local pref.:100          bgp keep alive:60 sec
bgp metric:10                bgp local as:1
bgp ext. distance:20         bgp int. distance:200
System Parameters      Default      Maximum      Current
ip-arp                  4000        64000       64000
ip-static-arp            512         6000        6000
multicast-route          64          8192        8192
pim-mcache               1024        4096        4096
igmp-max-group-addr     4096        8192        8192
ip-cache                 10000       32768       32768
ip-filter-port            4093        4093        4093
ip-filter-sys             2048        4096        4096
l3-vlan                  32          1024        1024
ip-qos-session            1024        16000       16000
mac                      32768       32768       32768
ip-route                 12000       16100       16100
ip-static-route           64          2048        2048
vlan                     64          4095        4095
spanning-tree             32          255         255
mac-filter-port            16          256         256
mac-filter-sys             32          512         512
ip-subnet-port             24          128         128
session-limit              8192        16384       16384
view                      10          65535       65535
virtual-interface          255         512         512
rmon-entries               1024        32768       32768
mld-max-group-addr        8192        32768       32768
igmp-snoop-mcache          512         8192        8192
mld-snoop-mcache           512         8192        8192
hw-ip-route-tcam           16384       16384       16384
```

The following shows an example output on a FastIron X Series IPv6 device running the Layer 3 software image.

```
device#show default values
sys log buffers:50          mac age time:300 sec      telnet sessions:5
ip arp age:10 min           bootp relay max hops:4   ip ttl:64 hops
ip addr per intf:24
when multicast enabled :
igmp group memb.:260 sec    igmp query:125 sec
when ospf enabled :
ospf dead:40 sec            ospf hello:10 sec
ospf transit delay:1 sec
```

Basic Layer 2 Features

```

when bgp enabled :
bgp local pref.:100      bgp keep alive:60 sec      bgp hold:180 sec
bgp metric:10            bgp local as:1           bgp cluster id:0
bgp ext. distance:20    bgp int. distance:200     bgp local distance:200
System Parameters Default Maximum Current
ip-arp      6000   64000  64000
ip-static-arp 512    6000   6000
multicast-route 64    8192   8192
dvmrp-route  2048   32000  32000
dvmrp-mcache 512    4096   4096
pim-mcache  1024   4096   4096
igmp-max-group-addr 4096  8192   8192
ip-cache    10000  32768  32768
ip-filter-port 1015  1015   1015
ip-filter-sys 2048   8192   4096
13-vlan     32     1024   1024
ip-qos-session 1024  16000  16000
mac         16384  32768  32768
ip-route    262144 524288 524288
ip-static-route 64    2048   2048
vlan        64     4095   4095
spanning-tree 32     255    255
mac-filter-port 16     256    256
mac-filter-sys 32     512    512
ip-subnet-port 24     128    128
session-limit 65536  160000 65537
view        10     65535  65535
virtual-interface 255   512    512
hw-ip-next-hop 2048   6144   2481
hw-traffic-condition 50    1024   52
rmon-entries 2048   32768  32768
mld-max-group-addr 8192  32768  32768
igmp-snoop-mcache 512    8192   8192
mld-snoop-mcache 512    8192   8192
ip6-route   32768  65536  65536
ip6-static-route 64    512    512
ip6-cache   65536  131072 131072
msdp-sa-cache 4096  8192   8192
gre-tunnels 16     64     64

```

The following shows an example output on a ICX 7750 device.

```

device# show default values
sys log buffers:50          mac age time:300 sec      telnet sessions:5
                           bootp relay max hops:4      ip ttl:64 hops
ip arp age:10 min          ip addr per intf:24
                           igmp group memb.:260 sec      igmp query:125 sec      hardware drop: enabled
                           ospf dead:40 sec          ospf hello:10 sec      ospf retrans:5 sec
                           ospf transit delay:1 sec
                           when bgp enabled :
                           bgp local pref.:100      bgp keep alive:60 sec      bgp hold:180 sec
                           bgp metric:10            bgp local as:1           bgp cluster id:0
                           bgp ext. distance:20    bgp int. distance:200     bgp local distance:200
System Parameters Default Maximum Current Configured
ip-arp      8192   64000  64000  64000
ip-static-arp 512    1024   512    512
ip-cache    8192   32768  32768  32768
ip-filter-port 2047  2047   2047  2047
ip-filter-sys 3072  8192   3072  3072
13-vlan     32     1024   32     32
ip-qos-session 1024  16000  1024  1024
mac         32768  32768  32768  32768
ip-route    98304  131072 98304  98304
ip-static-route 64    2048   64     64
vlan        64     4095   4095  4095
spanning-tree 128   254    254   254
mac-filter-port 32     256    32     32
mac-filter-sys 64     512    64     64
ip-subnet-port 24     128    24     24
session-limit 65536  160000 65536  65536
view        10     65535  10     10

```

virtual-interface	255	512	255	255
hw-ip-next-hop	17408	17408	17408	17408
hw-traffic-condition	50	1024	50	50
rmon-entries	2048	32768	2048	2048
igmp-snoop-mcache	512	6144	6144	6144
mld-snoop-mcache	512	6144	6144	6144
ip6-route	5120	7168	5120	5120
ip6-static-route	64	1024	64	64
ip6-cache	1024	2048	1024	1024
msdp-sa-cache	1024	4096	1024	1024
gre-tunnels	16	64	16	16
ip-vrf	128	128	128	128
ip-route-default-vrf	65536	131072	10000	10000
ip6-route-default-vr	2048	7168	310	310
ip-route-vrf	4096	131072	1500	1500
ip6-route-vrf	1024	7168	800	800
pim-hw-mcache	1024	6144	6144	6144
pim6-hw-mcache	512	2048	1024	1024
igmp-snoop-group-add	4096	8192	8192	8192
mld-snoop-group-addr	4096	8192	8192	8192
mac-notification-buf	4000	16000	4000	4000

The following table defines the system parameters in the **show default values** command output.

TABLE 5 System parameters in show default values command

Parameter	Definition
dvmrp-mcache	PIM and DVMRP multicast cache flows stored in CAM
dvmrp-route	DVMRP routes
hw-ip-mcast-mll	Multicast output interfaces (clients)
hw-ip-next-hop	IP next hops and routes, including unicast next hops and multicast route entries
hw-logical-interface	Hardware logical interface pairs (physical port and VLAN pairs)
hw-traffic-conditioner	Traffic policies
ip-arp	ARP entries
ip-cache	IP forwarding cache entries
ip-filter-port	IP ACL entries per port
ip-filter-sys	IP ACL entries per system
ip-qos-session	Layer 4 session table entries
ip-route	Learned IP routes
ip-static-arp	Static IP ARP entries
ip-static-route	Static IP routes
ip-subnet-port	IP subnets per port
l3-vlan	Layer 3 VLANs
mac	MAC entries

TABLE 5 System parameters in show default values command (Continued)

Parameter	Definition
mac-filter-port	MAC address filter entries per port
mac-filter-sys	MAC address filter entries per system
multicast-route	Multicast routes
pim-mcache	PIM multicast cache entries
rmon-entries	RMON control table entries
session-limit	Session entries
spanning-tree	Spanning tree instances
view	SNMP views
virtual-interface	Virtual routing interfaces
vlan	VLANs
mld-max-group-addr	MLD group limit
igmp-snoop-mcache	IGMP snooping cache entries
mld-snoop-mcache	MLD snooping cache entries

Modifying system parameter default values

Information for the configurable tables appears under the columns that are shown in bold type in the above examples. To simplify configuration, the command parameter you enter to configure the table is used for the table name. For example, to increase the capacity of the IP route table, enter the following commands.

```
device(config)#system-max ip-route 120000
device(config)#write memory
device(config)#exit
device#reload
```

Syntax: **system-max ip-route num**

The *num* parameter specifies the maximum number of routes in the IP route table. The minimum value is 4096. The maximum value is 524288 (subject to route patterns for FSX). The default is 80000 IP routes. For ICX 6650, the minimum value is 2048. The maximum value is 7168. The default is 5120 IP routes.

NOTE

If you accidentally enter a value that is not within the valid range of values, the CLI will display the valid range for you.

To increase the number of IP subnet interfaces you can configure on each port on a device running Layer 3 code from 24 to 64, enter the following commands.

```
device(config) #system-max ip-subnet-port 64
device(config) #write memory
device(config) #exit
device#reload
```

Syntax: system-max ip-subnet-port num

The *num* parameter specifies the maximum number of subnet addresses per port and can be from 24 - 128. The default is 24.

Cut-through switching

Brocade devices operate in cut-through switching mode, meaning it starts forwarding a frame even before the whole frame has been received. The amount of time the device takes to start forwarding the packet (referred to as the switch's latency) is on the order of a few microseconds only, regardless of the packet size. The Table provides the latency details.

TABLE 6 Cut-through latency

Packet size in bytes	10G latency in microseconds (10G to 10G)	40G latency in microseconds(40G to 40G)
64	1.41	1.26
128	1.47	1.27
256	1.55	1.31
512	1.75	1.36
1024	1.73	1.46
1516	1.73	1.55
5000	1.73	1.66
9212	1.73	1.66

- If there is any over-subscription on the egress port, either due to speed mismatch or network topology, the device will buffer the packets and the forwarding behavior will be similar to store-and-forward mode.
- If an FCS error is determined when the packet is processed by the ingress pipe, it is dropped at the end of the ingress pipe. When an FCS error is determined after the packet transmission to the egress port has begun, it is transmitted with a faulty CRC. When an FCS error is determined during a packet transmission the packet truncated.
- Forwarding from fast speed ports to slower ports is equivalent to store-and-forward (has to be stored first). Forwarding from slower speed ports to faster ports is also equivalent to store-and-forward (to avoid underrun).
- Cut-through switching is not enabled on 1G ports.
- Cut-through minimum packet size is 128 bytes.
- Features that are based on the packet length are not supported since the packet is transmitted before being fully received.

Dynamic buffer allocation for QoS priorities for FastIron X Series devices

NOTE

This section applies to FastIron X Series devices only. To configure dynamic buffer limits Brocade FCX Series, and ICX devices, refer to [Dynamic buffer allocation for FCX and ICX devices](#) on page 48.

By default, Brocade IronWare software allocates a certain number of buffers to the outbound transmit queue for each port, based on QoS priority (traffic class). The buffers control the total number of packets permitted in the outbound transmit queue for the port. For each port, the Brocade device defines the maximum outbound transmit buffers, also called *queue depth limits*, as follows:

- Total Transmit Queue Depth Limit - The total maximum number of transmit buffers allocated for all outbound packets on a port. Packets are added to the port's outbound queue as long as the number of buffers currently in use is less than the total transmit queue depth limit. When this limit is reached, any new packets attempting to enter the port's transmit queue will be dropped until at least one buffer is freed.
- Transmit Queue Depth Limit for a Given Traffic Class - The maximum number of transmit buffers allocated for packets with a given traffic class (0 through 7) on a port. Packets with the specified traffic class are added to the port's outbound queue as long as the number of buffers currently in use for that traffic class is less than the transmit queue depth limit for the traffic class. When this limit is reached, any new packets with the specified traffic class attempting to enter the port's transmit queue will be dropped.

Except for the third-generation Interface modules, you can increase or decrease both of these queue depth limits per port. On the SX-FI48GPP, SX-FI-24GPP, SX-FI-24HF, SX-FI-2XG or SX-FI-8XG modules, to increase or decrease the queue depth limits for a port, you configure a buffer profile and then apply it to the port. A buffer profile can be tied to one or more ports. Therefore, if you change the configuration on one port, it will change the configuration on all ports associated with the same buffer profile.

Dynamic buffer allocation for QoS is useful in situations where applications have intermittent bursts of oversubscription. For example, by increasing the buffers on the egress port, the Brocade device will be able to forward oversubscribed packets instead of dropping them.

Default queue depth limits for FastIron X Series devices

The following table defines the default maximum queue depth values per port, per traffic class. The Brocade device drops the packets that cause the port to exceed these limits.

NOTE

The SX-FI48GPP Interface module supports 48 tri-speed (10/100/1000) ports. When the ports are configured at lower speeds, for example, 100 Mbps or 10 Mbps, the maximum queue depth is less than 256 when egress congestion occurs at the front-end Network Processor (NP) of the SX-FI48GPP module. If egress congestion occurs at the back-end NP of the SX-FI48GPP module, the maximum queue depth is 4095. The limit for buffer sharing is 4088.

TABLE 7 Default maximum queue depth

Port type	Maximum queue depth per port, per priority 0 - 6	Maximum queue depth per port, per priority 7	Total maximum queue depth
1 Gbps port	96	224	896
10 Gbps port without jumbo enabled	400	704	3504
10 Gbps port with jumbo enabled	352	640	3104

Configuring the total transmit queue depth limit for FastIron X Series devices

NOTE

To configure the total transmit queue depth limit on a third-generation module, refer to [Buffer profile configuration](#) on page 46.

To set the total transmit queue depth limit on a port, enter a command such as the following.

```
device(config)#qd 2 2049
```

This command sets the queue depth limit on port 2 to 2049. Packets are added to the port's outbound queue as long as the packets do not cause the port to exceed 2048 buffers. If the port reaches its queue depth limit of 2049, any new packets attempting to enter the port transmit queue will be dropped until at least one buffer is freed.

Syntax: `qd port limit`

The *limit* variable can be a value from 0 through 4095. [Table 7](#) on page 45 lists the default values.

Configuring the transmit queue depth limit for a given traffic class on FastIron X Series devices

NOTE

To configure transmit queue depth limits for an SX-FI48GPP module, refer to [Buffer profile configuration](#) on page 46.

To set the transmit queue depth limit on a port for a given traffic class, first enter the transmit queue depth limit for the traffic class, and then specify the traffic class.

```
device(config)#qd 2 200 7
```

This command sets the queue depth limit on port 2 to 200 for packets with a traffic class of 7. Packets with priority 7 are added to the outbound queue on port 2 as long as the packets do not exceed 199 buffers. When the port reaches its queue depth limit of 200, packets with the given traffic class will be dropped.

Syntax: `qd port limit traffic-class`

The limit variable can be a value from 0 through 4095 and cannot exceed the total transmit queue depth limit configured for the port. [Table 7](#) on page 45 lists the default values.

The sum of the queue depth limits for individual traffic classes on a port does not need to equal the total queue depth limit for the port:

- If the sum of the individual traffic class queue depth limits exceeds the total port limit and the total port limit is reached, any buffer that gets released can be used by any traffic class queue that has not reached its individual limit.
- If the sum of the individual traffic class queue depth limits is less than the total port limit, the remaining buffers can be used only by packets with a priority of 7.

The traffic-class variable can be a value from 0 through 7, where 7 is the highest priority queue.

Removing buffer allocation limits on FastIron X Series devices

You can remove buffer allocation limits on all ports and all Traffic Classes globally. This permits all available buffers in a port region to be used in a first-come-first-serve basis by any of its ports, regardless of priority. This can be done using the following command.

```
device (config) #buffer-sharing-full
```

Syntax: [no] buffer-sharing-full

The **buffer-sharing-full** command sets the total transmit queue depth limit and the transmit queue depth limits for each Traffic Class to 4095 for all ports of the device. The command overrides any existing individually configured queue depth limits.

ATTENTION

The **buffer-sharing-full** command should be used carefully. By entering this command, there is no limit on the number of buffers a port or a specific priority on a port can use. One port could potentially use up all the available buffers of its port region and cause starvation on other ports of the port region.

Buffer profile configuration

The following Interface modules support up to eight buffer profiles:

- SX-FI48GPP
- SX-FI-24GPP
- SX-FI-24HF
- SX-FI-2XG
- SX-FI-8XG

A buffer profile defines the *total transmit queue depth limit for a port* and the *transmit queue depth limit for a given traffic class*. On the listed supported Interface modules, each port is associated with a buffer profile. In contrast, Interface modules other than the SX-FI-48GPP, SX-FI-24GPP, SX-FI-24-HF, SX-FI-2XG, and the SX-FI-8XG support the configuration of transmit queue depth limits per port.

By default, each port on an SX-FI48GPP, SX-FI-24GPP, SX-FI-24HF, SX-FI-2XG or SX-FI-8XG Interface module is associated with buffer profile ID 1. Profile 0 is reserved for buffer sharing. Default queue depth limits are provided in [Default queue depth limits for FastIron X Series devices](#) on page 44. To change the queue depth limit for a port on the supported Interface modules, configure a buffer profile ID that is different from the default or configured value, and apply the buffer profile to the port. When you change the queue depth limit on a port, the configuration will apply to all ports associated with the same buffer profile ID. For example, if ports 1/1 through 1/5 are associated with buffer profile

ID 3 and you enter the `qd 1/1 1000` command, the queue depth limit of 1000 will apply to ports 1/1 through 1/5.

NOTE

In the case of SX-FI48GPP module, the queue depth configured using the `qd` command is not applicable when the port is set to speeds lower than 1Gbps. The `qd` command is applicable only when the port is set to speeds greater than 1Gbps.

Configuring a buffer profile and defining the queue depth limits

1. Create a buffer profile and assign it to a port. For example, to create buffer profile 2 and assign it to port 1/1, enter the following command.

```
device(config)#qd 1/1 profile-id 2
```

Syntax: [no] `qd slotnum/portnum profile-id number`

The number variable specifies the buffer profile ID associated with the slotnum / portnum . The buffer profile ID can be a number from 1 through 7. Profile 0 is reserved for buffer sharing.

2. Configure the total transmit queue depth limit associated with the buffer profile. For example, to change the total transmit queue depth limit for buffer profile 2 to 1000, enter the following command.

```
device(config)#qd 1/1 1000
```

Because port 1/1 is associated with buffer profile 2, this command sets the queue depth limit for buffer profile 2 to 1000. Packets are added to the outbound queue of a port as long as the packets do not cause the port to exceed 1000 buffers. If the port reaches its queue depth limit of 1000, any new packets attempting to enter the transmit queue will be dropped until at least one buffer is freed.

Syntax: [no] `qd slotnum/portnum limit`

The limit variable can be a value from 0 through 4095. The default is 4095.

3. If desired, configure the queue depth limit for a given traffic class. For example, to change the queue depth limit for buffer profile 2 to 300 for packets with a traffic class of 1, enter the following command.

```
device(config)#qd 1/1 300 1
```

For ports that use buffer profile 2, packets with priority 1 are added to the outbound queue as long as the packets do not exceed 299 buffers. When the port reaches its queue depth limit of 300, packets with the given traffic class will be dropped.

Syntax: [no] `qd slotnum/portnum limit traffic-class`

The limit variable can be a value from 0 through 4095 and cannot exceed the total transmit queue depth limit configured in the previous step.

The sum of the queue depth limits for individual traffic classes on a port does not need to equal the total queue depth limit for the port:

- If the sum of the individual traffic class queue depth limits exceeds the total port limit and the total port limit is reached, any buffer that gets released can be used by any traffic class queue that has not reached its individual limit.
- If the sum of the individual traffic class queue depth limits is less than the total port limit, the remaining buffers can be used only by packets with a priority of 7.

The traffic-class variable can be a value from 0 through 7, where 7 is the highest priority queue.

4. Enter the `write memory` command to save the configuration.

Displaying the buffer profile configuration

To display the buffer profile configuration for an SX-FI48GPP Interface module, use the show configuration command.

The following example shows that buffer profile 2 and its configured queue depth values apply to ports 1/1 and 1/2. Although the profile configuration was changed for port 1/1 only, port 1/2 has also changed to match the configuration.

```
device(config)#show configuration
qd 1/1 profile-id 2
qd 1/2 profile-id 2
!
qd 1/1 1000
qd 1/2 1000
qd 1/1 300 1
qd 1/2 300 1
```

Syntax: show configuration

Default queue depth limits for FastIron X Series devices

[Table 7](#) on page 45 defines the default maximum queue depth values per port, per traffic class. The Brocade device drops the packets that cause the port to exceed these limits.

NOTE

The SX-FI48GPP Interface module supports 48 tri-speed (10/100/1000) ports. When the ports are configured at lower speeds, for example 100 Mbps or 10 Mbps, the maximum queue depth is less than 256 when egress congestion occurs at the front-end Network Processor (NP) of the SX-FI48GPP module. If egress congestion occurs at the back-end NP of the SX-FI48GPP module, the maximum queue depth is 4095. The limit for buffer sharing is 4088.

Dynamic buffer allocation for FCX and ICX devices

By default, the traditional stack architecture allocates fixed buffers on a per-priority queue, per-packet processor basis. The buffers control the total number of packets that can be queued in the outbound transmit for the port. In instances of heavy traffic bursts to aggregation links, such as in stacking configurations or mixed-speed environments, momentary oversubscription of the buffers and descriptors may occur. A descriptor points to one or more packet buffers.

Brocade FastIron stackable devices provide the capability to allocate additional egress buffering and descriptors to handle momentary bursty traffic periods, especially when other priority queues may not be in use, or may not be experiencing heavy traffic. This allows users to allocate and fine-tune the depth of priority buffer queues for each packet processor.

Configuring buffer profiles

There are two different methods of allocating buffers and descriptors to the ports and its queues.

One method uses the qd-descriptor and qd-buffer CLI commands to allocate descriptors and buffers, respectively, to the port and its queues. This method is available on FCX and ICX devices.

The other method uses user-configurable buffer profiles. This method allows you to define a template of buffer allocations to be used on a per-port per-queue basis on the devices. When applied, this

buffer profile acts as if you created a series of the **qd** commands. This buffer profile is a simpler form of allocating descriptors and buffers to the port and its queues. This method is available on FCX and ICX devices.

Configuring buffer profiles with qd-descriptor and qd-buffer commands on FCX and ICX devices

The 48-port Brocade stackable switch has two packet processors. The 24-port Brocade stackable switch has a single packet processor. For devices in a traditional stack, each stack unit has the possibility of two packet processors, but the second processor for a 24-port stack unit cannot be configured. The number of actual available packet processors depends on the type and number of switches in the stack.

For example, for an 8-unit stack of 48 ports, the packet processor numbering scheme is as follows:

- Stack unit 1 - Packet processors 0 and 1
- Stack unit 2 - Packet processors 2 and 3
- Stack unit 3 - Packet processors 4 and 5
- Stack unit 4 - Packet processors 6 and 7
- Stack unit 5 - Packet processors 8 and 9
- Stack unit 6 - Packet processors 10 and 11
- Stack unit 7 - Packet processors 12 and 13
- Stack unit 8 - Packet processors 14 and 15

In this configuration, if stack unit 3 and stack unit 7 are 24-port devices, the odd-numbered packet processors 5 and 13 cannot be configured, and do not exist, although they are reserved.

Configuration steps for buffer profile with qd-descriptor and qd-buffer commands on FCX and ICX

The descriptor and buffer allocation process occurs in four sequential steps using the **qd-buffer** and **qd-descriptor** commands.

NOTE

For FCX devices, when you reset buffer values for the 10 Gbps ports, the buffer values for the rear-panel 10 Gbps and 16 Gbps ports are also reset.

1. Configure the allowable port descriptors. Port descriptors set the limit for the ports. The minimum limit for the port descriptors is 16. The maximum limit of the port descriptors depends on the hardware device. Port descriptor limits of different platforms are listed in the section [Buffer and descriptor maximum and default allocation values](#) on page 54. Configure the allowable port descriptors by entering a command similar to the following.

```
device# qd-descriptor 1 2 76
```

Syntax: [no] qd-descriptor devicenum porttypeval numdescriptors

The devicenum variable refers to the device in the stacking unit. The device number starts from 0.

The porttypeval variable is 1 for 1 Gbps ports and 2 for 10 Gbps ports.

The numdescriptors variable refers to the number of descriptors to allocate.

2. Configure the queue descriptors. The minimum limit for queue descriptors is 16. The system default queue descriptors for different platforms are listed in [Buffer and descriptor maximum and default allocation values](#) on page 54. Configure the queue descriptors for the queue by entering a command similar to the following.

```
device# qd-descriptor 1 2 76 2
```

Syntax: [no] **qd-descriptor** *devicenum porttypeval numdescriptors priorityqueue*

The devicenum variable refers to the device in the stacking unit. The device number starts from 0.

The porttypeval variable is 1 for 1 Gbps ports and 2 for 10 Gbps ports.

The numdescriptors variable refers to the number of descriptors to allocate.

The priorityqueue variable refers to the specific queue of the port from 0 through 7.

3. Configure the port buffers. The minimum limit for port buffers is 16. The maximum limit for the port buffer depends on the hardware device. Port buffer limits of different platforms are listed in [Buffer and descriptor maximum and default allocation values](#) on page 54. Configure the allowable packet buffers by entering a command similar to the following.

```
device# qd-buffer 1 2 76
```

Syntax: [no] **qd-buffer** *devicenum porttypeval numbuffers*

The devicenum variable refers to the device in the stacking unit. The device number starts from 0.

The porttypeval variable is 1 for 1 Gbps ports and 2 for 10 Gbps ports.

The numbuffers variable refers to the number of buffers to allocate.

4. Configure the queue buffers. The maximum limit of queue buffers depends on the hardware device. Queue buffer limits of different platforms are listed in [Buffer and descriptor maximum and default allocation values](#) on page 54. Configure the queue buffers by entering a command similar to the following.

```
device# qd-buffer 1 2 76 2
```

Syntax: [no] **qd-buffer** *devicenum porttypeval numbuffers priorityqueue*

The devicenum variable refers to the device in the stacking unit. The device number starts from 0.

The porttypeval variable is 1 for 1 Gbps ports and 2 for 10 Gbps ports.

The numbuffers variable refers to the number of buffers to allocate.

The priorityqueue variable refers to the specific queue of the port from 0 through 7.

Sample configuration for buffer profile with qd-descriptor and qd-buffer commands on FCX and ICX

This sample configuration assumes a four-unit stack with the following topology. Note that there is no packet processor 3 or 7, because stack units 2 and 4 are 24-port devices.

- Stack unit 1, 48 ports - Packet processor numbers 0 and 1
- Stack unit 2, 24 ports - Packet processor number 2
- Stack unit 3, 48 ports - Packet processors 4 and 5
- Stack unit 4, 24 ports - Packet processor number 6

The following commands allocate available buffers to be used by priority 0 queues in the four-unit stack.

```
qd-descriptor 0 1 4095
qd-descriptor 1 1 4095
qd-descriptor 2 1 4095
qd-descriptor 4 1 4095
qd-descriptor 5 1 4095
qd-descriptor 6 1 4095
qd-descriptor 0 2 4095
qd-descriptor 1 2 4095
qd-descriptor 2 2 4095
qd-descriptor 4 2 4095
qd-descriptor 5 2 4095
qd-descriptor 6 2 4095
qd-descriptor 0 1 4095 0
qd-descriptor 1 1 4095 0
qd-descriptor 2 1 4095 0
qd-descriptor 4 1 4095 0
qd-descriptor 5 1 4095 0
qd-descriptor 6 1 4095 0
qd-descriptor 0 2 4095 0
qd-descriptor 1 2 4095 0
qd-descriptor 2 2 4095 0
qd-descriptor 4 2 4095 0
qd-descriptor 5 2 4095 0
qd-descriptor 6 2 4095 0
qd-buffer 0 1 4095
qd-buffer 1 1 4095
qd-buffer 2 1 4095
qd-buffer 4 1 4095
qd-buffer 5 1 4095
qd-buffer 6 1 4095
qd-buffer 0 2 4095
qd-buffer 1 2 4095
qd-buffer 2 2 4095
qd-buffer 4 2 4095
qd-buffer 5 2 4095
qd-buffer 6 2 4095
qd-buffer 0 1 4095 0
qd-buffer 1 1 4095 0
qd-buffer 2 1 4095 0
qd-buffer 4 1 4095 0
qd-buffer 5 1 4095 0
qd-buffer 6 1 4095 0
qd-buffer 0 2 4095 0
qd-buffer 1 2 4095 0
qd-buffer 2 2 4095 0
qd-buffer 4 2 4095 0
qd-buffer 5 2 4095 0
qd-buffer 6 2 4095 0
```

Configuring user-configurable buffer profiles on FCX and ICX devices

A buffer profile is a mechanism to arbitrarily allocate the egress buffers and descriptors limits to be applied to a port and its queues. Users can define a limit for a port and its queues by configuring the buffer profiles on the device. Therefore, user-configurable buffer profiles provide a template to allocate

egress buffers and descriptors limits to the port and on its queues. This template is then applied to the device.

NOTE

Buffer profiles can be configured for 10 Gbps and 1 Gbps ports, but not for ICX 6610 40 Gbps ports. The 10 Gbps profile will apply to ICX 6430 and 6450 stacking ports, as well as FCX 16 Gbps stacking ports.

Configuring and applying a user-configurable buffer profile is a two-step process.

First, create a user-configurable buffer profile with the qd-buffer-profile CLI command. Define a name for the user-configurable buffer profile, assign buffer and descriptor limits at the port level, and then define buffer and descriptor limits per queue of that port.

Second, apply the user-configurable buffer profile you created to the device with the buffer-profile CLI command.

Configuring a user-configurable buffer profile

To configure a user-configurable buffer profile, complete the following steps.

1. Create a user-configurable buffer profile. For example, to create a user-configurable buffer profile named profile1 enter the following command.

```
device(config)# qd-buffer-profile profile1#
```

Syntax: [no] qd-buffer-profile *user-profile-name*

The *user-profile-name* variable specifies the name of the user-configurable buffer profile. The profile name can be up to 64 characters long.

2. Configure the port type for the user-configurable buffer profile.

The buffer profile port type can be configured for 1 Gbps ports, 10 Gbps ports, or for all the ports.

The **port-type** option should be in line with the qd-buffer legacy command; that is, 1 for 1 Gbps ports, 2 for 10 Gbps ports, 0 for all the ports. The default port type is set to 1 Gbps. To configure a user-configurable profile for 10 Gbps ports, the 10 Gbps port type must be explicitly provided by the **port-type** option. Modifications to buffers and descriptors of a port and its queues take effect dynamically.

When the profile type is configured as all 1 Gbps and 10 Gbps ports, the default buffers and descriptors will be set according to the port type; that is, all 1 Gbps ports use 1 Gbps defaults and 10 Gbps ports use 10 Gbps defaults. If you configure a port and its queue with egress buffer and descriptor limits, then the configured limits are used for both 1 Gbps and 10 Gbps ports.

To configure the port type 1 for the profile named "profile1", enter the following command at the profile configuration level.

```
device(qd-profile-profile1)# profile-config port-type 1
```

Syntax: [no] profile-config port-type { 0 | 1 | 2 }

- - **0** -- All 1 Gbps and 10 Gbps ports
- - **1** -- All 1 Gbps ports
- - **2** -- All 10 Gbps ports

NOTE

The no form of the command sets the profile port type to 1 Gbps.

Port type modification resets the profile to its default value. All the port and queue buffers and descriptors will be set to either 1 Gbps or 10 Gbps defaults as per the configuration, which means all the user configurations for the port and its queues will be lost.

NOTE

Port type modifications on an active profile are not allowed.

3. Configure the port buffers. Port buffer sets the maximum buffer limit for the ports. The maximum limit depends on the hardware device. Port buffer limits of different platforms are listed in the section [Buffer and descriptor maximum and default allocation values](#) on page 54. To configure the port buffers for the user-configurable buffer profile named "profile1", enter the following command at the profile configuration level.

```
device(qd-profile-profile1)# profile-config port-buffer 8000
```

Syntax: profile-config port-buffer decimal

The decimal variable refers to the number of buffers.

4. Configure the port descriptors. Port descriptors set the maximum descriptor limit for the ports. The maximum limit of port descriptors depends on the hardware device. Port descriptor limits of different platforms are listed in [Buffer and descriptor maximum and default allocation values](#) on page 54. To configure the port descriptors for the user-configurable buffer profile named "profile1", enter the following command at the profile configuration level.

```
device(qd-profile-profile1)# profile-config port-descriptor 8000
```

Syntax: profile-config port-descriptor decimal

The decimal variable refers to the number of descriptors.

5. Configure the queue buffers. Queue buffers set the maximum buffer limit for a specified queue. If a queue buffer limit is not set, then the default is used. The system default queue buffers for different platforms are listed in [Buffer and descriptor maximum and default allocation values](#) on page 54. To configure the queue buffers for the user-configurable buffer profile named "profile1", enter the following command at the profile configuration level.

```
device(qd-profile-profile1)# profile-config queue-buffers 2 600
```

Syntax: profile-config queue-buffers queue-num decimal

The queue-num variable refers to the number of the queue of a port from 0 through 7.

The decimal variable refers to the number of buffers.

6. Configure the queue descriptors. Queue descriptors set the maximum descriptor limit for the specified queue. If a queue descriptors limit is not set, then the default is used. The system default queue descriptors for different platforms are listed in [Buffer and descriptor maximum and default allocation values](#) on page 54. To configure the queue descriptors for the user-configurable buffer profile named "profile1", enter the following command at the profile configuration level.

```
device(qd-profile-profile1)# profile-config queue-descriptors 2 600
```

Syntax: profile-config queue-descriptors queue-num decimal

The queue-num variable refers to the number of the queue of a port from 0 through 7.

The decimal variable refers to the number of descriptors.

Applying a user-configurable buffer profile on the device

Once configured, a user-configurable buffer profile must be applied on the device. To apply the user-configurable buffer profile named "profile1" to a device, enter the following command.

```
device(config)# buffer-profile port-region 0 qd-buffer-profile profile1
```

Syntax: [no] buffer-profile port-region port-region qd-buffer-profile user-profile-name

The port-region variable is the device number on which the user-configurable buffer profile is applied.
 The user-profile-name variable is the name of the user-configured profile.

Buffer and descriptor maximum and default allocation values

This section lists the maximum and default buffers and descriptors values of a port and its queues on each hardware platform. The following tables are included:

- [Table 8](#) describes FCX devices.
- [Table 9](#) describes ICX 6610 devices.
- [Table 10](#) describes ICX 6430 devices.
- [Table 11](#) describes ICX 6450 devices.

NOTE

Values in the following tables are for software traffic classes (TCs) or QoS priority (qosp) levels.

TABLE 8 Port buffers and descriptors values on FCX devices

	1 Gbps buffers and descriptors	10 Gbps buffers and descriptors
Port Limit	8096	8096
TC0	512	768
TC1	192	320
TC2	192	320
TC3	192	320
TC4	192	320
TC5	256	384
TC6	256	384
TC7	256	384

NOTE

In FCX devices, the maximum buffer value that can be configured using the qd-buffer command is 7955.

TABLE 9 Port buffer and descriptors values on ICX 6610 devices

	1 Gbps buffers and descriptors	10 Gbps buffers and descriptors	40 Gbps buffers and descriptors
Port Limit	8096	8096	8096
TC0	128	160	256

TABLE 9 Port buffer and descriptors values on ICX 6610 devices (Continued)

	1 Gbps buffers and descriptors	10 Gbps buffers and descriptors	40 Gbps buffers and descriptors
TC1	32	48	64
TC2	32	48	64
TC3	32	48	64
TC4	32	48	64
TC5	64	96	144
TC6	64	96	144
TC7	64	96	144

TABLE 10 Port buffer and descriptors values on ICX 6430 devices¹

	1 Gbps buffers	10 Gbps buffers	1Gbps descriptors	10 Gbps descriptors
Port Limit	4032	NA	3854	NA
TC0	182	NA	182	NA
TC1	182	NA	182	NA
TC2	96	NA	96	NA
TC3	96	NA	96	NA
TC4	96	NA	96	NA
TC5	128	NA	128	NA
TC6	128	NA	128	NA
TC7	128	NA	128	NA

¹ Values are the same for stand-alone and stacking systems.

TABLE 11 Port buffer and descriptors values on ICX 6450 devices¹

	1 Gbps buffers	10 Gbps buffers	1Gbps descriptors	10 Gbps descriptors
Port Limit	6143	6143	5902	5092
TC0	256	384	256	384
TC1	144	192	144	192
TC2	144	192	144	192

TABLE 11 Port buffer and descriptors values on ICX 6450 devices¹ (Continued)

	1 Gbps buffers	10 Gbps buffers	1Gbps descriptors	10 Gbps descriptors
TC3	144	192	144	192
TC4	144	192	144	192
TC5	192	256	192	256
TC6	192	256	192	256
TC7	192	256	192	256

¹ Values are the same for stand-alone and stacking systems.

Configuring values for the ICX 6430

Port buffer and descriptor values in [Table 10](#) are default values for software traffic classes. For the ICX 6430, traffic classes are mapped to shared hardware queues (refer to "Queues for the ICX 6430 switch" section in the *FastIron Ethernet Switch Traffic Management Guide*. Refer to the following are considerations for configuring your own values and buffer profiles.

- User-defined values have precedence over default values. For example, if default values for TC2, TC3, and TC4 are 96 and you configure 80 for TC3, then the hardware TC1 will be 80 (software TC2, TC3, and TC4 are mapped to hardware TC1).
- The higher user-configured value has precedence. For example, if the software TC2 default value is 96 (which is also the value of TC3 and 4), if you configure 100 for TC2 and 120 for TC3, then the hardware TC1 value will be 120 (software TC2, TC3, and TC4 are mapped to hardware TC1).

Displaying the user-configurable buffer profile configuration on FCX and ICX devices

To display the specified user-configurable buffer profile configuration, use the `show qd-buffer-profile user-profile-name` command.

To display all the user-configurable buffer profiles configured on the device, use the `show qd-buffer-profile active-profile` command.

To display all the buffer profiles configured on the device, use the `show qd-buffer-profile all` command.

The following example shows that the user-configurable buffer profile "OneGigProfile" is configured for 1 Gbps ports with the number of buffers and descriptors allocated to each queue. Note that buffers and descriptors displayed in this example are not necessarily default values.

```
device(config)# show qd-buffer-profile OneGigProfile
User Buffer Profile: OneGigProfile      Port-type: 1Gig
Total Buffers = 8096    Total Descriptors = 8096
Per Queue details:    Buffers   Descriptors
Traffic Class 0        50          38
Traffic Class 1        50          38
Traffic Class 2        50          38
Traffic Class 3        50          38
Traffic Class 4        50          38
Traffic Class 5        50          38
Traffic Class 6       132         132
Traffic Class 7        20          20
```

Syntax: `show qd-buffer user-profile-name | all`

The following table defines the fields in the output of the `show qd-buffer all` command.

TABLE 12 Field definitions for the output of show qd-buffer-profile command

Field	Description
User Buffer Profile	The name of the user-configurable buffer profile
Port-type	The type of the port: 1 Gbps or 10 Gbps or All
Total Buffers	The total number of buffers allocated to the port
Total Descriptors	The total number of descriptors allocated to the port
Per Queue details	The names of the queues
Buffers	The total number of buffers allocated to the queue
Descriptors	The total number of descriptors allocated to the queue

Configuring buffer sharing on FCX and ICX devices

Network congestion can be caused by various reasons such as port shaping, flow control received on the link due to congestion on the peer, or oversubscription of the egress line rate.

To support priority queuing, FCX and ICX devices support a configurable amount of guaranteed buffers and descriptors per (port, queue) pair. For information about guaranteed buffers and descriptors, refer to [Configuring buffer profiles](#) on page 48.

In addition, FCX and ICX devices support configurable shared buffer pools, which help absorb traffic bursts without packet loss. For a given (port, queue) pair, if its buffer usage exceeds the guaranteed limit, it will start using buffers in the sharing pool. The shared buffers are apportioned among the 1 Gbps, 10 Gbps, 16 Gbps, and stacking ports.

Buffer sharing is always enabled. You can configure buffer sharing by changing the buffer sharing level. To configure the buffer sharing level, enter a command similar to the following.

```
device(config)# qd-share-level 2
```

Syntax: **qd-share-level** /level

The level variable is the buffer sharing level, which is a decimal value /level/. The range of valid values for FCX is from 1 to 8. The range of valid values for an ICX 6610 is from 2 through 8.

For descriptions of the buffer sharing levels for FCX devices, refer to [FCX buffer sharing levels](#) on page 57.

For descriptions of the buffer sharing levels for ICX 6610 devices, refer to [ICX 6610 buffer sharing levels](#) on page 58.

For descriptions of the buffer sharing levels for ICX 6430 and 6450 devices, refer to [ICX 6430 and ICX 6450 buffer sharing levels](#) on page 59.

FCX buffer sharing levels

The FCX buffer sharing level configures the shared buffers on the device. The following table defines the FCX buffer sharing level settings. For information about configuring buffer sharing, refer to [Configuring buffer sharing on FCX and ICX devices](#) on page 57.

If you configure buffers at the port or queue level (using **qd** commands or buffer profiles), the buffer sharing level automatically changes to 1. You can change it manually.

TABLE 13 FCX buffer sharing level definitions

Buffer sharing level	Shared buffer limit (in buffers)	Shared buffer total (in kilobytes)
1	256	64
2	1024	250
3	1536	375
4	2048	500
5 (default)	2560	625
6	3072	750
7	3584	875
8	4096	1000

ICX 6610 buffer sharing levels

The ICX 6610 buffer sharing level configures the shared buffers on the device. The ICX 6610 shared buffers are divided into pools. Each of the following pools defines the buffer allocation for a set of traffic class (TC) queues:

- Pool 0 contains TCs 0 and 1.
- Pool 1 contains TCs 2, 3, and 4.
- Pool 2 contains TCs 5 and 6.
- Pool 3 contains TC 7.

For example, the buffers allocated to Pool 0 are shared between TCs 0 and 1. The following table defines the ICX 6610 buffer sharing level settings. Note that only the values of Pool 0 change. For information about configuring buffer sharing, refer to [Configuring buffer sharing on FCX and ICX devices](#) on page 57.

TABLE 14 ICX 6610 buffer sharing level definitions

Buffer sharing level	Shared buffer limit	Shared buffer total (in kilobytes)	Pool 0 sharing buffers (in kilobytes)
Pool 0 -TC 0, 1	Pool 1 - TC 2, 3, 4	Pool 2 - TC 5, 6	Pool 3 - TC 7
1	Not supported		
2	128	128	192
		192	250
3	256	128	192
		192	375
4	512	128	192
		192	500
			250

TABLE 14 ICX 6610 buffer sharing level definitions (Continued)

Buffer sharing level	Shared buffer limit					Shared buffer total (in kilobytes)	Pool 0 sharing buffers (in kilobytes)
5 (default)	768	128	192	192	625		375
6	1024	128	192	192	750		500
7	1280	128	192	192	875		625
8	1536	128	192	192	1000		750

ICX 6430 and ICX 6450 buffer sharing levels

The ICX 6430 and 6450 buffer sharing level configures the shared buffers on the device. [Table 13](#) on page 58 defines the FCX buffer sharing level settings. For information about configuring buffer sharing, refer to [Configuring buffer sharing on FCX and ICX devices](#) on page 57.

If you configure buffers at the port or queue level (using **qd** commands or buffer profiles), the buffer sharing level automatically changes to 1. You can change it manually.

TABLE 15 ICX 6430 and ICX 6450 buffer sharing level definitions

Buffer sharing level	Shared buffer limit (in buffers)	Shared buffer total (in kilobytes)
1	256	64
2 (default)	1024	250
3	1536	375
4	2048	500
5	2560	625
6	3072	750
7	3584	875
8	4096	1000

Displaying buffer sharing information

To display information about buffer sharing, enter the **show qd-share-level** command. Following is an example for an FCX device.

```
device# show qd-share-level
Sharing level: 1-64KB, 2-250KB, 3-375KB, 4-500KB, 5-625KB (default), 6-750KB,
7-875KB, 8-1000KB
Current qdsharing level 5
Device 0 Sharing buffers in use 0
Device 1 Sharing buffers in use 0
Device 1 Sharing buffers in use 0
Device 2 Sharing buffers in use 0
```

Following is an example for ICX 6610 devices.

```
ICX6610-48 Router# show qd-share-level
Sharing level: 1-64KB, 2-250KB, 3-375KB, 4-500KB, 5-625KB (default), 6-750KB,
7-875KB, 8-1000KB
Current qd sharing level 5
Sharing pools to Traffic Class (TC) map:
    Pool 0: TC 0,1   Pool 1: TC 2,3,4   Pool 2: TC 5,6   Pool 3: TC 7
Device 0 Sharing pool 0 buffers in use 0
Device 0 Sharing pool 1 buffers in use 0
Device 0 Sharing pool 2 buffers in use 0
Device 0 Sharing pool 3 buffers in use 0
Device 1 Sharing pool 0 buffers in use 0
Device 1 Sharing pool 1 buffers in use 0
Device 1 Sharing pool 2 buffers in use 0
Device 1 Sharing pool 3 buffers in use 0
Device 2 Sharing pool 0 buffers in use 0
Device 2 Sharing pool 1 buffers in use 0
Device 2 Sharing pool 2 buffers in use 0
Device 2 Sharing pool 3 buffers in use 0
```

Syntax: **show qd-share-level**

The command output displays the following information:

- Definitions of the buffer sharing levels
- The current buffer sharing level
- Mapping of traffic classes to sharing pools (ICX 6610 devices only)
- Buffer usage information

Removing buffer allocation limits on FCX and ICX

You can remove buffer allocation limits on all ports and all traffic classes globally. This permits all available buffers in a port region to be used on a first-come, first-served basis by any of its ports, regardless of priority. This can be done using the following command.

```
device(config)# buffer-sharing-full
```

Syntax: **[no] buffer-sharing-full**

The command overrides any existing configured queue depth limits and buffer allocation.

ATTENTION

Use the **buffer-sharing-full** command carefully. By entering this command, there is no limit to the number of buffers a port or a specific priority on a port can use. One port could potentially use up all the available buffers of its port region and cause starvation on other ports of the port region.

To prevent traffic loss during temporary network bursts, it is recommended that you use guaranteed (port, queue) buffers allocation or shared buffer allocation to adjust queue depth, rather than enabling the **buffer-sharing-full** command.

NOTE

The **buffer-sharing-full** command can create unpredictable behavior during traffic congestion or a blocking scenario, compromising network stability (by losing control packets), QoS, and stacking.

Buffer profiles for VoIP on FastIron stackable devices

NOTE

Configuring buffer profiles for VoIP traffic is not supported on FastIron X Series and ICX 6650 devices. Configuring VoIP buffer profiles adds buffer profiles for 1 GbE-to-100 Mbit traffic, simplifying configuration and improving performance.

VoIP profiles allows you to configure a pre-defined set of buffers and descriptors for the priority 0 and 7. These profiles support VoIP traffic that uses priority 7, with 10 Mbps or 100 Mbps uplink ports and 1000 Mbps downlink ports.

NOTE

In previous software versions, you could manually configure buffers and descriptors using **qd** commands. Buffer profiles for VoIP cannot coexist with **qd** commands. You may use buffer profiles for VoIP or **qd** command, but not both at the same time.

Configuring buffer profiles for VoIP

To configure predefined buffers, enter a command similar to the following.

```
device#buffer-profile port-region 0 voip downlink 100 uplink 1000
```

Syntax: [no] buffer-profile port-region num voip downlink 100 uplink 1000

NOTE

The port-region num can be either 0 (ports 0/1/1 to 0/1/24) or 1 (ports 0/1/25 to 0/1/48).

NOTE

FCX and ICX models (except ICX 6650) do not support this command.

Deleting buffer profiles for VoIP

To delete an existing buffer profile configuration, use the no form of the command.

```
device#no buffer-profile port-region 0 voip downlink 100 uplink 1000
```

Syntax: [no] buffer-profile port-region num voip downlink 100 uplink 1000

Buffer allocation for ICX 6650

Buffer allocation details for ICX 6650

By default, the architecture allocates fixed buffers on a per-priority queue, per-packet processor basis. The buffers control the total number of packets that can be queued in the outbound transmit for the port. In instances of heavy traffic bursts to aggregation links, such as in stacking configurations or mixed-speed environments, momentary oversubscription of the buffers and descriptors may occur. A descriptor points to one or more packet buffers.

The ingress descriptors are total of 16K buffers. Each buffer is 512 bytes. The 16K buffers are divided into 8 cores of 2K each.

The egress descriptors are divided into two pools. Pool 1 is shared by ports 1/1/9 to 1/1/56, and pool 2 is shared by the rest of ports. Each pool is 8K. Frames targetting ports that belongs to core 0-3 uses descriptors from the first pool. Frames targetting ports that belong to 4-7 uses descriptors from the second pool.

Buffer and descriptor maximum and default allocation values for ICX 6650

The maximum and default buffers and descriptor values of a port on ICX 6650

The table below lists the maximum and default buffers and descriptors values of a port and its queues on Brocade ICX 6650. Values in the following tables are for software traffic classes (TCs) or QoS priority (qosp) levels.

TABLE 16 Port buffer and descriptors values

	1 Gbps buffers and descriptors	10 Gbps buffers and descriptors	40 Gbps buffers and descriptors
Port Limit	8096	8096	8096
TC0	128	160	256
TC1	32	48	64
TC2	32	48	64
TC3	32	48	64
TC4	32	48	64
TC5	64	96	144
TC6	64	96	144
TC7	64	96	144

ICX 6650 buffer sharing levels

This section details the buffer sharing levels for ICX 6650.

For ICX 6650, the buffer sharing level configures the shared buffers on the device. The shared buffers are divided into pools. Each of the following pools defines the buffer allocation for a set of traffic class (TC) queues

- Pool 0 contains TCs 0 and 1.
- Pool 1 contains TCs 2, 3, and 4.
- Pool 2 contains TCs 5 and 6.
- Pool 3 contains TC 7.

TABLE 17 ICX 6650 buffer sharing level definitions

Shared buffer limit					
Pool 0 –TC 0, 1	Pool 1 – TC 2, 3, 4	Pool 2 – TC 5, 6	Pool 3 – TC 7	Shared buffer total (in kilobytes)	Pool 0 sharing buffers (in kilobytes)
768	128	192	192	625	375

Displaying buffer sharing information on the ICX 6650

To display information about buffer sharing, enter the **show qd-share-level** command.

```
device# show qd-share-level
Sharing pools to Traffic Class (TC) map:
Pool 0: TC 0,1 Pool 1: TC 2,3,4 Pool 2: TC 5,6 Pool 3: TC 7
Device 0 pool 0 configured buffer limit 768
Device 0 pool 1 configured buffer limit 128
Device 0 pool 2 configured buffer limit 192
Device 0 pool 3 configured buffer limit 192
Device 0 Hemisphere 0 Sharing pool 0 buffers in use 0
Device 0 Hemisphere 1 Sharing pool 0 buffers in use 0
Device 0 Hemisphere 0 Sharing pool 1 buffers in use 0
Device 0 Hemisphere 1 Sharing pool 1 buffers in use 0
Device 0 Hemisphere 0 Sharing pool 2 buffers in use 0
Device 0 Hemisphere 1 Sharing pool 2 buffers in use 0
Device 0 Hemisphere 0 Sharing pool 3 buffers in use 0
```

Ethernet loopback

The Ethernet loopback functionality provides a means to gauge the network continuity and performance of an Ethernet port.

The testing of network continuity is achieved by enabling the remote Ethernet device to swap the source MAC address with the destination MAC address and send the incoming frames back to the source. The looping of the incoming traffic back to the source allows to verify the maximum rate of frame transmission without any frame loss.

By enabling Ethernet loopback on multiple remote devices, the network performance of an entire Metro Ethernet Network (MEN) can be analyzed using a single traffic generator device installed at the network core. However, the loopback support is limited to a LAN segment.

Ethernet loopback operational modes

The Ethernet loopback functionality can be enabled on an interface and can be bound either to a specific interface port or to a port and one or more associated VLANs.

Ethernet loopback can be configured in the following modes:

- VLAN-unaware mode
- VLAN-aware mode

In VLAN-unaware mode, the Ethernet loopback configuration is at the interface level and all the frames received on the ports are looped back irrespective of any VLAN. The port does not need to be explicitly assigned as a member of any VLAN. In VLAN-aware mode, the ports must be a part of the associated

VLAN and all the frames received on the ports that are associated with a specific VLAN are looped back. The VLANs to which the port is not associated with the loopback function will continue to process traffic normally, allowing non-disruptive loopback testing.

A classification of the traffic flow can also be configured in VLAN-aware and VLAN-unaware modes. The loopback can be configured as flow-aware by specifying the source MAC address and destination MAC address on the interface. In the flow-aware configuration, only the frames received with a specific source MAC address and destination MAC address are looped back. During the loopback, the source MAC address and destination MAC address of the packets are swapped.

Ethernet loopback-enabled ports can send the incoming frames back to the source in the flow-unaware mode also. If the source MAC address and destination MAC address are not specified, all the frames received on the port are looped back and the port does not distinguish between control and data traffic and Ethernet address types (unicast, multicast, or broadcast). This makes the flow-unaware mode disruptive because control traffic is also looped back and affects other services operating on this port. However, this mode is effective when the traffic source device is directly connected to the port .

Ethernet loopback can be configured in the following combinations:

- VLAN-unaware and flow-unaware
- VLAN-unaware and flow-aware
- VLAN-aware and flow-unaware
- VLAN-aware and flow-aware

NOTE

The flow-unaware configuration is not supported on the Brocade ICX 7750, Brocade ICX 7450, and Brocade ICX 7250.

Ethernet loopback configuration considerations

The configuration considerations for Ethernet loopback are as follows:

- An interface port cannot be configured in both flow-aware and flow-unaware modes simultaneously.
- An interface port cannot be configured in both VLAN-aware and VLAN-unaware modes simultaneously.
- The source MAC address and destination MAC address which define the flow-aware configuration must be unicast MAC addresses.
- The source MAC address configured in the flow-aware configuration must be unique across the network.
- Ports can be added or removed in different Ethernet loopback modes.
- A flow-aware configuration can be added on an in-service Ethernet loopback port.
- A flow-aware configuration on a port cannot be removed from an in-service Ethernet loopback port.
- The Ethernet loopback configuration is persistent across reboots if the configuration is saved. This will help to measure switching time at reload time from a remote device.
- Ethernet loopback cannot be enabled when one or more of the following features are configured:
 - ACL
 - 802.1X port security
 - Traffic shaping
 - Dual mode
 - Rate limiting

- Ethernet loopback depends on ACL entry availability because it uses ACL resources.
- MAC learning is not supported for a packet that is looped back in devices other than the Brocade ICX 7750, Brocade ICX 7450, and Brocade ICX 7250.
- Static MAC configuration is not allowed globally when Ethernet loopback is configured in the system.
- When Ethernet loopback is enabled, the packets are looped back at the rate received. However, the packets can be dropped potentially when the device is oversubscribed.
- Ethernet loopback is supported on the physical interface and LAG interface.
- Ethernet loopback can be enabled only on an existing LAG.
- An Ethernet loopback-enabled LAG cannot be undeployed.
- An Ethernet loopback-enabled port cannot be added to an existing LAG.
- VLAN priority remarking is not allowed on an Ethernet loopback-enabled port.
- The state of the port (up or down) does not affect the Ethernet loopback functionality.
- Ethernet loopback configuration is not allowed on mult-range VLAN (MVLAN), VLAN Group, or VLAN Range.
- Ethernet loopback cannot be configured on a set of VLANs that share a Layer 2 topology (Topology Group).
- Ethernet loopback must be configured in a loop-free network for better results.
- Configuring Ethernet loopback on an MCT ICL port is not recommended as it may impact MCT operations.

Configuring Ethernet loopback in VLAN-unaware mode

The following steps configure Ethernet loopback in VLAN-unaware mode.

1. Enter the **configure terminal** command to enter global configuration mode.
device# configure terminal
2. Enter the **interface ethernet** command to enter interface configuration mode.
device(config)# interface ethernet 1/1/1
3. (Optional) Enter the **ethernet loopback test-mac** command to configure the port as flow-aware.

Once configured and when Ethernet loopback is enabled, only the frames received with the specific source MAC address and destination MAC address are looped back. Skip this step to configure flow-unaware mode.

NOTE

On Brocade ICX 7750, Brocade ICX 7450, and Brocade ICX 7250 devices, configuring the **ethernet loopback test-mac** command is mandatory because these devices support only flow-aware mode.

```
device(config-if-e1000-1/1/1)# ethernet loopback test-mac 1111.2222.3333
4444.5555.5555
```

4. Enter the **ethernet loopback** command to enable Ethernet loopback.
device(config-if-e1000-1/1/1)# ethernet loopback

The following example configures Ethernet loopback in VLAN-unaware mode as flow-aware.

```
device# configure terminal
device(config)# interface ethernet 1/1/1
device(config-if-e1000-1/1/1)# ethernet loopback test-mac 1111.2222.3333
4444.5555.5555
device(config-if-e1000-1/1/1)# ethernet loopback
```

The following example configures Ethernet loopback in VLAN-unaware mode as flow-unaware.

```
device# configure terminal  
device(config)# interface ethernet 1/1/1  
device(config-if-e1000-1/1/1)# ethernet loopback
```

Configuring Ethernet loopback in VLAN-aware mode

The following steps configure Ethernet loopback in VLAN-aware mode.

1. Enter the **configure terminal** command to enter global configuration mode.

```
device# configure terminal
```

2. (Optional) Enter the **ethernet loopback test-mac** command from interface configuration mode to configure the port as flow-aware and exit interface configuration mode.

Once configured and when Ethernet loopback is enabled, only the frames received with the specific source MAC address and destination MAC address are looped back. Skip this step to configure flow-unaware mode.

NOTE

On Brocade ICX 7750, Brocade ICX 7450, and Brocade ICX 7250 devices, configuring the **ethernet loopback test-mac** command is mandatory because these devices support only flow-aware mode. In other supported platforms, the **ethernet loopback test-mac** command is optional because you can configure flow-aware or flow-unaware mode.

```
device(config)# interface ethernet 1/1/1  
device(config-if-e1000-1/1/1)# ethernet loopback test-mac 1111.2222.3333  
4444.5555.5555  
device(config-if-e1000-1/1/1)# exit
```

3. Enter the VLAN configuration mode using the **vlan** command.

```
device(config)# vlan 100
```

4. Enter the **ethernet loopback** command by specifying the Ethernet interface to enable Ethernet loopback on one or a set of ports in a specific VLAN (VLAN-aware mode).

```
device(config-vlan-100)# ethernet loopback ethernet 1/1/1
```

The following example configures Ethernet loopback in VLAN-aware mode as flow-aware.

```
device(config)# interface ethernet 1/1/1  
device(config-if-e1000-1/1/1)# ethernet loopback test-mac 1111.2222.3333  
4444.5555.5555  
device(config-if-e1000-1/1/1)# exit  
device(config)# vlan 100  
device(config-vlan-100)# ethernet loopback ethernet 1/1/1
```

The following example configures Ethernet loopback in VLAN-aware mode as flow-unaware.

```
device(config)# vlan 100  
device(config-vlan-100)# ethernet loopback ethernet 1/1/1
```

The following example configures Ethernet loopback in VLAN-aware mode as flow-unaware on a set of ports.

```
device(config)# vlan 100  
device(config-vlan-100)# ethernet loopback ethernet 1/1/1 to 1/1/10
```

Ethernet loopback syslog messages

The syslog messages in the following table are generated when Ethernet loopback is configured or unconfigured.

TABLE 18 Ethernet loopback syslog messages

Event	Syslog output
Ethernet loopback enabled in the VLAN-aware mode	<14>0d00h56m26s:BROCADE-6430 PORT: 1/1/7 VLAN 10 enabled for ethernet loop back
Ethernet loopback disabled in the VLAN-unaware mode	<14>0d00h56m26s:BROCADE-6430 PORT: 1/1/7 VLAN N/A enabled for ethernet loop back

Remote Fault Notification on 1Gbps fiber connections

NOTE

Remote fault notification (RFN) is only available for 1 Gbps Ethernet Fiber ports. It is not available for 10/100 ports and Gbps Ethernet Copper ports. This feature is not available on ICX 6430 and ICX 6450 devices.

For fiber-optic connections, you can optionally configure a transmit port to notify the receive port on the remote device whenever the transmit port becomes disabled.

When you enable this feature, the transmit port notifies the remote port whenever the fiber cable is either physically disconnected or has failed. When this occurs and the feature is enabled, the device disables the link and turns OFF both LEDs associated with the ports.

By default, RFN is enabled.

You can configure RFN as follows:

- Globally, on the entire device
- On a trunk group
- On an individual interface

Enabling and disabling remote fault notification

RFN is ON by default. To disable RFN, use the following command.

```
device(config)#interface e 0/1/1
device(config-if-e1000-0/1/1)#gig-default neg-off
```

To re-enable RFN, use the following command.

```
device(config)#interface e 0/1/1
device(config-if-e1000-0/1/1)#gig-default auto-gig
```

Syntax: **gig-default { neg-off | auto-gig }**

For more information about the parameters supported with the **gig-default** command, see "Changing the Gbps fiber negotiation mode" section in the *FastIron Ethernet Switch Administration Guide*.

Link Fault Signaling for 10Gbps Ethernet devices

Link Fault Signaling (LFS) is a physical layer protocol that enables communication on a link between two 10 Gbps Ethernet devices. When configured on a Brocade 10 Gbps Ethernet port, the port can detect and report fault conditions on transmit and receive ports. Brocade recommends enabling LFS on both ends of a link.

LFS is disabled by default on all Brocade FastIron devices except Brocade ICX 6650.

NOTE

Enable LFS on any device prior to connecting that device with FastIron platforms. Any connecting device must have LFS currently enabled to ensure interoperability.

When LFS is enabled on an interface, the following Syslog messages are generated when the link goes up or down, or when the TX or RX fiber is removed from one or both sides of the link that has LFS enabled.

```
Interface ethernet1/1, state down - link down
Interface ethernet1/1, state up
```

When a link fault occurs, the Link and Activity LEDs turn OFF.

The Link and Activity LEDs turn ON when there is traffic traversing the link after the fiber is installed.

Enabling Link Fault Signaling

To enable Link Fault Signaling (LFS) between two 10 Gbps Ethernet devices, enter commands such as the following on both ends of the link.

```
device(config)#interface e 1/1
device(config-if-e1000-1/1)#link-fault-signal
```

Syntax: [no] link-fault-signal

Use the **no** form of the command to disable LFS.

Viewing the status of LFS-enabled links

The status of an LFS-enabled link is shown in the output of the **show interface** and **show interface brief** commands, as shown in the following examples.

```
device#show interface e 10/1
10GigabitEthernet10/1 is down (remote fault), line protocol is down
  Hardware is 10GigabitEthernet, address is 0000.0027.79d8 (bia 0000.0027.79d8)
  Configured speed 10Gbit, actual unknown, configured duplex fdx, actual unknown
  Member of L2 VLAN ID 1, port is untagged, port state is BLOCKING
  BPDU guard is Disabled, ROOT protect is Disabled
  Link Fault Signaling is Enabled, Link Error Dampening is Disabled
  STP configured to ON, priority is level0
  Flow Control is disabled
  mirror disabled, monitor disabled
Truncated for brevity...
```

The bold text in the above output shows that the LFS-enabled link (port 10/1) is down because of an error on the remote port, as indicated by **remote fault**.

Syntax: show interface ethernet port

```
device#show interfaces brief
Port Link State Dupl Speed Trunk Tag Pvid Pri MAC Name
10/1 Err-LFS
None None None None No 1 0 0000.0027.79d8
```

The bold text in the above output indicates that there is an error on the LFS-enabled link on port 10/1 and the link is down.

Syntax: show interfaces brief

Jumbo frame support

Ethernet traffic moves in units called frames. The maximum size of frames is called the Maximum Transmission Unit (MTU). When a network device receives a frame larger than its MTU, the data is either fragmented or dropped. Historically, Ethernet has a maximum frame size of 1500 bytes, so most devices use 1500 as their default MTU.

Jumbo frames are Ethernet frames with more than 1,500 bytes MTU. Conventionally, jumbo frames can carry up to 10200 bytes MTU. Brocade FastIron devices support Layer 2 jumbo frames on 10/100, 100/100/1000, and 10GbE ports.

Brocade ICX 6650 devices support Layer 2 jumbo frames on 10/100, 100/100/1000, 40GbE and 10GbE ports. Conventionally, jumbo frames can carry up to 9,000 bytes MTU. In cut-through mode, in jumbo mode, the MTU is 10200 which uses 20 buffers. In non-jumbo mode MTU is 1522 which uses 3 buffers.

Packet InError Detection

Packet InError Detection identifies links that receive more number of bad frames than configured threshold and disables them to avoid instability in the network. For instance, if a network has redundant uplinks, usually only one link is in forwarding state and the rest are redundant and blocked. If one of the redundant links becomes faulty, it may drop the PDUs and become a forwarding link. This can cause loops in the network. Packet InError Detection detects the faults in the link and disables the link to prevent loops in the network.

Packet InError Detection counts an ingress frame that has one or more of the following errors as an inError packet:

- Alignment error
- CRC error
- Oversized frame error
- Internal received MAC address error (Errors that do not fall in the above 3 types)
- Symbol error (includes the fragmented, short, or undersized frames)

You can configure the number of inError packets allowed per port in a specified sampling interval. If the port receives more than the configured number of inError packets in two consecutive sampling intervals, then the port becomes error-disabled. The output of the **show interface ethernet** command for the affected port will show the status of the port as "ERR-DISABLED (packet-inerror)".

NOTE

It is recommended to use Packet InError Detection only on required ports. If you enable this on a large number of ports in a device and use a very short sampling interval, it may lead to heavy CPU usage.

NOTE

The inError count configured on the primary port of a LAG is inherited by other member ports of the LAG. However, the LAG ports are individually sampled for inError packets. Therefore, inError packets on a port disable only that port and not the entire LAG.

NOTE

Executing commands that clear the packet counters, such as the **clear statistics** command may interfere with the proper functioning of Packet InError Detection because these commands reset the inError packet count.

Configuring Packet InError Detection

1. Run the **errdisable packet-inerror-detect** command in global configuration mode to enable the feature and to define the sampling time interval.
2. Run the **packet-inerror-detect** command in interface configuration mode of the port that you want to monitor for inError packets.
3. (*Optional*) If you want the ports to automatically recover from the error-disabled state after the expiry of a configured recovery timer, run the **errdisable recovery cause** and **errdisable recovery interval** commands in global configuration mode. For more details, see “[Enabling an error-disabled port automatically](#)” on page 15 on page 317 and [Setting the recovery interval](#) on page 318.

The following example shows the configuration of Packet InError Detection on a device and its Ethernet interface 1/1/1.

```
Brocade(config)# errdisable packet-inerror-detect interval 3
Brocade(config)# errdisable recovery cause packet-inerror-detect
Brocade(config)# errdisable recovery interval 20
Brocade(config)# interface ethernet 1/1/1
Brocade(config-if-e1000-1/1/1)# packet-inerror-detect 10
```

The ethernet interface 1/1/1 becomes disabled if more than 10 inError packets are received in each of two consecutive 3-second intervals. After the interface is disabled, it automatically recovers to the enabled state after 20 seconds.

Syslog message for error-disabled port due to inError packets

The following syslog message is generated when a port is error-disabled because of inError packets.

```
0d01h38m44s:I:PORT: 1/1/37 is ERR-DISABLED due to number of packet inErrors exceeded the threshold
```

Metro Features

• Topology groups	71
• Metro Ring Protocol	74
• VSRP	95

Topology groups

A topology group is a named set of VLANs that share a Layer 2 topology. Topology groups simplify configuration and enhance scalability of Layer 2 protocols by allowing you to run a single instance of a Layer 2 protocol on multiple VLANs.

You can use topology groups with the following Layer 2 protocols:

- STP/RSTP
- MRP
- VSRP
- 802.1W

Topology groups simplify Layer 2 configuration and provide scalability by enabling you to use the same instance of a Layer 2 protocol for multiple VLANs. For example, if a Brocade device is deployed in a Metro network and provides forwarding for two MRP rings that each contain 128 VLANs, you can configure a topology group for each ring. If a link failure in a ring causes a topology change, the change is applied to all the VLANs in the ring topology group. Without topology groups, you would need to configure a separate ring for each VLAN.

Master VLAN and member VLANs

Each topology group contains a master VLAN and can contain one or more member VLANs and VLAN groups:

- Master VLAN - The master VLAN contains the configuration information for the Layer 2 protocol. For example, if you plan to use the topology group for MRP, the topology group master VLAN contains the ring configuration information.
- Member VLANs - The member VLANs are additional VLANs that share ports with the master VLAN. The Layer 2 protocol settings for the ports in the master VLAN apply to the same ports in the member VLANs. A change to the master VLAN Layer 2 protocol configuration or Layer 2 topology affects all the member VLANs. Member VLANs do not independently run a Layer 2 protocol.
- Member VLAN groups - A VLAN group is a named set of VLANs. The VLANs within a VLAN group have the same ports and use the same values for other VLAN parameters.

When a Layer 2 topology change occurs on a port in the master VLAN, the same change is applied to that port in all the member VLANs that contain the port. For example, if you configure a topology group whose master VLAN contains ports 1/1 and 1/2, a Layer 2 state change on port 1/1 applies to port 1/1 in all the member VLANs that contain that port. However, the state change does not affect port 1/1 in VLANs that are not members of the topology group.

Control ports and free ports

A port that is in a topology group can be a control port or a free port:

- Control port - A control port is a port in the master VLAN, and is therefore controlled by the Layer 2 protocol configured in the master VLAN. The same port in all the member VLANs is controlled by the master VLAN Layer 2 protocol. Each member VLAN must contain all of the control ports and can contain additional ports.
- Free port - A free port is not controlled by the master VLAN Layer 2 protocol. The master VLAN can contain free ports. (In this case, the Layer 2 protocol is disabled on those ports.) In addition, any ports in the member VLANs that are not also in the master VLAN are free ports.

NOTE

Since free ports are not controlled by the master port Layer 2 protocol, they are assumed to always be in the Forwarding state.

Topology group configuration considerations

- You must configure the master VLAN and member VLANs or member VLAN groups before you configure the topology group.
- You can configure up to 30 (256 for ICX 6650) topology groups. Each group can control up to 4096 VLANs. A VLAN cannot be controlled by more than one topology group.
- The topology group must contain a master VLAN and can also contain individual member VLANs, VLAN groups, or a combination of individual member VLANs and VLAN groups.
- If you add a new master VLAN to a topology group that already has a master VLAN, the new master VLAN replaces the older master VLAN. All member VLANs and VLAN groups follow the Layer 2 protocol settings of the new master VLAN.
- If you remove the master VLAN (by entering **no master-vlan *vlan-id***), the software selects the new master VLAN from member VLANs. A new candidate master VLAN will be in configured order to a member VLAN so that the first added member VLAN will be a new candidate master VLAN. Once you save and reload, a member-vlan with the youngest VLAN ID will be the new candidate master. The new master VLAN inherits the Layer 2 protocol settings of the older master VLAN.
- The topology group will be deleted if the master is deleted and there are no member VLANs. This is true even if the topology group has member-groups.
- Once you add a VLAN as a member of a topology group, all the Layer 2 protocol information on the VLAN is deleted.
- A default VLAN cannot be a member of a topology group.

Configuring a topology group

To configure a topology group, enter commands such as the following.

```
device(config)#topology-group 2
device(config-topo-group-2)#master-vlan 2
device(config-topo-group-2)#member-vlan 3
device(config-topo-group-2)#member-vlan 4
device(config-topo-group-2)#member-vlan 5
device(config-topo-group-2)#member-group 2
```

These commands create topology group 2 and add the following:

- Master VLAN 2
- Member VLANs 2, 3, and 4
- Member VLAN group 2

Syntax: [no] topology-group *group-id*

The *group-id* parameter specifies the topology group ID and can be from 1 - 256.

Syntax: [no] master-vlan *vlan-id*

This command adds the master VLAN. The VLAN must already be configured. Make sure all the Layer 2 protocol settings in the VLAN are correct for your configuration before you add the VLAN to the topology group. A topology group can have only one master VLAN.

NOTE

When removing the master VLAN from the topology group, spanning tree is disabled on the master VLAN.

NOTE

If you remove the master VLAN (by entering **no master-vlan *vlan-id***), the software selects the new master VLAN from member VLANs. For example, if you remove master VLAN 2 from the example above, the CLI converts member VLAN 3 into the new master VLAN. The new master VLAN inherits the Layer 2 protocol settings of the older master VLAN.

NOTE

If you add a new master VLAN to a topology group that already has a master VLAN, the new master VLAN replaces the older master VLAN. All member VLANs and VLAN groups follow the Layer 2 protocol settings of the new master VLAN.

Syntax: [no] member-vlan *vlan-id*

The *vlan-id* parameter specifies a VLAN ID. The VLAN must already be configured.

Syntax: [no] member-group *num*

The *num* specifies a VLAN group ID. The VLAN group must already be configured.

NOTE

Once you add a VLAN or VLAN group as a member of a topology group, all the Layer 2 protocol configuration information for the VLAN or group is deleted. For example, if STP is configured on a VLAN and you add the VLAN to a topology group, the STP configuration is removed from the VLAN. Once you add the VLAN to a topology group, the VLAN uses the Layer 2 protocol settings of the master VLAN. If you remove a member VLAN or VLAN group from a topology group, you will need to reconfigure the Layer 2 protocol information in the VLAN or VLAN group.

Displaying topology group information

The following sections show how to display STP information and topology group information for VLANS.

Displaying STP information

To display STP information for a VLAN, enter a command such as the following.

```
device#show span vlan 4
VLAN 4 BPDU cam_index is 14344 and the Master DMA Are(HEX) 18 1A
STP instance owned by VLAN 2
```

This example shows STP information for VLAN 4. The line shown in bold type indicates that the VLAN STP configuration is controlled by VLAN 2. This information indicates that VLAN 4 is a member of a topology group and VLAN 2 is the master VLAN in that topology group.

Displaying topology group information

To display topology group information, enter the following command.

```
device#show topology-group
Topology Group 3
=====
master-vlan 2
member-vlan none
Common control ports      L2 protocol
ethernet 1/1               MRP
ethernet 1/2               MRP
ethernet 1/5               VSRP
ethernet 2/22              VSRP
Per vlan free ports
ethernet 2/3               Vlan 2
ethernet 2/4               Vlan 2
ethernet 2/11              Vlan 2
ethernet 2/12              Vlan 2
```

Syntax: show topology-group [group-id]

This display shows the following information.

TABLE 19 CLI display of topology group information

Field	Description
master-vlan	The master VLAN for the topology group. The settings for STP, MRP, or VSRP on the control ports in the master VLAN apply to all control ports in the member VLANs within the topology group.
member-vlan	The member VLANs in the topology group.
Common control ports	The master VLAN ports that are configured with Layer 2 protocol information. The Layer 2 protocol configuration and state of these ports in the master VLAN applies to the same port numbers in all the member VLANs.
L2 protocol	The Layer 2 protocol configured on the control ports. The Layer 2 protocol can be one of the following: <ul style="list-style-type: none"> • MRP • STP • VSRP
Per vlan free ports	The ports that are not controlled by the Layer 2 protocol information in the master VLAN.

Metro Ring Protocol

Metro Ring Protocol (MRP) is a Brocade proprietary protocol that prevents Layer 2 loops and provides fast reconvergence in Layer 2 ring topologies. It is an alternative to STP and is especially useful in Metropolitan Area Networks (MANs) where using STP has the following drawbacks:

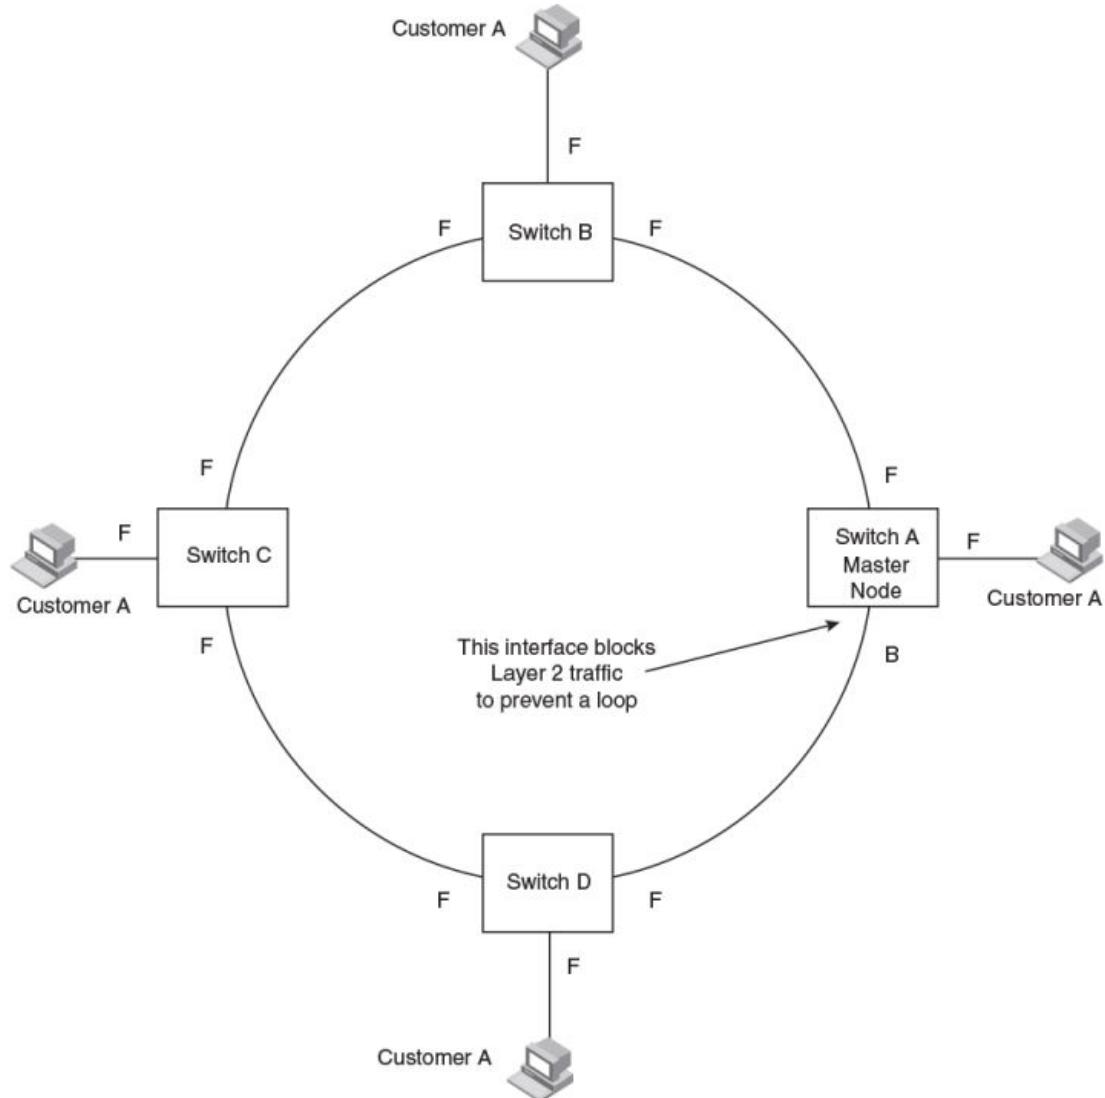
- STP allows a maximum of seven nodes. Metro rings can easily contain more nodes than this.
- STP has a slow reconvergence time, taking many seconds or even minutes. MRP can detect and heal a break in the ring in sub-second time.

NOTE

MRP portocol is not supported on the Brocade ICX-6430-C devices.

The following figure shows an example of an MRP metro ring.

FIGURE 1 Metro ring - normal state



The ring in this example consists of four MRP nodes (Brocade switches). Each node has two interfaces with the ring. Each node also is connected to a separate customer network. The nodes forward Layer 2 traffic to and from the customer networks through the ring. The ring interfaces are all in one port-based VLAN. Each customer interface can be in the same VLAN as the ring or in a separate VLAN.

One node is configured as the master node of the MRP ring. One of the two interfaces on the master node is configured as the primary interface; the other is the secondary interface. The primary interface

originates Ring Health Packets (RHPs), which are used to monitor the health of the ring. An RHP is forwarded on the ring to the next interface until it reaches the secondary interface of the master node. The secondary interface blocks the packet to prevent a Layer 2 loops.

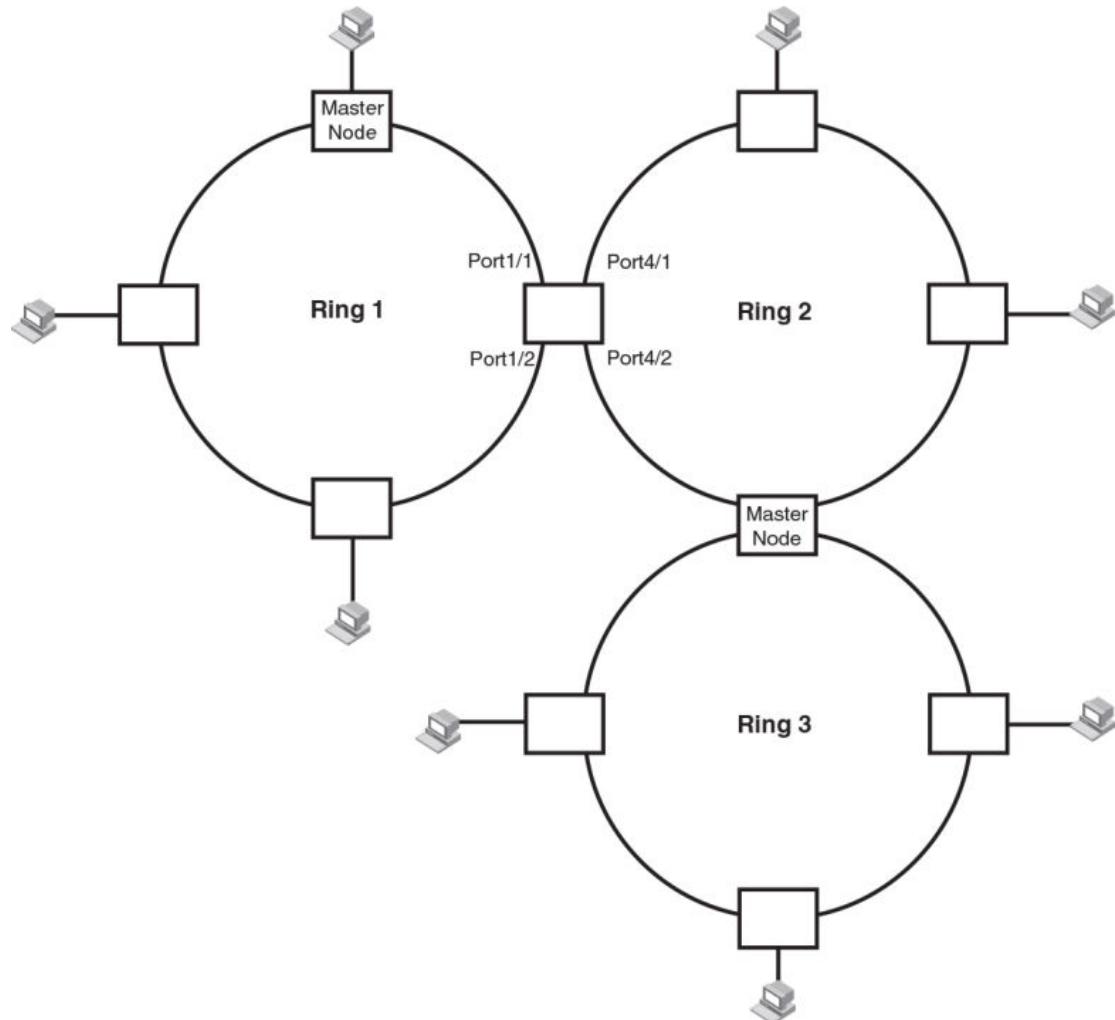
Metro Ring Protocol configuration notes

- When you configure Metro Ring Protocol (MRP), Brocade recommends that you disable one of the ring interfaces before beginning the ring configuration. Disabling an interface prevents a Layer 2 loop from occurring while you are configuring MRP on the ring nodes. Once MRP is configured and enabled on all the nodes, you can re-enable the interface.
- The above configurations can be configured as MRP masters or MRP members (for different rings).
- Brocade does not recommend configuring more than 15 MRP instances on FCX and ICX 6650 devices. Also, due to hardware limitations on this platforms, configuring 40 or more MRP instances may cause errors.
- If you configure MRP on a device running Layer 3 software, then restart the device running Layer 2 software, the MRP configuration gets deleted.

MRP rings without shared interfaces (MRP Phase 1)

MRP Phase 1 allows you to configure multiple MRP rings, as shown in [Figure 2](#), but the rings cannot share the same link. For example, you cannot configure ring 1 and ring 2 to each have interfaces 1/1 and 1/2.

Also, when you configure an MRP ring, any node on the ring can be designated as the master node for the ring. A master node can be the master node of more than one ring. (Refer to [Figure 2](#).) Each ring is an independent ring and RHP packets are processed within each ring.

FIGURE 2 Metro ring - multiple rings

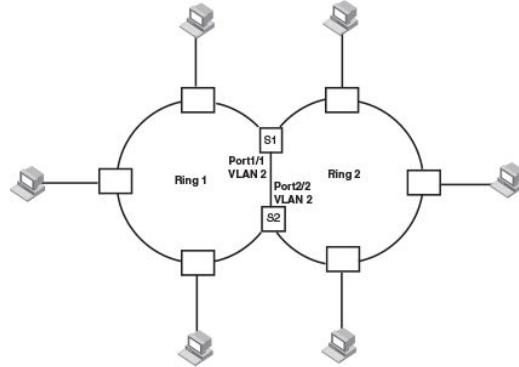
In this example, two nodes are each configured with two MRP rings. Any node in a ring can be the master for its ring. A node also can be the master for more than one ring.

MRP rings with shared interfaces (MRP Phase 2)

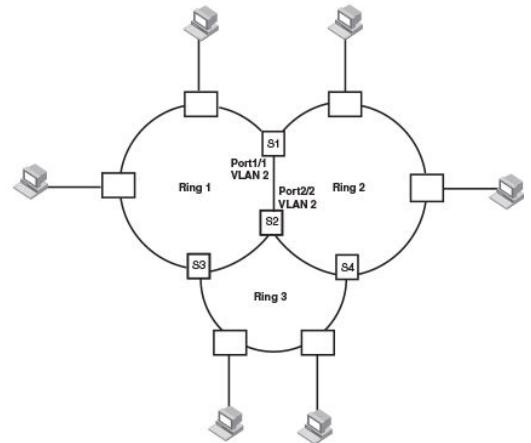
With MRP Phase 2, MRP rings can be configured to share the same interfaces as long as the interfaces belong to the same VLAN. [Figure 3](#) shows examples of multiple MRP rings that share the same interface.

FIGURE 3 Examples of multiple rings sharing the same interface - MRP Phase 2

Example 1

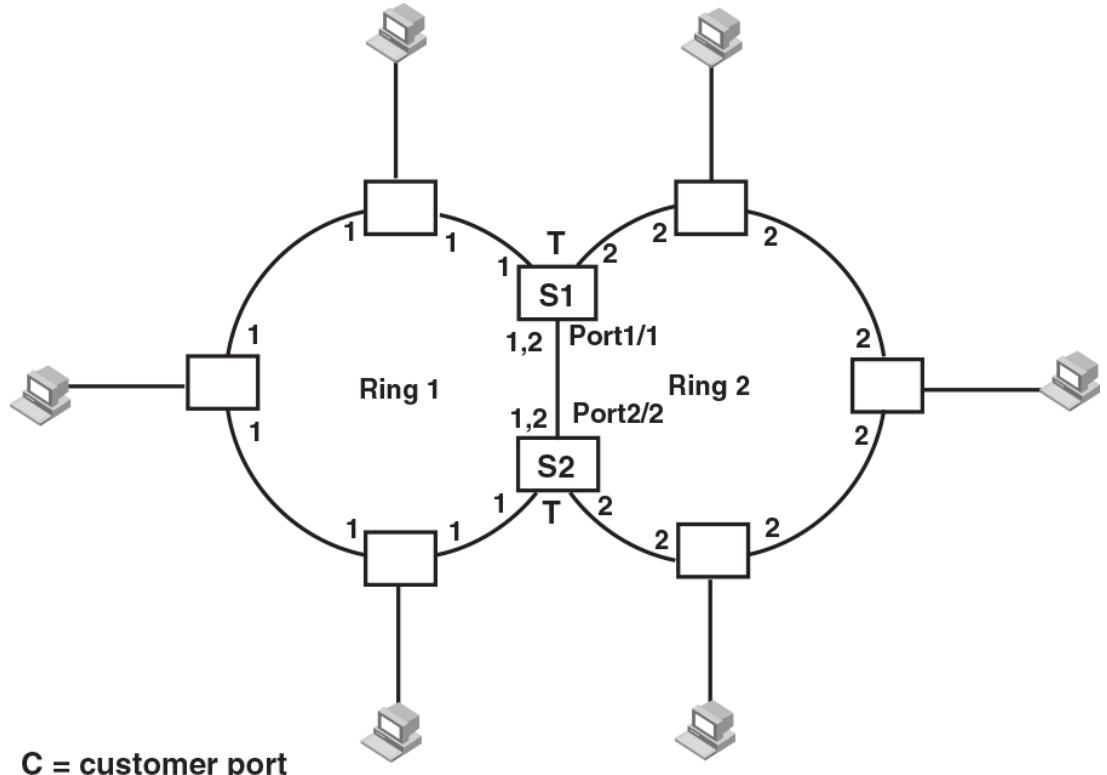


Example 2



On each node that will participate in the ring, you specify the ring ID and the interfaces that will be used for ring traffic. In a multiple ring configuration, a ring ID determines its priority. The lower the ring ID, the higher priority of a ring.

A ring ID is also used to identify the interfaces that belong to a ring.

FIGURE 4 Interface IDs and types

For example, in [Figure 4](#), the ID of all interfaces on all nodes on Ring 1 is 1 and all interfaces on all nodes on Ring 2 is 2. Port 1/1 on node S1 and Port 2/2 on S2 have the IDs of 1 and 2 since the interfaces are shared by Rings 1 and 2.

The ring ID is also used to determine an interface priority. Generally, a ring ID is also the ring priority and the priority of all interfaces on that ring. However, if the interface is shared by two or more rings, then the highest priority (lowest ID) becomes the priority of the interface. For example, in [Figure 4](#), all interfaces on Ring 1, except for Port 1/1 on node S1 and Port 2/2 on node S2 have a priority of 1. Likewise, all interfaces on Ring 2, except for Port 1/1 on node S1 and Port 2/2 on node S2 have a priority of 2. Port 1/1 on S1 and Port 2/2 on S2 have a priority of 1 since 1 is the highest priority (lowest ID) of the rings that share the interface.

If a node has interfaces that have different IDs, the interfaces that belong to the ring with the highest priority become regular ports. Those interfaces that do not belong to the ring with the highest priority become tunnel ports. In [Figure 4](#), nodes S1 and S2 have interfaces that belong to Rings 1 and 2. Those interfaces with a priority of 1 are regular ports. The interfaces with a priority of 2 are the tunnel ports since they belong to Ring 2, which has a lower priority than Ring 1.

Selection of master node

Allowing MRP rings to share interfaces limits the nodes that can be designated as the master node. Any node on an MRP ring that does not have a shared interface can be designated as the ring master node. However, if all nodes on the ring have shared interfaces, nodes that do not have tunnel ports can be designated as the master node of that ring. If none of the nodes meet these criteria, you must change the rings' priorities by reconfiguring the rings' ID.

NOTE

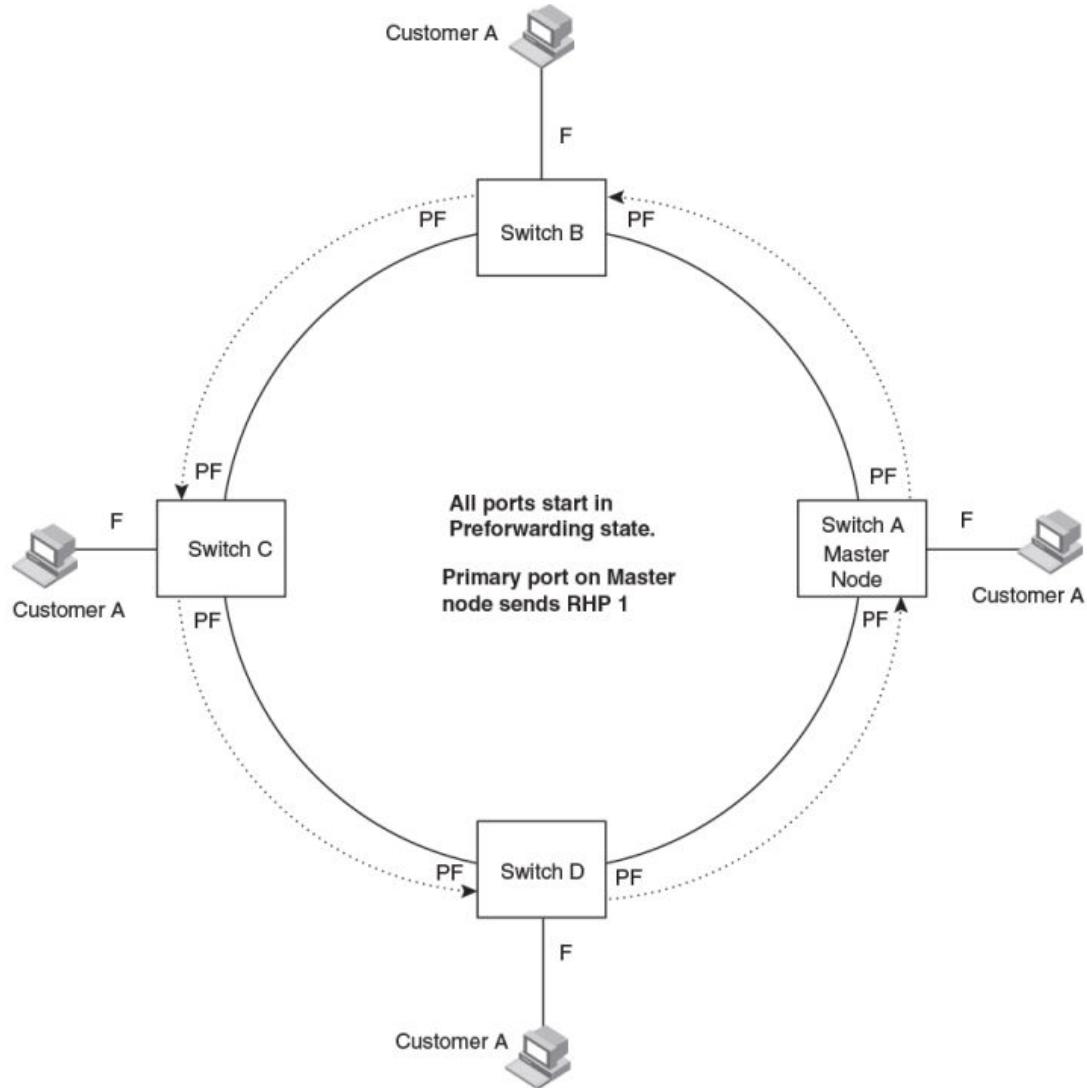
Any node on an MRP ring that has two shared interfaces cannot be elected as the master node.

In [Figure 4](#) on page 78, any of the nodes on Ring 1, even S1 or S2, can be a master node since none of its interfaces are tunnel ports. However in Ring 2, neither S1 nor S2 can be a master node since these nodes contain tunnel ports.

Ring initialization

The ring shown in [Figure 1](#) on page 75 shows the port states in a fully initialized ring without any broken links. [Figure 5](#) shows the initial state of the ring, when MRP is first enabled on the ring switches. All ring interfaces on the master node and member nodes begin in the Preforwarding state (PF).

FIGURE 5 Metro ring - initial state



MRP uses Ring Health Packets (RHPs) to monitor the health of the ring. An RHP is an MRP protocol packet. The source address is the MAC address of the master node and the destination MAC address is a protocol address for MRP. The Master node generates RHPs and sends them on the ring. The state of a ring port depends on the RHPs.

RHP processing in MRP Phase 1

A ring interface can have one of the following MRP states:

- Preforwarding (PF) - The interface can forward RHPS but cannot forward data. All ring ports begin in this state when you enable MRP.
- Forwarding (F) - The interface can forward data as well as RHPS. An interface changes from Preforwarding to Forwarding when the port preforwarding time expires. This occurs if the port does not receive an RHP from the Master, or if the forwarding bit in the RHPS received by the port is off. This indicates a break in the ring. The port heals the ring by changing its state to Forwarding. The

preforwarding time is the number of milliseconds the port will remain in the Preforwarding state before changing to the Forwarding state, even without receiving an RHP.

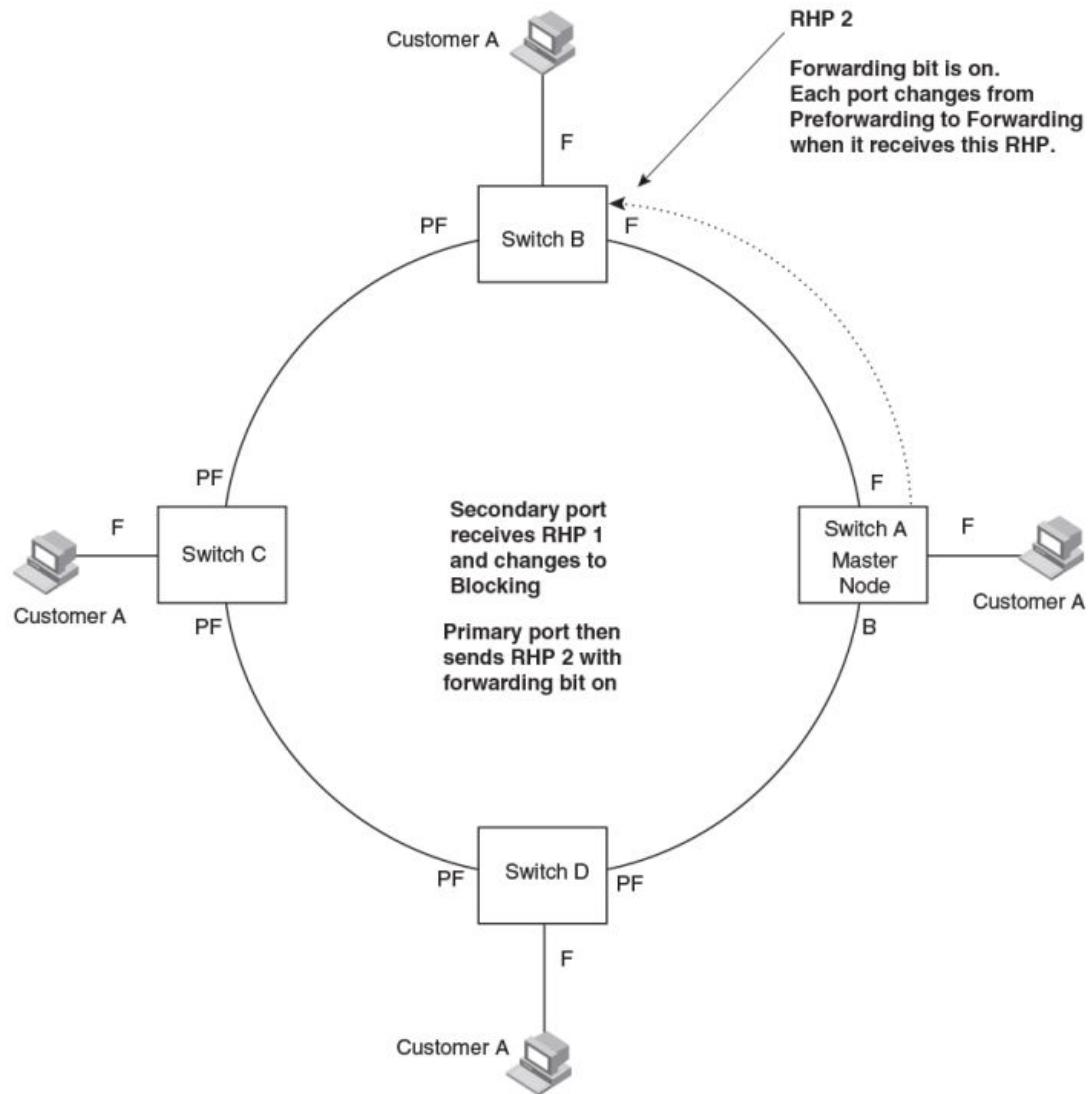
- Blocking (B) - The interface cannot forward data. Only the secondary interface on the Master node can be Blocking.

When MRP is enabled, all ports begin in the Preforwarding state. The primary interface on the Master node, although it is in the Preforwarding state like the other ports, immediately sends an RHP onto the ring. The secondary port on the Master node listens for the RHP.

- If the secondary port receives the RHP, all links in the ring are up and the port changes its state to Blocking. The primary port then sends another MRP with its forwarding bit set on. As each of the member ports receives the RHP, the ports changes their state to Forwarding. Typically, this occurs in sub-second time. The ring very quickly enters the fully initialized state.
- If the secondary port does not receive the RHP by the time the preforwarding time expires, a break has occurred in the ring. The port changes its state to Forwarding. The member ports also change their states from Preforwarding to Forwarding as their preforwarding timers expire. The ring is not intact, but data can still travel among the nodes using the links that are up.

The following figure shows an example.

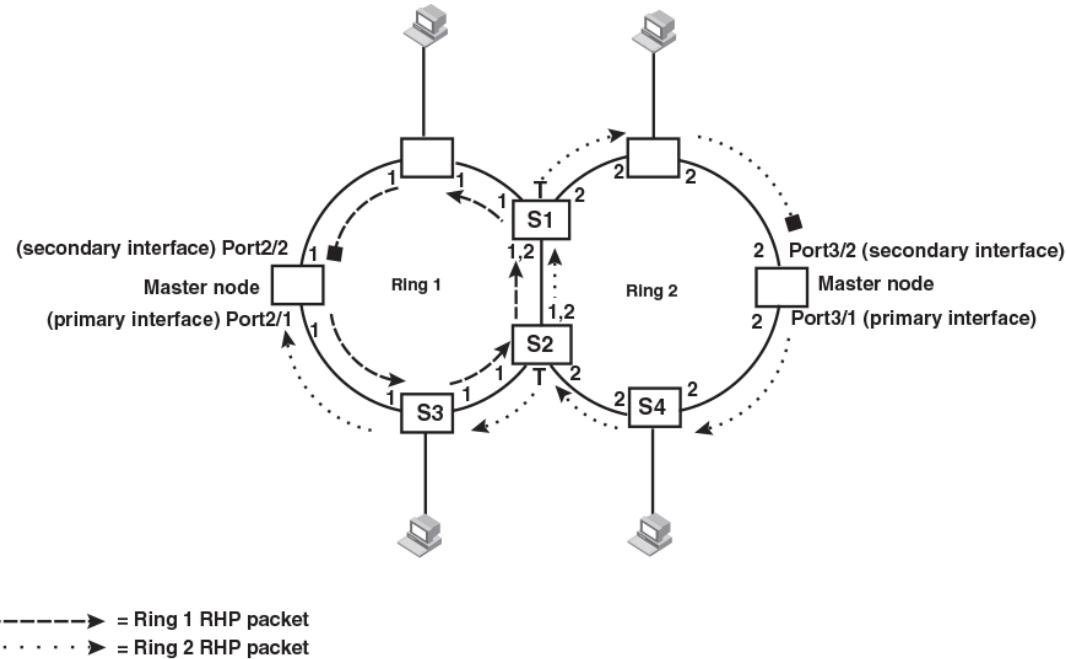
FIGURE 6 Metro ring - from preforwarding to forwarding



Each RHP also has a sequence number. MRP can use the sequence number to determine the round-trip time for RHPs in the ring. Refer to [Metro Ring Protocol diagnostics](#) on page 89.

RHP processing in MRP Phase 2

[Figure 7](#) shows an example of how RHP packets are processed normally in MRP rings with shared interfaces.

FIGURE 7 Flow of RHP packets on MRP rings with shared interfaces

Port 2/1 on Ring 1 master node is the primary interface of the master node. The primary interface forwards an RHP packet on the ring. Since all the interfaces on Ring 1 are regular ports, the RHP packet is forwarded to all the interfaces until it reaches Port 2/2, the secondary interface of the master node. Port 2/2 then blocks the packet to complete the process.

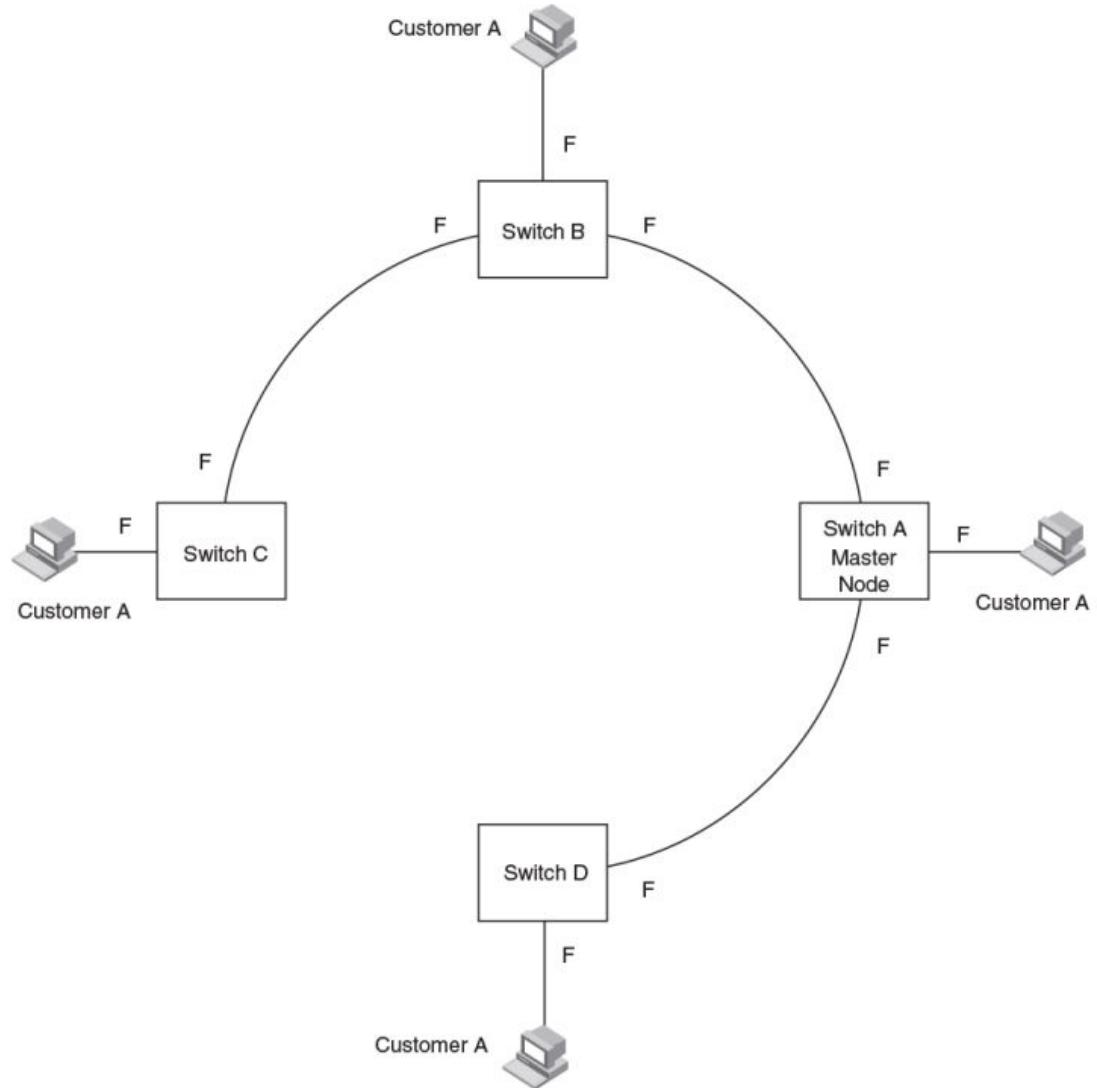
On Ring 2, Port 3/1, is the primary interface of the master node. It sends an RHP packet on the ring. Since all ports on S4 are regular ports, the RHP packet is forwarded on those interfaces. When the packet reaches S2, the receiving interface is a tunnel port. The port compares the packet priority to its priority. Since the packet priority is the same as the tunnel port priority, the packet is forwarded up the link shared by Rings 1 and 2.

When the RHP packet reaches the interface on node S2 shared by Rings 1 and 2, the packet is forwarded since its priority is less than the interface priority. The packet continues to be forwarded to node S1 until it reaches the tunnel port on S1. That tunnel port determines that the RHP packet priority is equal to the port priority and forwards the packet. The RHP packet is forwarded to the remaining interfaces on Ring 2 until it reaches port 3/2, the secondary interface of the master node. Port 3/2 then blocks the packet to prevent a loop.

When the RHP packet from Ring 2 reached S2, it was also forwarded from S2 to S3 on Ring 1 since the port on S2 has a higher priority than the RHP packet. The packets are forwarded around Ring 1 until it reaches port 2/2, Ring 1 the secondary port. The RHP packet is then blocked by that port.

How ring breaks are detected and healed

Figure 8 shows ring interface states following a link break. MRP quickly heals the ring and preserves connectivity among the customer networks.

FIGURE 8 Metro ring - ring break

If a break in the ring occurs, MRP heals the ring by changing the states of some of the ring interfaces:

- **Blocking interface** - The Blocking interface on the Master node has a dead timer. If the dead time expires before the interface receives one of its ring RHPs, the interface changes state to Preforwarding. Once the secondary interface changes state to Preforwarding:
 - If the interface receives an RHP, the interface changes back to the Blocking state and resets the dead timer.
 - If the interface does not receive an RHP for its ring before the Preforwarding time expires, the interface changes to the Forwarding state, as shown in [Figure 8](#).
- **Forwarding interfaces** - Each member interface remains in the Forwarding state.

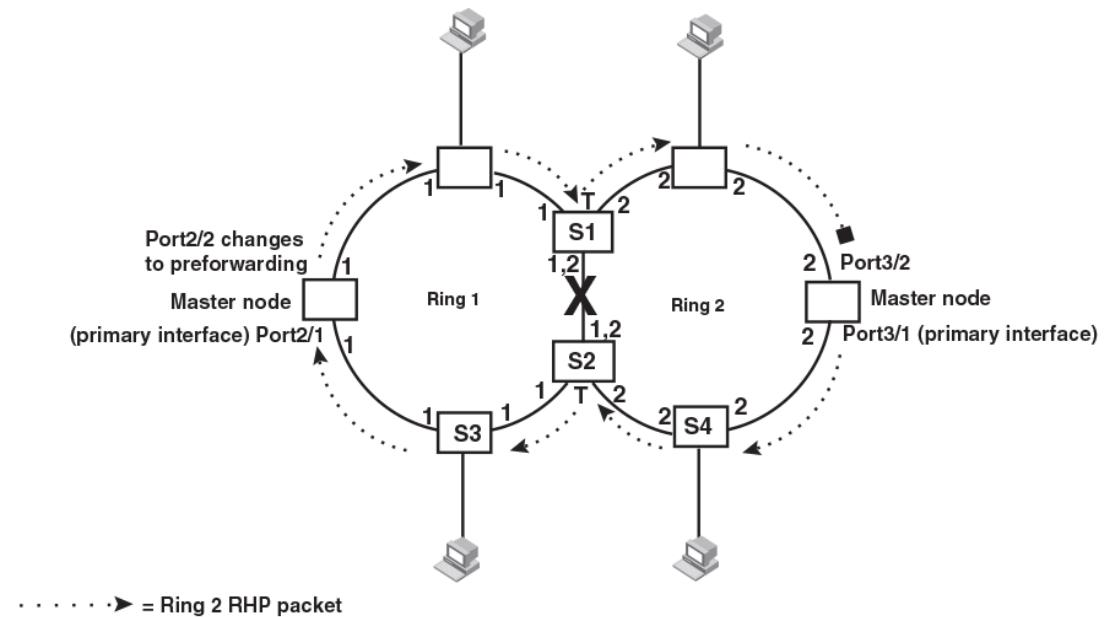
When the broken link is repaired, the link interfaces come up in the Preforwarding state, which allows RHPs to travel through the restored interfaces and reach the secondary interface on the Master node:

- If an RHP reaches the Master node secondary interface, the ring is intact. The secondary interface changes to Blocking. The Master node sets the forwarding bit on in the next RHP. When the restored interfaces receive this RHP, they immediately change state to Forwarding.
- If an RHP does not reach the Master node secondary interface, the ring is still broken. The Master node does not send an RHP with the forwarding bit on. In this case, the restored interfaces remain in the Preforwarding state until the preforwarding timer expires, then change to the Forwarding state.

If the link between shared interfaces breaks (Figure 9), the secondary interface on Ring 1 master node changes to a preforwarding state. The RHP packet sent by port 3/1 on Ring 2 is forwarded through the interfaces on S4, then to S2. The packet is then forwarded through S2 to S3, but not from S2 to S1 since the link between the two nodes is not available. When the packet reaches Ring 1 master node, the packet is forwarded through the secondary interface since it is currently in a preforwarding state. A secondary interface in preforwarding mode ignores any RHP packet that is not from its ring. The secondary interface changes to blocking mode only when the RHP packet forwarded by its primary interface is returned.

The packet then continues around Ring 1, through the interfaces on S1 to Ring 2 until it reaches Ring 2 master node. Port 3/2, the secondary interface on Ring 2 changes to blocking mode since it received its own packet, then blocks the packet to prevent a loop.

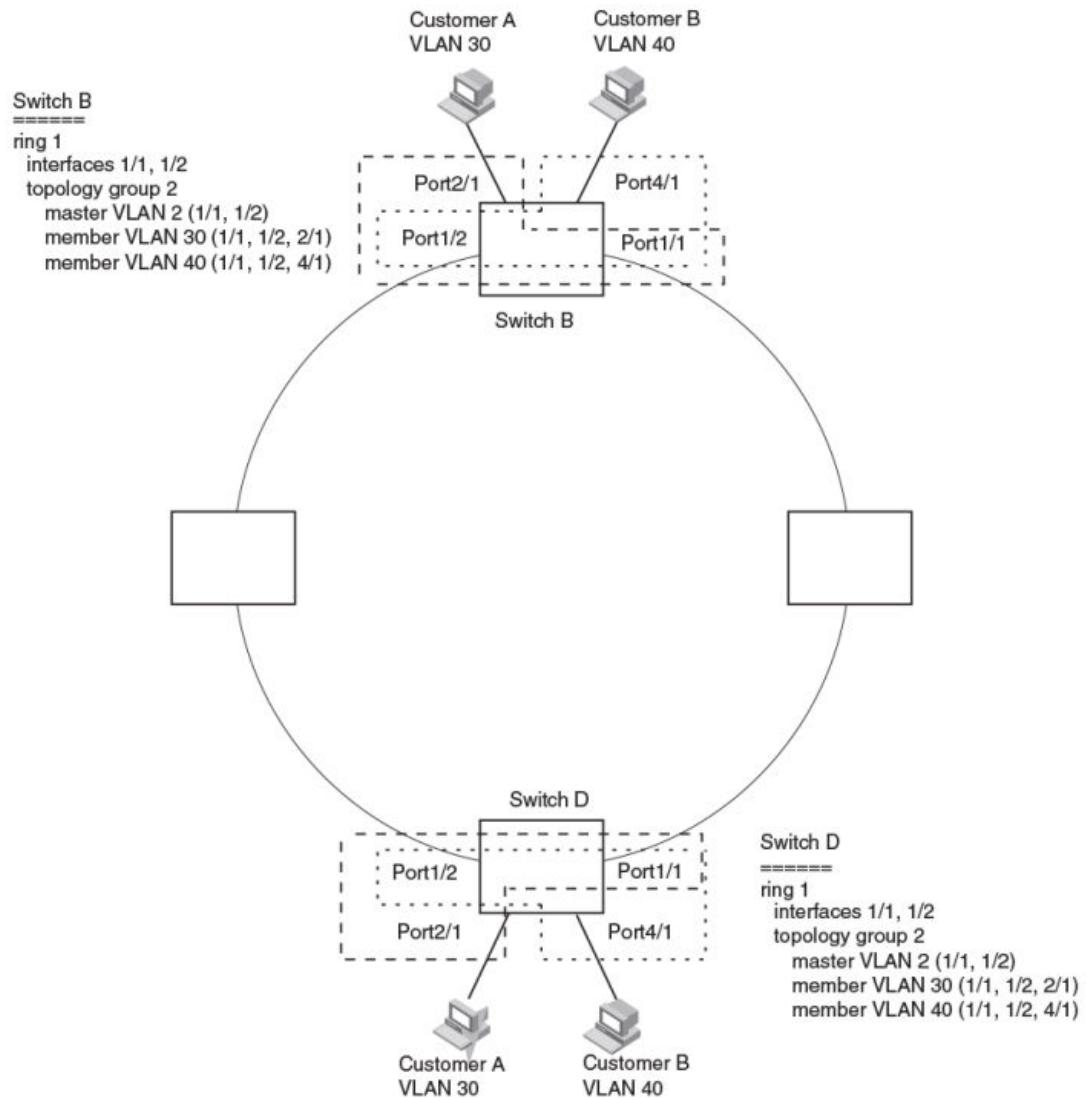
FIGURE 9 Flow of RHP packets when a link for shared interfaces breaks



RHP packets follow this flow until the link is restored; then the RHP packet returns to its normal flow as shown in Figure 7 on page 83.

Master VLANs and customer VLANs

All the ring ports must be in the same VLAN. Placing the ring ports in the same VLAN provides Layer 2 connectivity for a given customer across the ring. The following figure shows an example.

FIGURE 10 Metro ring - ring VLAN and customer VLANs

Notice that each customer has their own VLAN. Customer A has VLAN 30 and Customer B has VLAN 40. Customer A host attached to Switch D can reach the Customer A host attached to Switch B at Layer 2 through the ring. Since Customer A and Customer B are on different VLANs, they will not receive each other traffic.

You can configure MRP separately on each customer VLAN. However, this is impractical if you have many customers. To simplify configuration when you have a lot of customers (and therefore a lot of VLANs), you can use a topology group.

A topology group enables you to control forwarding in multiple VLANs using a single instance of a Layer 2 protocol such as MRP. A topology group contains a master VLAN and member VLANs. The master VLAN contains all the configuration parameters for the Layer 2 protocol (STP, MRP, or VSRP). The member VLANs use the Layer 2 configuration of the master VLAN.

In [Figure 10](#), VLAN 2 is the master VLAN and contains the MRP configuration parameters for ring 1. VLAN 30 and VLAN 40, the customer VLANs, are member VLANs in the topology group. Since a topology group is used, a single instance of MRP provides redundancy and loop prevention for both the customer VLANs.

If you use a topology group:

- The master VLAN must contain the ring interfaces. The ports must be tagged, since they will be shared by multiple VLANs.
- The member VLAN for a customer must contain the two ring interfaces and the interfaces for the customer. Since these interfaces are shared with the master VLAN, they must be tagged. Do not add another customer interfaces to the VLAN.

For more information about topology groups, refer to [Topology groups](#) on page 71.

Refer to [MRP CLI example](#) on page 93 for the configuration commands required to implement the MRP configuration shown in [Figure 10](#).

Metro Ring Protocol configuration

To configure Metro Ring Protocol (MRP), perform the following tasks. You need to perform the first task on only one of the nodes. Perform the remaining tasks on all the nodes.

NOTE

There are no new commands or parameters to configure MRP with shared interfaces (MRP Phase 2).

- Disable one of the ring interfaces. This prevents a Layer 2 loop from occurring while you are configuring the devices for MRP.
- Add an MRP ring to a port-based VLAN. When you add a ring, the CLI changes to the configuration level for the ring, where you can perform the following tasks.
 - Optionally, specify a name for the ring.
 - On the master node only, enable the device to be the master for the ring. Each ring can have only one master node.
 - Specify the MRP interfaces. Each device has two interfaces to an MRP ring.
 - Optionally, change the hello time and the preforwarding time. These parameters control how quickly failover occurs following a change in the state of a link in the ring.
 - Enable the ring.
- Optionally, add the ring VLAN to a topology group to add more VLANs to the ring. If you use a topology group, make sure you configure MRP on the group master VLAN. Refer to [Topology groups](#) on page 71.
- Re-enable the interface you disabled to prevent a Layer 2 loop. Once MRP is enabled, MRP will prevent the Layer 2 loop.
- On FCX devices, when configuring MRP-1 or MRP-2 rings on a VLAN, using the **metro-rings** command in addition to the **metro-ring** command is highly recommended. Since these devices do not support mac-range filtering, the **metro-rings** command greatly reduces the number of FDB entries.

Adding an MRP ring to a VLAN

To add an MRP ring to a VLAN, enter commands such as the following.

NOTE

If you plan to use a topology group to add VLANs to the ring, make sure you configure MRP on the topology group master VLAN.

```
device(config)#vlan 2
device(config-vlan-2)#metro-ring 1
device(config-vlan-2-mrp-1)#name CustomerA
```

```
device(config-vlan-2-mrp-1)#master
device(config-vlan-2-mrp-1)#ring-interface ethernet 1/1 ethernet 1/2
device(config-vlan-2-mrp-1)#enable
```

These commands configure an MRP ring on VLAN 2. The ring ID is 1, the ring name is CustomerA, and this node (this Brocade device) is the master for the ring. The ring interfaces are 1/1 and 1/2. Interface 1/1 is the primary interface and 1/2 is the secondary interface. The primary interface will initiate RHPs by default. The ring takes effect in VLAN 2.

```
device(config)#vlan 2
device(config-vlan-2)#metro-ring 1
device(config-vlan-2-mrp-1)#name CustomerA
device(config-vlan-2-mrp-1)#ring-interface ethernet 1/1 ethernet 1/2
device(config-vlan-2-mrp-1)#enable
device(config-vlan-2-mrp-1)#metro-ring 2
device(config-vlan-2-mrp-2)#name CustomerB
device(config-vlan-2-mrp-2)#ring-interface ethernet 1/1 ethernet 1/2
device(config-vlan-2-mrp-2)#enable
```

Syntax: [no] metro-ring *ring id*

The *ring-id* parameter specifies the ring ID. The *ring-id* can be from 1 - 1023; ID 256 is reserved for VSRP.

On FCX and ICX devices, enter the **metro-rings** in addition to the **metro-ring** command as shown below.

```
device(config)#vlan 2
device(config-vlan-2)#metro-rings 1 2
device(config-vlan-2)#metro-ring 1
device(config-vlan-2-mrp-1)#name CustomerA
device(config-vlan-2-mrp-1)#ring-interface ethernet 1/1 ethernet 1/2
device(config-vlan-2-mrp-1)#enable
device(config-vlan-2-mrp-1)#metro-ring 2
device(config-vlan-2-mrp-2)#name CustomerB
device(config-vlan-2-mrp-2)#ring-interface ethernet 1/1 ethernet 1/2
device(config-vlan-2-mrp-2)#enable
```

Syntax: [no] metro-rings *ringid ringid ...*

The *ring id* variables identify the metro rings you want to configure on the VLAN.

Syntax: [no] name *string*

The *string* parameter specifies a name for the ring. The name is optional, but it can be up to 20 characters long and can include blank spaces. If you use a name that has blank spaces, enclose the name in double quotation marks (for example: "Customer A").

Syntax: [no] master

Configures this node as the master node for the ring. Enter this command only on one node in the ring. The node is a member (non-master) node by default.

Syntax: [no] ring-interface ethernet *primary-if* ethernet *secondary-if*

The **ethernet primary-if** parameter specifies the primary interface. On the master node, the primary interface is the one that originates RHPs. Ring control traffic and Layer 2 data traffic will flow in the outward direction from this interface by default. On member nodes, the direction of traffic flow depends on the traffic direction selected by the master node. Therefore, on a member node, the order in which you enter the interfaces does not matter.

The **ethernet secondary-if** parameter specifies the secondary interface.

NOTE

To take advantage of every interface in a Metro network, you can configure another MRP ring and either configure a different Master node for the ring or reverse the configuration of the primary and

secondary interfaces on the Master node. Configuring multiple rings enables you to use all the ports in the ring. The same port can forward traffic one ring while blocking traffic for another ring.

Syntax: [no] enable

The **enable** command enables the ring.

Changing the hello and preforwarding times

You also can change the RHP hello time and preforwarding time. To do so, enter commands such as the following.

```
device(config-vlan-2-mrp-1)#hello-time 200
device(config-vlan-2-mrp-1)#preforwarding-time 400
```

These commands change the hello time to 200 ms and change the preforwarding time to 400 ms.

Syntax: [no] hello-time *ms*

Syntax: [no] preforwarding-time *ms*

The *ms* specifies the number of milliseconds. For the hello time, you can specify from 100 - 1000 (one second). The default hello time is 100 ms. The preforwarding time can be from 200 - 5000 ms, but must be at least twice the value of the hello time and must be a multiple of the hello time. The default preforwarding time is 300 ms. A change to the hello time or preforwarding time takes effect as soon as you enter the command.

Configuration notes for changing the hello and preforwarding times

- The preforwarding time must be at least twice the value of the hello time and must be a multiple of the hello time.
- If UDLD is also enabled on the device, Brocade recommends that you set the MRP preforwarding time slightly higher than the default of 300 ms; for example, to 400 or 500 ms.
- You can use MRP ring diagnostics to determine whether you need to change the hello time and preforwarding time. Refer to [Metro Ring Protocol diagnostics](#) on page 89.

Metro Ring Protocol diagnostics

The Metro Ring Protocol (MRP) diagnostics feature calculates how long it takes for RHP packets to travel through the ring. When you enable MRP diagnostics, the software tracks RHP packets according to their sequence numbers and calculates how long it takes an RHP packet to travel one time through the entire ring. When you display the diagnostics, the CLI shows the average round-trip time for the RHP packets sent since you enabled diagnostics. The calculated results have a granularity of 1 microsecond.

Enabling MRP diagnostics

To enable MRP diagnostics for a ring, enter the following command on the Master node, at the configuration level for the ring.

```
device(config-vlan-2-mrp-1)#diagnostics
```

Syntax: [no] diagnostics

NOTE

This command is valid only on the master node.

Displaying MRP diagnostics

To display MRP diagnostics results, enter the following command on the Master node.

```
device#show metro 1 diag
Metro Ring 1 - CustomerA
=====
diagnostics results
Ring      Diag      RHP average   Recommended   Recommended
id        state     time(microsec) hello time(ms) Prefwing time(ms)
2         enabled    125          100           300
Diag frame sent   Diag frame lost
1230          0
```

Syntax: show metro ring-id diag

This display shows the following information.

TABLE 20 CLI display of MRP ring diagnostic information

Field	Description
Ring id	The ring ID.
Diag state	The state of ring diagnostics.
RHP average time	The average round-trip time for an RHP packet on the ring. The calculated time has a granularity of 1 microsecond.
Recommended hello time	The hello time recommended by the software based on the RHP average round-trip time.
Recommended Prefwing time	The preforwarding time recommended by the software based on the RHP average round-trip time.
Diag frame sent	The number of diagnostic RHPs sent for the test.
Diag frame lost	The number of diagnostic RHPs lost during the test.

If the recommended hello time and preforwarding time are different from the actual settings and you want to change them, refer to [Metro Ring Protocol configuration](#) on page 87.

Displaying MRP information

You can display the following MRP information:

- Topology group configuration information
- Ring configuration information and statistics

Displaying topology group information

To display topology group information, enter the following command.

Syntax: **show topology-group [group-id]**

Refer to [Displaying topology group information](#) on page 73 for more information.

Displaying ring information

To display ring information, enter the following command.

```
device#show metro
Metro Ring 1
=====
Ring      State     Ring      Master      Topo      Hello      Prefwing
id        role      vlan       group      time(ms)   time(ms)
2         enabled   member    2          not conf   100        300
Ring interfaces      Interface role      Forwarding state      Active interface
Interface Type
ethernet 1/1      primary      disabled
ethernet 1/2      secondary
RHPs sent          RHPs rcvd
3                  0           0
                                         0
                                         4
                                         State changes
                                         Regular Tunnel
```

Syntax: **show metro [ring-id]**

This display shows the following information.

TABLE 21 CLI display of MRP ring information

Field	Description
Ring id	The ring ID
State	The state of MRP. The state can be one of the following: <ul style="list-style-type: none"> • enabled - MRP is enabled • disabled - MRP is disabled
Ring role	Whether this node is the master for the ring. The role can be one of the following: <ul style="list-style-type: none"> • master • member
Master vlan	The ID of the master VLAN in the topology group used by this ring. If a topology group is used by MRP, the master VLAN controls the MRP settings for all VLANs in the topology group.
NOTE	
The topology group ID is 0 if the MRP VLAN is not the master VLAN in a topology group. Using a topology group for MRP configuration is optional.	
Topo group	The topology group ID.
Hello time	The interval, in milliseconds, at which the Forwarding port on the ring master node sends Ring Hello Packets (RHPs).

TABLE 21 CLI display of MRP ring information (Continued)

Field	Description
Prefwing time	<p>The number of milliseconds an MRP interface that has entered the Preforwarding state will wait before changing to the Forwarding state.</p> <p>If a member port in the Preforwarding state does not receive an RHP within the Preforwarding time (Prefwing time), the port assumes that a topology change has occurred and changes to the Forwarding state.</p> <p>The secondary port on the Master node changes to Blocking if it receives an RHP, but changes to Forwarding if the port does not receive an RHP before the preforwarding time expires.</p>
NOTE	
A member node Preforwarding interface also changes from Preforwarding to Forwarding if it receives an RHP whose forwarding bit is on.	
Ring interfaces	The device two interfaces with the ring.
NOTE	
If the interfaces are trunk groups, only the primary ports of the groups are listed.	
Interface role	<p>The interface role can be one of the following:</p> <ul style="list-style-type: none"> • primary <ul style="list-style-type: none"> - Master node - The interface generates RHPs. - Member node - The interface forwards RHPs received on the other interface (the secondary interface). • secondary - The interface does not generate RHPs. <ul style="list-style-type: none"> - Master node - The interface listens for RHPs. - Member node - The interface receives RHPs.
Forwarding state	<p>Whether MRP Forwarding is enabled on the interface. The forwarding state can be one of the following:</p> <ul style="list-style-type: none"> • blocking - The interface is blocking Layer 2 data traffic and RHPs • disabled - The interface is down • forwarding - The interface is forwarding Layer 2 data traffic and RHPs • preforwarding - The interface is listening for RHPs but is blocking Layer 2 data traffic
Active interface	The physical interfaces that are sending and receiving RHPs.
NOTE	
If a port is disabled, its state is shown as "disabled".	
NOTE	
If an interface is a trunk group, the member port which comes up first is listed.	

TABLE 21 CLI display of MRP ring information (Continued)

Field	Description
Interface Type	Shows if the interface is a regular port or a tunnel port.
RHPs sent	The number of RHPs sent on the interface.
NOTE	
	This field applies only to the master node. On non-master nodes, this field contains 0. This is because the RHPs are forwarded in hardware on the non-master nodes.
RHPs rcvd	The number of RHPs received on the interface.
NOTE	
	On most Brocade devices, this field applies only to the master node. On non-master nodes, this field contains 0. This is because the RHPs are forwarded in hardware on the non-master nodes. However, on the FastIron devices, the RHP received counter on non-master MRP nodes increment. This is because, on FastIron devices, the CPU receives a copy of the RHPs forwarded in hardware.
TC RHPs rcvd	The number of Topology Change RHPs received on the interface. A Topology Change RHP indicates that the ring topology has changed.
State changes	The number of MRP interface state changes that have occurred. The state can be one of the states listed in the Forwarding state field.
Interface Type	Shows if the interface is a regular port or a tunnel port.

MRP CLI example

The following examples show the CLI commands required to implement the MRP configuration shown in [Figure 10](#) on page 86.

NOTE

For simplicity, the figure shows the VLANs on only two switches. The CLI examples implement the ring on all four switches.

MRP commands on Switch A (master node)

The following commands configure a VLAN for the ring. The ring VLAN must contain both of the node interfaces with the ring. Add these interfaces as tagged interfaces, since the interfaces also must be in each of the customer VLANs configured on the node.

```
device(config)#vlan 2
device(config-vlan-2)#tag ethernet 1/1 to 1/2
device(config-vlan-2)#metro-ring 1
device(config-vlan-2-mrp-1)#name "Metro A"
device(config-vlan-2-mrp-1)#master
device(config-vlan-2-mrp-1)#ring-interface ethernet 1/1 ethernet 1/2
```

```
device(config-vlan-2-mrp-1)#enable
device(config-vlan-2-mrp-1)#exit
device(config-vlan-2)#exit
```

The following commands configure the customer VLANs. The customer VLANs must contain both the ring interfaces as well as the customer interfaces.

```
device(config)#vlan 30
device(config-vlan-30)#tag ethernet 1/1 to 1/2
device(config-vlan-30)#tag ethernet 2/1
device(config-vlan-30)#exit
device(config)#vlan 40
device(config-vlan-40)#tag ethernet 1/1 to 1/2
device(config-vlan-40)#tag ethernet 4/1
device(config-vlan-40)#exit
```

The following commands configure topology group 1 on VLAN 2. The master VLAN is the one that contains the MRP configuration. The member VLANs use the MRP parameters of the master VLAN. The control interfaces (the ones shared by the master VLAN and member VLAN) also share MRP state.

```
device(config)#topology-group 1
device(config-topo-group-1)#master-vlan 2
device(config-topo-group-1)#member-vlan 30
device(config-topo-group-1)#member-vlan 40
```

MRP commands on Switch B

The commands for configuring Switches B, C, and D are similar to the commands for configuring Switch A, with two differences: the nodes are not configured to be the ring master. Omitting the **master** command is required for non-master nodes.

```
device(config)#vlan 2
device(config-vlan-2)#tag ethernet 1/1 to 1/2
device(config-vlan-2)#metro-ring 1
device(config-vlan-2-mrp-1)#name "Metro A"
device(config-vlan-2-mrp-1)#ring-interface ethernet 1/1 ethernet 1/2
device(config-vlan-2-mrp-1)#enable
device(config-vlan-2)#exit
device(config)#vlan 30
device(config-vlan-30)#tag ethernet 1/1 to 1/2
device(config-vlan-30)#tag ethernet 2/1
device(config-vlan-30)#exit
device(config)#vlan 40
device(config-vlan-40)#tag ethernet 1/1 to 1/2
device(config-vlan-40)#tag ethernet 4/1
device(config-vlan-40)#exit
device(config)#topology-group 1
device(config-topo-group-1)#master-vlan 2
device(config-topo-group-1)#member-vlan 30
device(config-topo-group-1)#member-vlan 40
```

MRP commands on Switch C

```
device(config)#vlan 2
device(config-vlan-2)#tag ethernet 1/1 to 1/2
device(config-vlan-2)#metro-ring 1
device(config-vlan-2-mrp-1)#name "Metro A"
device(config-vlan-2-mrp-1)#ring-interface ethernet 1/1 ethernet 1/2
device(config-vlan-2-mrp-1)#enable
device(config-vlan-2)#exit
device(config)#vlan 30
device(config-vlan-30)#tag ethernet 1/1 to 1/2
device(config-vlan-30)#tag ethernet 2/1
device(config-vlan-30)#exit
device(config)#vlan 40
device(config-vlan-40)#tag ethernet 1/1 to 1/2
device(config-vlan-40)#tag ethernet 4/1
```

```

device(config-vlan-40)#exit
device(config)#topology-group 1
device(config-topo-group-1)#master-vlan 2
device(config-topo-group-1)#member-vlan 30
device(config-topo-group-1)#member-vlan 40

```

MRP commands on Switch D

```

device(config)#vlan 2
device(config-vlan-2)#tag ethernet 1/1 to 1/2
device(config-vlan-2)#metro-ring 1
device(config-vlan-2-mrp-1)#name "Metro A"
device(config-vlan-2-mrp-1)#ring-interface ethernet 1/1 ethernet 1/2
device(config-vlan-2-mrp-1)#enable
device(config-vlan-2)#exit
device(config)#vlan 30
device(config-vlan-30)#tag ethernet 1/1 to 1/2
device(config-vlan-30)#tag ethernet 2/1
device(config-vlan-30)#exit
device(config)#vlan 40
device(config-vlan-40)#tag ethernet 1/1 to 1/2
device(config-vlan-40)#tag ethernet 4/1
device(config-vlan-40)#exit
device(config)#topology-group 1
device(config-topo-group-1)#master-vlan 2
device(config-topo-group-1)#member-vlan 30
device(config-topo-group-1)#member-vlan 40

```

VSRP

Virtual Switch Redundancy Protocol (VSRP) is a Brocade proprietary protocol that provides redundancy and sub-second failover in Layer 2 and Layer 3 mesh topologies. Based on the Brocade Virtual Router Redundancy Protocol Extended (VRRP-E), VSRP provides one or more backups for a Layer 2 Switch or Layer 3 Switch. If the active Layer 2 Switch or Layer 3 Switch becomes unavailable, one of the backups takes over as the active device and continues forwarding traffic for the network.

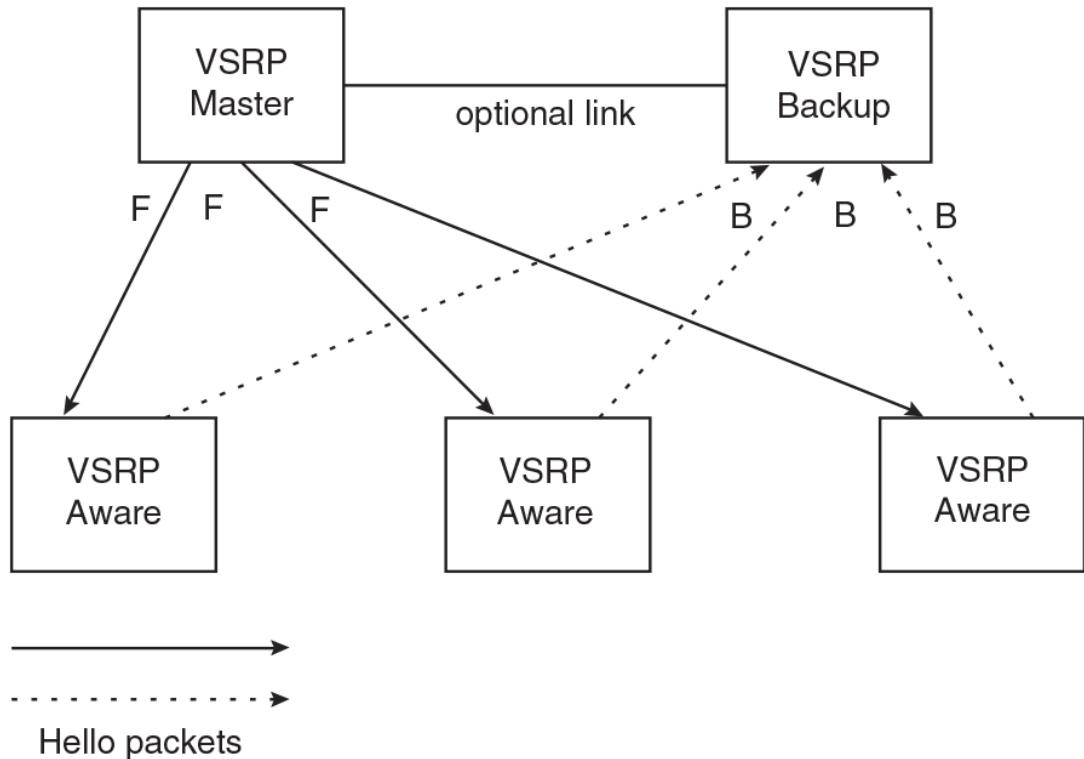
The FastIron family of switches support full VSRP as well as VSRP-awareness . A Brocade device that is not itself configured for VSRP but is connected to a Brocade device that is configured for VSRP, is **VSRP aware** .

NOTE

VSRP protocol is not supported on the Brocade ICX 6430-C devices. Only VSRP-awareness is supported.

You can use VSRP for Layer 2, Layer 3, or for both layers. On Layer 3 Switches, Layer 2 and Layer 3 share the same VSRP configuration information. On Layer 2 Switches, VSRP applies only to Layer 2.

The following example shows an example of a VSRP configuration.

FIGURE 11 VSRP mesh - redundant paths for Layer 2 and Layer 3 traffic

In this example, two Brocade devices are configured as redundant paths for VRID 1. On each of the devices, a Virtual Router ID (VRID) is configured on a port-based VLAN. Since VSRP is primarily a Layer 2 redundancy protocol, the VRID applies to the entire VLAN. However, you can selectively remove individual ports from the VRID if needed.

Following Master election (described below), one of the Brocade devices becomes the Master for the VRID and sets the state of all the VLAN ports to Forwarding. The other device is a Backup and sets all the ports in its VRID VLAN to Blocking.

If a failover occurs, the Backup becomes the new Master and changes all its VRID ports to the Forwarding state.

NOTE

The link between VSRP Master and VSRP Backup is "optional" in Figure 11. However, if the VSRP-aware device is a Brocade FastIron family device, such as ICX 6610, FCX/SX/ ICX 6430/6450 and ICX 6650, this link is required and recommended. This is due to the need for interoperability between devices of these two platforms having different default timers. The link between the VSRP Master and Backup guarantees that the VSRP Hello message is flowing between the VSRP Master and the VSRP Standby directly to cause VSRP transition instead of relying on VSRP-Aware devices to forward and risk missing the VSRP Hello message.

Other Brocade devices can use the redundant paths provided by the VSRP devices. In this example, three Brocade devices use the redundant paths. A Brocade device that is not itself configured for VSRP but is connected to a Brocade device that is configured for VSRP, is VSRP aware. In this example, the three Brocade devices connected to the VSRP devices are VSRP aware. A Brocade device that is VSRP aware can failover its link to the new Master in sub-second time, by changing the MAC address associated with the redundant path.

When you configure VSRP, make sure each of the non-VSRP Brocade devices connected to the VSRP devices has a separate link to each of the VSRP devices.

VSRP configuration notes and feature limitations

- VSRP and 802.1Q-n-Q tagging are not supported together on the same device.
- VSRP and Super Aggregated VLANs are not supported together on the same device.

NOTE

802.1Q-n-Q tagging and Aggregated VLANs are not supported on the Brocade ICX 6430-C devices.

- The VLAN supports IGMP snooping version 2 and version 3 when VSRP or VSRP-aware is configured on a VLAN.

Layer 2 and Layer 3 redundancy

You can configure VSRP to provide redundancy for Layer 2 only or also for Layer 3:

- Layer 2 only - The Layer 2 links are backed up but specific IP addresses are not backed up.
- Layer 2 and Layer 3 - The Layer 2 links are backed up and a specific IP address is also backed up. Layer 3 VSRP is the same as VRRP-E. However, using VSRP provides redundancy at both layers at the same time.

Layer 2 Switches support Layer 2 VSRP only. Layer 3 Switches support Layer 2 and Layer 3 redundancy. You can configure a Layer 3 Switch for either Layer 2 only or Layer 2 and Layer 3. To configure for Layer 3, specify the IP address you are backing up.

NOTE

If you want to provide Layer 3 redundancy only, disable VSRP and use VRRP-E.

Master election and failover

Each VSRP device advertises its VSRP priority in Hello messages. During Master election, the VSRP device with the highest priority for a given VRID becomes the Master for that VRID. After Master election, the Master sends Hello messages at regular intervals to inform the Backups that the Master is healthy.

If there is a tie for highest VSRP priority, the tie is resolved as follows:

- Layer 2 Switches - The Layer 2 Switch with the higher management IP address becomes the Master.
 - Switches with management IP addresses are preferred over switches without management IP addresses.
 - If neither of the switches has a management IP address, then the switch with the higher MAC address becomes the Master. (VSRP compares the MAC addresses of the ports configured for the VRID, not the base MAC addresses of the switches.)
- Layer 3 Switches - The Layer 3 Switch whose virtual routing interface has a higher IP address becomes the master.

VSRP failover

Each Backup listens for Hello messages from the Master. The Hello messages indicate that the Master is still available. If the Backups stop receiving Hello messages from the Master, the election process occurs again and the Backup with the highest priority becomes the new Master.

Each Backup waits for a specific period of time, the Dead Interval, to receive a new Hello message from the Master. If the Backup does not receive a Hello message from the Master by the time the Dead Interval expires, the Backup sends a Hello message of its own, which includes the Backup's VSRP priority, to advertise the Backup's intent to become the Master. If there are multiple Backups for the VRID, each Backup sends a Hello message.

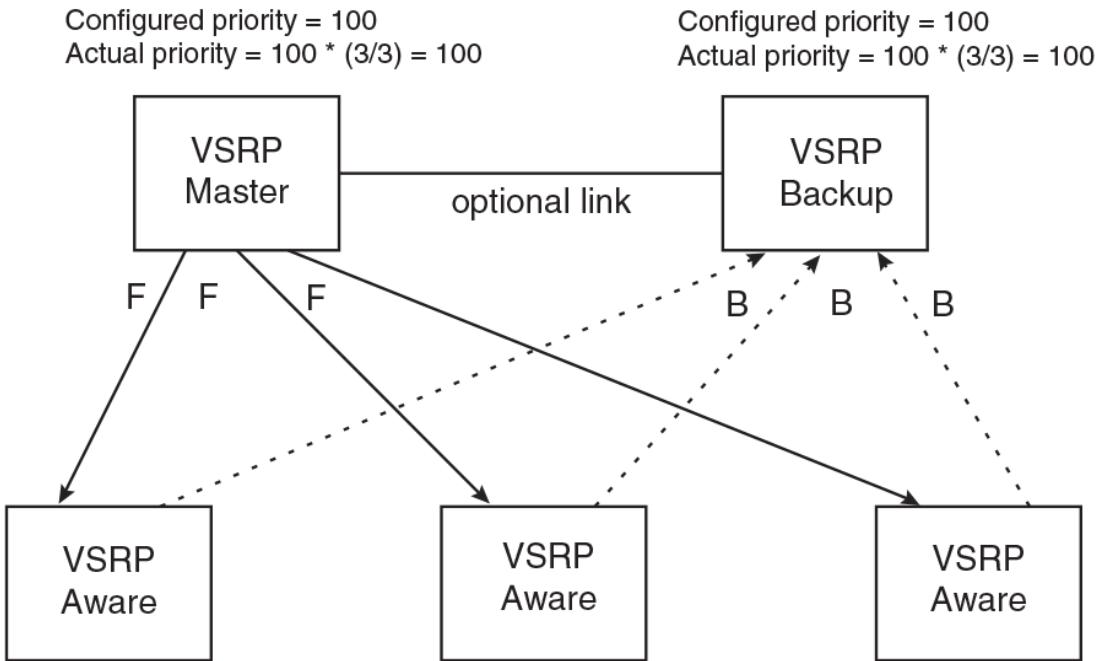
When a Backup sends a Hello message announcing its intent to become the Master, the Backup also starts a hold-down timer. During the hold-down time, the Backup listens for a Hello message with a higher priority than its own.

- If the Backup receives a Hello message with a higher priority than its own, the Backup resets its Dead Interval and returns to normal Backup status.
- If the Backup does not receive a Hello message with a higher priority than its own by the time the hold-down timer expires, the Backup becomes the new Master and starts forwarding Layer 2 traffic on all ports.

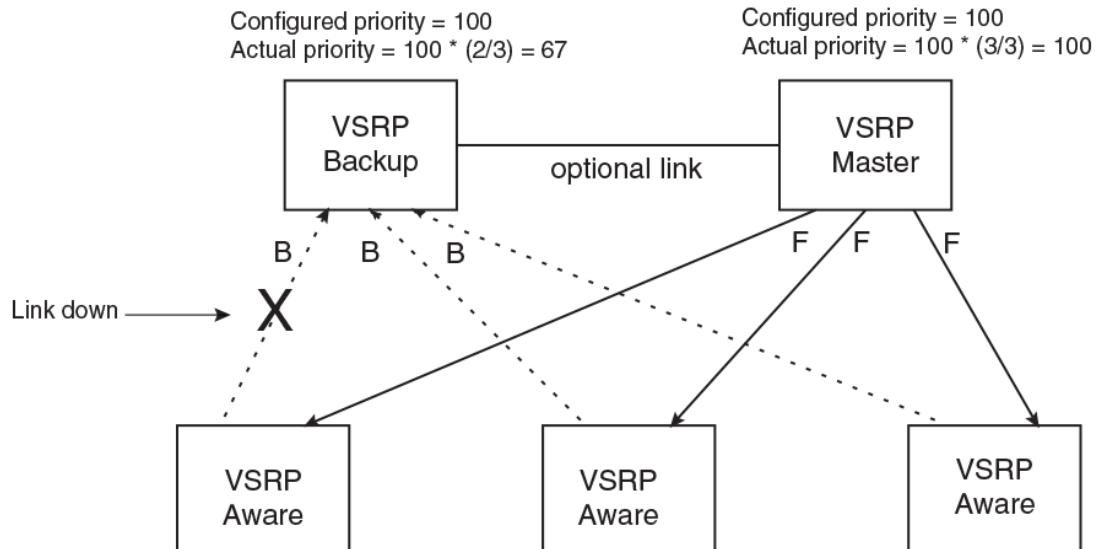
If you increase the timer scale value, each timer value is divided by the scale value. To achieve sub-second failover times, you can change the scale to a value up to 10. This shortens all the VSRP timers to 10 percent of their configured values.

VSRP priority calculation

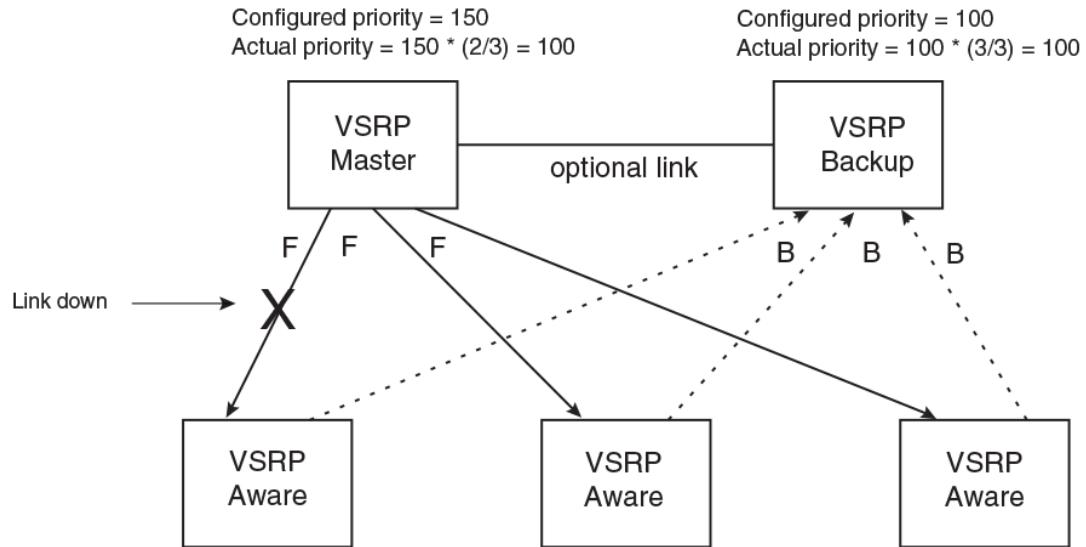
Each VSRP device has a VSRP priority for each VRID and its VLAN. The VRID is used during Master election for the VRID. By default, a device VSRP priority is the value configured on the device (which is 100 by default). However, to ensure that a Backup with a high number of up ports for a given VRID is elected, the device reduces the priority if a port in the VRID VLAN goes down. For example, if two Backups each have a configured priority of 100, and have three ports in VRID 1 in VLAN 10, each Backup begins with an equal priority, 100. This is shown in the following figure.

FIGURE 12 VSRP priority

However, if one of the VRID ports goes down on one of the Backups, that Backup priority is reduced. If the Master priority is reduced enough to make the priority lower than a Backup priority, the VRID fails over to the Backup. The following figure shows an example.

FIGURE 13 VSRP priority recalculations

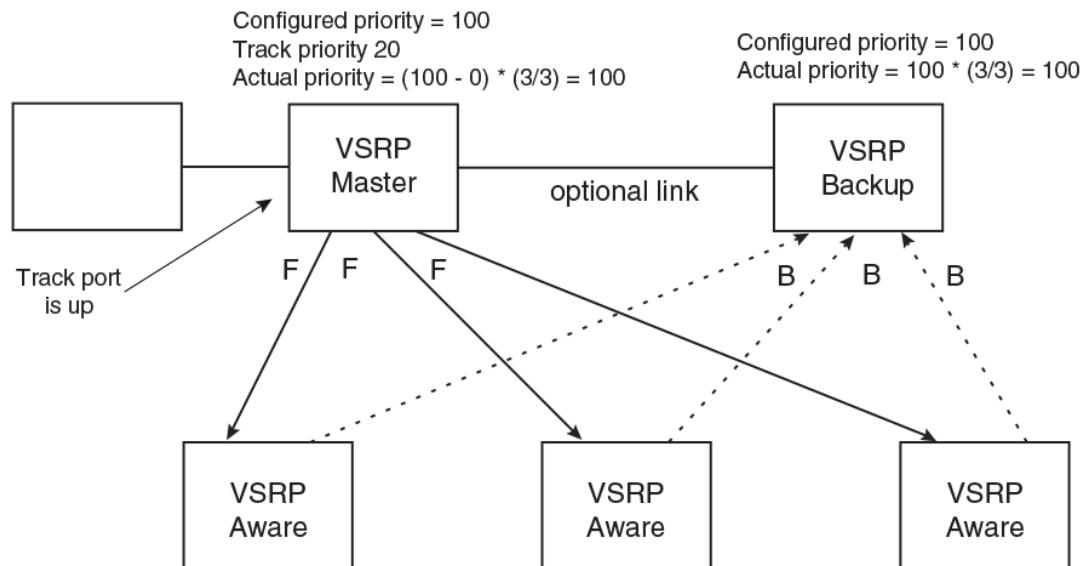
You can reduce the sensitivity of a VSRP device to failover by increasing its configured VSRP priority. For example, you can increase the configured priority of the VSRP device on the left in [Figure 13](#) to 150. In this case, failure of a single link does not cause failover. The link failure caused the priority to be reduced to 100, which is still equal to the priority of the other device. This is shown in the following figure.

FIGURE 14 VSRP priority bias

Track ports

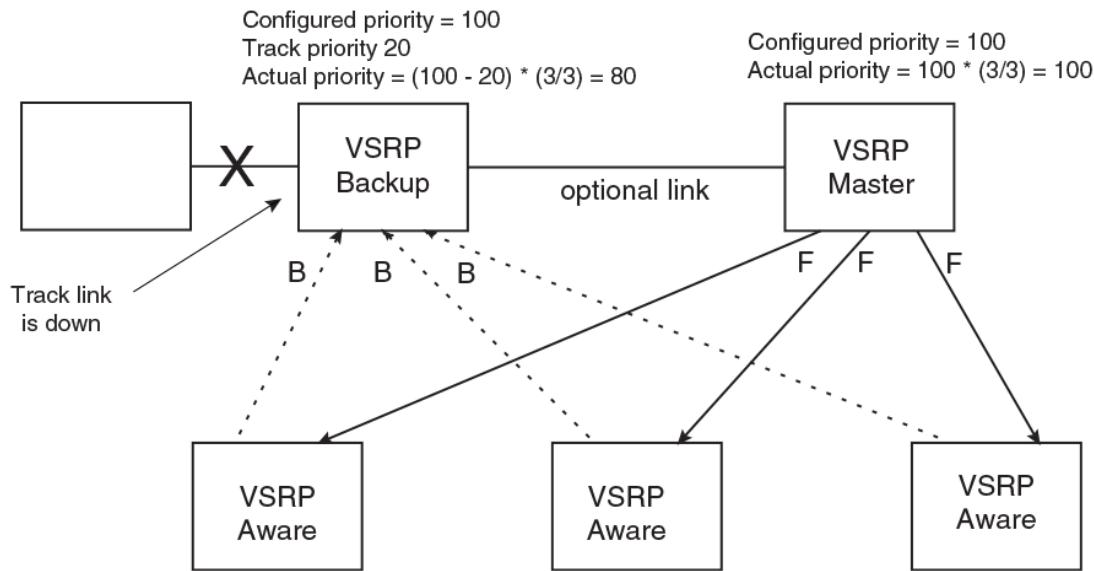
Optionally, you can configure track ports to be included during VSRP priority calculation. In VSRP, a track port is a port that is not a member of the VRID VLAN, but whose state is nonetheless considered when the priority is calculated. Typically, a track port represents the exit side of traffic received on the VRID ports. By default, no track ports are configured.

When you configure a track port, you assign a priority value to the port. If the port goes down, VSRP subtracts the track port priority value from the configured VSRP priority. For example, if you configure a track port with priority 20 and the configured VSRP priority is 100, the software subtracts 20 from 100 if the track port goes down, resulting in a VSRP priority of 80. The new priority value is used when calculating the VSRP priority. The following figure shows an example.

FIGURE 15 Track port priority

In Figure 15, the track port is up. Since the port is up, the track priority does not affect the VSRP priority calculation. If the track port goes down, the track priority does affect VSRP priority calculation, as shown in the following figure.

FIGURE 16 Track port priority subtracted during priority calculation



MAC address failover on VSRP-aware devices

VSRP-aware devices maintain a record of each VRID and its VLAN. When the device has received a Hello message for a VRID in a given VLAN, the device creates a record for that VRID and VLAN and includes the port number in the record. Each subsequent time the device receives a Hello message for the same VRID and VLAN, the device checks the port number:

- If the port number is the same as the port that previously received a Hello message, the VSRP-aware device assumes that the message came from the same VSRP Master that sent the previous message.
- If the port number does not match, the VSRP-aware device assumes that a VSRP failover has occurred to a new Master, and moves the MAC addresses learned on the previous port to the new port.

The VRID records age out if unused. This can occur if the VSRP-aware device becomes disconnected from the Master. The VSRP-aware device will wait for a Hello message for the period of time equal to the following.

$$\text{VRID Age} = \text{Dead Interval} + \text{Hold-down Interval} + (3 \times \text{Hello Interval})$$

The values for these timers are determined by the VSRP device sending the Hello messages. If the Master uses the default timer values, the age time for VRID records on the VSRP-aware devices is as follows.

$$3 + 3 + (3 \times 1) = 9 \text{ seconds}$$

In this case, if the VSRP-aware device does not receive a new Hello message for a VRID in a given VLAN, on any port, the device assumes the connection to the Master is unavailable and removes the VRID record.

VSRP interval timers

The VSRP Hello interval, Dead interval, Backup Hello interval, and Hold-down interval timers are individually configurable. You also can easily change all the timers at the same time while preserving the ratios among their values. To do so, change the timer scale. The *timer scale* is a value used by the software to calculate the timers. The software divides a timer value by the timer scale value. By default, the scale is 1. This means the VSRP timer values are the same as the values in the configuration.

VSRP-aware security features

This feature protects against unauthorized VSRP hello packets by enabling you to configure VSRP-aware security parameters. Without VSRP-aware security, a VSRP-aware device passively learns the authentication method conveyed by the received VSRP hello packet. The VSRP-aware device then stores the authentication method until it ages out with the aware entry.

The VSRP-aware security feature enables you to perform the following:

- Define the specific authentication parameters that a VSRP-aware device will use on a VSRP backup switch. The authentication parameters that you define will not age out.
- Define a list of ports that have authentic VSRP backup switch connections. For ports included in the list, the VSRP-aware switch will process VSRP hello packets using the VSRP-aware security configuration. Conversely, for ports not included in the list, the VSRP-aware switch will not use the VSRP-aware security configuration.

If VSRP hello packets do not meet the acceptance criteria, the VSRP-aware device forwards the packets normally, without any VSRP-aware security processing.

To configure VSRP-Aware Security features, refer to [Configuring security features on a VSRP-aware device](#) on page 108.

VSRP parameters

The following table lists the VSRP parameters.

TABLE 22 VSRP parameters

Parameter	Description	Default	For more information
Protocol	VSRP state	Enabled	Disabling or re-enabling VSRP on page 106
NOTE On a Layer 3 Switch, you must disable VSRP to use VRRP-E or VRRP.			
Virtual Router ID (VRID)	The ID of the virtual switch you are creating by configuring multiple devices as redundant links. You must configure the same VRID on each device that you want to use to back up the links.	None	Configuring basic VSRP parameters on page 106

TABLE 22 VSRP parameters (Continued)

Parameter	Description	Default	For more information
Timer scale	The value used by the software to calculate all VSRP timers. Increasing the timer scale value decreases the length of all the VSRP timers equally, without changing the ratio of one timer to another.	1	Changing the timer scale on page 107
Interface parameters			
Authentication type	<p>The type of authentication the VSRP devices use to validate VSRP packets. On Layer 3 Switches, the authentication type must match the authentication type the VRID port uses with other routing protocols such as OSPF.</p> <ul style="list-style-type: none"> No authentication - The interfaces do not use authentication. Simple - The interface uses a simple text-string as a password in packets sent on the interface. If the interface uses simple password authentication, the VRID configured on the interface must use the same authentication type and the same password. 	No authentication	Configuring authentication on page 108
NOTE MD5 is not supported.			
VSRP-Aware Security Parameters			
VSRP-Aware Authentication type	<p>The type of authentication the VSRP-aware devices will use on a VSRP backup switch:</p> <ul style="list-style-type: none"> No authentication - The device does not accept incoming packets that have authentication strings. Simple - The device uses a simple text-string as the authentication string for accepting incoming packets. 	Not configured	Configuring security features on a VSRP-aware device on page 108
VRID parameters			
VSRP device type	<p>Whether the device is a VSRP Backup for the VRID.</p> <p>All VSRP devices for a given VRID are Backups.</p>	Not configured	Configuring basic VSRP parameters on page 106
VSRP ports	The ports in the VRID VLAN that you want to use as VRID interfaces. You can selectively exclude individual ports from VSRP while allowing them to remain in the VLAN.	All ports in the VRID VLAN	Removing a port from the VRID VLAN on page 109

TABLE 22 VSRP parameters (Continued)

Parameter	Description	Default	For more information
VRID IP address	A gateway address you are backing up. Configuring an IP address provides VRRP-E Layer 3 redundancy in addition to VSRP Layer 2 redundancy. The VRID IP address must be in the same subnet as a real IP address configured on the VSRP interface, but cannot be the same as a real IP address configured on the interface.	None	Configuring a VRID IP address on page 109
NOTE This parameter is valid only on Layer 3 Switches.			
Backup priority	A numeric value that determines a Backup preferability for becoming the Master for the VRID. During negotiation, the device with the highest priority becomes the Master. In VSRP, all devices are Backups and have the same priority by default. If two or more Backups are tied with the highest priority, the Backup with the highest IP address becomes the Master for the VRID.	100 for all Backups	Changing the backup priority on page 110
Preference of timer source	When you save a Backup configuration, the software can save the configured VSRP timer values or the VSRP timer values received from the Master. Saving the current timer values instead of the configured ones helps ensure consistent timer usage for all the VRID devices.	Configured timer values are saved	Saving the timer values received from the master on page 110
NOTE The Backup always gets its timer scale value from the Master.			
Time-to-Live (TTL)	The maximum number of hops a VSRP Hello packet can traverse before being dropped. You can specify from 1 - 255.	2	Changing the TTL setting on page 110
Hello interval	The amount of time between Hello messages from the Master to the Backups for a given VRID. The interval can be from 1 - 84 seconds.	One second	Changing the hello interval setting on page 111
Dead interval	The amount of time a Backup waits for a Hello message from the Master for the VRID before determining that the Master is no longer active. If the Master does not send a Hello message before the dead interval expires, the Backups negotiate (compare priorities) to select a new Master for the VRID.	Three times the Hello Interval	Changing the dead interval setting on page 111

TABLE 22 VSRP parameters (Continued)

Parameter	Description	Default	For more information
Backup Hello state and interval	<p>The amount of time between Hello messages from a Backup to the Master.</p> <p>The message interval can be from 60 - 3600 seconds.</p> <p>You must enable the Backup to send the messages. The messages are disabled by default on Backups. The current Master sends Hello messages by default.</p>	Disabled 60 seconds when enabled	Changing the backup hello state and interval setting on page 112
Hold-down interval	<p>The amount of time a Backup that has sent a Hello packet announcing its intent to become Master waits before beginning to forward traffic for the VRID. The hold-down interval prevents Layer 2 loops from occurring during VSRP rapid failover.</p> <p>The interval can from 1 - 84 seconds.</p>	3 seconds	Changing the hold-down interval setting on page 112
Track priority	A VSRP priority value assigned to the tracked ports. If a tracked port link goes down, the VRID port VSRP priority is reduced by the amount of the tracked port priority.	5	Changing the default track priority setting on page 112
Track port	<p>A track port is a port or virtual routing interface that is outside the VRID but whose link state is tracked by the VRID. Typically, the tracked interface represents the other side of VRID traffic flow through the device.</p> <p>If the link for a tracked interface goes down, the VSRP priority of the VRID interface is changed, causing the devices to renegotiate for Master.</p>	None	Specifying a track port setting on page 113
Backup preempt mode	Prevents a Backup with a higher VSRP priority from taking control of the VRID from another Backup that has a lower priority but has already assumed control of the VRID.	Enabled	Disabling or re-enabling backup pre-emption setting on page 113
VRID active state	The active state of the VSRP VRID.	Disabled	Configuring basic VSRP parameters on page 106
RIP parameters			
Suppression of RIP advertisements	A Layer 3 Switch that is running RIP normally advertises routes to a backed up VRID even when the Layer 3 Switch is not currently the active Layer 3 Switch for the VRID. Suppression of these advertisements helps ensure that other Layer 3 Switches do not receive invalid route paths for the VRID.	Disabled (routes are advertised)	Suppressing RIP advertisement from backups on page 114
NOTE <hr/> <p>This parameter is valid only on Layer 3 Switches.</p> <hr/>			

Configuring basic VSRP parameters

To configure VSRP, perform the following required tasks:

- Configure a port-based VLAN containing the ports for which you want to provide VSRP service.

NOTE

If you already have a port-based VLAN but only want to use VSRP on a sub-set of the VLANs ports, you can selectively remove ports from VSRP service in the VLAN. Refer to [Removing a port from the VRID VLAN](#) on page 109.

- Configure a VRID:

- Specify that the device is a backup. Since VSRP, like VRRP-E, does not have an "owner", all VSRP devices are backups. The active device for a VRID is elected based on the VRID priority, which is configurable.
- Activate the VRID.

The following example shows a simple VSRP configuration.

```
device(config)#vlan 200
device(config-vlan-200)#tag ethernet 1/1 to 1/8
device(config-vlan-200)#vsrp vrid 1
device(config-vlan-200-vrid-1)#backup
device(config-vlan-200-vrid-1)#activate
```

Syntax: [no] vsrp vrid num

The num parameter specifies the VRID and can be from 1 - 255.

Syntax: [no] backup [priority value] [track-priority value]

This command is required. In VSRP, all devices on which a VRID are configured are Backups. The Master is then elected based on the VSRP priority of each device. There is no "owner" device as there is in VRRP.

For information about the command optional parameters, refer to the following:

- [Changing the backup priority](#) on page 110
- [Changing the default track priority setting](#) on page 112

Syntax: [no] activate

or

Syntax: enable | disable

Configuring optional VSRP parameters

The following sections describe how to configure optional VSRP parameters.

Disabling or re-enabling VSRP

VSRP is enabled by default on Layer 2 Switches and Layer 3 Switches. On a Layer 3 Switch, if you want to use VRRP or VRRP-E for Layer 3 redundancy instead of VSRP, you need to disable VSRP first. To do so, enter the following command at the global CONFIG level.

```
device(config)#no router vsrp
router vsrp is disabled. All vsrp config data will be lost when writing to flash
```

To re-enable the protocol, enter the following command.

```
device(config)#router vsrcp
```

Syntax: [no] router vsrcp

Since VRRP and VRRP-E do not apply to Layer 2 Switches, there is no need to disable VSRP and there is no command to do so. The protocol is always enabled.

Changing the timer scale

To achieve sub-second failover times, you can shorten the duration of all scale timers for VSRP, VRRP, and VRRP-E by adjusting the timer scale. The timer scale is a value used by the software to calculate the timers. By default, the scale value is 1. If you increase the timer scale, each timer value is divided by the scale value. Using the timer scale to adjust timer values enables you to easily change all the timers while preserving the ratios among their values. Here is an example.

Timer	Timer scale	Timer value
Hello interval	1	1 second
	2	0.5 seconds
Dead interval	1	3 seconds
	2	1.5 seconds
Backup Hello interval	1	60 seconds
	2	30 seconds
Hold-down interval	1	3 seconds
	2	1.5 second

NOTE

The Backups always use the value of the timer scale received from the Master, and the value from the Master will be written in the configuration file.

To change the timer scale, enter a command such as the following at the global CONFIG level of the CLI.

```
device(config)# scale-timer 2
```

This command changes the scale to 2. All VSRP, VRRP, and VRRP-E timer values will be divided by 2.

Syntax: [no] scale-timer num

The num parameter specifies the multiplier. You can specify a timer scale from 1 - 10.

Configuring authentication

If the interfaces on which you configure the VRID use authentication, the VSRP packets on those interfaces also must use the same authentication. VSRP supports the following authentication types:

- No authentication - The interfaces do not use authentication.
- Simple - The interfaces use a simple text-string as a password in packets sent on the interface. If the interfaces use simple password authentication, the VRID configured on the interfaces must use the same authentication type and the same password.

To configure a simple password, enter a command such as the following at the VLAN configuration level.

```
device(config-vlan-10)#vsrp auth-type simple-text-auth ourpwd
```

This command configures the simple text password "ourpwd".

Syntax: [no] vsrp auth-type { no-auth | simple-text-auth auth-data }

The **auth-type no-auth** parameter indicates that the VRID and the interface it is configured on do not use authentication.

The **auth-type simple-text-auth auth-data** parameter indicates that the VRID and the interface it is configured on use a simple text password for authentication. The auth-data value is the password, and can be up to eight characters. If you use this parameter, make sure all interfaces on all the devices supporting this VRID are configured for simple password authentication and use the same password.

Configuring security features on a VSRP-aware device

This section shows how to configure security features on a VSRP-aware device. For an overview of this feature, refer to [VSRP-aware security features](#) on page 102.

Specifying an authentication string for VSRP hello packets

The following configuration defines **pri-key** as the authentication string for accepting incoming VSRP hello packets. In this example, the VSRP-aware device will accept all incoming packets that have this authorization string.

```
device(config)#vlan 10
device(config-vlan-10)#vsrp-aware vrid 3 simple-text-auth pri-key
```

Syntax: vsrp-aware vrid vridnumber simple-text-auth string

Specifying no authentication for VSRP hello packets

The following configuration specifies no authentication as the preferred VSRP-aware security method. In this case, the VSRP device will not accept incoming packets that have authentication strings.

```
device(config)#vlan 10
device(config-vlan-10)#vsrp-aware vrid 2 no-auth
```

Syntax: vsrp-aware vrid vridnumber no-auth

The following configuration specifies no authentication for VSRP hello packets received on ports 1/1, 1/2, 1/3, and 1/4 in VRID 4. For these ports, the VSRP device will not accept incoming packets that have authentication strings.

```
device(config)#vlan 10
device(config-vlan-10)#vsrp-aware vrid 4 no-auth port-list ethe 1/1 to 1/4
```

Syntax: [no] vsrp-aware vrid vridnumber no-auth port-list portrange

vrid number is a valid VRID (from 1 to 255).

no-auth specifies no authentication as the preferred VSRP-aware security method. The VSRP device will not accept incoming packets that have authentication strings.

simple-text-auth string specifies the authentication string for accepting VSRP hello packets, where string can be up to 8 characters.

port-list port range specifies the range of ports to include in the configuration.

Removing a port from the VRID VLAN

By default, all the ports on which you configure a VRID are interfaces for the VRID. You can remove a port from the VRID while allowing it to remain in the VLAN.

Removing a port is useful in the following cases:

- There is no risk of a loop occurring, such as when the port is attached directly to an end host.
- You plan to use a port in an MRP ring.

To remove a port from a VRID, enter a command such as the following at the configuration level for the VRID.

```
device(config-vlan-200-vrid-1)#no include-port ethernet 1/2
```

Syntax: [no] include-port ethernet [slotnum]|portnum

The slotnum parameter is required on chassis devices.

The portnum parameter specifies the port you are removing from the VRID. The port remains in the VLAN but its forwarding state is not controlled by VSRP. If you are configuring a chassis device, specify the slot number as well as the port number (slotnum / portnum).

Configuring a VRID IP address

If you are configuring a Layer 3 Switch for VSRP, you can specify an IP address to back up. When you specify an IP address, VSRP provides redundancy for the address. This is useful if you want to back up the gateway address used by hosts attached to the VSRP Backups.

VSRP does not require you to specify an IP address. If you do not specify an address, VSRP provides Layer 2 redundancy. If you do specify an address, VSRP provides Layer 2 and Layer 3 redundancy.

The Layer 3 redundancy support is the same as VRRP-E support. For information, refer to the *FastIron Ethernet Switch Layer 3 Routing Configuration Guide*.

NOTE

The VRID IP address must be in the same subnet as a real IP address configured on the VSRP interface, but cannot be the same as a real IP address configured on the interface.

NOTE

Failover applies to both Layer 2 and Layer 3.

To specify an IP address to back up, enter a command such as the following at the configuration level for the VRID.

```
device(config-vlan-200-vrid-1)#ip-address 10.10.10.1
```

Syntax: [no] ip-address *ip-addr*

Changing the backup priority

When you enter the backup command to configure the device as a VSRP Backup for the VRID, you also can change the backup priority and the track priority:

- The backup priority is used for election of the Master. The VSRP Backup with the highest priority value for the VRID is elected as the Master for that VRID. The default priority is 100. If two or more Backups are tied with the highest priority, the Backup with the highest IP address becomes the Master for the VRID.
- The track priority is used with the track port feature. Refer to [VSRP priority calculation](#) on page 98 and [Changing the default track priority setting](#) on page 112.

To change the backup priority, enter a command such as the following at the configuration level for the VRID.

```
device(config-vlan-200-vrid-1)#backup priority 75
```

Syntax: [no] **backup** [**priority** *value*] [**track-priority** *value*]

The **priority** value parameter specifies the backup priority for this interface and VRID. Specify a value as follows:

- For VRRP, specify a value from 3 - 254. The default is 100.
- For VSRP and VRRP-E, specify a value from 6 - 255. The default is 100.

For a description of the **track-priority** value parameter, refer to [Changing the default track priority setting](#) on page 112.

Saving the timer values received from the master

The Hello messages sent by a VRID master contain the VRID values for the following VSRP timers:

- Hello interval
- Dead interval
- Backup Hello interval
- Hold-down interval

The Backups always use the value of the timers received from the Master.

To configure a Backup to save the VSRP timer values received from the Master instead of the timer values configured on the Backup, enter the following command. Saving the current timer values instead of the configured ones helps ensure consistent timer usage for all the VRID devices.

```
device(config-vlan-200-vrid-1)#save-current-values
```

Syntax: [no] **save-current-values**

Changing the TTL setting

A VSRP Hello packet time to live (TTL) specifies how many hops the packet can traverse before being dropped. A hop can be a Layer 3 Switch or a Layer 2 Switch. You can specify from 1 - 255. The default TTL is 2. When a VSRP device (Master or Backup) sends a VSRP HELLO packet, the device subtracts one from the TTL. Thus, if the TTL is 2, the device that originates the Hello packet sends it out with a TTL of 1. Each subsequent device that receives the packet also subtracts one from the packet TTL. When the packet has a TTL of 1, the receiving device subtracts 1 and then drops the packet because the TTL is zero.

NOTE

An MRP ring is considered to be a single hop, regardless of the number of nodes in the ring.

To change the TTL for a VRID, enter a command such as the following at the configuration level for the VRID.

```
device(config-vlan-200-vrid-1)#initial-ttl 5
```

Syntax: [no] initial-ttl num

The num parameter specifies the TTL and can be from 1 - 255. The default TTL is 2.

Changing the hello interval setting

The Master periodically sends Hello messages to the Backups. To change the Hello interval, enter a command such as the following at the configuration level for the VRID.

```
device(config-vlan-200-vrid-1)#hello-interval 10
```

Syntax: [no] hello-interval num

The num parameter specifies the interval and can be from 1 - 84 seconds. The default is 1 second.

NOTE

The default Dead interval is three times the Hello interval plus one-half second. Generally, if you change the Hello interval, you also should change the Dead interval on the Backups.

NOTE

If you change the timer scale, the change affects the actual number of seconds.

Changing the dead interval setting

The Dead interval is the number of seconds a Backup waits for a Hello message from the Master before determining that the Master is dead. The default is 3 seconds. This is three times the default Hello interval.

To change the Dead interval, enter a command such as the following at the configuration level for the VRID.

```
device(config-vlan-200-vrid-1)#dead-interval 30
```

Syntax: [no] dead-interval num

The num parameter specifies the interval and can be from 1 - 84 seconds. The default is 3 seconds.

NOTE

If you change the timer scale, the change affects the actual number of seconds.

Changing the backup hello state and interval setting

By default, Backups do not send Hello messages to advertise themselves to the Master. You can enable these messages if desired and also change the message interval.

To enable a Backup to send Hello messages to the Master, enter a command such as the following at the configuration level for the VRID.

```
device(config-vlan-200-vrid-1)#advertise backup
```

Syntax: [no] advertise backup

When a Backup is enabled to send Hello messages, the Backup sends a Hello message to the Master every 60 seconds by default. You can change the interval to be up to 3600 seconds.

To change the Backup Hello interval, enter a command such as the following at the configuration level for the VRID.

```
device(config-vlan-200-vrid-1)#backup-hello-interval 180
```

Syntax: [no] backup-hello-interval num

The num parameter specifies the message interval and can be from 60 - 3600 seconds. The default is 60 seconds.

NOTE

If you change the timer scale, the change affects the actual number of seconds.

Changing the hold-down interval setting

The hold-down interval prevents Layer 2 loops from occurring during failover, by delaying the new Master from forwarding traffic long enough to ensure that the failed Master is really unavailable.

To change the Hold-down interval, enter a command such as the following at the configuration level for the VRID.

```
device(config-vlan-200-vrid-1)#hold-down-interval 4
```

Syntax: [no] hold-down-intervalnum

The num parameter specifies the hold-down interval and can be from 1 - 84 seconds. The default is 3 seconds.

NOTE

If you change the timer scale, the change affects the actual number of seconds.

Changing the default track priority setting

When you configure a VRID to track the link state of other interfaces, if one of the tracked interface goes down, the software changes the VSRP priority of the VRID interface.

The software reduces the VRID priority by the amount of the priority of the tracked interface that went down. For example, if the VSRP interface priority is 100 and a tracked interface with track priority 60 goes down, the software changes the VSRP interface priority to 40. If another tracked interface goes down, the software reduces the VRID priority again, by the amount of the tracked interface track priority.

The default track priority for all track ports is 5. You can change the default track priority or override the default for an individual track port.

- To change the default track priority, use the **backup priority** value **track-priority** value command, described below.
- To override the default track priority for a specific track port, use the **track-port** command. Refer to [Specifying a track port setting](#) on page 113.

To change the track priority, enter a command such as the following at the configuration level for the VRID.

```
device(config-vlan-200-vrid-1)#backup priority 100 track-priority 2
```

Syntax: [no] **backup** [**priority value**] [**track-priority value**]

Specifying a track port setting

You can configure the VRID on one interface to track the link state of another interface on the device. This capability is useful for tracking the state of the exit interface for the path for which the VRID is providing redundancy. Refer to [VSRP priority calculation](#) on page 98.

To configure a VRID to track an interface, enter a command such as the following at the configuration level for the VRID.

```
device(config-vlan-200-vrid-1)#track-port ethernet 2/4
```

Syntax: [no] **track-port ethernet** [**slotnum**] **portnum** | **ve num** [**priority num**]

The **priority num** parameter changes the VSRP priority of the interface. If this interface goes down, the VRID VSRP priority is reduced by the amount of the track port priority you specify here.

NOTE

The **priority num** option changes the priority of the specified interface, overriding the default track port priority. To change the default track port priority, use the **backup track-priority num** command.

Disabling or re-enabling backup pre-emption setting

By default, a Backup that has a higher priority than another Backup that has become the Master can preempt the Master, and take over the role of Master. If you want to prevent this behavior, disable preemption.

Preemption applies only to Backups and takes effect only when the Master has failed and a Backup has assumed ownership of the VRID. The feature prevents a Backup with a higher priority from taking over as Master from another Backup that has a lower priority but has already become the Master of the VRID.

Preemption is especially useful for preventing flapping in situations where there are multiple Backups and a Backup with a lower priority than another Backup has assumed ownership, because the Backup with the higher priority was unavailable when ownership changed.

If you enable the non-preempt mode (thus disabling the preemption feature) on all the Backups, the Backup that becomes the Master following the disappearance of the Master continues to be the Master. The new Master is not preempted.

To disable preemption on a Backup, enter a command such as the following at the configuration level for the VRID.

```
device(config-vlan-200-vrid-1)#non-preempt-mode
```

Syntax: [no] non-preempt-mode

Suppressing RIP advertisement from backups

Normally, for Layer 3 a VSRP Backup includes route information for a backed up IP address in RIP advertisements. As a result, other Layer 3 Switches receive multiple paths for the backed up interface and might sometimes unsuccessfully use the path to the Backup rather than the path to the Master.

You can prevent the Backups from advertising route information for the backed up interface by enabling suppression of the advertisements.

NOTE

This parameter applies only if you specified an IP address to back up and is valid only on Layer 3 Switches.

To suppress RIP advertisements, enter the following commands.

```
device(config)#router rip  
device(config-rip-router)#use-vrrp-path
```

Syntax: [no] use-vrrp-path

VSRP-aware interoperability

The **vsrp-aware tc-vlan-flush** command should be used in network configurations in which the Brocade switch operates as the VSRP-Aware device connecting to a FastIron platforms configured as a VSRP Master.

The command is available at the VLAN level, and is issued per a specific VRID, as shown here for VRID 11.

```
device(config-vlan-10)#vsrp-aware vrid 11 tc-vlan-flush
```

Syntax: vsrp-aware vrid num tc-vlan-flush

When this command is enabled, MAC addresses will be flushed at the VLAN level, instead of at the port level. MAC addresses will be flushed for every topology change (TC) received on the VSRP-aware ports.

When you configure the **vsrp-aware tc-vlan-flush** command on a VSRP-aware device, and the device receives VSRP hello packets from the VSRP master, VSRP authentication is automatically configured. However, if the VSRP-aware device does not receive VSRP hello packets from the VSRP master when the **vsrp-aware tc-vlan-flush** command is configured, you must manually configure VSRP authentication. For more information on configuring VSRP authentication, refer to [Configuring authentication](#) on page 108.

When this command is enabled, the results of the **show vsrp-aware vlan** command resemble the following.

```
device(config-vlan-10)#vsrp-aware vrid 11 tc-vlan-flush  
device(config-vlan-10)#show vsrp aware vlan 10  
Aware Port Listing  
VLAN ID VRID Last Port Auth Type Mac-Flush Age  
10 11 N/A no-auth Configured Enabled 00:00:00.0
```

Displaying VSRP information

You can display the following VSRP information:

- Configuration information and current parameter values for a VRID or VLAN
- The interfaces on a VSRP-aware device that are active for the VRID

Displaying VRID information

To display VSRP information, enter the following command.

```
device#show vsrp vrid 1
Total number of VSRP routers defined: 2
VLAN 200
auth-type no authentication
VRID 1
  State      Administrative-status Advertise-backup Preempt-mode save-current
  standby     enabled             disabled        true       false

  Parameter      Configured Current   Unit
  priority       100          80        (100-0)*(4.0/5.0)
  hello-interval 1            1         sec/1
  dead-interval 3            3         sec/1
  hold-interval 3            3         sec/1
  initial-ttl   2            2         hops
  next hello sent in 00:00:00.8
  Member ports: ethe 1/1 to 1/5
  Operational ports: ethe 1/1 to 1/4
  Forwarding ports: ethe 1/1 to 1/4
```

Syntax: **show vsrp [vrid num | vlan vlan-id]**

This display shows the following information when you use the **vrid num** or **vlan vlan-id** parameter. For information about the display when you use the **aware** parameter, refer to [Displaying the active interfaces for a VRID](#) on page 118.

TABLE 23 CLI display of VSRP VRID or VLAN information

Field	Description
Total number of VSRP routers defined	The total number of VRIDs configured on this device.
VLAN	The VLAN on which VSRP is configured.
auth-type	The authentication type in effect on the ports in the VSRP VLAN.
VRID parameters	
VRID	The VRID for which the following information is displayed.

TABLE 23 CLI display of VSRP VRID or VLAN information (Continued)

Field	Description
state	<p>This device VSRP state for the VRID. The state can be one of the following:</p> <ul style="list-style-type: none"> initialize - The VRID is not enabled (activated). If the state remains "initialize" after you activate the VRID, make sure that the VRID is also configured on the other routers and that the routers can communicate with each other. <hr/> <p>NOTE</p> <p>If the state is "initialize" and the mode is incomplete, make sure you have specified the IP address for the VRID.</p>
	<ul style="list-style-type: none"> standby - This device is a Backup for the VRID. master - This device is the Master for the VRID.
Administrative-status	<p>The administrative status of the VRID. The administrative status can be one of the following:</p> <ul style="list-style-type: none"> disabled - The VRID is configured on the interface but VSRP or VRRP-E has not been activated on the interface. enabled - VSRP has been activated on the interface.
Advertise-backup	<p>Whether the device is enabled to send VSRP Hello messages when it is a Backup. This field can have one of the following values:</p> <ul style="list-style-type: none"> disabled - The device does not send Hello messages when it is a Backup. enabled - The device does send Hello messages when it is a Backup.
Preempt-mode	<p>Whether the device can be pre-empted by a device with a higher VSRP priority after this device becomes the Master. This field can have one of the following values:</p> <ul style="list-style-type: none"> disabled - The device cannot be pre-empted. enabled - The device can be pre-empted.
save-current	<p>The source of VSRP timer values preferred when you save the configuration. This field can have one of the following values:</p> <ul style="list-style-type: none"> false - The timer values configured on this device are saved. true - The timer values most recently received from the Master are saved instead of the locally configured values. <hr/> <p>NOTE</p> <p>For the following fields:</p>
	<ul style="list-style-type: none"> Configured - indicates the parameter value configured on this device. Current - indicates the parameter value received from the Master. Unit - indicates the formula used for calculating the VSRP priority and the timer scales in effect for the VSRP timers. A timer true value is the value listed in the Configured or Current field divided by the scale value.

TABLE 23 CLI display of VSRP VRID or VLAN information (Continued)

Field	Description
priority	The device preferability for becoming the Master for the VRID. During negotiation, the Backup with the highest priority becomes the Master. If two or more Backups are tied with the highest priority, the Backup interface with the highest IP address becomes the Master for the VRID.
hello-interval	The number of seconds between Hello messages from the Master to the Backups for a given VRID.
dead-interval	The configured value for the dead interval. The dead interval is the number of seconds a Backup waits for a Hello message from the Master for the VRID before determining that the Master is no longer active. If the Master does not send a Hello message before the dead interval expires, the Backups negotiate (compare priorities) to select a new Master for the VRID.
NOTE If the value is 0, then you have not configured this parameter.	
hold-interval	The number of seconds a Backup that intends to become the Master will wait before actually beginning to forward Layer 2 traffic for the VRID. If the Backup receives a Hello message with a higher priority than its own before the hold-down interval expires, the Backup remains in the Backup state and does not become the new Master.
initial-ttl	The number of hops a Hello message can traverse after leaving the device before the Hello message is dropped.
NOTE An MRP ring counts as one hop, regardless of the number of nodes in the ring.	
next hello sent in	The amount of time until the Master dead interval expires. If the Backup does not receive a Hello message from the Master by the time the interval expires, either the IP address listed for the Master will change to the IP address of the new Master, or this Layer 3 Switch itself will become the Master.
NOTE This field applies only when this device is a Backup.	
Member ports	The ports in the VRID.
Operational ports	The member ports that are currently up.
Forwarding ports	The member ports that are currently in the Forwarding state. Ports that are forwarding on the Master are listed. Ports on the Standby, which are in the Blocking state, are not listed.

Displaying the active interfaces for a VRID

On a VSRP-aware device, you can display VLAN and port information for the connections to the VSRP devices (Master and Backups).

To display the active VRID interfaces, enter the following command on the VSRP-aware device.

```
device#show vsrp aware
Aware port listing
VLAN ID VRID Last Port
100      1     3/2
200      2     4/1
```

Syntax: show vsrp aware

This display shows the following information when you use the **aware** parameter. For information about the display when you use the **vrid num** or **vlan vlan-id** parameter, refer to [Displaying VRID information](#) on page 115.

TABLE 24 CLI display of VSRP-aware information

Field	Description
VLAN ID	The VLAN that contains the VSRP-aware device connection with the VSRP Master and Backups.
VRID	The VRID.
Last Port	The most recent active port connection to the VRID. This is the port connected to the current Master. If a failover occurs, the VSRP-aware device changes the port to the port connected to the new Master. The VSRP-aware device uses this port to send and receive data through the backed up node.

VSRP fast start

VSRP fast start allows non-Brocade or non-VSRP aware devices that are connected to a Brocade device that is the VSRP Master to quickly switchover to the new Master when a VSRP failover occurs.

This feature causes the port on a VSRP Master to restart when a VSRP failover occurs. When the port shuts down at the start of the restart, ports on the non-VSRP aware devices that are connected to the VSRP Master flush the MAC address they have learned for the VSRP master. After a specified time, the port on the previous VSRP Master (which now becomes the Backup) returns back online. Ports on the non-VSRP aware devices switch over to the new Master and learn its MAC address.

Configuring VSRP fast start

The VSRP fast start feature can be enabled on a VSRP-configured Brocade device, either on the VLAN to which the VRID of the VSRP-configured device belongs (globally) or on a port that belongs to the VRID.

To globally configure a VSRP-configured device to shut down its ports when a failover occurs, then restart after five seconds, enter the following command.

```
device(config)#vlan 100
device(config-vlan-100)#vsrp vrid 1
device(config-vlan-100-vrid-1)#restart-ports 5
```

Syntax: [no] restart-ports seconds

This command shuts down all the ports that belong to the VLAN when a failover occurs. All the ports will have the specified VRID.

To configure a single port on a VSRP-configured device to shut down when a failover occurs, then restart after a period of time, enter the following command.

```
device(config)#interface ethernet 1/1
device(config-if-e1000-1/1)#restart-vsdp-port 5
```

Syntax: [no] **restart-vsdp-port** seconds

In both commands, the seconds parameter instructs the VSRP Master to shut down its port for the specified number of seconds before it starts back up. Enter a value between 1 - 120 seconds. The default is 1 second.

Displaying ports that Have the VSRP fast start feature enabled

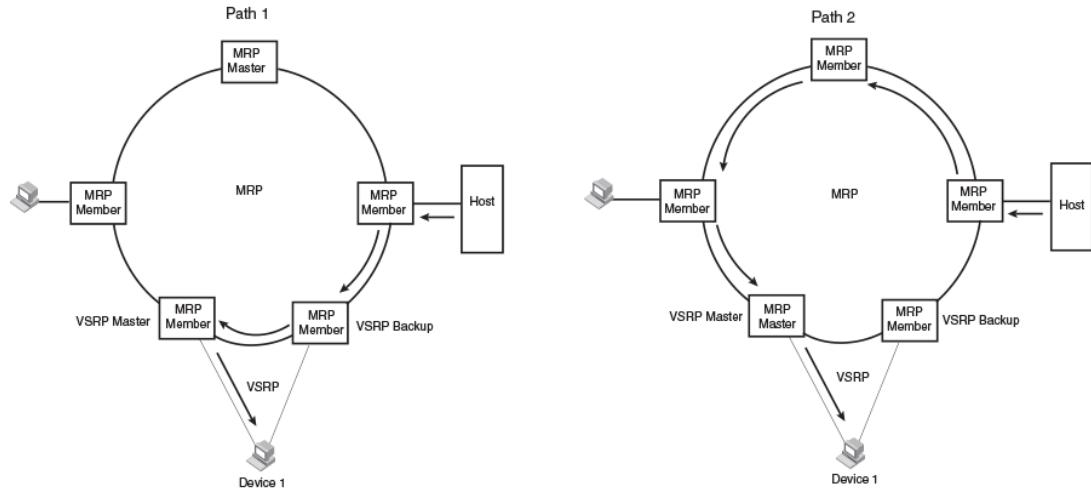
The **show vsrp vrid** command shows the ports on which the VSRP fast start feature is enabled.

```
device#show vsrp vrid 100
VLAN 100
  auth-type no authentication
  VRID 100
  =====
    State      Administrative-status Advertise-backup Preempt-mode save-current
    master     enabled           disabled        true       false
    Parameter   Configured   Current   Unit/Formula
    priority    100          50        (100-0)*(2.0/4.0)
    hello-interval 1            1        sec/1
    dead-interval 3            3        sec/1
    hold-interval 3            3        sec/1
    initial-ttl  2            2        hops
    next hello sent in 00:00:00.3
    Member ports: ethe 2/5 to 2/8
    Operational ports: ethe 2/5 ethe 2/8
    Forwarding ports: ethe 2/5 ethe 2/8
    Restart ports:   2/5(1) 2/6(1) 2/7(1) 2/8(1)
```

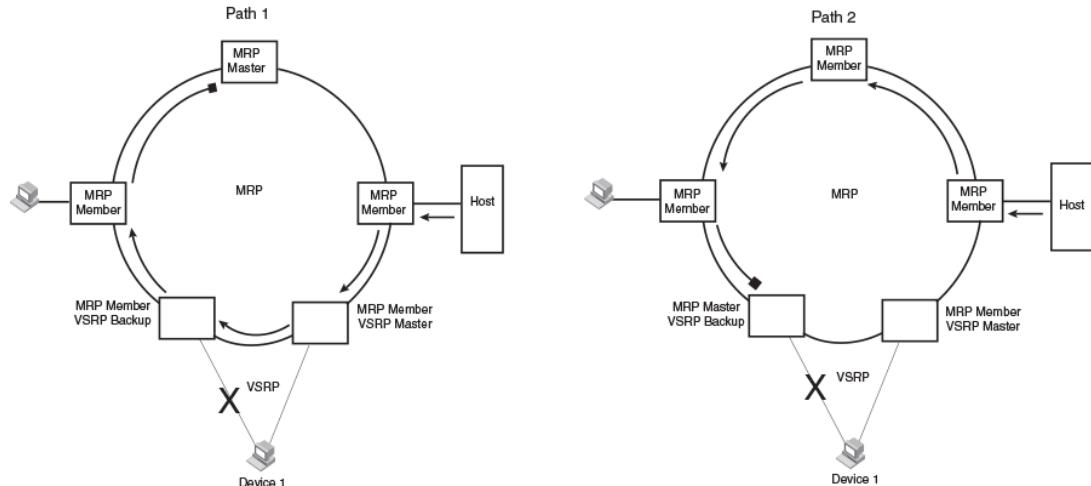
The "Restart ports:" line lists the ports that have the VSRP fast start enabled, and the downtime for each port. Refer to [Table 23](#) on page 115 to interpret the remaining information on the display.

VSRP and MRP signaling

A device may connect to an MRP ring through VSRP to provide a redundant path between the device and the MRP ring. VSRP and MRP signaling ensures rapid failover by flushing MAC addresses appropriately. The host on the MRP ring learns the MAC addresses of all devices on the MRP ring and VSRP link. From these MAC addresses, the host creates a MAC database (table), which is used to establish a data path from the host to a VSRP-linked device. The following figure below shows two possible data paths from the host to Device 1.

FIGURE 17 Two data paths from host on an MRP ring to a VSRP-linked device

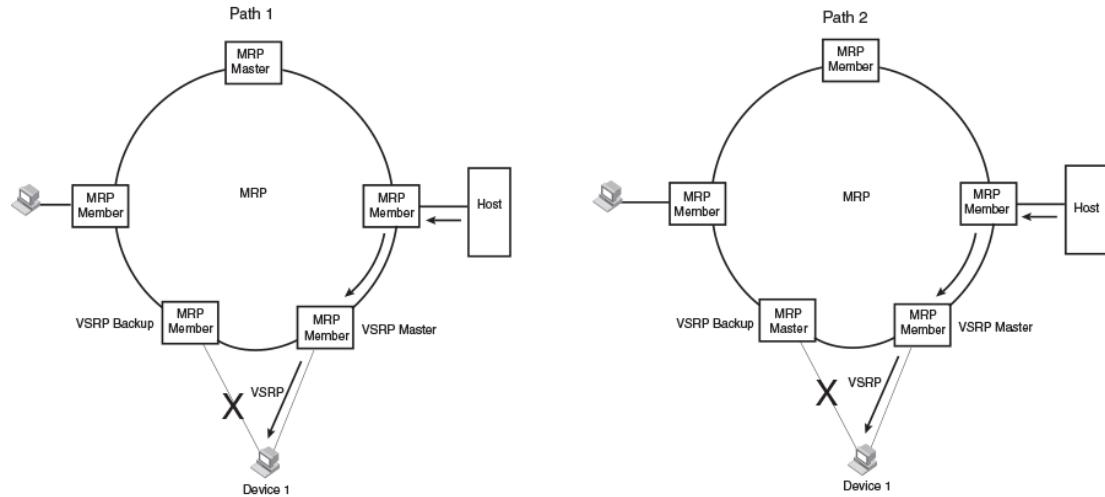
If a VSRP failover from master to backup occurs, VSRP needs to inform MRP of the topology change; otherwise, data from the host continues along the obsolete learned path and never reach the VSRP-linked device, as shown in the following figure.

FIGURE 18 VSRP on MRP rings that failed over

A signaling process for the interaction between VSRP and MRP ensures that MRP is informed of the topology change and achieves convergence rapidly. When a VSRP node fails, a new VSRP master is selected. The new VSRP master finds all MRP instances impacted by the failover. Then each MRP instance does the following:

- The MRP node sends out an MRP PDU with the mac-flush flag set three times on the MRP ring.
- The MRP node that receives this MRP PDU empties all the MAC entries from its interfaces that participate on the MRP ring.
- The MRP node then forwards the MRP PDU with the mac-flush flag set to the next MRP node that is in forwarding state.

The process continues until the Master MRP node secondary (blocking) interface blocks the packet. Once the MAC address entries have been flushed, the MAC table can be rebuilt for the new path from the host to the VSRP-linked device as shown in the following figure.

FIGURE 19 New path established

There are no CLI commands used to configure this process.

UDLD and Protected Link Groups

- UDLD overview..... 123
- Protected link groups..... 128

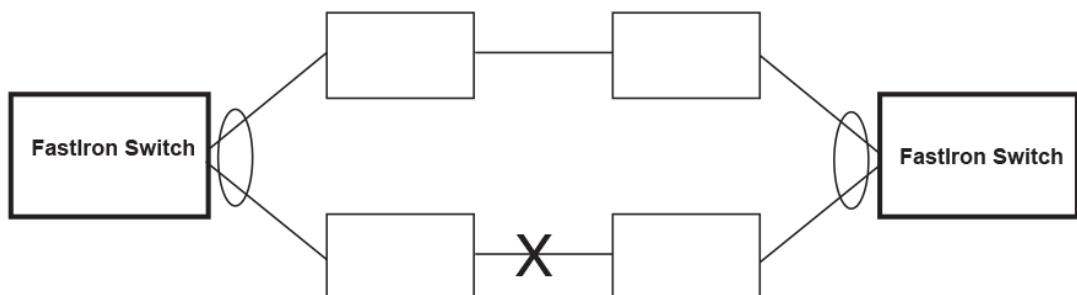
UDLD overview

Uni-Directional Link Detection (UDLD) monitors a link between two Brocade devices and brings the ports on both ends of the link down if the link goes down at any point between the two devices. This feature is useful for links that are individual ports and for trunk links. The following figure shows an example.

FIGURE 20 UDLD example

Without link keepalive, the FastIron ports remain enabled. Traffic continues to be load balanced to the ports connected to the failed link.

When link keepalive is enabled, the feature brings down the FastIron ports connected to the failed link.



Normally, a Brocade device load balances traffic across the ports in a trunk group. In this example, each Brocade device load balances traffic across two ports. Without the UDLD feature, a link failure on a link that is not directly attached to one of the Brocade devices is undetected by the Brocade devices. As a result, the Brocade devices continue to send traffic on the ports connected to the failed link.

When UDLD is enabled on the trunk ports on each Brocade device, the devices detect the failed link, disable the ports connected to the failed link, and use the remaining ports in the trunk group to forward the traffic.

Ports enabled for UDLD exchange proprietary health-check packets once every second (the keepalive interval). If a port does not receive a health-check packet from the port at the other end of the link within the keepalive interval, the port waits for two more intervals. If the port still does not receive a health-check packet after waiting for three intervals, the port concludes that the link has failed and takes the port down.

UDLD for tagged ports

The default implementation of UDLD sends the packets untagged, even across tagged ports. If the untagged UDLD packet is received by a third-party switch, that switch may reject the packet. As a

result, UDLD may be limited only to Brocade devices, since UDLD may not function on third-party switches.

To solve this issue, you can configure ports to send out UDLD control packets that are tagged with a specific VLAN ID. This feature also enables third party switches to receive the control packets that are tagged with the specified VLAN. For tagged operation, all of the following conditions must be met:

- A VLAN is specified when UDLD is configured
- The port belongs to the configured VLAN as tagged member
- All the devices across the UDLD link are in the same VLAN

For configuration details, refer to [Enabling UDLD for tagged ports](#) on page 125.

Configuration notes and feature limitations for UDLD

- UDLD is supported only on Ethernet ports.
- UDLD can be enabled on only one VLAN for tagged port.
- To configure UDLD on a trunk group, you must enable and configure the feature on each port of the group individually. Configuring UDLD on a trunk group primary port enables the feature on that port only.
- Low UDLD **link-keepalive** interval and retry options are not recommended as they are more sensitive and prone to flaps.
- When UDLD is enabled on a trunk port, trunk threshold is not supported.
- Dynamic trunking is not supported. If you want to configure a trunk group that contains ports on which UDLD is enabled, you must remove the UDLD configuration from the ports. After you create the trunk group, you can re-add the UDLD configuration.
- If MRP is also enabled on the device, Brocade recommends that you set the MRP preforwarding time slightly higher than the default of 300 ms; for example, to 400 or 500 ms. Refer to [Changing the hello and preforwarding times](#) on page 89.

Enabling UDLD

NOTE

This section shows how to configure UDLD for untagged control packets. To configure UDLD for tagged control packets, refer to [Enabling UDLD for tagged ports](#) on page 125.

To enable UDLD on a port, enter a command such as the following at the global CONFIG level of the CLI.

```
device(config)#link-keepalive ethernet 0/1/1
```

To enable the feature on a trunk group, enter commands such as the following.

```
device(config)#link-keepalive ethernet 0/1/1 ethernet 0/1/2
device(config)#link-keepalive ethernet 0/1/3 ethernet 0/1/4
```

Syntax: [no] link-keepalive ethernet port [to port | ethernet port]

This command is not supported if you downgrade the device to FCX 6.0 or FSX5.1. In this case, use the following command to configure multiple ports:

Syntax: [no] link-keepalive ethernet port [ethernet port]

Enabling UDLD for tagged ports

To enable ports to receive and send UDLD control packets tagged with a specific VLAN ID, enter commands such as the following.

```
device(config)#link-keepalive ethernet 1/18 vlan 22
```

This command enables UDLD on port 1/18 and allows UDLD control packet tagged with VLAN 22 to be received and sent on port 1/18.

Syntax: [no] link-keepalive ethernet port [vlan *vlan-ID*]

For the *vlan-ID* variable, enter the ID of the VLAN that the UDLD control packets can contain to be received and sent on the port. If a VLAN ID is not specified, then UDLD control packets are sent out of the port as untagged packets.

NOTE

You must configure the same VLANs that will be used for UDLD on all devices across the network; otherwise, the UDLD link cannot be maintained.

Changing the Keepalive interval

By default, ports enabled for UDLD send a link health-check packet once every 500 ms. You can change the interval to a value from 1 - 60, where 1 is 100 ms, 2 is 200 ms, and so on. To change the interval, enter a command such as the following.

```
device(config)#link-keepalive interval 4
```

Syntax: [no] link-keepalive interval *num*

The *num* parameter specifies how often the ports send a UDLD packet. You can specify from 1 - 60, in 100 ms increments. The default is 5 (500 ms).

Changing the Keepalive retries

By default, a port waits one second to receive a health-check reply packet from the port at the other end of the link. If the port does not receive a reply, the port tries four more times by sending up to four more health-check packets. If the port still does not receive a reply after the maximum number of retries, the port goes down.

You can change the maximum number of keepalive attempts to a value from 3 - 64. To change the maximum number of attempts, enter a command such as the following.

```
device(config)#link-keepalive retries 10
```

Syntax: [no] link-keepalive retries *num*

The *num* parameter specifies the maximum number of times the port will try the health check. You can specify a value from 3 - 64. The default is 7.

Displaying UDLD information

This section describes the commands used to display information about a UDLD configuration.

Displaying information for all ports

To display UDLD information for all ports, enter the following command.

```
device#show link-keepalive
Total link-keepalive enabled ports: 4
Keepalive Retries: 3      Keepalive Interval: 1 Sec.
Port    Physical Link   Logical Link   State           Link-vlan
4/1     up              up             FORWARDING    3
4/2     up              up             FORWARDING
4/3     down            down           DISABLED
4/4     up              down           DISABLED
```

Syntax: **show link-keepalive**

TABLE 25 CLI display of UDLD information

Field	Description
Total link-keepalive enabled ports	The total number of ports on which UDLD is enabled.
Keepalive Retries	The number of times a port will attempt the health check before concluding that the link is down.
Keepalive Interval	The number of seconds between health check packets.
Port	The port number.
Physical Link	The state of the physical link. This is the link between the Brocade port and the directly connected device.
Logical Link	The state of the logical link. This is the state of the link between this Brocade port and the Brocade port on the other end of the link.
State	The traffic state of the port.
Link-vlan	The ID of the tagged VLAN in the UDLD packet.

If a port is disabled by UDLD, the change also is indicated in the output of the **show interfaces brief** command. An example is given below.

```
device#show interfaces brief
Port  Link State      Dupl Speed Trunk Tag Priori MAC          Name
1/1   Up   LK-DISABLE
None  None  None  No  level0 0000.00a9.bb00
1/2   Down None       None None  None  No  level0 0000.00a9.bb01
1/3   Down None       None None  None  No  level0 0000.00a9.bb02
1/4   Down None       None None  None  No  level0 0000.00a9.bb03
```

If the port was already down before you enabled UDLD for the port, the port state is listed as None.

Syntax: **show interfaces brief**

Displaying information for a single port

To display detailed UDLD information for a specific port, enter a command such as the following.

```
device#show link-keepalive ethernet 4/1
Current State      : up          Remote MAC Addr  : 0000.00d2.5100
Local Port         : 4/1        Remote Port     : 2/1
Local System ID   : e0927400  Remote System ID : e0d25100
Packets sent      : 254       Packets received : 255
Transitions        : 1          Link-vlan      : 100
```

Syntax: show link-keepalive [ethernet [slotnum] portnum]

TABLE 26 CLI display of detailed UDLD information

Field	Description
Current State	The state of the logical link. This is the link between this Brocade port and the Brocade port on the other end of the link.
Remote MAC Addr	The MAC address of the port or device at the remote end of the logical link.
Local Port	The port number on this Brocade device.
Remote Port	The port number on the Brocade device at the remote end of the link.
Local System ID	A unique value that identifies this Brocade device. The ID can be used by Brocade technical support for troubleshooting.
Remote System ID	A unique value that identifies the Brocade device at the remote end of the link.
Packets sent	The number of UDLD health-check packets sent on this port.
Packets received	The number of UDLD health-check packets received on this port.
Transitions	The number of times the logical link state has changed between up and down.
Link-vlan	The ID of the tagged VLAN in the UDLD packet.

The **show interface ethernet** command also displays the UDLD state for an individual port. In addition, the line protocol state listed in the first line will say "down" if UDLD has brought the port down. An example is given below.

```
device#show interface ethernet 1/1
FastEthernet1/1 is down, line protocol is up, link keepalive is enabled
Hardware is FastEthernet, address is 0000.00a9.bbca (bia 0000.00a9.bbca)
Configured speed auto, actual unknown, configured duplex fdx, actual unknown
Member of L2 VLAN ID 1, port is untagged, port state is DISABLED
STP configured to ON, priority is level10, flow control enabled
mirror disabled, monitor disabled
Not member of any active trunks
Not member of any configured trunks
No port name
300 second input rate: 0 bits/sec, 0 packets/sec, 0.00% utilization
300 second output rate: 0 bits/sec, 0 packets/sec, 0.00% utilization
0 packets input, 0 bytes, 0 no buffer
Received 0 broadcasts, 0 multicasts, 0 unicasts
0 input errors, 0 CRC, 0 frame, 0 ignored
0 runts, 0 giants, DMA received 0 packets
```

```
19 packets output, 1216 bytes, 0 underruns  
Transmitted 0 broadcasts, 19 multicasts, 0 unicasts  
0 output errors, 0 collisions, DMA transmitted 19 packets
```

In this example, the port has been brought down by UDLD. Notice that in addition to the information in the first line, the port state on the fourth line of the display is listed as DISABLED.

Clearing UDLD statistics

To clear UDLD statistics, enter the following command.

```
device#clear link-keepalive statistics
```

Syntax: clear link-keepalive statistics

This command clears the Packets sent, Packets received, and Transitions counters in the **show link keepalive ethernet [slotnum [/]portnum** display.

Protected link groups

NOTE

Protected Link group is not supported on ICX 7750, ICX 7250 and ICX 7450 devices.

A protected link group minimizes disruption to the network by protecting critical links from loss of data and power. In a protected link group, one port in the group acts as the primary or active link, and the other ports act as secondary or standby links. The active link carries the traffic. If the active link goes down, one of the standby links takes over.

During normal operation, the active port in a protected link group is enabled and the standby ports are logically disabled. If the active port fails, the Brocade device immediately enables one of the standby ports, and switches traffic to the standby port. The standby port becomes the new, active port.

About active ports

When you create a protected link group, you can optionally specify which port in the protected link group is the active port. If you do not explicitly configure an active port, the Brocade device dynamically assigns one. A dynamic active port is the first port in the protected link group that comes up (usually the lowest numbered port in the group).

Static and dynamic active ports operate as follows:

- A static active port (an active port that you explicitly configured) pre-empts other ports in the protected link group. So, if a static active link comes back up after a failure, the Brocade device will revert to this link as the active link.
- A dynamic active port (an active port assigned by the software) is non-pre-emptive. Therefore, if a dynamic active link comes back up after a failure, the Brocade device does not revert to this link, but continues carrying traffic on the current active link.

Using UDLD with protected link groups

You can use UDLD with protected link groups to detect uni-directional link failures and to improve the speed at which the device detects a failure in the link. Use UDLD and protected link groups

simultaneously when the FastIron X Series device is connected to a device with slower link detection times.

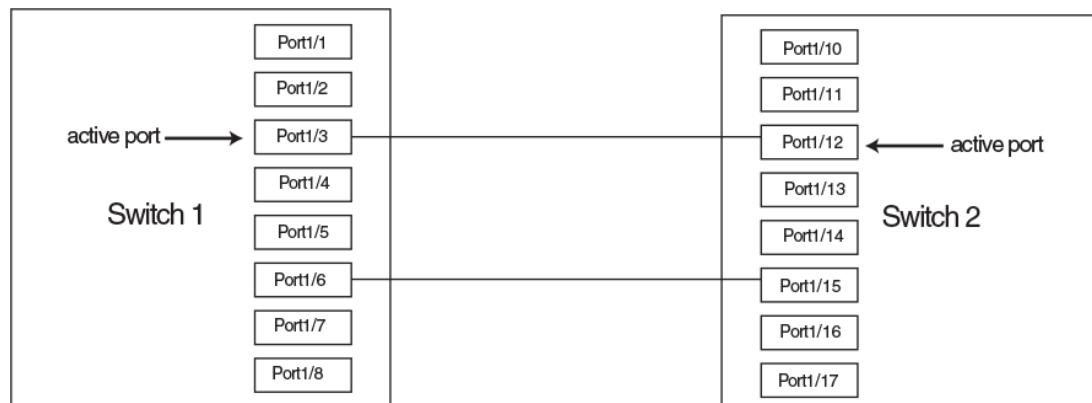
NOTE

When UDLD and protected links are configured on a port and the link goes down, protected links will not come up after UDLD becomes "healthy" again without first physically disabling then re-enabling the link.

UDLD with protected link groups configuration notes

- You can configure a maximum of 32 protected link groups.
- There is no restriction on the number of ports in a protected link group.
- Each port can belong to one protected link group at a time.
- On FastIron X Series devices, there is no restriction on the type of ports in a protected link group. FastIron X Series devices support protected link groups consisting of 10-GbE ports, Gbps fiber ports, 10/100/1000 copper ports, and 10/100 ports, or any combination thereof.
- This feature is supported with tagged and untagged ports.
- This feature is supported with trunk ports.
- The protected link groups feature is not supported with LACP.
- There is no restriction on the properties of ports in a protected link group. For example, member ports can be in the same VLAN or in different VLANs.
- When two switches are connected together with links in a protected link group, and the ports connecting the switches together are part of a protected link group, you must configure two connecting ports (one port on each switch) as active ports of the protected link group. The following example illustrates this scenario.

FIGURE 21



The configuration for the above illustration is as follows.

Switch 1

```
device(config)# protected-link-group 1 e 1/3 e 1/6
device(config)# protected-link-group 1 active-port 1/3
```

Switch 2

```
device(config)# protected-link-group 1 e 1/12 e 1/15
device(config)# protected-link-group 1 active-port 1/12
```

Creating a protected link group and assigning an active port

Follow the steps given below to create a protected link group.

1. Specify the member ports in the protected link group. Enter a command such as the following.

```
device(config)#protected-link-group 10 e 1 to 4
```

2. Optionally specify which port will be the active port for the protected link group. Enter a command such as the following.

```
device(config)#protected-link-group 10 active-port 1
```

NOTE

If you do not explicitly configure an active port, the Brocade device automatically assigns one as the first port in the protected link group to come up.

These commands configure port e1 as the active port and ports e2 - e4 as standby ports. If port 1 goes down, the Brocade device enables the first available standby port, and switches the traffic to that port. Since the above configuration consists of a statically configured active port, the active port pre-empts other ports in the protected link group. Refer to [About active ports](#) on page 128.

Syntax: [no] protected-link-group *group-ID* ethernet *port to port*

The *group-ID* parameter specifies the protected link group number. Enter a number from 1 - 32.

Syntax: [no] protected-link-group *group-ID* active-port ethernet *port*

The *group-ID* parameter specifies the protected link group number. Enter a number from 1 - 32.

Viewing information about protected link groups

You can use the following **show** commands to view information about protected link groups:

- **show protected-link-group**
- **show interface brief**
- **show interface**

The following shows example output for the **show protected-link-group** command.

```
device#show protected-link-group
Group ID: 1
Member Port(s): ethe 1 to 7
Configured Active Port: 7
Current Active Port: 7
Standby Port(s): ethe 5
Total Number of Protected Link Groups: 1
```

Syntax: show protected-link-group [*group-ID*]

TABLE 27 CLI display of protected link group information

Field	Description
Group ID	The ID number of the protected link group.
Member Port(s)	The ports that are members of the protected link group.
Configured Active Port	The statically configured active port. If you do not statically configure an active port, this value will be "None".
Current Active Port	The current active port for the protected link group. If all member ports are down, this value will be "None".

TABLE 27 CLI display of protected link group information (Continued)

Field	Description
Standby Port(s)	The member ports that are on standby.

The **show interface brief** command also displays information about protected link groups.

```
device#show int brief e 3 to 4
Port Link State Dupl Speed Trunk Tag Priori MAC Name
3 Up Inactive Full Auto None Yes level0 0000.00a8.7140
4 Up Forward Full 1G None Yes level0 0000.00a8.7140
```

In the above output, the **State** of port 3 is **Inactive**, which means port 3 is an inactive port in a protected link group. For active ports in a protected link group, the **State** will be **Active**.

Syntax: **show interface brief ethernet port**

The **show interface** command also displays information about protected link groups.

```
device#show int e 3
GigabitEthernet3 is up, line protocol is up, link keepalive is enabled
Port up for 60 seconds
    Hardware is GigabitEthernet, address is 0000.00a8.7140 (bia 0000.00a8.7142)
    Configured speed auto, actual 1Gbit, configured duplex fdx, actual fdx
    Configured mdi mode AUTO, actual MDIX
    Member of 3 L2 VLANs, port is tagged, port state is protected-link-inactive
    BPDU guard is Disabled, ROOT protect is Disabled
    Link Error Dampening is Disabled
    STP configured to ON, priority is level0
    ...
some lines omitted for brevity
```

In the above output, the port state is **protected-link-inactive** which means port 3 is an inactive port in a protected link group.

Syntax: **show interface ethernet port**

Link Aggregation

• Overview of link aggregation.....	133
• LAG formation rules.....	134
• Configuration notes for FastIron devices in a traditional stack.....	135
• Maximum number of LAGs.....	137
• Migrating from a previous release to 08.0.00a LAG or LACP configuration.....	138
• Downgrade considerations.....	139
• LAG Load Sharing.....	139
• LAG hashing on stacking products	140
• Configuring a LAG.....	141
• Deploying a LAG.....	146
• Preboot eXecution Environment boot support.....	157

Overview of link aggregation

This chapter describes how to configure Link Aggregation Groups (LAG). Beginning with FastIron 08.0.00a, you can use a single interface to configure any of the following LAG types:

Static LAGs - These LAGs are manually-configured aggregate links containing multiple ports.

Dynamic LAGs - This LAG type uses the Link Aggregation Control Protocol (LACP), to maintain aggregate links over multiple port. LACP PDUs are exchanged between ports on each device to determine if the connection is still active. The LAG then shuts down ports whose connection is no longer active.

NOTE

The LAG functionality was referred to as Trunk Groups in previous releases.

Keep Alive LAGs - In a Keep Alive LAG a single connection between a single port on 2 Brocade devices is established. In a keep alive LAG, LACP PDUs are exchanged between the 2 ports to determine if the connection between the devices is still active. If it is determined that the connection is no longer active, the ports are blocked.

NOTE

The Keep Alive LAG functionality was referred to as Single Link LACP in previous releases.

The new LAG configuration procedures supersede the previous configurations procedures for LAGs and Dynamic Link Aggregation. When a Brocade device is upgraded to 08.0.00a any configurations for LAGs or Dynamic Link Aggregation defined in releases prior to 08.0.00a will be converted to a 08.0.00a (and later) compatible LAG configuration. Details about how this conversion is performed are described in [Migrating from a previous release to 08.0.00a LAG or LACP configuration](#) on page 138. LAG formation rules

LAG formation rules

- A port can be member of only a single LAG which can be static, dynamic, or keep-alive LAG.
- Flexible LAG membership: Any number or combination of ports between 1 and 16 within the device can be used to configure a LAG. FSX third generation modules in an all third generation system can have 1 to 12 ports. For Brocade ICX 6430, Brocade ICX 6450, Brocade ICX 6610, and Brocade ICX 6650 devices, the port range is between 1 and 8. Brocade ICX 7750, Brocade ICX 7450, and Brocade ICX 7250 devices can have 1 to 16 ports. The maximum number of LAG ports is checked when adding ports to a LAG. Ports in a LAG could be on different linecard modules in a chassis or could be on different units in a stack.
- All ports configured in a LAG must be of equal bandwidth. For example all 10 G ports.
- All ports configured in a LAG must be configured with the same port attributes.
- LAG formation rules are checked when a static or dynamic LAG is deployed.
- A LAG must have its primary port selected before it can be deployed.
- All ports configured in a LAG must be configured in the same VLAN.
- Layer 2 requirements:

The LAG is rejected if the LAG ports:

- Do not have the same untagged VLAN component.
- Do not share the same VLAN membership or do not share the same uplink VLAN membership
- Do not share the same protocol-VLAN configuration
- Are configured as mainly primary and secondary interfaces
- Layer 3 requirements:

The LAG is rejected if any of the secondary LAG port has any Layer 3 configurations, such as IPv4 or IPv6 address, OSPF, RIP, RIPNG, IS-IS, and so on.

- Layer 4 (ACL) requirements:

All LAG ports must have the same ACL configurations; otherwise, the LAG is rejected.

- All LAG member properties must match the primary port of the LAG with respect to the following parameters:

- Port tag type (untagged or tagged port)
- Port dual-mode
- Default port speed and duplex
- Configured port speed and duplex
- TOS-based Configuration - During deployment the configuration on the primary port is replicated to all ports and on undeployment, each port inherits the same TOS-based QoS configuration.

To change port parameters, you must change them on the primary port. The software automatically applies the changes to the other ports in the LAG.

- Make sure the device on the other end of the LAG link can support the same number of ports in the link.
- Legacy ports and 48 Gbps copper ports cannot be members of the same LAG in hardware configurations such as the following:
 - 48-port 10/100/1000 Mbps (RJ45) Ethernet PoE interface module (SX-FI48GPP) and IPv4/IPv6 interface modules or management modules with user ports.
- Combination of different generations of ports cannot be members of the same LAG under the following hardware configurations:
 - 24-port fiber and copper Ethernet PoE/Fiber interfaces (SX-FI24GPP, SX-FI24GF) and IPv4/IPv6 interface modules or management modules with user ports
 - 2-port 10G and 8-port 10G interfaces (SX-FI8XG, SX-FI2XG) and the IPv4/IPv6 interface modules or management modules with user ports

- Combination of copper and fibre ports, even if they are of the same speed, cannot be members of the same LAG.
- The FastIron SX chassis supports up to 12 ports per LAG and a maximum of 255 LAGs only when the following third generation interface modules are installed in the chassis:
 - SX-FI48GPP--48-port 10/100/1000 Mbps Ethernet PoE interface module
 - SX-FI24GPP--24-port Gigabit Ethernet copper interface module
 - SX-FI24HF--24-port Gigabit Ethernet fiber interface module
 - SX-FI2XG--2-port 10 Gigabit Ethernet interface module
 - SX-FI8XG--8-port 10 Gigabit Ethernet interface module
- Cross module LAG is only supported between the same generation SX cards.
- LAG is supported on 1-GbE, 10-GbE or 40-GbE ports.
- You cannot combine 1-GbE and 10-GbE ports in the same LAG.
- Port assignment on a module need not be consecutive. The port range can contain gaps. For example, you can configure ports 1, 3, and 4 (excluding 2).
- Although the FastIron devices have port ranges, they do not apply to LAGs.
- You can select any port to be the primary port of the LAG.
- Make sure the device on the other end of the LAG can support the same number of ports in the link. For example, if you configure a 3-port LAG on the FSX and the other end is a different type of switch, make sure the other switch can support a 3-port LAG.
- All the ports must be connected to the same physical or logical device at the other end.
- The sFlow configuration enabled on the primary port of a LAG is applicable to all the LAG ports. Disabling the sFlow on the primary port of a LAG removes the configuration from all the LAG ports.

Configuration notes for FastIron devices in a traditional stack

In a Brocade traditional stack system, a LAG may have port members distributed across multiple stack units. Both static and dynamic LAGs are supported.

NOTE

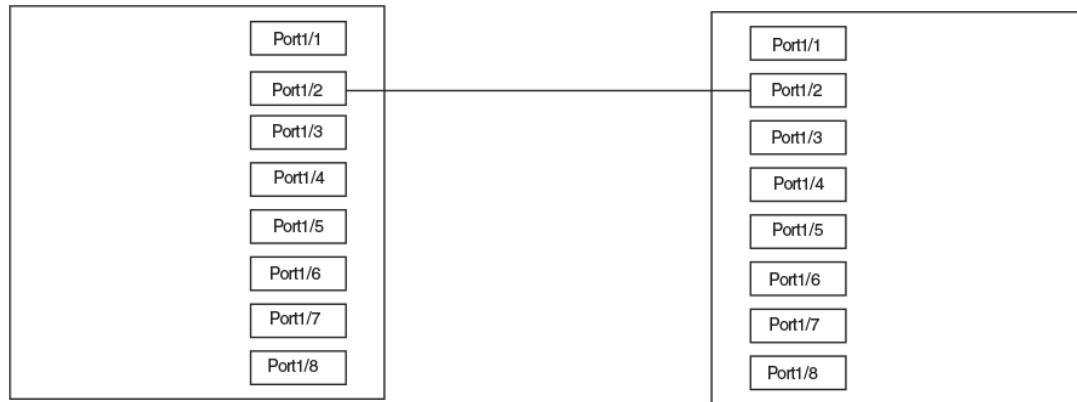
Cascaded LAG between stack units are supported on Brocade ICX devices only.

The following notes apply to FastIron stackable devices that are part of a traditional stack:

- If a stack unit fails or is removed from the stack, its static LAG configuration becomes a reserved configuration on the Active Controller. Any remaining ports of the static LAG in the traditional stack continue to function.
- When a new stack unit is added to a traditional stack, the new unit receives running configuration and LAG information, including a list of ports that are up and are members of a LAG, from the Active Controller.
- Before merging two traditional stack devices, make sure that there are no static LAGs configured between them. This can result in self-looped ports.
- You cannot configure LAG between cross units in a mixed stack.
- You cannot create a LAG between devices that belong to different product lines in a mixed stack. For example, between an ICX 6610 and an ICX 6450.
- When a traditional stack device with static LAG partitions into multiple traditional stacks, loops and forwarding errors may occur. In these cases, user intervention is required to remove the loops.
- 10 Gbps links support up to eight ports in a LAG for stackable units. On Brocade ICX 7750, Brocade ICX 7450, and Brocade ICX 7250 devices, 16 ports are supported in a LAG.

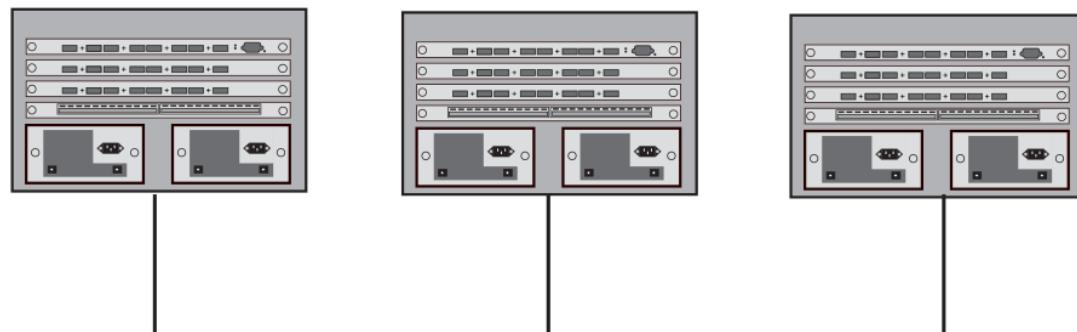
The following figure displays an example of a valid, Keep ALIVE LAG link between two devices. This configuration does not aggregate ports but uses the LACP PDUs to maintain the connection status between the two ports.

FIGURE 22 Example of a 1-port keep alive LAG

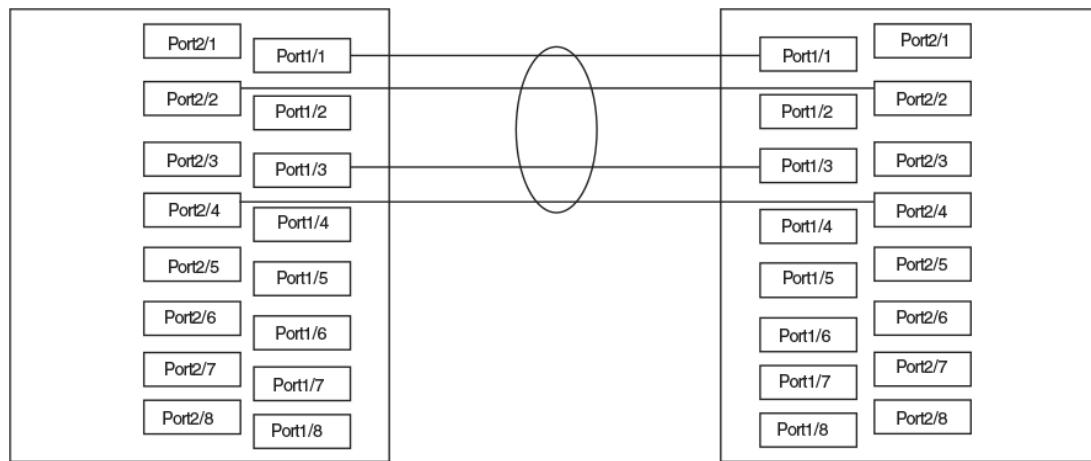


The following figure shows an example of a valid 2-port LAG link between devices where the ports on each end are on the same interface module. Ports in a valid 2-port LAG on one device are connected to two ports in a valid 2-port LAG on another device.

FIGURE 23 Example of 2-port LAG



The following figure shows an example of two devices connected over a 4 port LAG where the ports on each end of the LAG are on different interface modules.

FIGURE 24 Examples of multi-slot, multi-port LAG

Maximum number of LAGs

The following table lists the maximum number of LAGs you can configure on a Brocade device and the valid number of ports in a LAG. The table applies to static and LACP ports.

TABLE 28 Maximum number of LAGs

Model	Maximum number of LAGs		Valid number of ports in a group
	Static	LACP	
FCX624	124	124	1, 2, 3, 4, 5, 6, 7, or 8
FCX648			
FSX 800	31 (Supported on first and second generation modules)	31 (Supported on first and second generation modules)	First Generation IPv4 devices: 1, 2, 3, or 4
FSX 1600	255 (Supported on third generation modules)	64	Second Generation IPv6 devices: 1, 2, 3, 4, 5, 6, 7, or 8 Third Generation devices: 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, or 12
ICX 6430	29	29	1, 2, 3, 4, 5, 6, 7, or 8
ICX6430C-12	16	16	
ICX 6450	124	124	
ICX 6610	120	120	

TABLE 28 Maximum number of LAGs (Continued)

Model	Maximum number of LAGs		Valid number of ports in a group
	Static	LACP	
ICX 6650	124	124	1, 2, 3, 4, 5, 6, 7, or 8 For 40 Gbps ports, the valid number ports in a group are 2, 3, and 4.
ICX 7750	124 ¹	124 ¹	1 to 16
ICX 7450			
ICX 7250 ¹			

Migrating from a previous release to 08.0.00a LAG or LACP configuration

If you are upgrading from a version of the software prior to 08.0.00a and have either LAGs or LACP configured, the previous configuration will be automatically updated with the new commands to form an LAG that is equivalent to the previous configuration. To accomplish this, the old trunk and **link-aggregation** commands are maintained during startup configuration parsing, but disabled during normal configurations.

The following are the major differences between in LAG configuration in 08.0.00a and prior releases:

- A LAG is not created until a LAG is deployed using the **deploy** command.
- LACP is not started until a dynamic LAG is deployed.
- The number of LAG ports can range between 1 and 16 on Brocade ICX 7750, Brocade ICX 7450, and Brocade ICX 7250 devices. For FSX 3rd generation modules, the port range is between 1 and 12. For Brocade ICX 6430, Brocade ICX 6450, Brocade ICX 6610, and Brocade ICX 6650 devices, the port range is between 1 and 8. A LAG is created even if a static or dynamic LAG has only one port.

The following process is followed during the conversion of the trunk and **link-aggregation** to the new LAG commands.

1. For any static lag configured using the trunk **ethernet stack/slot/port to stack/slot/port** command, the following conversion procedure is followed.
 - a) A static LAG is created containing the port list specified in the trunk command. This LAG is then automatically deployed.
 - b) The lowest-numbered port from the original trunk list is selected as the primary port of the LAG.
 - c) The converted LAG is named "LAG_x", where "x" is a unique number assigned by the system starting from 1.
2. For any dynamic link aggregation (LACP) group configured using the port-level **link-aggregate** commands, the following conversion procedure is followed.

¹ In the Brocade ICX 7250G, the maximum number of LAGs that can be configured is 24.

- a) A dynamic LAG is created by grouping all ports in the original configuration having the same link-aggregation key.
- b) If **link-aggregate active/pассив** is configured originally, the converted dynamic LAG is configured as deployed, otherwise is not be converted because such ports were originally not operating under LACP.
- c) If the original mode is passive, the converted dynamic LAG will be configured as **deploy passive**. Otherwise active mode is the default.
- d) The timeout configuration set by the command **link-aggregate configure timeout** will be converted to the **lacp-timeout** command.
- e) The value of the link-aggregate configure key command is used in the conversion in determining the set of ports that form an LAG, so prior to upgrade the key must be configured on all the link-agg groups. In the new LAG user interface, there is no need for a user to explicitly configure a key. Each dynamic LAG will automatically select a unique key for the system. Hence the original configured key will not be retained.

NOTE

You cannot copy configurations to the running configurations from a TFTP server. It is not supported when you upgrade from a software version earlier to 08.0.00a to 8.0 configuration. The configuration must be saved on flash as startup configuration and reload is required without write-memory.

- f) The command **link-aggregate configure system-priority** is retired and will not be directly converted. This value is currently not in use by the system's LACP protocol processing, and will maintain a default value of 1.
- g) The lowest-numbered port will be selected as the primary port of the LAG.
- h) Port names configured in the original interface configuration will be converted to port names within the LAG.
- i) The converted LAG will be named "LAG_x", where "x" is a unique number assigned by the system starting from 1.

Downgrade considerations

When you downgrade to an earlier software version from 08.0.00a, all the existing LAG configuration will be lost.

LAG Load Sharing

Brocade devices load-share across the ports in the LAG group. The method used for the load sharing depends on the device type and traffic type (Layer 2 or Layer 3).

Support for IPv6 when sharing traffic across a LAG group

Brocade devices that support IPv6 take the IPv6 address for a packet into account when sharing traffic across a LAG group. The load sharing is performed in the same way it is for IPv4 addresses; that is, LAG types with a traffic load that is shared based on IPv4 address information can now use IPv6 addresses to make the load sharing decision.

Load balancing for unknown unicast, multicast, and broadcast traffic

Brocade devices load balance unknown unicast, multicast, and broadcast traffic based on the source port and VLAN ID and not on any source or destination information in the packet.

For example, when the switch receives unknown unicast, multicast, and broadcast packets, and the packets are from the same source port, the packets are forwarded to the same port of the LAG group. Conversely, when the switch receives unknown unicast, multicast, and broadcast packets, and the packets are from different source ports, the packets are load balanced across all the ports of the LAG group.

Note that this does not apply to known unicast traffic, which is always load balanced across all the ports of a LAG group based on the traffic's Layer 2 and Layer 3 source and destination parameters.

How LAG load sharing works

The load balancing method for bridged traffic varies depending on the traffic type. Load balancing for routed traffic is always based on the source and destination IP addresses and protocol field (applicable only for Brocade ICX 7750, Brocade ICX 7450, and Brocade ICX 7250 stackable devices).

TABLE 29 LAG Load sharing on FastIron devices

Traffic type	Load balancing method
Layer 2 Bridged non-IP	Source and destination MAC addresses
Layer 2 Bridged TCP/UDP	Source and destination MAC addresses, source and destination IP addresses, and source and destination TCP/UDP ports.
Layer 2 Bridged IP (non-TCP/UDP)	Source and destination MAC addresses, and source and destination IP addresses.
Layer 3 Routed traffic	Source and destination IP addresses and protocol field.
Layer 2 Bridged IPv4 TCP/UDP	Source and destination IP addresses, and source and destination TCP/UDP ports.
Layer 2 Bridged IPv4 Non-TCP/UDP	Source and destination IP addresses.
Layer 2 Bridged IPv6 TCP/UDP	Source and destination IP addresses, source and destination TCP and UDP ports, and flow label.
Layer 2 Bridged IPv6 Non-TCP/UDP	Source and destination TCP and UDP ports, and flow label.
Layer 3 Routed traffic	Source and destination IP addresses and protocol field.

LAG hashing on stacking products

This configuration is required when multicast routing is configured on a tunnel interface and if the IP multicast packets are tunnel terminating. For example, when `ip pim`, `ip pim-sparse`, `ip igmp proxy` multicast routing commands are configured on a tunnel interface.

Removing Layer 2 information from LAG hash output

NOTE

Removing Layer 2 information from LAG hash output is not supported on FastIron X Series devices and on Brocade ICX 7750, Brocade ICX 7450, and Brocade ICX 7250 devices.

FastIron stackable devices support the option to exclude Layer 2 information in the LAG hash calculation for IP packets.

```
Brocade(config) # lag-hash exclude-l2mac
```

The **lag-hash exclude-l2mac** command removes Layer 2 information (text in bold) to the following load-balancing parameters:

- Non-IP: Source and destination MAC addresses
- IPv4 TCP/UDP: Source and destination IP addresses, and source and destination TCP/UDP ports, **Source MAC , Destination MAC**
- IPv4 Non-TCP/UDP: Source and destination IP addresses, **Source MAC , Destination MAC**
- IPv6 TCP/UDP: Source and destination IP addresses, source and destination TCP and UDP ports, and flow label, **Source MAC , Destination MAC**
- IPv6 Non-TCP/UDP: Source and destination TCP and UDP ports, and flow label, **Source MAC , Destination MAC**

Configuring a LAG

The following configuration procedures are used to configure a LAG. Depending upon whether you are configuring a static, dynamic or keep-alive LAG, the configuration procedures may or may not apply as described:

- Creating a Link Aggregation Group - Required for all static, dynamic or keep alive LAGs.
- Adding Ports to a LAG - Required for all static, dynamic, or keep alive LAGs. A keep alive LAG contains only one port while static and dynamic LAGs can have 1 to 12 ports.
- Configuring the Primary Port for a LAG - Required for all static and dynamic LAGs. Since a keep alive LAG contains only one port, it is unnecessary to configure this parameter.
- Configuring the Load Sharing Type - Optional for all static and dynamic LAGs. Since a keep alive LAG contains only one port, it is unnecessary to configure this parameter.
- Specifying the LAG Threshold for a LAG Group - Optional for static and dynamic LAGs. Since a keep alive LAG contains only one port, it is unnecessary to configure this parameter.
- Configuring an LACP Timeout - Optional for dynamic and keep alive LAGs.

Creating a Link Aggregation Group (LAG)

Before setting-up ports or configuring any other aspects of a LAG, you must create it as shown in the following:

```
device(config) # lag blue static
device(config-lag-blue) #
```

Syntax: [no] lag *lag-name* { static | dynamic | keep-alive }

The **static** option specifies that the LAG with the name specified by the *lag-name* variable will be configured as a static LAG.

The **dynamic** option specifies that the LAG with the name specified by the *lag-name* variable will be configured as a dynamic LAG.

The **keep-alive** option specifies that the LAG with the name specified by the *lag-name* variable will be configured as a keep-alive LAG. The keep-alive LAG configuration is a new configuration option to configure a LAG for use in keep alive applications similar to the UDLD feature.

Creating a Link Aggregation Group (LAG) using the LAG ID option

Before setting-up ports or configuring any other aspects of a LAG, you must create it first.

You can either assign a LAG **ID** explicitly or it will be automatically generated by the system. The LAG **ID** remains the same across system reload and hitless upgrade.

The command to configure LAGs allows explicit configuration of the LAG ID for static and dynamic LAGs.

To create a LAG with the LAG ID option, enter a command such as the following.

```
device(config)# lag blue static id 1  
device(config-lag-blue)#[/pre]
```

Syntax: [no] **lag** *lag-name* [{ **static** | **dynamic** } [**id** *number*]]

The *lag-name* is an ASCII string and can have a maximum of 64 characters.

The **id** parameter is optional. The value of the **id** parameter that you can enter is from 1 to 2047. If you do not enter a LAG **ID**, the system will generate one automatically. Once the LAG **ID** is generated the system will save it in the configuration file along with the LAG name, therefore the value will stay the same across system reload.

NOTE

The LAG **id** parameter is for static and dynamic LAGs only. No explicit configuration of a LAG ID is allowed on keepalive LAGs.

The **static** parameter specifies that the LAG with the name specified by the *lag-name* variable will be configured as a static LAG.

The **dynamic** option specifies that the LAG with the name specified by the *lag-name* variable will be configured as a dynamic LAG.

Configuration considerations

LAG IDs are unique for each LAG in the system. The same LAG ID cannot be assigned to two or more different LAGs. If a LAG ID is already used, the CLI will reject the new LAG configuration and display an error message that suggests the next available LAG ID that can be used.

```
device(config)#lag lag3 static id 123  
Error: LAG id 123 is already used. The next available LAG id is 2  
.[/pre]
```

NOTE

If you upgrade from an earlier version to a version with the LAG ID configuration feature, the old configuration file will be parsed correctly and each LAG configured will get a LAG ID automatically.

```
!  
lag lag1 static id 124  
ports ethernet 1/1/2 to 1/1/3  
primary-port 1/1/3  
deploy  
![/pre]
```

: **show lag** command and the output.

```
device(config)# show lag
```

```

Total number of LAGs:      5
Total number of deployed LAGs: 3
Total number of trunks created:2 (253 available)
LACP System Priority / ID:    1 / 0024.3889.3b00
LACP Long timeout:           120, default: 120
LACP Short timeout:          3, default: 3
==== LAG "test" ID 35 (static Deployed) ====
LAG Configuration:
  Ports:          e 3/10
  Port Count:     1
  Primary Port:   3/10
  Trunk Type:    hash-based
Deployment: HW Trunk ID 1
Port   Link   State   Dupl Speed Trunk Tag Pvid Pri MAC          Name
3/10  Down   None    None None  35   No   1    0 0024.3889.3b09
==== LAG "test2" ID 1 (static Deployed) ====
LAG Configuration:
  Ports:          e 3/11
  Port Count:     1
  Primary Port:   3/11
  Trunk Type:    hash-based
Deployment: HW Trunk ID 2
Port   Link   State   Dupl Speed Trunk Tag Pvid Pri MAC          Name
3/11  Down   None    None None  1   No   1    0 0024.3889.3b0a
==== LAG "test3" (keep-alive Deployed) ====
LAG Configuration:
  Ports:          e 3/12
  Port Count:     1
  Primary Port:   3/12
  Trunk Type:    hash-based
  LACP Key:      9860
Deployment:
Port   Link   State   Dupl Speed Trunk Tag Pvid Pri MAC          Name
3/12  Down   None    None None  No  1    0 0024.3889.3b0b
Port   [Sys P] [Port P] [ Key ] [Act][Tio][Agg][Syn][Col][Dis][Def][Exp][Ope]
3/12      1       1    9860   Yes   S   Agg   Syn  No  No  Def  No  Dwn
Partner Info and PDU Statistics
Port      Partner      Partner      LACP      LACP
          System MAC      Key      Rx Count  Tx Count
3/12      0000.0000.0000  139      0         0
==== LAG "test4" (keep-alive Not Deployed) ====
LAG Configuration:
  Ports:          e 3/13
  Port Count:     1
  Primary Port:   3/13
  Trunk Type:    hash-based
  LACP Key:      0
==== LAG "test5" ID 2 (static Not Deployed) ====
LAG Configuration:
  Ports:          e 3/14
  Port Count:     1
  Primary Port:   none
  Trunk Type:    hash-based
  Hardware failover mode: all-ports

```

Creating a keepalive LAG

To create a **keep-alive** LAG, enter the following.

```
device(config)# lag lag1 keep-alive
```

Syntax: [no] lag lag-name [keep-alive]

The **keep-alive** option specifies that the LAG with the name specified by the lag-name variable will be configured a keep-alive LAG. The keep-alive LAG option allows you to configure a LAG for use in keep alive applications similar to the UDLD feature.

Adding Ports to a LAG or Deleting Ports from a LAG

A static or dynamic LAG can consist of 1 to 8, 1 to 12 or 1 to 16 ports (depending on the device you are using) of the same type and speed that are on any interface module within the Brocade chassis. A keep alive LAG consists of only one port.

To configure the static LAG named "blue" with two ports, use the following command:

```
device(config)# lag blue static  
device(config-lag-blue)# ports ethernet 1/3/1 ethernet 1/3/2
```

Syntax: [no] ports ethernet stack[slot]port [to stack[slot]port] [ethernet stack[slot]port]

The ports added to a LAG can be of type **ethernet** as specified for the **stack[slot]port** where they reside. The ports can be added to the LAG sequentially as shown in the following example:

```
device(config-lag-blue)# ports ethernet 1/3/1 ethernet 1/7/2 ethernet 1/4/3 ethernet 1/3/4
```

A range of ports from a single interface module can be specified. In the following example, Ethernet ports 1, 2, 3 and 4 on the interface module in slot 3 are configured in a single LAG:

```
device(config-lag-blue)# ports ethernet 1/3/1 to 1/3/4
```

Additionally, you can mix a range of ports from one interface module with individual ports from other interface modules to form a LAG as shown in the following:

```
device(config-lag-blue)# ports ethernet 1/3/1 to 1/3/4 ethernet 1/2/2
```

Using the **no** option allows you to remove ports from a LAG. For example, you can remove port 1/3/4 from the LAG created above, as shown in the following:

```
device(config-lag-blue)# no ports ethernet 1/3/4
```

Ports can be added to an undeployed LAG or to currently deployed LAG using the commands described. For special considerations when adding ports to or deleting ports from a currently deployed LAG, refer to the following sections:

- [Adding a Port to Currently Deployed LAG](#) on page 148
- [Deleting a Port from a Currently Deployed LAG](#) on page 148

Configuring the primary port for a LAG

The primary port must be explicitly assigned using the **primary-port** command.

To designate the primary port for the static LAG "blue", use the following command.

```
device(config)# lag blue static  
device(config-lag-blue)# primary-port 1/3/2
```

Syntax: [no] primary-port stack[slot]port

Once a primary port has been configured for a LAG, all configurations that apply to the primary port are applied to the other ports in the LAG.

NOTE

This configuration is only applicable for configuration of a static or dynamic LAGs.

Specifying the LAG threshold for a LAG group

You can configure the Brocade device to disable all of the ports in a LAG group when the number of active member ports drops below a specified threshold value. When a LAG is shut down because the number of ports drops below the configured threshold, the LAG is kept intact and it is re-enabled if enough ports become active to reach the threshold. For example, if a LAG group has 8 ports, and the threshold for the LAG group is 5, then the LAG group is disabled if the number of available ports in the LAG group drops below 5. If the LAG group is disabled, then traffic is forwarded over a different link or LAG group.

NOTE

This configuration is only applicable for only the configuration of static LAGs.

For example, the following commands establish a LAG group consisting of 4 ports, then establish a threshold for this LAG group of 3 ports.

```
device(config)# lag blue static
device(config-lag-blue)# ports ethernet 1/3/1 to 1/3/4
device(config-lag-blue)# trunk-threshold 3
```

In this example, if the number of active ports drops below 3, then all the ports in the LAG group are disabled.

Syntax: [no] trunk-threshold *number*

You can specify a threshold from 1 (the default) up to the number of ports in the LAG group.

When a LAG is shut down because the number of ports drops below the configured threshold, the LAG is kept intact and it is re-enabled if enough ports become active to reach the threshold.

NOTE

The trunk-threshold command should be configured only at one end of the trunk. If it is set on both sides, link failures will result in race-conditions and the will not function properly.

NOTE

The trunk-threshold command cannot be used in conjunction with protected link groups.

NOTE

Use a short LACP timeout when setting the trunk-threshold value equal to the number of links in the LAG or connecting to third party devices. See [Configuring an LACP timeout](#) on page 145.

Configuring an LACP timeout

In a dynamic or keep-alive LAG, a port's timeout can be configured as short (3 seconds) or long (90 seconds). After you configure a port timeout, the port remains in that timeout mode whether it is up or down and whether or not it is part of a LAG.

All the ports in a LAG should have the same timeout mode. This requirement is checked when the LAG is enabled on the ports. For example, to configure a port for a short LACP timeout, use the following command.

```
device(config)# lag blue dynamic
device(config-lag-blue)# lacp-timeout short
```

Syntax: [no] lacp-timeout [long | short]

To delete the configuration, use the **no** form of this command.

The **long** keyword configures the port for the long timeout mode--120 seconds. With the long timeout, an LACPDU is sent every 30 seconds. If no response comes from its partner after 3 LACPDUs are sent, a timeout event occurs, and the LACP state machine transition to the appropriate state based on its current state.

The **short** keyword configures the port for the short timeout mode--3 seconds. In the short timeout configuration, an LACPDU is sent every second. If no response comes from its partner after 3 LACPDUs are sent, a timeout event occurs, and the LACP state machine transitions to the appropriate state based on its current state.

If you specify neither **long** nor **short**, the state machine operates based on the standard IEEE specification as its default behavior. The original IEEE specification says that the state machine starts with short the timeout and moves to the long timeout after the LAG is established. However, sometimes a vendor's implementation always uses either the short timeout or the long timeout without changing the timeout. Brocade provides this command so that you can configure Brocade devices to interoperate with other vendor's devices.

NOTE

This configuration is applicable to the configuration of dynamic or keep-alive LAGs only.

Deploying a LAG

After configuring a LAG, you must explicitly enable it before it begins aggregating traffic. This task is accomplished by executing the **deploy** command within the LAG configuration. After the **deploy** command runs, the LAG is in the aggregating mode. Only the primary port within the LAG is available at the individual interface level. All the secondary ports should have the same IP directed-broadcast configuration as the primary port. Any configuration performed on the primary port applies to all ports within the LAG. The running configuration will no longer display deployed LAG ports other than the primary port.

To deploy a LAG, at least one port must be in the LAG and the primary port must be specified for non keep-alive LAGs. After a non keep-alive LAG is deployed, a LAG is formed. If there is only one port in the LAG, a single port LAG is formed. For a dynamic LAG, LACP is started for each LAG port. For a keep-alive LAG, no LAG is formed and LACP is started on the LAG port.

You can deploy a LAG as shown in the following for the "blue" LAG.

```
device(config)# lag blue static
device(config-lag-blue)# deploy
```

Syntax: [no] deploy [passive]

When the **deploy** command is executed:

For dynamic LAGs, LACP is activated on all LAG ports. When activating LACP, use active mode if **passive** is not specified; otherwise, use **passive** mode.

For a keep-alive LAGs, no LAG is formed, and LACP is started on the LAG port.

Once the **deploy** command is issued, all LAG ports will behave like a single port.

If the **no deploy** command is executed, the LAG is removed. For dynamic LAGs, LACP is de-activated on all of the LAG ports. All the secondary ports are disabled automatically and there will be no changes to the primary port.

Commands available under LAG once it is deployed

Once a LAG has been deployed, the following configurations can be performed on the deployed LAG:

- Disabling Ports within a LAG
- Enabling Ports within a LAG
- Monitoring and Individual LAG Port
- Assigning a name to a port within a LAG
- Enabling sFlow Forwarding on a port within a LAG
- Setting the sFlow Sampling Rate for a port within a LAG
- IP assignment within a LAG
- Renaming an existing LAG

Disabling ports within a LAG

You can disable an individual port within a LAG using the disable command within the LAG configuration as shown in the following.

```
device(config)# lag blue static
device(config-lag-blue)# disable ethernet 1/3/1
```

Syntax: [no] disable { ethernet stack/slot/port [to stack/slot/port] [ethernet stack/slot/port] | port-name name }

Use the **ethernet** option with the appropriate **stack/slot/port** variable to specify a Ethernet port within the LAG that you want to disable.

Use the **port-name** option with the appropriate **name** variable to specify a named port within the LAG that you want to disable.

To disable a port belonging to a keep-alive LAG, you need to configure from the interface level.

```
Brocade(config-lag-test)#interface e 7/8
Brocade(config-if-e1000-7/8)#disable
Brocade(config-if-e1000-7/8) #
```

Enabling ports within a LAG

You can enable an individual port within a LAG using the enable command within the LAG configuration as shown in the following.

```
device(config)# lag blue static
device(config-lag-blue)# deploy
device(config-lag-blue)# enable ethernet 1/3/1
```

Syntax: [no] enable { ethernet stack/slot/port [to stack/slot/port] [ethernet stack/slot/port] | port-name name }

Use the **ethernet** option with the appropriate **stack/slot/port** variable to specify a Ethernet port within the LAG that you want to enable.

Use the **port-name** option with the appropriate **name** variable to specify a named port within the LAG that you want to enable.

To enable a port belonging to a keep-alive LAG, you need to configure from the interface level.

```
Brocade(config-lag-test)#interface e 7/8
Brocade(config-if-e1000-7/8)#enable
Brocade(config-if-e1000-7/8) #
```

Adding a Port to Currently Deployed LAG

Ports can be added to a currently deployed LAG. Adding a port to a deployed LAG uses the same procedures as described in [Adding Ports to a LAG or Deleting Ports from a LAG](#) on page 144. When you add ports to a deployed LAG, the MAC address of the port being added is changed to that of the primary port of the LAG to which it is being added.

When you add a new secondary port to a currently deployed LAG, the IP directed-broadcast configuration and all other configuration of the new port should be the same as that of the primary port of the LAG.

NOTE

In an operational dynamic LAG, adding or removing a port causes port flapping for all LAG ports. This may cause loss of traffic.

Deleting a Port from a Currently Deployed LAG

Ports can be deleted from a currently deployed LAG. Deleting a port in a currently deployed LAG uses the same procedures as described in [Adding Ports to a LAG or Deleting Ports from a LAG](#) on page 144. However, when deleting ports from a currently deployed LAG you must consider the following:

- The primary port cannot be removed.
- If removal of a port will result in the trunk threshold value becoming greater than the number of ports in the LAG, the port deletion will be rejected.
- When you remove a port from a deployed LAG, the port is disabled automatically.

To delete port 1/3/1 which is in the "enabled" state from a currently deployed LAG named "blue", use the following command:

```
device(config)# lag blue static
device(config-lag-blue)# no ports ethernet 1/3/1
```

Syntax: [no] ports ethernet *stack/slot/port* [to *stack/slot/port*] [ethernet *stack/slot/port*]

NOTE

When a port is deleted from a currently deployed LAG, the MAC address of the port is changed back to its original value.

NOTE

In an operational dynamic LAG, removing an operational port causes port flapping for all LAG ports. This may cause loss of traffic.

Monitoring an individual LAG port

By default, when you monitor the primary port in a LAG group, aggregated traffic for all the ports in the LAG is copied to the mirror port. You can configure the device to monitor individual ports in a LAG including Ethernet, or named ports. You can monitor the primary port or another member port individually. Once a LAG is deployed and a primary port is specified using the primary-port command, monitoring across all ports of the LAG can be configured at the primary port. If a new port is added to a deployed LAG and if the entire LAG is monitored, the new port will also be mirrored by the same port monitoring traffic across the entire LAG.

NOTE

You can use only one mirror port for each monitored LAG port. You cannot configure mirroring on an undeployed LAG.

To monitor traffic on an individual port in a LAG group, run the following commands.

```
device(config)# lag blue static
device(config-lag-blue)# ports ethernet 1/1/1 ethernet 1/1/47
device(config-lag-blue)# primary-port 1/1/1
device(config-lag-blue)# deploy
device(config-lag-blue)# monitor ethe-port-monitored 1/1/47 ethernet 1/1/15 output
```

Syntax: [no] **monitor** { **ethe-port-monitored** *stack/slot/port* | **named-port-monitored** *name* } [**ethernet** [*stack/slot/port*]] { **input** | **output** | **both** }

Use the **ethe-port-monitored** option with the appropriate [*stack/slot/port*] variable to specify a Ethernet port within the LAG that you want to monitor.

Use the **named-port-monitored** option with the appropriate **name** variable to specify a named port within the LAG that you want monitor.

The **ethernet** *stack/slot/port* parameter specifies the port to which the traffic analyzer is attached.

The **input**, **output**, and **both** parameters specify the traffic direction to be monitored.

Assigning a name to a port within a LAG

You can assign a name to an individual port within a LAG using the **port-name** command within the LAG configuration as shown in the following.

```
device(config)#lag "test" dynamic id 1
device(config-lag-test)#ports ethernet 1/1/1 to 1/1/3
device(config-lag-test)#port-name "Brocade lag" ethernet 1/1/1
device(config-lag-test)#primary-port 1/1/1
device(config-lag-test)#deploy
```

Syntax: [no] **port-name** *name* **ethernet** *stack/slot/port*

The *name* variable specifies the port name. The name can be up to 255 characters long.

NOTE

Port name with space must be enclosed within double quotation marks.

Use the **ethernet** option with the appropriate *stack/slot/port* variable to apply the specified name to an Ethernet port within the LAG.

Allowable characters for LAG names

When creating a LAG name, you can use spaces in a file or subdirectory name if you enclose the name in double quotes. For example, to specify a subdirectory name that contains spaces, enter a string such as the following: "a long subdirectory name". The maximum length for a string is 64 characters.

The following characters are valid in file names:

- All upper and lowercase letters
- All digits

Any of the following special characters are valid:

- \$
- %
- '
- -
- —
- .
- @
- ~
- `
- !
- (
-)
- {
- }
- ^
- #
- &

Enabling sFlow forwarding on a port in a LAG

You can enable sFlow forwarding on an individual port within a LAG using the **sflow-forwarding** command within the LAG configuration as shown in the following.

```
device(config)# lag blue static
device(config-lag-blue)# deploy
device(config-lag-blue)# sflow forwarding ethernet 1/3/1
```

Syntax: [no] sflow forwarding { ethernet stack/slot/port | port-name name }

Use the **ethernet** option with the appropriate *stack/slot/port* variable to specify a Ethernet port within the LAG that you want to enable sFlow forwarding for.

Use the **port-name** option with the appropriate *name* variable to specify a named port within the LAG that you want to enable sFlow forwarding for.

For a keep-alive LAG, sFlow can be enabled only at the interface level and not at a lag context. To configure sFlow for an interface belonging to the keep-alive lag, configure directly under the interface.

```
Brocade(config-lag-test)#interface e 7/8
Brocade(config-if-e1000-7/8)#sflow forwarding
Brocade(config-if-e1000-7/8)#+
```

Setting the sFlow sampling rate for a port in a LAG

You can set the sFlow sampling rate for an individual port within a LAG using the **sflow-subsampling** command within the LAG configuration as shown in the following.

```
device(config)# lag blue static
device(config-lag-blue)# deploy
device(config-lag-blue)# sflow sample 512
```

Syntax: [no] sflow sample number

The *number* variable specifies the average number of packets from which each sample will be taken. The software rounds the value you enter up to the next odd power of 2. This can be a value between 8 - 1048576.

For a keep-alive LAG, you need to configure sFlow sampling at the interface level and not within the LAG configuration.

```
Brocade(config-lag-test)#interface e 7/8
Brocade(config-if-e1000-7/8)#sflow sample 512
Brocade(config-if-e1000-7/8)#+
```

IP assignment within a LAG

Layer 3 static or dynamic LAG support IP assignment. All the configurations has to be done on the primary port of the LAG.

The following is a sample configuration:

```
lag lag_dist_a 1 dynamic id 15
  ports ethe 1/1 to 1/12
  primary-port 1/1
  deploy
!
router vrrp
!
interface eth1 1/1
ip address 192.168.10.1 255.255.255.0
ip vrrp vrid 1
  backup priority 50 track-priority 10
  ip-address 192.168.1.10
  activate
```

Renaming an existing LAG

You can change the name of an existing LAG without causing any impact on the functionality of the LAG.

You can rename the LAG using the **update-lag-name** command within the LAG configuration mode. The new name provided must be unique and unused. The LAG configuration mode will exit after successful name update.

```
device(config)# lag blue static
device(config-lag-blue)# update-lag-name blue1
INFORMATION: Lag blue is updated to new name blue1
```

Displaying LAG information

You can display LAG information for a Brocade device in either a **full** or **brief** mode.

The following example displays the **brief** option of the **show lag** command.

```
device# show lag brief
Total number of LAGs: 5
Total number of deployed LAGs: 3
Total number of trunks created: 2 (253 available)
LACP System Priority / ID: 1 / 0024.3889.3b00
LACP Long timeout: 120, default: 120
LACP Short timeout: 3, default: 3
LAG      Type Deploy Trunk Primary Port List
test    static   Y     35   3/10   e 3/10
test2   static   Y      1    3/11   e 3/11
test3   keep-al  Y    1153   3/12   e 3/12
test4   keep-al  N    1154   3/13   e 3/13
test5   static   N      2    none    e 3/14
```

Syntax: show lag brief

Table 30 describes the information displayed by the **show lag brief** command.

The following example displays the full option of the **show lag** command.

```
device# show lag
Total number of LAGs: 5
Total number of deployed LAGs: 3
Total number of trunks created: 2 (253 available)
LACP System Priority / ID: 1 / 0024.3889.3b00
LACP Long timeout: 120, default: 120
LACP Short timeout: 3, default: 3
== LAG "test" ID 35 (static Deployed) ==
LAG Configuration:
  Ports: e 3/10
  Port Count: 1
  Primary Port: 3/10
  Trunk Type: hash-based
Deployment: HW Trunk ID 1
Port Link State Dupl Speed Trunk Tag Pvid Pri MAC Name
3/10 Down None None None 35 No 1 0 0024.3889.3b09
== LAG "test2" ID 1 (static Deployed) ==
LAG Configuration:
  Ports: e 3/11
  Port Count: 1
  Primary Port: 3/11
  Trunk Type: hash-based
Deployment: HW Trunk ID 2
Port Link State Dupl Speed Trunk Tag Pvid Pri MAC Name
3/11 Down None None None 1 No 1 0 0024.3889.3b0a
== LAG "test3" (keep-alive Deployed) ==
LAG Configuration:
  Ports: e 3/12
  Port Count: 1
  Primary Port: 3/12
  Trunk Type: hash-based
  LACP Key: 9860
Deployment:
Port Link State Dupl Speed Trunk Tag Pvid Pri MAC Name
3/12 Down None None None None No 1 0 0024.3889.3b0b
Port [Sys P] [Port P] [Key] [Act] [Tio] [Agg] [Syn] [Col] [Dis] [Def] [Exp] [Ope]
3/12 1 1 9860 Yes S Agg Syn No No Def No Dwn
Partner Info and PDU Statistics
Port Partner Partner LACP LACP
System MAC Key Rx Count Tx Count
3/12 0000.0000.0000 139 0 0
== LAG "test4" (keep-alive Not Deployed) ==
LAG Configuration:
  Ports: e 3/13
  Port Count: 1
  Primary Port: 3/13
  Trunk Type: hash-based
  LACP Key: 0
== LAG "test5" ID 2 (static Not Deployed) ==
LAG Configuration:
  Ports: e 3/14
  Port Count: 1
  Primary Port: none
  Trunk Type: hash-based
```

Syntax: **show lag [lag-name | brief | deployed | dynamic | id | keep-alive | static]**

Using command this without options displays information for all LAGs configured on the device.

The **lag-name** variable allows you to limit the display to information for a specific LAG.

The **id** option displays the output for the LAG specified by the ID.

The **brief** displays a brief output.

The **deployed** option limits the display to LAGs that are currently deployed.

The **dynamic** option limits the display to dynamic LAGs.

The **keep-alive** option limits the display to keep alive LAGs.

The **static** option limits the display to static LAGs.

The following table describes the information displayed by the **show lag** command.

TABLE 30 Show LAG information

This field...	Displays...
Total number of LAGS	The total number of LAGs that have been configured on the device.
Total number of deployed LAGS	The total number of LAGs on the device that are currently deployed.
Total number of trunks created	The total number of LAGs that have been created on the LAG. The total number of LAGs available are shown also. Since keep-alive LAGs do not use a LAG ID, they are not listed and do not subtract for the number of LAGs available.
LACP System Priority /ID	The system priority configured for the device. The ID is the system priority which is the base MAC address of the device.
LACP Long timeout	The number of seconds used for the LACP Long timeout mode. This is only applicable for dynamic or keep-alive LAGs.
LACP Short timeout	The number of seconds used for the LACP Short timeout mode. This is only applicable for dynamic or keep-alive LAGs.

The following information is displayed per-LAG in the show lag brief command.

LAG	The name of the LAG, LAG ID number, the configured type of the LAG: static, dynamic, or keep-alive, status of LAG deployment: deployed or not
-----	---

The following information is displayed per-LAG the show lag command for each LAG configured.

LAG Configuration

- Ports: List of ports configured with the LAG.
- Port Count Number of ports configured on the LAG.
- Primary Port: The primary port configured on the LAG.
- Trunk Type: The load sharing method configured for the LAG: hash-based.
- LACP Key The link aggregation key for the LAG.

Deployment

- LAG ID The LAG ID number.
 - Active Primary The port within the LAG where most protocol packets are transmitted. This is not the same as the configured Primary Port of the LAG.
- | | |
|------|--|
| Port | The chassis slot and port number of the interface. |
|------|--|

TABLE 30 Show LAG information (Continued)

This field...	Displays...
Link	The status of the link which can be one of the following: <ul style="list-style-type: none"> • up • down
State	The L2 state for the port.
Dupl	The duplex state of the port, which can be one of the following: <ul style="list-style-type: none"> • Full • Half • None
Speed	The bandwidth of the interface.
Trunk	The LAG ID of the port.
Tag	Indicates whether the ports have 802.1q VLAN tagging. The value can be Yes or No.
Pri	Indicates the Quality of Service (QoS) priority of the ports. The priority can be a value from 0-7.
MAC	The MAC address of the port.
Name	The name (if any) configured for the port.
Sys P	Lists the system priority configured for the device.
Port P	Lists the port's link aggregation priority.
Key	Lists the link aggregation key.
Act	Indicates the link aggregation mode, which can be one of the following: <ul style="list-style-type: none"> • No - The mode is passive on the port. If link aggregation is enabled (and the mode is passive), the port can send and receive LACPDU messages to participate in negotiation of an aggregate link initiated by another port, but cannot search for a link aggregation port or initiate negotiation of an aggregate link. • Yes - The mode is active. The port can send and receive LACPDU messages.
Tio	Indicates the timeout value of the port. The timeout value can be one of the following: <ul style="list-style-type: none"> • L - Long. The LAG group has already been formed and the port is therefore using a longer message timeout for the LACPDU messages exchanged with the remote port. Typically, these messages are used as confirmation of the health of the aggregate link. • S - Short. The port has just started the LACPDU message exchange process with the port at the other end of the link. The S timeout value also can mean that the link aggregation information received from the remote port has expired and the ports are starting a new information exchange.

TABLE 30 Show LAG information (Continued)

This field...	Displays...
Agg	Indicates the link aggregation state of the port. The state can be one of the following: <ul style="list-style-type: none"> • Agg - Link aggregation is enabled on the port. • No - Link aggregation is disabled on the port.
Syn	Indicates the synchronization state of the port. The state can be one of the following: <ul style="list-style-type: none"> • No - The port is out of sync with the remote port. The port does not understand the status of the LACPDU process and is not prepared to enter a LAG link. • Syn - The port is in sync with the remote port. The port understands the status of the LACPDU message exchange process, and therefore knows the LAG group to which it belongs, the link aggregation state of the remote port, and so on.
Dis	Indicates the collection state of the port, which determines whether the port is ready to send traffic over the LAG link: <ul style="list-style-type: none"> • Col - The port is ready to send traffic over the LAG link. • No - The port is not ready to send traffic over the LAG link.
Col	Indicates the distribution state of the port, which determines whether the port is ready to receive traffic over the LAG link. <ul style="list-style-type: none"> • Dis - The port is ready to receive traffic over the LAG link. • No - The port is not ready to receive traffic over the LAG link.
Def	Indicates whether the port is using default link aggregation values. The port uses default values if it has not received link aggregation information through LACP from the port at the remote end of the link. This field can have one of the following values: <ul style="list-style-type: none"> • Def - The port has not received link aggregation values from the port at the other end of the link and is therefore using its default link aggregation LACP settings. • No - The port has received link aggregation information from the port at the other end of the link and is using the settings negotiated with that port.
Exp	Indicates whether the negotiated link aggregation settings have expired. The settings expire if the port does not receive an LACPDU message from the port at the other end of the link before the message timer expires. This field can have one of the following values: <ul style="list-style-type: none"> • Exp - The link aggregation settings this port negotiated with the port at the other end of the link have expired. The port is now using its default link aggregation settings. • No - The link aggregation values that this port negotiated with the port at the other end of the link have not expired. The port is still using the negotiated settings.
Ope	<ul style="list-style-type: none"> • Ope (operational) - The port is operating normally. • Blo (blocked) - The port is blocked because the adjacent port is not configured with link aggregation or because it is not able to join a LAG group. An LACP port is blocked until it becomes part of a LAG. Also, an LACP is blocked if its state becomes "default". To unblock the port and bring it to an operational state, enable link aggregation on the adjacent port and ensure that the ports have the same key. • Frc (force-up)- The port is in "force-up" mode. If you have configured the force-up ethernet command on the member port of a dynamic LAG, the port goes into "force-up" mode and is logically operational when the dynamic LAG is not operating.

Displaying information about LAG interface

You can view the details of the LAG interface including counters using the **show interfaces lag** command.

You can also view the details of a LAG by specifying the LAG name or LAG ID. If the specified LAG name or LAG ID is not available, a warning message is displayed.

```
device# show interfaces lag 1
Total number of LAGs: 1
Total number of deployed LAGs: 1
Total number of trunks created:1 (123 available)
LACP System Priority / ID: 1 / 748e.f8b1.66e0
LACP Long timeout: 120, default: 120
LACP Short timeout: 3, default: 3

==== LAG "test" ID 1 (dynamic Deployed) ====
LAG Configuration:
  Ports: e 1/1/1 to 1/1/2
  Port Count: 2
  Primary Port: 1/1/1
  Trunk Type: hash-based
  LACP Key: 20001
Deployment: HW Trunk ID 1
Port      Link     State   Dupl Speed Trunk Tag Pvid Pri MAC          Name
1/1/1     Up       Forward Full  1G    1      No   1    0
748e.f8b1.66e0
1/1/2     Up       Forward Full  1G    1      No   1    0    748e.f8b1.66e0
Port      [Sys P] [Port P] [Key] [Act][Tio][Agg][Syn][Col][Dis][Def][Exp][Ope]
1/1/1     1        1    20001  Yes  L  Agg  Syn  Col  Dis  No  No  Ope
1/1/2     1        1    20001  Yes  L  Agg  Syn  Col  Dis  No  No  Ope

  Partner Info and PDU Statistics
  Port      Partner      Partner      LACP      LACP
           System MAC     Key        Rx Count  Tx Count
  1/1/1     748e.f8b1.6020 20001      19        18
  1/1/2     748e.f8b1.6020 20001      18        19

  LAG test Counters:
    InOctets      91162279156      OutOctets      91155682034
    InPkts        171383016       OutPkts        171371929
    InBroadcastPkts 75449406      OutBroadcastPkts 75438497
    InMulticastPkts 10560       OutMulticastPkts 10553
    InUnicastPkts  95923050      OutUnicastPkts 95922879
    InBadPkts      0
    InFragments     0
    InDiscards      0      OutErrors      0
    CRC            0      Collisions      0
    InErrors        0      LateCollisions 0
    InGiantPkts     0
    InShortPkts     0
    InJabber        0
    InFlowCtrlPkts 0      OutFlowCtrlPkts 0
    InBitsPerSec   1931301848      OutBitsPerSec 1931301848
    InPktsPerSec   453126       OutPktsPerSec 453126
    InUtilization  100.00%      OutUtilization 100.00%
```

Enabling LAG hardware failover

LAG hardware failover reduces the time of packet loss if a LAG member is down, with minimal software intervention, using loopback on the down port. LAG hardware failover is disabled by default and is supported only on Brocade ICX 7750 devices.

NOTE

LAG should be undeployed to configure the **failover all** command.

Enter the **failover all** command in the LAG configuration mode to enable LAG hardware failover. **failover next** enables failover on the next port in LAG. .

In this example, the LAG **failover all** command is enabled on all ports.

```
device(config)#lag one dynamic
device(config-lag-one)#failover all
```

Preboot eXecution Environment boot support

The Preboot eXecution Environment (PXE), also known as Pre-Execution Environment, is an environment to boot devices using a network interface independent of data storage devices (such as hard disks) or installed operating systems. Consider an environment in which a PXE-capable host forms a dynamic LAG with a FastIron device. After the host successfully boots and runs an operating system, the LACP initiates negotiation to form the dynamic LAG for network access. To boot from the network, the host must be able to connect with the FastIron device initially without a dynamic LAG. To enable this, you can configure PXE boot support on one of the member ports of a dynamic LAG. This ensures that the port is logically operational as soon as you connect this port to the host, even when the dynamic LAG is not operating. At this stage, the port is in "force-up" mode and the **show lag** command shows the operational status "Ope" of this port as "Frc". Once the host successfully boots from the network using this port, the dynamic LAG can form to connect the host to the network with the LAG link. Even if the dynamic LAG fails later, this port is brought back to "force-up" mode and remains logically operational.

Enabling PXE boot support on a port

- The port should be an edge port on which you have not configured protocols such as STP, MRP, and UDLD.
- The dynamic LAG should be in an undeployed state.

NOTE

You can enable PXE boot support on only one member port of a dynamic LAG.

Run the **force-up ethernet** command in dynamic LAG configuration mode.

The following example shows PXE boot support enabled on member port 3/1/1 of a dynamic LAG R4-dyn.

```
Brocade(config)# lag R4-dyn
Brocade(config-lag-R4-dyn)# force-up ethernet 3/1/1
```


Multi-Chassis Trunking

• Multi-Chassis Trunking Overview.....	159
• Layer 2 behavior with MCT.....	174
• Displaying MCT information.....	184
• MCT configuration examples	189

Multi-Chassis Trunking Overview

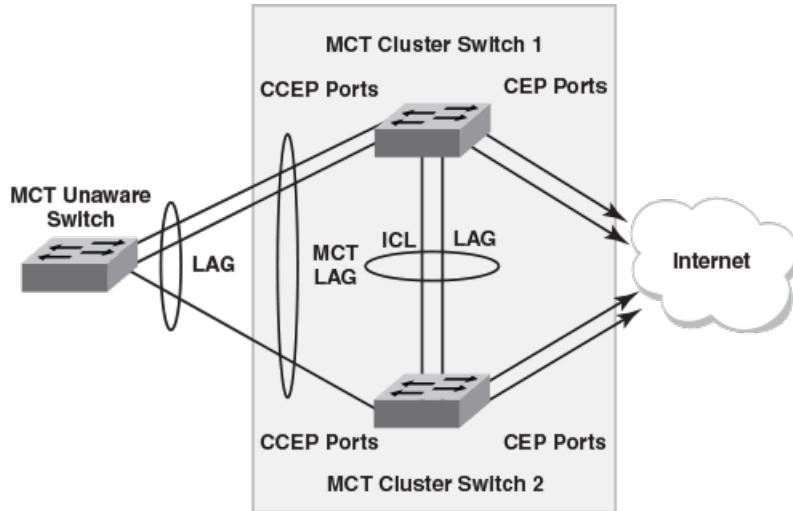
Multi-Chassis Trunking (MCT) is an alternative to spanning tree protocols. Spanning tree is a technology that protects the network against loops by blocking necessary ports, and having the network span to relearn topologies when one link fails in a network. MCT is a technology that allows two MCT-supporting switches to cluster together and appear as a single logical device. Trunking is a technology that allows multiple links of a device to appear as one logical link. The combination of MCT and trunking allows for creating a resilient network topology that utilizes all links in the network, creating an ideal network topology for latency sensitive applications.

Standard static or dynamic LACP trunks provide link-level redundancy and increased capacity. However, trunks do not provide device-level redundancy. If the device to which the trunk is attached fails, the entire trunk loses network connectivity. Two devices are needed for network resiliency with trunked links to both devices. With spanning tree, one of these trunks would be blocked from use until the failure of the other trunk is detected, taking from 1 to 30 seconds potentially adding latency and jitter, not only on the affected devices locally, but throughout the span topology. With MCT, member links of the trunk are split and connected to two clustered MCT-supporting switches. MCT has integrated loop detections, which allows all links to be active. If a failure is detected, traffic is dynamically allocated across the remaining links. The failure detection and allocation of traffic occur in sub-second time, without impact on the rest of the network.

MCT inherits all of the benefits of a trunk group and allows multiple physical links to act as a single logical link. The resulting available bandwidth is an aggregate of all the links in the group. Traffic is shared across the links in the group using dynamic flow-based load balancing, and traffic is moved to a remaining link group in sub-seconds if a failure occurs on one of the links. MCT eliminates the single point of failure that exists at the device level when all links of a trunk terminate on the same device without the overhead associated with spanning tree. MCT diverts a subset of the links to a second device to provide redundancy and sub-second fault detection at the device level.

How MCT works

The following table shows a basic MCT configuration. The MCT originates at a single MCT-unaware server or switch and terminates at two MCT-aware devices.

FIGURE 25 How MCT works

The MCT process involves the following processes:

- Sub-second failover occurs if a link, module, switch fabric, control plane, or device fails.
- Sub-second failover operates at the physical level.
- Layer 2 and Layer 3 forwarding (when using fast path forwarding) is done at the first hop regardless of VRRP-E state.
- Load balancing is flow based (it does not involve VLANs sharing across network links).
- Resiliency is supported regardless of the traffic type (Layer 3, Layer 2, or non-IP legacy protocols).
- Interaction with Metro Ring Protocol (MRP) builds larger resilient Layer 2 domains.
- Device-level redundancy is provided in addition to link and modular redundancy.
- Traffic received from an ICL port is not forwarded to the Cluster Client Edge Ports (CCEPs) if the MCT peer device has the ability to reach the same cluster client.
- Traffic received from non-ICL ports is forwarded the same way as non-MCT devices.
- Known unicast traffic received on Cluster Edge Ports (CEP) or ICL ports is forwarded to the destination port.
- For unknown unicast, multicast, and broadcast traffic received on ICL ports, the forwarding behavior depends on the peer MCT device's ability to reach the same client.
- Unknown unicast, multicast, and broadcast traffic received from CCEP is forwarded as usual, by default, flooding the entire VLAN.
- The cluster ID must be unique when there are multiple clusters interconnected in a topology. For example, in a cascaded Stage 2 MCT cluster, the cluster ID on a stage 1 pair of switches should be different from the cluster ID on a stage 2 pair of switches.

MCT terminology

- Cluster Client Edge Port (CCEP): A physical port or trunk group interface on an MCT cluster device that is connected to client devices.
- Cluster Edge Port (CEP): A port on an MCT cluster device that belongs to the MCT VLAN and connects to an upstream core switch/router but is neither a CCEP nor an ICL.
- Cluster Communication Protocol (CCP): A Brocade proprietary protocol that provides reliable, point-to-point transport to synchronize information between MCT cluster devices. It provides the default MCT control path between the two peer devices. CCP comprises two main components: CCP peer management and CCP client management. CCP peer management deals with establishing and

maintaining a TCP transport session between peers, while CCP client management provides event-based, reliable packet transport to CCP peers.

- Inter-Chassis Link (ICL): A single-port or multi-port 1 GbE, 10 GbE, or 40 GbE LAG between the two MCT cluster devices. It provides the control path for CCP for the cluster and also serves as the data path between the two devices.
- MCT cluster: A pair of devices (switches) that is clustered together using MCT to appear as a single logical device. The devices are connected as peers through an Inter-Chassis Link (ICL).
- MCT cluster client: A device that connects with MCT cluster devices through static or dynamic trunks. It can be a switch or an endpoint server host in the single-level MCT topology or another pair of MCT devices in a multi-tier MCT topology.
- MCT cluster device: One of the two devices in an MCT cluster.
- MCT peer device: From the perspective of an MCT cluster device, the other device in the MCT cluster.
- MCT VLANs: VLANs on which MCT cluster clients are operating. Any VLAN that has an ICL port is an MCT VLAN, even if it does not have any clients.
 - MCT keep-alive VLAN: The VLAN that provides a backup control path if the ICL goes down.
 - MCT session VLANs: The VLAN used by the MCT cluster for control operations. CCP protocol runs over this VLAN. The interface can be a single link or a trunk group port. If it is a trunk group port, it should be the primary port of the trunk group. The MCT session VLAN subnet is not distributed in routing protocols using redistribute commands.
- RBridgeID: RBridgeID is a value assigned to MCT cluster devices and clients that uniquely identifies them and helps associate the source MAC address with an MCT device.

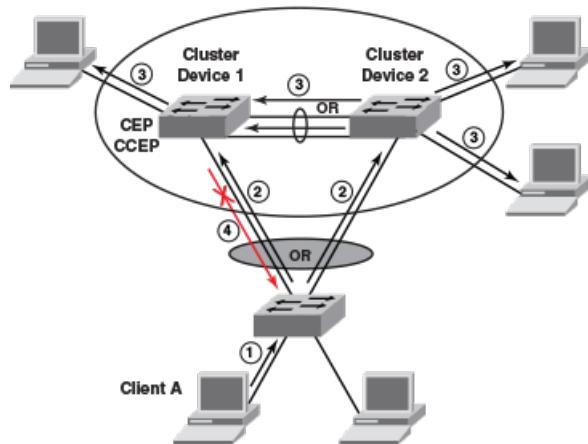
MCT data flow

MCT can be deployed in a single-level configuration that includes two MCT cluster devices or in a cascading configuration, where a pair of MCT cluster devices operate as switches, and another pair of cluster devices operates as routers. Refer to [Single-level MCT example](#) on page 189 for a single-level illustration and configuration example, and [Two-level MCT example](#) on page 192 for a two-level or cascading configuration example.

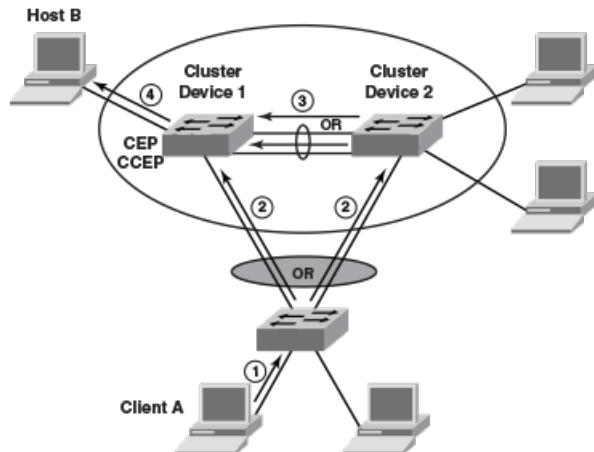
Basic MCT data flow works as follows.

Broadcast, unknown unicast, and multicast (BUM) traffic from a client through a CCEP

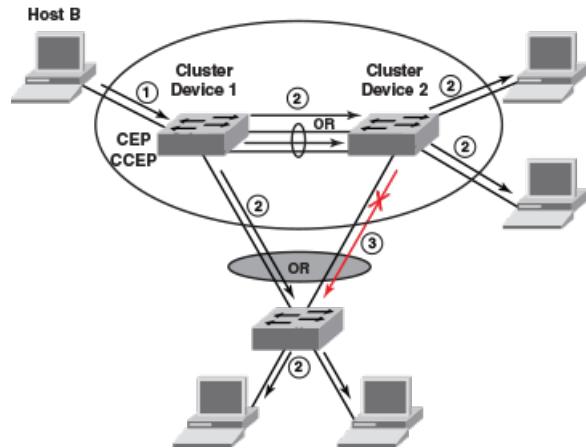
1. Traffic originates at the client.
2. Because the link between the client switch and the MCT cluster is a trunk, the traffic travels over one physical link. In the example shown in the following figure, the traffic travels over the link toward cluster device 2. The traffic enters the MCT cluster through the CCEP of cluster device 2.
3. The traffic is sent to any local CEPs and CCEPs. It passes to the peer cluster device over the ICL link, where it is sent to the peer device's local CEPs.
4. Traffic does not pass back down to the client through the CCEP.

FIGURE 26 MCT data flow - BUM traffic from CCEP***Unicast traffic from a client through a CCEP to a CEP***

1. Traffic originates at the client.
2. Because the link between the client switch and the MCT cluster is a trunk, the traffic travels over one physical link. As shown in the following figure, the traffic travels over the link toward cluster device 2. The traffic enters the MCT cluster through the CCEP of cluster device 2.
3. Depending on the destination, the traffic may pass over the ICL link to the other cluster device. In the following figure, the destination is on cluster device 1, so the traffic is forwarded out to the ICL port.
4. The traffic passes out to the destination.

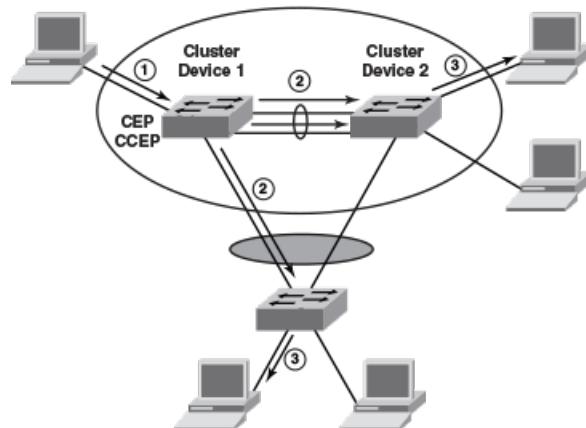
FIGURE 27 MCT data flow - unicast traffic from CCEP***Broadcast, unknown unicast, and multicast (BUM) traffic from a client through a CEP***

1. Traffic originates at the client and enters one of the MCT cluster devices through a CEP.
2. As shown in the following figure, the traffic is sent to the peer cluster device through the ICL link and is also sent to any local CCEPs and CEPs. Once traffic is received on the peer cluster device, it will be sent to its local CEPs.
3. Traffic does not pass back down to the client through the CCEP.

FIGURE 28 MCT data flow - BUM traffic from a CEP

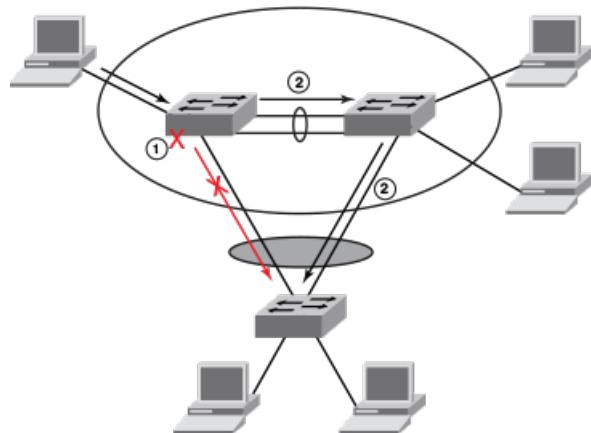
Unicast traffic from a client through a CEP to another CEP or a CCEP

1. Traffic originates at the client and enters one of the cluster devices through the CEP as shown in the following figure.
2. Depending on the destination, the traffic may pass over the ICL link to the other cluster device, or it may be sent to a local CCEP.
3. The traffic passes out to the destination.

FIGURE 29 MCT data flow - unicast traffic from a CEP

Port failure on the cluster device

1. A CCEP on the cluster device that received the unicast or BUM traffic fails.
2. As shown in the following figure, the traffic is automatically redirected to the other MCT cluster device over the ICL and on to its destinations through CCEPs.

FIGURE 30 MCT data flow with port failure

MCT and VLANs

MCT relies on the following VLAN types:

- Session VLAN: Provides the control channel for CCP. Brocade recommends keeping only ICL ports in the session VLAN. A virtual interface must be configured on the session VLAN for the router image.
- Keep-alive VLAN: Provides a backup control path if the ICL goes down (optional, but strongly recommended).
- MCT VLAN: Serves the customer data traffic. An ICL must belong to every MCT VLAN to provide a data path between two cluster devices. When an ICL is added to a VLAN, it becomes an MCT VLAN.

Cluster client automatic configuration

Client configuration includes setting the client name, client RBridgeID (unique identification for each client), client interface (CCEP), and deployment settings on both MCT cluster devices. With up to 150 clients per cluster, manual configuration can take a considerable amount of time.

Cluster client automatic configuration saves the time that would be required to complete the entire configuration manually.

The following limitations apply to cluster client automatic configuration:

- Cluster client automatic configuration is designed for generating new clients, not for updating an existing client.
- A single client span across multiple devices is not supported (cascading MCT). For example, the configuration of cascading MCT through cluster client automatic configuration is not supported.
- Multiple clients on the same device are not supported.
- LACP client interface auto-detection is supported only for devices running release 7.4 software and later on FastIron platforms.
- RBridgeID collision: When hash collisions occur, cluster client automatic configuration reports errors, and manual intervention is required.

For cluster client automatic configuration to work, the following prerequisites are required on the cluster side:

- The cluster must be configured on both MCT cluster devices.
- An MCT VLAN must be configured on both MCT cluster devices.
- The trunk group configuration must be removed from the client interfaces.
- The client interfaces must be up and operational.
- The cluster ID must be unique when there are multiple clusters interconnected in a topology. For example, in a cascaded Stage 2 MCT cluster, the cluster ID on a stage 1 pair of switch should be different from the cluster ID on a stage 2 pair of switches.

The following prerequisites are required on the client side:

- VLAN and trunk group configuration must be completed.
- Link Level Discovery Protocol (LLDP) must be enabled.

Refer to [Setting up cluster client automatic configuration](#) on page 170 for detailed instructions on the cluster client automatic configuration process.

MCT feature interaction

The following FastIron features are supported with MCT. All security features are locally significant and are not synchronized across an MCT cluster.

- LACP on the CCEP.
- VRRP on the CCEP.
- MRP and MRP II, with the restriction that the ICL port cannot be the secondary port of the MRP ring.
- Flooding features (such as VLAN CPU protection and multicast flooding) on MCT VLANs.
- Unidirectional Link Detection (UDLD) as independent boxes (configured independently).
- ARP as independent boxes (configured independently).
- STP and RSTP.
- Ingress ACLs on all MCT ports. Egress ACLs are supported only on MCT CEPs or ICL ports. Egress ACLs are not supported on MCT CCEPs.
- QoS and MAC filters and profiles with the same configuration on both cluster devices.
- IPv4 ACLs and rate limits. If the rules are applied on the CCEPs, the same rules must be applied to the CCEP ports on both cluster devices.
- Layer 3 Routing. VE with IP address assignment is supported on CCEPs for VRRP.
- Static multi-port MAC.
- Port MAC security, multi-port authentication, and 802.1X, only on CEPs.
- Static MAC address configuration. Static MAC addresses are programmed on both local and remote peers as static entries.
- DAI and DHCP snooping for clients connected through CEPs. They must be configured independently on both cluster devices.
 - If the trusted ports are off the CCEP, the **arp inspection trust** or **dhcp snoop trust** command must be used on the CCEPs and ICL ports.
 - DHCP and ARP entries are created on both MCT cluster devices if the flow traverses both the CCEP and ICL.
- Hitless failover. If the failover operation is performed with a cluster configuration, the TCP session is reestablished. The MAC addresses from the cluster peer devices are revalidated and programmed accordingly.
- Hitless upgrade. If the upgrade operation is performed with a cluster configuration, the TCP session is reestablished. The MAC addresses from the cluster peer devices are revalidated and programmed accordingly.

The following FastIron features are not supported with MCT:

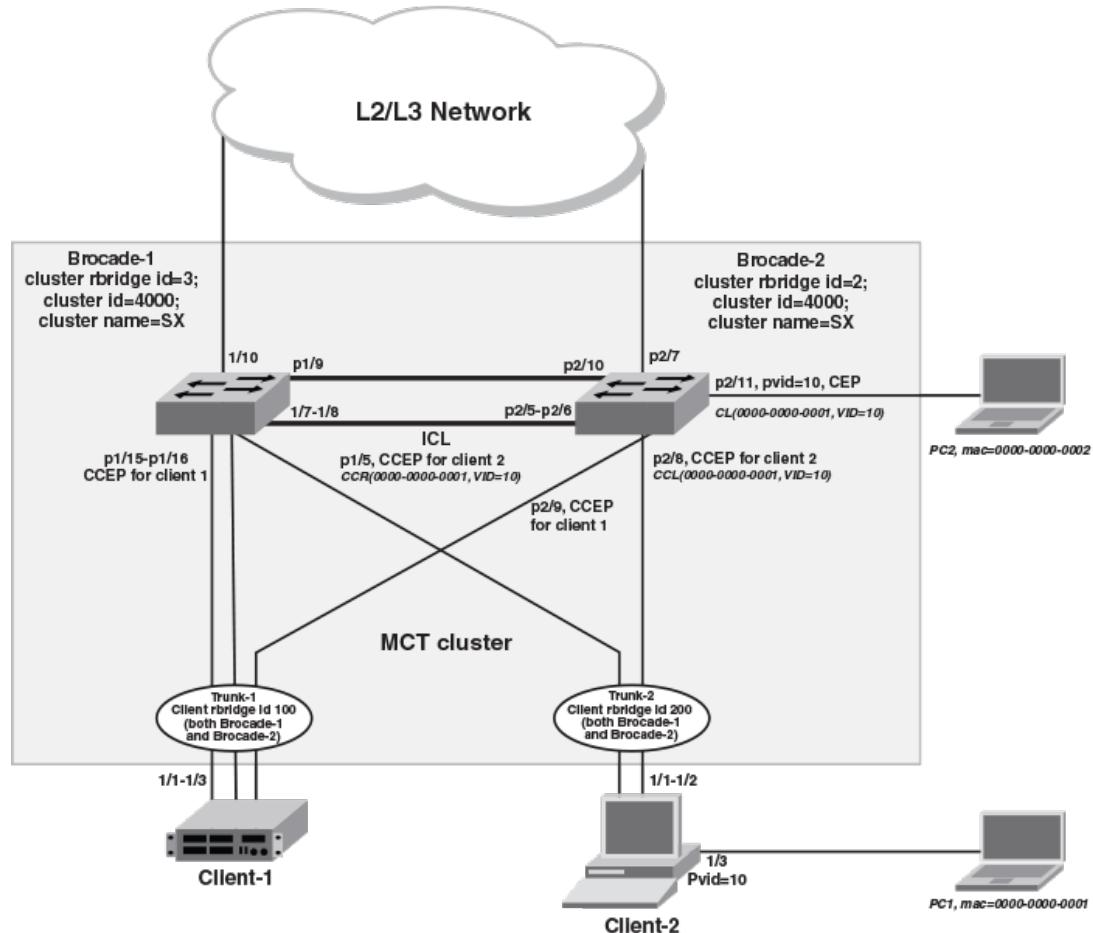
- LACP on ICL.
- MSTP, VSRP, RIP.

- Layer 3 multicast features. For example, PIM, MSDP, anycast RP, embedded RP etc.
- IPv6, VRRP-E (IPv6), and VRRPV3.
- GRE on the ICL VE interfaces.
- DAI on the CCEPs.
- Host security features (port MAC security, multi-port authentication, 802.1X, DAI, DHCP snooping) on CCEPs.
- Multi-port ARP on ICL or CCEPs.
- Web authentication on MCT VLANs.

Basic MCT configuration

This section describes how to set up a basic MCT configuration. The following figure shows a basic MCT topology, which applies to Layer 2 and Layer 3. MCT can also be supported with VRRP or VRRP-E.

FIGURE 31 Basic MCT configuration



MCT configuration considerations

- Configuring flow-based MAC address learning and MCT on the same device is not supported.
- When running STP, the STP state should be the same on both cluster devices. For additional information on running STP with MCT, refer to "STP/RSTP" under MCT Layer 2 protocols and to related configuration examples.
- One ICL can be configured per device, and a device can be in only one cluster.
- The software version in both cluster devices must be exactly the same for the cluster to function.
- An ICL port should not be an untagged member of any VLAN.
- It is recommended that you set up ICL as a static LAG with at least two ports. This provides port-level redundancy and higher bandwidth for cluster communication.
- ICL ports must be part of MCT VLANs and session VLANs.
- An ICL cannot be a regular port link or an LACP trunk. It must be a single or multiple ports static LAG.
- MAC learning is disabled on ICL ports for all VLANs.
- MDUP synchronizes all MAC entries for VLANs served by an ICL link.
- In any MCT configuration, there are two different cluster-related IDs, the Cluster ID and the Cluster RBridge ID. The Cluster ID uniquely identifies a cluster. All cluster devices in the same MCT cluster have the same Cluster ID. The Cluster RBridge ID uniquely identifies a cluster device within the cluster. To avoid conflicts, ensure that the Cluster ID and the Cluster RBridge ID are unique within an MCT configuration and cannot be confused with each other.
- The cluster ID should be the same on both cluster devices.
- The cluster RBridgeID should not conflict with any client RBridgeID or with the peer RBridgeID.
- The client RBridgeID is unique and should be the same on cluster devices.
- Brocade recommends keeping only ICL ports in the session VLAN during operation.
- MCT can support up to 16 members per trunk group, depending on the software version and Switch type.
- An ICL interface cannot be configured as the CCEP in any client.
- BPDU guard and root guard configuration should be identical on both cluster devices.
- Because Egress PCL is configured on CCEPs, egress ACL cannot be configured on them. All types of ingress ACLs, DoS attack prevention, and so on can still be configured on those ports.
- Brocade recommends that you configure a keep-alive VLAN as a separate link (not ICL). The keep-alive VLAN provides a backup control path when CCP goes down.
- 48GC ports should not be used as MCT trunks or CCEP ports.

Differences in configuring MCT for the switch and router image

There are some differences in the MCT configuration for the switch image versus the router image:

- On a switch image, STP is by default enabled for all the VLANs; however, for MCT, Layer 2 protocols such as STP and RSTP should not be enabled on the session VLAN. Therefore, STP must be disabled explicitly for the session VLAN. STP is automatically disabled in the router image.
- Virtual Ethernet (VE) cannot be configured on a session VLAN in a switch image, but an IP address is needed for the cluster devices to communicate via CCP. Therefore, in a switch image, the configured management IP address is used to establish communication between the cluster devices.
- The management IP addresses in each of the cluster devices should be configured in the same subnet. If the IP addresses are in different subnets, ARP does not resolve the addresses, and MCT may not work. ARP for the peer cluster devices is always learned on the ICL port or trunk, so any management traffic between the two devices always goes through the ICL ports.

NOTE

CLI may vary somewhat among different Brocade platforms. These variations are not documented in the configuration examples provided in this chapter.

Configuring MCT

This section provides basic configuration steps, which should be completed in the specified order.

[Step 1: Configure LAG](#) on page 168

[Step 2: Configure the MCT VLAN, MCT session VLAN, and recommended MCT keep-alive VLAN](#) on page 168

[Step 3: Configure the cluster](#) on page 169

[Step 4: Configure clients](#) on page 170

After completing these steps, you can verify the configuration by running the **show cluster** command. Refer to [Displaying peer and client states](#) on page 184.

Step 1: Configure LAG

You can configure a static or dynamic LAG. Static LAG groups are manually configured aggregate links containing multiple ports. Dynamic LAGs use Link Aggregation Control Protocol (LACP) to maintain aggregate links over multiple port. LACP PDUs are exchanged between ports on each device to determine if the connection is still active. The LAG then shuts down any port whose connection is no longer active.

NOTE

ICL LAGs only support static trunks.

Syntax: [no] lag *lag-name* [{ static | dynamic } [id *number*]]

To configure an ICL static LAG, enter the following commands.

```
device-1(config)# lag MCT lag1 static id 2
Brocade-1(config-lag-MCT_lag1)# ports ethernet 1/7 to 1/8
Brocade-1(config-lag-MCT_lag1)# primary-port 1/7
Brocade-1(config-lag-MCT_lag1)# deploy
```

Step 2: Configure the MCT VLAN, MCT session VLAN, and recommended MCT keep-alive VLAN

To create the MCT session VLAN and recommended MCT keep-alive VLAN for Brocade-1 in the topology of [Figure 31](#) on page 166, enter the following commands.

```
device-1(config)# vlan 3001 name MCT-keep-alive
device-1(config-vlan-3001)# tagged ethernet 1/9
device-1(config-vlan-3001)# exit
device-1(config)# vlan 3000 name Session-VLAN
device-1(config-vlan-3000)# tagged ether 1/7 to 1/8
device-1(config-vlan-3000)# no spanning-tree
```

For routers, add the following commands.

```
device-1(config-vlan-3000)# router-interface ve 3000
device-1(config)# interface ve 3000
device-1(config-vif-3000)# ip address 10.1.1.3/24
```

For switches, add the following commands.

```
device-1(config)# ip address 10.1.1.3/24
```

To create a session VLAN and keep-alive VLAN for device-2, enter the following commands.

```
device-2(config)# vlan 3001 name MCT-keep-alive
device-2(config-vlan-3001)# tagged ethernet 2/10
device-2(config-vlan-3001)# exit
device-2(config)# vlan 3000 name Session-VLAN
device-2(config-vlan-3000)# tagged ether 2/5 to 2/6
device-2(config-vlan-3000)# no spanning-tree
```

For routers, add the following commands.

```
device-2(config-vlan-3000)# router-interface ve 3000
device-2(config)#interface ve 3000
device-2(config-vif-3000)# ip address 10.1.1.2/24
```

For switches, add the following commands.

```
device-2(config)# ip address 10.1.1.2/24
```

To implicitly configure the MCT VLAN and add the ICL as a tagged member of the VLAN, enter the following commands.

```
device-1(config)# vlan 1000 name MCT-VLAN-example
device-1(config-vlan-1000)# tagged ether 1/15 to 1/16 e 1/7 to 1/8
```

Step 3: Configure the cluster

Cluster local configuration uses the cluster ID and RBridgeID for the local switch or router.

Syntax: [no] **cluster** [*cluster-name*] *cluster-id*

Syntax: [no] **rbridge-id** *id*

Configuration of the peer device involves the peer's IP address, RBridgeID, and ICL specification. The *cluster-name* variable is optional; the device auto-generates the cluster name as CLUSTER-X when only the cluster ID is specified. The *cluster-id* variable must be the same on both cluster devices.

Syntax: [no] **peer** *peer-ip* **rbridge-id** *peer-rbridge* **icl** *map-icl*

The RBridgeID must be different from the cluster RBridge and any other client in the cluster. The MCT member VLAN is defined as any VLAN of which the ICL is a member.

To configure Brocade-1 for the cluster in the topology of [Figure 31](#) on page 166, enter the following commands.

```
device-1(config)#cluster SX 4000
device-1(config-cluster-SX)#rbridge-id 3
device-1(config-cluster-SX)#session-vlan 3000
device-1(config-cluster-SX)#keep-alive-vlan 3001
device-1(config-cluster-SX)#icl SX-MCT ethernet 1/7
device-1(config-cluster-SX)#peer 10.1.1.2 rbridge-id 2 icl SX-MCT
device-1(config-cluster-SX)#deploy
```

To configure Brocade-2 for the cluster in the topology of [Figure 31](#) on page 166, enter the following commands.

```
device-2(config)# cluster SX 4000
device-2(config-cluster-SX)#rbridge-id 2
device-2(config-cluster-SX)#session-vlan 3000
device-2(config-cluster-SX)#keep-alive-vlan 3001
device-2(config-cluster-SX)#icl SX-MCT ethernet 2/5
device-2(config-cluster-SX)#peer 10.1.1.3 rbridge-id 3 icl SX-MCT
device-2(config-cluster-SX)#deploy
```

Step 4: Configure clients

This section describes how to configure clients manually. For instructions on automatic client configuration, refer to [Setting up cluster client automatic configuration](#) on page 170.

Client configuration requires the client name, RBridgeID, and CCEP. In the network shown in the [Figure 31](#) on page 166, Client-1 has a three-port LACP trunk (1/1-1/3), while Client-2 has a two-port static trunk (1/1-1/2) towards the MCT cluster.

The client name can be different on the different cluster devices. To configure the client name, enter the following command.

Syntax: [no] client *client-name*

The client RBridgeID must be identical on both of the cluster devices. To configure the client RBridgeID, use the following command.

Syntax: [no] rbridge-id *id*

To configure the physical port or static trunk as the client CCEP, use the following command.

Syntax: [no] client-interface ethernet *slot/port*

To configure Client-2 on Brocade-1 in the topology of [Figure 31](#) on page 166, enter the following command.

```
device-1(config-cluster-SX)# client client-2
device-1(config-cluster-SX-client-1)#rbridge-id 200
device-1(config-cluster-SX-client-1)#client-interface ether 1/5
device-1(config-cluster-SX-client-1)#deploy
```

To configure Client-2 on Brocade-2 in the topology of [Figure 31](#) on page 166, enter the following command.

```
device-2(config-cluster-SX)# client client-2
device-2(config-cluster-SX-client-2)#rbridge-id 200
device-2(config-cluster-SX-client-2)#client-interface ether 2/8
device-2(config-cluster-SX-client-2)#deploy
```

Setting up cluster client automatic configuration

Complete the following steps to configure cluster client automatic configuration.

1. Enable the client auto-detect ports on both MCT devices.

```
device-1(config-cluster-SX)# client-auto-detect ethernet 1/15-1/16
```

In the port list, specify all the CCEPs for all potential clients.

2. Start the client auto-detect process on both cluster devices.

```
device-1(config-cluster-SX)# client-auto-detect start
```

Within one minute, the system reports information and errors (if there are mismatches such as an LACP configuration mismatch). You can fix the mismatch while the process is running.

3. Check and fix the automatically detected clients.

```
device-1(config-cluster-SX)# show cluster cluster-SX client-auto-detect
cluster cluster-SX 4000
  rbridge-id 3
  session-vlan 3000
  icl SX-MCT ethernet 1/7
  peer 10.1.1.2 rbridge-id 2 icl SX-MCT
  client-auto-config ethe 1/15 to 1/16 ethe 8/5 ethe 8/7 eth 8/9
  client-auto-config start
  deploy
  client AUTO-FCX624-Router002438769e00
  rbridge-id 3593
```

```
client-interface ethe 1/15
!
```

NOTE

At this point, the client configuration does not appear in the running configuration and cannot be modified. Static trunk and LACP configuration are not effective yet.

4. Configure automatically detected clients into the running configuration.

```
device-1(config-cluster-SX) # client-auto-detect config
```

All automatically configured client information is now published into the running configuration, and the static trunk configuration is generated, created, and deployed. LACP is started. By default, clients are in the non-deployed state and the CCEPs are put into the disabled state. Ports that are successfully programmed as CCEP are removed from the autoconfig-enabled port list. If the port list is empty, which means all ports are configured into clients successfully, the automatic configuration process stops. The original LLDP configuration is restored. Otherwise, the automatic configuration process continues only on the ports still left in the list.

Other cluster client automatic configuration commands

You can use the following commands as an alternative to the step-by-step procedure in [Cluster client automatic configuration](#) on page 164.

Use the following command to enable or disable cluster client automatic configuration on a range of ports.

Syntax: [no] client-auto-detect Ethernet x [to y]

Use the following command as an alternative to **client-auto-detect config**. This command also configures automatically detected clients into the running configuration and deploys all of the automatically detected clients.

Syntax: client-auto-detect config deploy-all

Use the following command to start the cluster client automatic configuration. Within one minute of the time that each client is discovered, the client is automatically configured and deployed into the running configuration.

Make sure that the network connection and configuration are in place before using this command.

Syntax: client-auto-detect start [config-deploy-all]

Use the following command to stop the current running cluster client automatic configuration process. All auto-detected but unconfigured clients will be cleared.

Syntax: client-auto-detect stop

MCT failover scenarios

The following scenarios describe what happens if specific elements in the MCT configuration fail.

- Client interface on one of the MCT cluster devices goes down.
 - Traffic switches to the other cluster device with minimal traffic loss.
- MCT cluster device goes down.
 - When an MCT cluster device goes down (for example, due to a power failure), the traffic fails over to the other MCT cluster device.
- Hitless failover occurs.

- The MCT CCEPs stay up during hitless switchover, failover, or upgrade. Link protocols such as UDLD and LACP on CCEPs do not flap. Traffic disruption is minimal (sub-second). The MCT CCP connection flaps once, and MAC is re-synced between the peer devices.
- The CCP goes down and comes back up again once the hitless failover is completed.
- ICL interface or CCP goes down (keep-alive is configured).
 - If a keep-alive VLAN is used, the devices in the cluster can communicate even if the ICL goes down. If the peer device is reachable over the keep-alive VLAN, the MCT peers perform the master/slave negotiation per client. After negotiation, the slave shuts down its client ports, and the master client ports continue to forward the traffic.
 - The master/slave negotiation is performed per MCT client on the basis of RBridgeliD and client Local or Remote accessibility. If the client is reachable from both MCT devices, the lower RBridgeliD becomes the master. If the client can be accessed only from one of the MCT devices, the cluster device on which it is reachable becomes the master.
 - If the peer device cannot be reached over the keep-alive VLAN, then both cluster devices keep forwarding.

NOTE

Brocade recommends using keep-alive VLANs with the MCT configurations. This provides alternative access if the ICL interface goes down. However, a keep-alive VLAN should not be configured when **bpdu-flood-enable** is configured. Refer to [BPDU forwarding](#) on page 179.

- ICL interface or CCP goes down (keep-alive is not configured).
 - When the keep-alive VLAN is not configured, both cluster devices keep forwarding. Use the **client-isolation strict** command to disable the client interface as soon as the ICL link goes down to completely isolate the client.
- Double failures occur (for example, the ICL goes down and the client interface goes down on one of the MCT cluster devices).
 - Multiple failures could cause traffic to drop, even if there is a physical path available.

NOTE

The keep-alive link is supported on the ICX 7750 in an Sstp or MST environment, even though ICX 7750 has the BPDU flood-enable feature built in. This is because the BPDU flood-enable in the ICX 7750 is hardware enabled.

NOTE

With the ICX 6650 and SX chassis, if the environment uses Sstp or MST in an MCT configuration, then the keep-alive link is not supported with the BPDU flood enable command. This is because the command is software-based in the ICX 6650 and SX chassis.

Cluster failover mode

The following failover modes can be configured with MCT:

- Fast-failover (default) - As soon as the ICL interface goes down, the CCP goes down. All the remote MAC addresses are flushed.
- Slow-failover - Even if the ICL interface goes down, the CCP waits for the hold-time before taking the CCP down. Remote MAC addresses are flushed only when the CCP is down.

To disable the fast-failover mode, enter a command such as the following.

```
device-1(config-cluster-SX) # peer 10.1.1.3 disable-fast-failover
```

Syntax: [no] peer peer-ip disable-fast-failover

Client isolation mode

NOTE

You must create the same isolation mode on both cluster devices. The CLI will allow modification of the client isolation mode on MCT cluster devices even when the cluster is deployed.

MCT cluster devices can operate in two modes. Both peer devices should be configured in the same mode.

Loose mode (default): When the CCP goes down, the peer device performs the master/slave negotiation. After negotiation, the slave shuts down its peer ports, but the master peer ports continue to forward traffic if a keep-alive VLAN is configured.

If a keep-alive VLAN is not configured, both peer devices become masters, and both of the client ports stay up.

```
device-1(config-cluster-SX) # client-isolation loose
```

Strict mode: When the CCP goes down, the interfaces on both the cluster devices are administratively shut down. In this mode, the client is completely isolated from the network if the CCP is not operational.

```
device-1(config-cluster-SX) # client-isolation strict
```

Syntax: [no] client-isolation strict

Shutting down all client interfaces

Use the **client-interfaces shutdown** command when performing a hitless upgrade operation. This command can be used to shut down all the local client interfaces in the cluster, which results in failover of traffic to the peer device.

```
device-1(config-cluster-SX) # client-interfaces shutdown
```

Syntax: [no] client-interfaces shutdown

Using the keep-alive VLAN

CCRR messages are used to exchange information between peer devices. When the CCP is up, CCRR messages are sent over the CCP. When the CCP client cannot be reached or the ICL is down, you can use the **keep-alive-vlan** command under the cluster context so CCRR messages are periodically sent over the keep-alive VLAN. Only one VLAN can be configured as a keep-alive VLAN. The keep-alive VLAN cannot be a member VLAN of the MCT, and this VLAN can be tagged or untagged.

NOTE

Keep-alive VLAN configuration is not allowed when the client isolation mode is strict. When a keep-alive VLAN is configured, client isolation mode cannot be configured as strict.

```
device-1(config-cluster-SX) # keep-alive-vlan 10
```

Syntax: [no] keep-alive-vlan *vlan-id*

The *vlan_id* variable specifies the VLAN range. Possible values are from 1 to 4089.

When the CCP is down, the following results occur.

- If the keep-alive VLAN is configured, CCRR messages are sent every second over that VLAN.
- When CCP is down and a keep-alive VLAN is configured, master/slave selection is based on the following criteria:
 - If one device's CCEPs are up and the peer's CCEPs are down, the peer with the local CCEPs down becomes the slave.
 - Otherwise, the device with the higher RBridgeID becomes the slave.
- If no packets are received from the peer device for a period of three seconds, the peer is considered down.
- If a keep-alive VLAN is not configured and both the peer devices are up, both peers keep forwarding traffic independently.

Setting keep-alive timers and hold-time

To specify the keep-alive timers and hold time for the peer devices, enter a command such as the following.

```
device-1(config-cluster-SX))# peer 10.1.1.3 timers keep-alive 40 hold-time 120
```

Syntax: [no] peer *peer-ip* timers keep-alive *keep-alive-time* *hold-time* *hold-time*

The *peer-ip* parameter should be in the same subnet as the cluster management interface.

The *keep-alive-time* variable can be from 0 to 21845 seconds. The default is 10 seconds.

The *hold-time* variable can be from 3 to 65535 seconds and must be at least 3 times the keep-alive time. The default is 90 seconds.

NOTE

The keep-alive VLAN and keep-alive timers are not related. The keep-alive timer is used by CCP.

Layer 2 behavior with MCT

This section describes the Layer 2 behavior when MCT is configured.

MAC operations

This section describes configuration operations related to MAC addresses.

MAC Database Update

Each MAC address is advertised with a cost. Low-cost MAC addresses are given preference over high-cost addresses. MAC addresses that are learned locally are given the highest priority, or the cost of 0, so that they are always selected as the best MAC address.

If a MAC address moves from a CCEP port to a CEP port, a MAC move message is sent to the peer, and the peer moves the MAC address from its CCEP ports to the ICL links.

If two MAC addresses have the same cost, the address learned from the lower RBridgeID wins and is installed in the FDB.

MAC addresses in MCT VLANs are updated across the cluster using MDUP messages.

Cluster MAC types

Cluster Local MAC (CL): MAC addresses that are learned on the MCT VLAN and on CEPs locally. MAC addresses are synchronized to the cluster peer device and are subject to aging.

Cluster Remote MAC (CR): MAC addresses that are learned via MDUP messages from the peer device (CL on the peer). The MAC addresses are always programmed on the ICL port and do not age. The CR is deleted only when the CL is deleted from the peer. An MDB entry is created for these MAC addresses with a cost of 1 and is associated with the peer RBridgeID.

Cluster Client Local MAC (CCL): MAC addresses that are learned on the MCT VLAN and on CCEPs.

The MAC addresses are synchronized to the cluster peer device and are subject to aging. An MDB entry with a cost of 0 is created for these addresses, and they are associated with the client and cluster RBridgeIDs.

Cluster Client Remote MAC (CCR): MAC addresses that are learned via MDUP message from the peer device (CCL on the peer). The MAC addresses are always programmed on the corresponding CCEP port and do not age. The CCR is deleted only when the CCL is deleted from the peer. An MDB entry with the cost of 1 is created for the MAC addresses, and they are associated with the client and peer RBridgeIDs.

Cluster Multi-Destination Local MAC (CML): A static MAC entry that is configured locally on the MCT VLAN. Any static MAC address configured on MCT VLAN will have the ICL added by default. Consequently, the address automatically becomes a multi-destination MAC entry. The local configuration generates a local MDB. Any CML entry can still have up to 2 associated MDBs, one local and one remote. The remote MDB contains the remote static configuration for the same MAC and VLAN. If the dynamic MAC and static configuration co-exist, the dynamic MAC address is removed, whether it is learned locally or from MDUP. The port list of a CML entry contains an ICL port, the client ports from the client list in the local configuration and the remote configuration (if it exists), and all locally configured CEP ports.

Cluster Multi-Destination Remote MAC (CMR): A static MAC entry that is configured on the MCT VLAN on the peer side and has no associated local configuration. The CMR entry has only the information from the remote MDB. The port list of a CMR entry contains an ICL port and all the client ports from the client list in the remote configuration. When there is a local configuration for the same entry, the CMR is converted to the CML.

MAC aging

Only the local MAC entries are aged on a cluster device. The remote MAC address entries are aged based on explicit MDUP messages only.

The remote MAC addresses learned through MDUP messages are dynamic addresses, but they never age from the FDB.

MAC flush

If the CEP is down, the MAC addresses are flushed, and individual MAC deletion messages are sent to the peer device.

If the CCEP local port is down, the MAC addresses are flushed locally, and individual MAC deletion messages are sent to the peer device.

If the **clear mac** command is given, all the MDB and FDB are rebuilt.

If the **clear mac vlan** command is given, all the local MDB and FDB are rebuilt for the VLAN.

MAC movement happens normally on the local device.

CEP to CCEP MAC movement - MAC movement happens normally on the local device, and it deletes all the other MDBs from the peer to create a new local MDB.

MAC show commands

To display all local MAC address entries for a cluster, use the **show mac cluster** command.

```
device# show mac-address cluster 1000
Total Cluster Enabled(CL+CR+CCL+CCR) MACs: 1
Total Cluster Local(CL) MACs: 1
CCL: Cluster Client Local CCR:Cluster Client Remote CL:Local CR:Remote
Total active entries from all ports = 1
Total static entries from all ports = 3
MAC-Address      Port          Type        Index   MCT-Type  VLAN
0000.0022.3333  8/1           Static      4254    CML       20
0000.0022.3333  8/3           Static      4254    CML       20
0000.0022.3333  8/13          Static      4254    CML       20
```

Syntax: **show mac-address [cluster { id|name } local| remote]**

MAC clear commands

To clear all MAC addresses in the system, enter the following command.

```
device# clear mac
```

Syntax: **clear mac**

NOTE

On SXL with authentication protocols and high traffic, the clear mac and mac flush operations log a lot of new address messages, which results in high CPU utilization for a few minutes.

Clearing cluster-specific MAC addresses

To clear cluster-specific MAC addresses in the system, enter a command such as the following.

```
device# clear mac cluster AGG-1 local
```

Syntax: **clear mac cluster { cluster-id |cluster-name } { local | remote }**

Clearing client-specific MAC addresses

To clear client-specific MAC addresses in the system, enter a command such as the following.

```
device# clear mac cluster AGG-1 client 1 local
```

Syntax: **clear mac cluster { cluster-id | cluster-name } client client-name { local | remote }**

Clearing VLAN-specific MAC addresses

To clear VLAN-specific MAC addresses in the system, enter a command such as the following.

```
device# clear mac vlan 2
```

Syntax: `clear mac vlan vlan_id`

Clearing MCT VLAN-specific MAC addresses

To clear MCT VLAN-specific MAC addresses in the system, enter a command such as the following.

```
device# clear mac cluster AGG-1 vlan 1 local
```

Syntax: `clear mac cluster { cluster_id | cluster-name } vlan vlan_id { local | remote }`

Clearing cluster client vlan-specific MACs

To clear cluster client-specific MAC addresses in the system, enter a command such as the following.

```
device# clear mac cluster AGG-1 vlan 2 client 1 local
```

Syntax: `clear mac cluster {cluster_id | cluster-name} vlan vlan_id client client_name { local | remote }`

Displaying MDUP packet statistics

To display the statistics of MDUP packets, enter a command such as the following.

```
device#show mac mdup-stats
MDUP Information
=====
MDUP Data buffers in queue : 0
MDUP Statistics
=====
MDUP Update Messages sent: 7
Add Mac sent: 20
Del Mac sent: 0
Move Mac sent: 0
MDUP Mac Info Messages sent: 1
MDUP Flush Messages sent: 1
MDUP Synch Messages sent: 0
MDUP Update Messages received: 3
Add Mac received: 40
Del Mac received: 0
Move Mac received: 0
MDUP Mac Info Messages received: 0
MDUP Flush Messages received: 0
MDUP Synch Messages received: 0
```

Syntax: `show mac mdup-stats`

Syncing router MAC addresses to peer MCT devices

The MCT cluster device uses a router MAC address to identify the packets that are addressed to the switch. Such packets may be received by a peer cluster device. The peer device switches packets over the ICL to the local MCT device to be routed properly.

Dynamic trunks

The MCT client creates a single dynamic trunk group toward the MCT cluster devices. The dynamic trunk group consists of two trunk groups, each of which is configured on one of the MCT devices. A dynamic trunk group runs Link Aggregation Control Protocol (LACP).

For the two dynamic trunk groups of the MCT to behave as a single trunk group from the MCT client's perspective, both of the dynamic trunk groups should have the same LACP system ID and key, referred to as the MCT system ID and MCT key.

NOTE

The LAG IDs are only significant locally and need not match on the two ends of a LAG.

The LACP system ID in the MCT-supporting device normally comes from the port MAC address. To support LACP over MCT, the ID must be obtained in another way. MCT uses a pre-defined algorithm to obtain the ID.

NOTE

Each MCT cluster device has a unique cluster ID and one MCT client ID. The LACP key is predefined from the client ID and cluster ID. The user cannot change the key.

MCT does not involve stacking, and control protocol synchronization is minimal. The LACP runs independently on the cluster devices.

Port loop detection

Loop detection can be used in an MCT topology to detect Layer 2 loops that occur due to misconfigurations, for example, on the client side when MCT links are not configured as trunk links on the MCT-unaware client.

In MCT, ICL links should be up at all times to prevent the cluster from going down. These links should not be shut down when a loop is detected in a network. Instead, other available ports (CCEPs) should be shut down. If loop detection BDPUs are received on the ICL port, instead of shutting down the ICL links, all CCEPs are error-disabled, and the user is notified with the following log message.

Loop-detection: Packet received on ICL port <port_number> for vlan <vlan_id>. Errdisable CCEPs.

Strict mode loop detection can be enabled on ICL ports. In strict mode, a port is disabled only if a packet is looped back to that same port. Strict mode overcomes specific hardware issues where packets are echoed back to the input port. This process assists in detecting hardware faults on ICL ports.

Loop-detection can be enabled on MCT and non-MCT VLANs simultaneously. There is no change in loop detection behavior when it is enabled on non-MCT VLANs.

The following example shows how to configure loop detection on MCT and non-MCT VLANs.

```
device(config)# vlan 1905
device(config-vlan-1905)# loop-detection
device(config-vlan-1905)# end
```

MCT Layer 2 protocols

Keep the following information in mind when configuring Layer 2 protocols with MCT.

MRP

- An ICL interface cannot be configured as an MRP secondary interface or vice versa because the ICL cannot be BLOCKING.
- MRP cannot be enabled on MCT CCEP port or vice versa.

STP/RSTP

Configuring STP on MCT VLANs at MCT cluster devices is not recommended. By default, the spanning tree is disabled in the MCT VLANs. If the network topology may create Layer 2 loops through external connections, STP may be enabled on switches outside the MCT cluster to prevent the Layer 2 loop. The MCT cluster devices then performs a pass-through forwarding of STP BPDUs received through its ports in the MCT VLAN.

- In rare cases in which the network topology consists of Layer 2 loops outside the MCT cluster that require STP/RSTP to be enabled on MCT VLANs in the cluster, the CCEPs are always in the spanning tree disabled state. Refer to [MCT configuration examples using STP](#) on page 199 to view deployment scenarios where STP is used in an MCT configuration to prevent Layer 2 loops.
- The STP/RSTP algorithms have been modified so that the ICL never enters blocking state. The ICL guard mechanism ensures that if the ICL is about to go into a blocking state, the port on which the superior BPDUs are being received is moved to blocking state, and the ICL guard timer starts running on it. This timer runs as long as superior BPDUs are received on this interface. As long as this timer runs on an interface, the superior BPDUs are dropped.
- The new BLK_BY_ICL STP state indicates that superior BPDUs were received on this interface, which could have led to blocking of the ICL interface, with the result that the CL port guard mechanism has been triggered on this port.
- In an 802.1s MSTP deployment, Brocade recommends disabling spanning tree on MCT cluster devices at the global level. MSTP cannot be configured on individual cluster devices.
- An MCT cluster can support up to 32 spanning tree instances.

BPDU forwarding

If the network deploys single STP or IEEE 802.1s (MSTP), the MCT cluster devices must be configured using the **bpdu-flood-enable** command to flood the single STP/MSTP BPDUs in the SSTP/MSTP domain (that is, to forward to all of the ports in the cluster switch, irrespective of VLAN.)

Syntax: [no] bpdu-flood-enable

When **bpdu-flood-enable** is configured, only the ICL should connect the two MCT cluster devices. (The keep-alive VLAN link should not connect them.) If there is an additional link, the flooded BPDU will cause a loop and high CPU utilization.

NOTE

The **bpdu-flood-enable** command is not supported on the Brocade ICX 7750.

Protocol-based VLANs

Protocol and subnet VLANs can be configured on MCT VLANs; however, the ICL and CCEPs cannot be configured as dynamic members of protocol-based VLANs (and vice versa). The ICL and CCEP can either be excluded from protocol-based VLANs, or they can be configured as static members. CEPs can be configured as dynamic or static members of protocol-based VLANs, or they can be excluded from them.

In a cluster, both cluster devices should have exactly the same protocol VLAN membership configuration for the ICL and CCEPs. The ICL and CCEPs should be configured with the same type of protocol/VLAN membership, although the CLI does not impose any restriction.

NOTE

This feature is not supported on the Brocade ICX 7750, Brocade ICX 7450, or Brocade ICX 7250.

Uplink switch

Uplink switch capability is supported on MCT VLANs. ICLs and CCEPs can be configured as uplink-switch ports. Both cluster devices should have exactly the same uplink-switch port memberships configured for the ICL and CCEPs.

NOTE

This feature is not supported on the Brocade ICX 7750, Brocade ICX 7450, or Brocade ICX 7250.

Layer 2 multicast snooping over MCT

To support multicast snooping over MCT, the ICL port uses MDUP to synchronize the following information between the cluster devices:

- MAC-forward entries (mcache entries on MCT VLAN)
- IGMP/MLD Join/Leave (control packets on MCT VLAN)
- PIM-SM/PIM6-SM Join/Prune (control packets on MCT VLAN)
- IGMP/MLD dynamic router ports on MCT VLAN

IGMP/MLD snooping

Snooping can be configured globally or at the VLAN level. Each cluster device in the MCT VLAN can be configured as active or passive. There is no restriction for cluster devices to run active-active or passive-passive configurations.

The following commands show configuration commands for the VLAN level (IGMP and MLD), the global level (IGMP/MLD), and for PIM-SM and PIM6-SM.

VLAN level (IGMP)

```
device(config)# vlan 100
device(config-vlan-100)# multicast active/passive
```

VLAN level (MLD)

```
device(config-vlan-100)# multicast6 active/passive
```

Global Level (IGMP/MLD)

```
device(config)# ip multicast active/passive
device(config)# ipv6 multicast active/passive
```

PIM-SM snooping (configured only on a VLAN and requires IGMP snooping to run in a passive mode):

```
device(config)# vlan 100
device(config-vlan-100)# multicast passive
device(config-vlan-100)# multicast pimsm-snooping
```

PIM6-SM snooping (configured only on a VLAN and requires MLD snooping to run in a passive mode):

```
device(config)# vlan 100
device(config-vlan-100)# multicast6 passive
device(config-vlan-100)# multicast6 pimsm-snooping
```

IGMP/MLD snooping behavior on MCT cluster devices

- Local information is synchronized to the MCT peer device using CCP. The information includes Mcache/FDB entry (on arrival of data traffic), joins/leaves, dynamic router ports, and PIM-SM snooping joins/prunes.
- Native control packets (joins/leaves) that are received are processed by protocol code and are forwarded if necessary.
- All control and data traffic is received on the ICL. The traffic is forwarded out of a CCEP only if the remote CCEP is down; otherwise, it is dropped by the egress filters on the CCEP.
- The ICL is added as outgoing interface (OIF) by default whenever the CCEP is a source or a receiver. This provides faster convergence during MCT failover.
- For IGMP/MLD joins/leaves:
 - Only control packets received on a CCEP are synced to the MCT peer using CCP.
 - Control packets received on a CEP are not synced to the MCT peer.
- Static groups and static router ports configured on a CCEP are not synced across to the MCT peer. For these features to work correctly, they must be manually configured on the respective CCEP of both the cluster nodes.

How failovers are handled for Layer 2 multicast over MCT

The following failover scenarios may occur. Refer to [MCT failover scenarios](#) on page 171 for other types of failover scenarios.

- Local CCEP Down EVENT:
 - Outgoing traffic on local CCEP will now go through the ICL and out of the remote CCEP.
 - Incoming traffic on local CCEP will now ingress through the remote CCEP, and then ingress through the ICL locally.
- Local CCEP Up EVENT:
 - Outgoing traffic on a remote CCEP (after egressing through the local ICL) will now start going out of the local CCEP.
 - Incoming traffic from a client through the ICL (after ingressing on remote CCEP) will now switch back to the local CCEP (this is true only if the client trunk hashing sends the traffic toward the local CCEP).
- CCP (Cluster communication protocol) Down EVENT:
 - All related information (IGMP/MLD group, mcache, dynamic router port, pim-sm snooping entry) that was synced from the peer device will now be marked for aging locally.
- CCP (Cluster communication protocol) Up EVENT:
 - All related information (IGMP/MLD group, mcache, dynamic router port, pim-sm snooping entry) that was learned locally will be synced to the peer device.

PIM-SM and PIM6-SM snooping over MCT

- PIM-SM snooping can be configured only on a VLAN. It requires IGMP snooping to be running in passive mode. IPv6 snooping is supported.
- PIM6-SM snooping can be configured only on a VLAN. It requires MLD snooping to be running in passive mode.

- Router ports can be configured on a VLAN or globally. They can be learned dynamically on the port where the query is received or configured statically.
- Both MCT1 devices must run pimsm-snoop.
- PIM messages are forwarded via the hardware.
- PIM join/prune is synced to the peer cluster device using CCP.
- PIM prune is processed only if indicated by the peer cluster device.
- PIM join/prune received natively on ICL is ignored.
- PIM hello is not synced but is received natively on ICL.
- PIM port/source information is refreshed on both cluster devices by syncing PIM messages. The information ages out if not refreshed.

Forwarding entries for PIM-SM and PIM6-SM multicast snooping

Table 31 and Table 32 list the forwarding entries for PIM-SM and PIM6-SM multicast snooping.

TABLE 31 Forwarding entries (*,G)^a

Event	MCT-1	MCT-2
No-Join	(*,G)->blackhole	(*,G)->blackhole
(S,G) Join on (MCT-1) CEP	(*,G)->CEP [s] ^b	(*,G)->ICL [s]
(S,G) Join on (MCT-2) CEP	(*,G)->ICL [s]	(*,G)->CEP [s]
(S,G) Join on (MCT-1) CCEP	(*,G)->CCEP [s], ICL [s]	(*,G)->CCEP [s], ICL [s]
(S,G) Join on (MCT-2) CCEP	(*,G)->CCEP[s], ICL [s]	(*,G)->CCEP [s], ICL [s]

a.) *ICL: The ICL port is added as default whenever CCEP is in OIF. The data traffic received from the ICL port will be filtered out by egress filters dynamically programmed on CCEPs.

b.) [s]: denotes sources maintained on port hash-list.

TABLE 32 Forwarding entries (S,G)^a

Event	MCT-1	MCT-2
No-Join	(S,G)->blackhole	(S,G)->blackhole
Join (MCT-1) CEP	(S,G)->CEP	(S,G)->ICL
Join (MCT-2) CEP	(S,G)->ICL	(S,G)->CEP
Join (MCT-1) CCEP	(S,G)->CCEP, ICL	(S,G)->CCEP, ICL
Join (MCT-2) CCEP	(S,G)->CCEP, ICL	(S,G)->CCEP, ICL

a.) *ICL: The ICL port is added as default whenever CCEP is in OIF. The data traffic received from the ICL port will be filtered out by egress filters dynamically programmed on CCEPs.

Forcing a port up in a basic MCT configuration

In a static trunk environment, Preboot eXecution Environment (PXE) images are too small for most operating systems to leverage LACP during the boot process. As a result, during a PXE build process, traffic sent by the server is dropped, and the build process can fail.

To correct this situation, a port on an ICX 6650 device connected to a server that is configured as an MCT client can be set to a “force-up” state so that even if the LACPDU is not received from the server, the connected port is up and forwards packets.

NOTE

When multiple ports from the same server are connected to an ICX 6650, the port on the ICX 6650 connected to the PXE-capable port on the server is the port that must be configured to the force-up state. The PXE-capable port varies from server to server.

Keep the following points in mind when configuring a port to a force-up state:

- A port can only be configured as the force-up port before the client is deployed.
- Only one port in an LACP link aggregation group can be configured as the force-up port. If you configure multiple ports as force-up, this error message is displayed: `Error: port portno is already configured as force-up port.`
- When a port is configured for force-up and the server boots for the first time, the port does not wait for any LACPDU but immediately begins to forward packets.
- If the port receives an LACPDU, it bundles with other ports and forms a link aggregation group. The server is operational.

- If the force-up port goes down while in a link aggregation group, the port continues to perform as a normal LACP trunk, and the server remains operational, with some ports down.
 - If the force-up port stops receiving LACPDU, the port ignores the time-out and remains operational.
- To configure the LACP client in a force-up state, use the **client-interface link-aggregation force-up ethernet** command at the Client level.

Syntax: [no] **client-interface link-aggregation force-up ethernet unit/slotnum/portnum**

The following example shows the link aggregation information for a port configured to a force-up state.

```
ICX6650-64 Router# show lag id 163
Total number of LAGs: 11
Total number of deployed LAGs: 11
Total number of trunks created:11 (113 available)
LACP System Priority / ID: 1 / 748e.f88f.2222
LACP Long timeout: 120, default: 120
LACP Short timeout: 3, default: 3

==== LAG "CCEP-163" ID 163 (dynamic Deployed) ===
LAG Configuration:
  Ports: e 1/1/47 to 1/1/48
  Port Count: 2
  Primary Port: 1/1/47
  Trunk Type: hash-based
  LACP Key: 20163
Deployment: HW Trunk ID 3
This is a Multi Chassis Trunk: (System Id: 0180.c200.0001, Key: 30163)

Port Link State Dupl Speed Trunk Tag Pvid Pri MAC Name
1/1/47 Up Forward Full 1G 163 Yes N/A 0 748e.f88f.2222
1/1/48 Down None None None 163 Yes N/A 0 748e.f88f.2222

Port [Sys P] [Port P] [ Key ] [Act] [Tio] [Agg] [Syn] [Col] [Dis] [Def] [Exp] [Ope]
1/1/47 1 1 20163 Yes L Agg Syn Col Dis Def No Frc
1/1/48 1 1 20163 Yes L Agg Syn No No Def No Dwn

  Partner Info and PDU Statistics
  Port Partner Partner LACP LACP
  System MAC Key Rx Count Tx Count
  1/1/47 0000.0000.0000 46 5475 5558
  1/1/48 0000.0000.0000 47 5477 5487
```

Displaying MCT information

This section describes the commands available to display information about MCT configuration and operation.

Displaying peer and client states

Use the **show cluster config** command to display the peer device and client states.

```
device# show cluster SXR122 config
cluster SXR122 100
rbridge-id 100
session-vlan 1
keep-alive-vlan 3
icl SXR122-MCT ethernet 1/1
peer 172.17.0.2 rbridge-id 101 icl SXR122-MCT
deploy
```

```

client KL134
rbridge-id 14
client-interface ethernet 1/23
deploy
client AGG131
rbridge-id 10
client-interface ethernet 12/2
deploy
client FOX135
rbridge-id 15
client-interface ethernet 1/25
deploy

```

Syntax: `show cluster cluster-name/cluster-id config`

Displaying state machine information

Use the **show cluster client** command to display additional state machine information, including the reason a local CCEP has gone down. You can specify an individual cluster and client as an option.

```

device# show cluster 1 client
Cluster 1
=====
Rbridge Id: 101, Session Vlan: 3999, Keep-Alive Vlan: 4001
Cluster State: Deploy
Client Isolation Mode: Loose
Configured Member Vlan Range: 100 to 105
Active Member Vlan Range: 100 to 105
MCT Peer's Reachability status using Keep-Alive Vlan: Peer Reachable
Client Info:
-----
Client: c1, rbridge-id: 300, Deployed
Client Port: 3/11
State: Up
Number of times Local CCEP down: 0
Number of times Remote CCEP down: 0
Number of times Remote Client undeployed: 0
Total CCRR packets sent: 4
Total CCRR packets received: 3

```

Syntax: `show cluster cluster_name/cluster_id client [client_name/client_RbridgeID]`

The following table shows the messages that may be displayed to explain why the local CCEP is down.

TABLE 33 Reasons for Local CCEP down

Reason for Local CCEP down	Meaning
client-interfaces shutdown	Command is configured.
client-isolation strict	Command is configured.
Deploy mismatch	Client is not deployed remotely.
Slave state	Client is in slave state when CCP is down.
cluster and client undeployed	Neither the cluster nor client is deployed.
cluster undeployed	Cluster is not deployed.
client undeployed	Client is not deployed.

Displaying cluster, peer, and client states

Use the **show cluster ccp peer** command to display cluster, peer device, and client states. As an option, you can specify an individual cluster and request additional details.

```
device# show cluster 1 ccp peer
...
PEER IP ADDRESS      STATE      UP TIME
-----
10.1.1.1          OPERATIONAL 0 days: 2 hr:25 min:16 sec
device (config-cluster-SX_1)# show cluster 1 ccp peer detail
*****Peer Session Details*****
IP address of the peer           10.1.1.1
Rbridge ID of the peer           100
Session state of the peer        OPERATIONAL
Next message ID to be send       287
Keep Alive interval in seconds   30
Hold Time Out in seconds         90
Fast Failover is enable for the session
UP Time                         0 days: 2 hr:22 min:58 sec
Number of  tcp packet allocations failed 0
Message   Init     Keepalive   Notify      Application  Badmessages
Send      3        2421        2            53           0
Receive   3        2415        0            37           0
TCP connection is up
TCP connection is initiated by   10.1.1.2
TCP connection tcbHandle not pending
TCP connection packets not received
*****TCP Connection Details*****
TCP Connection state: ESTABLISHED Maximum segment size: 1436
Local host: 10.1.1.2, Local Port: 12203
Remote host: 10.1.1.1, Remote Port: 4175
ISentSeq: 1867652277 SendNext: 1867660731 TotUnAck: 0
TotSent: 8454 ReTrans: 9 UnAckSeq: 1867660731
IRcvSeq: 3439073167 RcvNext: 3439078415 SendWnd: 16384
TotalRcv: 5248 DupliRcv: 16 RcvWnd: 16384
SendQue: 0 RcvQue: 0 CngstWnd: 1452
```

Syntax: **show cluster [cluster_name/cluster-id] ccp peer [details]**

Displaying information about Ethernet interfaces

Use the **show interface ethernet** command to display information about Ethernet interfaces. The MCT-related information is shown in bold in the following example.

```
device# show interface ethernet 7/1
...
GigabitEthernet7/1 is disabled, line protocol is down
Hardware is GigabitEthernet, address is 0024.3822.8260 (bia 0024.3822.8260)
Configured speed auto, actual unknown, configured duplex fdx, actual unknown
Configured mdi mode AUTO, actual unknown
Member of L2 VLAN ID 1, port is untagged, port state is DISABLED
BPDU guard is Disabled, ROOT protect is Disabled
Link Error Dampening is Disabled
STP configured to ON, priority is level0
Flow Control is config enabled, oper disabled, negotiation disabled
Mirror disabled, Monitor disabled
Not member of any active trunks
Not member of any configured trunks
No port name
IPG MII 96 bits-time, IPG GMII 96 bits-time
MTU 1500 bytes, encapsulation Ethernet
ICL port for icl1 in cluster id 1
300 second input rate: 0 bits/sec, 0 packets/sec, 0.00% utilization
300 second output rate: 0 bits/sec, 0 packets/sec, 0.00% utilization
0 packets input, 0 bytes, 0 no buffer
Received 0 broadcasts, 0 multicasts, 0 unicasts
0 input errors, 0 CRC, 0 frame, 0 ignored
0 runts, 0 giants
0 packets output, 0 bytes, 0 underruns
Transmitted 0 broadcasts, 0 multicasts, 0 unicasts
```

```

0 output errors, 0 collisions
Relay Agent Information option: Disabled
show interface ethernet 7/3
GigabitEthernet7/3 is disabled, line protocol is down
    Hardware is GigabitEthernet, address is 0024.3822.8262 (bia 0024.3822.8262)
    Configured speed auto, actual unknown, configured duplex fdx, actual unknown
    Configured mdi mode AUTO, actual unknown
    Member of L2 VLAN ID 1, port is untagged, port state is DISABLED
    BPDU guard is Disabled, ROOT protect is Disabled
    Link Error Dampening is Disabled
    STP configured to ON, priority is level0
    Flow Control is config enabled, oper disabled, negotiation disabled
    Mirror disabled, Monitor disabled
    Not member of any active trunks
    Not member of any configured trunks
    No port name
    IPG MII 96 bits-time, IPG GMII 96 bits-time
    MTU 1500 bytes, encapsulation Ethernet
CCEP for client c149_150 in cluster id 1
    300 second input rate: 0 bits/sec, 0 packets/sec, 0.00% utilization
    300 second output rate: 0 bits/sec, 0 packets/sec, 0.00% utilization
    0 packets input, 0 bytes, 0 no buffer
    Received 0 broadcasts, 0 multicasts, 0 unicasts
    0 input errors, 0 CRC, 0 frame, 0 ignored
    0 runts, 0 giants
    0 packets output, 0 bytes, 0 underruns
    Transmitted 0 broadcasts, 0 multicasts, 0 unicasts
    0 output errors, 0 collisions
    Relay Agent Information option: Disabled

```

Syntax: show interface ethernet x/y

Displaying STP information

Use the **show span** command to display STP information for an entire device.

```

device# show span
...
STP instance owned by VLAN 90

Global STP (IEEE 802.1D) Parameters:
VLAN Root          Root Root   Prio Max He- Ho- Fwd Last     Chg Bridge
ID    ID           Cost Port   rity Age llo ld dly Chang   cnt Address
                         Hex sec sec sec sec sec sec
90 8000748ef8f9739d 0     Root    8000 20 2 1 15 259968 1 748ef8f9739d

Port STP Parameters:
Port  Prio Path State      Fwd     Design  Designated      Designated
Num   rity Cost          Trans   Cost   Root      Bridge
Hex
1/2/1 80   1    FORWARDING 1      0      8000748ef8f9739d 8000748ef8f9739d
1/3/5 80   0    DISABLED   0      0      0000000000000000 0000000000000000
1/3/6 80   0    DISABLED   0      0      0000000000000000 0000000000000000

```

Syntax: show span [vlan vlan-id] | [pvst-mode] | [num] | [detail [vlan vlan-id [Ethernet [stack-unit[slotnum] portnum] | num]]]

Displaying information for multicast snooping

Use the **show ip pimsm-snooping cache** command to display (*,g), (s,g) and oif information learned via PIM join/prune messages.

```

Device(config)# show ip pimsm-snooping cache
OIF Info:
TR - OIF Belongs to Trunk/LAG, Primary port is displayed
SG - (*,g)/(s,g) downstream fsm state:

```

```

NI : No Info, J : Join, PP : Prune Pending, CLEAN : cleanup in progress
RPT - (s,g,rpt) downstream fsm state:
NI : No Info, P : Pruned, PP : Prune Pending, Px : Temp step in (*,G)
join processing, PPx : Temp State in (*,G) processing, CLEAN : cleanup
in progress.

PIMSM Snoop cache for vlan 503
1    (* 225.0.0.1) Up Time: 1d 19:41:48
OIF: 1
TR(e3/13) G : J(194) ET: 210, Up Time: 1d 19:41:48 , ICL, Remote

2    (* 225.1.1.1) Up Time: 5d 18:43:56
OIFs: 2
TR(e3/10) G : J(167) ET: 210, Up Time: 5d 18:43:56 , CCEP, Local
TR(e3/13) G : J(200) ET: 210, Up Time: 1d 19:41:48 , ICL, Remote

```

Syntax: show ip pimsm-snooping cache

You can also use the **show ip pimsm-snooping cache** command to display the MCT information if the VLAN is an MCT member.

In the following example, YES indicates that reports/leaves were received locally (processing native control packets).

```

Device(config)# show ip multicast cluster group
p-physical, ST:static, QR:querier, EX:exclude, IN:include, Y:yes, N:no
VL100 : 1 groups, 1 group-port
group p-port          ST  QR  life mode source local
1     225.1.1.1  e5/5  no   no  200   EX      0   YES
2     225.1.1.1  e5/10 no   no  200   EX      0   YES

```

In the following example, NO indicates that reports/leaves were received remotely. In this case, a join was received on the CCEP of the MCT peer device. Native control packets were processed by the peer device, and then the entries were synched over MDUP to this cluster device.

```

Device(config)# show ip multicast cluster group
p-physical, ST:static, QR:querier, EX:exclude, IN:include, Y:yes, N:no
VL100 : 1 groups, 1 group-port
group p-port          ST  QR  life mode source local
1     225.1.1.1  e1/10 no   no  200   EX      0   NO
2     225.1.1.1  e1/10 no   no  200   EX      0   NO

```

The following example displays information about the IGMP multicast mcache. It is used to verify if FDB is programmed when a data packet arrives.

```

Device(config-vlan-101)# show ip multicast cluster mcache
Example: (S G) cnt=: cnt is number of SW processed packets
          OIF: e1/22 TR(e1/32), TR is trunk, e1/32 primary
          [1,10]: [1 - has local oif, 10 - ICL due to CCEP count]
vlan 101, 1 caches. use 1 VIDX
1    (* 230.1.2.23) cnt=2
      OIF: TR(e5/4) tag TR(e5/5)
      age=37s up-time=37s, change=37s vidx=7405 (ref-cnt=1)

```

The following example displays status about the IGMP router port.

```

Device(config)# show ip multicast cluster vlan 100
Version=2, Intervals: Query=125, Group Age=260, Max Resp=10, Other Qr=260
VL100: cfg V3, vlan cfg passive, 1 grp, 2 (SG) cache, rtr ports,
router ports: e5/9(260) 100.100.100.1 (local:1, mct peer:0),
e5/4    has 1 groups,
This interface is non-Querier (passive)
default V3 trunk
(local:1, mct peer:0)

```

Syntax: show ip multicast cluster { group | mcache | vlan vlan-id }

Use the **show ip multicast cluster pimsm-snooping** command to display detailed information about OIFs added via a pimsm-snooping module.

```
Device(config)# show ip multicast cluster pimsm-snooping
```

```
Example: Port: 7/3 (age, port type, ref_count, owner flag, pruned flag)
source: 7/3 has 1 src: 11.0.0.5(age, ref_count, owner flag, pruned flag)
owner flag: 0x0: local, 0x1 remote cep, 0x2 remote ccep
vlan 100, has 1 caches.
1 (* 224.10.10.10) has 2 pim join ports out of 2 OIF
7/3 (1,ICL), 7/5 (1, CCEP)
```

Syntax: show ip multicast cluster pimsm-snooping [group | vlan]

Refer to the *FastIron Command Reference* for information on **show ipv6 multicast cache** and other MLD multicast commands.

MCT configuration examples

The examples in this section show the topology and configuration for a single-level MCT deployment, VRRP/VRRP-E over MCT, OSPF over MCT and BGP over MCT.

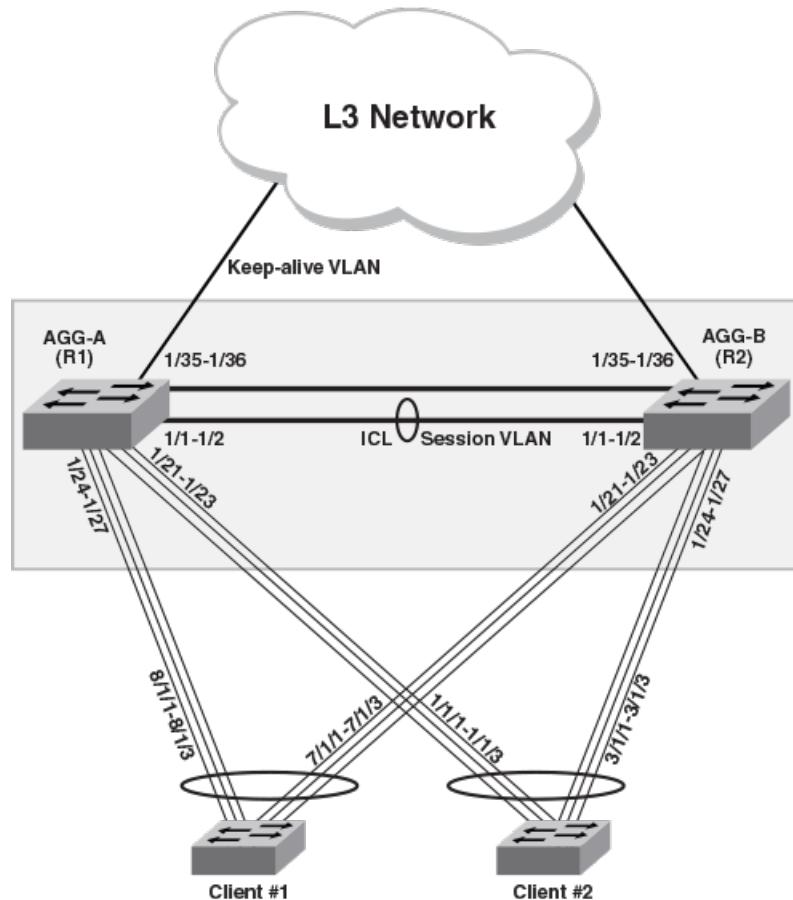
Single-level MCT example

The following figure depicts a single-level MCT configuration. The clients can be server hosts or networking devices. The associated configuration follows.

NOTE

The LAG IDs are locally significant only and need not be matching on the two ends of a LAG.

FIGURE 32 Single level MCT configuration



Client 1 - Configuration

If client 1 is a Brocade switch in [Figure 32](#) on page 190, you can configure it as follows:

```
!
vlan 1905 by port
  tagged ethe 7/1/1 to 7/1/3 ethe 8/1/1 to 8/1/3
  spanning-tree
!
!
lag lag_client1_1 dynamic id 100
  ports ethe 7/1/1 to 7/1/3 ethe 8/1/1 to 8/1/3
  primary-port 7/1/1
  deploy
!
```

Client 2- Configuration

If client 2 is a Brocade switch in [Figure 32](#) on page 190, you can configure it as follows:

```
!
vlan 1905 by port
  tagged ethe 1/1/1 to 1/1/3 ethe 3/1/1 to 3/1/3
  spanning-tree
!
lag lag_client2_1 dynamic id 200
```

```

ports ethe 1/1/1 to 1/1/3 ethe 3/1/1 to 3/1/3
primary-port 1/1/1
deploy
!
```

AGG-A (R1) - Configuration

This section presents the configuration for the AGG-A (R1) cluster device in [Figure 32](#) on page 190.

```

lag lag_agg_a_1 static id 103
ports ethe 1/1/1 to 1/2
primary-port 1/1
deploy
!
lag lag_agg_a_2 dynamic id 104
ports ethe 1/2/4 to 1/2/7
primary-port 1/24
deploy
!
lag lag_agg_a_3 dynamic id 105
ports ethe 1/2/1 to 1/2/3
primary-port 1/21
deploy
!
vlan 2 name session-vlan by port
tagged ethe 1/1 to 1/2
router-interface ve 2
!
vlan 3 name keep-alive-vlan by port
tagged ethe 1/35-1/36
router-interface ve 3
!
vlan 1905 by port
tagged ethe 1/2/4 to 1/2/7 ethe 1/2/1 ethe 1/2/3 ethe 1/1 to 1/2
!
hostname R1
!
interface ve 2
ip address 10.1.1.1 255.255.255.0
!
interface ve 3
ip address 10.1.2.1 255.255.255.0
!
!
cluster MCT1 1
rbridge-id 1
session-vlan 2
keep-alive-vlan 3
icl BH1 ethernet 1/1
peer 10.1.1.2 rbridge-id 2 icl BH1
deploy
client client-1
rbridge-id 1901
client-interface ethe 1/24
deploy
client client-2
rbridge-id 1902
client-interface ethe 1/21
deploy
!
```

AGG-B (R2) - Configuration

This section presents the configuration for the AGG-B (R2) cluster device in [Figure 32](#) on page 190.

```

lag lag_agg_b_1 static id 103
ports ethe 1/1/1 to 1/2
primary-port 1/1
deploy
!
lag lag_agg_b_2 dynamic id 105
```

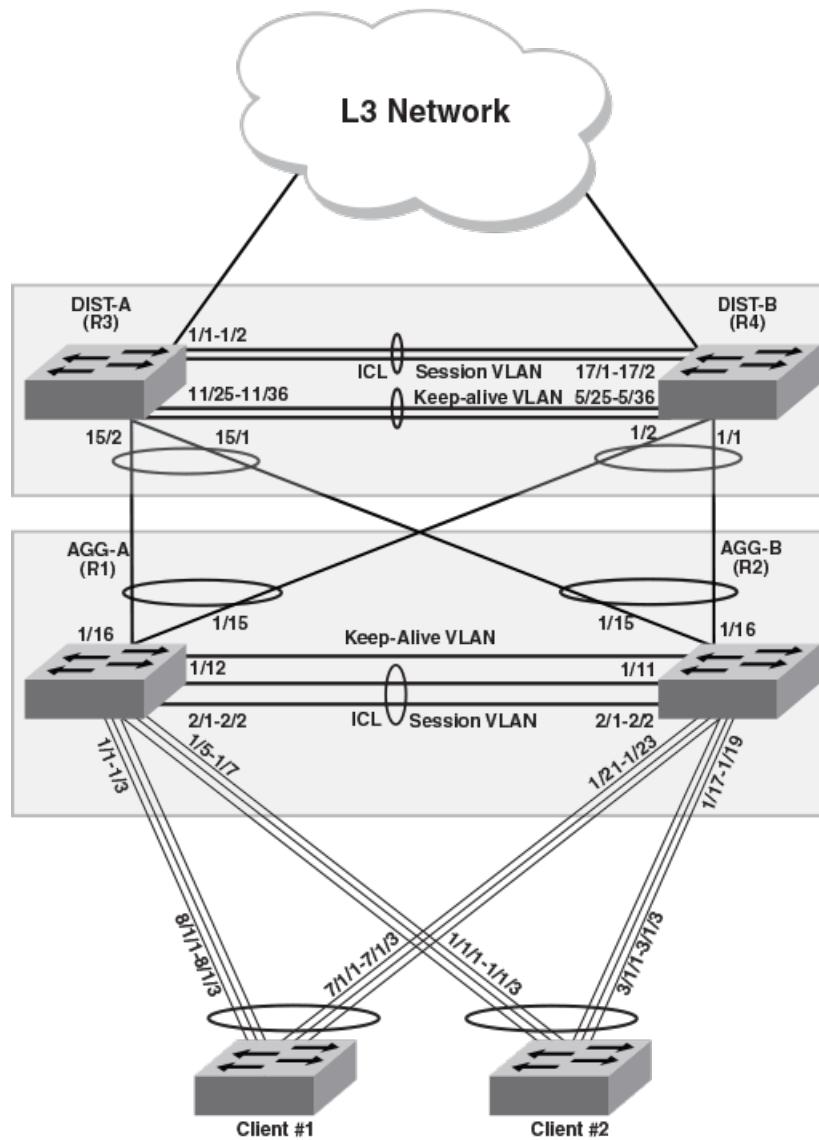
```
ports ethe 1/24 to 1/27
primary-port 1/24
deploy
!
lag lag_agg_b 3 dynamic id 104
ports éthe 1/21 to 1/23
primary-port 1/21
deploy
!
vlan 2 name session-vlan by port
tagged ethe 1/1 to 1/2
router-interface ve 2
!
vlan 3 by port
tagged ethe 1/35-1/36
router-interface ve 3
!
!
vlan 1905 by port
tagged ethe 1/24 to 1/27 ethe 1/21 to 1/23 ethe 1/1 to 1/2
!
hostname R2
!
interface ve 2
ip address 10.1.1.2 255.255.255.0
!
interface ve 3
ip address 10.1.2.2 255.255.255.0
!
cluster MCT1 1
rbridge-id 2
session-vlan 2
keep-alive-vlan 3
icl BH1 ethernet 1/1
peer 10.1.1.1 rbridge-id 1 icl BH1
deploy
client client-1
rbridge-id 1901
client-interface ethe 1/21
deploy
client client-2
rbridge-id 1902
client-interface ethe 1/24
!
```

Two-level MCT example

The following figure depicts a two-level MCT configuration. The clients can be server hosts or networking devices. The associated configuration follows.

NOTE

The LAG IDs are locally significant only and need not be matching on the two ends of a LAG.

FIGURE 33 Two-level MCT configuration

NOTE

In a two-level MCT configuration using dynamic LAGs, ensure that the upper and lower clusters have different Cluster IDs because the Cluster LACP module uses the Cluster ID as part of the LACPDUs's system ID.

The client configuration is the same as in the single-level example (refer to [Single-level MCT example](#) on page 189).

AGG-A (R1) - Configuration

This example presents the configuration for the AGG-A (R1) cluster device in [Figure 33](#) on page 193.

```
lag lag_agg_a_1 static id 103
```

```

ports ethe 2/1 to 2/2
primary-port 2/1
deploy
!
lag lag_agg_a_2 dynamic id 104
ports ethe 1/7 to 1/3
primary-port 1/1
deploy
!
lag lag_agg_a_3 dynamic id 105
ports ethe 1/7 to 1/7
primary-port 1/5
deploy
!
lag lag_agg_a_4 dynamic id 106
ports ethe 1/7 to 1/16
primary-port 1/15
deploy
!
vlan 2 name session-vlan by port
tagged ethe 2/1 to 2/2
router-interface ve 2
!
vlan 3 name keep-alive-vlan by port
tagged ethe 1/12
router-interface ve 3
!
!
vlan 1905 by port
tagged ethe 1/1 to 1/3 ethe 1/5 to 1/7 ethe 1/15 to 1/16 ethe 2/1 to 2/2
!
hostname R1
!
interface ve 2
ip address 10.1.1.1 255.255.255.0
!
interface ve 3
ip address 10.1.2.1 255.255.255.0
!
!
cluster MCT1 1
rbridge-id 1
session-vlan 2
keep-alive-vlan 3
icl BH1 ethernet 2/1
peer 10.1.1.2 rbridge-id 2 icl BH1
deploy
client client-1
rbridge-id 1901
client-interface ethe 1/1
deploy
client client-2
rbridge-id 1902
client-interface ethe 1/5
deploy
client DIST_Cluster
rbridge-id 1903
client-interface ethe 1/15
deploy
!

```

AGG-B (R2) - Configuration

This example presents the configuration for the AGG-B (R2) cluster device in [Figure 33](#) on page 193.

```

lag lag_agg_b_1 static id 106
ports ethe 2/7 to 2/2
primary-port 2/1
deploy
!
lag lag_agg_b_2 dynamic id 107
ports ethe 1/7 to 1/19
primary-port 1/17
deploy
!
```

```

lag lag_agg_b_3 dynamic id 108
  ports ethe 1/21 to 1/23
  primary-port 1/21
  deploy
!
lag lag_agg_b_4 dynamic id 109
  ports ethe 1/15 to 1/16
  primary-port 1/15
  deploy
!
vlan 2 name session-vlan by port
  tagged ethe 2/1 to 2/2
  router-interface ve 2
!
vlan 3 name keep-alive-vlan by port
  tagged ethe 1/11
  router-interface ve 3
!
vlan 1905 by port
  tagged ethe 1/15 to 1/19 ethe 1/21 to 1/23 ethe 2/1 to 2/2
!
hostname R2
!
interface ve 2
  ip address 10.1.1.2 255.255.255.0
!
interface ve 3
  ip address 10.1.2.2 255.255.255.0
!
cluster MCT1 1
  rbridge-id 2
  session-vlan 2
  keep-alive-vlan 3
  icl BH1 ethernet 2/1
  peer 10.1.1.1 rbridge-id 1 icl BH1
  deploy
  client client-1
    rbridge-id 1901
    client-interface ethe 1/21
    deploy
  client client-2
    rbridge-id 1902
    client-interface ethe 1/17
    deploy
  client DIST_Cluster
    rbridge-id 1903
    client-interface ethe 1/15
    deploy
!
```

DIST-A (R3) - Configuration

This example presents the configuration for the DIST-A (R3) cluster device in [Figure 33](#) on page 193.

```

!
lag lag_dist_a_1 static id 15
  ports ethe 1/1 to 1/2
  primary-port 1/1
  deploy
lag lag_dist_a_2 dynamic id 16
  ports ethe 15/1 to 15/2
  primary-port 15/1
  deploy
!
lag keep-alive static id 200
  ports ether 11/25 to 11/36
  primary-port 11/25
  deploy
!
vlan 5 name session-vlan by port
  tagged ethe 1/1 to 1/2
  router-interface ve 5
!
vlan 6 name keep-alive-vlan by port
```

```

tagged ethe 11/25 to 11/36
router-interface ve 6
spanning-tree
!
vlan 1905 by port
tagged ethe 1/1 to 1/2 ethe 15/1 to 15/2
!
hostname R3
hitless-failover enable
!
interface ve 5
ip address 10.2.1.1 255.255.255.0
!
interface ve 6
ip address 10.2.2.1 255.255.255.0
!
cluster MCT2 2
rbridge-id 3
session-vlan 5
keep-alive-vlan 6
icl BH3 ethernet 1/1
peer 10.2.1.2 rbridge-id 4 icl BH3
deploy
client AGG_Cluster
rbridge-id 1801
client-interface ethe 15/1
deploy

```

DIST-B (R4) - Configuration

This example presents the configuration for the DIST-B (R4) cluster device in [Figure 33](#) on page 193.

```

lag lag dist b 1 static id 40
ports ethe 17/1 to 17/2
primary-port 17/1
deploy
!
lag dist b_2 dynamic id 41
ports ethe 1/1 to 1/2
primary-port 1/1
deploy
!
lag keep-alive static id 201
ports ethe 5/25 to 5/36
primary-port 5/25
deploy
!
vlan 5 name session-vlan by port
tagged ethe 17/1 to 17/2
router-interface ve 5
!
vlan 6 name keep-alive-vlan by port
tagged ethe 5/25 to 5/36
router-interface ve 6
spanning-tree
!
vlan 1905 by port
tagged ethe 1/1 to 1/2 ethe 17/1 to 17/2
!
hostname R4
hitless-failover enable
!
interface ve 5
ip address 10.2.1.2 255.255.255.0
!
interface ve 6
ip address 10.2.2.2 255.255.255.0
!
cluster MCT2 2
rbridge-id 4
session-vlan 5
keep-alive-vlan 6
icl BH3 ethernet 17/1
peer 10.2.1.1 rbridge-id 3 icl BH3
deploy

```

```

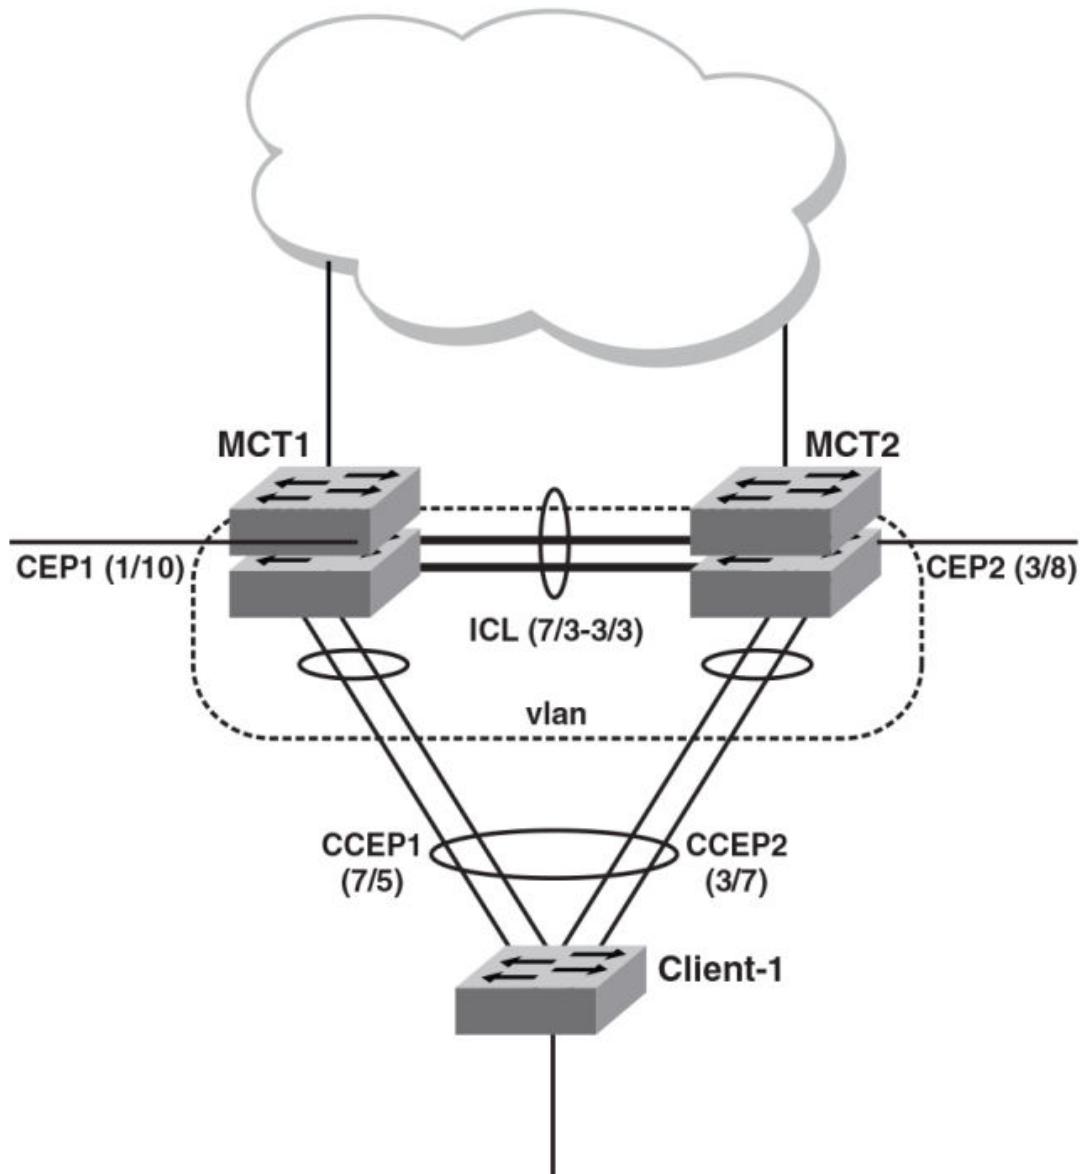
client AGG_Cluster
rbridge-id 1801
client-interface ethe 1/1
deploy

```

Multicast snooping configuration example

The following figure depicts a multicast snooping configuration. Sample configurations follow.

FIGURE 34 Multicast snooping over MCT



The following example shows the configuration for multicast snooping for the MCT1 cluster device in the previous figure.

```
vlan 100 by port
```

```

tagged ethe 7/3
untagged ethe 7/5 ethe 7/6
multicast passive
multicast pimsm-snooping
multicast6 passive
multicast6 pimsm-snooping
!
vlan 3000 name session by port
tagged ethe 7/3
router-interface ve 3000
vlan 3001 name keep-alive-vlan
tagged eth 7/4
interface ve 3000
ip address 10.1.1.2 255.255.255.0
!
cluster SX 3000
rbridge-id 2
session-vlan 3000
keep-alive-vlan 3001
icl SX-MCT ethernet 7/3
peer 10.1.1.3 rbridge-id 3 icl SX-MCT
deploy
client client-1
    rbridge-id 100
    client-interface ethernet 7/5
    deploy
!

```

The following example shows the configuration for multicast snooping for the MCT2 cluster device in [Figure 34](#).

```

!
vlan 100 by port
tagged ethe 3/3
untagged ethe 3/7 ethe 3/8
multicast passive
multicast pimsm-snooping
multicast6 passive
multicast6 pimsm-snooping
!
vlan 3000 name session by port
tagged ethe 3/3
router-interface ve 3000
vlan 3001 name keep-alive-vlan
tagged eth 3/4
interface ve 3000
ip address 10.1.1.3 255.255.255.0
!
cluster SX 3000
rbridge-id 3
session-vlan 3000
keep-alive-vlan 3001
icl SX-MCT ethernet 3/3
peer 10.1.1.2 rbridge-id 2 icl SX-MCT
deploy
client client-1
    rbridge-id 100
    client-interface ethernet 3/7
    deploy
!
```

The following example shows the global configuration for multicast snooping for the MCT1 cluster device in [Figure 34](#).

```

vlan 100 by port
tagged ethe 7/3
untagged ethe 7/5 ethe 7/6
!
vlan 3000 name session by port
tagged ethe 7/3
router-interface ve 3000
vlan 3001 name keep-alive-vlan
tagged eth 7/4
ip multicast active
interface ve 3000
ip address 10.1.1.2 255.255.255.0

```

```
!
cluster SX 3000
rbridge-id 2
session-vlan 3000
keep-alive-vlan 3001
icl SX-MCT ethernet 7/3
peer 10.1.1.3 rbridge-id 3 icl SX-MCT
deploy
client client-1
rbridge-id 100
client-interface ethernet 7/5
deploy
!
```

The following example shows the global configuration for multicast snooping for the MCT2 cluster device in [Figure 34](#).

```
!
vlan 100 by port
tagged ethe 3/3
untagged ethe 3/7 ethe 3/8
!
vlan 3000 name session by port
tagged ethe 3/3
router-interface ve 3000
vlan 3001 name keep-alive-vlan
tagged eth 3/4
ip multicast passive
interface ve 3000
ip address 10.1.1.3 255.255.255.0
!
cluster SX 3000
rbridge-id 3
session-vlan 3000
keep-alive-vlan 3001
icl SX-MCT ethernet 3/3
peer 10.1.1.2 rbridge-id 2 icl SX-MCT
deploy
client client-1
rbridge-id 100
client-interface ethernet 3/7
deploy
```

MCT configuration examples using STP

Although MCT is considered an alternative to Spanning Tree, Spanning Tree protocols can be enabled in an MCT configuration as an added protection for any Layer 2 loops. The following use-case scenarios demonstrate the use of Spanning Tree protocols in an MCT configuration:

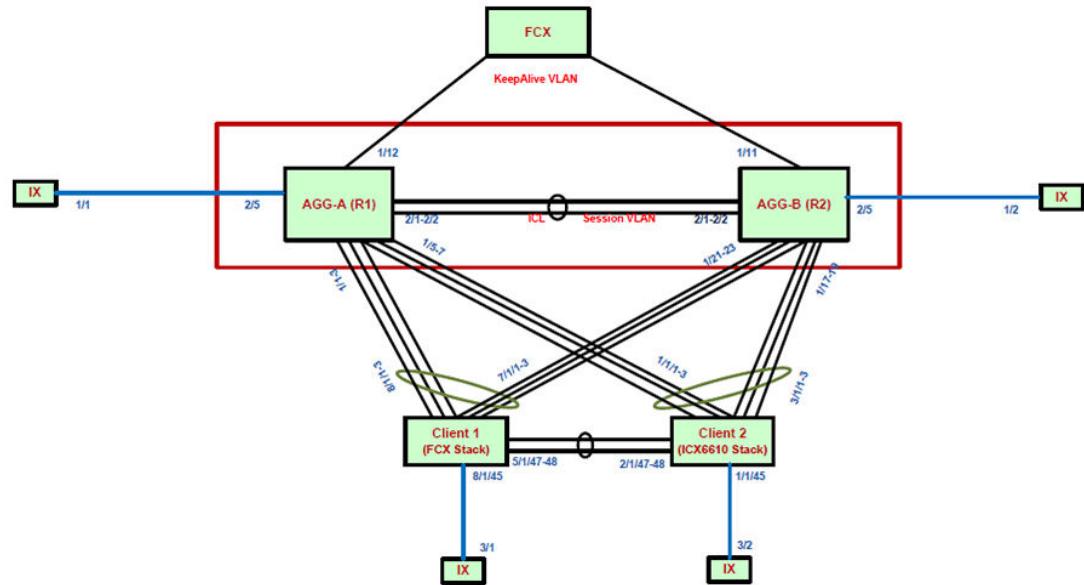
- [Example 1: Configure the Per-VLAN Spanning Tree on the MCT Clients](#) on page 202
- [Example 2: Configure Single Spanning Tree \(SSTP\) on the MCT Clients](#) on page 202
- [Example 3: Configure Multiple Spanning Tree \(MSTP\) on the MCT Clients](#) on page 203

The examples show STP elements enabled on MCT clients. It is recommended that you enable STP only on MCT clients, NOT on MCT cluster devices.

The following figure shows the base configuration of the MCT cluster, MCT clients, and the associated link-aggregation. The scenarios are depicted based on this network topology.

NOTE

The LAG IDs are locally significant only and need not be matching on the two ends of a LAG.

FIGURE 35 Sample network topology - Using STP in an MCT configuration

Router-1 configuration

```

!
lag "1" static id 1
  ports ethernet 1/1 ethernet 1/3
  primary-port 1/1
  deploy

lag "1" static id 2
  ports ethernet 1/5 ethernet 1/7
  primary-port 1/5
  deploy

lag "1" static id 3
  ports ethernet 2/1 ethernet 2/2
  primary-port 2/1
  deploy
!

vlan 2 name session-vlan by port
tagged ethe 2/1 to 2/2
router-interface ve 2
!

vlan 3 name keep-alive-vlan by port
tagged ethe 1/12
router-interface ve 3
!

vlan 1905 name MAC-scaling-vlan by port
tagged ethe 1/1 to 1/3 ethe 1/5 to 1/7 ethe 1/15 to
1/16 ethe 2/1 to 2/2 ethe 2/4 ethe 3/1 ethe 3/3
!
hostname R1
!
link-keepalive ethe 1/1 to 1/3 ethe 1/5 to 1/7
hitless-failover enable
!
interface ve 2
ip address 21.1.1.1 255.255.255.0
!
interface ve 3

```

AGG-B (R2) - Configuration

```
!
lag "1" static id 1
  ports ethernet 1/17 ethernet 1/19
  primary-port 1/17
  deploy

lag "1" static id 2
  ports ethernet 1/21 ethernet 1/23
  primary-port 1/21
  deploy

lag "1" static id 3
  ports ethernet 2/1 ethernet 2/3
  primary-port 2/1
  deploy
!
vlan 2 name session-vlan by port
tagged ethe 2/1 to 2/2
router-interface ve 2
!
vlan 3 by port
tagged ethe 1/11
router-interface ve 3
!
vlan 1905 name MAC-scaling-vlan by port
tagged ethe 1/15 to 1/19 ethe 1/21
```

Client-1 - Configuration

```
!
trunk ethe 7/1/1 to 7/1/3 ethe 8/1/1 to 8/1/3
!
vlan 1905 by port
  tagged ethe 5/1/47 to 5/1/48 ethe 7/1/1 to 7/1/3 ethe 8/1/1 to 8/1/3 ethe 8/1/45
!
link-keepalive ethe 7/1/1 to 7/1/3 ethe 8/1/1 to 8/1/3
!
lag lag1 dynamic id 1
  ports ethernet 5/1/47 to 5/1/48
  primary-port 5/1/47
  deploy
!
lldp run
end
```

Client-2 - Configuration

```
!
trunk ethe 1/1/1 to 1/1/3 ethe 3/1/1 to 3/1/3
!
vlan 1905 name MAC-scaling-vlan by port
  tagged ethe 1/1/1 to 1/1/3 ethe 1/1/45 ethe 2/1/47 to 2/1/48 ethe 3/1/1 to 3/1/3
  2/1/48 ethe 3/1/1 to 3/1/3
!
link-keepalive ethe 1/1/1 to 1/1/3 ethe 3/1/1 to 3/1/3
!
lag lag1 dynamic id 1
  ports ethernet 2/1/47 to 2/1/48
  primary-port 2/1/47
  deploy
!
lldp run
end
```

Example 1: Configure the Per-VLAN Spanning Tree on the MCT Clients

External connections between clients other than the links in an MCT cluster can cause Layer 2 loops. Use Spanning Tree on the MCT clients so that the MCT cluster forwards Spanning Tree Bridge Protocol Data Units (BPDU) as if the cluster were in a pass-through mode.

Configure per-VLAN Spanning Tree on the two MCT VLANS 1901 and 1905 to have Rapid Spanning Tree (RSTP/802.1w). This example is based on the network topology shown in [Figure 35](#) on page 200.

Client-1 Configuration

```
Client-1(config)# vlan 1901 1905
Client-1(config-mvlan-1901*1905)# spanning-tree 802-1w
Client-1(config-mvlan-1901*1905)# end
```

Client-2 Configuration

```
Client-2(config)# vlan 1901
Client-2(config-vlan-1901)# spanning-tree 802-1w
Client-2(config-vlan-1901)# spanning-tree 802-1w priority 4095
Client-2(config-vlan-1901)# vlan 1905
Client-2(config-vlan-1901)# spanning-tree 802-1w
Client-2(config-vlan-1901)# end
```

The MCT cluster switches do not have the spanning tree configured, but the BPDU's are passed through, and the Spanning Tree on the clients converges.

Use the **show 802-1w vlan** command to display the RSTP information for the specified port-based VLAN.

Example 2: Configure Single Spanning Tree (SSTP) on the MCT Clients

In a network where MCT clients have Single 802.1d or Single 802-1w elements of the Spanning Tree protocol enabled, configure the central processing unit (CPU) of the MCT peer switches so that it performs BPDU forwarding to avoid Layer 2 loops.

Using SSTP on MCT clients allows you to run a separate spanning tree on each port-based VLAN, which you can enable or disable on an individual basis. As an alternative, you can run a single spanning tree across all ports and VLANs on the device.

Enabling BDPU flooding can increase the CPU usage. When BDPU flooding is enabled, do not create redundant links between the MCT cluster devices or cascade multiple MCT clusters.

To enable the CPU to perform BDPU forwarding, use the **bpdu-flood-enable** command. This example is based on the network topology shown in [Figure 35](#) on page 200.

Router-1 configuration

```
Router-1(config)# bpdu-flood-enable
Warning - Any received untagged BPDU's will now be flooded to all the ports.
```

Router-2 configuration

```
Router-2(config)# bpdu-flood-enable
Warning - Any received untagged BPDU's will now be flooded to all the ports.
```

Client-1 configuration

```

Client-1(config)# spanning-tree single 802-1w
Client-1(config)# show 802-1w vlan 1905
Single spanning tree is enabled. use "show 802-1w" command.
VLAN is a member of global SSTP - IEEE 802-1w
  PORT 5/1/47 - FORWARDING
  PORT 5/1/48 - FORWARDING
  PORT 7/1/1 - FORWARDING
  PORT 7/1/2 - FORWARDING
  PORT 7/1/3 - FORWARDING
  PORT 8/1/1 - FORWARDING
  PORT 8/1/2 - FORWARDING
  PORT 8/1/3 - FORWARDING
  PORT 8/1/45 - FORWARDING
Client-1(config)#

```

Client-2 configuration

```

Client-2(config)# spanning-tree single 802-1w
Client-2(config)# end
Client-2(config)# show 802-1w vlan 1905
Single spanning tree is enabled. use "show 802-1w" command.
VLAN is a member of global SSTP - IEEE 802-1w
  PORT 1/1/1 - FORWARDING
  PORT 1/1/2 - FORWARDING
  PORT 1/1/3 - FORWARDING
  PORT 2/1/47 - BLOCKING
  PORT 2/1/48 - BLOCKING
  PORT 3/1/1 - FORWARDING
  PORT 3/1/2 - FORWARDING
  PORT 3/1/3 - FORWARDING
Client-2(config)#

```

Example 3: Configure Multiple Spanning Tree (MSTP) on the MCT Clients

MSTP (802.1s) allows multiple VLANs to be managed by a single STP instance, and several VLANs can be mapped to a reduced number of spanning-tree instances. Use MSTP on MCT clients to ensure loop-free topology for one or more VLANs that have a similar Layer 2 topology.

MSTP requires that BPDU flooding be enabled on the MCT Cluster devices. This example is based on the network topology shown in [Figure 35](#) on page 200.

Router-1 configuration

```

Router-1(config)# bpdu-flood-enable
Warning - Any received untagged BPDUs will now be flooded to all the ports.

```

Router-2 configuration

```

Router-2(config)# bpdu-flood-enable
Warning - Any received untagged BPDUs will now be flooded to all the ports.

```

Client-1 configuration

```

Client-1(config)# mstp scope all
Enter MSTP scope would remove STP and topology group related configuration for system
Are you sure? (enter 'y' or 'n'): y
'MSTP Start' need to be entered in order to activate this MSTP feature
Client-1(config)# mstp start
Client-1(config)# mstp instance 1 vlan 1901

```

```
Client-1(config)# mstp instance 1 vlan 1905
Client-1(config)#
```

Client-2 configuration

```
Client-2(config)# mstp scope all
Enter MSTP scope would remove STP and topology group related configuration for system
Are you sure? (enter 'y' or 'n'): y
'MSTP Start" need to be entered in order to activate this MSTP feature
Client-2(config)# mstp start
Client-2(config)# mstp instance 1 vlan 1901
Client-2(config)# mstp instance 1 vlan 1905
Client-2(config)
```

GVRP

• GVRP overview.....	205
• GVRP application examples.....	205
• VLAN names created by GVRP.....	208
• Configuration notes for GVRP.....	208
• GVRP configuration.....	209
• Converting a VLAN created by GVRP into a statically-configured VLAN.....	213
• Displaying GVRP information.....	214
• Clearing GVRP statistics.....	220
• GVRP CLI examples.....	221

GVRP overview

GARP VLAN Registration Protocol (GVRP) is a Generic Attribute Registration Protocol (GARP) application that provides VLAN registration service by means of dynamic configuration (registration) and distribution of VLAN membership information.

A Brocade device enabled for GVRP can do the following:

- Learn about VLANs from other Brocade devices and configure those VLANs on the ports that learn about the VLANs. The device listens for GVRP Protocol Data Units (PDUs) from other devices, and implements the VLAN configuration information in the PDUs.
- Advertise VLANs configured on the device to other Brocade devices. The device sends GVRP PDUs advertising its VLANs to other devices. GVRP advertises statically configured VLANs and VLANs learned from other devices through GVRP.

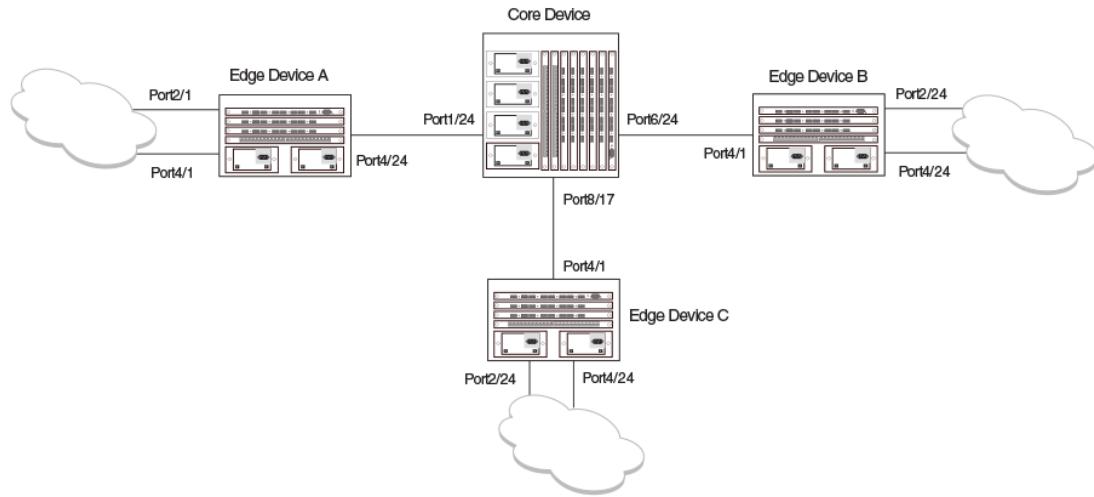
GVRP enables a Brocade device to dynamically create 802.1Q-compliant VLANs on links with other devices that are running GVRP. GVRP reduces the chances for errors in VLAN configuration by automatically providing VLAN ID consistency across the network. You can use GVRP to propagate VLANs to other GVRP-aware devices automatically, without the need to manually configure the VLANs on each device. In addition, if the VLAN configuration on a device changes, GVRP automatically changes the VLAN configurations of the affected devices.

The Brocade implementation of GARP and GVRP is based on the following standards:

- ANSI/IEEE standard 802.1D, 1998 edition
- IEEE standard 802.1Q, 1998 edition; approved December 8, 1998
- IEEE draft P802.1w/D10, March 26, 2001
- IEEE draft P802.1u/D9, November 23, 2000
- IEEE draft P802.1t/D10, November 20, 2000

GVRP application examples

The following figure shows an example of a network that uses GVRP. This section describes various ways you can use GVRP in a network such as this one. [GVRP CLI examples](#) on page 221 lists the CLI commands to implement the applications of GVRP described in this section.

FIGURE 36 Example of GVRP

In this example, a core device is attached to three edge devices. Each of the edge devices is attached to other edge devices or host stations (represented by the clouds).

The effects of GVRP in this network depend on which devices the feature is enabled on, and whether both learning and advertising are enabled. In this type of network (a core device and edge devices), you can have the following four combinations:

- Dynamic core and fixed edge
- Dynamic core and dynamic edge
- Fixed core and dynamic edge
- Fixed core and fixed edge

Dynamic core and fixed edge

In this configuration, all ports on the core device are enabled to learn and advertise VLAN information. The edge devices are configured to advertise their VLAN configurations on the ports connected to the core device. GVRP learning is disabled on the edge devices.

Core device	Edge device A	Edge device B	Edge device C
<ul style="list-style-type: none"> • GVRP is enabled on all ports. • Both learning and advertising are enabled. 	<ul style="list-style-type: none"> • GVRP is enabled on port 4/24. Learning is disabled. • VLAN 20 • Port 2/1 (untagged) • Port 4/24 (tagged) • VLAN 40 • Port 4/1 (untagged) • Port 4/24 (tagged) 	<ul style="list-style-type: none"> • GVRP is enabled on port 4/1. Learning is disabled. • VLAN 20 • Port 2/24 (untagged) • Port 4/1 (tagged) • VLAN 30 • Port 4/24 (untagged) • Port 4/1 (tagged) 	<ul style="list-style-type: none"> • GVRP is enabled on port 4/1. Learning is disabled. • VLAN 30 • Port 2/24 (untagged) • Port 4/1 (tagged) • VLAN 40 • Port 4/24 (untagged) • Port 4/1 (tagged)
NOTE Since learning is disabled on all the edge devices, advertising on the core device has no effect in this configuration.			

In this configuration, the edge devices are statically (manually) configured with VLAN information. The core device dynamically configures itself to be a member of each of the edge device VLANs. The operation of GVRP on the core device results in the following VLAN configuration on the device:

- VLAN 20
 - 1/24 (tagged)
 - 6/24 (tagged)
- VLAN 30
 - 6/24 (tagged)
 - 8/17 (tagged)
- VLAN 40
 - 1/24 (tagged)
 - 8/17 (tagged)

VLAN 20 traffic can now travel through the core between edge devices A and B. Likewise, VLAN 30 traffic can travel between B and C and VLAN 40 traffic can travel between A and C. If an edge device is moved to a different core port or the VLAN configuration of an edge device is changed, the core device automatically reconfigures itself to accommodate the change.

Notice that each of the ports in the dynamically created VLANs is tagged. All GVRP VLAN ports configured by GVRP are tagged, to ensure that the port can be configured for additional VLANs.

NOTE

This example assumes that the core device has no static VLANs configured. However, you can have static VLANs on a device that is running GVRP. GVRP can dynamically add other ports to the statically configured VLANs but cannot delete statically configured ports from the VLANs.

Dynamic core and dynamic edge

GVRP is enabled on the core device and on the edge devices. This type of configuration is useful if the devices in the edge clouds are running GVRP and advertise their VLANs to the edge devices. The edge devices learn the VLANs and also advertise them to the core. In this configuration, you do not need to statically configure the VLANs on the edge or core devices, although you can have statically configured VLANs on the devices. The devices learn the VLANs from the devices in the edge clouds.

Fixed core and dynamic edge

GVRP learning is enabled on the edge devices. The VLANs on the core device are statically configured, and the core device is enabled to advertise its VLANs but not to learn VLANs. The edge devices learn the VLANs from the core.

Fixed core and fixed edge

The VLANs are statically configured on the core and edge devices. On each edge device, VLAN advertising is enabled but learning is disabled. GVRP is not enabled on the core device. This configuration enables the devices in the edge clouds to learn the VLANs configured on the edge devices.

VLAN names created by GVRP

The **show vlans** command lists VLANs created by GVRP as "GVRP_VLAN_vlan-id". VLAN names for statically configured VLANs are not affected. To distinguish between statically-configured VLANs that you add to the device and VLANs that you convert from GVRP-configured VLANs into statically-configured VLANs, the **show vlans** command displays a converted VLAN name as "STATIC_VLAN_vlan-id".

Configuration notes for GVRP

- If you disable GVRP, all GVRP configuration information is lost if you save the configuration change (**write memory** command) and then reload the software. However, if you reload the software without first saving the configuration change, the GVRP configuration is restored following a software reload.
- The maximum number of VLANs supported on a device enabled for GVRP is the same as the maximum number on a device that is not enabled for GVRP.
 - To display the maximum number of VLANs allowed on your device, enter the **show default values** command. See the "vlan" row in the System Parameters section. Make sure you allow for the default VLAN (1), the GVRP base VLAN (4093), and the Single STP VLAN (4094). These VLANs are maintained as "Registration Forbidden" in the GVRP database. Registration Forbidden VLANs cannot be advertised or learned by GVRP.
 - To increase the maximum number of VLANs supported on the device, enter the **system-max vlannum** command at the global CONFIG level of the CLI, then save the configuration and reload the software. The maximum number you can specify is listed in the Maximum column of the **show default values** display.
- The default VLAN (VLAN 1) is not advertised by the Brocade implementation of GVRP. The default VLAN contains all ports that are not members of statically configured VLANs or VLANs enabled for GVRP.

NOTE

The default VLAN has ID 1 by default. You can change the VLAN ID of the default VLAN, but only before GVRP is enabled. You cannot change the ID of the default VLAN after GVRP is enabled.

- Single STP must be enabled on the device. Brocade implementation of GVRP requires Single STP. If you do not have any statically configured VLANs on the device, you can enable Single STP as follows.

```
device(config)#vlan 1
device(config-vlan-1)#exit
device(config)#span
device(config)#span single
```

These commands enable configuration of the default VLAN (VLAN 1), which contains all the device ports, and enable STP and Single STP.

- All VLANs that are learned dynamically through GVRP are added to the single spanning tree.
- All ports that are enabled for GVRP become tagged members of the GVRP base VLAN (4093). If you need to use this VLAN ID for another VLAN, you can change the GVRP VLAN ID. Refer to [Changing the GVRP base VLAN ID](#) on page 209. The software adds the GVRP base VLAN to the single spanning tree.
- All VLAN ports added by GVRP are tagged.

- GVRP is supported only for tagged ports or for untagged ports that are members of the default VLAN. GVRP is not supported for ports that are untagged and are members of a VLAN other than the default VLAN.
- To configure GVRP on a trunk group, enable the protocol on the primary port in the trunk group. The GVRP configuration of the primary port is automatically applied to the other ports in the trunk group.
- You can use GVRP on a device even if the device has statically configured VLANs. GVRP does not remove any ports from the statically configured VLANs, although GVRP can add ports to the VLANs. GVRP advertises the statically configured VLANs. Ports added by GVRP do not appear in the running-config and will not appear in the startup-config file when save the configuration. You can manually add a port to make the port a permanent member of the VLAN. After you manually add the port, the port will appear in the running-config and be saved to the startup-config file when you save the configuration.
- VLANs created by GVRP do not support virtual routing interfaces or protocol-based VLANs. virtual routing interfaces and protocol-based VLANs are still supported on statically configured VLANs even if GVRP adds ports to those VLANs.
- You cannot manually configure any parameters on a VLAN that is created by GVRP. For example, you cannot change STP parameters for the VLAN.
- The GVRP timers (Join, Leave, and Leaveall) must be set to the same values on all the devices that are exchanging information using GVRP.
- If the network has a large number of VLANs, the GVRP traffic can use a lot of CPU resources. If you notice high CPU utilization after enabling GVRP, set the GVRP timers to longer values. In particular, set the Leaveall timer to a longer value. Refer to [Changing the GVRP timers](#) on page 211.
- The feature is supported only on Ethernet ports.

NOTE

If you plan to change the GVRP base VLAN ID (4093) or the maximum configurable value for the Leaveall timer (300000 ms by default), you must do so before you enable GVRP.

GVRP configuration

To configure a device for GVRP, globally enable support for the feature, then enable the feature on specific ports. Optionally, you can disable VLAN learning or advertising on specific interfaces.

You can also change the protocol timers and the GVRP base VLAN ID.

Changing the GVRP base VLAN ID

By default, GVRP uses VLAN 4093 as a base VLAN for the protocol. All ports that are enabled for GVRP become tagged members of this VLAN. If you need to use VLAN ID 4093 for a statically configured VLAN, you can change the GVRP base VLAN ID.

NOTE

If you want to change the GVRP base VLAN ID, you must do so before enabling GVRP.

To change the GVRP base VLAN ID, enter a command such as the following at the global CONFIG level of the CLI.

```
device(config)#gvrp-base-vlan-id 1001
```

This command changes the GVRP VLAN ID from 4093 to 1001.

Syntax: [no] gvrp-base-vlan-id *vlan-id*

The *vlan-id* parameter specifies the new VLAN ID. You can specify a VLAN ID from 2 - 4092 or 4095.

Increasing the maximum configurable value of the Leaveall timer

By default, the highest value you can specify for the Leaveall timer is 300000 ms. You can increase the maximum configurable value of the Leaveall timer to 1000000 ms.

NOTE

You must enter this command before enabling GVRP. Once GVRP is enabled, you cannot change the maximum Leaveall timer value.

NOTE

This command does not change the default value of the Leaveall timer itself. The command only changes the maximum value to which you can set the Leaveall timer.

To increase the maximum value you can specify for the Leaveall timer, enter a command such as the following at the global CONFIG level of the CLI.

```
device(config)#gvrp-max-leaveall-timer 1000000
```

Syntax: [no] gvrp-max-leaveall-timer *ms*

The *ms* parameter specifies the maximum number of ms to which you can set the Leaveall timer. You can specify from 300000 - 1000000 (one million) ms. The value must be a multiple of 100 ms. The default is 300000 ms.

Enabling GVRP

To enable GVRP, enter commands such as the following at the global CONFIG level of the CLI.

```
device(config)#gvrp-enable  
device(config-gvprp)#enable all
```

The first command globally enables support for the feature and changes the CLI to the GVRP configuration level. The second command enables GVRP on all ports on the device.

The following command enables GVRP on ports 1/24, 2/24, and 4/17.

```
device(config-gvprp)#enable ethernet 1/24 ethernet 2/24 ethernet 4/17
```

Syntax: [no] gvrp-enable

Syntax: [no] enable all | ethernet *port* [ethernet *port* | to *port*]

The **all** keyword enables GVRP on all ports.

To specify a list of ports, enter each port as **ethernet *port*** followed by a space. For example, **ethernet 1/24 ethernet 6/24 ethernet 8/17**

To specify a range of ports, enter the first port in the range as **ethernet *port*** followed by the last port in the range. For example, **ethernet 1/1 to 1/8**.

You can combine lists and ranges in the same command. For example: **enable ethernet 1/1 to 1/8 ethernet 1/24 ethernet 6/24 ethernet 8/17**.

Disabling VLAN advertising

To disable VLAN advertising on a port enabled for GVRP, enter a command such as the following at the GVRP configuration level.

```
device(config-gvrp)#block-applicant ethernet 1/24 ethernet 6/24 ethernet 8/17
```

This command disables advertising of VLAN information on ports 1/24, 6/24, and 8/17.

Syntax: [no] **block-applicant all | ethernet port [ethernet port | to port]**

NOTE

Leaveall messages are still sent on the GVRP ports.

The **all** keyword disables VLAN advertising on all ports enabled for GVRP.

To specify a list of ports, enter each port as **ethernet port** followed by a space. For example, **ethernet 1/24 ethernet 6/24 ethernet 8/17**

To specify a range of ports, enter the first port in the range as **ethernet port** followed by the last port in the range. For example, **ethernet 1/1 to 1/8**.

You can combine lists and ranges in the same command. For example: **enable ethernet 1/1 to 1/8 ethernet 1/24 ethernet 6/24 ethernet 8/17**.

Disabling VLAN learning

To disable VLAN learning on a port enabled for GVRP, enter a command such as the following at the GVRP configuration level.

```
device(config-gvrp)#block-learning ethernet 6/24
```

This command disables learning of VLAN information on port 6/24.

NOTE

The port still advertises VLAN information unless you also disable VLAN advertising.

Syntax: [no] **block-learning all | ethernet port [ethernet port | to port]**

The **all** keyword disables VLAN learning on all ports enabled for GVRP.

To specify a list of ports, enter each port as **ethernet port** followed by a space. For example, **ethernet 1/24 ethernet 6/24 ethernet 8/17**

To specify a range of ports, enter the first port in the range as **ethernet port** followed by the last port in the range. For example, **ethernet 1/1 to 1/8**.

You can combine lists and ranges in the same command. For example: **enable ethernet 1/1 to 1/8 ethernet 1/24 ethernet 6/24 ethernet 8/17**.

Changing the GVRP timers

GVRP uses the following timers:

- Join - The maximum number of milliseconds (ms) a device GVRP interfaces wait before sending VLAN advertisements on the interfaces. The actual interval between Join messages is randomly calculated to a value between 0 and the maximum number of milliseconds specified for Join

messages. You can set the Join timer to a value from 200 - one third the value of the Leave timer. The default is 200 ms.

- Leave - The number of ms a GVRP interface waits after receiving a Leave message on the port to remove the port from the VLAN indicated in the Leave message. If the port receives a Join message before the Leave timer expires, GVRP keeps the port in the VLAN. Otherwise, the port is removed from the VLAN. When a port receives a Leave message, the port GVRP state is changed to Leaving. Once the Leave timer expires, the port GVRP state changes to Empty. You can set the Leave timer to a value from three times the Join timer - one fifth the value of the Leaveall timer. The default is 600 ms.

NOTE

When all ports in a dynamically created VLAN (one learned through GVRP) leave the VLAN, the VLAN is immediately deleted from the device's VLAN database. However, this empty VLAN is still maintained in the GVRP database for an amount of time equal to the following. (number-of-GVRP-enabled-up-ports) * (2 * join-timer) While the empty VLAN is in the GVRP database, the VLAN does not appear in the **show vlans** display but does still appear in the **show gvrp vlan all** display.

- Leaveall - The minimum interval at which GVRP sends Leaveall messages on all GVRP interfaces. Leaveall messages ensure that the GVRP VLAN membership information is current by aging out stale VLAN information and adding information for new VLAN memberships, if the information is missing. A Leaveall message instructs the port to change the GVRP state for all its VLANs to Leaving, and remove them unless a Join message is received before the Leave timer expires. By default, you can set the Leaveall timer to a value from five times the Leave timer - maximum value allowed by software (configurable from 300000 - 1000000 ms). The default is 10000.

NOTE

The actual interval is a random value between the Leaveall interval and 1.5 * the Leaveall time or the maximum Leaveall time, whichever is lower.

NOTE

You can increase the maximum configurable value of the Leaveall timer from 300000 ms up to 1000000 ms using the **gvrp-max-leaveall-timer** command. (Refer to [Increasing the maximum configurable value of the Leaveall timer](#) on page 210.)

Timer configuration requirements

- All timer values must be in multiples of 100 ms.
- The Leave timer must be $\geq 3^*$ the Join timer.
- The Leaveall timer must be $\geq 5^*$ the Leave timer.
- The GVRP timers must be set to the same values on all the devices that are exchanging information using GVRP.

Changing the Join, Leave, and Leaveall timers

The same CLI command controls changes to the Join, Leave, and Leaveall timers. To change values to the timers, enter a command such as the following.

```
device(config-gvpr) #join-timer 1000 leave-timer 3000 leaveall-timer 15000
```

This command changes the Join timer to 1000 ms, the Leave timer to 3000 ms, and the Leaveall timer to 15000.

Syntax: [no] **join-timer** *ms* **leave-timer** *ms* **leaveall-timer** *ms*

NOTE

When you enter this command, all the running GVRP timers are canceled and restarted using the new times specified by the command.

Resetting the timers to their defaults

To reset the Join, Leave, and Leaveall timers to their default values, enter the following command.

```
device(config-gvrp)#default-timers
```

Syntax: default-timers

This command resets the timers to the following values:

- Join - 200 ms
- Leave - 600 ms
- Leaveall - 10000 ms

Converting a VLAN created by GVRP into a statically-configured VLAN

You cannot configure VLAN parameters on VLANs created by GVRP. Moreover, VLANs and VLAN ports added by GVRP do not appear in the running-config and cannot be saved in the startup-config file.

To be able to configure and save VLANs or ports added by GVRP, you must convert the VLAN ports to statically-configured ports.

To convert a VLAN added by GVRP into a statically-configured VLAN, add the ports using commands such as the following.

```
device(config)#vlan 22
device(config-vlan-222)#tagged ethernet 1/1 to 1/8
```

These commands convert GVRP-created VLAN 22 containing ports 1/1 through 1/8 into statically-configured VLAN 22.

Syntax: [no] vlan *vlan-id***Syntax: [no] tagged ethernet port [to port | ethernet port]**

Use the same commands to statically add ports that GVRP added to a VLAN.

NOTE

You cannot add the VLAN ports as untagged ports.

NOTE

After you convert the VLAN, the VLAN name changes from "GVRP_VLAN_<vlan-id>" to "STATIC_VLAN_<vlan-id>".

To specify a list of ports, enter each port as **etherne~~t~~port** followed by a space. For example, **etherne~~t~~1/24 etherne~~t~~6/24 etherne~~t~~8/17**

To specify a range of ports, enter the first port in the range as **etherenetport** followed by the last port in the range. For example, **ethernet 1/1 to 1/8**.

You can combine lists and ranges in the same command. For example: **enable ethernet 1/1 to 1/8 ethernet 1/24 ethernet 6/24 ethernet 8/17**.

Displaying GVRP information

You can display the following GVRP information:

- GVRP configuration information
- GVRP VLAN information
- GVRP statistics
- CPU utilization statistics
- GVRP diagnostic information

Displaying GVRP configuration information

To display GVRP configuration information, enter a command such as the following.

```
device#show gvrp
GVRP is enabled on the system
GVRP BASE VLAN ID      : 4093
GVRP MAX Leaveall Timer : 300000 ms
GVRP Join Timer        : 200 ms
GVRP Leave Timer       : 600 ms
GVRP Leave-all Timer   : 10000 ms
=====
Configuration that is being used:
block-learning ethe 1/3
block-applicant ethe 2/7 ethe 2/11
enable ethe 1/1 to 1/7 ethe 2/1 ethe 2/7 ethe 2/11
=====
Spanning Tree: SINGLE SPANNING TREE
Dropped Packets Count: 0
=====
Number of VLANs in the GVRP Database: 15
Maximum Number of VLANs that can be present: 4095
=====
```

Syntax: **show gvrp [ethernet port]**

This display shows the following information.

TABLE 34 CLI display of summary GVRP information

Field	Description
Protocol state	The state of GVRP. The display shows one of the following: <ul style="list-style-type: none"> • GVRP is disabled on the system • GVRP is enabled on the system
GVRP BASE VLAN ID	The ID of the base VLAN used by GVRP.

TABLE 34 CLI display of summary GVRP information (Continued)

Field	Description
GVRP MAX Leaveall Timer	The maximum number of ms to which you can set the Leaveall timer.
NOTE To change the maximum value, refer to Increasing the maximum configurable value of the Leaveall timer on page 210.	
GVRP Join Timer	The value of the Join timer.
NOTE For descriptions of the Join, Leave, and Leaveall timers or to change the timers, refer to Changing the GVRP timers on page 211.	
GVRP Leave Timer	The value of the Leave timer.
GVRP Leave-all Timer	The value of the Leaveall timer.
Configuration that is being used	The configuration commands used to enable GVRP on individual ports. If GVRP learning or advertising is disabled on a port, this information also is displayed.
Spanning Tree	The type of STP enabled on the device.
NOTE The current release supports GVRP only with Single STP.	
Dropped Packets Count	The number of GVRP packets that the device has dropped. A GVRP packet can be dropped for either of the following reasons: <ul style="list-style-type: none"> • GVRP packets are received on a port on which GVRP is not enabled.
NOTE If GVRP support is not globally enabled, the device does not drop the GVRP packets but instead forwards them at Layer 2.	
<ul style="list-style-type: none"> • GVRP packets are received with an invalid GARP Protocol ID. The protocol ID must always be 0x0001. 	
Number of VLANs in the GVRP Database	The number of VLANs in the GVRP database.
NOTE This number includes the default VLAN (1), the GVRP base VLAN (4093), and the single STP VLAN (4094). These VLANs are not advertised by GVRP but are maintained as "Registration Forbidden".	

TABLE 34 CLI display of summary GVRP information (Continued)

Field	Description
Maximum Number of VLANs that can be present	The maximum number of VLANs that can be configured on the device. This number includes statically configured VLANs, VLANs learned through GVRP, and VLANs 1, 4093, and 4094. To change the maximum number of VLANs the device can have, use the system-max vlan num command. Refer to Displaying and modifying system parameter default settings on page 36 Displaying GVRP configuration information .

To display detailed GVRP information for an individual port, enter a command such as the following.

```
device#show gvrp ethernet 2/1
Port 2/1 -
  GVRP Enabled    : YES
  GVRP Learning   : ALLOWED
  GVRP Applicant  : ALLOWED
  Port State      : UP
  Forwarding      : YES

  VLAN Membership:      [VLAN-ID]          [MODE]
                        1                 FORBIDDEN
                        2                 FIXED
                        1001              NORMAL
                        1003              NORMAL
                        1004              NORMAL
                        1007              NORMAL
                        1009              NORMAL
                        1501              NORMAL
                        2507              NORMAL
                        4001              NORMAL
                        4093              FORBIDDEN
                        4094              FORBIDDEN
```

This display shows the following information.

TABLE 35 CLI display of detailed GVRP information for a port

Field	Description
Port number	The port for which information is being displayed.
GVRP Enabled	Whether GVRP is enabled on the port.
GVRP Learning	Whether the port can learn VLAN information from GVRP.
GVRP Applicant	Whether the port can advertise VLAN information into GVRP.
Port State	The port link state, which can be UP or DOWN.
Forwarding	Whether the port is in the GVRP Forwarding state: <ul style="list-style-type: none"> • NO - The port is in the Blocking state. • YES - The port is in the Forwarding state.

TABLE 35 CLI display of detailed GVRP information for a port (Continued)

Field	Description
VLAN Membership	The VLANs of which the port is a member. For each VLAN, the following information is shown: <ul style="list-style-type: none"> • VLAN ID - The VLAN ID. • Mode - The type of VLAN, which can be one of the following: • FIXED - The port will always be a member of this VLAN and the VLAN will always be advertised on this port by GVRP. A port becomes FIXED when you configure the port as a tagged member of a statically configured VLAN. • FORBIDDEN - The VLAN is one of the special VLANs that is not advertised or learned by GVRP. In the current release, the following VLANs are forbidden: the default VLAN (1), the GVRP base VLAN (4093), or the Single STP VLAN (4094). • NORMAL - The port became a member of this VLAN after learning about the VLAN through GVRP. The port membership in the VLAN depends on GVRP. If the VLAN is removed from the ports that send GVRP advertisements to this device, then the port will stop being a member of the VLAN.

Displaying GVRP VLAN information

To display information about all the VLANs on the device, enter the following command.

```
device#show gvrp vlan brief
Number of VLANs in the GVRP Database: 7
Maximum Number of VLANs that can be present: 4095
[VLAN-ID] [MODE] [VLAN-INDEX]
 1 STATIC-DEFAULT 0
 7 STATIC 2
 11 STATIC 4
 1001 DYNAMIC 7
 1003 DYNAMIC 8
 4093 STATIC-GVRP-BASE-VLAN 6
 4094 STATIC-SINGLE-SPAN-VLAN 5
=====
```

Syntax: show gvrp vlan all | brief | *vlan-id*

This display shows the following information.

TABLE 36 CLI display of summary VLAN information for GVRP

Field	Description
Number of VLANs in the GVRP Database	The number of VLANs in the GVRP database.
NOTE This number includes the default VLAN (1), the GVRP base VLAN (4093), and the single STP VLAN (4094). These VLANs are not advertised by GVRP but are included in the total count.	

TABLE 36 CLI display of summary VLAN information for GVRP (Continued)

Field	Description
Maximum Number of VLANs that can be present	The maximum number of VLANs that can be configured on the device. This number includes statically configured VLANs, VLANs learned through GVRP, and VLANs 1, 4093, and 4094. To change the maximum number of VLANs the device can have, use the system-max vlan num command. Refer to Displaying and modifying system parameter default settings on page 36.
VLAN-ID	The VLAN ID.
MODE	The type of VLAN, which can be one of the following: <ul style="list-style-type: none"> • STATIC - The VLAN is statically configured and cannot be removed by GVRP. This includes VLANs you have configured as well as the default VLAN (1), base GVRP VLAN (4093), and Single STP VLAN (4094). • DYNAMIC - The VLAN was learned through GVRP.
VLAN-INDEX	A number used as an index into the internal database.

To display detailed information for a specific VLAN, enter a command such as the following.

```
device#show gvrp vlan 1001
VLAN-ID: 1001, VLAN-INDEX: 7, STATIC: NO, DEFAULT: NO, BASE-VLAN: NO
Timer to Delete Entry Running: NO
Legend: [S=Slot]
Forbidden Members: None
Fixed Members: None
Normal(Dynamic) Members: (S2) 1
```

This display shows the following information.

TABLE 37 CLI display of summary VLAN information for GVRP

Field	Description
VLAN-ID	The VLAN ID.
VLAN-INDEX	A number used as an index into the internal database.
STATIC	Whether the VLAN is a statically configured VLAN.
DEFAULT	Whether this is the default VLAN.
BASE-VLAN	Whether this is the base VLAN for GVRP.
Timer to Delete Entry Running	Whether all ports have left the VLAN and the timer to delete the VLAN itself is running. The timer is described in the note for the Leave timer in Changing the GVRP timers on page 211.
Legend	The meanings of the letter codes used in other parts of the display.

TABLE 37 CLI display of summary VLAN information for GVRP (Continued)

Field	Description
Forbidden Members	The ports that cannot become members of a VLAN advertised or learned by GVRP.
Fixed Members	The ports that are statically configured members of the VLAN. GVRP cannot remove these ports.
Normal(Dynamic) Members	The ports that were added by GVRP. These ports also can be removed by GVRP.
MODE	<p>The type of VLAN, which can be one of the following:</p> <ul style="list-style-type: none"> • STATIC - The VLAN is statically configured and cannot be removed by GVRP. This includes VLANs you have configured as well as the default VLAN (1), base GVRP VLAN (4093), and Single STP VLAN (4094). • DYNAMIC - The VLAN was learned through GVRP.

To display detailed information for all VLANs, enter the **show gvrp vlan all** command.

Displaying GVRP statistics

To display GVRP statistics for a port, enter a command such as the following.

```
device#show gvrp statistics ethernet 2/1
PORT 2/1 Statistics:
Leave All Received : 147
Join Empty Received : 4193
Join In Received : 599
Leave Empty Received : 0
Leave In Received : 0
Empty Received : 588
Leave All Transmitted : 157
Join Empty Transmitted : 1794
Join In Transmitted : 598
Leave Empty Transmitted : 0
Leave In Transmitted : 0
Empty Transmitted : 1248
Invalid Messages/Attributes Skipped : 0
Failed Registrations : 0
```

Syntax: **show gvrp statistics all | ethernet port**

This display shows the following information for the port.

TABLE 38 CLI display of GVRP statistics

Field	Description
Leave All Received	The number of Leaveall messages received.
Join Empty Received	The number of Join Empty messages received.
Join In Received	The number of Join In messages received.
Leave Empty Received	The number of Leave Empty messages received.

TABLE 38 CLI display of GVRP statistics (Continued)

Field	Description
Leave In Received	The number of Leave In messages received.
Empty Received	The number of Empty messages received.
Leave All Transmitted	The number of Leaveall messages sent.
Join Empty Transmitted	The number of Join Empty messages sent.
Join In Transmitted	The number of Join In messages sent.
Leave Empty Transmitted	The number of Leave Empty messages sent.
Leave In Transmitted	The number of Leave In messages sent.
Empty Transmitted	The number of Empty messages sent.
Invalid Messages/Attributes Skipped	The number of invalid messages or attributes received or skipped. This can occur in the following cases: <ul style="list-style-type: none"> • The incoming GVRP PDU has an incorrect length. • "End of PDU" was reached before the complete attribute could be parsed. • The Attribute Type of the attribute that was being parsed was not the GVRP VID Attribute Type (0x01). • The attribute that was being parsed had an invalid attribute length. • The attribute that was being parsed had an invalid GARP event. • The attribute that was being parsed had an invalid VLAN ID. The valid range is 1 - 4095.
Failed Registrations	The number of failed registrations that have occurred. A failed registration can occur for the following reasons: <ul style="list-style-type: none"> • Join requests were received on a port that was blocked from learning dynamic VLANs (GVRP Blocking state). • An entry for a new GVRP VLAN could not be created in the GVRP database.

To display GVRP statistics for all ports, enter the **show gvrp statistics all** command.

Clearing GVRP statistics

To clear the GVRP statistics counters, enter the **clear gvrp statistics all** command.

```
device#clear gvrp statistics all
```

This command clears the counters for all ports. To clear the counters for a specific port only, enter a command such as the following.

```
device#clear gvrp statistics ethernet 2/1
```

Syntax: clear gvrp statistics all | ethernet *port*

GVRP CLI examples

The following sections show the CLI commands for implementing the applications of GVRP described in [GVRP application examples](#) on page 205.

NOTE

Although some of the devices in these configuration examples do not have statically configured VLANs, this is not a requirement. You always can have statically configured VLANs on a device that is running GVRP.

Dynamic core and fixed edge

In this configuration, the edge devices advertise their statically configured VLANs to the core device. The core device does not have any statically configured VLANs but learns the VLANs from the edge devices.

Enter the following commands on the core device.

```
device> enable
device#configure terminal
device(config)#gvrp-enable
device(config-gvprp)#enable all
```

These commands globally enable GVRP support and enable the protocol on all ports.

Enter the following commands on edge device A.

```
device> enable
device#configure terminal
device(config)#vlan 20
device(config-vlan-20)#untag ethernet 2/1
device(config-vlan-20)#tag ethernet 4/24
device(config-vlan-20)#vlan 40
device(config-vlan-40)#untag ethernet 2/1
device(config-vlan-40)#tag ethernet 4/24
device(config-vlan-40)#exit
device(config)#gvrp-enable
device(config-gvprp)#enable ethernet 4/24
device(config-gvprp)#block-learning ethernet 4/24
```

These commands statically configure two port-based VLANs, enable GVRP on port 4/24, and block GVRP learning on the port. The device will advertise the VLANs but will not learn VLANs from other devices.

Enter the following commands on edge device B.

```
device> enable
device#configure terminal
device(config)#vlan 20
device(config-vlan-20)#untag ethernet 2/24
device(config-vlan-20)#tag ethernet 4/1
device(config-vlan-20)#vlan 30
device(config-vlan-30)#untag ethernet 4/24
device(config-vlan-30)#tag ethernet 4/1
device(config-vlan-30)#exit
device(config)#gvrp-enable
device(config-gvprp)#enable ethernet 4/1
device(config-gvprp)#block-learning ethernet 4/1
```

Enter the following commands on edge device C.

```
device> enable
device#configure terminal
device(config)#vlan 30
device(config-vlan-30)#untag ethernet 2/24
device(config-vlan-30)#tag ethernet 4/1
device(config-vlan-20)#vlan 40
device(config-vlan-40)#untag ethernet 4/24
device(config-vlan-40)#tag ethernet 4/1
device(config-vlan-40)#exit
device(config)#gvrp-enable
device(config-gvrp)#enable ethernet 4/1
device(config-gvrp)#block-learning ethernet 4/1
```

Dynamic core and dynamic edge

In this configuration, the core and edge devices have no statically configured VLANs and are enabled to learn and advertise VLANs. The edge and core devices learn the VLANs configured on the devices in the edge clouds. To enable GVRP on all the ports, enter the following command on each edge device and on the core device.

```
device> enable
device#configure terminal
device(config)#gvrp-enable
device(config-gvrp)#enable all
```

Fixed core and dynamic edge

In this configuration, GVRP learning is enabled on the edge devices. The VLANs on the core device are statically configured, and the core device is enabled to advertise its VLANs but not to learn VLANs. The edge devices learn the VLANs from the core.

Enter the following commands on the core device.

```
device> enable
device#configure terminal
device(config)#vlan 20
device(config-vlan-20)#tag ethernet 1/24
device(config-vlan-20)#tag ethernet 6/24
device(config-vlan-20)#vlan 30
device(config-vlan-30)#tag ethernet 6/24
device(config-vlan-30)#tag ethernet 8/17
device(config-vlan-30)#vlan 40
device(config-vlan-40)#tag ethernet 1/5
device(config-vlan-40)#tag ethernet 8/17
device(config-vlan-40)#vlan 50
device(config-vlan-50)#untag ethernet 6/1
device(config-vlan-50)#tag ethernet 1/11
device(config-vlan-50)#exit
device(config)#gvrp-enable
device(config-gvrp)#enable ethernet 1/24 ethernet 6/24 ethernet 8/17
device(config-gvrp)#block-learning ethernet 1/24 ethernet 6/24 ethernet 8/17
```

These VLAN commands configure VLANs 20, 30, 40, and 50. The GVRP commands enable the protocol on the ports that are connected to the edge devices, and disable VLAN learning on those ports. All the VLANs are advertised by GVRP.

Enter the following commands on edge devices A, B, and C.

```
device> enable
device#configure terminal
device(config)#gvrp-enable
device(config-gvrp)#enable all
device(config-gvrp)#block-applicant all
```

Fixed core and fixed edge

The VLANs are statically configured on the core and edge devices. On each edge device, VLAN advertising is enabled but learning is disabled. GVRP is not configured on the core device. This configuration enables the devices in the edge clouds to learn the VLANs configured on the edge devices.

This configuration does not use any GVRP configuration on the core device.

The configuration on the edge device is the same as in [Dynamic core and fixed edge](#) on page 221.

Port mirroring and Monitoring

● Port mirroring and monitoring overview.....	225
● Port mirroring and monitoring configuration.....	225
● Mirroring configuration on a traditional stack.....	228
● ACL-based inbound mirroring.....	229
● MAC address filter-based mirroring.....	233
● VLAN-based mirroring.....	234

Port mirroring and monitoring overview

Port mirroring is a method of monitoring network traffic that forwards a copy of each incoming or outgoing packet from one port on a network switch to another port where the packet can be analyzed. Port mirroring can be used as a diagnostic tool or debugging feature, especially for preventing attacks. Port mirroring can be managed locally or remotely.

You can configure port mirroring, by assigning a port (known as the Monitor port), from which the packets are copied and sent to a destination port (known as the Mirror port). All packets received on the Monitor port or issued from it, are forwarded to the second port. You next attach a protocol analyzer on the mirror port to monitor each segment separately. The analyzer captures and evaluates the data without affecting the client on the original port.

The mirror port may be a port on the same switch with an attached RMON probe, a port on a different switch in the same hub, or the switch processor.

Port mirroring and monitoring configuration

To configure port monitoring, first specify the mirror port, then enable monitoring on the monitored port.

The *mirror port* is the port to which the monitored traffic is copied. Attach your protocol analyzer to the mirror port. The monitored port is the port with the traffic you want to monitor.

The following table lists the number of mirror and monitor ports supported on the Brocade devices. For more information about port regions, refer to [About port regions](#) on page 17.

TABLE 39 Number of mirror and monitored ports supported

Maximum number supported			
Port Type	FastIron X Series	FCX	ICX (6450, 6430, 6610, 6650, 7450, 7250, 7750)
Ingress mirror ports	1 per port region	1 per port region	1 per port region
Egress mirror ports	1 per port region	1 per port region	1 per port region
Ingress monitored ports	No limit	No limit	No limit

TABLE 39 Number of mirror and monitored ports supported (Continued)

Maximum number supported	Egress monitored ports	8	8

NOTE

FastIron X Series devices support multiple ingress and egress mirror ports. For 1 Gbps ports, ports in groups of 12 share one ingress mirror port and one egress mirror port. Therefore, ports 1 and 2 cannot have different mirror ports, but ports 1 and 13 can (port 25, for third-generation modules). Each 10 Gbps port (or each third-generation 10 GbE module) can have one ingress mirror port and one egress mirror port.

NOTE

For devices other than the FastIron X Series, it is possible to configure more than eight egress ports, although only the first eight are operational. This is also true for mirrored VLANs - more than eight can be configured, but only the first eight are operational.

Configuration notes for port mirroring and monitoring

Refer to the following guidelines when configuring port mirroring and monitoring:

- If you configure both ACL mirroring and ACL-based rate limiting on the same port, then all packets that match are mirrored, including the packets that exceed the rate limit.
- ICX and FCX Series devices support sFlow and port monitoring together on the same port.
- FastIron X Series devices support port monitoring and sFlow together on the same device. The caveat is that port monitoring and sFlow cannot be configured together within the same port region. Refer to [About port regions](#) on page 17 for a list of valid port ranges on these devices. This restriction only applies to first- and second-generation modules.
- You can configure a mirror port specifically as an ingress port, an egress port, or both.
- Mirror ports can run at any speed and are not related to the speed of the ingress or egress monitored ports.
- The same port cannot be both a monitored port and the mirror port.
- The same port can be monitored by one mirror port for ingress traffic and another mirror port for egress traffic.
- The mirror port cannot be a trunk port.
- The monitored port and its mirror port do not need to belong to the same port-based VLAN:
 - If the mirror port is in a *different* VLAN from the monitored port, the packets are tagged with the monitor port VLAN ID. This does not apply if the mirror port resides on the SX-FI48GPP module. In this case, mirrored packets are not tagged with a monitor port VLAN ID.
 - If the mirror port is in the *same* VLAN as the monitored port, the packets are tagged or untagged, depending on the mirror port configuration.
- More than one monitored port can be assigned to the same mirror port.
- If the primary interface of a trunk is enabled for monitoring, the entire trunk is monitored. You can also enable an individual trunk port for monitoring using the **config-trunk-ind** command.
- For *stacked* devices, if the ingress and egress analyzer ports are always network ports on the local device, each device may configure the ingress and egress analyzer port independently. However, if you need to mirror to a remote port, then only one ingress and one egress analyzer port are supported for the entire system.
- For ingress ACL mirroring, the ingress rule for stacked devices also applies. The analyzer port setting command **acl-mirror-port** must be specified for each port, even though the hardware only

supports one port per device. This applies whether the analyzer port is on the local device or on a remote device. For example, when port mirroring is set to a remote device, any mirroring-enabled ports (ACL, MAC address filter, or VLAN) enabled ports are set globally to a single analyzer port, as shown in the following example.

```
device(config)# mirror ethernet 1/1/24
device(config)# mirror ethernet 2/1/48
device(config)# interface ethernet 1/1/1
device(config-if-e1000-1/1/1)# monitor ethernet 2/1/48 both
```

The analyzer port (2/1/48) is set to all devices in the system.

```
device(config)# interface ethernet 1/1/2
device(config-if-e1000-1/1/2)# ip access-group 101 in
device(config-if-e1000-1/1/2)# interface ethernet 1/1/1
device(config-if-e1000-1/1/1)# acl-mirror-port ethernet 2/1/48
```

The previous command is required even though the analyzer port is already set globally by the port mirroring command.

```
device(config)# interface ethernet 1/1/3
device(config-if-e1000-1/1/3)# ip access-group 101 in
device(config-if-e1000-1/1/3)# acl-mirror-port ethernet 2/1/48
device(config-if-e1000-1/1/3)# ip access-group 102 in
```

- Because of hardware limitations on the SX-FI48GPP interface module, if the monitored port is on the SX-FI48GPP module, mirrored packets vary slightly from original (monitored) packets, depending on the type of management or interface module on which the mirror port (analyzer) is configured:
 - When ingress or egress mirroring is enabled between a monitored port on an SX-FI48GPP interface module and a mirror port on an IPv4 interface module, mirrored packets contain CRC errors. All other data in the original and mirrored packets is the same. Port mirroring on SX chassis between third generation and first generation modules is not supported. If the first generation module has the mirror port and the third generation module is being monitored, the mirrored packets will contain CRC errors.
 - When ingress or egress mirroring is enabled between a monitored port on an SX-FI48GPP interface module and a mirror port on a 0 port XL, SX-FIZMR-XL/XL-PREM6 management modules, IPv6 interface module, or SX-FI48GPP module, mirrored packets are larger than the original packet by 4 bytes, but do not contain CRC errors. All other data in the original and mirrored packets is the same.

Command syntax for port mirroring and monitoring

This section describes how to configure port mirroring and monitoring.

Monitoring a port

To configure port monitoring on an individual port on a Brocade device, enter commands similar to the following.

```
device(config)#mirror-port ethernet 1/2/4
device(config)#interface ethernet 1/2/11
device(config-if-e1000-11)#monitor ethernet 1/2/4 both
```

Syntax: [no] mirror-port ethernet port [input | output]

Syntax:[no] monitor ethernet port both | in | out

The port variable for **mirror-port ethernet** specifies the port to which the monitored traffic is copied. The port variable for **monitor ethernet** specifies the port on which traffic is monitored.

The **input** and **output** parameters configure the mirror port exclusively for ingress or egress traffic. If you do not specify one, both types of traffic apply.

The **both**, **in**, and **out** parameters specify the traffic direction you want to monitor on the mirror port. There is no default.

To display the port monitoring configuration, enter the **show monitor** and **show mirror** commands.

Monitoring an individual LAG port

You can monitor the traffic on an individual port of a static LAG group, and on an individual port of an LACP LAG group.

By default, when you monitor the primary port in a LAG group, aggregated traffic for all the ports in the LAG group is copied to the mirror port. You can configure the device to monitor individual ports in a LAG group. You can monitor the primary port or a secondary port individually.

To configure port monitoring on an individual port in a LAG group, enter commands such as the following.

```
device(config)#lag automation static id 1
device(config-lag-automation)#ports e 1/1/2 to 1/1/9
device(config-lag-automation)#primary-port 1/1/2
device(config-lag-automation)#deploy
device(config)#mirror-port ethernet 1/1/1
device(config)#lag automation
device(config-lag-automation)#monitor ethe-port-monitored 1/1/2 ethernet 1/1/1 both

device#sh mirror
Mirror port 1/1/1
  Input monitoring      : (U1/M1)    1
  Output monitoring     : (U1/M1)    1

device#sh mirror ethernet 1/1/1
Mirror port 1/1/1
  Input monitoring      : (U1/M1)    1
  Output monitoring     : (U1/M1)    1

device#sh run | i mirror
mirror-port ethernet 1/1/1
device#sh run | i monitor ethe
monitor ethe-port-monitored 1/1/2 ethe 1/1/1 both
```

Traffic on LAG port e 1/1/2 is monitored, and the monitored traffic is copied to port e 1/1/1, the mirror port.

Mirroring configuration on a traditional stack

You can configure mirroring on a Brocade traditional stack. A traditional stack consists of up to eight FCX and ICX devices. The stack operates as a chassis. The following examples show how to configure mirroring for ports that are on different members of a stack, and for ports that are on the same stack member as the mirror port.

Configuration notes for traditional stack mirroring

The following mirroring configuration information applies to all FCX and ICX devices connected in a traditional stack topology:

- The input or output mirroring port can be on different ports.
- All FastIron devices can have one mirroring port that monitors multiple ports, but cannot have multiple mirror ports for one monitored port.

- If the mirror port and the monitored ports are on different stack units, only one active mirror port is allowed for the entire traditional stack.
- If the mirror port and the monitored ports are on the same port region, multiple active mirror ports are allowed for the entire traditional stack. Devices in a traditional stack support 24 ports per port region.
- The maximum number of monitored VLANs on a traditional stack is 8.

Configuring mirroring for ports on different members in a traditional stack example

In this example, although two ports are configured as active ports, only one active mirror port (port 1/1/24) is allowed for the entire stack because the mirror ports and the monitored ports are on different stack members.

```
device(config)#mirror-port ethernet 1/1/24
device(config)#mirror-port ethernet 2/1/24
device(config)#interface ethernet 1/1/1
device(config-if-e1000-1/1/1)#monitor ethernet 1/1/24 both
device(config)#interface ethernet 2/1/1
device(config-if-e1000-2/1/1)#monitor ethernet 1/1/24 both
device(config)#interface ethernet 4/1/1
device(config-if-e1000-4/1/1)#monitor ethernet 1/1/24 both
```

Configuring mirroring for ports on the same stack member in a traditional stack example

In this example, the mirror ports are assigned to different monitor ports.

```
device(config)#mirror-port ethernet 1/1/24
device(config)#mirror-port ethernet 2/1/24
device(config)#mirror-port ethernet 3/1/24
device(config)#mirror-port ethernet 4/1/24
device(config)#interface ethernet 1/1/1
device(config-if-e1000-1/1/1)#monitor ethernet 1/1/24 both
device(config)#interface ethernet 2/1/1
device(config-if-e1000-2/1/1)#monitor ethernet 2/1/24 both
device(config)#interface ethernet 4/1/1
device(config-if-e1000-4/1/1)#monitor ethernet 4/1/24 both
```

ACL-based inbound mirroring

This section describes ACL-based inbound mirroring for FastIron devices.

Creating an ACL-based inbound mirror clause for FCX and ICX devices

The following example shows how to configure an ACL-based inbound mirror clause.

1. Configure the mirror port.

```
device(config)#mirror-port ethernet 1/1/2
```

2. Configure the ACL-based inbound mirror clause.

```
device(config)#access-list 101 permit ip any any mirror
```

3. Apply the ACL-based inbound clause to the monitor port.

```
device(config)#interface e 1/1/5
device(config-if-e1000-1/1/5)#ip access-group 101 in
```

4. Create the ACL mirror port.

```
device(config-if-e1000-1/1/5)#acl-mirror-port ethernet 1/1/2
```

To display ACL mirror settings, enter the **show access-list all** command.

```
device#show access-list all
Extended IP access list 101
permit ip any any mirror
```

ACL-based inbound mirror clauses for FastIron X Series devices

The **mirror** parameter in an ACL clause causes the system to direct traffic that meets the clause to be sent to a mirror port. Consider the following example.

```
device(config)#access-list 101 permit ip any any mirror
```

The **mirror** parameter directs selected traffic to the mirrored port. Traffic mirroring is only supported on Layer 3 ACLs for FastIron X Series devices.

You can select traffic to be mirrored using a **permit** or **deny** clause on ports on the following interface modules:

- SX-FI-24GPP
- SX-FI-24HF
- SX-FI-2XG
- SX-FI-8XG

On all other interface modules, you can select traffic to be mirrored using only a **permit** clause.

Destination mirror port

You can specify physical ports or a trunk to mirror traffic. If you complete the rest of the configuration but do not specify a destination mirror port, the port-mirroring ACL is non-operational. This can be useful if you want to be able to mirror traffic by a set criteria on demand. With this configuration, you configure a destination mirror port whenever you want the port-mirroring ACL to become operational.

The following sections describe how to specify a destination port for a port or a trunk, as well as the special considerations required when mirroring traffic from a virtual interface.

Specifying the destination mirror port for physical ports

When you want traffic that has been selected by ACL-based inbound mirroring to be mirrored, you must configure a destination mirror port. This configuration is performed at the interface configuration level of the port with the traffic you are mirroring. The destination port must be the same for all ports in a port region as described in [Ports from a port region must be mirrored to the same destination mirror port](#).

In the following example, ACL mirroring traffic from port 1/1 is mirrored to port 1/3.

```
device(config)#interface ethernet 1/1
device(config-if-e1000e-1/1)#ACL-mirror-port ethernet 1/3
```

Syntax: [no] ACL-mirror-port ethernet port

The port variable specifies the mirror port to which the monitored port traffic is copied.

Ports from a port region must be mirrored to the same destination mirror port

Port regions, as described in [About port regions](#) on page 17, are important when defining a destination mirror port. This is because all traffic mirrored from any single port in a port region is mirrored to the

same destination mirror port as traffic mirrored from any other port in the same port region. For example, ports 1/1 to 1/12 are in the same port region. If you configure ports 1/1 and 1/2 to mirror their traffic, they should use the same destination mirror port as shown in the following configuration.

```
device(config)#interface ethernet 1/1
device(config-if-e10000-1/1)#ACL-mirror-port ethernet 2/3
device(config)#interface ethernet 1/2
device(config-if-e10000-1/2)#ACL-mirror-port ethernet 2/3
```

If ports within the same port region are mirrored to different destination ports, the configuration is disallowed, and an error message is generated, as shown in the following example.

```
device(config)#interface ethernet 1/1
device(config-if-e10000-1/1)#ACL-mirror-port ethernet 4/3
device(config)#interface ethernet 1/2
device(config-if-e10000-1/2)#ACL-mirror-port ethernet 4/7
Error - Inbound Mirror port 4/3 already configured for port region 1/1 - 1/12
```

When a destination port is configured for any port within a port region, traffic from any ACL with a mirroring clause assigned to any port in that port region is mirrored to that destination port. This will occur even if a destination port is not explicitly configured for the port with the ACL configured. In the following example, an ACL with a mirroring clause (101) is applied to a port (1/1). Another port in the same region (1/3) has a destination port set (4/3). In this example, traffic generated from operation of ACL 101 is mirrored to port 4/3 even though a destination port has not explicitly been defined for traffic from port 1/1.

```
device(config)#interface ethernet 1/1
device(config-if-e10000-1/1)#ip access-group 101 in
device(config)#interface ethernet 1/3
device(config-if-e10000-1/3)#ACL-mirror-port ethernet 4/3
```

NOTE

If a destination mirror port is not configured for any ports within the port region where the port-mirroring ACL is configured, the ACL does not mirror the traffic but the ACL is applied to traffic on the port.

Specifying the destination mirror port for trunk ports

You can mirror the traffic that has been selected by ACL-based inbound mirroring from a trunk by configuring a destination port for the primary port within the trunk configuration, as shown in the following example.

```
device(config)#trunk ethernet 1/1 to 1/4
device(config)#interface ethernet 1/1
device(config-if-e10000-1/1)#ACL-mirror-port ethernet 1/8
```

Using this configuration, all trunk traffic is mirrored to port 1/8.

Limitations when configuring ACL-based mirroring with trunks

The **config-trunk-ind** command cannot operate with ACL-based mirroring:

- If a trunk is configured with the **config-trunk-ind** command, ACL-based mirroring will not be allowed.
- If the **config-trunk-ind** command is added to a trunk, any ports that are configured for ACL-based mirroring will have monitoring removed and the following message is displayed.

Trunk port monitoring, if any, has been removed.

If an individual port is configured for ACL-based mirroring, you cannot add it to a trunk. If you try to add a port that is configured for ACL-based mirroring to a trunk, the following message appears.

Note – ACL-mirror-port configuration is removed from port 2 in new trunk.

NOTE

If you want to add a port configured for ACL-based mirroring to a trunk, you must first remove the **ACL-mirror-port** command from the port configuration. You can then add the port to a trunk that can then be configured for ACL-based trunk mirroring.

Behavior of ACL-based mirroring when deleting trunks

If you delete a trunk that has ACL-based mirroring configured, the ACL-based mirroring configuration is configured on the individual ports that made up the trunk.

For example, if a trunk is configured as shown in the following example and is then deleted from the configuration as shown, each of the ports that previously was contained in the trunk is configured for ACL-based mirroring.

```
device(config)#lag test static id 111
device(config-lag-test)#ports ethernet 1/1/1 to 1/1/2
device(config-lag-test)#primary-port 1/1/1
device(config-lag-test)#deploy
device(config-if-e-1/1/1)#acl-mirror-port ethernet 1/1/38
```

To delete the trunk, enter the following command.

```
device(config)#no lag test
```

Configuring ACL-based mirroring for ACLs bound to virtual interfaces

For configurations that have an ACL configured for ACL-based mirroring bound to a virtual interface, you must use the **ACL-mirror-port** command on a physical port that is a member of the same VLAN as the virtual interface. Additionally, only traffic that arrives at ports that belong to the same port group as the physical port where the **ACL-mirror-port** command has been used is mirrored. This follows the same rules described in [Ports from a port region must be mirrored to the same destination mirror port](#) on page 230.

For example, in the following configuration, ports 4/1, 4/2, and 5/3 are in VLAN 10 with ve 10. Ports 4/1 and 4/2 belong to the same port group, while port 5/3 belongs to another port group.

```
device(config)#vlan 10
device(config-vlan-10)#tagged ethernet 4/1 to 4/2
device(config-vlan-10)#tagged ethernet 5/3
device(config-vlan-10)#router-interface ve 10
device(config)#interface ethernet 4/1
device(config-if-e1000e-4/1)#ACL-mirror-port ethernet 5/1
device(config)#interface ve 10
device(config-vif-10)#ip address 10.10.10.254/24
device(config-vif-10)#ip access-group 102 in
device(config)#access-list 102 permit ip any any mirror
```

In this configuration, the **ACL-mirror-port** command is applied to port 4/1, which is a member of ve 10. Because of this, ACL-based mirroring will only apply to VLAN 10 traffic that arrives on ports 4/1 and 4/2. It will not apply to VLAN 10 traffic that arrives on port 5/3 because that port belongs to a port group different from ports 4/1 and 4/2. This is because if you apply ACL-based mirroring on an entire VE, and enable mirroring in only one port region, traffic that is in the same VE but on a port in a different port region will not be mirrored.

To make the configuration apply ACL-based mirroring to VLAN 10 traffic arriving on port 5/3, you must add the following commands to the configuration.

```
device(config)#interface ethernet 5/3
device(config-if-e10000-5/3)#ACL-mirror-port ethernet 5/1
```

If a port is in both mirrored and non-mirrored VLANs, only traffic on the port from the mirrored VLAN is mirrored. For example, the following configuration adds VLAN 20 to the previous configuration. In this example, ports 4/1 and 4/2 are in both VLAN 10 and VLAN 20. ACL-based mirroring is only applied to VLAN 10. Consequently, traffic that is on ports 4/1 and 4/2 that belongs to VLAN 20 will not be mirrored.

```
device(config)#vlan 10
device(config-vlan-10)#tagged ethernet 4/1 to 4/2
device(config-vlan-10)#tagged ethernet 5/3
device(config-vlan-10)#router-interface ve 10
device(config)#vlan 20
device(config-vlan-20)#tagged ethernet 4/1 to 4/2
device(config)#interface ethernet 4/1
device(config-if-e10000-4/1)#ACL-mirror-port ethernet 5/1
device(config)#interface ve 10
device(config-vif-10)#ip address 10.10.10.254/24
device(config-vif-10)#ip access-group 102 in
device(config)#access-list 102 permit ip any any mirror
```

MAC address filter-based mirroring

NOTE

The MAC address filter-based mirroring feature is not supported on FastIron X Series devices.

This feature allows traffic entering an ingress port to be monitored from a mirror port connected to a data analyzer, based on specific source and destination MAC addresses. This feature supports mirroring of inbound traffic only. Outbound mirroring is not supported.

MAC-filter-based mirroring allows a user to specify a particular stream of data for mirroring using a filter. This eliminates the need to analyze all incoming data to the monitored port. To configure MAC-filter-based mirroring, the user must perform three steps:

1. Define a mirror port
2. Create a MAC address filter with a mirroring clause
3. Apply the MAC address filter to an interface

Configuring MAC address filter-based mirroring

Complete the following steps to configure MAC address filter-based mirroring.

1. Defining a mirror port

To activate mirroring on a port, use the **mirror** command in global configuration mode.

```
device(config)#mirror ethernet 1/1/14
```

Configuration notes for defining a mirror port

- If there is no input mirror port configured, MAC-filter based mirroring does not take effect. It remains in the configuration, but is not activated.
- MAC-filter-based mirroring can be enabled on a port at the same time as either port-based mirroring or VLAN-based mirroring. When port-based mirroring and MAC-filter-based mirroring are enabled on a port at the same time, the preference order is port-based mirroring followed by MAC-based filtering. When VLAN-based mirroring and MAC-filter-based mirroring are enabled on a port at the same time, the preference order is VLAN-based mirroring and MAC-filter-based mirroring.

NOTE

Port-based mirroring and VLAN-based mirroring can not be enabled on a port at the same time.

2. Creating a MAC address filter with a mirroring clause

The **mirror** keyword is added to MAC address filter clauses to direct desired traffic to the mirror port. In the following example, the MAC address filter directs traffic to a mirror port.

```
device(config)#mac filter 1 permit 0000.0011.2222 ffff.ffff.ffff 0000.0022.3333  
ffff.ffff.fff mirror
```

In this example, any flow matching the source address (SA) 0000.0011.2222 and the destination address (DA) 0000.0022.3333 is mirrored. Other flows are not mirrored.

3. Applying the MAC address filter to an interface

Apply the MAC address filter to an interface using the **mac-filter-group** command.

```
device(config)#interface ethernet 0/1/1  
device(config-if-e10000-0/1/1)#mac filter-group 1
```

4. Configuring the monitor port to use the mirror port

```
device(config)#interface ethernet 0/1/5  
device(config-if-e10000-0/1/5)#acl-mirror-port ethernet 0/1/14
```

VLAN-based mirroring

NOTE

VLAN-based mirroring is supported on FastIron X Series devices equipped with third generation or later modules, ICX 6650, ICX 7750, ICX 7450 and ICX 7250 devices. Refer to [VLAN-based mirroring](#) on page 236.

The VLAN-based mirroring feature allows users to monitor all incoming traffic in one or more VLANs by sending a mirror image of that traffic to a configured mirror port. This feature meets the requirements of CALEA (Communications Assistance for Law Enforcement Act of 1994).

Configuring VLAN-based mirroring

Configure VLAN-based mirroring using the **monitor ethernet** command in VLAN configuration mode. For example, to enable mirroring on VLANs 10 and 20, to mirror port e 1/1/21, enter the following commands.

```
device(config)#mirror-port ethernet 1/1/21 input
device(config)#vlan 10
device(config-VLAN-10)#monitor ethernet 1/1/21
device(config-VLAN-10)#exit
device(config)#vlan 20
device(config-VLAN-20)#monitor ethernet 1/1/21
device(config-VLAN-20)#end
```

Syntax:[no] **monitor ethernet port**

NOTE

For FCX and ICX 6650 devices, because it is possible to have multiple mirror ports, monitor ports must specify which mirror port they are monitoring.

To disable mirroring on VLAN 20, enter the following commands.

```
device(config)#vlan 20
device(config-VLAN-20)#no monitor ethernet 1/1/21
device(config-VLAN-20)#end
```

Displaying VLAN-based mirroring status

The **show vlan** command displays the VLAN-based mirroring status.

```
device#show vlan
Total PORT-VLAN entries: 4
Maximum PORT-VLAN entries: 4060
Legend: [Stk=Stack-Unit, S=Slot]
PORT-VLAN 1, Name DEFAULT-VLAN, Priority level0, Spanning tree On
Untagged Ports: (Stk0/S1) 3 4 5 6 7 8 9 10 11 12 13 14
Untagged Ports: (Stk0/S1) 15 16 17 18 19 20 21 22 23 24 25 26
Untagged Ports: (Stk0/S1) 27 28 29 30 31 32 33 34 35 36 37 38
Untagged Ports: (Stk0/S1) 39 40 41 42 43 44 45 46 47 48
Untagged Ports: (Stk0/S2) 1 2
    Tagged Ports: None
    Uplink Ports: None
DualMode Ports: None
Mac-Vlan Ports: None
    Monitoring: Disabled
PORT-VLAN 10, Name [None], Priority level0, Spanning tree On
Untagged Ports: (Stk0/S1) 1
    Tagged Ports: None
    Uplink Ports: None
DualMode Ports: None
Mac-Vlan Ports: None
    Monitoring: Enabled
PORT-VLAN 20, Name [None], Priority level0, Spanning tree On
Untagged Ports: (Stk0/S1) 2
    Tagged Ports: None
    Uplink Ports: None
DualMode Ports: None
Mac-Vlan Ports: None
    Monitoring: Disabled
```

Configuration notes for VLAN-based mirroring

The following guidelines apply to VLAN-based mirroring configurations:

- A VLAN must have at least one port member configured before monitoring can be configured.
- Multiple VLANs can have monitoring enabled at the same time, and the maximum number of monitor-configured VLANs is 8.
- The mirror port is subject to the same scheduling and bandwidth management as the other ports in the system. If the amount of traffic being sent to the mirror port exceeds the available bandwidth, some of that traffic may be dropped.
- All incoming traffic (tagged and untagged) in the VLAN is mirrored. Mirroring is "as-is", and is not affected by the configuration of the mirror port itself. Incoming tagged traffic is sent out tagged and incoming untagged traffic is sent out untagged, regardless of which VLANs the mirror port belongs to, and whether the mirror port is tagged or untagged.
- VLAN-based mirroring is supported on Layer 2 and Layer 3 images.

VLAN-based mirroring

With the new FastIron X Series of modules, the sFlow processing has been separated from the packet mirroring functionality. This allows for support of VLAN-based mirroring on the FastIron X Series devices. The packet processor on the FastIron X Series of modules also allows for egress VLAN-based mirroring. In order for VLAN-based monitoring to function, the FastIron X Series must have only the following SX modules installed.

The following interface modules are new to the FastIron X Series:

- SX-FI-24GPP
- SX-FI-24HF
- SX-FI-2XG
- SX-FI-8XG
- SX-FI48GPP

NOTE

Egress VLAN-based mirroring is not currently supported on the stacking platforms.

The FastIron X Series of modules are capable of monitoring 4096 VLANs. In a chassis environment, this introduces restrictions to the number of ports that can be configured as mirror ports. Because a single VLAN can contain 384 untagged ports (24 per slot) if that VLAN is configured for monitoring, every device must have an identical number of corresponding analyzer ports. However, the egress mirror-port and ingress mirror-port do not have to be the same. You can use two separate ports.

This introduces restrictions on port-based mirroring coexisting with VLAN-based mirroring. Port-based mirroring allows for multiple ports to be configured as mirror-ports. However, once a particular port belonging to a particular device is configured for monitoring to a specific mirror-port, no other mirror-port can be used to monitor any other port on that device. This restriction has been extended to VLAN-based mirroring, with one caveat: only one mirror-port in either direction at a time can exist within the system. Refer to [About port regions](#) on page 17 for a list of valid port ranges on these devices.

VLAN-based mirroring is also supported on ICX 6650 devices.

Restrictions and capabilities of VLAN-based mirroring

The following is a list of restrictions and capabilities:

- Only the modules that support VLAN-based mirroring should be installed.
- There can be only one input or output mirror-port configured in the system at a time.
- The amount of traffic mirrored is limited by the bandwidth of the mirror-port.
- The maximum amount of egress traffic that can be mirrored is further limited by the bandwidth of the loopback port, which is 10 Gbps.

- The monitored VLAN must be created in hardware.
- An ingress or egress mirror-port must be configured when monitoring the ingress or egress VLAN traffic.
- A maximum of 4096 VLANs can be monitored at a time.
- A VLAN can be monitored for ingress and egress traffic concurrently.
- Port mirroring can be configured concurrently with VLAN-based mirroring, but only one mirror-port can be used for both.
- sFlow can be enabled concurrently with VLAN-based mirroring and port mirroring.
- VLAN-based mirroring is supported on the default VLAN. If the default VLAN is changed dynamically, the configuration is not lost.
- VLAN-based mirroring on VLAN groups is not supported, but it is supported on topology groups.
- In the case of enabling VLAN-based monitoring on the interface modules in an MCT-enabled chassis, the VLAN configuration is not synced across the cluster. Each chassis in the cluster is configured independently for VLAN configuration.

One of the concerns about VLAN-based mirroring is the effects of ingress and egress ACLs, as well as rate shaping and rate limiting, on mirrored packets:

- Ingress VLAN-based mirroring: Any packets that are coming in from the network on the VLAN should be mirrored out. Any ingress ACL actions or rate limiting actions do not take precedence in this case.
- Egress VLAN-based mirroring: Any packets that are sent out onto the network are not affected by egress ACLs or rate shaping.

See the following table for a summary of the effects of ACLs and rate limiting.

TABLE 40 ACL and rate limiting effects

ACL profile	Ingress result	Egress result
Ingress ACL on port	Packets ingress mirrored at expected (sent) rate	Packets egress mirrored at expected (sent) rate
Egress ACL on port	Packets ingress mirrored at expected (sent) rate	Packets egress mirrored at expected (sent) rate
Ingress rate limiting on port	Packets ingress mirrored at expected (sent) rate	Packets egress mirrored at the limited rate
Egress rate shaping on port	Packets ingress mirrored at expected (sent) rate	Packets egress mirrored at expected (sent) rate

Tagged versus untagged ports in VLANs

The following table describes how VLAN-based mirroring coexists with port mirroring, assuming a VLAN consisting of one ingress and one egress port.

TABLE 41 VLAN-based mirroring and port mirroring effects

Packets sent	VLAN-based mirroring direction	Which port in VLAN is monitored	Ingress mirror traffic count expectation	Egress mirror traffic count expectation
10000	Ingress	Ingress	~10000	~0
10000	Egress	Egress	~0	~20000

TABLE 41 VLAN-based mirroring and port mirroring effects (Continued)

Packets sent	VLAN-based mirroring direction	Which port in VLAN is monitored	Ingress mirror traffic count expectation	Egress mirror traffic count expectation
10000	Ingress	Egress	~10000	~10000
10000	Egress	Ingress	~10000	~10000

The above table can be summarized into the following two rules, assuming that VLAN-based mirroring and port mirroring are operating concurrently.

- If the VLAN is ingress monitored, and ports belonging to the VLAN are also ingress monitored, the ingress traffic is only mirrored once and there are no duplicated mirrored packets.
- If the VLAN is egress monitored and ports belonging to the VLAN are also egress monitored, the egress traffic is mirrored for each egress port, as well as the VLAN, resulting in several duplicated mirrored packets. The count of duplicate packets is computed as (1 + Number of egress mirrored ports in the VLAN) * Number of egress packets.

In addition to the previously mentioned rules, the behavior of VLAN-based mirroring is affected by the tag type of the ports belonging to that VLAN, as shown in the following table.

TABLE 42 VLAN-based mirroring behavior: Tagged versus untagged ports

Ingress tag type	Egress tag type	VLAN-based mirroring direction	Mirrored traffic tag type
Untagged	Untagged	Ingress	Untagged
Untagged	Untagged	Egress	Tagged
Untagged	Tagged	Ingress	Untagged
Untagged	Tagged	Egress	Tagged
Tagged	Untagged	Ingress	Tagged
Tagged	Untagged	Egress	Tagged
Tagged	Tagged	Ingress	Tagged
Tagged	Tagged	Egress	Tagged

As illustrated in the above table, regardless of the egress port tag type, if a VLAN is egress mirrored, the mirrored traffic is always tagged. This functionality is built into the hardware and is the expected behavior, as the tag acts as an identifier for monitored traffic on the mirror-port (in case the mirror-port belongs to some other VLAN).

Configuring VLAN-based mirroring

Configure VLAN-based mirroring using the **monitor ethernet** command in VLAN configuration mode. For example, to enable incoming mirroring on VLANs 10 and 20, to mirror port e 6/24, enter the following commands.

```
device(config)#mirror-port ethernet 6/24
device(config)#vlan 10
device(config-VLAN-10)#monitor ethernet 6/24 input
device(config-VLAN-10)#exit
```

```
device(config)#vlan 20
device(config-VLAN-20)#monitor ethernet 6/24 input
device(config-VLAN-20)#end
```

For example, to enable outgoing mirroring on VLANs 10 and 20, to mirror port e 6/24, enter the following commands.

```
device(config)#mirror-port ethernet 6/24
device(config)#vlan 10
device(config-VLAN-10)#monitor ethernet 6/24 output
device(config-VLAN-10)#exit
device(config)#vlan 20
device(config-VLAN-20)#monitor ethernet 6/24 output
device(config-VLAN-20)#end
```

For example, to enable bidirectional mirroring on VLANs 10 and 20, to mirror port e 6/24, enter the following commands.

```
device(config)#mirror-port ethernet 6/24
device(config)#vlan 10
device(config-VLAN-10)#monitor ethernet 6/24 both
device(config-VLAN-10)#exit
device(config)#vlan 20
device(config-VLAN-20)#monitor ethernet 6/24 both
device(config-VLAN-20)#end
```

To disable mirroring on VLAN 20, enter the following commands.

```
device(config)#vlan 20
device(config-VLAN-20)#no monitor ethernet 6/24
device(config-VLAN-20)#end
```

Displaying VLAN-based mirroring status

The **show vlan** command displays the VLAN-based mirroring status.

```
device(config-if-e1000-5/2)#show vlan 10
Total PORT-VLAN entries: 2
Maximum PORT-VLAN entries: 64
Legend: [Stk=Stack-Id, S=Slot]
PORT-VLAN 10, Name To-SXR1600, Priority level0, Spanning tree Off
Untagged Ports: (S5) 2 4
    Tagged Ports: None
    Uplink Ports: None
DualMode Ports: None
    RX Monitoring: Disabled
    TX Monitoring: Enabled
```

Syntax: [no] monitor ethernet port [input | output | both]

Displaying VLAN-based mirroring status

Spanning Tree Protocol

• STP overview.....	241
• Standard STP parameter configuration.....	241
• STP feature configuration.....	255
• PVST/PVST+ compatibility.....	306
• PVRST compatibility.....	312
• BPDU guard.....	312
• Root guard.....	315
• Designated Protection.....	316
• Error disable recovery.....	317
• 802.1s Multiple Spanning Tree Protocol.....	319

STP overview

The Spanning Tree Protocol (STP) eliminates Layer 2 loops in networks, by selectively blocking some ports and allowing other ports to forward traffic, based on global (bridge) and local (port) parameters you can configure.

STP related features, such as RSTP and PVST, extend the operation of standard STP, enabling you to fine-tune standard STP and avoid some of its limitations.

You can enable or disable STP on a global basis (for the entire device), a port-based VLAN basis (for the individual Layer 2 broadcast domain), or an individual port basis.

Configuration procedures are provided for the standard STP bridge and port parameters.

Standard STP parameter configuration

Brocade Layer 2 Switches and Layer 3 Switches support standard STP as described in the IEEE 802.1D specification. STP is enabled by default on Layer 2 Switches but disabled by default on Layer 3 Switches.

By default, each port-based VLAN on a Brocade device runs a separate spanning tree (a separate instance of STP). A Brocade device has one port-based VLAN (VLAN 1) by default that contains all the device ports. Thus, by default each Brocade device has one spanning tree. However, if you configure additional port-based VLANs on a Brocade device, then each of those VLANs on which STP is enabled and VLAN 1 all run separate spanning trees.

If you configure a port-based VLAN on the device, the VLAN has the same STP state as the default STP state on the device. Thus, on Layer 2 Switches, new VLANs have STP enabled by default. On Layer 3 Switches, new VLANs have STP disabled by default. You can enable or disable STP in each VLAN separately. In addition, you can enable or disable STP on individual ports.

STP parameters and defaults

The following table lists the default STP states for Brocade devices.

TABLE 43 Default STP states

Device type	Default STP type	Default STP state	Default STP state of new VLANs ¹
Switch image default	STP 802.1d	Enabled	Enabled
Router image default	No span	Enabled	Disabled
Base L3 image default	No span	Disabled	Disabled
ICX 6650			
Layer 2 switch	MSTP ²	Enabled	Enabled
Layer 3 switch	MSTP	Disabled	Disabled

1. When you create a port-based VLAN, the new VLAN STP state is the same as the default STP state on the device. The new VLAN does not inherit the STP state of the default VLAN. The Single Spanning Tree Protocol (SSTP) is another type of STP. SSTP includes all VLANs on which STP is enabled in a single spanning tree. Refer to [Single Spanning Tree \(SSTP\)](#) on page 300.

2. MSTP stands for "Multiple Spanning Tree Protocol". In this type of STP, each port-based VLAN, including the default VLAN, has its own spanning tree. References in this documentation to "STP" apply to MSTP.

The following table lists the default STP bridge parameters. The bridge parameters affect the entire spanning tree. If you are using MSTP, the parameters affect the VLAN. If you are using SSTP, the parameters affect all VLANs that are members of the single spanning tree.

TABLE 44 Default STP bridge parameters

Parameter	Description	Default and valid values
Forward Delay	The period of time spent by a port in the listening and learning state before moving on to the learning or forwarding state, respectively. The forward delay value is also used for the age time of dynamic entries in the filtering database, when a topology change occurs.	15 seconds Possible values: 4 - 30 seconds
Maximum Age	The interval a bridge will wait for a configuration BPDU from the root bridge before initiating a topology change.	20 seconds Possible values: 6 - 40 seconds
Hello Time	The interval of time between each configuration BPDU sent by the root bridge.	2 seconds Possible values: 1 - 10 seconds
Priority	A parameter used to identify the root bridge in a spanning tree (instance of STP). The bridge with the lowest value has the highest priority and is the root. A higher numerical value means a lower priority; thus, the highest priority is 0.	32768 Possible values: 0 - 65535

NOTE

If you plan to change STP bridge timers, Brocade recommends that you stay within the following ranges, from section 8.10.2 of the IEEE STP specification. $2 * (\text{forward_delay} - 1) \geq \text{max_age}$
 $\text{max_age} \geq 2 * (\text{hello_time} + 1)$

The following table lists the default STP port parameters. The port parameters affect individual ports and are separately configurable on each port.

TABLE 45 Default STP port parameters

Parameter	Description	Default and valid values
Priority	The preference that STP gives this port relative to other ports for forwarding traffic out of the spanning tree. A higher numerical value means a lower priority.	128 Possible values: 0 - 240 (configurable in increments of 16)
Path Cost	The cost of using the port to reach the root bridge. When selecting among multiple links to the root bridge, STP chooses the link with the lowest path cost and blocks the other paths. Each port type has its own default STP path cost.	10 Mbps - 100 100 Mbps - 19 Gbps - 4 10 Gbps - 2 Possible values are 0 - 65535

Enabling or disabling the Spanning Tree Protocol (STP)

STP is *enabled* by default on devices running Layer 2 code. STP is *disabled* by default on devices running Layer 3 code.

You can enable or disable STP on the following levels:

- Globally - Affects all ports and port-based VLANs on the device.
- Port-based VLAN - Affects all ports within the specified port-based VLAN. When you enable or disable STP within a port-based VLAN, the setting overrides the global setting. Thus, you can enable STP for the ports within a port-based VLAN even when STP is globally disabled, or disable the ports within a port-based VLAN when STP is globally enabled.
- Individual port - Affects only the individual port. However, if you change the STP state of the primary port in a trunk group, the change affects all ports in the trunk group.

NOTE

The CLI converts the STP groups into topology groups when you save the configuration. For backward compatibility, you can still use the STP group commands. However, the CLI converts the commands into the topology group syntax. Likewise, the **show stp-group** command displays STP topology groups.

Configuration modes for STP

The following configuration modes apply while configuring STP.

- Spanning-tree single - This configuration can be enabled on systems running IEEE 802.1D. The single spanning tree controls all the 4000 VLANs. You can opt in and out of this single spanning tree using the **spanning-tree** command under the VLAN prompt.
- Spanning-tree single 802.1w - This configuration can be enabled on systems running IEEE 802.1w. The single rapid spanning tree controls all the 4000 VLANs. The VLAN can opt in and out of this single rapid spanning tree using the **spanning-tree** command under the VLAN prompt. If there is a

“spanning-tree” configuration under the VLAN, that VLAN will be with that single 802.1w instance’s control, which implies that the VLAN traffic is subject to blocking or forwarding by that spanning tree instance.

- Per VLAN spanning tree - In this configuration mode you can turn on 802.1D or 802.1w (Rapid Spanning Tree) at the VLAN level individually.

Enabling or disabling STP globally

Use the following method to enable or disable STP on a device on which you have not configured port-based VLANs.

NOTE

When you configure a VLAN, the VLAN inherits the global STP settings. However, once you begin to define a VLAN, you can no longer configure standard STP parameters globally using the CLI. From that point on, you can configure STP only within individual VLANs.

To enable STP for all ports in all VLANs on a Brocade device, enter the **spanning-tree** command.

```
device(config)# spanning-tree
```

The **spanning-tree** command enables a separate spanning tree in each VLAN, including the default VLAN.

Syntax: [no] spanning-tree

To set system maximum value for number of spanning tree instances, enter the command such as the following:

```
Brocade (config)# system-max spanning-tree 254
```

Syntax: system-max spanning-tree number

NOTE

The number of spanning tree instances ranges from 1 - 254, with 32 as the default.

Enabling or disabling STP in a port-based VLAN

Use the following procedure to disable or enable STP on a device on which you have configured a port-based VLAN. Changing the STP state in a VLAN affects only that VLAN.

To enable STP for all ports in a port-based VLAN, enter commands such as the following.

```
device(config)#vlan 10
device(config-vlan-10)#spanning-tree
```

Syntax: [no] spanning-tree

Enabling or disabling STP on an individual port

Use the following procedure to disable or enable STP on an individual port.

NOTE

If you change the STP state of the primary port in a trunk group, it affects all ports in the trunk group.

To enable STP on an individual port, enter commands such as the following.

```
device(config)#interface 1/1
device(config-if-e1000-1/1)#spanning-tree
```

Syntax: [no] spanning-tree

Changing STP bridge and port parameters

[Table 44](#) on page 242 and [Table 45](#) on page 243 list the default STP parameters. If you need to change the default value for an STP parameter, use the following procedures.

Changing STP bridge parameters

NOTE

If you plan to change STP bridge timers, Brocade recommends that you stay within the following ranges, from section 8.10.2 of the IEEE STP specification. $2 * (\text{forward_delay} - 1) \geq \text{max_age}$
 $\text{max_age} \geq 2 * (\text{hello_time} + 1)$

To change a STP bridge priority on a Brocade device to the highest value to make the device the root bridge, enter the following command.

```
device(config)#spanning-tree priority 0
```

The command in this example changes the priority on a device on which you have not configured port-based VLANs. The change applies to the default VLAN. If you have configured a port-based VLAN on the device, you can configure the parameters only at the configuration level for individual VLANs. Enter commands such as the following.

```
device(config)#vlan 20
device(config-vlan-20)#spanning-tree priority 0
```

To make this change in the default VLAN, enter the following commands.

```
device(config)#vlan 1
device(config-vlan-1)#spanning-tree priority 0
```

Syntax: [no] spanning-tree [**forward-delay** *value*] | [**hello-time** *value*] | [**maximum-age** *value*] | [**priority** *value*]

The **forward-delay** *value* parameter specifies the forward delay and can be a value from 4 - 30 seconds. The default is 15 seconds.

NOTE

You can configure a Brocade device for faster convergence (including a shorter forward delay) using Fast Span or Fast Uplink Span. Refer to [STP feature configuration](#) on page 255.

The **hello-time** *value* parameter specifies the hello time and can be a value from 1 - 10 seconds. The default is 2 seconds.

NOTE

This parameter applies only when this device or VLAN is the root bridge for its spanning tree.

The **maximum-age** *value* parameter specifies the amount of time the device waits for receipt of a configuration BPDU from the root bridge before initiating a topology change. You can specify from 6 - 40 seconds. The default is 20 seconds.

The **priority** *value* parameter specifies the priority and can be a value from 0 - 65535. A higher numerical value means a lower priority. Thus, the highest priority is 0. The default is 32768.

You can specify some or all of these parameters on the same command line. If you specify more than one parameter, you must specify them in the order shown above, from left to right.

Changing STP port parameters

To change the path and priority costs for a port, enter commands such as the following.

```
device(config)#vlan 10
device(config-vlan-10)#spanning-tree ethernet 5 path-cost 15 priority 64
```

Syntax: **spanning-tree ethernet** *port* **path-cost** *value* | **priority** *value* | **disable** | **enable**

The **path-cost** *value* parameter specifies the port cost as a path to the spanning tree root bridge. STP prefers the path with the lowest cost. You can specify a value from 0 - 65535.

The default depends on the port type:

- 10 Mbps - 100
- 100 Mbps - 19
- Gbps - 4
- 10 Gbps - 2
- The **priority***value* parameter specifies the preference that STP gives this port relative to other ports for forwarding traffic out of the spanning tree. If you are upgrading a device that has a configuration saved under an earlier software release, and the configuration contains a value from 0 - 7 for a port STP priority, the software changes the priority to the default when you save the configuration while running the new release.

The **disable** and **enable** parameter disables or re-enables STP on the port. The STP state change affects only this VLAN. The port STP state in other VLANs is not changed.

STP protection enhancement

STP protection provides the ability to prohibit an end station from initiating or participating in an STP topology change.

The 802.1W Spanning Tree Protocol (STP) detects and eliminates logical loops in a redundant network by selectively blocking some data paths (ports) and allowing only the best data paths to forward traffic.

In an STP environment, switches, end stations, and other Layer 2 devices use Bridge Protocol Data Units (BPDUs) to exchange information that STP will use to determine the best path for data flow. When a Layer 2 device is powered ON and connected to the network, or when a Layer 2 device goes down, it sends out an STP BPDU, triggering an STP topology change.

In some instances, it is unnecessary for a connected device, such as an end station, to initiate or participate in an STP topology change. In this case, you can enable the STP Protection feature on the Brocade port to which the end station is connected. STP Protection disables the connected device ability to initiate or participate in an STP topology change, by dropping all BPDUs received from the connected device.

Enabling STP protection

You can enable STP Protection on a per-port basis.

To prevent an end station from initiating or participating in STP topology changes, enter the following command at the Interface level of the CLI.

```
device#(config) interface ethernet 2
device#(config-if-e1000-2)#stp-protect
```

This command causes the port to drop STP BPDU sent from the device on the other end of the link.

Syntax: [no] stp-protect

Enter the **no** form of the command to disable STP protection on the port.

Clearing BPDU drop counters

For each port that has STP Protection enabled, the Brocade device counts and records the number of dropped BPDU. You can use CLI commands to clear the BPDU drop counters for all ports on the device, or for a specific port on the device.

To clear the BPDU drop counters for all ports on the device that have STP Protection enabled, enter the following command at the Global CONFIG level of the CLI.

```
device(config)#clear stp-protect-statistics
```

To clear the BPDU drop counter for a specific port that has STP Protection enabled, enter the following command at the Global CONFIG level of the CLI.

```
device#clear stp-protect-statistics e 2
```

Syntax: clear stp-protect-statistics [ethernet port]

Viewing the STP Protection Configuration

You can view the STP Protection configuration for all ports on a device, or for a specific port only. The **show stp-protect** command output shows the port number on which STP Protection is enabled, and the number of BPDU dropped by each port.

To view the STP Protection configuration for all ports on the device, enter the following command at any level of the CLI.

```
device#show stp-protect
Port ID          BPDU Drop Count
  3              478
  5              213
  6              0
  12             31
```

To view STP Protection configuration for a specific port, enter the following command at any level of the CLI.

```
device#show stp-protect e 3
STP-protect is enabled on port 3.  BPDU drop count is 478
```

If you enter the **show stp-protect** command for a port that does not have STP protection enabled, the following message displays on the console.

```
device#show stp-protect e 4
STP-protect is not enabled on port 4.
```

Syntax: `show stp-protect [ethernet port]`

Displaying STP information

You can display the following Spanning Tree Protocol (STP) information:

- All the global and interface STP settings
- CPU utilization statistics
- Detailed STP information for each interface
- STP state information for a port-based VLAN
- STP state information for an individual interface

Displaying STP information for an entire device

To display STP information, enter the following command at any level of the CLI.

```
device#show span
VLAN 1 BPDU cam_index is 3 and the Master DMA Are (HEX)
STP instance owned by VLAN 1
Global STP (IEEE 802.1D) Parameters:
VLAN Root          Root Root Prio Max He- Ho- Fwd Last      Chg  Bridge
ID    ID           Cost Port rity Age llo ld dly Chang   cnt Address
                           Hex sec sec sec sec sec
                           1   800000e0804d4a00 0   Root 8000 20 2 1 15 689   1   00e0804d4a00

Port STP Parameters:
Port Prio Path State      Fwd Design Designated Designated
Num  rity Cost          Trans Cost Root Bridge
Hex
1    80   19  FORWARDING 1   0   800000e0804d4a00 800000e0804d4a00
2    80   0   DISABLED   0   0   0000000000000000 0000000000000000
3    80   0   DISABLED   0   0   0000000000000000 0000000000000000
4    80   0   DISABLED   0   0   0000000000000000 0000000000000000
5    80   19  FORWARDING 1   0   800000e0804d4a00 800000e0804d4a00
6    80   19  BLOCKING   0   0   800000e0804d4a00 800000e0804d4a00
7    80   0   DISABLED   0   0   0000000000000000 0000000000000000

<lines for remaining ports excluded for brevity>
```

Syntax: `show span [vlan vlan-id] | [pvst-mode] | [num] | [detail [vlan vlan-id [ethernet [port] | num]]]`

The `vlan``vlan-id` parameter displays STP information for the specified port-based VLAN.

The `pvst-mode` parameter displays STP information for the device Per VLAN Spanning Tree (PVST+) compatibility configuration. Refer to [PVST/PVST+ compatibility](#) on page 306

The `num` parameter displays only the entries after the number you specify. For example, on a device with three port-based VLANs, if you enter 1, then information for the second and third VLANs is displayed, but information for the first VLAN is not displayed. Information is displayed according to VLAN number, in ascending order. The entry number is not the same as the VLAN number. For example, if you have port-based VLANs 1, 10, and 2024, then the command output has three STP entries. To display information for VLANs 10 and 2024 only, enter `show span 1`.

The `detail` parameter and its additional optional parameters display detailed information for individual ports. Refer to [Displaying detailed STP information for each interface](#) on page 251.

The `show span` command shows the following information.

TABLE 46 CLI display of STP information

Field	Description
Global STP parameters	
VLAN ID	The port-based VLAN that contains this spanning tree (instance of STP). VLAN 1 is the default VLAN. If you have not configured port-based VLANs on this device, all STP information is for VLAN 1.
Root ID	The ID assigned by STP to the root bridge for this spanning tree.
Root Cost	The cumulative cost from this bridge to the root bridge. If this device is the root bridge, then the root cost is 0.
Root Port	The port on this device that connects to the root bridge. If this device is the root bridge, then the value is "Root" instead of a port number.
Priority Hex	This device or VLAN STP priority. The value is shown in hexadecimal format.
NOTE If you configure this value, specify it in decimal format. Refer to Changing STP bridge parameters on page 245.	
Max age sec	The number of seconds this device or VLAN waits for a configuration BPDU from the root bridge before deciding the root has become unavailable and performing a reconvergence.
Hello sec	The interval between each configuration BPDU sent by the root bridge.
Hold sec	The minimum number of seconds that must elapse between transmissions of consecutive Configuration BPDUs on a port.
Fwd dly sec	The number of seconds this device or VLAN waits following a topology change and consequent reconvergence.
Last Chang sec	The number of seconds since the last time a topology change occurred.
Chg cnt	The number of times the topology has changed since this device was reloaded.
Bridge Address	The STP address of this device or VLAN.
NOTE If this address is the same as the Root ID, then this device or VLAN is the root bridge for its spanning tree.	
Port STP parameters	
Port Num	The port number.

TABLE 46 CLI display of STP information (Continued)

Field	Description
Priority Hex	The port STP priority, in hexadecimal format.
NOTE	
If you configure this value, specify it in decimal format. Refer to Changing STP port parameters on page 246 .	
Path Cost	The port STP path cost.
State	<p>The port STP state. The state can be one of the following:</p> <ul style="list-style-type: none"> • BLOCKING - STP has blocked Layer 2 traffic on this port to prevent a loop. The device or VLAN can reach the root bridge using another port, whose state is FORWARDING. When a port is in this state, the port does not transmit or receive user frames, but the port does continue to receive STP BPDUs. • DISABLED - The port is not participating in STP. This can occur when the port is disconnected or STP is disabled on the port. • FORWARDING - STP is allowing the port to send and receive frames. • LISTENING - STP is responding to a topology change and this port is listening for a BPDU from neighboring bridges in order to determine the new topology. No user frames are transmitted or received during this state. • LEARNING - The port has passed through the LISTENING state and will change to the FORWARDING state, depending on the results of STP reconvergence. The port does not transmit or receive user frames during this state. However, the device can learn the MAC addresses of frames that the port receives during this state and make corresponding entries in the MAC table. • DESIGNATED INCONSISTENT - This shows as DESI-INCONS in the output. You can disallow the designated forwarding state on a port in STP 802.1d or 802.1w with the spanning-tree designated-protect command. If STP tries to put this port into the designated forwarding role, the device would put this port into a designated inconsistent STP state. This is effectively equivalent to the listening state in STP in which a port cannot transfer any user traffic. When STP no longer marks this port as a designated port, the device automatically removes the port from the designated inconsistent state.
Fwd Trans	The number of times STP has changed the state of this port between BLOCKING and FORWARDING.
Design Cost	The cost to the root bridge as advertised by the designated bridge that is connected to this port. If the designated bridge is the root bridge itself, then the cost is 0. The identity of the designated bridge is shown in the Design Bridge field.
Designated Root	The root bridge as recognized on this port. The value is the same as the root bridge ID listed in the Root ID field.
Designated Bridge	The designated bridge to which this port is connected. The designated bridge is the device that connects the network segment on the port to the root bridge.

Displaying the STP state of a port-based VLAN

When you display information for a port-based VLAN, that information includes the STP state of the VLAN.

To display information for a port-based VLAN, enter a command such as the following at any level of the CLI. The STP state is shown in bold type in this example.

```
device#show vlans
Total PORT-VLAN entries: 2
Maximum PORT-VLAN entries: 16
legend: [S=Slot]
PORT-VLAN 1, Name DEFAULT-VLAN, Priority level0, Spanning tree On
  Untagged Ports: (S3) 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16
  Untagged Ports: (S3) 17 18 19 20 21 22 23 24
  Untagged Ports: (S4) 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17
  Untagged Ports: (S4) 18 19 20 21 22 23 24
    Tagged Ports: None
    Uplink Ports: None
PORT-VLAN 2, Name greenwell, Priority level0, Spanning tree Off
  Untagged Ports: (S1) 1 2 3 4 5 6 7 8
  Untagged Ports: (S4) 1
    Tagged Ports: None
    Uplink Ports: None
```

Syntax: **show vlan [vlan-id | ethernet port]**

The *vlan-id* parameter specifies a VLAN for which you want to display the configuration information.

Displaying detailed STP information for each interface

To display the detailed STP information, enter the following command at any level of the CLI.

```
device#show span detail
=====
VLAN 1 - MULTIPLE SPANNING TREE (MSTP) ACTIVE
=====
Bridge identifier - 0x800000e0804d4a00
Active global timers - Hello: 0
Port 1/1 is FORWARDING
  Port - Path cost: 19, Priority: 128, Root: 0x800000e052a9bb00
  Designated - Bridge: 0x800000e052a9bb00, Interface: 1, Path cost: 0
  Active Timers - None
  BPDU's - Sent: 11, Received: 0
Port 1/2 is DISABLED
Port 1/3 is DISABLED
Port 1/4 is DISABLED
<lines for remaining ports excluded for brevity>
```

NOTE

The line in the above output, VLAN 1 - MULTIPLE SPANNING TREE (MSTP) ACTIVE, is not the 802.1s standard. It is the same Global STP (IEEE 802.1D) type as shown in the output of the show span CLI command.

If a port is disabled, the only information shown by this command is "DISABLED". If a port is enabled, this display shows the following information.

Syntax: **show span detail [vlan *vlan-id* | ethernet *port* | num]**

The **vlan** *vlan-id* parameter specifies a VLAN.

The *num* parameter specifies the number of VLANs you want the CLI to skip before displaying detailed STP information. For example, if the device has six VLANs configured (VLAN IDs 1, 2, 3, 99, 128, and 256) and you enter the command **show span detail 4**, detailed STP information is displayed for VLANs 128 and 256 only.

NOTE

If the configuration includes VLAN groups, the **show span detail** command displays the master VLANs of each group but not the member VLANs within the groups. However, the command does indicate that the VLAN is a master VLAN. The **show span detail vlan-id** command displays the information for the VLAN even if it is a member VLAN. To list all the member VLANs within a VLAN group, enter the **show vlan-group [group-id]** command.

The **show span detail** command shows the following information.

TABLE 47 CLI display of detailed STP information for ports

Field	Description
Active Spanning Tree protocol	<p>The VLAN that contains the listed ports and the active Spanning Tree protocol.</p> <p>The STP type can be one of the following:</p> <ul style="list-style-type: none"> • MULTIPLE SPANNING TREE (MSTP) • GLOBAL SINGLE SPANNING TREE (SSTP)
Bridge identifier	<p>NOTE</p> <p>If STP is disabled on a VLAN, the command displays the following message instead: "Spanning-tree of port-vlan vlan-id is disabled."</p>
Active global timers	<p>The global STP timers that are currently active, and their current values. The following timers can be listed:</p> <ul style="list-style-type: none"> • Hello - The interval between Hello packets. This timer applies only to the root bridge. • Topology Change (TC) - The amount of time during which the topology change flag in Hello packets will be marked, indicating a topology change. This timer applies only to the root bridge. • Topology Change Notification (TCN) - The interval between Topology Change Notification packets sent by a non-root bridge toward the root bridge. This timer applies only to non-root bridges.

TABLE 47 CLI display of detailed STP information for ports (Continued)

Field	Description
Port number and STP state	<p>The internal port number and the port STP state.</p> <p>The internal port number is one of the following:</p> <ul style="list-style-type: none"> • The port interface number, if the port is the designated port for the LAN. • The interface number of the designated port from the received BPDU, if the interface is not the designated port for the LAN. <p>The state can be one of the following:</p> <ul style="list-style-type: none"> • BLOCKING - STP has blocked Layer 2 traffic on this port to prevent a loop. The device or VLAN can reach the root bridge using another port, whose state is FORWARDING. When a port is in this state, the port does not transmit or receive user frames, but the port does continue to receive STP BPDUs. • DISABLED - The port is not participating in STP. This can occur when the port is disconnected or STP is administratively disabled on the port. • FORWARDING - STP is allowing the port to send and receive frames. • LISTENING - STP is responding to a topology change and this port is listening for a BPDU from neighboring bridges in order to determine the new topology. No user frames are transmitted or received during this state. • LEARNING - The port has passed through the LISTENING state and will change to the BLOCKING or FORWARDING state, depending on the results of STP reconvergence. The port does not transmit or receive user frames during this state. However, the device can learn the MAC addresses of frames that the port receives during this state and make corresponding entries in the MAC table.
NOTE	
	If the state is DISABLED, no further STP information is displayed for the port.
Port Path cost	The STP path cost for the port.
Port Priority	This STP priority for the port. The value is shown as a hexadecimal number.
Root	The ID assigned by STP to the root bridge for this spanning tree.
Designated Bridge	The MAC address of the designated bridge to which this port is connected. The designated bridge is the device that connects the network segment on the port to the root bridge.
Designated Port	The port number sent from the designated bridge.
Designated Path Cost	The cost to the root bridge as advertised by the designated bridge that is connected to this port. If the bridge is the root bridge itself, then the cost is 0. The identity of the designated bridge is shown in the Designated Bridge field.

TABLE 47 CLI display of detailed STP information for ports (Continued)

Field	Description
Active Timers	The current values for the following timers, if active: <ul style="list-style-type: none"> Message age - The number of seconds this port has been waiting for a hello message from the root bridge. Forward delay - The number of seconds that have passed since the last topology change and consequent reconvergence. Hold time - The number of seconds that have elapsed since transmission of the last Configuration BPDU.
BPDUs Sent and Received	The number of BPDUs sent and received on this port since the software was reloaded.

Displaying detailed STP information for a single port in a specific VLAN

Enter a command such as the following to display STP information for an individual port in a specific VLAN.

```
device#show span detail vlan 1 ethernet 7/1
Port 7/1 is FORWARDING
  Port - Path cost: 19, Priority: 128, Root: 0x80000e052a9bb00
  Designated - Bridge: 0x80000e052a9bb00, Interface: 7, Path cost: 0
  Active Timers - None
  BPDUs - Sent: 29, Received: 0
```

Syntax: **show span detail [**vlan** **vlan-id** **etherent** **port**]**

Displaying STP state information for an individual interface

To display STP state information for an individual port, you can use the methods in [Displaying STP information for an entire device](#) on page 248 or [Displaying detailed STP information for each interface](#) on page 251. You also can display STP state information for a specific port using the following method.

To display information for a specific port, enter a command such as the following at any level of the CLI.

```
device#show interface ethernet 3/11
FastEthernet3/11 is up, line protocol is up
Port up for 1 hour 50 minutes 30 seconds
  Hardware is FastEthernet, address is 0000.00a9.bb49 (bia 0000.00a9.bb49)
  Configured speed auto, actual 100Mbit, configured duplex fdx, actual fdx
  Member of L2 VLAN ID 1, port is untagged, port state is FORWARDING
  STP configured to ON
  , priority is level0, flow control enabled
  mirror disabled, monitor disabled
  Not member of any active trunks
  Not member of any configured trunks
  No port name
  MTU 1518 bytes, encapsulation ethernet
  5 minute input rate: 352 bits/sec, 0 packets/sec, 0.00% utilization
  5 minute output rate: 0 bits/sec, 0 packets/sec, 0.00% utilization
  1238 packets input, 79232 bytes, 0 no buffer
  Received 686 broadcasts, 0 runts, 0 giants
  0 input errors, 0 CRC, 0 frame, 0 ignored
  529 multicast
  918 packets output, 63766 bytes, 0 underruns
  0 output errors, 0 collisions
```

The STP information is shown in bold type in this example.

Syntax: show interfaces [ethernet port] | [loopback num] | [slot slot-num] | [ve num] | [brief]

You also can display the STP states of all ports by entering the show interface brief command such as the following, which uses the **brief** parameter.

```
device#show interface brief
Port Link State Dupl Speed Trunk Tag Priori MAC Name
1/1 Down None None None No level0 0000.00a9.bb00
1/2 Down None None None No level0 0000.00a9.bb01
1/3 Down None None None No level0 0000.00a9.bb02
1/4 Down None None None No level0 0000.00a9.bb03
1/5 Down None None None No level0 0000.00a9.bb04
1/6 Down None None None No level0 0000.00a9.bb05
1/7 Down None None None No level0 0000.00a9.bb06
1/8 Down None None None No level0 0000.00a9.bb07
.
.
.
some rows omitted for brevity
.
.
3/10 Down None None None No level0 0000.00a9.bb4a
3/11 Up Forward Full 100M None No level0 0000.00a9.bb49
```

In the example above, only one port, 3/11, is forwarding traffic toward the root bridge.

STP feature configuration

Spanning Tree Protocol (STP) features extend the operation of standard STP, enabling you to fine-tune standard STP and avoid some of its limitations.

This section describes how to configure these parameters on Brocade Layer 3 Switches using the CLI.

Fast port span

When STP is running on a device, message forwarding is delayed during the spanning tree recalculation period following a topology change. The STP forward delay parameter specifies the period of time a bridge waits before forwarding data packets. The forward delay controls the listening and learning periods of STP reconvergence. You can configure the forward delay to a value from 4 - 30 seconds. The default is 15 seconds. Thus, using the standard forward delay, convergence requires 30 seconds (15 seconds for listening and an additional 15 seconds for learning) when the default value is used.

This slow convergence is undesirable and unnecessary in some circumstances. The Fast Port Span feature allows certain ports to enter the forwarding state in four seconds. Specifically, Fast Port Span allows faster convergence on ports that are attached to end stations and thus do not present the potential to cause Layer 2 forwarding loops. Because the end stations cannot cause forwarding loops, they can safely go through the STP state changes (blocking to listening to learning to forwarding) more quickly than is allowed by the standard STP convergence time. Fast Port Span performs the convergence on these ports in four seconds (two seconds for listening and two seconds for learning).

In addition, Fast Port Span enhances overall network performance in the following ways:

- Fast Port Span reduces the number of STP topology change notifications on the network. When an end station attached to a Fast Span port comes up or down, the Brocade device does not generate a topology change notification for the port. In this situation, the notification is unnecessary since a change in the state of the host does not affect the network topology.
- Fast Port Span eliminates unnecessary MAC cache aging that can be caused by topology change notifications. Bridging devices age out the learned MAC addresses in their MAC caches if the addresses are unrefreshed for a given period of time, sometimes called the MAC aging interval. When STP sends a topology change notification, devices that receive the notification use the value of

the STP forward delay to quickly age out their MAC caches. For example, if a device normal MAC aging interval is 5 minutes, the aging interval changes temporarily to the value of the forward delay (for example, 15 seconds) in response to an STP topology change.

In normal STP, the accelerated cache aging occurs even when a single host goes up or down. Because Fast Port Span does not send a topology change notification when a host on a Fast Port Span port goes up or down, the unnecessary cache aging that can occur in these circumstances under normal STP is eliminated.

Fast Port Span is a system-wide parameter and is enabled by default. Thus, when you boot a device, all the ports that are attached only to end stations run Fast Port Span. For ports that are not eligible for Fast Port Span, such as ports connected to other networking devices, the device automatically uses the normal STP settings. If a port matches any of the following criteria, the port is ineligible for Fast Port Span and uses normal STP instead:

- The port is 802.1Q tagged
- The port is a member of a trunk group
- The port has learned more than one active MAC address
- An STP Configuration BPDU has been received on the port, thus indicating the presence of another bridge on the port.

You also can explicitly exclude individual ports from Fast Port Span if needed. For example, if the only uplink ports for a wiring closet switch are Gbps ports, you can exclude the ports from Fast Port Span.

Disabling and re-enabling fast port span

Fast Port Span is a system-wide parameter and is enabled by default. Therefore, all ports that are eligible for Fast Port Span use it.

To disable or re-enable Fast Port Span, enter the following commands.

```
device(config)#no fast port-span  
device(config)#write memory
```

Syntax: [no] fast port-span

NOTE

The **fast port-span** command has additional parameters that let you exclude specific ports. These parameters are shown in the following section.

To re-enable Fast Port Span, enter the following commands.

```
device(config)#fast port-span  
device(config)#write memory
```

Excluding specific ports from fast port span

To exclude a port from Fast Port Span while leaving Fast Port Span enabled globally, enter commands such as the following.

```
device(config)#fast port-span exclude ethernet 1  
device(config)#write memory
```

To exclude a set of ports from Fast Port Span, enter commands such as the following.

```
device(config)#fast port-span exclude ethernet 1 ethernet 2 ethernet 3  
device(config)#write memory
```

To exclude a contiguous (unbroken) range of ports from Fast Span, enter commands such as the following.

```
device(config)#fast port-span exclude ethernet 1 to 24
device(config)#write memory
```

Syntax: [no] fast port-span [exclude ethernet port [ethernet port] | to [port]]

To re-enable Fast Port Span on a port, enter a command such as the following.

```
device(config)#no fast port-span exclude ethernet 1
device(config)#write memory
```

This command re-enables Fast Port Span on port 1 only and does not re-enable Fast Port Span on other excluded ports. You also can re-enable Fast Port Span on a list or range of ports using the syntax shown above this example.

To re-enable Fast Port Span on all excluded ports, disable and then re-enable Fast Port Span by entering the following commands.

```
device(config)#no fast port-span
device(config)#fast port-span
device(config)#write memory
```

Disabling and then re-enabling Fast Port Span clears the exclude settings and thus enables Fast Port Span on all eligible ports. To make sure Fast Port Span remains enabled on the ports following a system reset, save the configuration changes to the startup-config file after you re-enable Fast Port Span. Otherwise, when the system resets, those ports will again be excluded from Fast Port Span.

Fast Uplink Span

The Fast Port Span feature described in the previous section enhances STP performance for end stations. The Fast Uplink Span feature enhances STP performance for wiring closet switches with redundant uplinks. Using the default value for the standard STP forward delay, convergence following a transition from an active link to a redundant link can take 30 seconds (15 seconds for listening and an additional 15 seconds for learning).

You can use the Fast Uplink Span feature on a Brocade device deployed as a wiring closet switch to decrease the convergence time for the uplink ports to another device to just one second. The new Uplink port directly goes to forward mode (bypassing listening and learning modes). The wiring closet switch must be a Brocade device but the device at the other end of the link can be a Brocade device or another vendor's switch.

Configuration of the Fast Uplink Span feature takes place entirely on the Brocade device. To configure the Fast Uplink Span feature, specify a group of ports that have redundant uplinks on the wiring closet switch (Brocade device). If the active link becomes unavailable, the Fast Uplink Span feature transitions the forwarding to one of the other redundant uplink ports in just one second. All Fast Uplink Span-enabled ports are members of a single Fast Uplink Span group.

NOTE

To avoid the potential for temporary bridging loops, Brocade recommends that you use the Fast Uplink feature only for wiring closet switches (switches at the edge of the network cloud). In addition, enable the feature only on a group of ports intended for redundancy, so that at any given time only one of the ports is expected to be in the forwarding state.

NOTE

When the wiring closet switch (Brocade device) first comes up or when STP is first enabled, the uplink ports still must go through the standard STP state transition without any acceleration. This behavior guards against temporary routing loops as the switch tries to determine the states for all the ports. Fast Uplink Span acceleration applies only when a working uplink becomes unavailable.

Active uplink port failure

The active uplink port is the port elected as the root port using the standard STP rules. All other ports in the group are redundant uplink ports. If an active uplink port becomes unavailable, Fast Uplink Span transitions the forwarding of traffic to one of the redundant ports in the Fast Uplink Span group in one second bypassing listening and learning port states.

Switchover to the active uplink port

When a failed active uplink port becomes available again, switchover from the redundant port to the active uplink port is delayed by 30 seconds. The delay allows the remote port to transition to forwarding mode using the standard STP rules. After 30 seconds, the blocked active uplink port begins forwarding in just one second and the redundant port is blocked.

NOTE

Use caution when changing the spanning tree priority. If the switch becomes the root bridge, Fast Uplink Span will be disabled automatically.

Fast Uplink Span Rules for Trunk Groups

If you add a port to a Fast Uplink Span group that is a member of a trunk group, the following rules apply:

- If you add the primary port of a trunk group to the Fast Uplink Span group, all other ports in the trunk group are automatically included in the group. Similarly, if you remove the primary port in a trunk group from the Fast Uplink Span group, the other ports in the trunk group are automatically removed from the Fast Uplink Span group.
- You cannot add a subset of the ports in a trunk group to the Fast Uplink Span group. All ports in a trunk group have the same Fast Uplink Span property, as they do for other port properties.
- If the working trunk group is partially down but not completely down, no switch-over to the backup occurs. This behavior is the same as in the standard STP feature.
- If the working trunk group is completely down, a backup trunk group can go through an accelerated transition only if the following are true:
 - The trunk group is included in the fast uplink group.
 - All other ports except those in this trunk group are either disabled or blocked. The accelerated transition applies to all ports in this trunk group.

When the original working trunk group comes back (partially or fully), the transition back to the original topology is accelerated if the conditions listed above are met.

Configuring a Fast Uplink Port Group

To configure a group of ports for Fast Uplink Span, enter the following commands:

```
device(config)# fast uplink-span ethernet 4/1 to 4/4
device(config)# write memory
```

Syntax: [no] fast uplink-span [ethernet port [ethernet port...| to port]]

This example configures four ports, 4/1 - 4/4, as a Fast Uplink Span group. In this example, all four ports are connected to a wiring closet switch. Only one of the links is expected to be active at any time. The other links are redundant. For example, if the link on port 4/1 is the active link on the wiring closet switch but becomes unavailable, one of the other links takes over. Because the ports are configured in a Fast Uplink Span group, the STP convergence takes one second instead of taking at least 30 seconds using the standard STP forward delay.

You can add ports to a Fast Uplink Span group by entering the fast uplink-span command additional times with additional ports. The device can have only one Fast Uplink Span group, so all the ports you identify as Fast Uplink Span ports are members of the same group.

To remove a Fast Uplink Span group or to remove individual ports from a group, use "no" in front of the appropriate fast uplink-span command. For example, to remove ports 4/3 and 4/4 from the Fast Uplink Span group configured above, enter the following commands:

```
device(config)# no fast uplink-span ethernet 4/3 to 4/4
device(config)# write memory
```

To check the status of ports with Fast Uplink Span enabled.

```
device(config)# show span fast-uplink-span
STP instance owned by VLAN 1
Global STP (IEEE 802.1D) Parameters:
VLAN Root          Root Root    Prio Max He- Ho- Fwd Last    Chg Bridge
ID     ID          Cost Port    rity Age llo ld dly Chang   cnt Address
                           Hex sec sec
                           1 000000c100000001 2      1/3/1  8000 20 2 1 15 65      15  000011111111
Port STP Parameters:
Port  Prio Path State        Fwd    Design  Designated      Designated
Num   rity Cost          Trans  Cost   Root       Bridge
                           Hex
                           1/1/2 80 0  DISABLED  0  0 0000000000000000 0000000000000000
                           1/1/3 80 0  DISABLED  0  0 0000000000000000 0000000000000000
                           1/1/4 80 4  FORWARDING 1  2 000000c100000001 8000000011111111
                           1/1/5 80 0  DISABLED  0  0 0000000000000000 0000000000000000
                           1/1/6 80 0  DISABLED  0  0 0000000000000000 0000000000000000
                           1/1/7 80 0  DISABLED  0  0 0000000000000000 0000000000000000
                           1/1/8 80 0  DISABLED  0  0 0000000000000000 0000000000000000
                           1/1/9 80 0  DISABLED  0  0 0000000000000000 0000000000000000
```

Syntax: show span fast-uplink-span

Configuring Fast Uplink Span within a VLAN

You can also configure Fast Uplink Span on the interfaces within a VLAN.

To configure Fast Uplink Span for a VLAN, enter command such as the following.

```
device(config)#vlan 10
device(config-vlan-10)#untag ethernet 8/1 to 8/2
device(config-vlan-10)#fast uplink-span ethernet 8/1 to 8/2
```

Syntax:[no] fast uplink-span ethernet port-no

To check the status of Fast Uplink Span for a specified VLAN.

```
device(config-vlan-2)#show span vlan 2 fast-uplink-span
```

```

STP instance owned by VLAN 2
Global STP (IEEE 802.1D) Parameters:
VLAN Root      Root Root    Prio Max He- Ho- Fwd Last    Chg Bridge
ID           ID       Cost Port   rity Age llo ld dly Chang   cnt Address
                           Hex sec sec sec sec sec
                           Hex sec sec sec sec sec
2 8000000011111111 0     Root    8000 20 2 1 15 29596 0 000011111111
Port STP Parameters:
Port Prio Path State      Fwd Design Designated      Designated
Num rity Cost          Trans Cost Root Bridge
Hex
1/1/1 80 4 LISTENING 0 0 800000011111111 8000000011111111

```

Syntax: show span vlan *vlan-id* fast-uplink-span

The VLAN *vlan-id* parameter displays Fast Uplink Span information for the specified VLAN.

802.1W Rapid Spanning Tree (RSTP)

Earlier implementation by Brocade of Rapid Spanning Tree Protocol (RSTP), which was 802.1W Draft 3, provided only a subset of the IEEE 802.1W standard; whereas the 802.1W RSTP feature provides the full standard. The implementation of the 802.1W Draft 3 is referred to as RSTP Draft 3.

RSTP Draft3 will continue to be supported on Brocade devices for backward compatibility. However, customers who are currently using RSTP Draft 3 should migrate to 802.1W.

The 802.1W feature provides rapid traffic reconvergence for point-to-point links within a few milliseconds (0 - 500 milliseconds), following the failure of a bridge or bridge port. This reconvergence occurs more rapidly than the reconvergence provided by the 802.1D Spanning Tree Protocol (STP) or by RSTP Draft 3.

NOTE

This rapid convergence will not occur on ports connected to shared media devices, such as hubs. To take advantage of the rapid convergence provided by 802.1W, make sure to explicitly configure all point-to-point links in a topology.

The convergence provided by the standard 802.1W protocol occurs more rapidly than the convergence provided by previous spanning tree protocols because of the following:

- Classic or legacy 802.1D STP protocol requires a newly selected Root port to go through listening and learning stages before traffic convergence can be achieved. The 802.1D traffic convergence time is calculated using the following formula.

$$2 \times FORWARD_DELAY + BRIDGE_MAX_AGE.$$

If default values are used in the parameter configuration, convergence can take up to 50 seconds. (In this document STP will be referred to as 802.1D.)

- RSTP Draft 3 works only on bridges that have Alternate ports, which are the precalculated "next best root port". (Alternate ports provide back up paths to the root bridge.) Although convergence occurs from 0 - 500 milliseconds in RSTP Draft 3, the spanning tree topology reverts to the 802.1D convergence if an Alternate port is not found.
- Convergence in 802.1w bridge is not based on any timer values. Rather, it is based on the explicit handshakes between Designated ports and their connected Root ports to achieve convergence in less than 500 milliseconds.

Bridges and bridge port roles

A bridge in an 802.1W rapid spanning tree topology is assigned as the root bridge if it has the highest priority (lowest bridge identifier) in the topology. Other bridges are referred to as non-root bridges.

Unique roles are assigned to ports on the root and non-root bridges. Role assignments are based on the following information contained in the Rapid Spanning Tree Bridge Packet Data Unit (RST BPDU):

- Root bridge ID
- Path cost value
- Transmitting bridge ID
- Designated port ID

The 802.1W algorithm uses this information to determine if the RST BPDU received by a port is superior to the RST BPDU that the port transmits. The two values are compared in the order as given above, starting with the Root bridge ID. The RST BPDU with a lower value is considered superior. The superiority and inferiority of the RST BPDU is used to assign a role to a port.

If the value of the received RST BPDU is the same as that of the transmitted RST BPDU, then the port ID in the RST BPDUs are compared. The RST BPDU with the lower port ID is superior. Port roles are then calculated appropriately.

The port role is included in the BPDU that it transmits. The BPDU transmitted by an 802.1W port is referred to as an RST BPDU, while it is operating in 802.1W mode.

Ports can have one of the following roles:

- **Root** - Provides the lowest cost path to the root bridge from a specific bridge
- **Designated** - Provides the lowest cost path to the root bridge from a LAN to which it is connected
- **Alternate** - Provides an alternate path to the root bridge when the root port goes down
- **Backup** - Provides a backup to the LAN when the Designated port goes down
- **Disabled** - Has no role in the topology

Assignment of port roles

At system start-up, all 802.1W-enabled bridge ports assume a Designated role. Once start-up is complete, the 802.1W algorithm calculates the superiority or inferiority of the RST BPDU that is received and transmitted on a port.

On a root bridge, each port is assigned a Designated port role, except for ports on the same bridge that are physically connected together. In these type of ports, the port that receives the superior RST BPDU becomes the Backup port , while the other port becomes the Designated port .

On non-root bridges, ports are assigned as follows:

- The port that receives the RST BPDU with the lowest path cost from the root bridge becomes the Root port .
- If two ports on the same bridge are physically connected, the port that receives the superior RST BPDU becomes the Backup port , while the other port becomes the Designated port .
- If a non-root bridge already has a Root port, then the port that receives an RST BPDU that is superior to those it can transmit becomes the Alternate port .
- If the RST BPDU that a port receives is inferior to the RST BPDUs it transmits, then the port becomes a Designated port .
- If the port is down or if 802.1W is disabled on the port, that port is given the role of Disabled port . Disabled ports have no role in the topology. However, if 802.1W is enabled on a port with a link down and the link of that port comes up, then that port assumes one of the following port roles: Root, Designated, Alternate, or Backup.

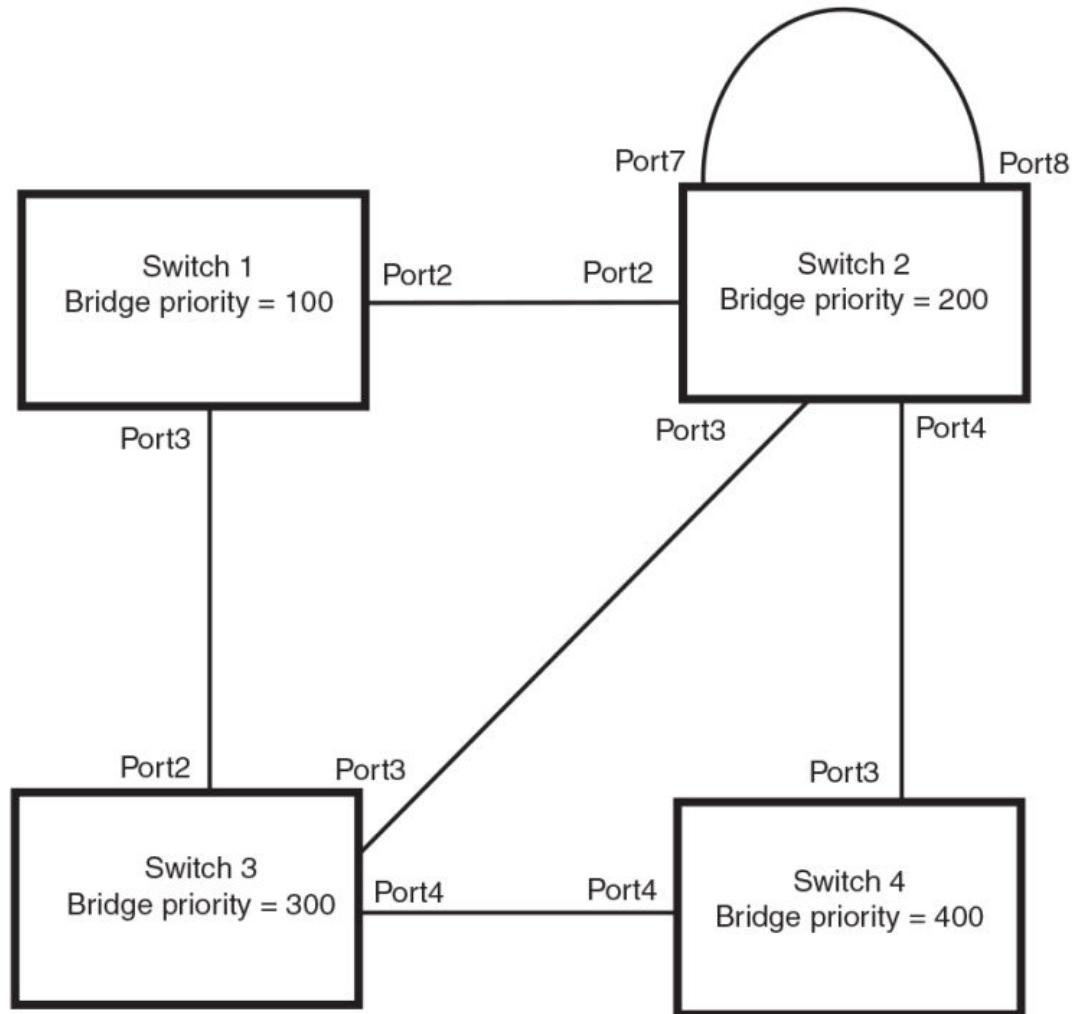
The following example ([Figure 37](#)) explains role assignments in a simple RSTP topology.

NOTE

All examples in this document assume that all ports in the illustrated topologies are point-to-point links and are homogeneous (they have the same path cost value) unless otherwise specified.

The topology in the following figure contains four bridges. Switch 1 is the root bridge since it has the lowest bridge priority. Switch 2 through Switch 4 are non-root bridges.

FIGURE 37 Simple 802.1W topology



Assignment of ports on Switch 1

All ports on Switch 1, the root bridge, are assigned Designated port roles.

Assignment of ports on Switch 2

Port2 on Switch 2 directly connects to the root bridge; therefore, Port2 is the Root port.

The bridge priority value on Switch 2 is superior to that of Switch 3 and Switch 4; therefore, the ports on Switch 2 that connect to Switch 3 and Switch 4 are given the Designated port role.

Furthermore, Port7 and Port8 on Switch 2 are physically connected. The RST BPDUs transmitted by Port7 are superior to those Port8 transmits. Therefore, Port8 is the Backup port and Port7 is the Designated port.

Assignment of ports on Switch 3

Port2 on Switch 3 directly connects to the Designated port on the root bridge; therefore, it assumes the Root port role.

The root path cost of the RST BPDUs received on Port4/Switch 3 is inferior to the RST BPDUs transmitted by the port; therefore, Port4/Switch 3 becomes the Designated port.

Similarly Switch 3 has a bridge priority value inferior to Switch 2. Port3 on Switch 3 connects to Port 3 on Switch 2. This port will be given the Alternate port role, since a Root port is already established on this bridge.

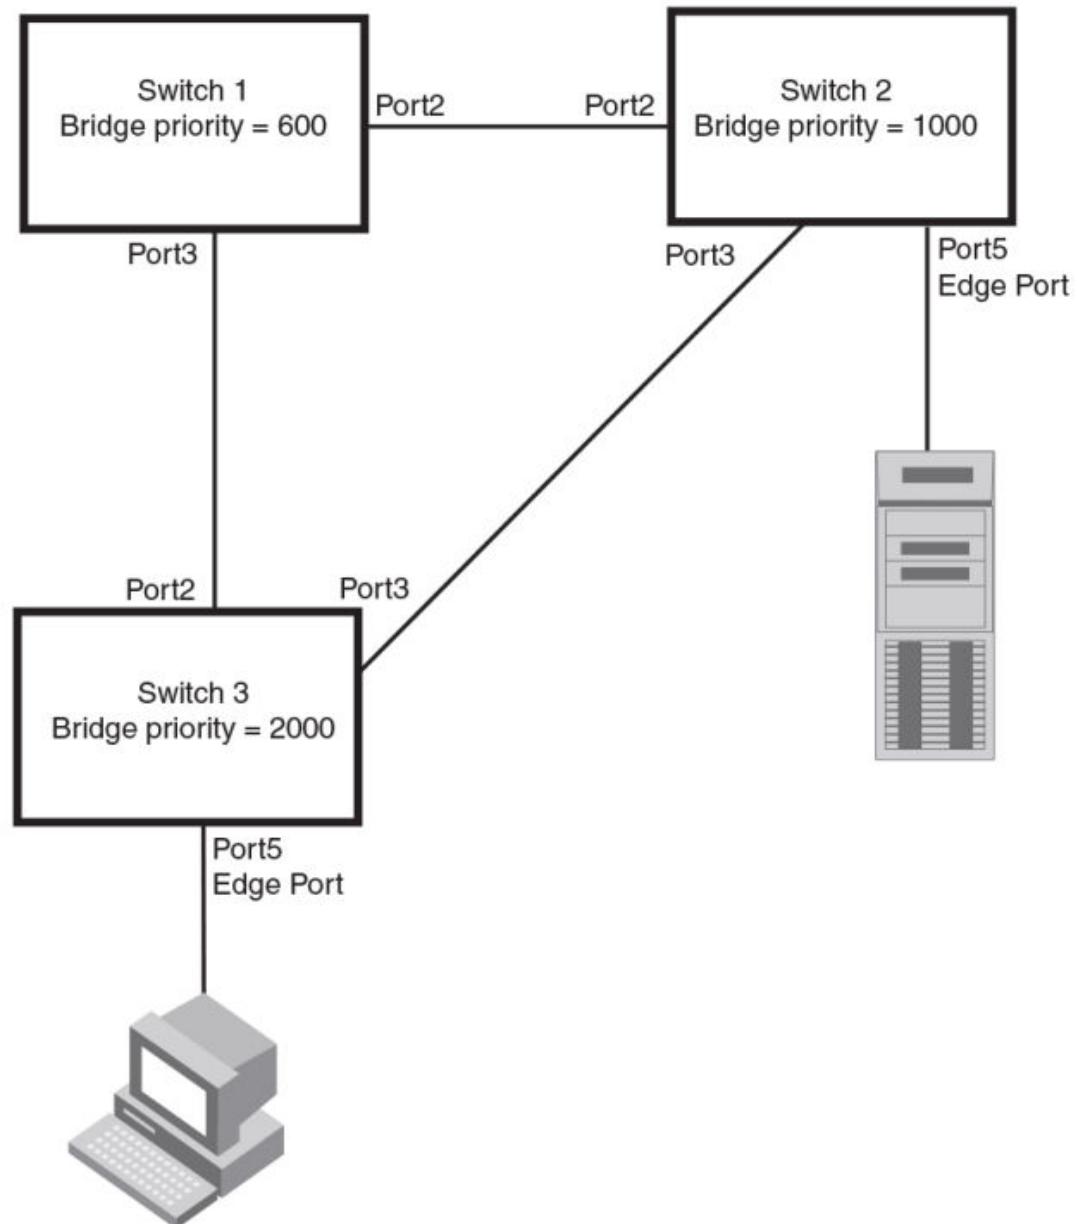
Assignment of ports on Switch 4

Switch 4 is not directly connected to the root bridge. It has two ports with superior incoming RST BPDUs from two separate LANs: Port3 and Port4. The RST BPDUs received on Port3 are superior to the RST BPDUs received on port 4; therefore, Port3 becomes the Root port and Port4 becomes the Alternate port.

Edge ports and edge port roles

The Brocade implementation of 802.1W allows ports that are configured as Edge ports to be present in an 802.1W topology. ([Figure 38](#)). Edge ports are ports of a bridge that connect to workstations or computers. Edge ports do not register any incoming BPDU activities.

Edge ports assume Designated port roles. Port flapping does not cause any topology change events on Edge ports since 802.1W does not consider Edge ports in the spanning tree calculations.

FIGURE 38 Topology with edge ports

However, if any incoming RST BPDU is received from a previously configured Edge port, 802.1W automatically makes the port as a non-edge port. This is extremely important to ensure a loop free Layer 2 operation since a non-edge port is part of the active RSTP topology.

The 802.1W protocol can auto-detect an Edge port and a non-edge port. An administrator can also configure a port to be an Edge port using the CLI. It is recommended that Edge ports are configured explicitly to take advantage of the Edge port feature, instead of allowing the protocol to auto-detect them.

Point-to-point ports

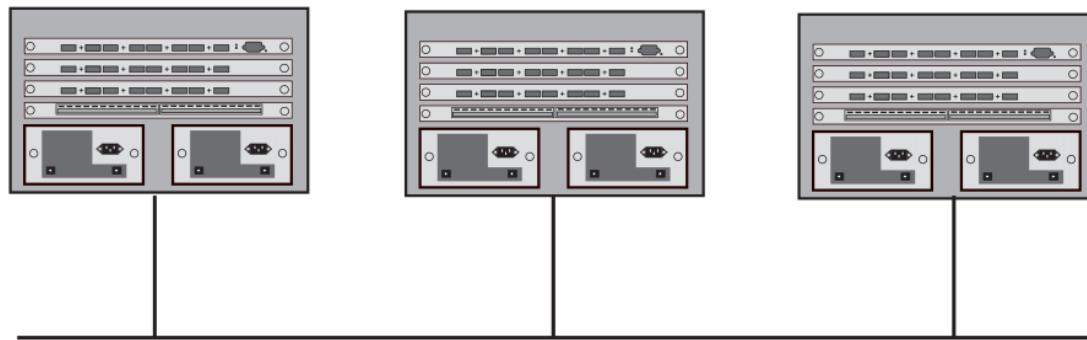
To take advantage of the 802.1W features, ports on an 802.1W topology should be explicitly configured as point-to-point links using the CLI. Shared media should not be configured as point-to-point links.

NOTE

Configuring shared media or non-point-to-point links as point-to-point links could lead to Layer 2 loops.

The topology in the following figure is an example of shared media that should not be configured as point-to-point links. In this figure, a port on a bridge communicates or is connected to at least two ports.

FIGURE 39 Example of shared media



Bridge port states

Ports roles can have one of the following states:

- Forwarding - 802.1W is allowing the port to send and receive all packets.
- Discarding - 802.1W has blocked data traffic on this port to prevent a loop. The device or VLAN can reach the root bridge using another port, whose state is forwarding. When a port is in this state, the port does not transmit or receive data frames, but the port does continue to receive RST BPDUs. This state corresponds to the listening and blocking states of 802.1D.
- Learning - 802.1W is allowing MAC entries to be added to the filtering database but does not permit forwarding of data frames. The device can learn the MAC addresses of frames that the port receives during this state and make corresponding entries in the MAC table.
- Disabled - The port is not participating in 802.1W. This can occur when the port is disconnected or 802.1W is administratively disabled on the port.

A port on a non-root bridge with the role of Root port is always in a forwarding state. If another port on that bridge assumes the Root port role, then the old Root port moves into a discarding state as it assumes another port role.

A port on a non-root bridge with a Designated role starts in the discarding state. When that port becomes elected to the Root port role, 802.1W quickly places it into a forwarding state. However, if the Designated port is an Edge port, then the port starts and stays in a forwarding state and it cannot be elected as a Root port.

A port with an Alternate or Backup role is always in a discarding state. If the port role changes to Designated, then the port changes into a forwarding state.

If a port on one bridge has a Designated role and that port is connected to a port on another bridge that has an Alternate or Backup role, the port with a Designated role cannot be given a Root port role until two instances of the forward delay timer expires on that port.

Edge port and non-edge port states

As soon as a port is configured as an Edge port using the CLI, it goes into a forwarding state instantly (in less than 100 msec).

When the link to a port comes up and 802.1W detects that the port is an Edge port, that port instantly goes into a forwarding state.

If 802.1W detects that port as a non-edge port, the port state is changed as determined by the result of processing the received RST BPDU. The port state change occurs within four seconds of link up or after two hello timer expires on the port.

Changes to port roles and states

To achieve convergence in a topology, a port role and state changes as it receives and transmits new RST BPDUs. Changes in a port role and state constitute a topology change. Besides the superiority and inferiority of the RST BPDU, bridge-wide and per-port state machines are used to determine a port role as well as a port state. Port state machines also determine when port role and state changes occur.

Port Role Selection state machines

The bridge uses the Port Role Selection state machine to determine if port role changes are required on the bridge. This state machine performs a computation when one of the following events occur:

- New information is received on any port on the bridge
- The timer expires for the current information on a port on the bridge

Each port uses the following state machines:

- Port Information - This state machine keeps track of spanning-tree information currently used by the port. It records the origin of the information and ages out any information that was derived from an incoming BPDU.
- Port Role Transition - This state machine keeps track of the current port role and transitions the port to the appropriate role when required. It moves the Root port and the Designated port into forwarding states and moves the Alternate and Backup ports into discarding states.
- Port Transmit - This state machine is responsible for BPDU transmission. It checks to ensure only the maximum number of BPDUs per hello interval are sent every second. Based on what mode it is operating in, it sends out either legacy BPDUs or RST BPDUs. In this document legacy BPDUs are also referred to as STP BPDUs.
- Port Protocol Migration - This state machine deals with compatibility with 802.1D bridges. When a legacy BPDU is detected on a port, this state machine configures the port to transmit and receive legacy BPDUs and operate in the legacy mode.
- Topology Change - This state machine detects, generates, and propagates topology change notifications. It acknowledges Topology Change Notice (TCN) messages when operating in 802.1D mode. It also flushes the MAC table when a topology change event takes place.
- Port State Transition - This state machine transitions the port to a discarding, learning, or forwarding state and performs any necessary processing associated with the state changes.
- Port Timers - This state machine is responsible for triggering any of the state machines described above, based on expiration of specific port timers.

In contrast to the 802.1D standard, the 802.1W standard does not have any bridge specific timers. All timers in the CLI are applied on a per-port basis, even though they are configured under bridge parameters.

802.1W state machines attempt to quickly place the ports into either a forwarding or discarding state. Root ports are quickly placed in forwarding state when both of the following events occur:

- It is assigned to be the Root port.
- It receives an RST BPDU with a proposal flag from a Designated port. The proposal flag is sent by ports with a Designated role when they are ready to move into a forwarding state.

When a the role of Root port is given to another port, the old Root port is instructed to reroot. The old Root port goes into a discarding state and negotiates with its peer port for a new role and a new state. A peer port is the port on the other bridge to which the port is connected. For example, in [Figure 40](#), Port1 of Switch 200 is the peer port of Port2 of Switch 100.

A port with a Designated role is quickly placed into a forwarding state if one of the following occurs:

- The Designated port receives an RST BPDU that contains an agreement flag from a Root port
- The Designated port is an Edge port

However, a Designated port that is attached to an Alternate port or a Backup port must wait until the forward delay timer expires twice on that port while it is still in a Designated role, before it can proceed to the forwarding state.

Backup ports are quickly placed into discarding states.

Alternate ports are quickly placed into discarding states.

A port operating in 802.1W mode may enter a learning state to allow MAC entries to be added to the filtering database; however, this state is transient and lasts only a few milliseconds, if the port is operating in 802.1W mode and if the port meets the conditions for rapid transition.

Handshake mechanisms

To rapidly transition a Designated or Root port into a forwarding state, the Port Role Transition state machine uses handshake mechanisms to ensure loop free operations. It uses one type of handshake if no Root port has been assigned on a bridge, and another type if a Root port has already been assigned.

Handshake when no root port is elected

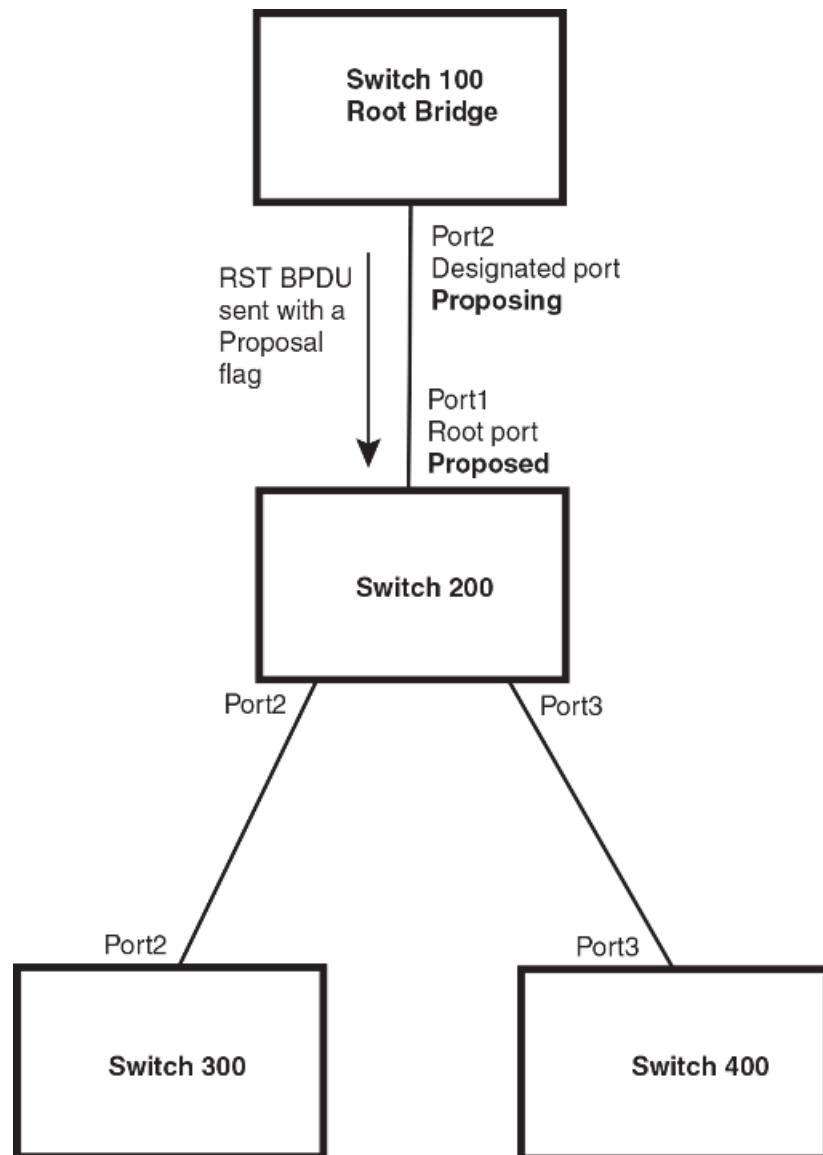
If a Root port has not been assigned on a bridge, 802.1W uses the Proposing -> Proposed -> Sync -> Synced -> Agreed handshake:

- Proposing - The Designated port on the root bridge sends an RST BPDU packet to its peer port that contains a proposal flag. The proposal flag is a signal that indicates that the Designated port is ready to put itself in a forwarding state ([Figure 40](#)). The Designated port continues to send this flag in its RST BPDU until it is placed in a forwarding state ([Figure 43](#)) or is forced to operate in 802.1D mode. (Refer to [Compatibility of 802.1W with 802.1D](#) on page 287).
- Proposed - When a port receives an RST BPDU with a proposal flag from the Designated port on its point-to-point link, it asserts the Proposed signal and one of the following occurs ([Figure 40](#)):
 - If the RST BPDU that the port receives is superior to what it can transmit, the port assumes the role of a Root port. (Refer to the section on [Bridges and bridge port roles](#) on page 260.)
 - If the RST BPDU that the port receives is inferior to what it can transmit, then the port is given the role of Designated port.

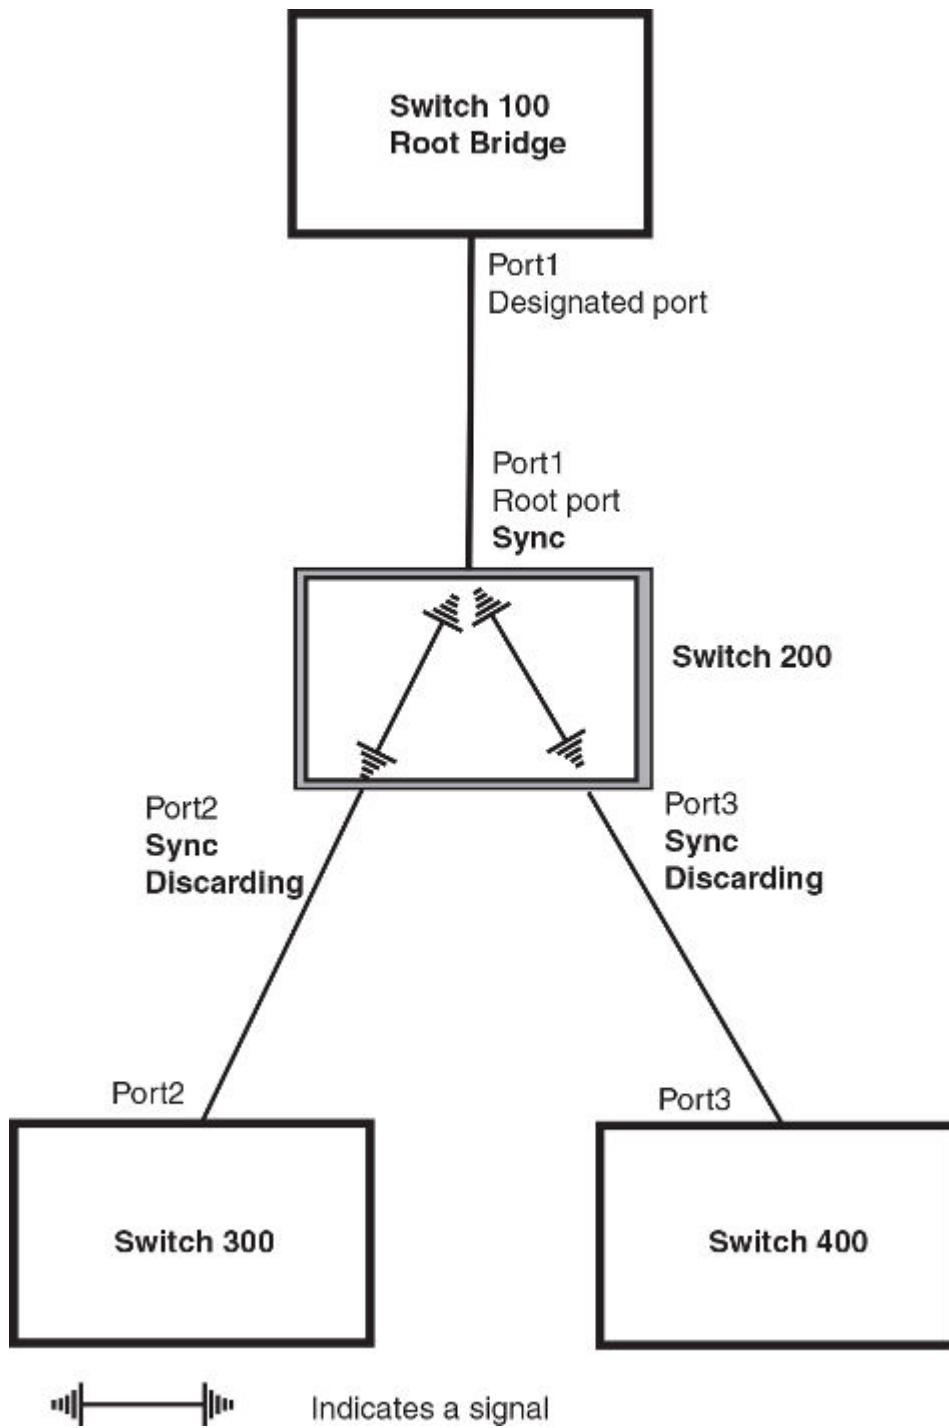
NOTE

Proposed will never be asserted if the port is connected on a shared media link.

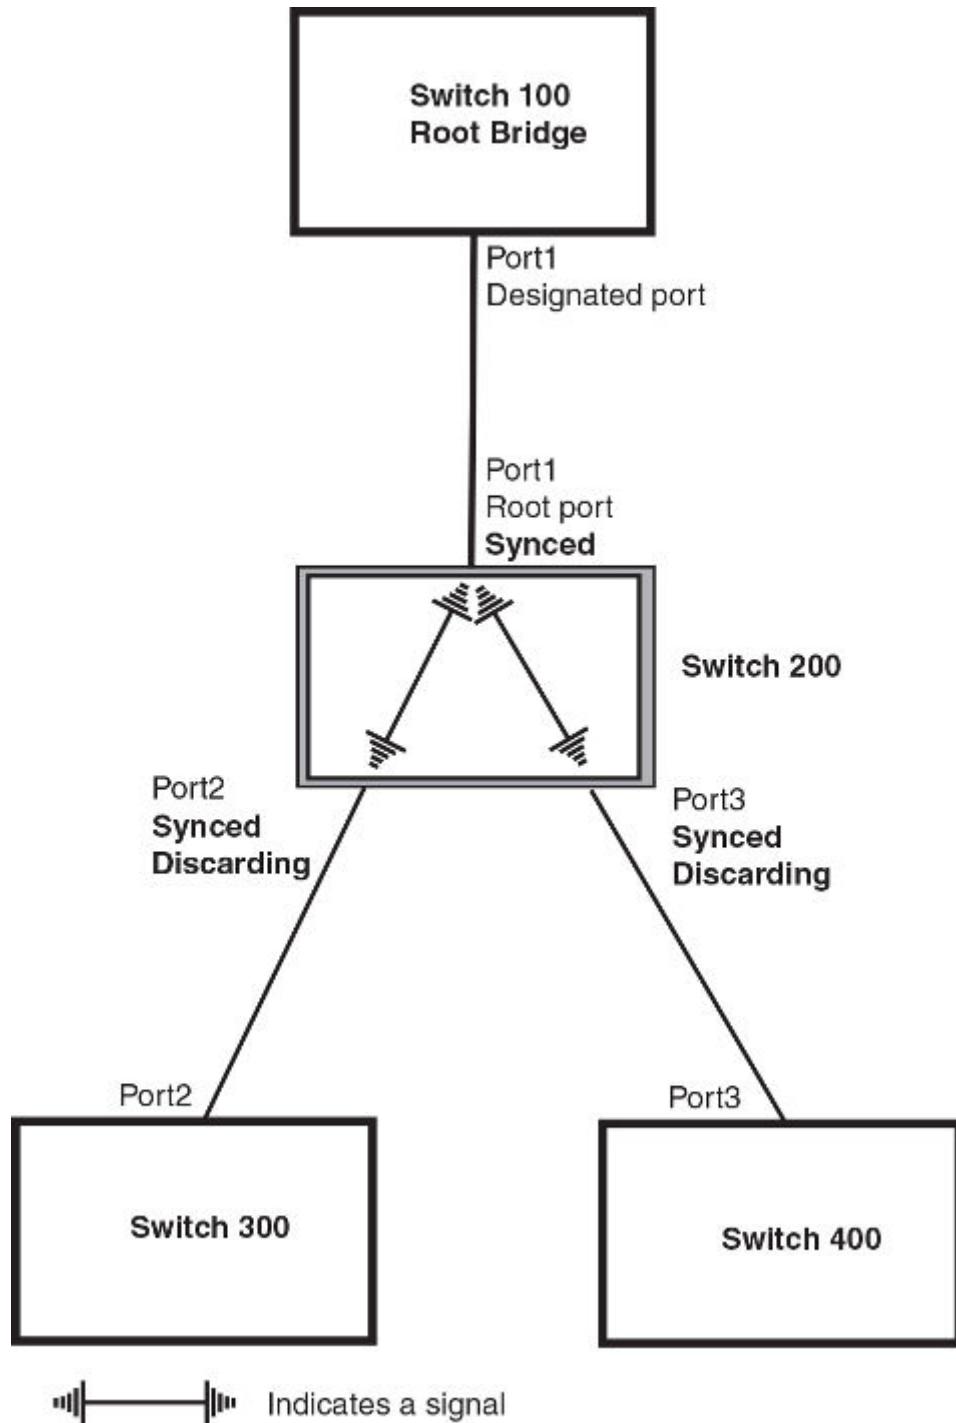
In the following figure, Port3/Switch 200 is elected as the Root port.

FIGURE 40 Proposing and proposed stage

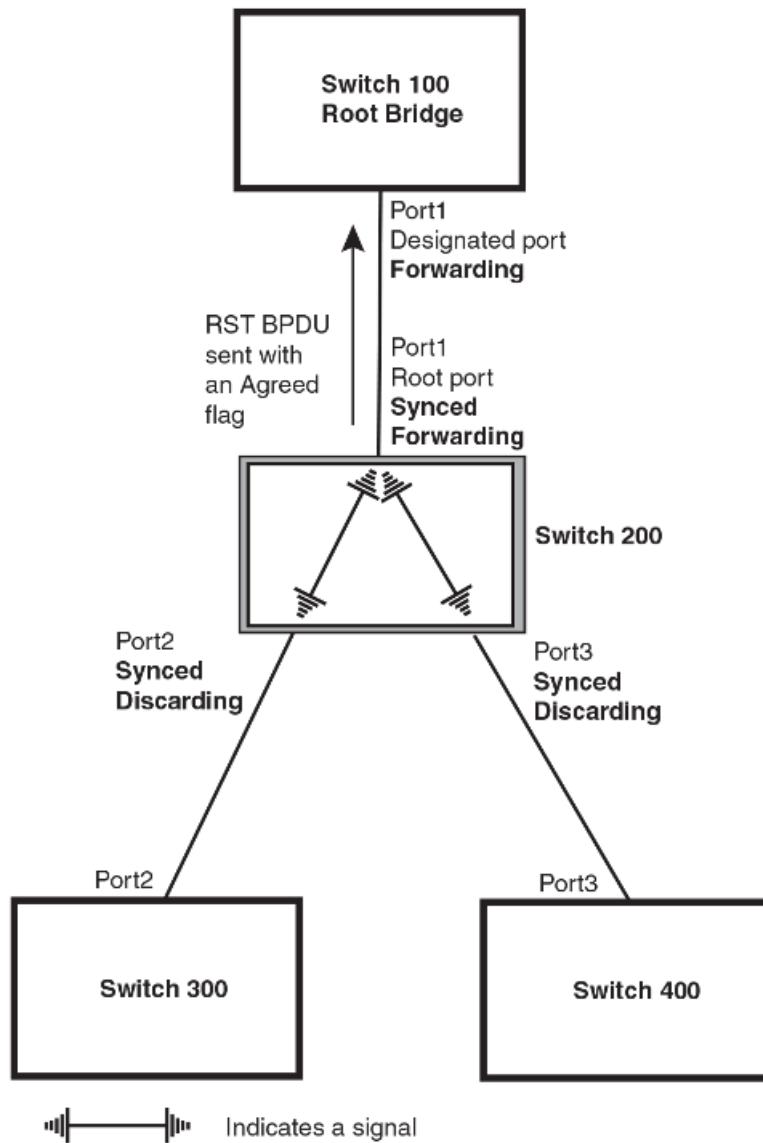
- Sync - Once the Root port is elected, it sets a sync signal on all the ports on the bridge. The signal tells the ports to synchronize their roles and states ([Figure 41](#)). Ports that are non-edge ports with a role of Designated port change into a discarding state. These ports have to negotiate with their peer ports to establish their new roles and states.

FIGURE 41 Sync stage

- Synced - Once the Designated port changes into a discarding state, it asserts a synced signal. Immediately, Alternate ports and Backup ports are synced. The Root port monitors the synced signals from all the bridge ports. Once all bridge ports asserts a synced signal, the Root port asserts its own synced signal as shown in the following figure.

FIGURE 42 Synced stage

- Agreed - The Root port sends back an RST BPDU containing an agreed flag to its peer Designated port and moves into the forwarding state. When the peer Designated port receives the RST BPDU, it rapidly transitions into a forwarding state.

FIGURE 43 Agree stage

At this point, the handshake mechanism is complete between Switch 100, the root bridge, and Switch 200.

Switch 200 updates the information on the Switch 200 Designated ports (Port2 and Port3) and identifies the new root bridge. The Designated ports send RST BPDUs, containing proposal flags, to their downstream bridges, without waiting for the hello timers to expire on them. This process starts the handshake with the downstream bridges.

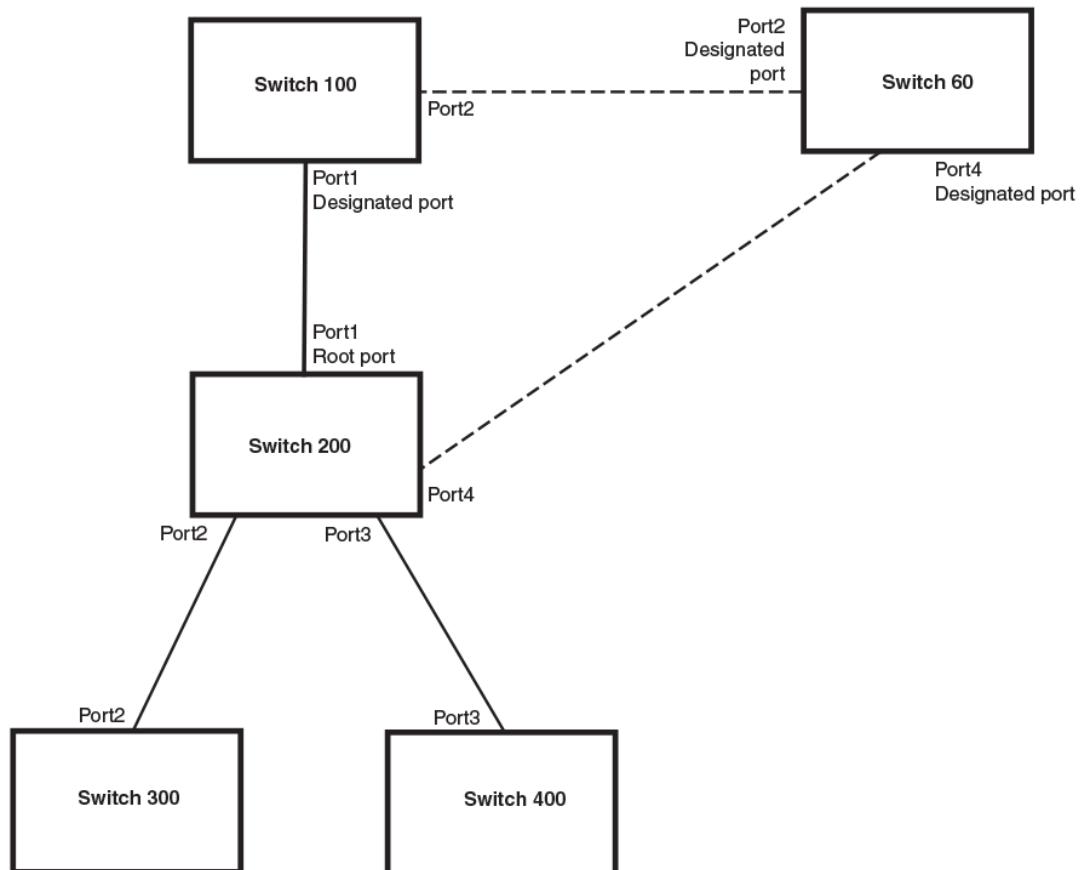
For example, Port2/Switch 200 sends an RST BPDU to Port2/Switch 300 that contains a proposal flag. Port2/Switch 300 asserts a proposed signal. Ports in Switch 300 then set sync signals on the ports to synchronize and negotiate their roles and states. Then the ports assert a synced signal and when the Root port in Switch 300 asserts its synced signal, it sends an RST BPDU to Switch 200 with an agreed flag.

This handshake is repeated between Switch 200 and Switch 400 until all Designated and Root ports are in forwarding states.

Handshake when a root port has been elected

If a non-root bridge already has a Root port, 802.1W uses a different type of handshake. For example, in the following figure, a new root bridge is added to the topology.

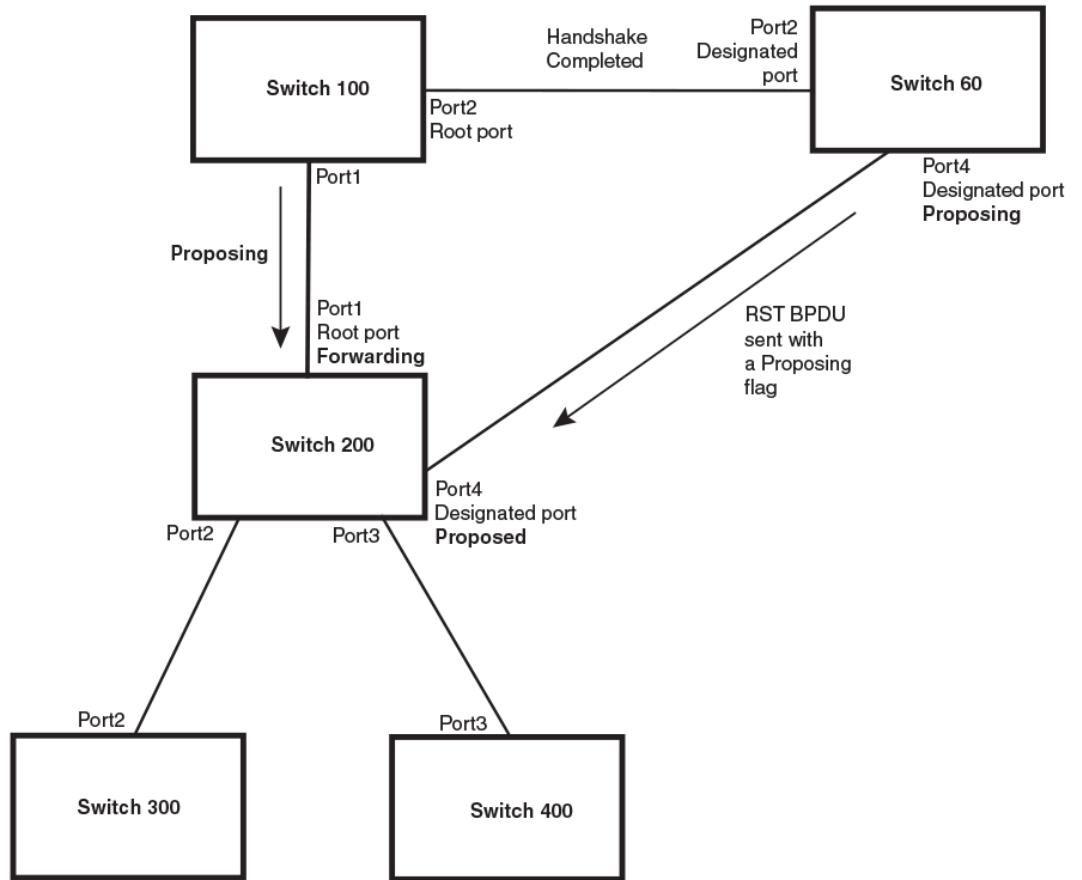
FIGURE 44 Addition of a new root bridge



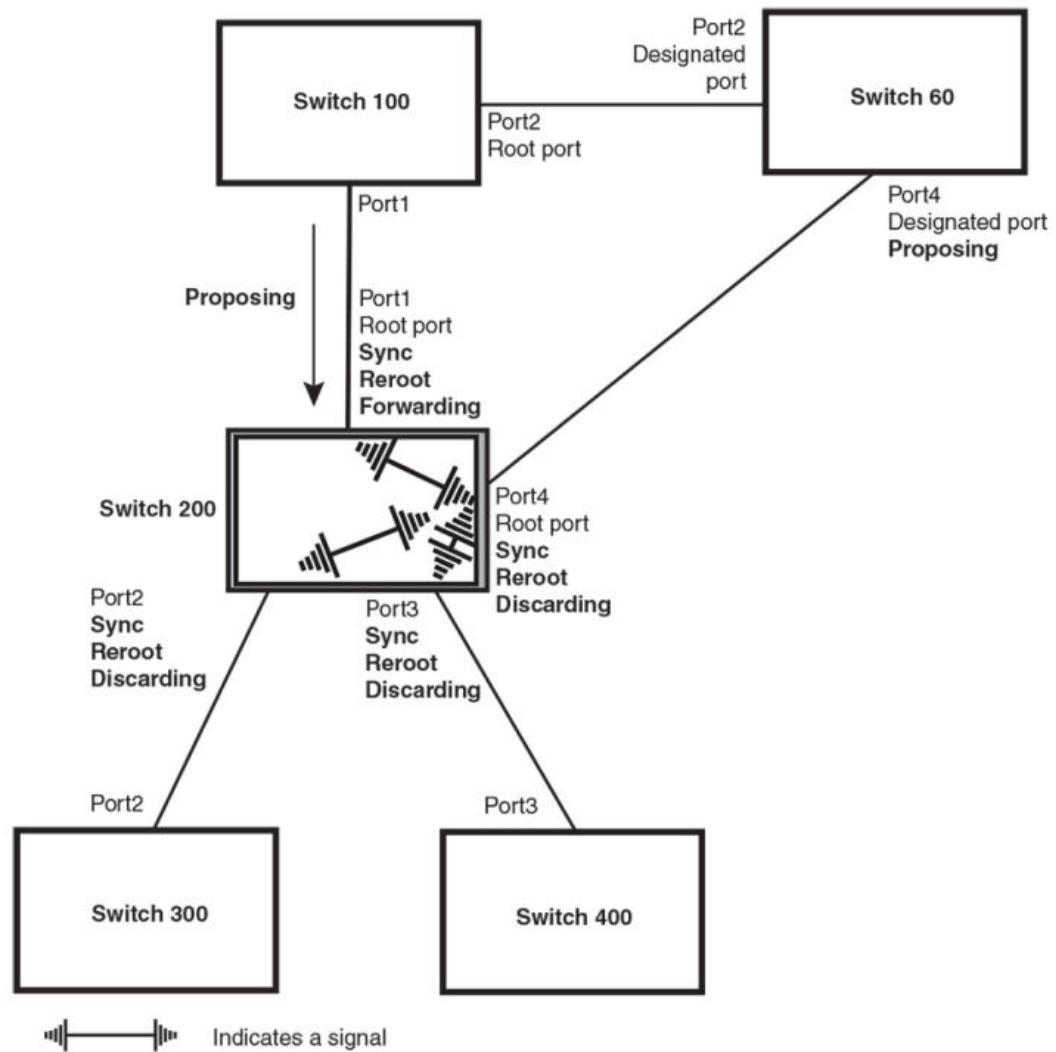
The handshake that occurs between Switch 60 and Switch 100 follows the one described in the previous section ([Handshake when no root port is elected](#)). The former root bridge becomes a non-root bridge and establishes a Root port ([Figure 45](#)).

However, since Switch 200 already had a Root port in a forwarding state, 802.1W uses the Proposing -> Proposed -> Sync and Reroot -> Sync and Rerooted -> Rerooted and Synced -> Agreed handshake:

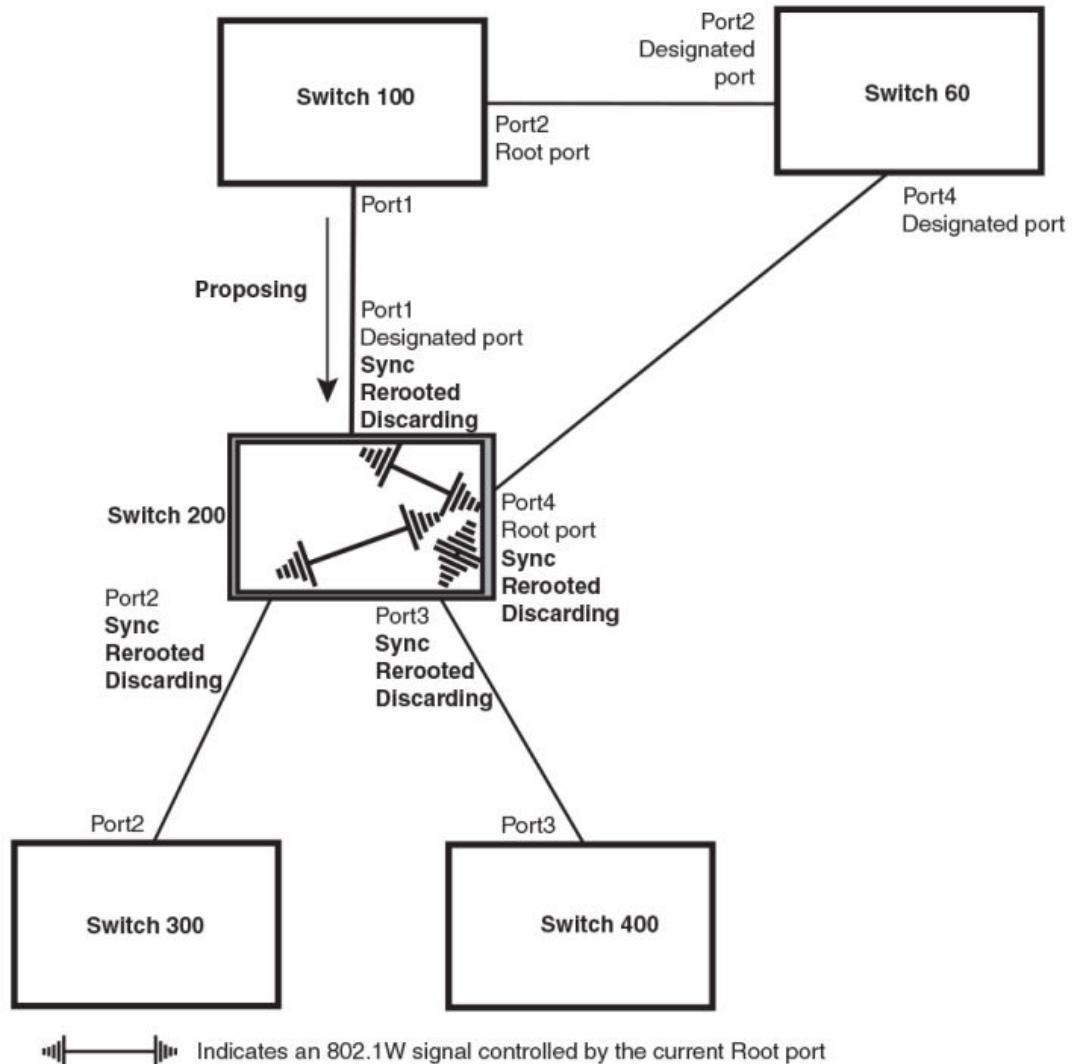
- Proposing and Proposed - The Designated port on the new root bridge (Port4/Switch 60) sends an RST BPDU that contains a proposing signal to Port4/Switch 200 to inform the port that it is ready to put itself in a forwarding state ([Figure 45](#)). 802.1W algorithm determines that the RST BPDU that Port4/Switch 200 received is superior to what it can generate, so Port4/Switch 200 assumes a Root port role.

FIGURE 45 New root bridge sending a proposal flag

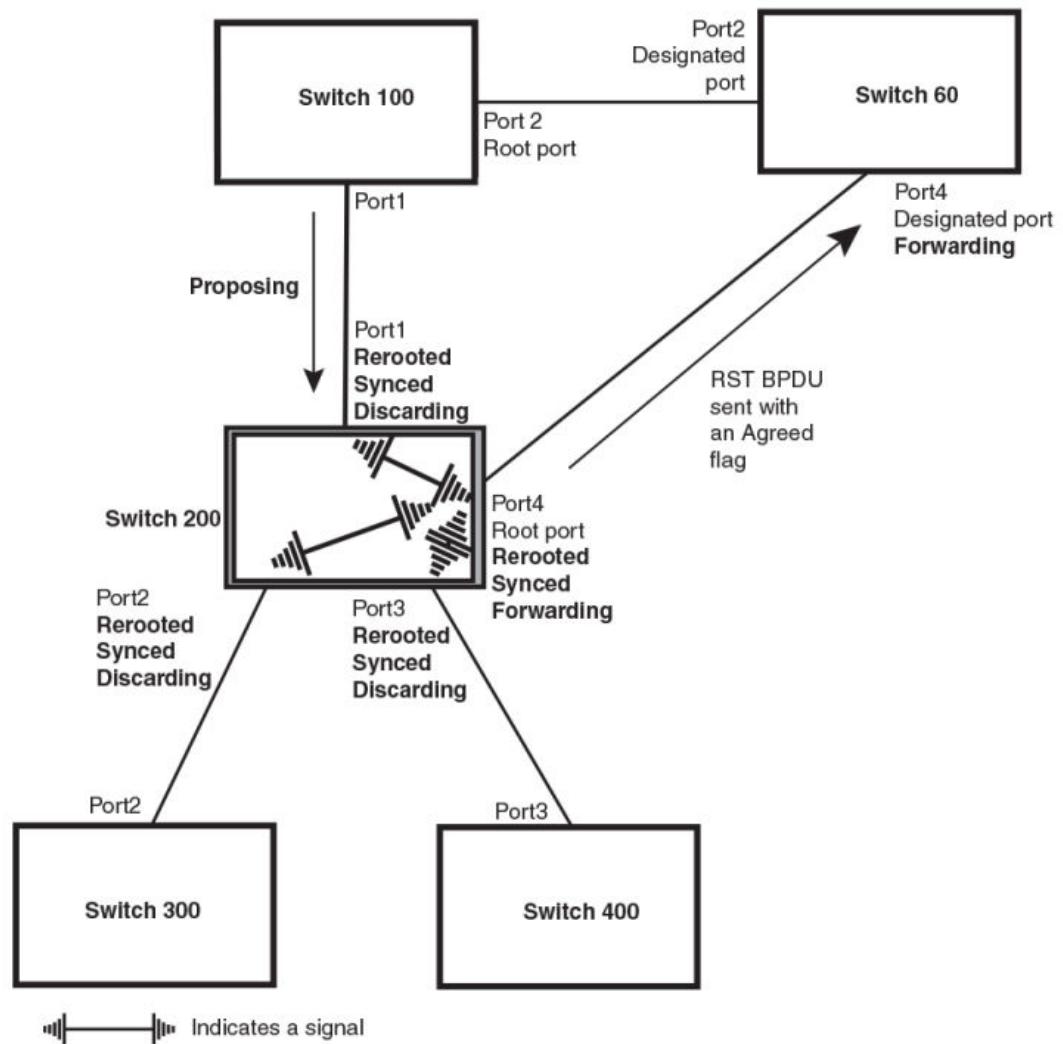
- Sync and Reroot - The Root port then asserts a sync and a reroot signal on all the ports on the bridge. The signal tells the ports that a new Root port has been assigned and they are to renegotiate their new roles and states. The other ports on the bridge assert their sync and reroot signals. Information about the old Root port is discarded from all ports. Designated ports change into discarding states as shown in the following figure.

FIGURE 46 Sync and reroot

- Sync and Rerooted - When the ports on Switch 200 have completed the reroot phase, they assert their rerooted signals and continue to assert their sync signals as they continue in their discarding states. They also continue to negotiate their roles and states with their peer ports as shown in the following figure.

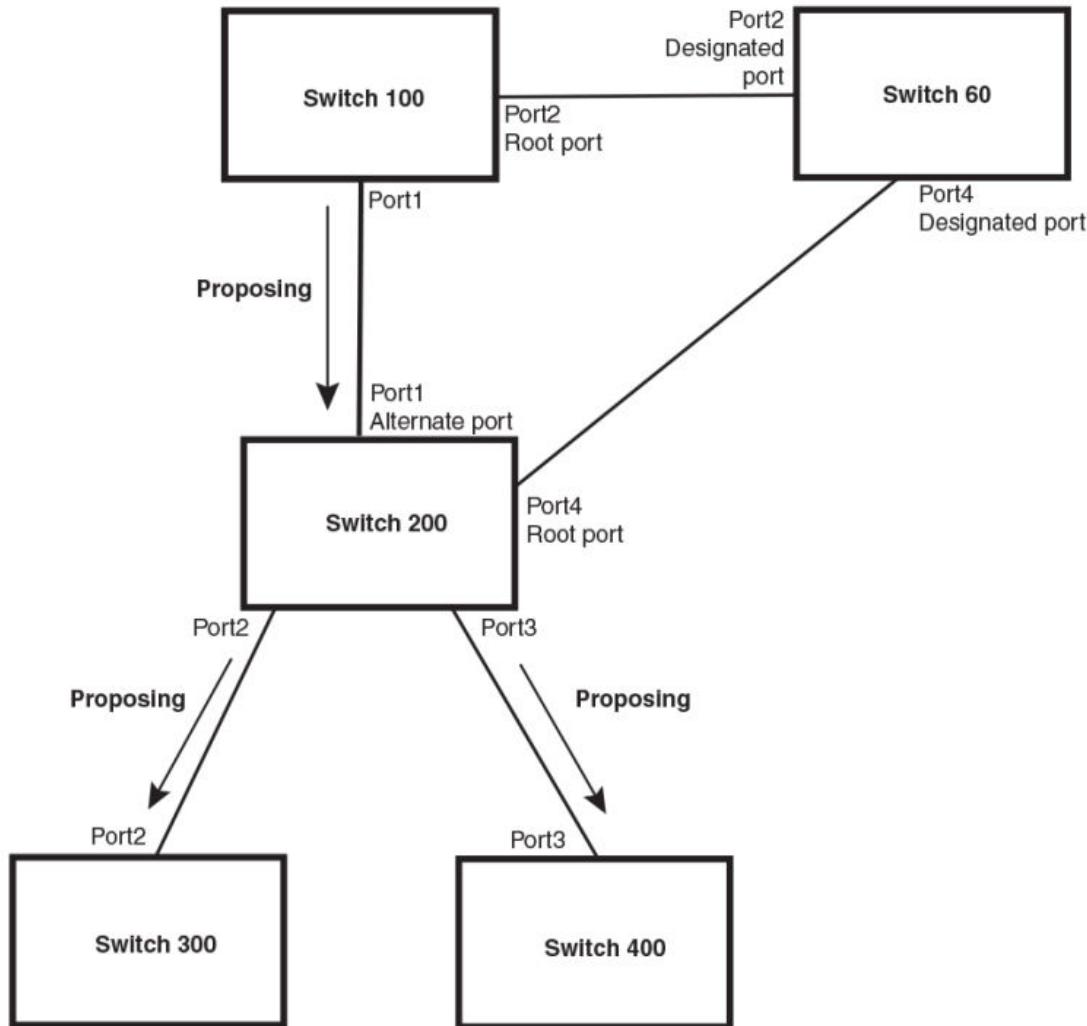
FIGURE 47 Sync and rerooted

- Synced and Agree - When all the ports on the bridge assert their synced signals, the new Root port asserts its own synced signal and sends an RST BPDU to Port4/Switch 60 that contains an agreed flag as shown in the following figure. The Root port also moves into a forwarding state.

FIGURE 48 Rerooted, synced, and agreed

The old Root port on Switch 200 becomes an Alternate Port as shown in the following figure. Other ports on that bridge are elected to appropriate roles.

The Designated port on Switch 60 goes into a forwarding state once it receives the RST BPDU with the agreed flag.

FIGURE 49 Handshake completed after election of new root port

Recall that Switch 200 sent the agreed flag to Port4/Switch 60 and not to Port1/Switch 100 (the port that connects Switch 100 to Switch 200). Therefore, Port1/Switch 100 does not go into forwarding state instantly. It waits until two instances of the forward delay timer expires on the port before it goes into forwarding state.

At this point the handshake between the Switch 60 and Switch 200 is complete.

The remaining bridges (Switch 300 and Switch 400) may have to go through the reroot handshake if a new Root port needs to be assigned.

802.1W convergence in a simple topology

The examples in this section illustrate how 802.1W convergence occurs in a simple Layer 2 topology at start-up.

NOTE

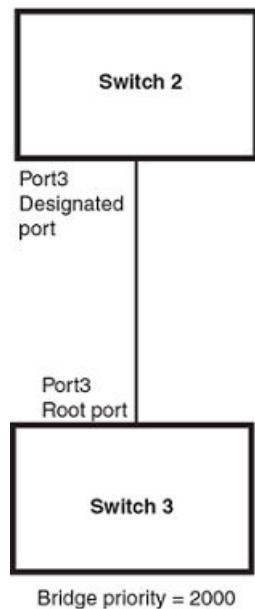
The remaining examples assume that the appropriate handshake mechanisms occur as port roles and states change.

Convergence at start up

In the following figure, two bridges Switch 2 and Switch 3 are powered up. There are point-to-point connections between Port3/Switch 2 and Port3/Switch 3.

FIGURE 50 Convergence between two bridges

Bridge priority = 1500



Bridge priority = 2000

At power up, all ports on Switch 2 and Switch 3 assume Designated port roles and are at discarding states before they receive any RST BPDU.

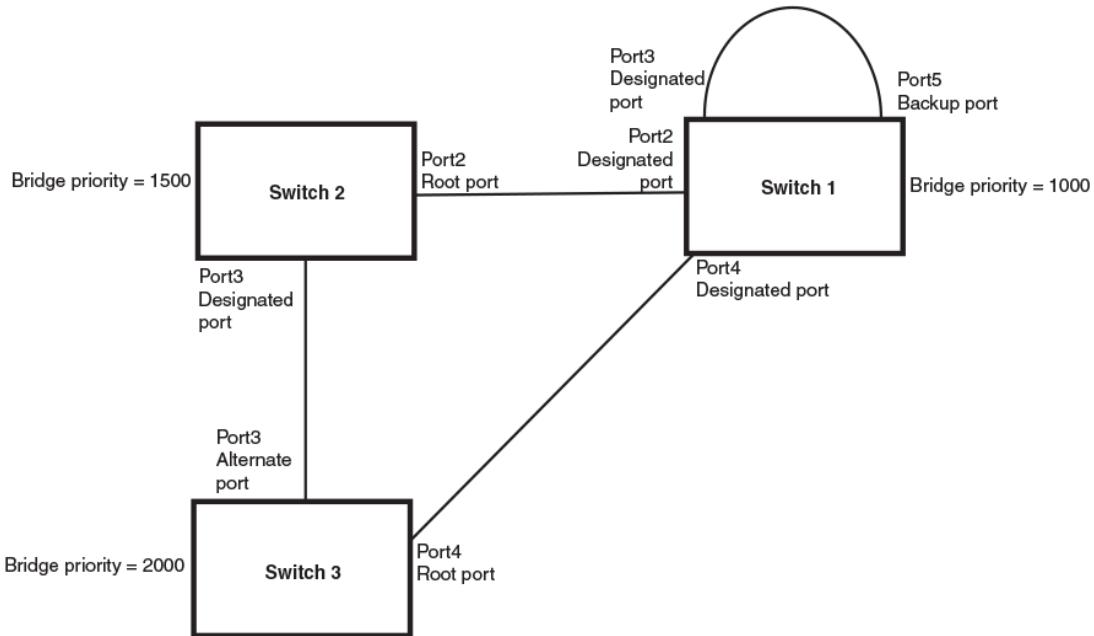
Port3/Switch 2, with a Designated role, transmits an RST BPDU with a proposal flag to Port3/Switch 3. A ports with a Designated role sends the proposal flag in its RST BPDU when they are ready to move to a forwarding state.

Port3/Switch 3, which starts with a role of Designated port, receives the RST BPDU and finds that it is superior to what it can transmit; therefore, Port3/Switch 3 assumes a new port role, that of a Root port. Port3/Switch 3 transmits an RST BPDU with an agreed flag back to Switch 2 and immediately goes into a forwarding state.

Port3/Switch 2 receives the RST BPDU from Port3/Switch 3 and immediately goes into a forwarding state.

Now 802.1W has fully converged between the two bridges, with Port3/Switch 3 as an operational root port in forwarding state and Port3/Switch 2 as an operational Designated port in forwarding state.

Next, Switch 1 is powered up. See the following figure.

FIGURE 51 Simple Layer 2 topology

The point-to-point connections between the three bridges are as follows:

- Port2/Switch 1 and Port2/Switch 2
- Port4/Switch 1 and Port4/Switch 3
- Port3/Switch 2 and Port3/Switch 3

Ports 3 and 5 on Switch 1 are physically connected.

At start up, the ports on Switch 1 assume Designated port roles, which are in discarding state. The ports begin sending RST BPDUUs with proposal flags. The flags indicate the ID of the bridge that the ports belong to, and the bridge that the ports understand to be the root bridge. The switch that eventually becomes the downstream neighbor is the only switch that sends a BPDU with the agreement bit set.

When Port4/Switch 3 receives these RST BPDUUs 802.1W algorithm determines that they are better than the RST BPDUUs that were previously received on Port3/Switch 3. Port4/Switch 3 is now selected as Root port. This new assignment signals Port3/Switch 3 to begin entering the discarding state and to assume an Alternate port role. As it goes through the transition, Port3/Switch 3 negotiates a new role and state with its peer port, Port3/Switch 2.

Port4/Switch 3 sends an RST BPDU with an agreed flag to Port4/Switch 1. Both ports go into forwarding states.

Port2/Switch 2 receives an RST BPDU. The 802.1W algorithm evaluates the BPDU and determines that it is superior to any BPDU that any other port on Switch 2 can transmit. Port2/Switch 2 assumes the role of a Root port.

The new Root port then signals all ports on the bridge to start synchronization. Since none of the ports are Edge ports, they all enter the discarding state and assume the role of Designated ports. Port3/Switch 2, which previously had a Designated role with a forwarding state, starts the discarding state. They also negotiate port roles and states with their peer ports. Port3/Switch 2 also sends an RST BPDU to Port3/Switch 3 with a proposal flag to request permission go into a forwarding state.

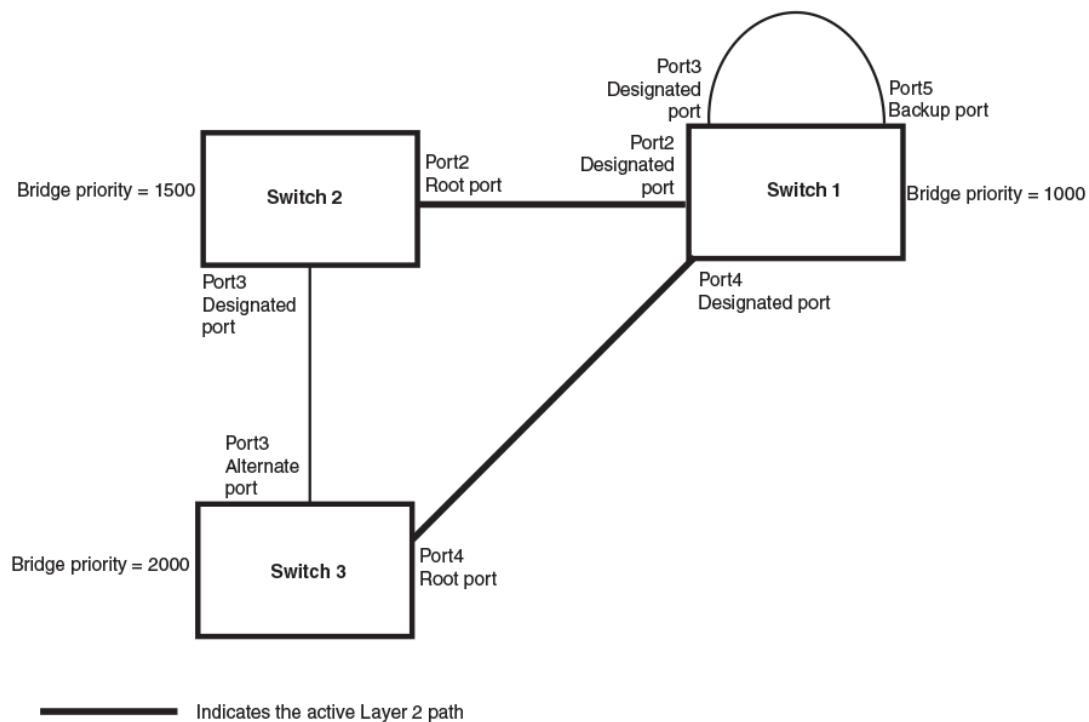
The Port2/Switch 2 bridge also sends an RST BPDU with an agreed flag to Port2/Switch 1 that Port2 is the new Root port. Both ports go into forwarding states.

Now, Port3/Switch 3 is currently in a discarding state and is negotiating a port role. It received RST BPDUs from Port3/Switch 2. The 802.1W algorithm determines that the RST BPDUs Port3/Switch 3 received are superior to those it can transmit; however, they are not superior to those that are currently being received by the current Root port (Port4). Therefore, Port3 retains the role of Alternate port.

Ports 3/Switch 1 and Port5/Switch 1 are physically connected. Port5/Switch 1 received RST BPDUs that are superior to those received on Port3/Switch 1; therefore, Port5/Switch 1 is given the Backup port role while Port3 is given the Designated port role. Port3/Switch 1, does not go directly into a forwarding state. It waits until the forward delay time expires twice on that port before it can proceed to the forwarding state.

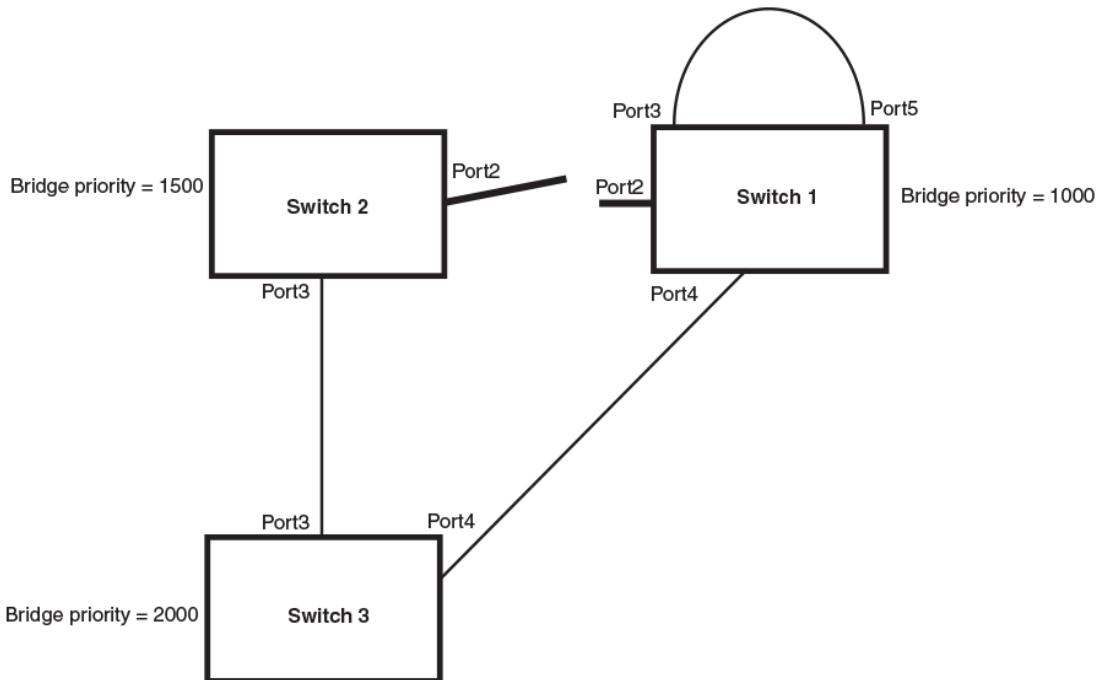
Once convergence is achieved, the active Layer 2 forwarding path converges as shown in the following figure.

FIGURE 52 Active Layer 2 path



Convergence after a link failure

The following figure illustrates a link failure in the 802.1W topology. In this example, Port2/Switch, which is the port that connects Switch 2 to the root bridge (Switch 1), failed and both Switch 2 and Switch 1 are affected by the topology change.

FIGURE 53 Link failure in the topology

Switch 1 sets its Port2 into a discarding state.

At the same time, Switch 2 assumes the role of a root bridge since its root port failed and it has no operational Alternate port. Port3/Switch 2, which currently has a Designated port role, sends an RST BPDU to Switch 3. The RST BPDU contains a proposal flag and a bridge ID of Switch 2 as its root bridge ID.

When Port3/Switch 3 receives the RST BPDUs, 802.1W algorithm determines that they are inferior to those that the port can transmit. Therefore, Port3/Switch 3 is given a new role, that of a Designated port. Port3/Switch 3 then sends an RST BPDU with a proposal flag to Switch 2, along with the new role information. However, the root bridge ID transmitted in the RST BPDU is still Switch 1.

When Port3/Switch 2 receives the RST BPDU, 802.1W algorithm determines that it is superior to the RST BPDU that it can transmit; therefore, Port3/Switch 2 receives a new role; that of a Root port. Port3/Switch 2 then sends an RST BPDU with an agreed flag to Port3/Switch 3. Port3/Switch 2 goes into a forwarding state.

When Port3/Switch 3 receives the RST BPDU that Port3/Switch 2 sent, Port3/Switch 3 changes into a forwarding state, which then completes the full convergence of the topology.

Convergence at link restoration

When Port2/Switch 2 is restored, both Switch 2 and Switch 1 recognize the change. Port2/Switch 1 starts assuming the role of a Designated port and sends an RST BPDU containing a proposal flag to Port2/Switch 2.

When Port2/Switch 2 receives the RST BPDUs, 802.1W algorithm determines that the RST BPDU received are better than those received on Port3/Switch 3; therefore, Port2/Switch 2 is given the role of a Root port. All the ports on Switch 2 are informed that a new Root port has been assigned which then signals all the ports to synchronize their roles and states. Port3/Switch 2, which was the previous Root port, enters a discarding state and negotiates with other ports on the bridge to establish its new role and state, until it finally assumes the role of a Designated port.

Next, the following happens:

- Port3/Switch 2, the Designated port, sends an RST BPDU, with a proposal flag to Port3/Switch 3.
- Port2/Switch 2 also sends an RST BPDU with an agreed flag to Port2/Switch 1 and then places itself into a forwarding state.

When Port2/Switch 1 receives the RST BPDU with an agreed flag sent by Port2/Switch 2, it puts that port into a forwarding state. The topology is now fully converged.

When Port3/Switch 3 receives the RST BPDU that Port3/Switch 2 sent, 802.1W algorithm determines that these RST BPDUs are superior to those that Port3/Switch 3 can transmit. Therefore, Port3/Switch 3 is given a new role, that of an Alternate port. Port3/Switch 3 immediately enters a discarding state.

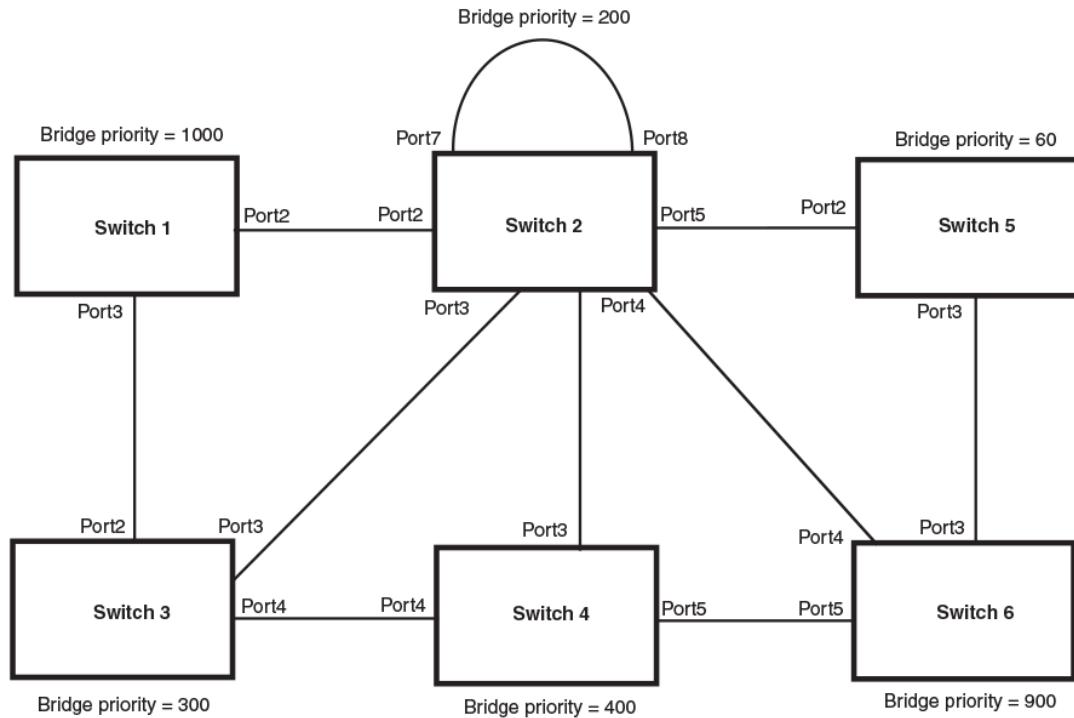
Now Port3/Switch 2 does not go into a forwarding state instantly like the Root port. It waits until the forward delay timer expires twice on that port while it is still in a Designated role, before it can proceed to the forwarding state. The wait, however, does not cause a denial of service, since the essential connectivity in the topology has already been established.

When fully restored, the topology is the same as that shown on [Figure 51](#) on page 279.

Convergence in a complex 802.1W topology

The following figure illustrates a complex 802.1W topology.

FIGURE 54 Complex 802.1W topology



In the above figure, Switch 5 is selected as the root bridge since it is the bridge with the highest priority. Lines in the figure show the point-to-point connection to the bridges in the topology.

Switch 5 sends an RST BPDU that contains a proposal flag to Port5/Switch 2. When handshakes are completed in Switch 5, Port5/Switch 2 is selected as the Root port on Switch 2. All other ports on Switch 2 are given Designated port role with discarding states.

Port5/Switch 2 then sends an RST BPDU with an agreed flag to Switch 5 to confirm that it is the new Root port and the port enters a forwarding state. Port7 and Port8 are informed of the identity of the new Root port. 802.1W algorithm selects Port7 as the Designated port while Port8 becomes the Backup port.

Port3/Switch 5 sends an RST BPDU to Port3/Switch 6 with a proposal flag. When Port3/Switch 5 receives the RST BPDU, handshake mechanisms select Port3 as the Root port of Switch 6. All other ports are given a Designated port role with discarding states. Port3/Switch 6 then sends an RST BPDU with an agreed flag to Port3/Switch 5 to confirm that it is the Root port. The Root port then goes into a forwarding state.

Now, Port4/Switch 6 receives RST BPDUs that are superior to what it can transmit; therefore, it is given the Alternate port role. The port remains in discarding state.

Port5/Switch 6 receives RST BPDUs that are inferior to what it can transmit. The port is then given a Designated port role.

Next Switch 2 sends RST BPDUs with a proposal flag to Port3/Switch 4. Port3 becomes the Root port for the bridge; all other ports are given a Designated port role with discarding states. Port3/Switch 4 sends an RST BPDU with an agreed flag to Switch 2 to confirm that it is the new Root port. The port then goes into a forwarding state.

Now Port4/Switch 4 receives an RST BPDU that is superior to what it can transmit. The port is then given an Alternate port role, and remains in discarding state.

Likewise, Port5/Switch 4 receives an RST BPDU that is superior to what it can transmit. The port is also given an Alternate port role, and remains in discarding state.

Port2/Switch 2 transmits an RST BPDU with a proposal flag to Port2/Switch 1. Port2/Switch 1 becomes the Root port. All other ports on Switch 1 are given Designated port roles with discarding states.

Port2/Switch 1 sends an RST BPDU with an agreed flag to Port2/Switch 2 and Port2/Switch 1 goes into a forwarding state.

Port3/Switch 1 receives an RST BPDUs that is inferior to what it can transmit; therefore, the port retains its Designated port role and goes into forwarding state only after the forward delay timer expires twice on that port while it is still in a Designated role.

Port3/Switch 2 sends an RST BPDU to Port3/Switch 3 that contains a proposal flag. Port3/Switch 3 becomes the Root port, while all other ports on Switch 3 are given Designated port roles and go into discarding states. Port3/Switch 3 sends an RST BPDU with an agreed flag to Port3/Switch 2 and Port3/Switch 3 goes into a forwarding state.

Now, Port2/Switch 3 receives an RST BPDUs that is superior to what it can transmit so that port is given an Alternate port state.

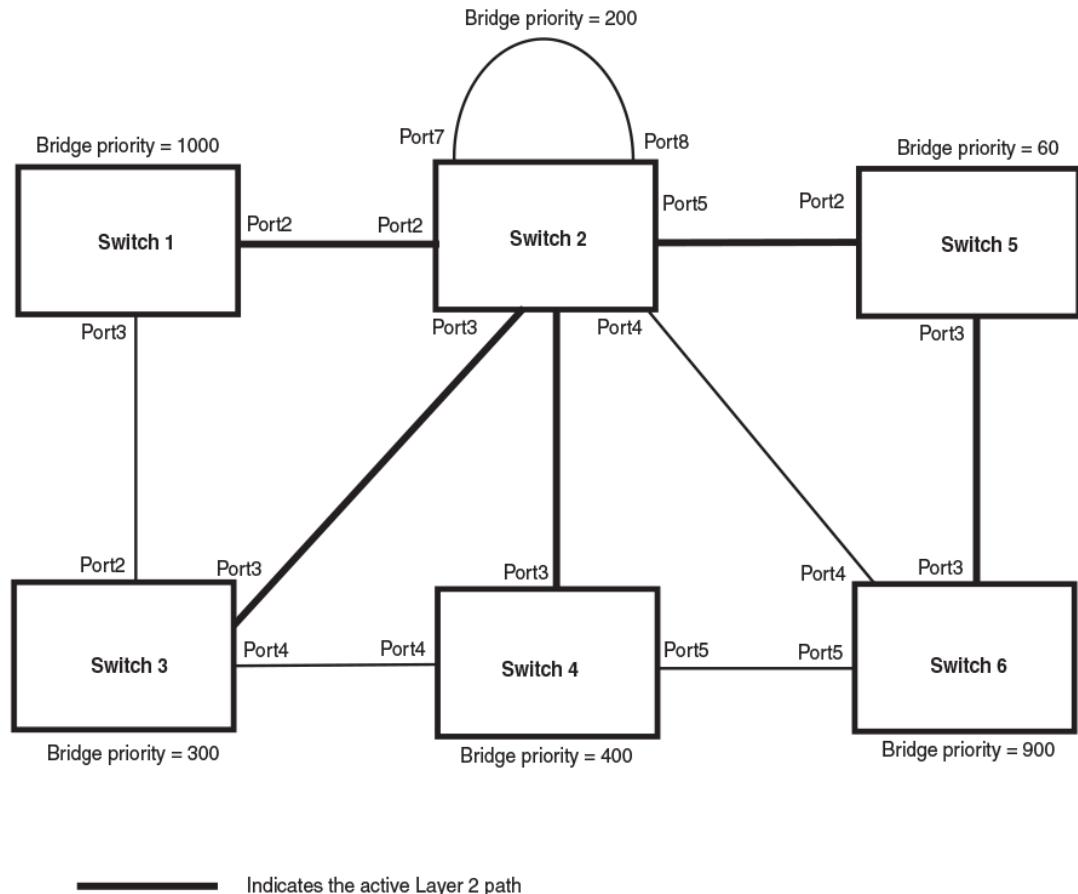
Port4/Switch 3 receives an RST BPDU that is inferior to what it can transmit; therefore, the port retains its Designated port role.

Ports on all the bridges in the topology with Designated port roles that received RST BPDUs with agreed flags go into forwarding states instantly. However, Designated ports that did not receive RST BPDUs with agreed flags must wait until the forward delay timer expires twice on those port. Only then will these port move into forwarding states.

The entire 802.1W topology converges in less than 300 msec and the essential connectivity is established between the designated ports and their connected root ports.

After convergence is complete, the following figure shows the active Layer 2 path of the topology in [Figure 54](#).

FIGURE 55 Active Layer 2 path in complex topology



Propagation of topology change

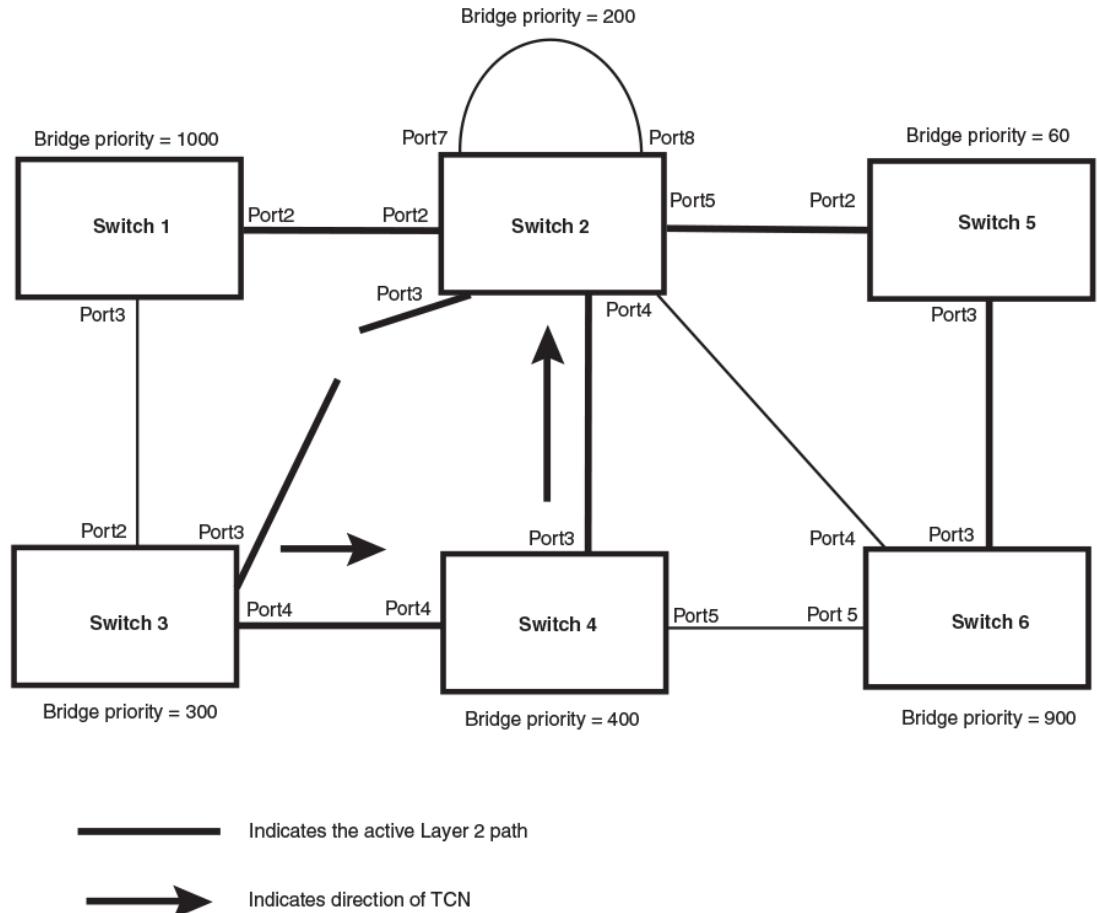
The Topology Change state machine generates and propagates the topology change notification messages on each port. When a Root port or a Designated port goes into a forwarding state, the Topology Change state machine on those ports send a topology change notice (TCN) to all the bridges in the topology to propagate the topology change.

NOTE

Edge ports, Alternate ports, or Backup ports do not need to propagate a topology change.

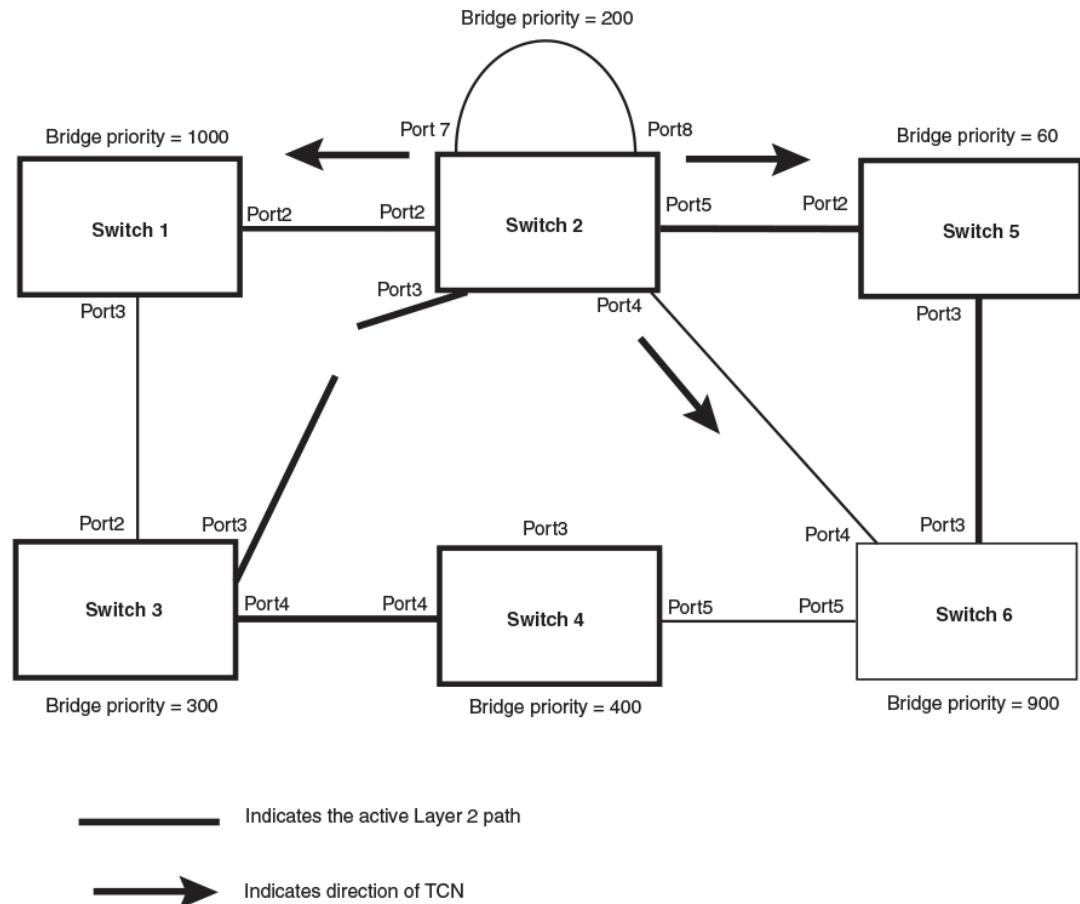
The TCN is sent in the RST BPDU that a port sends. Ports on other bridges in the topology then acknowledge the topology change once they receive the RST BPDU, and send the TCN to other bridges until all the bridges are informed of the topology change.

For example, Port3/Switch 2 in the following figure, fails. Port4/Switch 3 becomes the new Root port. Port4/Switch 3 sends an RST BPDU with a TCN to Port4/Switch 4. To propagate the topology change, Port4/Switch 4 then starts a TCN timer on itself, on the bridge Root port, and on other ports on that bridge with a Designated role. Then Port3/Switch 4 sends RST BPDU with the TCN to Port4/Switch 2. (Note the new active Layer 2 path in the following figure.)

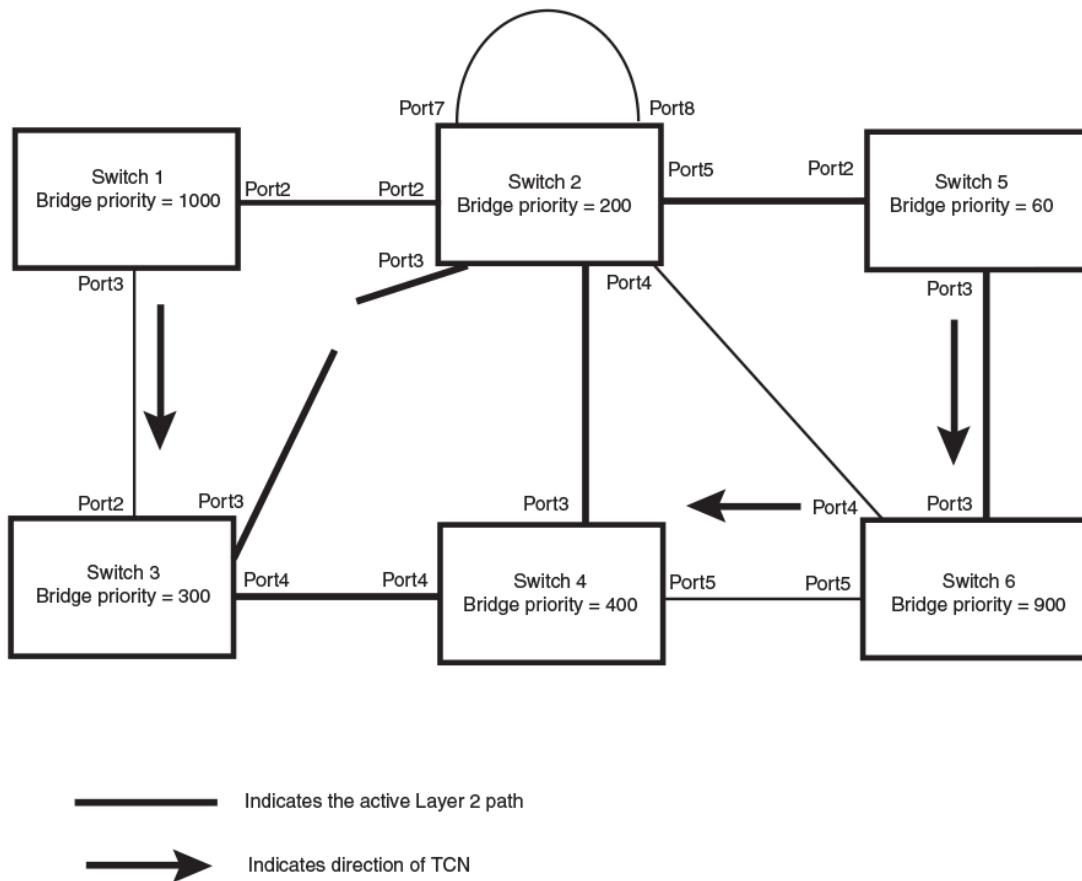
FIGURE 56 Beginning of topology change notice

Switch 2 then starts the TCN timer on the Designated ports and sends RST BPDUs that contain the TCN as follows ([Figure 57](#)):

- Port5/Switch 2 sends the TCN to Port2/Switch 5
- Port4/Switch 2 sends the TCN to Port4/Switch 6
- Port2/Switch 2 sends the TCN to Port2/Switch 1

FIGURE 57 Sending TCN to bridges connected to Switch 2

Then Switch 1, Switch 5, and Switch 6 send RST BPDUs that contain the TCN to Switch 3 and Switch 4 to complete the TCN propagation as shown in the following figure.

FIGURE 58 Completing the TCN propagation

Compatibility of 802.1W with 802.1D

802.1W-enabled bridges are backward compatible with IEEE 802.1D bridges. This compatibility is managed on a per-port basis by the Port Migration state machine. **However, intermixing the two types of bridges in the network topology is not advisable if you want to take advantage of the rapid convergence feature.**

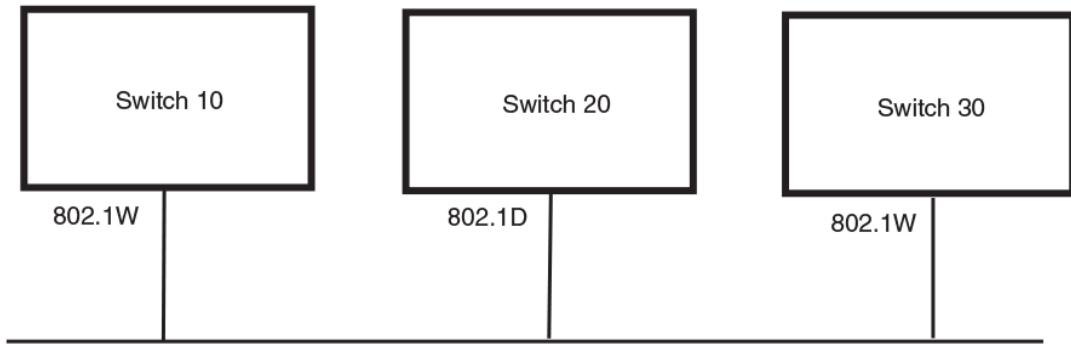
Compatibility with 802.1D means that an 802.1W-enabled port can send BPDUs in the STP or 802.1D format when one of the following events occur:

- The port receives a legacy BPDU. A legacy BPDU is an STP BPDU or a BPDU in an 802.1D format. The port that receives the legacy BPDU automatically configures itself to behave like a legacy port. It sends and receives legacy BPDUs only.
- The entire bridge is configured to operate in an 802.1D mode when an administrator sets the bridge parameter to zero at the CLI, forcing all ports on the bridge to send legacy BPDUs only.

Once a port operates in the 802.1D mode, 802.1D convergence times are used and rapid convergence is not realized.

For example, in the following figure, Switch 10 and Switch 30 receive legacy BPDUs from Switch 20. Ports on Switch 10 and Switch 30 begin sending BPDUs in STP format to allow them to operate transparently with Switch 20.

FIGURE 59 802.1W bridges with an 802.1D bridge



Once Switch 20 is removed from the LAN, Switch 10 and Switch 30 receive and transmit BPDUs in the STP format to and from each other. This state will continue until the administrator enables the **force-migration-check** command to force the bridge to send RSTP BPDU during a migrate time period. If ports on the bridges continue to hear only STP BPDUs after this migrate time period, those ports will return to sending STP BPDUs. However, when the ports receive RST BPDUs during the migrate time period, the ports begin sending RST BPDUs. The migrate time period is non-configurable. It has a value of three seconds.

NOTE

The IEEE standards state that 802.1W bridges need to interoperate with 802.1D bridges. IEEE standards set the path cost of 802.1W bridges to be between 1 and 200,000,000; whereas path cost of 802.1D bridges are set between 1 and 65,535. In order for the two bridge types to be able to interoperate in the same topology, the administrator needs to configure the bridge path cost appropriately. Path costs for either 802.1W bridges or 802.1D bridges need to be changed; in most cases, path costs for 802.1W bridges need to be changed.

Configuring 802.1W parameters on a Brocade device

The remaining 802.1W sections explain how to configure the 802.1W protocol in a Brocade device.

NOTE

With RSTP running, enabling static trunk on ports that are members of VLAN 4000 will keep the system busy for 20 to 25 seconds.

Brocade devices are shipped from the factory with 802.1W disabled. Use the following methods to enable or disable 802.1W. You can enable or disable 802.1W at the following levels:

- Port-based VLAN - Affects all ports within the specified port-based VLAN. When you enable or disable 802.1W within a port-based VLAN, the setting overrides the global setting. Thus, you can enable 802.1W for the ports within a port-based VLAN even when 802.1W is globally disabled, or disable the ports within a port-based VLAN when 802.1W is globally enabled.
- Individual port - Affects only the individual port. However, if you change the 802.1W state of the primary port in a trunk group, the change affects all ports in the trunk group.

Enabling or disabling 802.1W in a port-based VLAN

Use the following procedure to disable or enable 802.1W on a device on which you have configured a port-based VLAN. Changing the 802.1W state in a VLAN affects only that VLAN.

To enable 802.1W for all ports in a port-based VLAN, enter commands such as the following.

```
device(config)#vlan 10
device(config-vlan-10)#spanning-tree 802-1w
```

Syntax: [no] spanning-tree 802-1w

Note regarding pasting 802.1W settings into the running configuration

If you paste 802.1W settings into the running configuration, and the pasted configuration includes ports that are already up, the ports will initially operate in STP legacy mode before operating in 802.1W RSTP mode. For example, the following pasted configuration will cause ports e 1 and e 2 to temporarily operate in STP legacy mode, because these ports are already up and running.

```
conf t
vlan 120
tag e 1 to e 2
spanning-tree 802-1w
spanning-tree 802-1w priority 1001
end
```

To avoid this issue, 802.1W commands/settings that are pasted into the configuration should be in the following order.

1. Ports that are not yet connected
2. 802.1W RSTP settings
3. Ports that are already up

Example

```
conf t
vlan 120
untag e 3
spanning-tree 802-1w
spanning-tree 802-1w priority 1001
tag e 1 to 2
end
```

In the above configuration, untagged port e3 is added to VLAN 120 *before* the 802.1W RSTP settings, and ports e1 and e2 are added *after* the 802.1W RSTP settings. When these commands are pasted into the running configuration, the ports will properly operate in 802.1W RSTP mode.

Enabling or disabling 802.1W on a single spanning tree

To enable 802.1W for all ports of a single spanning tree, enter a command such as the following.

```
device(config-vlan-10)#spanning-tree single 802-1w
```

Syntax: [no] spanning-tree single 802-1w

Disabling or enabling 802.1W on an individual port

The **spanning-tree 802-1w** or **spanning-tree single 802-1w** command must be used to initially enable 802.1W on ports. Both commands enable 802.1W on all ports that belong to the VLAN or to the single spanning tree.

Once 802.1W is enabled on a port, it can be disabled on individual ports. 802.1W that have been disabled on individual ports can then be enabled as required.

NOTE

If you change the 802.1W state of the primary port in a trunk group, the change affects all ports in that trunk group.

To disable or enable 802.1W on an individual port, enter commands such as the following.

```
device(config)#interface e 1
device(config-if-e1000-1)#no spanning-tree
```

Syntax: [no] spanning-tree**Changing 802.1W bridge parameters**

When you make changes to 802.1W bridge parameters, the changes are applied to individual ports on the bridge. To change 802.1W bridge parameters, use the following methods.

To designate a priority for a bridge, enter a command such as the following.

```
device(config)#spanning-tree 802-1w priority 10
```

The command in this example changes the priority on a device on which you have not configured port-based VLANs. The change applies to the default VLAN. If you have configured a port-based VLAN on the device, you can configure the parameters only at the configuration level for individual VLANs. Enter commands such as the following.

```
device(config)#vlan 20
device(config-vlan-20)#spanning-tree 802-1w priority 0
```

To make this change in the default VLAN, enter the following commands.

```
device(config)#vlan 1
device(config-vlan-1)#spanning-tree 802-1w priority 0
```

Syntax: spanning-tree 802-1w [forward-delay value] | [hello-time value] | [max-age time] | [force-version value] | [priority value]

The **forward-delay** value parameter specifies how long a port waits before it forwards an RST BPDU after a topology change. This can be a value from 4 - 30 seconds. The default is 15 seconds.

The **hello-time** value parameter specifies the interval between two hello packets. This parameter can have a value from 1 - 10 seconds. The default is 2 seconds.

The **max-age** value parameter specifies the amount of time the device waits to receive a hello packet before it initiates a topology change. You can specify a value from 6 - 40 seconds. The default is 20 seconds.

The value of **max-age** must be greater than the value of **forward-delay** to ensure that the downstream bridges do not age out faster than the upstream bridges (those bridges that are closer to the root bridge).

The **force-version** value parameter forces the bridge to send BPDUs in a specific format. You can specify one of the following values:

- 0 - The STP compatibility mode. Only STP (or legacy) BPDUs will be sent.
- 2 - The default. RST BPDUs will be sent unless a legacy bridge is detected. If a legacy bridge is detected, STP BPDUs will be sent instead.

The default is 2.

The **priority** value parameter specifies the priority of the bridge. You can enter a value from 0 - 65535. A lower numerical value means the bridge has a higher priority. Thus, the highest priority is 0. The default is 32768.

You can specify some or all of these parameters on the same command line. If you specify more than one parameter, you must specify them in the order shown above, from left to right.

Changing port parameters

The 802.1W port commands can be enabled on individual ports or on multiple ports, such as all ports that belong to a VLAN.

The 802.1W port parameters are preconfigured with default values. If the default parameters meet your network requirements, no other action is required.

You can change the following 802.1W port parameters using the following method.

```
device(config)#vlan 10
device(config-vlan-10)#spanning-tree 802-1w ethernet 5 path-cost 15 priority 64
```

Syntax: **spanning-tree 802-1w ethernet port path-cost value | priority value | [admin-edge-port] | [admin-pt2pt-mac] | [force-migration-check]**

The **path-cost value** parameter specifies the cost of the port path to the root bridge. 802.1W prefers the path with the lowest cost. You can specify a value from 1 - 20,000,000. The following table shows the recommended path cost values from the IEEE standards.

TABLE 48 Recommended path cost values of 802.1W

Link speed	Recommended (Default) 802.1W path cost values	Recommended 802.1W patch cost range
Less than 100 kilobits per second	200,000,000	20,000,000 – 200,000,000
1 Megabit per second	20,000,000	2,000,000 – 200,000,000
10 Megabits per second	2,000,000	200,000 – 200,000,000
100 Megabits per second	200,000	20,000 – 200,000,000
1 Gbps per second	20,000	2,000 – 200,000,000
10 Gbps per second	2,000	200 – 20,000
100 Gbps per second	200	20 – 2,000
1 Terabits per second	20	2 – 200
10 Terabits per second	2	1 – 20

- You can specify a value from 0 – 240, in increments of 16. If you enter a value that is not divisible by 16, the software returns an error message. The default value is 128. A higher numerical value means a lower priority; thus, the highest priority is 0.
- Set the **admin-edge-port** to enabled or disabled. If set to enabled, then the port becomes an edge port in the domain.

Set the **admin-pt2pt-mac** to enabled or disabled. If set to enabled, then a port is connected to another port through a point-to-point link. The point-to-point link increases the speed of convergence. This parameter, however, does not auto-detect whether or not the link is a physical point-to-point link.

The **force-migration-check** parameter forces the specified port to sent one RST BPDU. If only STP BPDUs are received in response to the sent RST BPDU, then the port will go return to sending STP BPDUs.

Example

Suppose you want to enable 802.1W on a system with no active port-based VLANs and change the hello-time from the default value of 2 to 8 seconds. Additionally, suppose you want to change the path and priority costs for port 5 only. To do so, enter the following commands.

```
device(config)#spanning-tree 802-1w hello-time 8
device(config)#spanning-tree 802-1w ethernet 5 path-cost 15 priority 64
```

Displaying information about 802-1W

To display a summary of 802-1W, use the following command.

```
device#show
 802-1w
--- VLAN 1 [ STP Instance owned by VLAN 1 ] -----
VLAN 1 BPDU cam index is 2 and the IGC and DMA master Are(HEX) 0 1 2 3
Bridge IEEE 802.1W Parameters:
Bridge          Bridge   Bridge   Force   tx
Identifier      MaxAge   Hello    FwdDly  Version Hold
hex             sec      sec     sec      cnt
800000e080541700 20       2       15      Default 3
RootBridge      RootPath DesignatedBri- Root   Max   Fwd   Hel
Identifier      Cost     dge     Identifier Port   Age   Dly   lo
hex             hex      hex      hex      sec   sec   sec
800000e0804c9c00 200000  800000e0804c9c00 1     20   15   2
Port IEEE 802.1W Parameters:
    <-- Config Params -->|<----- Current state ----->
Port  Pri PortPath P2P Edge Role      State      Designa- Designated
Num   Cost      Mac Port          ted cost bridge
1     128 200000 F   F   ROOT        FORWARDING 0      800000e0804c9c00
2     128 200000 F   F   DESIGNATED FORWARDING 200000 800000e080541700
3     128 200000 F   F   DESIGNATED FORWARDING 200000 800000e080541700
4     128 200000 F   F   BACKUP      DISCARDING 200000 800000e080541700
```

Syntax: show 802-1w [vlan *vlan-id*]

The **vlan *vlan-id*** parameter displays 802.1W information for the specified port-based VLAN.

The **show 802-1w** command shows the information listed in the following figure.

TABLE 49 CLI display of 802.1W summary

Field	Description
VLAN ID	The port-based VLAN that owns the STP instance. VLAN 1 is the default VLAN. If you have not configured port-based VLANs on this device, all 802.1W information is for VLAN 1.
Bridge IEEE 802.1W parameters	
Bridge Identifier	The ID of the bridge.
Bridge Max Age	The configured max age for this bridge. The default is 20.
Bridge Hello	The configured hello time for this bridge. The default is 2.
Bridge FwdDly	The configured forward delay time for this bridge. The default is 15.
Force-Version	The configured force version value. One of the following value is displayed: <ul style="list-style-type: none"> • 0 - The bridge has been forced to operate in an STP compatibility mode. • 2 - The bridge has been forced to operate in an 802.1W mode. (This is the default.)

TABLE 49 CLI display of 802.1W summary (Continued)

Field	Description
txHoldCnt	The number of BPDUs that can be transmitted per Hello Interval. The default is 3.
Root Bridge Identifier	ID of the Root bridge that is associated with this bridge
Root Path Cost	The cost to reach the root bridge from this bridge. If the bridge is the root bridge, then this parameter shows a value of zero.
Designated Bridge Identifier	The bridge from where the root information was received. It can be from the root bridge itself, but it could also be from another bridge.
Root Port	The port on which the root information was received. This is the port that is connected to the Designated Bridge.
Max Age	<p>The max age is derived from the Root port. An 802.1W-enabled bridge uses this value, along with the hello and message age parameters to compute the effective age of an RST BPDU.</p> <p>The message age parameter is generated by the Designated port and transmitted in the RST BPDU. RST BPDUs transmitted by a Designated port of the root bridge contains a message value of zero.</p> <p>Effective age is the amount of time the Root port, Alternate port, or Backup port retains the information it received from its peer Designated port. Effective age is reset every time a port receives an RST BPDU from its Designated port. If a Root port does not receive an RST BPDU from its peer Designated port for a duration more than the effective age, the Root port ages out the existing information and recomputes the topology.</p> <p>If the port is operating in 802.1D compatible mode, then max age functionality is the same as in 802.1D (STP).</p>
Fwd Dly	<p>The number of seconds a non-edge Designated port waits until it can apply any of the following transitions, if the RST BPDU it receives does not have an agreed flag:</p> <ul style="list-style-type: none"> • Discarding state to learning state • Learning state to forwarding state <p>When a non-edge port receives the RST BPDU it goes into forwarding state within 4 seconds or after two hello timers expire on the port.</p> <p>Fwd Dly is also the number of seconds that a Root port waits for an RST BPDU with a proposal flag before it applies the state transitions listed above.</p> <p>If the port is operating in 802.1D compatible mode, then forward delay functionality is the same as in 802.1D (STP).</p>
Hello	The hello value derived from the Root port. It is the number of seconds between two Hello packets.
Port IEEE 802.1W parameters	
Port Num	The port number shown in a <i>slot#/port#</i> format.

TABLE 49 CLI display of 802.1W summary (Continued)

Field	Description
Pri	The configured priority of the port. The default is 128 or 0x80.
Port Path Cost	The configured path cost on a link connected to this port.
P2P Mac	Indicates if the point-to-point-mac parameter is configured to be a point-to-point link: <ul style="list-style-type: none"> • T - The link is configured as a point-to-point link. • F - The link is not configured as a point-to-point link. This is the default.
Edge port	Indicates if the port is configured as an operational Edge port: <ul style="list-style-type: none"> • T - The port is configured as an Edge port. • F - The port is not configured as an Edge port. This is the default.
Role	The current role of the port: <ul style="list-style-type: none"> • Root • Designated • Alternate • Backup • Disabled <p>Refer to Bridges and bridge port roles on page 260 for definitions of the roles.</p>
State	The port current 802.1W state. A port can have one of the following states: <ul style="list-style-type: none"> • Forwarding • Discarding • Learning • Disabled <p>Refer to Bridge port states on page 265 and Edge port and non-edge port states on page 266.</p>
Designated Cost	The best root path cost that this port received, including the best root path cost that it can transmit.
Designated Bridge	The ID of the bridge that sent the best RST BPDU that was received on this port.

To display detailed information about 802.1W, enter the 802-1w command.

```
device#show 802-1w detail
=====
VLAN 1 - MULTIPLE SPANNING TREE (MSTP - IEEE 802.1W) ACTIVE
=====
BridgeId 800000e080541700, forceVersion 2, txHoldCount 3
Port 1 - Role: ROOT - State: FORWARDING
  PathCost 200000, Priority 128, AdminOperEdge F, AdminPt2PtMac F
  DesignatedPriority - Root: 0x800000e0804c9c00, Bridge: 0x800000e080541700
  ActiveTimers - rrWhile 4 rcvdInfoWhile 4
  MachineStates - PIM: CURRENT, PRT: ROOT_PORT, PST: FORWARDING
  TCM: ACTIVE, PPM: SENDING_STP, PTX: TRANSMIT_IDLE
  Received - RST_BPDUs 0, Config_BPDUs 1017, TCN_BPDUs 0
Port 2 - Role: DESIGNATED - State: FORWARDING
  PathCost 200000, Priority 128, AdminOperEdge F, AdminPt2PtMac F
  DesignatedPriority - Root: 0x800000e0804c9c00, Bridge: 0x800000e080541700
  ActiveTimers - helloWhen 0
```

```

MachineStates - PIM: CURRENT, PRT: DESIGNATED_PORT, PST: FORWARDING
TCM: ACTIVE, PPM: SENDING_RSTP, PTX: TRANSMIT_IDLE
Received - RST BPDU 0, Config BPDU 0, TCN BPDU 0

```

Syntax: show 802-1w detail [*vlan/vlan-id*]

The **vlan *vlan-id*** parameter displays 802.1W information for the specified port-based VLAN.

The **show spanning-tree 802.1W** command shows the following information.

TABLE 50 CLI display of show spanning-tree 802.1W

Field	Description
VLAN ID	ID of the VLAN that owns the instance of 802.1W and whether or not it is active.
Bridge ID	ID of the bridge.
forceVersion	the configured version of the bridge: <ul style="list-style-type: none"> • 0 - The bridge has been forced to operate in an STP compatible mode. • 2 - The bridge has been forced to operate in an 802.1W mode.
txHoldCount	The number of BPDU that can be transmitted per Hello Interval. The default is 3.
Port	ID of the port in slot#/port#format.
Role	The current role of the port: <ul style="list-style-type: none"> • Root • Designated • Alternate • Backup • Disabled Refer to Bridges and bridge port roles on page 260 for definitions of the roles.
State	The port current 802.1W state. A port can have one of the following states: <ul style="list-style-type: none"> • Forwarding • Discarding • Learning • Disabled Refer to Bridge port states on page 265 and Edge port and non-edge port states on page 266.
Path Cost	The configured path cost on a link connected to this port.
Priority	The configured priority of the port. The default is 128 or 0x80.
AdminOperEdge	Indicates if the port is an operational Edge port. Edge ports may either be auto-detected or configured (forced) to be Edge ports using the CLI: <ul style="list-style-type: none"> • T - The port is an Edge port. • F - The port is not an Edge port. This is the default.

TABLE 50 CLI display of show spanning-tree 802.1W (Continued)

Field	Description
AdminP2PMac	Indicates if the point-to-point-mac parameter is configured to be a point-to-point link: <ul style="list-style-type: none"> • T - The link is a point-to-point link • F - The link is not a point-to-point link. This is the default.
DesignatedPriority	Shows the following: <ul style="list-style-type: none"> • Root - Shows the ID of the root bridge for this bridge. • Bridge - Shows the ID of the Designated bridge that is associated with this port.
ActiveTimers	Shows what timers are currently active on this port and the number of seconds they have before they expire: <ul style="list-style-type: none"> • rrWhile - Recent root timer. A non-zero value means that the port has recently been a Root port. • rcvdlInfoWhile - Received information timer. Shows the time remaining before the information held by this port expires (ages out). This timer is initialized with the effective age parameter. (Refer to the Max Age field in the Table 49.) • rbWhile - Recent backup timer. A non-zero value means that the port has recently been a Backup port. • helloWhen - Hello period timer. The value shown is the amount of time between hello messages. • tcWhile - Topology change timer. The value shown is the interval when topology change notices can be propagated on this port. • fdWhile - Forward delay timer. • mdelayWhile - Migration delay timer. The amount of time that a bridge on the same LAN has to synchronize its migration state with this port before another BPDU type can cause this port to change the BPDU that it transmits.
Machine States	The current states of the various state machines on the port: <ul style="list-style-type: none"> • PIM - State of the Port Information state machine. • PRT - State of the Port Role Transition state machine. • PST - State of the Port State Transition state machine. • TCM - State of the Topology Change state machine. • PPM - State of the Port Protocol Migration. • PTX - State of the Port Transmit state machine. <p>Refer to the "Port Role Selection state machines" section in Changes to port roles and states on page 266 for details on state machines.</p>
Received	Shows the number of BPDU types the port has received: <ul style="list-style-type: none"> • RST BPDU - BPDU in 802.1W format. • Config BPDU - Legacy configuration BPDU (802.1D format). • TCN BPDU - Legacy topology change BPDU (802.1D format).

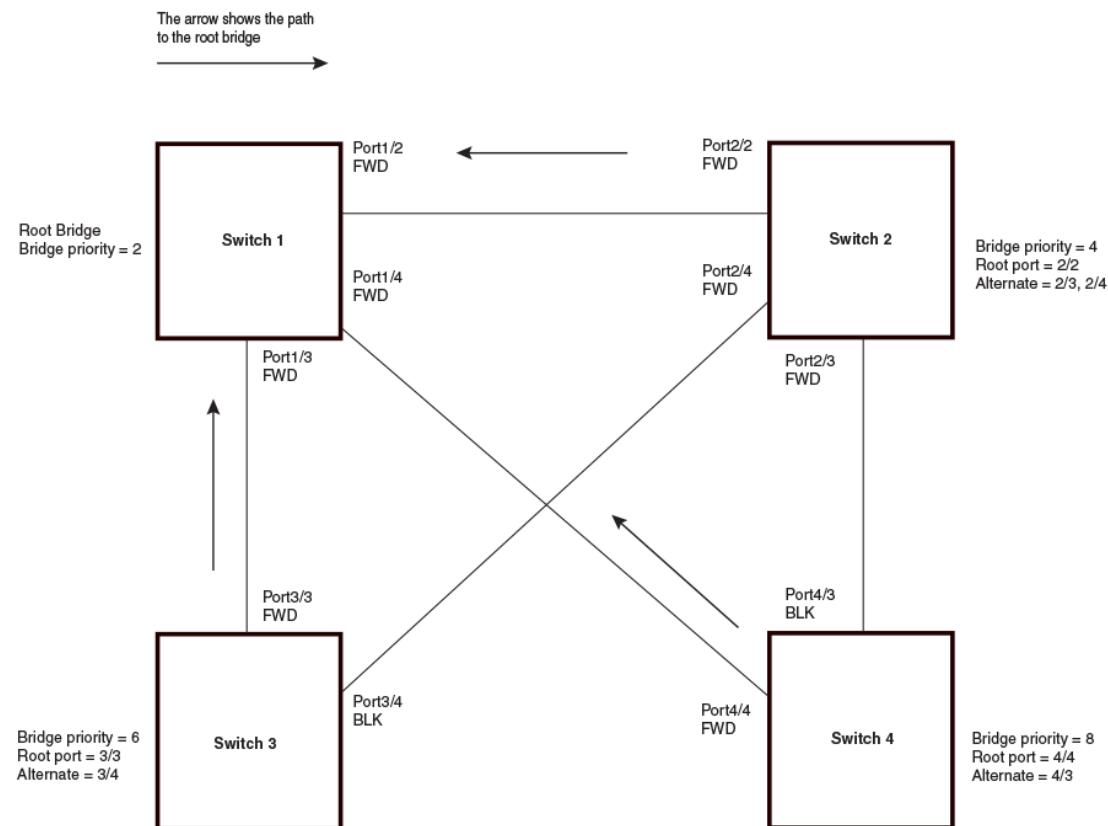
802.1W Draft 3

As an alternative to full 802.1W, you can configure 802.1W Draft 3. 802.1W Draft 3 provides a subset of the RSTP capabilities described in the 802.1W STP specification.

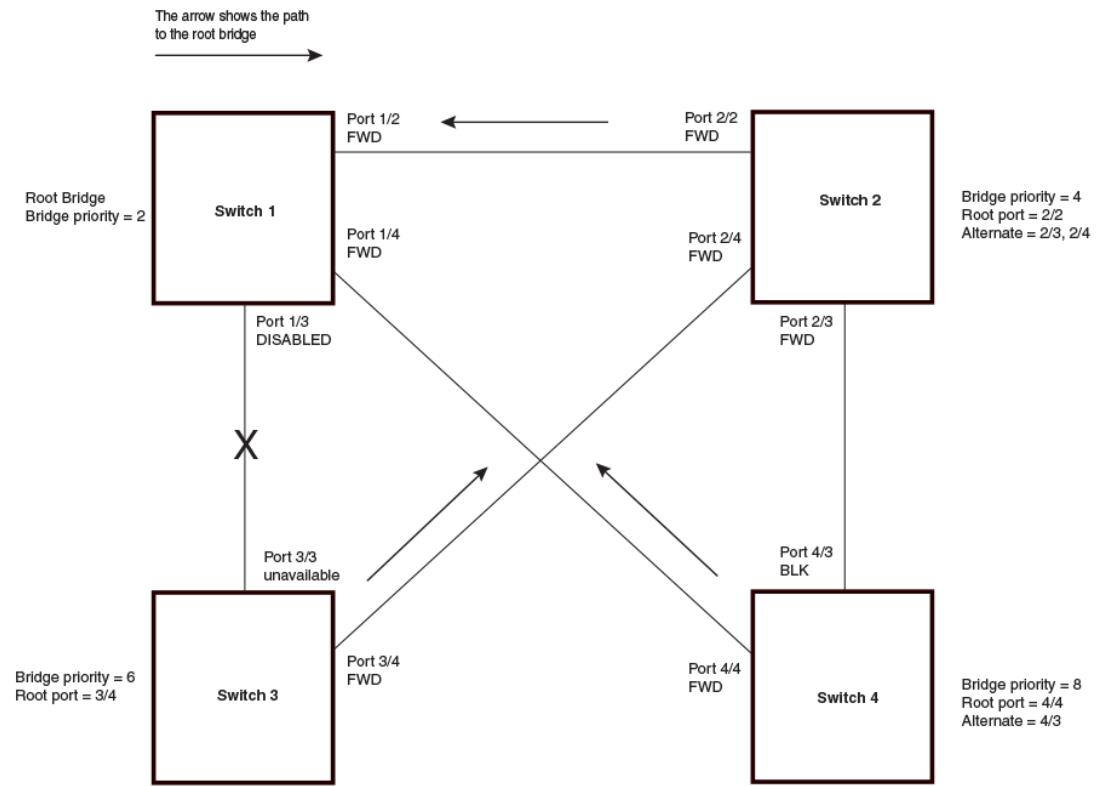
802.1W Draft 3 support is disabled by default. When the feature is enabled, if a root port on a Brocade device that is not the root bridge becomes unavailable, the device can automatically Switch over to an alternate root port, without reconvergence delays. 802.1W Draft 3 does not apply to the root bridge, since all the root bridge ports are always in the forwarding state.

The following figure shows an example of an optimal STP topology. In this topology, all the non-root bridges have at least two paths to the root bridge (Switch 1 in this example). One of the paths is through the root port. The other path is a backup and is through the alternate port. While the root port is in the forwarding state, the alternate port is in the blocking state.

FIGURE 60 802.1W Draft 3 RSTP ready for failover



If the root port on a Switch becomes unavailable, 802.1W Draft 3 immediately fails over to the alternate port, as shown in the following figure.

FIGURE 61 802.1W Draft 3 RSTP failover to alternate root port

In this example, port 3/3 on Switch 3 has become unavailable. In standard STP (802.1D), if the root port becomes unavailable, the Switch must go through the listening and learning stages on the alternate port to reconverge with the spanning tree. Thus, port 3/4 must go through the listening and learning states before entering the forwarding state and thus reconverging with the spanning tree.

802.1W Draft 3 avoids the reconvergence delay by calculating an alternate root port, and immediately failing over to the alternate port if the root port becomes unavailable. The alternate port is in the blocking state as long as the root port is in the forwarding state, but moves immediately to the active state if the root port becomes unavailable. Thus, using 802.1W Draft 3, Switch 3 immediately fails over to port 3/4, without the delays caused by the listening and learning states.

802.1W Draft 3 selects the port with the next-best cost to the root bridge. For example, on Switch 3, port 3/3 has the best cost to the root bridge and thus is selected by STP as the root port. Port 3/4 has the next-best cost to the root bridge, and thus is selected by 802.1W Draft 3 as the alternate path to the root bridge.

Once a failover occurs, the Switch no longer has an alternate root port. If the port that was an alternate port but became the root port fails, standard STP is used to reconverge with the network. You can minimize the reconvergence delay in this case by setting the forwarding delay on the root bridge to a lower value. For example, if the forwarding delay is set to 15 seconds (the default), change the forwarding delay to a value from 3 - 10 seconds.

During failover, 802.1W Draft 3 flushes the MAC addresses leaned on the unavailable root port, selects the alternate port as the new root port, and places that port in the forwarding state. If traffic is flowing in both directions on the new root port, addresses are flushed (moved) in the rest of the spanning tree automatically.

Spanning tree reconvergence time

Spanning tree reconvergence using 802.1W Draft 3 can occur within one second.

After the spanning tree reconverges following the topology change, traffic also must reconverge on all the bridges attached to the spanning tree. This is true regardless of whether 802.1W Draft 3 or standard STP is used to reconverge the spanning tree.

Traffic reconvergence happens after the spanning tree reconvergence, and is achieved by flushing the Layer 2 information on the bridges:

- Following 802.1W Draft 3 reconvergence of the spanning tree, traffic reconvergence occurs in the time it takes for the bridge to detect the link changes plus the STP maximum age set on the bridge.
- If standard STP reconvergence occurs instead, traffic reconvergence takes two times the forward delay plus the maximum age.

NOTE

802.1W Draft 3 does not apply when a failed root port comes back up. When this happens, standard STP is used.

802.1W configuration considerations

802.1W Draft 3 is disabled by default. To ensure optimal performance of the feature before you enable it, do the following:

- Configure the bridge priorities so that the root bridge is one that supports 802.1W Draft 3. (Use a Brocade device or third-party device that supports 802.1W Draft 3.)
- Change the forwarding delay on the root bridge to a value lower than the default 15 seconds. Brocade recommends a value from 3 - 10 seconds. The lower forwarding delay helps reduce reconvergence delays in cases where 802.1W Draft 3 is not applicable, such as when a failed root port comes back up.
- Configure the bridge priorities and root port costs so that each device has an active path to the root bridge if its root port becomes unavailable. For example, port 3/4 is connected to port 2/4 on Switch 2, which has the second most favorable bridge priority in the spanning tree.

NOTE

If reconvergence involves changing the state of a root port on a bridge that supports 802.1D STP but not 802.1W Draft 3, then reconvergence still requires the amount of time it takes for the ports on the 802.1D bridge to change state to forwarding (as needed), and receive BPDUs from the root bridge for the new topology.

Enabling 802.1W Draft 3

802.1W Draft 3 is disabled by default. The procedure for enabling the feature differs depending on whether single STP is enabled on the device.

NOTE

STP must be enabled before you can enable 802.1W Draft 3.

Enabling 802.1W Draft 3 when single STP is not enabled

By default, each port-based VLAN on the device has its own spanning tree. To enable 802.1W Draft 3 in a port-based VLAN, enter commands such as the following.

```
device(config)#vlan 10  
device(config-vlan-10)#spanning-tree rstp
```

Syntax: [no] spanning-tree rstp

This command enables 802.1W Draft 3. You must enter the command separately in each port-based VLAN in which you want to run 802.1W Draft 3.

NOTE

This command does not also enable STP. To enable STP, first enter the **spanning-tree** command without the **rstp** parameter. After you enable STP, enter the **spanning-tree rstp** command to enable 802.1W Draft 3.

To disable 802.1W Draft 3, enter the following command.

```
device(config-vlan-10)#no spanning-tree rstp
```

Enabling 802.1W Draft 3 when single STP is enabled

To enable 802.1W Draft 3 on a device that is running single STP, enter the following command at the global CONFIG level of the CLI.

```
device(config)#spanning-tree single rstp
```

Syntax: [no] spanning-tree single rstp

This command enables 802.1W Draft 3 on the whole device.

NOTE

This command does not also enable single STP. To enable single STP, first enter the **spanning-tree single** command without the **rstp** parameter. After you enable single STP, enter the **spanning-tree single rstp** command to enable 802.1W Draft 3.

To disable 802.1W Draft 3 on a device that is running single STP, enter the following command.

```
device(config)#no spanning-tree single rstp
```

Single Spanning Tree (SSTP)

By default, each port-based VLAN on a Brocade device runs a separate spanning tree, which you can enable or disable on an individual VLAN basis.

Alternatively, you can configure a Brocade device to run a single spanning tree across all ports and VLANs on the device. The Single STP feature (SSTP) is especially useful for connecting a Brocade device to third-party devices that run a single spanning tree in accordance with the 802.1Q specification.

SSTP uses the same parameters, with the same value ranges and defaults, as the default STP support on Brocade devices. Refer to [STP parameters and defaults](#) on page 242.

SSTP defaults

SSTP is disabled by default. When you enable the feature, all VLANs on which STP is enabled become members of a single spanning tree. All VLANs on which STP is disabled are excluded from the single spanning tree.

To add a VLAN to the single spanning tree, enable STP on that VLAN. To remove a VLAN from the single spanning tree, disable STP on that VLAN.

When you enable SSTP, all the ports that are in port-based VLANs with STP enabled become members of a single spanning tree domain. Thus, the ports share a single BPDU broadcast domain. The Brocade device places all the ports in a non-configurable VLAN, 4094, to implement the SSTP domain. However, this VLAN does not affect port membership in the port-based VLANs you have configured. Other broadcast traffic is still contained within the individual port-based VLANs. Therefore, you can use SSTP while still using your existing VLAN configurations without changing your network. In addition, SSTP does not affect 802.1Q tagging. Tagged and untagged ports alike can be members of the single spanning tree domain.

NOTE

When SSTP is enabled, the BPDUs on tagged ports go out untagged.

If you disable SSTP, all VLANs that were members of the single spanning tree run MSTP instead. In MSTP, each VLAN has its own spanning tree. VLANs that were not members of the single spanning tree were not enabled for STP. Therefore, STP remains disabled on those VLANs.

Enabling SSTP

To enable SSTP, use one of the following methods.

NOTE

If the device has only one port-based VLAN (the default VLAN), then the device is already running a single instance of STP. In this case, you do not need to enable SSTP. You need to enable SSTP only if the device contains more than one port-based VLAN and you want all the ports to be in the same STP broadcast domain.

To configure the Brocade device to run a single spanning tree, enter the following command at the global CONFIG level.

```
device(config)#spanning-tree single
```

NOTE

If the device has only one port-based VLAN, the CLI command for enabling SSTP is not listed in the CLI. The command is listed only if you have configured a port-based VLAN.

To change a global STP parameter, enter a command such as the following at the global CONFIG level.

```
device(config)# spanning-tree single priority 2
```

This command changes the STP priority for all ports to 2.

To change an STP parameter for a specific port, enter commands such as the following.

```
device(config) spanning-tree single ethernet 1 priority 10
```

The commands shown above override the global setting for STP priority and set the priority to 10 for port 1/1.

Here is the syntax for the global STP parameters.

Syntax: [no] spanning-tree single [forward-delay value] [hello-time value] | [maximum-age time] | [priority value]

Here is the syntax for the STP port parameters.

Syntax: [no] spanning-tree single [ethernet port path-cost value | priority value]

NOTE

Both commands listed above are entered at the global CONFIG level.

Displaying SSCP information

To verify that SSCP is in effect, enter the following commands at any level of the CLI.

```
device#show span
```

Syntax: show span [vlan vlan-id] | [pvst-mode] | [num] | [detail [vlan vlan-id [ethernet port] | num]]

The **vlan** *vlan-id* parameter displays STP information for the specified port-based VLAN.

The **pvst-mode** parameter displays STP information for the device Per VLAN Spanning Tree (PVST+) compatibility configuration. Refer to [PVST/PVST+ compatibility](#) on page 306.

The **num** parameter displays only the entries after the number you specify. For example, on a device with three port-based VLANs, if you enter 1, then information for the second and third VLANs is displayed, but information for the first VLAN is not displayed. Information is displayed according to VLAN number, in ascending order. The entry number is not the same as the VLAN number. For example, if you have port-based VLANs 1, 10, and 2024, then the command output has three STP entries. To display information for VLANs 10 and 2024 only, enter **show span 1**.

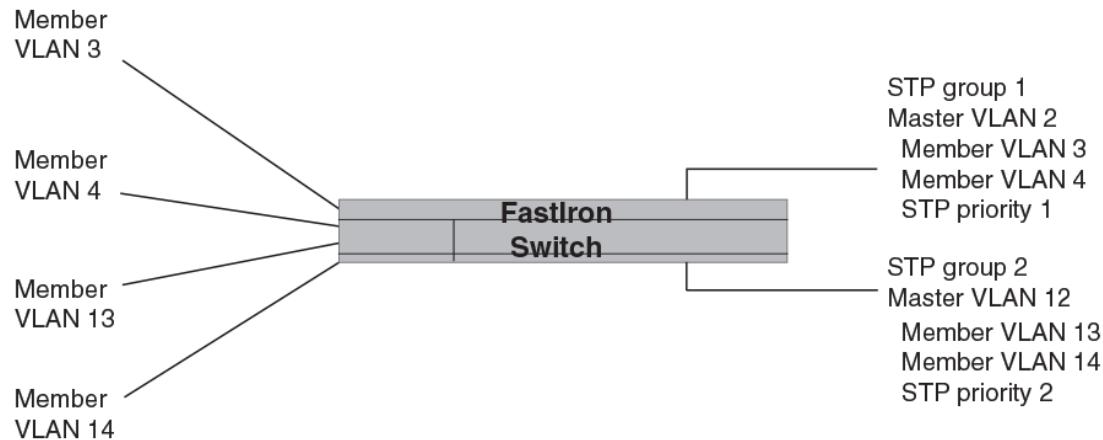
The **detail** parameter and its additional optional parameters display detailed information for individual ports. Refer to [Displaying detailed STP information for each interface](#) on page 251.

STP per VLAN group

STP per VLAN group is an STP enhancement that provides scalability while overcoming the limitations of the following scalability alternatives:

- Standard STP - You can configure up to 254 instances of standard STP on a Brocade device. It is possible to need more instances of STP than this in large configurations. Using STP per VLAN group, you can aggregate STP instances.
- Single STP - Single STP allows all the VLANs to run STP, but each VLAN runs the same instance of STP, resulting in numerous blocked ports that do not pass any Layer 2 traffic. STP per VLAN group uses all available links by load balancing traffic for different instances of STP on different ports. A port that blocks traffic for one spanning tree forwards traffic for another spanning tree.

STP per VLAN group allows you to group VLANs and apply the same STP parameter settings to all the VLANs in the group. The following figure shows an example of a STP per VLAN group implementation.

FIGURE 62 STP per VLAN group example

A master VLAN contains one or more member VLANs. Each of the member VLANs in the STP Group runs the same instance of STP and uses the STP parameters configured for the master VLAN. In this example, the FastIron switch is configured with VLANs 3, 4, 13, and 14. VLANs 3 and 4 are grouped in master VLAN 2, which is in STP group 1. VLANs 13 and 14 are grouped in master VLAN 12, which is in STP group 2. The VLANs in STP group 1 all share the same spanning tree. The VLANs in STP group 2 share a different spanning tree.

All the portss are tagged. The ports must be tagged so that they can be in both a member VLAN and the member's master VLAN. For example, ports 1/1 - 1/4 are in member VLAN 3 and also in master VLAN 2 (since master VLAN 2 contains member VLAN 3).

STP load balancing

Notice that the STP groups each have different STP priorities. In configurations that use the STP groups on multiple devices, you can use the STP priorities to load balance the STP traffic. By setting the STP priorities for the same STP group to different values on each device, you can cause each of the devices to be the root bridge for a different STP group. This type of configuration distributes the traffic evenly across the devices and also ensures that ports that are blocked in one STP group spanning tree are used by another STP group spanning tree for forwarding. Refer to [Configuration example for STP load sharing](#) on page 304 for an example using STP load sharing.

Configuring STP per VLAN group

To configure STP per VLAN group, perform the following tasks:

1. Configure the member VLANs.
2. Optionally, configure master VLANs to contain the member VLANs. This is useful when you have a lot of member VLANs and you do not want to individually configure STP on each one. Each of the member VLANs in the STP group uses the STP settings of the master VLAN.
3. Configure the STP groups. Each STP group runs a separate instance of STP.

The following CLI commands implement the STP per VLAN group configuration shown in [Figure 62](#) on page 303. The following commands configure the member VLANs (3, 4, 13, and 14) and the master VLANs (2 and 12). Notice that changes to STP parameters are made in the master VLANs only, not in the member VLANs.

```
device(config)#vlan 2
device(config-vlan-2)#spanning-tree priority 1
device(config-vlan-2)#tagged ethernet 1/1 to 1/4
```

```
device(config-vlan-2)#vlan 3
device(config-vlan-3)#tagged ethernet 1/1 to 1/4
device(config-vlan-3)#vlan 4
device(config-vlan-4)#tagged ethernet 1/1 to 1/4
device(config-vlan-4)#vlan 12
device(config-vlan-12)#spanning-tree priority 2
device(config-vlan-12)#tagged ethernet 1/1 to 1/4
device(config-vlan-13)#vlan 13
device(config-vlan-13)#tagged ethernet 1/1 to 1/4
device(config-vlan-14)#vlan 14
device(config-vlan-14)#tagged ethernet 1/1 to 1/4
device(config-vlan-14)#exit
```

The following commands configure the STP groups.

```
device(config)#stp-group 1
device(config-stp-group-1)#master-vlan 2
device(config-stp-group-1)#member-vlan 3 to 4
device(config-stp-group-1)#exit
device(config)#stp-group 2
device(config-stp-group-2)#master-vlan 12
device(config-stp-group-2)#member-vlan 13 to 14
```

Syntax: [no] stp-group num

This command changes the CLI to the STP group configuration level. The following commands are valid at this level. The num parameter specifies the STP group ID and can be from 1 - 32.

Syntax: [no] master-vlan num

This command adds a master VLAN to the STP group. The master VLAN contains the STP settings for all the VLANs in the STP per VLAN group. The num parameter specifies the VLAN ID. An STP group can contain one master VLAN.

If you delete the master VLAN from an STP group, the software automatically assigns the first member VLAN in the group to be the new master VLAN for the group.

Syntax: [no] member-vlan num [to num]

This command adds additional VLANs to the STP group. These VLANs also inherit the STP settings of the master VLAN in the group.

Syntax: [no] member-group num

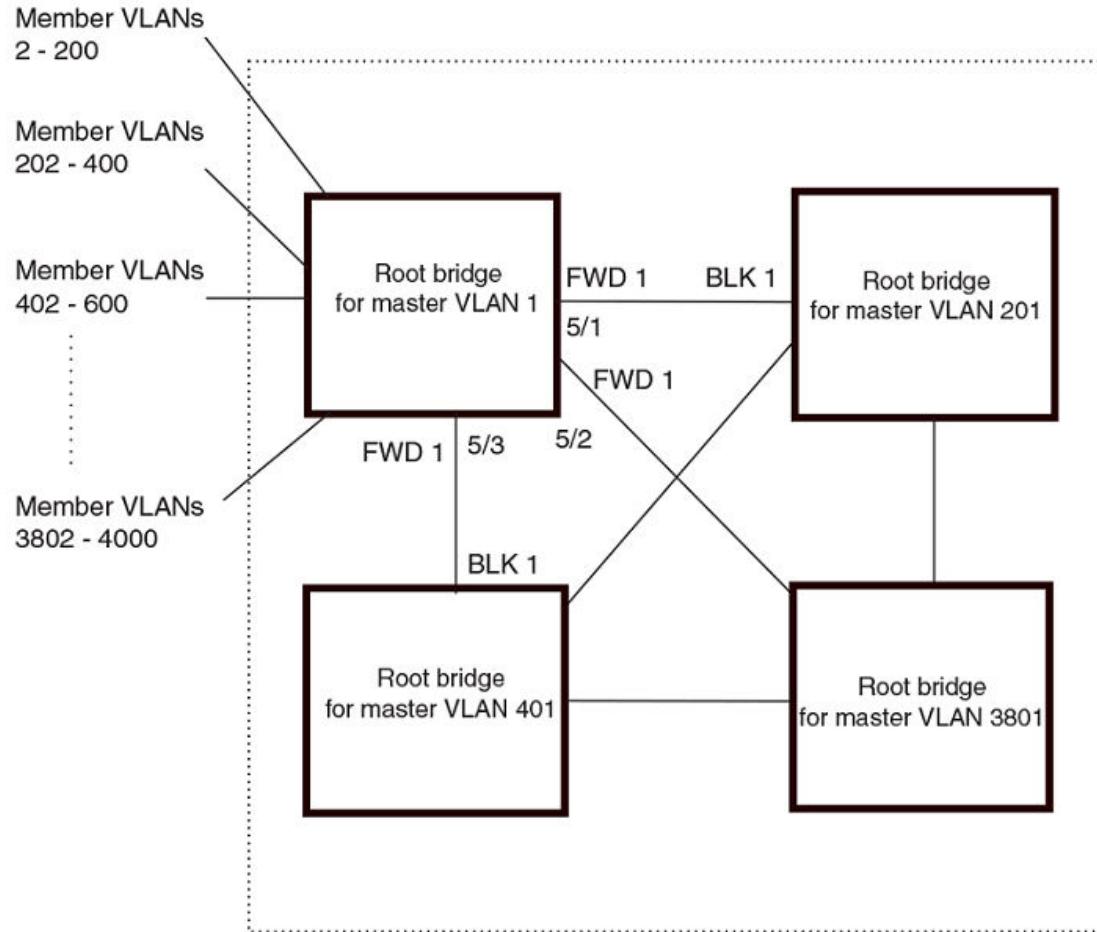
This command adds a member group (a VLAN group) to the STP group. All the VLANs in the member group inherit the STP settings of the master VLAN in the group. The num parameter specifies the VLAN group ID.

NOTE

This command is optional and is not used in the example above. For an example of this command, refer to [Configuration example for STP load sharing](#) on page 304.

Configuration example for STP load sharing

The following figure shows another example of a STP per VLAN group implementation.

FIGURE 63 More complex STP per VLAN group example

In this example, each of the devices in the core is configured with a common set of master VLANs, each of which contains one or more member VLANs. Each of the member VLANs in an STP group runs the same instance of STP and uses the STP parameters configured for the master VLAN.

The STP group ID identifies the STP instance. All VLANs within an STP group run the same instance of STP. The master VLAN specifies the bridge STP parameters for the STP group, including the bridge priority. In this example, each of the devices in the core is configured to be the default root bridge for a different master VLAN. This configuration ensures that each link can be used for forwarding some traffic. For example, all the ports on the root bridge for master VLAN 1 are configured to forward BPDU for master VLAN spanning tree. Ports on the other devices block or forward VLAN 1 traffic based on STP convergence. All the ports on the root bridge for VLAN 2 forward VLAN 2 traffic, and so on.

All the portss are tagged. The ports must be tagged so that they can be in both a member VLAN and the member's master VLAN. For example, port 1/1 - and ports 5/1, 5/2, and 5/3 are in member VLAN 2 and master VLAN 1 (since master VLAN a contains member VLAN 2).

Here are the commands for configuring the root bridge for master VLAN 1 in figure [Figure 62](#) on page 303 for STP per VLAN group. The first group of commands configures the master VLANs. Notice that the STP priority is set to a different value for each VLAN. In addition, the same VLAN has a different STP priority on each device. This provides load balancing by making each of the devices a root bridge for a different spanning tree.

```
device(config)#vlan 1
device(config-vlan-1)#spanning-tree priority 1
```

```

device(config-vlan-1)#tag ethernet 1/1 ethernet 5/1 to 5/3
device(config-vlan-1)#vlan 201
device(config-vlan-201)#spanning-tree priority 2
device(config-vlan-201)#tag ethernet 1/2 ethernet 5/1 to 5/3
device(config-vlan-201)#vlan 401
device(config-vlan-401)#spanning-tree priority 3
device(config-vlan-401)#tag ethernet 1/3 ethernet 5/1 to 5/3
...
device(config-vlan-3601)#vIan 3801
device(config-vlan-3801)#spanning-tree priority 20
device(config-vlan-3801)#tag ethernet 1/20 ethernet 5/1 to 5/3
device(config-vlan-3801)#exit

```

The next group of commands configures VLAN groups for the member VLANs. Notice that the VLAN groups do not contain the VLAN numbers assigned to the master VLANs. Also notice that no STP parameters are configured for the groups of member VLANs. Each group of member VLANs will inherit its STP settings from its master VLAN.

Set the bridge priority for each master VLAN to the highest priority (1) on one of the devices in the STP per VLAN group configuration. By setting the bridge priority to the highest priority, you make the device the default root bridge for the spanning tree. To ensure STP load balancing, make each of the devices the default root bridge for a different master VLAN.

```

device(config)#vIan-group 1 vIan 2 to 200
device(config-vlan-group-1)#tag ethernet 1/1 ethernet 5/1 to 5/3
device(config-vlan-group-1)#vIan-group 2 vIan 202 to 400
device(config-vlan-group-2)#tag ethernet 1/2 ethernet 5/1 to 5/3
device(config-vlan-group-2)#vIan-group 3 vIan 402 to 600
device(config-vlan-group-2)#tag ethernet 1/3 ethernet 5/1 to 5/3
...
device(config-vlan-group-19)#vIan-group 20 vIan 3082 to 3282
device(config-vlan-group-20)#tag ethernet 1/20 ethernet 5/1 to 5/3
device(config-vlan-group-20)#exit

```

The following group of commands configures the STP groups. Each STP group in this configuration contains one master VLAN, which contains a VLAN group. This example shows that an STP group also can contain additional VLANs (VLANs not configured in a VLAN group).

```

device(config)#stp-group 1
device(config-stp-group-1)#master-vlan 1
device(config-stp-group-1)#member-group 1
device(config-stp-group-1)#member-vlan 4001 4004 to 4010
device(config-stp-group-1)#stp-group 2
device(config-stp-group-2)#master-vlan 201
device(config-stp-group-2)#member-group 2
device(config-stp-group-2)#member-vlan 4002 4003 4011 to 4015
device(config-stp-group-2)#stp-group 3
device(config-stp-group-3)#master-vlan 401
device(config-stp-group-3)#member-group 3
...
device(config-stp-group-19)#stp-group 20
device(config-stp-group-20)#master-vlan 3081
device(config-stp-group-20)#member-group 20

```

PVST/PVST+ compatibility

The FastIron family of switches support Cisco's Per VLAN Spanning Tree plus (PVST+), by allowing the device to run multiple spanning trees (MSTP) while also interoperating with IEEE 802.1Q devices¹.

NOTE

Brocade ports automatically detect PVST+ BPDUs and enable support for the BPDUs once detected. You do not need to perform any configuration steps to enable PVST+ support. However, to support the IEEE 802.1Q BPDUs, you might need to enable dual-mode support.

Support for Cisco's Per VLAN Spanning Tree plus (PVST+), allows a Brocade device to run multiple spanning trees (MSTP) while also interoperating with IEEE 802.1Q devices. Brocade ports automatically detect PVST+ BPDUs and enable support for the BPDUs once detected. The enhancement allows a port that is in PVST+ compatibility mode due to auto-detection to revert to the default MSTP mode when one of the following events occurs:

- The link is disconnected or broken
- The link is administratively disabled
- The link is disabled by interaction with the link-keepalive protocol

This enhancement allows a port that was originally interoperating with PVST+ to revert to MSTP when connected to a Brocade device.

1 Cisco user documentation for PVST/PVST+ refers to the IEEE 802.1Q spanning tree as the Common Spanning Tree (CST).

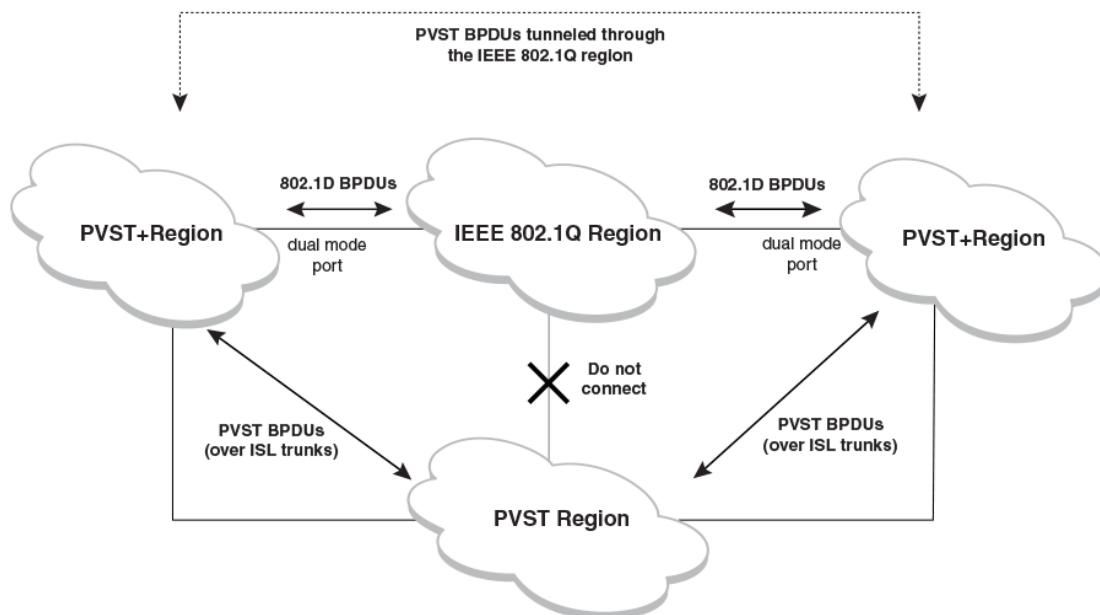
Overview of PVST and PVST+

Per VLAN Spanning Tree (PVST) is a Cisco proprietary protocol that allows a Cisco device to have multiple spanning trees. The Cisco device can interoperate with spanning trees on other PVST devices but cannot interoperate with IEEE 802.1Q devices. An IEEE 802.1Q device has all its ports running a single spanning tree. PVST+ is an extension of PVST that allows a Cisco device to also interoperate with devices that are running a single spanning tree (IEEE 802.1Q).

Enhanced PVST+ support allows a Brocade device to interoperate with PVST spanning trees and the IEEE 802.1Q spanning tree at the same time.

IEEE 802.1Q and PVST regions cannot interoperate directly but can interoperate indirectly through PVST+ regions. PVST BPDUs are tunneled through 802.1Q regions, while PVST BPDUs for VLAN 1 (the IEEE 802.1Q VLAN) are processed by PVST+ regions. The following figure shows the interaction of IEEE 802.1Q, PVST, and PVST+ regions.

FIGURE 64 Interaction of IEEE 802.1Q, PVST, and PVST+ regions



VLAN tags and dual mode

The dual-mode feature enables a port to send and receive both tagged and untagged frames. When the dual-mode feature is enabled on a port, the port is an untagged member of one of its VLANs and is at the same time a tagged member of all its other VLANs. The untagged frames are supported on the port Port Native VLAN .

The dual-mode feature must be enabled on a Brocade port in order to interoperate with another vendor device. Some vendors use VLAN 1 by default to support the IEEE 802.1Q-based standard spanning tree protocols, such as 802.1d and 802.1w for sending untagged frames on VLAN 1. On Brocade switches, by default, the Port Native VLAN is the same as the Default VLAN , which is VLAN 1. Thus, to support IEEE 802.1Q in a typical configuration, a port must be able to send and receive untagged frames for VLAN 1 and tagged frames for the other VLANs, and interoperate with other vendor devices using VLAN 1.

If you want to use tagged frames on VLAN 1, you can change the default VLAN ID to an ID other than 1. You also can specify the VLAN on which you want the port to send and receive untagged frames (the Port Native VLAN). The Port Native VLAN ID does not need to be the same as the default VLAN. Make sure that the untagged (native) VLAN is also changed on the interoperating vendor side to match that on the Brocade side.

To support the IEEE 802.1Q with non-standard proprietary protocols such as PVST and PVST+, a port must always send and receive untagged frames on VLAN 1 on both sides. In this case, enable the dual-mode 1 feature to allow untagged BPDUs on VLAN 1 and use Native VLAN 1 on the interoperating vendor side. You should not use VLAN 1 for tagged frames in this case.

Configuring PVST+ support

PVST+ support is automatically enabled when the port receives a PVST BPDU. You can manually enable the support at any time or disable the support if desired.

If you want a tagged port to also support IEEE 802.1Q BPDUs, you need to enable the dual-mode feature on the port. The dual-mode feature is disabled by default and must be enabled manually.

A port that is in PVST+ compatibility mode due to auto-detection reverts to the default MSTP mode when one of the following events occurs:

- The link is disconnected or broken
- The link is administratively disabled
- The link is disabled by interaction with the link-keepalive protocol

This allows a port that was originally interoperateing with PVST+ to revert to MSTP when connected to a Brocade device.

Enabling PVST+ support manually

To immediately enable PVST+ support on a port, enter commands such as the following.

```
device(config)#interface ethernet 1/1
device(config-if-1/1)#pvst-mode
```

Syntax: [no] pvst-mode

NOTE

If you disable PVST+ support, the software still automatically enables PVST+ support if the port receives a BPDU with PVST+ format.

NOTE

If 802.1W and pvst-mode (either by auto-detection or by explicit configuration) are enabled on a tagged VLAN port, 802.1W will treat the PVST BPDUs as legacy 802.1D BPDUs.

Enabling dual-mode support

To enable the dual-mode feature on a port, enter the following command at the interface configuration level for the port.

```
device(config-if-1/1)#dual-mode
```

Syntax: [no] dual-mode [vlan-id]

The *vlan-id* specifies the port Port Native VLAN. This is the VLAN on which the port will support untagged frames. By default, the Port Native VLAN is the same as the default VLAN (which is VLAN 1 by default).

For more information about the dual-mode feature, refer to [Dual-mode VLAN ports](#) on page 409.

Displaying PVST+ support information

To display PVST+ information for ports on a Brocade device, enter the following command at any level of the CLI.

```
device#show span pvst-mode
PVST+ Enabled on:
Port      Method
1/1      Set by configuration
1/2      Set by configuration
2/10     Set by auto-detect
3/12     Set by configuration
4/24     Set by auto-detect
```

Syntax: show span pvst-mode

This command displays the following information.

TABLE 51 CLI display of PVST+ information

Field	Description
Port	The Brocade port number.
NOTE	
The command lists information only for the ports on which PVST+ support is enabled.	
Method	The method by which PVST+ support was enabled on the port. The method can be one of the following: <ul style="list-style-type: none"> • Set by configuration - You enabled the support. • Set by auto-detect - The support was enabled automatically when the port received a PVST+ BPDU.

PVST+ configuration examples

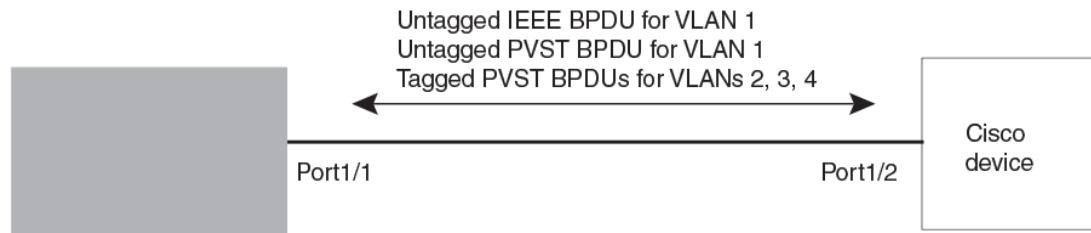
The following examples show configuration examples for two common configurations:

- Untagged IEEE 802.1Q BPDUs on VLAN 1 and tagged PVST+ BPDUs on other VLANs
- Tagged IEEE 802.1Q BPDUs on VLAN 1 and untagged BPDUs on another VLAN

Tagged port using default VLAN 1 as its port native VLAN

The following table shows an example of a PVST+ configuration that uses VLAN 1 as the untagged default VLAN and VLANs 2, 3, and 4 as tagged VLANs.

FIGURE 65 Default VLAN 1 for untagged BPDU



To implement this configuration, enter the following commands.

Commands on the Brocade Device

```

device(config)#vlan-group 1 vlan 2 to 4
device(config-vlan-group-1)#tagged ethernet 1/1
device(config-vlan-group-1)#exit
device(config)#interface ethernet 1/1
device(config-if-1/1)#dual-mode
device(config-if-1/1)#pvst-mode
  
```

These commands configure a VLAN group containing VLANs 2, 3, and 4, add port 1/1 as a tagged port to the VLANs, and enable the dual-mode feature and PVST+ support on the port. The dual-mode feature allows the port to send and receive untagged frames for the default VLAN (VLAN 1 in this case) in addition to tagged frames for VLANs 2, 3, and 4. Enabling the PVST+ support ensures that the port is ready to send and receive PVST+ BPDUs. If you do not manually enable PVST+ support, the support is not enabled until the port receives a PVST+ BPDU.

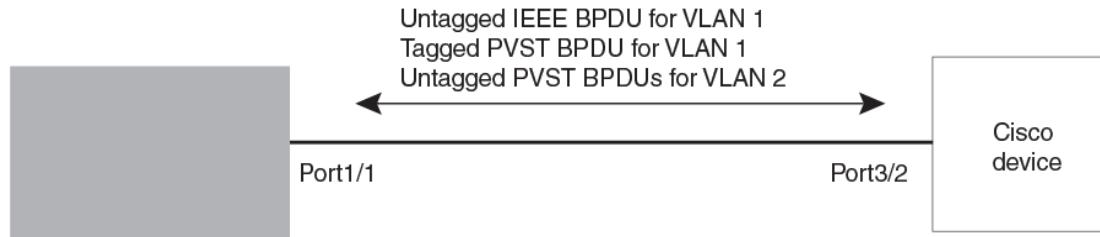
The configuration leaves the default VLAN and the port Port Native VLAN unchanged. The default VLAN is 1 and the port Port Native VLAN also is 1. The dual-mode feature supports untagged frames on the default VLAN only. Thus, port 1/1 can send and receive untagged BPDUs for VLAN 1 and can send and receive tagged BPDUs for the other VLANs.

Port 1/1 will process BPDUs as follows:

- Process IEEE 802.1Q BPDUs for VLAN 1.
- Process tagged PVST BPDUs for VLANs 2, 3, and 4.
- Drop untagged PVST BPDUs for VLAN 1.

Untagged port using VLAN 2 as port native VLAN

The following figure shows an example in which a port Port Native VLAN is not VLAN 1. In this case, VLAN 1 uses tagged frames and VLAN 2 uses untagged frames.

FIGURE 66 Port Native VLAN 2 for Untagged BPDUs

To implement this configuration, enter the following commands.

Commands on the Brocade Device

```
device(config)#default-vlan-id 4000
device(config)#vlan 1
device(config-vlan-1)#tagged ethernet 1/1
device(config-vlan-1)#exit
device(config)#vlan 2
device(config-vlan-2)#tagged ethernet 1/1
device(config-vlan-2)#exit
device(config)#interface ethernet 1/1
device(config-if-1/1)#dual-mode 2
device(config-if-1/1)#pvst-mode
device(config-if-1/1)#exit
```

These commands change the default VLAN ID, configure port 1/1 as a tagged member of VLANs 1 and 2, and enable the dual-mode feature and PVST+ support on port 1/1. Since VLAN 1 is tagged in this configuration, the default VLAN ID must be changed from VLAN 1 to another VLAN ID. Changing the default VLAN ID from 1 allows the port to process tagged frames for VLAN 1. VLAN 2 is specified with the **dual-mode** command, which makes VLAN 2 the port Port Native VLAN. As a result, the port processes untagged frames and untagged PVST BPDUs on VLAN 2.

NOTE

Although VLAN 2 becomes the port untagged VLAN, the CLI still requires that you add the port to the VLAN as a tagged port, since the port is a member of more than one VLAN.

Port 1/1 will process BPDUs as follows:

- Process IEEE 802.1Q BPDUs for VLAN 1.
- Process untagged PVST BPDUs for VLAN 2.
- Drop tagged PVST BPDUs for VLAN 1.

Note that when VLAN 1 is not the default VLAN, the ports must have the dual-mode feature enabled in order to process IEEE 802.1Q BPDUs.

For example, the following configuration is incorrect.

```
device(config)#default-vlan-id 1000
device(config)#vlan 1
device(config-vlan-1)#tagged ethernet 1/1 to 1/2
device(config-vlan-1)#exit
device(config)#interface ethernet 1/1
device(config-if-1/1)#pvst-mode
device(config-if-1/1)#exit
device(config)#interface ethernet 1/2
device(config-if-1/2)#pvst-mode
device(config-if-1/2)#exit
```

In the configuration above, all PVST BPDUs associated with VLAN 1 would be discarded. Since IEEE BPDUs associated with VLAN 1 are untagged, they are discarded because the ports in VLAN 1 are tagged. Effectively, the BPDUs are never processed by the Spanning Tree Protocol. STP assumes that

there is no better bridge on the network and sets the ports to FORWARDING. This could cause a Layer 2 loop.

The following configuration is correct.

```
device(config)#default-vlan-id 1000
device(config)#vlan 1
device(config-vlan-1)#tagged ethernet 1/1 to 1/2
device(config-vlan-1)#exit
device(config)#interface ethernet 1/1
device(config-if-1/1)#pvst-mode
device(config-if-1/1)#dual-mode
device(config-if-1/1)#exit
device(config)#interface ethernet 1/2
device(config-if-1/2)#pvst-mode
device(config-if-1/2)#dual-mode
device(config-if-1/2)#exit
```

Setting the ports as dual-mode ensures that the untagged IEEE 802.1Q BPDUs reach the VLAN 1 instance.

PVRST compatibility

PVRST, the "rapid" version of per-VLAN spanning tree (PVST), is a Cisco proprietary protocol. PVRST corresponds to the Brocade full implementation of IEEE 802.1w (RSTP). Likewise, PVST, also a Cisco proprietary protocol, corresponds to the Brocade implementation of IEEE 802.1D (STP). When a Brocade device receives PVRST BPDUs on a port configured to run 802.1w, it recognizes and processes these BPDUs and continues to operate in 802.1w mode.

PVRST compatibility is automatically enabled when a port receives a PVRST BPDU.

BPDU guard

In an STP environment, switches, end stations, and other Layer 2 devices use Bridge Protocol Data Units (BPDUs) to exchange information that STP will use to determine the best path for data flow.

The BPDU guard, an enhancement to STP, removes a node that reflects BPDUs back in the network. It enforces the STP domain borders and keeps the active topology predictable by not allowing any network devices behind a BPDU guard-enabled port to participate in STP.

In some instances, it is unnecessary for a connected device, such as an end station, to initiate or participate in an STP topology change. In this case, you can enable the STP BPDU guard feature on the Brocade port to which the end station is connected. STP BPDU guard shuts down the port and puts it into an errdisable state. This disables the connected device's ability to initiate or participate in an STP topology. A log message is then generated for a BPDU guard violation, and a CLI message is displayed to warn the network administrator of a severe invalid configuration. The BPDU guard feature provides a secure response to invalid configurations because the administrator must manually put the interface back in service if errdisable recovery is not enabled.

NOTE

BPDU guard is not supported on tagged ports. It can be configured on a tagged port, but the configuration will have no effect.

Enabling BPDU protection by port

You can enable STP BPDU guard on individual interfaces. The feature is disabled by default.

To enable STP BPDU guard on a specific port, enter a command such as the following.

NOTE

Spanning tree must be enabled on the corresponding VLAN.

```
device(config) interface eth1 2/1
device(config-if-e1000-2/1)#stp-bpdu-guard
```

Syntax: [no] stp-bpd-guard

The **no** parameter disables the BPDU guard on this interface.

You can also use the multiple interface command to enable this feature on multiple ports at once.

```
device(config)#interface ethernet 1/1 to 1/9
device(config-mif-1/1-1/9)#stp-bpdu-guard
device(config-mif-1/1-1/9) #
```

This will enable stp-bpdu-guard on ports 0/1/1 to 0/1/9

Re-enabling ports disabled by BPDU guard

When a BPDU Guard-enabled port is disabled by BPDU Guard, the Brocade device will place the port in **errdisabled** state and display a message on the console indicating that the port is errdisabled (refer to [BPDU guard status example console messages](#) on page 314). In addition, the **show interface** command output will indicate that the port is errdisabled.

```
device#show int e 2
Gigabit Ethernet2 is ERR-DISABLED (bpduguard), line protocol is down
```

To re-enable a port that is in **errdisabled** state, you must first disable the port then re-enable it. Enter commands such as the following.

```
device(config)#int e 2
device(config-if-e1000-2)#disable
device(config-if-e1000-2)#enable
```

If you attempt to enable an errdisabled port without first disabling it, the following error message will appear on the console.

```
device(config-if-e1000-2)#enable
Port 2 is errdisabled, do disable first and then enable to enable it
```

Displaying the BPDU guard status

To display the BPDU guard state, enter the **show running configuration** or the **show stp-bpdu-guard** command.

For FastIron X Series devices, enter the **stp-bpdu-guard** command.

```
device#show stp-bpdu-guard
BPDU Guard Enabled on:
Interface Violation
Port 1 No
```

```

Port 2 No
Port 3 No
Port 4 No
Port 5 No
Port 6 No
Port 7 No
Port 8 No
Port 9 No
Port 10 No
Port 11 No
Port 12 Yes
Port 13 No

```

BPDU guard status example configurations

The following example shows how to configure BPDU guard at the interface level and to verify the configuration by issuing the **show stp-bpdu-guard** and the **show interface** commands.

```

device(config)#interface ethernet 1
device(config-if-e1000-1)#stp-bpdu-guard
device(config-if-e1000-1)#
device(config-if-e1000-1)#show stp-bpdu-guard
BPDU Guard Enabled on:
Port
1
device(config-if-e1000-1)#
device(config-if-e1000-1)#show interfaces ethernet 1
GigabitEthernet1 is up, line protocol is up
Port up for 40 seconds
Hardware is GigabitEthernet, address is 0000.00a0.7100 (bia 0000.00a0.7100)
Configured speed auto, actual 100Mbit, configured duplex fdx, actual fdx
Configured mdi mode AUTO, actual MDI
Member of L2 VLAN ID 2, port is untagged, port state is FORWARDING
BPDU guard is Enabled
, ROOT protect is Disabled
STP configured to ON, priority is level0, flow control enabled
mirror disabled, monitor disabled
Not member of any active trunks
Not member of any configured trunks
No port name
Mac-notification is Enabled
IPG MII 96 bits-time, IPG GMII 96 bits-time
IP MTU 1500 bytes
300 second input rate: 8 bits/sec, 0 packets/sec, 0.00% utilization
300 second output rate: 256 bits/sec, 0 packets/sec, 0.00% utilization
88 packets input, 15256 bytes, 0 no buffer
Received 75 broadcasts, 13 multicasts, 0 unicasts
1 input errors, 0 CRC, 0 frame, 0 ignored
0 runts, 0 giants
4799 packets output, 313268 bytes, 0 underruns
Transmitted 90 broadcasts, 4709

```

NOTE

The port up/down time is required only for physical ports and not for loopback/ ve/ tunnel ports.

BPDU guard status example console messages

A console message such as the following is generated after a BPDU guard violation occurs on a system that is running MSTP.

```
device(config-if-e1000-23)#MSTP: Received BPDU on BPDU guard enabled Port
23,errdisable Port 23
```

A console message such as the following is generated after a BPDU guard violation occurs on a system that is running STP.

```
device(config)#STP: Received BPDU on BPDU guard enabled Port 23 (vlan=1), errdisable
Port 23
```

A console message such as the following is generated after a BPDU guard violation occurs on a system that is running RSTP.

```
device(config-vlan-1)#RSTP: Received BPDU on BPDU guard enabled Port 23
(vlan=1), errdisable Port 23
```

Root guard

The standard STP (802.1D), RSTP (802.1W) or 802.1S does not provide any way for a network administrator to securely enforce the topology of a switched layer 2 network. The forwarding topology of a switched network is calculated based on the root bridge position, along with other parameters. This means any switch can be the root bridge in a network as long as it has the lowest bridge ID. The administrator cannot enforce the position of the root bridge. A better forwarding topology comes with the requirement to place the root bridge at a specific predetermined location. Root Guard can be used to predetermine a root bridge location and prevent rogue or unwanted switches from becoming the root bridge.

When root guard is enabled on a port, it keeps the port in a designated role. If the port receives a superior STP Bridge Protocol Data Units (BPDU), it puts the port into a ROOT-INCONSISTANT state and triggers a log message and an SNMP trap. The ROOT-INCONSISTANT state is equivalent to the BLOCKING state in 802.1D and to the DISCARDING state in 802.1W. No further traffic is forwarded on this port. This allows the bridge to prevent traffic from being forwarded on ports connected to rogue or misconfigured STP bridges.

Once the port stops receiving superior BPDUs, root guard automatically sets the port back to learning, and eventually to a forwarding state through the spanning-tree algorithm.

Configure root guard on all ports where the root bridge should not appear. This establishes a protective network perimeter around the core bridged network, cutting it off from the user network.

NOTE

Root guard may prevent network connectivity if it is improperly configured. Root guard must be configured on the perimeter of the network rather than the core.

NOTE

Root guard is not supported when MSTP is enabled.

Enabling STP root guard

An STP root guard is configured on an interface by entering commands similar to the following.

```
device(config)#interface ethernet 5/5
device(config-if-e10000-5/5) spanning-tree root-protect
```

Syntax: [no] spanning-tree root-protect

Enter the **no** form of the command to disable STP root guard on the port.

Displaying the STP root guard

To display the STP root guard state, enter the **show running configuration** or the **show span root-protect** command.

```
device#show span root-protect
Root Protection Enabled on:
Port 1
```

Syntax: **show span root-protect**

Displaying the root guard by VLAN

You can display root guard information for all VLANs or for a specific VLAN. For example, to display root guard violation information for VLAN 7.

Syntax: **show spanning-tree [vlan-id]**

If you do not specify a *vlan-id*, information for all VLANs is displayed. For example, to display root guard violation information for VLAN 7.

```
device#show spanning-tree vlan 7
STP instance owned by VLAN 7
Global STP (IEEE 802.1D) Parameters:
VLAN Root Root Root Prio Max He- Ho- Fwd Last Chg Bridge
ID ID Cost Port rity Age llo ld dly Chang cnt Address
Hex sec sec sec sec sec
7 a00000001112220 0 Root a000 20 2 1 15 4 4 00001112220
Port STP Parameters:
Port Prio Path State Fwd Design Designated Designated
Num rity Cost Trans Cost Root Bridge
Hex
1 80 19 ROOT-INCONS 2 0 a00000001112220 a00000001112220
```

Designated Protection

Designated Protection ensures that a port cannot go to the designated forwarding state in STP 802.1d or 802.1w. For example, a fast uplink port should never become a designated port to avoid loops in a network topology. It should either be a root port in any STP state or a non-root port in a blocking state. You can enable Designated Protection on the port to ensure that it does not go to the designated forwarding state. If STP tries to put this port into the designated forwarding state, the device puts this port into a designated inconsistent STP state. This is effectively equivalent to the listening state in STP in which a port cannot transfer any user traffic. When STP no longer marks this port as a designated port, the port is automatically removed from the designated inconsistent state.

Designation Protection is a port-level feature, while the designated inconsistent state is a per-STP-instance, per-port state. In PVST, a port can belong to several VLANs where each VLAN runs a separate spanning tree instance. The designated inconsistent state in one spanning tree instance does not affect the traffic in other spanning tree instances.

For example, consider an interface eth 1 that is in VLAN 20 and VLAN 50. VLAN 20 runs one instance of STP and VLAN 50 runs another instance. Interface eth1 can be in the designated inconsistent state for VLAN 50 and block the VLAN 50 traffic while it is in root forwarding state for VLAN 20 and allow VLAN 20 traffic.

You can view the status of the Designated Protection feature on a port with the **show interface ethernet** command for that port.

NOTE

You cannot enable Designated Protection and Root Guard on the same port.

Enabling Designated Protection on a port

To disallow the designated forwarding state on a port in STP (802.1d or 802.1w), run the **spanning-tree designated-protect** command in interface configuration mode for that port.

The following example shows that the designated forwarding state is disallowed on Ethernet interface 1/1/1.

```
Brocade(config)# ethernet interface 1/1/1
Brocade(config-if-e1000-1/1/1)# spanning-tree designated-protect
```

Syslog message for a port in designated inconsistent state

The following syslog message is generated when a port is put in the designated inconsistent state.
5d19h00m12s:I:STP: VLAN 100 Designated-protect port 2/1/7, inconsistent, Put into Listening state

Error disable recovery

If a BPDU Guard violation or loop detection violation occurs, or the number of inError packets exceeds the configured threshold, or if an EFM-OAM enabled interface receives a critical event from the remote device, a port is placed into an error-disabled state, which is functionally equivalent to a disable state. Once in an error-disabled state, the port remains in that state until it is enabled either automatically or manually.

Enabling an error-disabled port automatically

To enable a port to recover automatically from the error-disabled state after the expiry of a configured error recovery timer, run the **errdisable recovery cause** command in global configuration mode.

For example, to enable error-disable recovery for BPDU guard, enter the following command:

```
Brocade(config)# errdisable recovery cause bpduguard
```

Syntax: errdisable recovery cause [bpduguard | loam-critical-event | loop-detection | packet-inerror-detect | all]

The **bpduguard** parameter allows the port to recover from the error-disabled state, if the state was caused by a BPDU guard violation.

The **all** parameter allows ports to recover from an error-disabled state caused by any reason, for example, a BPDU guard violation, due to reception of a critical event from the remote device in the case of EFM-OAM interface, the number of inError packets exceeding the configured threshold, or a loop-detection violation.

NOTE

When automatic recovery re-enables the port, the port is not in the error-disabled state, but it can remain down for other reasons, such as the Tx/Rx of the fibre optic not being seated properly. Thus, the port is not able to receive the signal from the other side. In this case, after the optic is inserted correctly, you should manually disable the port and then enable it.

Enabling an error-disabled port manually

To enable an error-disabled port manually, you must first run the **disable** command and then the **enable** command in interface configuration mode to disable the port and then enable the port respectively.

Setting the recovery interval

The **errdisable recovery interval** command allows you to configure a timeout for ports in the error-disabled state, after which the ports are re-enabled automatically. To set the error-disabled recovery timeout interval, enter the following command:

```
Brocade(config)# errdisable recovery interval 20
```

Syntax: [no] errdisable recovery interval seconds

The **seconds** variable allows you to set the timeout value for the recovery mechanism when the port is in an errdisabled state. Once this timeout value expires, the ports are automatically re-enabled. Valid values are from 10 through 65535 seconds.

Displaying the error disable recovery state by interface

The port status of errdisabled displays in the output of the **show interface** and the **show interface brief** commands. In this example, errdisable is enabled on interface ethernet 1 and errdisnable is enabled because of a BPDU guard violation.

```
device#show interfaces ethernet 1
GigabitEthernet1 is ERR-DISABLED (bpduguard),
line protocol is down
BPDU guard is Enabled, ROOT protect is Disabled
Port down for 2 hours 45 minutes 10 seconds
Hardware is GigabitEthernet, address is 0000.00a0.7100 (bia 0000.00a0.7100)
Configured speed auto, actual unknown, configured duplex fdx, actual unknown
Configured mdi mode AUTO, actual unknown
Member of L2 VLAN ID 2, port is untagged, port state is DISABLED
STP configured to ON, priority is level0, flow control enabled
mirror disabled, monitor disabled
Not member of any active trunks
Not member of any configured trunks
No port name
IPG MII 96 bits-time, IPG GMII 96 bits-time
IP MTU 1500 bytes
300 second input rate: 0 bits/sec, 0 packets/sec, 0.00% utilization
300 second output rate: 0 bits/sec, 0 packets/sec, 0.00% utilization
145 packets input, 23561 bytes, 0 no buffer
Received 124 broadcasts, 21 multicasts, 0 unicasts
1 input errors, 0 CRC, 0 frame, 0 ignored
0 runts, 0 giants
5067 packets output, 330420 bytes, 0 underruns
Transmitted 90 broadcasts, 4977 multicasts, 0 unicasts
0 output errors, 0 collisions
```

Displaying the recovery state for all conditions

Use the **show errdisable recovery** command to display all the default error disable recovery state for all possible conditions. In this example, port 6 is undergoing a recovery.

```
device#show errdisable recovery
ErrDisable Reason Timer Status
-----
all reason Disabled
bpduGuard Enabled
Timeout Value: 300 seconds
Interface that will be enabled at the next timeout:
Interface Errdisable reason Time left (sec)
-----
Port 6 bpduGuard 297
```

Syntax: **show errdisable recovery**

Displaying the recovery state by port number and cause

To see which ports are under an errdisabled state, use the **show errdisable summary** command. This command not only shows the port number, but also displays the reason why the port is in an errdisabled state and the method used to recover the port. In this example, port 6 is errdisabled for a BPDU guard violation.

```
device#show errdisable summary
Port 6 ERR_DisABLED for bpduGuard
```

Syntax: **show errdisable summary**

Errdisable Syslog messages

When the system places a port into an errdisabled state for BPDU guard, a log message is generated. When the errdisable recovery timer expires, a log message is also generated.

A Syslog message such as the following is generated after a port is placed into an errdisabled state for BPDU guard.

```
STP: VLAN 50 BPDU-guard port 3 detect (Received BPDU), putting into err-disable state
A Syslog message such as the following is generated after the recovery timer expires.
```

```
ERR_DISABLE: Interface ethernet 3, err-disable recovery timeout
```

802.1s Multiple Spanning Tree Protocol

Multiple Spanning Tree Protocol (MSTP), as defined in IEEE 802.1s, allows multiple VLANs to be managed by a single STP instance and supports per-VLAN STP. As a result, several VLANs can be mapped to a reduced number of spanning-tree instances. This ensures loop-free topology for one or more VLANs that have the similar layer-2 topology. The Brocade implementation supports up to 16 spanning tree instances in an MSTP enabled bridge which means that it can support up to 16 different Layer 2 topologies. The spanning tree algorithm used by MSTP is RSTP which provides quick convergence.

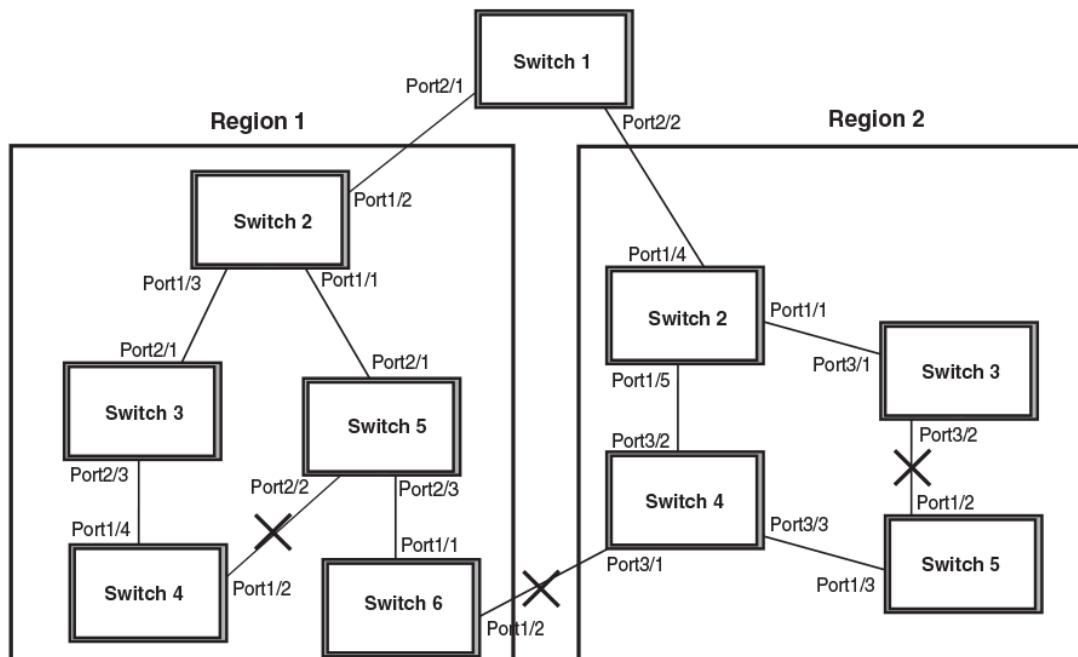
Multiple spanning-tree regions

Using MSTP, the entire network runs a common instance of RSTP. Within that common instance, one or more VLANs can be individually configured into distinct regions. The entire network runs the common spanning tree instance (CST) and the regions run a local instance. The local instance is known as Internal Spanning Tree (IST). The CST treats each instance of IST as a single bridge. Consequently, ports are blocked to prevent loops that might occur within an IST and also throughout the CST. With the exception of the provisions for multiple instances, MSTP operates exactly like RSTP.

For example, in [Figure 67](#) a network is configured with two regions: Region 1 and Region 2. The entire network is running an instance of CST. Each of the regions is running an instance of IST. In addition, this network contains Switch 1 running MSTP that is not configured in a region and consequently is running in the CIST instance. In this configuration, the regions are each regarded as a single bridge to the rest of the network, as is Switch 1. The CST prevents loops from occurring across the network. Consequently, a port is blocked at port 1/2 of switch 4.

Additionally, loops must be prevented in each of the IST instances. Within the IST Region 1, a port is blocked at port 1/2 of switch 4 to prevent a loop in that region. Within Region 2, a port is blocked at port 3/2 of switch 3 to prevent a loop in that region.

FIGURE 67 MSTP configured network



The following definitions describe the STP instances that define an MSTP configuration.

Common Spanning (CST) - CST is defined in 802.1q and assumes one spanning-tree instance for the entire bridged network regardless of the number of VLANs. In MSTP, an MSTP region appears as a virtual bridge that runs CST.

Internal Spanning Tree (IST) - IST is a new terminology introduced in 802.1s. An MSTP bridge must handle at least these two instances: one IST and one or more MSTIs (Multiple Spanning Tree Instances). Within each MST region, the MSTP maintains multiple spanning-tree instances. Instance 0 is a special instance known as IST, which extends CST inside the MST region. IST always exists if the switch runs MSTP. Besides IST, this implementation supports up to 15 MSTIs, numbered from 1 to 4094.

Common and Internal Spanning Trees (CIST) - CIST is a collection of the ISTs in each MST region and the CST that interconnects the MST regions and single spanning trees.

Multiple Spanning Tree Instance (MSTI) - The MSTI is identified by an MST identifier (MSTid) value between 1 and 4094.

MSTP Region - These are clusters of bridges that run multiple instances of the MSTP protocol. Multiple bridges detect that they are in the same region by exchanging their configuration (instance to VLAN mapping), name, and revision-level. Therefore, if you need to have two bridges in the same region, the two bridges must have identical configurations, names, and revision-levels. Also, one or more VLANs can be mapped to one MSTP instance (IST or MSTI) but a VLAN cannot be mapped to multiple MSTP instances.

NOTE

One or more VLANs can be mapped to one MSTP instance (IST or MSTI) but a VLAN cannot be mapped to multiple MSTP instances.

Configuration notes

When configuring MSTP, note the following:

- With MSTP running, enabling static trunk on ports that are members of many VLANs (4000 or more VLANs) will keep the system busy for 20 to 25 seconds.

Configuring MSTP mode and scope

With the introduction of MSTP, a system can be either under MSTP mode or not under MSTP mode. The default state is to not be under MSTP mode. MSTP configuration can only be performed in a system under MSTP mode.

With a system configured under MSTP mode, there is a concept called MSTP scope. MSTP scope defines the VLANs that are under direct MSTP control. You cannot run 802.1D or 802.1w on any VLAN (even outside of MSTP scope) and you cannot create topology groups when a system is under MSTP mode. While a VLAN group will still be supported when a system is under MSTP mode, the member VLAN should either be all in the MSTP scope or all out of the MSTP scope.

When a system is configured from non-MSTP mode to MSTP mode, the following changes are made to the system configuration:

- All 802.1D and 802.1w STP instances are deleted regardless of whether the VLAN is inside the MSTP scope or not.
- All topology groups are deleted.
- Any GVRP configuration is deleted.
- Any VSRP configuration is deleted.
- Single-span (if configured) is deleted.
- MRP running on a VLAN inside MSTP scope is deleted.
- The common and internal spanning trees (CIST) collection is created and all VLANs inside the MSTP scope are attached with the CIST.

Make sure that no physical Layer 2 loops exist prior to switching from non-MSTP mode to MSTP mode. If, for example, you have a Layer 2 loop topology configured as a redundancy mechanism before you perform the switch, a Layer 2 storm should be expected.

To configure a system into MSTP mode, use the following command at the Global Configuration level.

```
Device(config)#mstp scope all
```

Syntax: [no] **mstp scope { all | pvst }**

NOTE

MSTP is not operational however until the **mstp start** command is issued as described in the “Forcing ports to transmit an MSTP BPDU” section.

Once the system is configured into MSTP mode, CIST (sometimes referred to as “instance 0”) is created and all existing VLANs inside the MSTP scope are controlled by CIST. In addition, whenever you create a new VLAN inside MSTP scope, it is put under CIST control by default. In the Brocade MSTP implementation however, a VLAN ID can be pre-mapped to another MSTI as described in the “Configuring an MSTP instance” section. A VLAN whose ID is pre-mapped, will attach to the specified MSTI instead of to the CIST when created.

NOTE

Once under MSTP mode, CIST always controls all ports in the system. If you do not want a port to run MSTP, configure the **no spanning-tree** command under the specified interface configuration.

Configuring **no spanning-tree** command on a system that is configured for MSTP mode changes the system to non-MSTP mode. When this switch is made, all MSTP instances are deleted together with all MSTP configurations. ALL VLANs inside the original MSTP scope will not run any Layer 2 protocols after the switch.

Reduced occurrences of MSTP reconvergence

When a VLAN is deleted, the Brocade device retains the associated VLAN to MSTI mapping instead of deleting it from the configuration. This way, a VLAN can be pre-mapped to an MSTI and MSTP reconvergence may not be necessary when a VLAN is added to or deleted from the configuration. As long as the VLAN being created or deleted is pre-mapped to an MSTI, and the VLAN to MSTI mapping has not changed, MSTP reconvergence will not occur.

NOTE

MSTP reconvergence occurs when the VLAN to MSTI mapping is changed using the **mstp instance** command.

You can optionally remove VLAN to MSTI mappings from the configuration. Refer to [Deleting a VLAN to MSTI mapping](#) on page 323.

The following shows an example application.

Example application of MSTP reconvergence

The following example shows the running configuration file before and after deleting a VLAN from the configuration. The VLAN to MSTI mapping is retained in the running configuration, even after the VLAN is deleted.

```
device(config-vlan-20)#show run
Current configuration:
!
ver 04.2.00bT3e1
!
!
vlan 1 name DEFAULT-VLAN by port
  no spanning-tree
!
```

```

vlan 10 by port
  tagged ethe 1 to 2
  no spanning tree
!
vlan 20 by port           <----- VLAN 20 configuration
  tagged ethe 1 to 2
  no spanning-tree
!
mstp scope all
mstp instance 0 vlan 1
mstp instance 1 vlan 20
mstp start
some lines omitted for brevity...
device(config-vlan-20)#no vlan 20   <----- VLAN 20 deleted
device(config-vlan-20)#show run

Current configuration:
!
ver 04.2.00bT3e1
!
!
vlan 1 name DEFAULT-VLAN by port
  no spanning-tree
!
vlan 10 by port
  tagged ethe 1 to 2
  no spanning-tree
!
mstp scope all
mstp instance 0 vlan 1
mstp instance 1 vlan 10
mstp instance 1 vlan 20           <----- VLAN to MSTI mapping kept in
                                running configuration, even though
                                VLAN 20 was deleted
mstp start
some lines omitted for brevity...

```

Deleting a VLAN to MSTI mapping

You can optionally remove a VLAN to MSTI mapping using the **no mstp instance** command. To do so, enter a command such as the following.

```
device(config)#no mstp instance 7 vlan 4 to 7
```

This command deletes the VLAN to MSTI mapping from the running configuration and triggers an MSTP reconvergence.

Syntax: [no] mstp instance *instance-number* **vlan** *vlan-id* | **vlan-group** *group-id*

The **instance** parameter defines the number for the instance of MSTP that you are deleting.

The **vlan** parameter identifies one or more VLANs or a range of VLANs to the instance defined in this command.

The **vlan-group** parameter identifies one or more VLAN groups to the instance defined in this command.

Viewing the MSTP configuration digest

The MSTP Configuration Digest indicates the occurrence of an MSTP reconvergence. The Configuration Digest is recalculated whenever an MSTP reconvergence occurs. To view the Configuration Digest, use the **show mstp config** command. The following shows an example output.

```
Brocade (config-vlan-20)# show mstp config
MSTP CONFIGURATION
-----
Scope      : all system
Name      :
```

```

Revision      : 0
Version       : 3 (MSTP mode)
Config Digest: 0x9bbda9c70d91f633e1e145fbcbf8d321
Status        : Started
Instance VLANs
-----
0            1
1            10 20

```

Syntax: show mstp config

Configuring additional MSTP parameters

To configure a switch for MSTP, you could configure the name and the revision on each switch that is being configured for MSTP. You must then create an MSTP Instance and assign an ID. VLANs are then assigned to MSTP instances. These instances must be configured on all switches that interoperate with the same VLAN assignments. Port cost, priority and global parameters can then be configured for individual ports and instances. In addition, operational edge ports and point-to-point links can be created and MSTP can be disabled on individual ports.

Each of the commands used to configure and operate MSTP are described in the following:

- [Setting the MSTP name on page 324](#)
- [Setting the MSTP revision number on page 324](#)
- [Configuring an MSTP instance on page 325](#)
- [Configuring bridge priority for an MSTP instance on page 325](#)
- [Setting the MSTP global parameters on page 325](#)
- [Setting ports to be operational edge ports on page 326](#)
- [Setting automatic operational edge ports on page 326](#)
- [Setting point-to-point link on page 327](#)
- [Disabling MSTP on a port on page 328](#)
- [Forcing ports to transmit an MSTP BPDU on page 329](#)
- [Forcing ports to transmit an MSTP BPDU on page 329](#)

Setting the MSTP name

Each switch that is running MSTP is configured with a name. It applies to the switch which can have many different VLANs that can belong to many different MSTP regions.

To configure an MSTP name, use a command such as the following at the Global Configuration level.

```
device(config)#mstp name Brocade
```

Syntax: [no] mstp name *name*

The *name* parameter defines an ASCII name for the MSTP configuration. The default name is for the name variable to be blank.

Setting the MSTP revision number

Each switch that is running MSTP is configured with a revision number. It applies to the switch which can have many different VLANs that can belong to many different MSTP regions.

To configure an MSTP revision number, use a command such as the following at the Global Configuration level.

```
device(config)#mstp revision 4
```

Syntax: [no] mstp revision *revision-number*

The **revision** parameter specifies the revision level for MSTP that you are configuring on the switch. It can be a number from 0 and 65535. The default revision number is 0.

Configuring an MSTP instance

An MSTP instance is configured with an MSTP ID for each region. Each region can contain one or more VLANs. The Brocade implementation of MSTP allows you to assign VLANs or ranges of VLANs to an MSTP instance before or after they have been defined. If pre-defined, a VLAN will be placed in the MSTI that it was assigned to immediately when the VLAN is created. Otherwise, the default operation is to condition of assign all new VLANs to the CIST. VLANs assigned to the CIST by default can be moved later to a specified MSTI.

To configure an MSTP instance and map one or more VLANs to that MSTI, use a command such as the following at the Global Configuration level.

```
device(config)#mstp instance 7 vlan 4 to 7
```

Syntax: [no] **mstp instance instance-number [vlan vlan-id| vlan-group group-id]**

The **instance** parameter defines the number for the instance of MSTP that you are configuring. The value 0 (which identifies the CIST) cannot be used. You can have up to 15 instances, number 1 - 4094.

The **vlan** parameter assigns one or more VLANs or a range of VLANs to the instance defined in this command.

The **vlan-group** parameter assigns one or more VLAN groups to the instance defined in this command.

The **no** option moves a VLAN or VLAN group from its assigned MSTI back into the CIST.

NOTE

The system does not allow an MSTI without any VLANs mapped to it. Consequently, removing all VLANs from an MSTI, deletes the MSTI from the system. The CIST by contrast will exist regardless of whether or not any VLANs are assigned to it or not. Consequently, if all VLANs are moved out of a CIST, the CIST will still exist and functional.

Configuring bridge priority for an MSTP instance

Priority can be configured for a specified instance. To configure priority for an MSTP instance, use a command such as the following at the Global Configuration level.

```
device(config)#mstp instance 1 priority 8192
```

Syntax: [no] **mstp instance instance-number priority priority-value**

The *instance-number* variable is the number for the instance of MSTP that you are configuring.

You can set a **priority** to the instance that gives it forwarding preference over lower priority instances within a VLAN or on the switch. A higher number for the priority variable means a lower forwarding priority. Acceptable values are 0 - 61440 in increments of 4096. The default value is 32768.

Setting the MSTP global parameters

MSTP has many of the options available in RSTP as well as some unique options. To configure MSTP Global parameters for all instances on a switch.

```
device(config)#mstp force-version 0 forward-delay 10 hello-time 4 max-age 12 max-hops  
9
```

Syntax: [no] mstp force-version *mode-number* forward-delay *value* hello-time *value* max-age *value* max-hops *value*

The **force-version** parameter forces the bridge to send BPDUs in a specific format. You can specify one of the following mode-number values:

- 0 - The STP compatibility mode. Only STP BPDUs will be sent. This is equivalent to single STP.
- 2 - The RSTP compatibility mode. Only RSTP BPDUS will be sent. This is equivalent to single STP.
- 3 - MSTP mode. In this default mode, only MSTP BPDUS will be sent.

The **forward-delay value** specifies how long a port waits before it forwards an RST BPDU after a topology change. This can be a value from 4 - 30 seconds. The default is 15 seconds.

The **hello-time value** parameter specifies the interval between two hello packets. The parameter can have a value from 1 - 10 seconds. The default is 2 seconds.

The **max-age value** parameter specifies the amount of time the device waits to receive a hello packet before it initiates a topology change. You can specify a value from 6 - 40 seconds, where the value adheres to the following formula.

max age equal to or greater than $2 \times (\text{hello-time} + 1)$ AND max age equal to or greater than $2 \times (\text{forward-delay} - 1)$

The default max-age is 20 seconds.

The **max-hops value** parameter specifies the maximum hop count. You can specify a value from 1 - 40 hops. The default value is 20 hops.

Setting ports to be operational edge ports

You can define specific ports as edge ports for the region in which they are configured to connect to devices (such as a host) that are not running STP, RSTP, or MSTP. If a port is connected to an end device such as a PC, the port can be configured as an edge port. To configure ports as operational edge ports enter a command such as the following.

```
device(config)#mstp admin-edge-port ethernet 3/1
```

Syntax: [no] mstp admin-edge-port **ether**net *port*

Setting automatic operational edge ports

You can configure a Layer 3 switch to automatically set a port as an operational edge port if the port does not receive any BPDUs since link-up. If the port receives a BPDU later, it is automatically reset to become an operational non-edge port. This feature is set globally to apply to all ports on a router where it is configured. This feature is configured as shown in the following.

```
device(config)#mstp edge-port-auto-detect
```

Syntax: [no] mstp edge-port-auto-detect

NOTE

If this feature is enabled, it takes the port about 3 seconds longer to come to the enable state.

Setting point-to-point link

You can set a point-to-point link between ports to increase the speed of convergence. To create a point-to-point link between ports, use a command such as the following at the Global Configuration level.

```
device(config)#mstp admin-pt2pt-mac ethernet 2/5 ethernet 4/5
```

Syntax: [no] mstp admin-pt2pt-mac **ethernet port**

MSTP+ overview

The MSTP+ feature allows you to selectively include VLANs in the MSTP scope.

In the standard IEEE 802.1s MSTP all VLANs are automatically placed under CIST control so that the entire switch is controlled by the MSTP. The MSTP+ feature is an enhancement that allows you to exclude one or more VLANs from the MSTP scope and configure them in a non-MSTP topology. These VLANs are considered free VLANs and can run any Layer 2 protocols such as PVST/PVRST, MRP, VSRP, or any pure Layer 3 protocols.

You must ensure all the connected devices are properly configured, create the MSTP instances, and assign the VLANs to those instances. These instances must be configured on all devices that interoperate with the same VLAN assignments.

The following table lists the protocols that can run under free VLANs along with the MSTP+.

Protocol	Compatible with MSTP+
Spanning tree single (802.1D)	No
Rapid spanning tree single (802.1w)	No
MCT	No
Per-VLAN spanning tree (STP, RSTP)	Yes
VSRP	Yes
MRP	Yes
All Layer 3 protocols (pure Layer 3 network)	Yes

This means that you can create an independent Layer 3 topology even when on a switch that is configured with MSTP. The MSTP convergence does not affect the Layer 3 topology.

You can switch between non-MSTP, MSTP, and MSTP+ modes.

NOTE

Systems configured with MSTP+ may not interoperate properly with the systems on which standard MSTP is configured. It is recommended that you configure MSTP+ on both sides.

NOTE

Free VLANs must have their own means to break Layer 2 loops; MSTP+ cannot be relied on to do so.

Configuring MSTP+

Use the **mstp scope** command with the **pvst** keyword to configure MSTP+.

MSTP+ is not operational until you configure at least one MSTP instance and configure the **mstp start** command. You can create MSTP+ instances the same way you configure MSTP instances. See the “Configuring an MSTP instance” section for information on configuring MSTP.

1. Configure MSTP+.

```
Device(config)# mstp scope pvst
Enabling MSTP+ scope. MSTP instances need to be configured and 'mstp start'
need to be entered in order to activate this MSTP+ feature
Configures MSTP+. CIST is not automatically created and VLANs are not under MSTP scope
unless you explicitly configure the MSTP instances and attach the VLANs to them.
```

2. Create an MSTP instance.

```
Device(config)# mstp instance 1 vlan 4 to 7
Creates an MSTP instance on VLANs 4 to 7.
```

3. Start the MSTP+ protocol.

```
Device(config)# mstp start
Creates an MSTP instance on VLANs 4 to 7.
```

4. Remove the MSTP+ configuration.

```
Device(config)# no mstp scope pvst
Removes the MSTP+ configuration. The VLANs that were attached to MSTP+ are out of MSTP+
scope and there is no PVST under those VLANs. The non-MSTP VLANs are not affected.
```

Switching between non-MSTP, MSTP, and MSTP+ modes

Use the **mstp scope** command to switch between non-MSTP, MSTP, and MSTP+ modes. This allows you to move between modes without explicitly removing the current mode and reconfiguring the new mode.

When an MSTP instance is enabled, you can configure the **pvst** and **all** keywords to switch between modes.

1. When the **mstp scope all** command is configured and MSTP mode is active, change to MSTP+ mode.

```
Device(config)# mstp scope pvst
The mode is changed to MSTP+. You can remove the VLANs from MSTP+ instances. VLANs that
are removed from MSTP+ scope become free and other supported protocols can be configured.
```

2. When the **mstp scope pvst** command is configured and MSTP+ mode is active, change to MSTP mode.

```
Device(config)# mstp scope all
The mode is changed to MSTP. The VLANs that are already attached to MSTP+ are kept as is and
all the free VLANs are attached to a CIST instance. Any protocols configured under the free VLANs
are removed.
```

Disabling MSTP on a port

To disable MSTP on a specific port, use a command such as the following at the Global Configuration level.

```
device(config)#mstp disable ethernet 2/1
```

Syntax: [no] **mstp disable ethernet port**

When a port is disabled for MSTP, it behaves as blocking for all the VLAN traffic that is controlled by MSTIs and the CIST.

Forcing ports to transmit an MSTP BPDU

To force a port to transmit an MSTP BPDU, use a command such as the following at the Global Configuration level.

```
device(config)#mstp force-migration-check ethernet 3/1
```

Syntax: [no] mstp force-migration-check **etherent port**

Activating MSTP on a switch

MSTP scope must be enabled on the switch as described in [Configuring MSTP mode and scope](#) on page 321 before MSTP can be enabled.

To enable MSTP on your switch, use the following at the Global Configuration level.

```
device(config)#mstp start
```

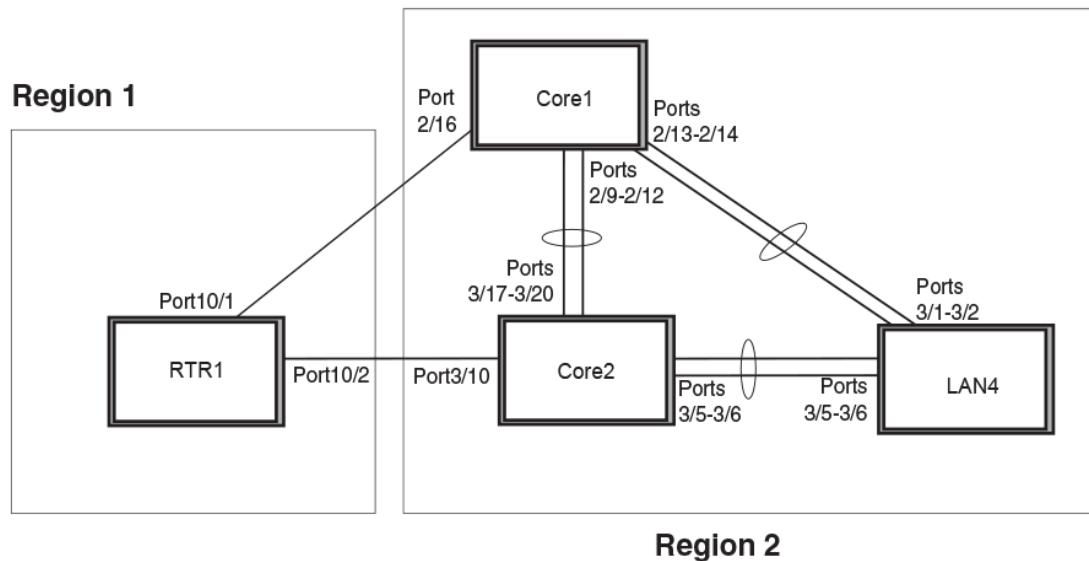
Syntax: [no] mstp start

The **no** option disables MSTP from operating on a switch.

Examples of an MSTP configuration

In the following figure, four Brocade device routers are configured in two regions. There are four VLANs in four instances in Region 2. Region 1 is in the CIST.

FIGURE 68 Sample MSTP configuration



RTR1 on MSTP configuration

```
device(config-vlan-4093)#tagged ethernet 10/1 to 10/2
device(config-vlan-4093)#exit
device(config)#mstp scope all
device(config)#mstp name Reg1
device(config)#mstp revision 1
device(config)#mstp admin-pt2pt-mac ethernet 10/1 to 10/2
device(config)#mstp start
device(config)#hostname RTR1
```

Core 1 on MSTP configuration

```

device(config)#trunk ethernet 2/9 to 2/12 ethernet 2/13 to 2/14
device(config-vlan-1)#name DEFAULT-VLAN by port
device(config-vlan-1)#exit
device(config)#vlan 20 by port
device(config-vlan-20)#tagged ethernet 2/9 to 2/14 ethernet 2/16
device(config-vlan-20)#exit
device(config)#vlan 21 by port
device(config-vlan-21)#tagged ethernet 2/9 to 2/14 ethernet 2/16
device(config-vlan-21)#exit
device(config)#vlan 22 by port
device(config-vlan-22)#tagged ethernet 2/9 to 2/14 ethernet 2/16
device(config-vlan-22)#exit
device(config)#vlan 23 by port
device(config)#mstp scope all
device(config)#mstp name HR
device(config)#mstp revision 2
device(config)#mstp instance 20 vlan 20
device(config)#mstp instance 21 vlan 21
device(config)#mstp instance 22 vlan 22
device(config)#mstp instance 0 priority 8192
device(config)#mstp admin-pt2pt-mac ethernet 2/9 to 2/14
device(config)#mstp admin-pt2pt-mac ethernet 2/16
device(config)#mstp disable ethernet 2/240.
device(config)#mstp start
device(config)#hostname CORE1

```

Core2 on MSTP configuration

```

device(config)#trunk ethernet 3/5 to 3/6 ethernet 3/17 to 3/20
device(config)#vlan 1 name DEFAULT-VLAN by port
device(config-vlan-1)#exit
device(config)#vlan 20 by port
device(config-vlan-20)#tagged ethernet 3/5 to 3/6 ethernet 3/17 to 3/20
device(config-vlan-20)#exit
device(config)#vlan 21 by port
device(config-vlan-21)#tagged ethernet 3/5 to 3/6 ethernet 3/17 to 3/20
device(config-vlan-21)#exit
device(config)#vlan 22 by port
device(config-vlan-22)#tagged ethernet 3/5 to 3/6 ethernet 3/17 to 3/20
device(config-vlan-22)#exit
device(config)#mstp scope all
device(config)#mstp name HR
device(config)#mstp revision 2
device(config)#mstp instance 20 vlan 20
device(config)#mstp instance 21 vlan 21
device(config)#mstp instance 22 vlan 22
device(config)#mstp admin-pt2pt-mac ethernet 3/17 to 3/20 ethernet 3/5 to 3/6
device(config)#mstp admin-pt2pt-mac ethernet 3/10
device(config)#mstp disable ethernet 3/7 ethernet 3/24
device(config)#mstp start
device(config)#hostname CORE2

```

LAN 4 on MSTP configuration

```

device(config)#trunk ethernet 3/5 to 3/6 ethernet 3/1 to 3/2
device(config)#vlan 1 name DEFAULT-VLAN by port
device(config-vlan-1)#exit
device(config)#vlan 20 by port
device(config-vlan-20)#tagged ethernet 3/1 to 3/2 ethernet 3/5 to 3/6
device(config)#exit
device(config)#vlan 21 by port
device(config-vlan-21)#tagged ethernet 3/1 to 3/2 ethernet 3/5 to 3/6
device(config-vlan-21)#exit
device(config)#vlan 22 by port
device(config-vlan-22)#tagged ethernet 3/1 to 3/2 ethernet 3/5 to 3/6
device(config)#mstp scope all
device(config)#mstp config name HR
device(config)#mstp revision 2
device(config)#mstp instance 20 vlan 20

```

```
device(config)#mstp instance 21 vlan 21
device(config)#mstp instance 22 vlan 22
device(config)#mstp admin-pt2pt-mac ethernet 3/5 to 3/6 ethernet 3/1 to 3/2
device(config)#mstp start
device(config)#hostname LAN4
```

Displaying MSTP statistics

MSTP statistics can be displayed using the commands shown below.

To display all general MSTP information, enter the following command.

```
device#show mstp
MSTP Instance 0 (CIST) - VLANs: 1
-----
Bridge      Bridge Bridge Bridge Bridge Root   Root   Root
Identifier  MaxAge Hello FwdDly Hop    MaxAge Hello FwdDly Hop
hex         sec   sec   sec   cnt   sec   sec   sec   cnt
8000000cdb80af01 20   2   15   20   20   2   15   19
Root        ExtPath RegionalRoot   IntPath Designated   Root
Bridge      Cost     Bridge       Cost     Bridge   Port
hex         hex
8000000480bb9876 2000   8000000cdb80af01 0   8000000480bb9876 3/1
Port Pri PortPath P2P Edge Role   State   Designat- Designat-
Num  Cost   Mac Port      Cost   ed cost  bridge
3/1  128 2000   T   F   ROOT   FORWARDING 0   8000000480bb9876
MSTP Instance 1 - VLANs: 2
-----
Bridge      Max RegionalRoot   IntPath Designated   Root   Root
Identifier  Hop   Bridge       Cost     Bridge   Port   Hop
hex         cnt hex
8001000cdb80af01 20   8001000cdb80af01 0   8001000cdb80af01 Root  20
Port Pri PortPath Role   State   Designat- Designat-
Num  Cost   Mac Port      Cost   ed cost  bridge
3/1  128 2000   MASTER   FORWARDING 0   8001000cdb80af01
```

Syntax: show mstp instance-number

The instance-number variable specifies the MSTP instance that you want to display information for.

TABLE 52 Output from Show MSTP

Field	Description
MSTP Instance	The ID of the MSTP instance whose statistics are being displayed. For the CIST, this number is 0.
VLANs	The number of VLANs that are included in this instance of MSTP. For the CIST this number will always be 1.
Bridge Identifier	The MAC address of the bridge.
Bridge MaxAge sec	Displays configured Max Age.
Bridge Hello sec	Displays configured Hello variable.
Bridge FwdDly sec	Displays configured FwdDly variable.
Bridge Hop cnt	Displays configured Max Hop count variable.
Root MaxAge sec	Max Age configured on the root bridge.

TABLE 52 Output from Show MSTP (Continued)

Field	Description
Root Hello sec	Hello interval configured on the root bridge.
Root FwdDly sec	FwdDly interval configured on the root bridge.
Root Hop Cnt	Maximum hop count left from the root bridge.
Root Bridge	Bridge identifier of the root bridge.
ExtPath Cost	The configured path cost on a link connected to this port to an external MSTP region.
Regional Root Bridge	The Regional Root Bridge is the MAC address of the Root Bridge for the local region.
IntPath Cost	The configured path cost on a link connected to this port within the internal MSTP region.
Designated Bridge	The MAC address of the bridge that sent the best BPDU that was received on this port.
Root Port	Port indicating shortest path to root. Set to "Root" if this bridge is the root bridge.
Port Num	The port number of the interface.
Pri	The configured priority of the port. The default is 128.
PortPath Cost	Configured or auto detected path cost for port.
P2P Mac	Indicates if the port is configured with a point-to-point link: <ul style="list-style-type: none"> • T - The port is configured in a point-to-point link • F - The port is not configured in a point-to-point link
Edge	Indicates if the port is configured as an operational edge port: <ul style="list-style-type: none"> • T - indicates that the port is defined as an edge port. • F - indicates that the port is not defined as an edge port
Role	The current role of the port: <ul style="list-style-type: none"> • Master • Root • Designated • Alternate • Backup • Disabled
State	The port current spanning tree state. A port can have one of the following states: <ul style="list-style-type: none"> • Forwarding • Discarding • Learning • Disabled

TABLE 52 Output from Show MSTP (Continued)

Field	Description
Designated Cost	Port path cost to the root bridge.
Max Hop cnt	The maximum hop count configured for this instance.

Displaying MSTP information for a specified instance

The following example displays MSTP information specified for an MSTP instance.

```
device#show mstp 1
MSTP Instance 1 - VLANs: 2
-----
Bridge      Max RegionalRoot      IntPath      Designated      Root      Root
Identifier   Hop Bridge          Cost         Bridge          Port      Hop
hex         cnt hex            hex          hex           cnt
8001000cdb80af01 20 8001000cdb80af01 0      8001000cdb80af01 Root    20
Port       Pri PortPath        Role        State        Designat- Designated
Num        Cost             Mac Port      FORWARDING  ed cost  bridge
3/1       128 2000          T F ROOT      0          8001000cdb80af01
```

Refer to [Table 52](#) on page 331 for details about the display parameters.

Displaying MSTP information for CIST instance 0

Instance 0 is the Common and Internal Spanning Tree Instance (CIST). When you display information for this instance there are some differences with displaying other instances. The following example displays MSTP information for CIST Instance 0.

```
device#show mstp 0
MSTP Instance 0 (CIST) - VLANs: 1
-----
Bridge      Bridge Bridge Bridge Bridge Root      Root      Root
Identifier   MaxAge Hello FwdDly Hop      MaxAge Hello FwdDly Hop
hex         sec   sec   sec   sec   sec   sec   sec   sec   cnt
8000000cdb80af01 20 2 15 20 20 2 15 19
Root       ExtPath  RegionalRoot      IntPath      Designated      Root
Bridge      Cost     Bridge          Cost         Bridge          Port
hex         hex      hex          hex          hex
8000000480bb9876 2000 8000000cdb80af01 0      8000000480bb9876 3/1
Port       Pri PortPath        P2P Edge Role      State        Designat- Designated
Num        Cost             Mac Port      FORWARDING  ed cost  bridge
3/1       128 2000          T F ROOT      0          8000000480bb9876
```

To display details about the MSTP configuration, enter the following command.

```
device#show mstp conf
MSTP CONFIGURATION
-----
Name      : Reg1
Revision : 1
Version  : 3 (MSTP mode)
Status   : Started
Instance VLANs
-----
0        4093
```

To display details about the MSTP that is configured on the device, enter the following command.

```
device#show mstp detail
MSTP Instance 0 (CIST) - VLANs: 4093
-----
Bridge: 800000b000c00000 [Priority 32768, SysId 0, Mac 00b000c00000]
```

```
FwdDelay 15, HelloTime 2, MaxHops 20, TxHoldCount 6
Port 6/54 - Role: DESIGNATED - State: FORWARDING
PathCost 20000, Priority 128, OperEdge T, OperPt2PtMac F, Boundary T
Designated - Root 800000b000c00000, RegionalRoot 800000b000c00000,
Bridge 800000b000c00000, ExtCost 0, IntCost 0
ActiveTimers - helloWhen 1
MachineState - PRX-DISCARD, PTX-IDLE, PPM-SENDING_RSTP, PIM-CURRENT
PRT-ACTIVE_PORT, PST-FORWARDING, TCM-INACTIVE
BPDUs - Rcvd MST 0, RST 0, Config 0, TCN 0
Sent MST 6, RST 0, Config 0, TCN 0
```

Refer to [Table 52](#) on page 331 for explanation about the parameters in the output.

Syntax: **show mstp [mstp-id | configuration | detail] [begin string | exclude string | include string]**

Enter an MSTP ID for *mstp-id*.

VLANs

● VLAN overview	335
● Routing between VLANs	360
● Configuring IP subnet, IPX network and protocol-based VLANs	365
● IP subnet, IPX network, and protocol-based VLANs within port-based VLANs	366
● IPv6 protocol VLAN configuration	369
● Routing between VLANs using virtual routing interfaces (Layer 3 Switches only)	370
● Configuring protocol VLANs with dynamic ports	375
● Configuring uplink ports within a port-based VLAN	379
● IP subnet address on multiple port-based VLAN configuration	379
● VLAN groups and virtual routing interface group	383
● Super-aggregated VLAN configuration	388
● 802.1ad tagging configuration	395
● Private VLAN configuration	399
● Dual-mode VLAN ports	409
● Displaying VLAN information	412

VLAN overview

The following sections provide details about the VLAN types and features supported on the FastIron family of switches.

Types of VLANs

This section describes the VLAN types supported on Brocade devices.

VLAN support on FastIron devices

You can configure the following types of VLANs on FastIron devices:

- Layer 2 port-based VLAN - a set of physical ports that share a common, exclusive Layer 2 broadcast domain
- Layer 3 protocol VLANs - a subset of ports within a port-based VLAN that share a common, exclusive broadcast domain for Layer 3 broadcasts of the specified protocol type
- IP subnet VLANs - a subset of ports in a port-based VLAN that share a common, exclusive subnet broadcast domain for a specified IP subnet
- IPv6 VLANs - a subset of ports in a port-based VLAN that share a common, exclusive network broadcast domain for IPv6 packets
- IPX network VLANs - a subset of ports in a port-based VLAN that share a common, exclusive network broadcast domain for a specified IPX network
- AppleTalk cable VLANs - a subset of ports in a port-based-based VLAN that share a common, exclusive network broadcast domain for a specified AppleTalk cable range

NOTE

IPX and AppleTalk VLANs are not present in the ICX 7xxx series devices.

When a FastIron device receives a packet on a port that is a member of a VLAN, the device forwards the packet based on the following VLAN hierarchy:

- If the port belongs to an IP subnet VLAN, IPX network VLAN, or AppleTalk cable VLAN and the packet belongs to the corresponding IP subnet, IPX network, or AppleTalk cable range, the device forwards the packet to all the ports within that VLAN.
- If the packet is a Layer 3 packet but cannot be forwarded as described above, but the port is a member of a Layer 3 protocol VLAN for the packet protocol, the device forwards the packet on all the Layer 3 protocol VLAN ports.
- If the packet cannot be forwarded based on either of the VLAN membership types listed above, but the packet can be forwarded at Layer 2, the device forwards the packet on all the ports within the receiving port port-based VLAN.

Protocol VLANs differ from IP subnet, IPX network, and AppleTalk VLANs in an important way.

Protocol VLANs accept any broadcast of the specified protocol type. An IP subnet, IPX network, or AppleTalk VLAN accepts only broadcasts for the specified IP subnet, IPX network, or AppleTalk cable range.

NOTE

Protocol VLANs are different from IP subnet, IPX network, and AppleTalk cable VLANs. A port-based VLAN cannot contain both an IP subnet, IPX network, or AppleTalk cable VLAN and a protocol VLAN for the same protocol. For example, a port-based VLAN cannot contain both an IP protocol VLAN and an IP subnet VLAN.

Layer 2 port-based VLANs

On all Brocade devices, you can configure port-based VLANs. A port-based VLAN is a subset of ports on a Brocade device that constitutes a Layer 2 broadcast domain.

By default, all the ports on a Brocade device are members of the default VLAN. Thus, all the ports on the device constitute a single Layer 2 broadcast domain. When you configure a port-based VLAN, the device automatically removes the ports you add to the VLAN from the default VLAN.

You can configure multiple port-based VLANs. You can configure up to 4094 port-based VLANs on a Layer 2 Switch or Layer 3 Switch. On both device types, valid VLAN IDs are 1 - 4095. You can configure up to the maximum number of VLANs within that ID range.

NOTE

VLAN IDs 4087, 4090, and 4093 are reserved for Brocade internal use only. VLAN 4094 is reserved for use by Single STP. If you want to use VLANs 4091 and 4092 as configurable VLANs, you can assign them to different VLAN IDs. For more information, refer to [Assigning different VLAN IDs to reserved VLANs 4091 and 4092](#) on page 362

NOTE

Each port-based VLAN can contain either tagged or untagged ports. A port cannot be a member of more than one port-based VLAN unless the port is tagged. 802.1Q tagging allows the port to add a four-byte tag field, which contains the VLAN ID, to each packet sent on the port. You also can configure port-based VLANs that span multiple devices by tagging the ports within the VLAN. The tag

enables each device that receives the packet to determine the VLAN the packet belongs to. 802.1Q tagging applies only to Layer 2 VLANs, not to Layer 3 VLANs.

Because each port-based VLAN is a separate Layer 2 broadcast domain, each VLAN can be configured to run a separate instance of the Spanning Tree Protocol (STP).

Layer 2 traffic is bridged within a port-based VLAN and Layer 2 broadcasts are sent to all the ports within the VLAN.

The following figure shows an example of a Brocade device on which a Layer 2 port-based VLAN has been configured.

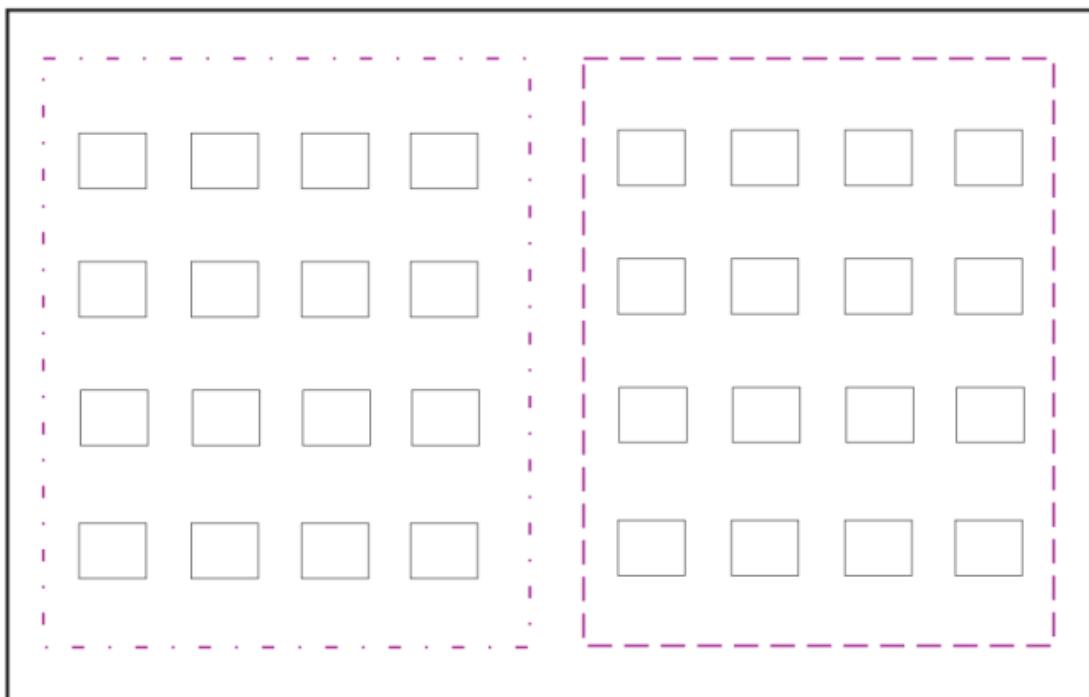
FIGURE 69 Brocade device containing user-defined Layer 2 port-based VLAN

DEFAULT-VLAN

VLAN ID = 1

Layer 2 Port-based VLAN

User-configured port-based VLAN



When you add a port-based VLAN,
the device removes all the ports in the
new VLAN from DEFAULT-VLAN.

Configuring port-based VLANs

Port-based VLANs allow you to provide separate spanning tree protocol (STP) domains or broadcast domains on a port-by-port basis.

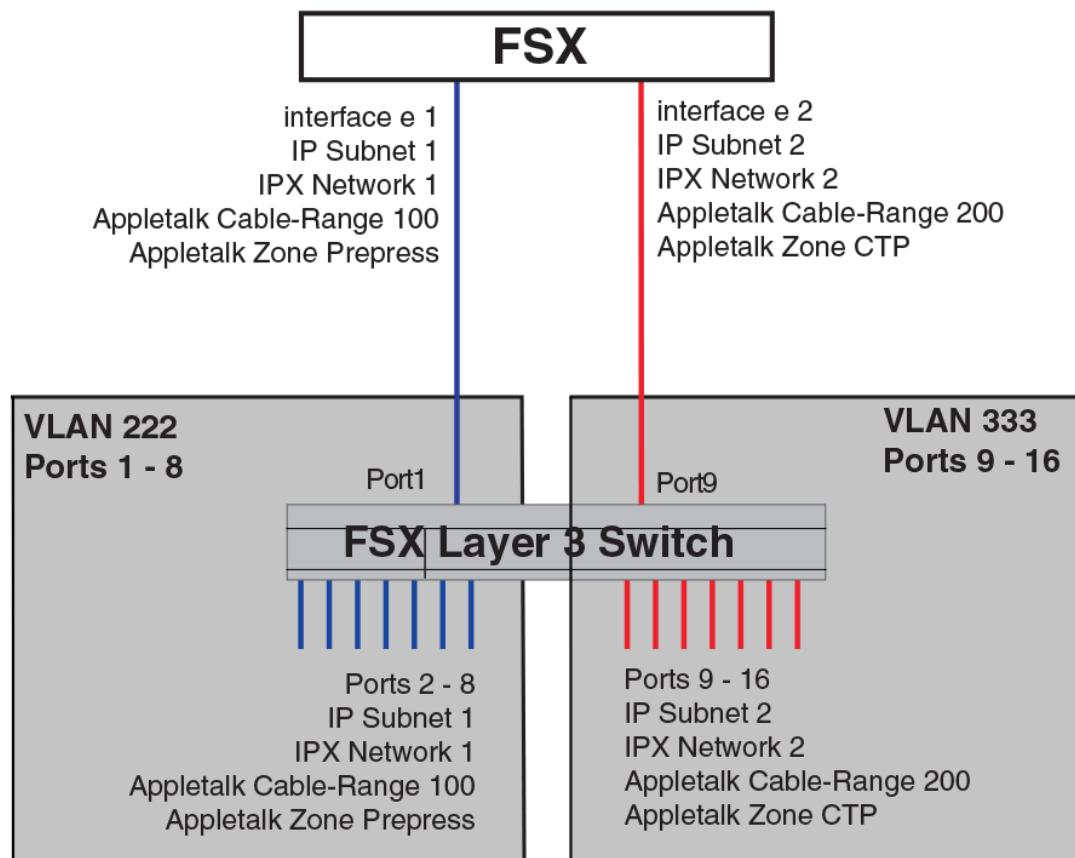
This section describes how to perform the following tasks for port-based VLANs using the CLI:

- Create a VLAN
- Delete a VLAN
- Modify a VLAN
- Change a VLAN priority
- Enable or disable STP on the VLAN

1--Simple port-based VLAN configuration

The following figure shows a simple port-based VLAN configuration using a single Brocade Layer 2 Switch. All ports within each VLAN are untagged. One untagged port within each VLAN is used to connect the Layer 2 Switch to a Layer 3 Switch (in this example, a FSX) for Layer 3 connectivity between the two port-based VLANs.

FIGURE 70 Port-based VLANs 222 and 333



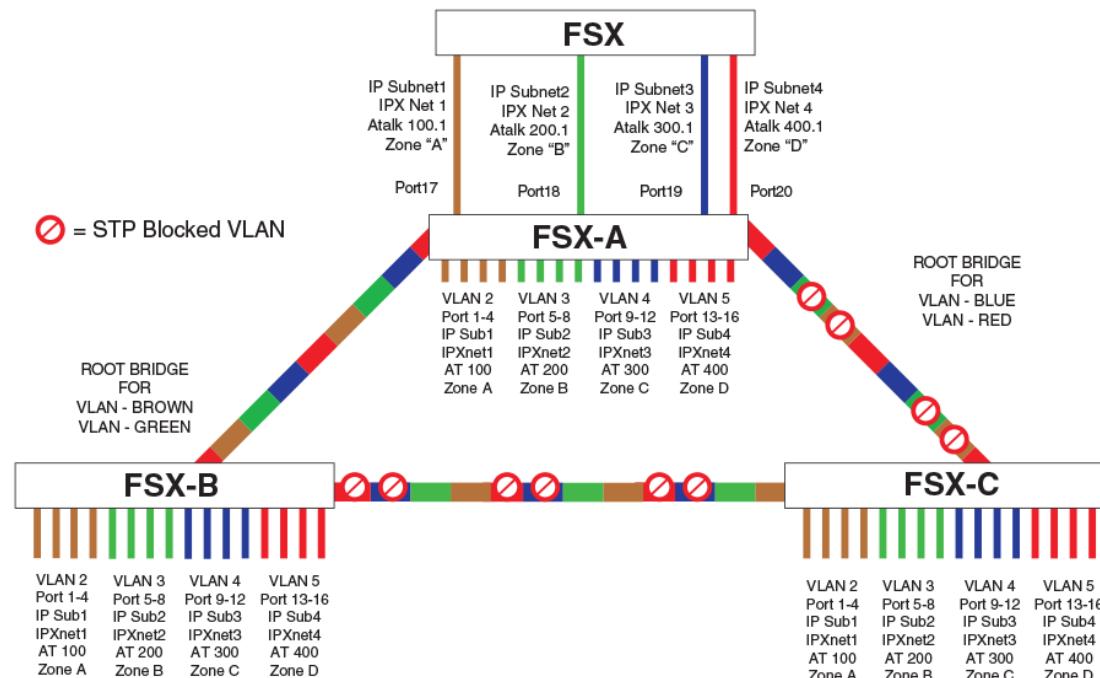
To create the two port-based VLANs shown in the above figure, enter the following commands.

```
device(config)# vlan 222 by port
device(config-vlan-222)# untagged ethernet 1 to 8
device(config-vlan-222)# vlan 333 by port
device(config-vlan-333)# untagged ethernet 9 to 16
```

Syntax: `vlan vlan-id by port`**Syntax: `untagged ethernet [slotnum/]portnum [to [slotnum/]portnum | ethernet [slotnum/]portnum]`**

2--More complex port-based VLAN configuration

The following figure shows a more complex port-based VLAN configuration using multiple Layer 2 Switches and IEEE 802.1Q VLAN tagging. The backbone link connecting the three Layer 2 Switches is tagged. One untagged port within each port-based VLAN on FSX-A connects each separate network wide Layer 2 broadcast domain to the router for Layer 3 forwarding between broadcast domains. The STP priority is configured to force FSX-A to be the root bridge for VLANs RED and BLUE. The STP priority on FSX-B is configured so that FSX-B is the root bridge for VLANs GREEN and BROWN.

FIGURE 71 More complex port-based VLAN

To configure the Port-based VLANs on the FSX Layer 2 Switches in the above figure, use the following method.

Configuring port-based VLANs on FSX-A

Enter the following commands to configure FSX-A.

```

device> enable
device# configure terminal
device(config)# hostname FSX-A
device-A(config)# vlan 2 name BROWN
device-A(config-vlan-2)# untagged ethernet 1 to 4 ethernet 17
device-A(config-vlan-2)# tagged ethernet 25 to 26
device-A(config-vlan-2)# spanning-tree
device-A(config-vlan-2)# vlan 3 name GREEN
device-A(config-vlan-3)# untagged ethernet 5 to 8 ethernet 18
device-A(config-vlan-3)# tagged ethernet 25 to 26
device-A(config-vlan-3)# spanning-tree
device-A(config-vlan-3)# vlan 4 name BLUE
device-A(config-vlan-4)# untagged ethernet 9 to 12 ethernet 19
device-A(config-vlan-4)# tagged ethernet 25 to 26
device-A(config-vlan-4)# spanning-tree

```

```
device-A(config-vlan-4)# spanning-tree priority 500
device-A(config-vlan-4)# vlan 5 name RED
device-A(config-vlan-5)# untagged ethernet 13 to 16 ethernet 20
device-A(config-vlan-5)# tagged ethernet 25 to 26
device-A(config-vlan-5)# spanning-tree
device-A(config-vlan-5)# spanning-tree priority 500
device-A(config-vlan-5)# end
device-A# write memory
```

Configuring port-based VLANs on FSX-B

Enter the following commands to configure FSX-B.

```
device> en
device# configure terminal
device(config)# hostname FSX-B
device-B(config)# vlan 2 name BROWN
device-B(config-vlan-2)# untagged ethernet 1 to 4
device-B(config-vlan-2)# tagged ethernet 25 to 26
device-B(config-vlan-2)# spanning-tree
device-B(config-vlan-2)# spanning-tree priority 500
device-B(config-vlan-2)# vlan 3 name GREEN
device-B(config-vlan-3)# untagged ethernet 5 to 8
device-B(config-vlan-3)# tagged ethernet 25 to 26
device-B(config-vlan-3)# spanning-tree
device-B(config-vlan-3)# spanning-tree priority 500
device-B(config-vlan-3)# vlan 4 name BLUE
device-B(config-vlan-4)# untagged ethernet 9 to 12
device-B(config-vlan-4)# tagged ethernet 25 to 26
device-B(config-vlan-4)# vlan 5 name RED
device-B(config-vlan-5)# untagged ethernet 13 to 16
device-B(config-vlan-5)# tagged ethernet 25 to 26
device-B(config-vlan-5)# end
device-B# write memory
```

Configuring port-based VLANs on FSX-C

Enter the following commands to configure FSX-C.

```
device> en
device# configure terminal
device(config)# hostname FSX-C
device-C(config)# vlan 2 name BROWN
device-C(config-vlan-2)# untagged ethernet 1 to 4
device-C(config-vlan-2)# tagged ethernet 25 to 26
device-C(config-vlan-2)# vlan 3 name GREEN
device-C(config-vlan-3)# untagged ethernet 5 to 8
device-C(config-vlan-3)# tagged ethernet 25 to 26
device-C(config-vlan-3)# vlan 4 name BLUE
device-C(config-vlan-4)# untagged ethernet 9 to 12
device-C(config-vlan-4)# tagged ethernet 25 to 26
device-C(config-vlan-4)# vlan 5 name RED
device-C(config-vlan-5)# untagged ethernet 13 to 16
device-C(config-vlan-5)# tagged ethernet 25 to 26
device-C(config-vlan-5)# end
device-C# write memory
```

Syntax: *vlan vlan-id by port*

Syntax: *untagged ethernet [slotnum/]portnum [to [slotnum/]portnum | ethernet [slotnum/]portnum]*

Syntax: *tagged ethernet [slotnum/]portnum [to <[slotnum/]portnum> | ethernet [slotnum/]portnum]*

Syntax: *[no] spanning-tree*

Syntax: *spanning-tree [ethernet [slotnum/]portnum path-cost value priority value] forward-delay value hello-time value maximum-age time priority value*

Modifying a port-based VLAN

You can make the following modifications to a port-based VLAN:

- Add or delete a VLAN port.
 - Enable or disable STP.

Removing a port-based VLAN

Suppose you want to remove VLAN 5 from the example in [Figure 71](#) on page 339. To do so, use the following procedure.

1. Access the global CONFIG level of the CLI on FSX-A by entering the following commands.

```
device-A> enable
```

No password has been assigned yet...

2. Enter the following command.

```
device-A(config)# no vlan 5  
device-A(config)#
```

3. Enter the following commands to exit the CONFIG level and save the configuration to the system-config file on flash memory.

```
device-A(config)#  
device-A(config)# end  
device-A# write memory  
device-A#
```

4. Repeat steps 1 - 3 on FSX-B.

Syntax: [no] **vlan** *vlan-id* **by port**

Removing a port from a VLAN

Suppose you want to remove port 11 from VLAN 4 on FSX-A shown in [Figure 71](#) on page 339. To do so, use the following procedure.

1. Access the global CONFIG level of the CLI on FSX-A by entering the following command.

```
device-A> enable  
No password has been assigned yet...
```

```
device-A# configure terminal  
device-A(config)#
```

2. Access the level of the CLI for configuring port-based VLAN 4 by entering the following command:

```
device-A(config)#
device-A(config)# vlan 4
device-A(config-vlan-4)#
```

3. Enter the following commands.

```
device-A(config-vlan-4)#
device-A(config-vlan-4)# no untagged ethernet 11
deleted port ethe 11 from port-vlan 4.
device-A(config-vlan-4)#
device-A(config-vlan-4)#
```

4. Enter the following commands to exit the VLAN CONFIG mode and save the configuration to the system-config file on flash memory.

```
device-A(config-vlan-4)#  
device-A(config-vlan-4)# end  
device-A# write memory
```

You can remove all the ports from a port-based VLAN without losing the rest of the VLAN configuration. However, you cannot configure an IP address on a virtual routing interface unless the VLAN contains ports. If the VLAN has a virtual routing interface, the virtual routing interface IP address is deleted when the ports associated with the interface are deleted. The rest of the VLAN configuration is retained.

Multi-range VLAN

The multi-range VLAN feature allows users to use a single command to create and configure multiple VLANs. These VLANs can be continuous, for example from 2 to 7 or discontinuous, for example, 2 4 7.

NOTE

The maximum number of VLANs you can create or configure with a single command is 64.

Creating a multi-range VLAN

To create more than one VLAN with a single command, you can specify the VLAN number and range.

Syntax: **[no] vlan num to num**

The *num* parameter specifies the VLAN ID.

To create a continuous range of VLANs, enter command such as the following.

```
device(config)#vlan 2 to 7
device(config-mvlan-2-7) #
```

Syntax: **[no] vlan num to num**

To create discontinuous VLANs, enter command such as the following.

```
device(config)#vlan 2 4 7
device(config-mvlan-2*7) #exit
```

Syntax: **[no] vlan num num num**

You can also create continuous and discontinuous VLANs. To create continuous and discontinuous VLANs, enter command such as the following.

```
device(config)#vlan 2 to 7 20 25
device(config-mvlan-2*25) #
```

Syntax: **[no] vlan num to num num**

Deleting a multi-range VLAN

You can also delete multiple VLANs with a single command.

To delete a continuous range of VLANs, enter command such as the following.

```
device(config)#no vlan 2 to 7
```

Syntax: **[no] vlan num to num**

To delete discontinuous VLANs, enter command such as the following.

```
device(config)#no vlan 2 4 7
```

Syntax: **[no] vlan num num num**

You can also delete continuous and discontinuous VLANs. To delete continuous and discontinuous VLANs, enter command such as the following.

```
device(config)#no vlan 2 to 7 20 25
```

Syntax: [no] vlan num to num num

If a single multi-range VLAN command contains more than 64 VLANs, the CLI does not add the VLAN IDs but instead displays an error message. An example is given below.

```
device(config)#vlan 100 to 356
ERROR -can't have more than 64 vlans at a time in a multi-range vlan command
```

Configuring a multi-range VLAN

You can configure multiple VLANs with a single command from the multi-range VLAN configuration level. For example, if you want to add tagged ethernet port 1/1/1 in the VLAN 16 17 20 21 22 23 24, enter the following commands.

```
device(config)#vlan 16 17 20 to 24
device(config-mvlan-16*24)#tag e 1/1/1
device(config-mvlan-16*24) #
```

The first command will take you to the multi-range VLAN configuration mode. The second command will add tagged ethernet port 1/1/1 in the specified VLANs, VLAN 16 17 20 21 22 23 and 24.

The following VLAN parameters can be configured with the specified VLAN range.

NOTE

The following commands are not supported on FastIron ICX 7xxx devices:

- **atalk-proto**
- **decnet-proto**
- **ipv6-proto**
- **netbios-proto**

Command	Explanation
atalk-proto	Set AppleTalk protocol VLAN
clear	Clear table/statistics/keys
decnet-proto	Set decnet protocol VLAN
end	End Configuration level and goto Privileged level
exit	Exit current level
ip-proto	Set IP protocol VLAN
ipv6-proto	Set IPv6 protocol VLAN
ipx-proto	Set IPX protocol VLAN
mac-vlan-permit	Define port to be used for MAC Based VLAN

Command	Explanation
monitor	Monitor Ingress Traffic on this VLAN(Enable VLAN)
multicast	IGMP snooping on this VLAN Mirroring)
netbios-proto	Set netbios protocol VLAN
no	Undo/disable commands
other-proto	Set other protocol VLAN
quit	Exit to User level
show	Show system information
spanning-tree	Set spanning tree for this VLAN
static-mac-address	Configure static MAC for this VLAN
tagged	802.1Q tagged port
uplink-switch	Define uplink ports and enable uplink switching

NOTE

This feature is not supported on the Brocade ICX 7750, Brocade ICX 7450, or Brocade ICX 7250.

vsrp	Configure VSRP
vsrp-aware	Configure VSRP Aware parameters
write	Write running configuration to flash or terminal

NOTE

In FSX platform, the **mac-vlan-permit** command is not available in the multi-range vlan configuration mode.

The VLAN parameters configured for the VLAN range are written in the configuration file of the individual VLANs. These VLAN parameters can also be removed or modified from the individual VLANs. In the following example, as the first step, create VLANs 16 17 20 21 22 23 24. Further, as the second step, add Ethernet port 1/1/1 in all the VLANs. As the third step, enabled 802.1w spanning tree on all these VLANs.

```
device(config)#vlan 16 17 20 to 24
device(config-mvlan-16*24)#tag e 1/1/1
device(config-mvlan-16*24)#
Added tagged port(s) ethe 1/1/1 to port-vlan16.
Added tagged port(s) ethe 1/1/1 to port-vlan 17.
Added tagged port(s) ethe 1/1/1 to port-vlan 20.
Added tagged port(s) ethe 1/1/1 to port-vlan 21.
Added tagged port(s) ethe 1/1/1 to port-vlan 22.
```

```
Added tagged port(s) ethe 1/1/1 to port-vlan 23.
Added tagged port(s) ethe 1/1/1 to port-vlan 24.
device(config-mvlan-16*24)#span 802-1w
```

The Ethernet port e 1/1/1 and spanning tree 802.1w is added to the database of each VLAN separately. You can verify the configuration with the **show running-config** command. See the example below.

```
device(config-mvlan-16*24)#show run
Current configuration:
!
!
output omitted
!
!
vlan 1 name DEFAULT-VLAN by port
!
vlan 16 by port
tagged ethe 1/1/1
spanning-tree 802-1w
!
vlan 17 by port
tagged ethe 1/1/1
spanning-tree 802-1w
!
vlan 20 by port
tagged ethe 1/1/1
spanning-tree 802-1w
!
vlan 21 by port
tagged ethe 1/1/1
spanning-tree 802-1w
!
vlan 22 by port
tagged ethe 1/1/1
spanning-tree 802-1w
!
vlan 23 by port
tagged ethe 1/1/1
spanning-tree 802-1w
!
vlan 24 by port
tagged ethe 1/1/1
spanning-tree 802-1w
!
!
output omitted
!
```

Now you can modify any one or some of the VLANs. See the example below.

In the following example, disable the spanning tree 802.1w on VLANs 22,23 and 24. And, verify with **show running-config** output that the spanning tree 802.1w is disabled on specified VLANs, VLAN 22, 23 and 24 and not on the VLANs 16, 17, 20 and 21.

```
device(config)#vlan 22 to 24
device(config-mvlan-22-24)#no span 8
device(config-mvlan-22-24)#exit
device(config)#show run
Current configuration:
output omitted
!
!
vlan 1 name DEFAULT-VLAN by port
!
vlan 16 by port
tagged ethe 1/1/1
spanning-tree 802-1w
!
vlan 17 by port
tagged ethe 1/1/1
spanning-tree 802-1w
!
```

```
vlan 20 by port
tagged ethe 1/1/1
spanning-tree 802-1w
!
vlan 21 by port
tagged ethe 1/1/1
spanning-tree 802-1w
vlan 22 by port
tagged ethe 1/1/1
!
vlan 23 by port
tagged ethe 1/1/1
!
vlan 24 by port
tagged ethe 1/1/1
output omitted
```

Multi-range VLAN show commands

This section describes the show commands for multi-range VLAN parameters.

In the multi-range VLAN mode, some of the Show commands are also available. The output of the Show commands in multi-range VLAN mode displays the information related to the specific VLANs only. See the example below.

In the following example, the first command will change the interface configuration level to the multi-range VLAN mode for the VLANs 4, 5 and 6. In the multi-range VLAN mode, enter the command **show 802.1w**. The output will display the information of STP for VLANs 4, 5 and 6

```
device(config)#vlan 4 to 6
device(config-mvlan-4-6)#show 802-1w
--- VLAN 4 [ STP Instance owned by VLAN 4 ] -----
Bridge IEEE 802.1W Parameters:
Bridge      Bridge Bridge Bridge Force   tx
Identifier  MaxAge Hello FwdDly Version Hold
hex         sec   sec   sec   cnt
8000002022227700 20    2     15    Default  3
RootBridge   RootPath DesignatedBri- Root   Max Fwd Hel
Identifier  Cost      dge Identifier Port   Age Dly lo
hex         hex      sec   sec   sec
8000002022227700 0       8000002022227700 Root   20  15  2
Port IEEE 802.1W Parameters:
      <--- Config Params --><----- Current state ----->
Port   Pri PortPath P2P Edge Role      State   Designa- Designated
Num    Cost   Mac Port                ted cost bridge
1/1/1  128 20000   F   DESIGNATED FORWARDING 0       8000002022227700
--- VLAN 5 [ STP Instance owned by VLAN 5 ] -----
Bridge IEEE 802.1W Parameters:
Bridge      Bridge Bridge Bridge Force   tx
Identifier  MaxAge Hello FwdDly Version Hold
hex         sec   sec   sec   cnt
8000002022227700 20    2     15    Default  3
RootBridge   RootPath DesignatedBri- Root   Max Fwd Hel
Identifier  Cost      dge Identifier Port   Age Dly lo
hex         hex      sec   sec   sec
8000002022227700 0       8000002022227700 Root   20  15  2
Port IEEE 802.1W Parameters:
      <--- Config Params --><----- Current state ----->
Port   Pri PortPath P2P Edge Role      State   Designa- Designated
Num    Cost   Mac Port                ted cost bridge
1/1/1  128 20000   F   DESIGNATED FORWARDING 0       8000002022227700
--- VLAN 6 [ STP Instance owned by VLAN 6 ] -----
Bridge IEEE 802.1W Parameters:
Bridge      Bridge Bridge Bridge Force   tx
Identifier  MaxAge Hello FwdDly Version Hold
hex         sec   sec   sec   cnt
8000002022227700 20    2     15    Default  3
RootBridge   RootPath DesignatedBri- Root   Max Fwd Hel
Identifier  Cost      dge Identifier Port   Age Dly lo
hex         hex      sec   sec   sec
8000002022227700 0       8000002022227700 Root   20  15  2
Port IEEE 802.1W Parameters:
      <--- Config Params --><----- Current state ----->
Port   Pri PortPath P2P Edge Role      State   Designa- Designated
```

Num	Cost	Mac	Port	Designated	Forwarding	ted	cost	bridge
1/1/1	128	20000	F	F	DESIGNATED FORWARDING	0	8000002022227700	

The following **show** parameters can be viewed for the specified VLAN range from the multi-range VLAN configuration mode. The output of these commands displays information about the specified VLANs only.

TABLE 53 VLAN show parameters

Command	Definition
802-1w	Rapid Spanning tree IEEE 802.1w status
mac-address	MAC address table
span	Spanning tree status
vlan	VLAN status
vsrp	Show VSRP commands

Layer 3 protocol-based VLANs

If you want some or all of the ports within a port-based VLAN to be organized according to Layer 3 protocol, you must configure a Layer 3 protocol-based VLAN within the port-based VLAN.

You can configure each of the following types of protocol-based VLAN within a port-based VLAN. All the ports in the Layer 3 VLAN must be in the same Layer 2 VLAN.

Layer 3 protocol-based VLANs are as follows:

- AppleTalk - The device sends AppleTalk broadcasts to all ports within the AppleTalk protocol VLAN.
- IP - The device sends IP broadcasts to all ports within the IP protocol VLAN.
- IPv6 - The device sends IPv6 broadcasts to all ports within the IPv6 protocol VLAN.
- IPX - The device sends IPX broadcasts to all ports within the IPX protocol VLAN.
- DECnet - The device sends DECnet broadcasts to all ports within the DECnet protocol VLAN.
- NetBIOS - The device sends NetBIOS broadcasts to all ports within the NetBIOS protocol VLAN.
- Other - The device sends broadcasts for all protocol types other than those listed above to all ports within the VLAN.

The following figure shows an example of Layer 3 protocol VLANs configured within a Layer 2 port-based VLAN.

FIGURE 72 Layer 3 protocol VLANs within a Layer 2 port-based VLAN

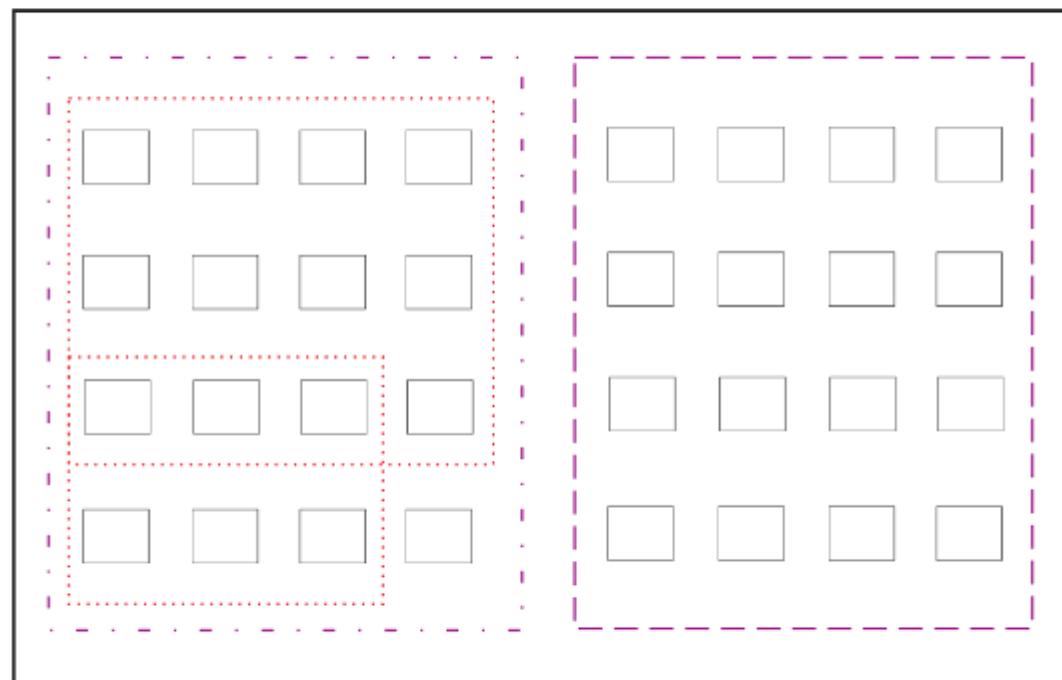
DEFAULT-VLAN

VLAN ID = 1

Layer 2 Port-based VLAN

User-configured port-based VLAN

User-configured protocol VLAN, IP sub-net VLAN,
IPX network VLAN, or AppleTalk cable VLAN



You can add Layer 3 protocol VLANs or
IP sub-net, IPX network, and AppleTalk
cable VLANs to port-based VLANs.

Layer 3 VLANs cannot span Layer 2 port-based
VLANs.

However, Layer 3 VLANs can overlap within
a Layer 2 port-based VLAN.

Integrated Switch Routing (ISR)

The Brocade Integrated Switch Routing (ISR) feature enables VLANs configured on Layer 3 Switches to route Layer 3 traffic from one protocol VLAN or IP subnet, IPX network, or AppleTalk cable VLAN to

another. Normally, to route traffic from one IP subnet, IPX network, or AppleTalk cable VLAN to another, you would need to forward the traffic to an external router. The VLANs provide Layer 3 broadcast domains for these protocols but do not in themselves provide routing services for these protocols. This is true even if the source and destination IP subnets, IPX networks, or AppleTalk cable ranges are on the same device.

ISR eliminates the need for an external router by allowing you to route between VLANs using virtual routing interfaces (ves). A virtual routing interface is a logical port on which you can configure Layer 3 routing parameters. You configure a separate virtual routing interface on each VLAN that you want to be able to route from or to. For example, if you configure two IP subnet VLANs on a Layer 3 Switch, you can configure a virtual routing interface on each VLAN, then configure IP routing parameters for the subnets. Thus, the Layer 3 Switch forwards IP subnet broadcasts within each VLAN at Layer 2 but routes Layer 3 traffic between the VLANs using the virtual routing interfaces.

NOTE

The Layer 3 Switch uses the lowest MAC address on the device (the MAC address of port 1 or 1/1) as the MAC address for all ports within all virtual routing interfaces you configure on the device.

The routing parameters and the syntax for configuring them are the same as when you configure a physical interface for routing. The logical interface allows the Layer 3 Switch to internally route traffic between the protocol-based VLANs without using physical interfaces.

All the ports within a protocol-based VLAN must be in the same port-based VLAN. The protocol-based VLAN cannot have ports in multiple port-based VLANs, unless the ports in the port-based VLAN to which you add the protocol-based VLAN are 802.1Q tagged.

You can configure multiple protocol-based VLANs within the same port-based VLAN. In addition, a port within a port-based VLAN can belong to multiple protocol-based VLANs of the same type or different types. For example, if you have a port-based VLAN that contains ports 1 - 10, you can configure port 5 as a member of an AppleTalk protocol VLAN, an IP protocol VLAN, and an IPX protocol VLAN, and so on.

IP subnet, IPX network, and AppleTalk cable VLANs

The protocol-based VLANs described in the previous section provide separate protocol broadcast domains for specific protocols. For IP, IPX, and AppleTalk, you can provide more granular broadcast control by instead creating the following types of VLAN:

- IP subnet VLAN - An IP subnet broadcast domain for a specific IP subnet.
- IPX network VLAN - An IPX network broadcast domain for a specific IPX network.
- AppleTalk cable VLAN - An AppleTalk broadcast domain for a specific cable range.

You can configure these types of VLANs on Layer 3 Switches only. The Layer 3 Switch sends broadcasts for the IP subnet, IPX network, or AppleTalk cable range to all ports within the IP subnet, IPX network, or AppleTalk cable VLAN at Layer 2.

The Layer 3 Switch routes packets between VLANs at Layer 3. To configure an IP subnet, IPX network, or AppleTalk cable VLAN to route, you must add a virtual routing interface to the VLAN, then configure the appropriate routing parameters on the virtual routing interface.

NOTE

The Layer 3 Switch routes packets between VLANs of the same protocol. The Layer 3 Switch cannot route from one protocol to another.

NOTE

IP subnet VLANs are not the same thing as IP protocol VLANs. An IP protocol VLAN sends all IP broadcasts on the ports within the IP protocol VLAN. An IP subnet VLAN sends only the IP subnet broadcasts for the subnet of the VLAN. You cannot configure an IP protocol VLAN and an IP subnet VLAN within the same port-based VLAN. This note also applies to IPX protocol VLANs and IPX network VLANs, and to AppleTalk protocol VLANs and AppleTalk cable VLANs.

Default VLAN

By default, all the ports on a FastIron device are in a single port-based VLAN. This VLAN is called the DEFAULT-VLAN and is VLAN number 1. FastIron devices do not contain any protocol VLANs or IP subnet, IPX network, or AppleTalk cable VLANs by default.

The following figure shows an example of the default Layer 2 port-based VLAN.

FIGURE 73 Default Layer 2 port-based VLAN

DEFAULT-VLAN

VLAN ID = 1

Layer 2 Port-based VLAN



By default, all ports belong to a single port-based VLAN, DEFAULT-VLAN.
Thus, all ports belong to a single Layer 2 broadcast domain.

When you configure a port-based VLAN, one of the configuration items you provide is the ports that are in the VLAN. When you configure the VLAN, the Brocade device automatically removes the ports that you place in the VLAN from DEFAULT-VLAN. By removing the ports from the default VLAN, the Brocade device ensures that each port resides in only one Layer 2 broadcast domain.

NOTE

Information for the default VLAN is available only after you define another VLAN.

Some network configurations may require that a port be able to reside in two or more Layer 2 broadcast domains (port-based VLANs). In this case, you can enable a port to reside in multiple port-based VLANs by tagging the port. Refer to the following section.

If your network requires that you use VLAN ID 1 for a user-configured VLAN, you can reassign the default VLAN to another valid VLAN ID. Refer to [Assigning a different VLAN ID to the default VLAN](#) on page 362.

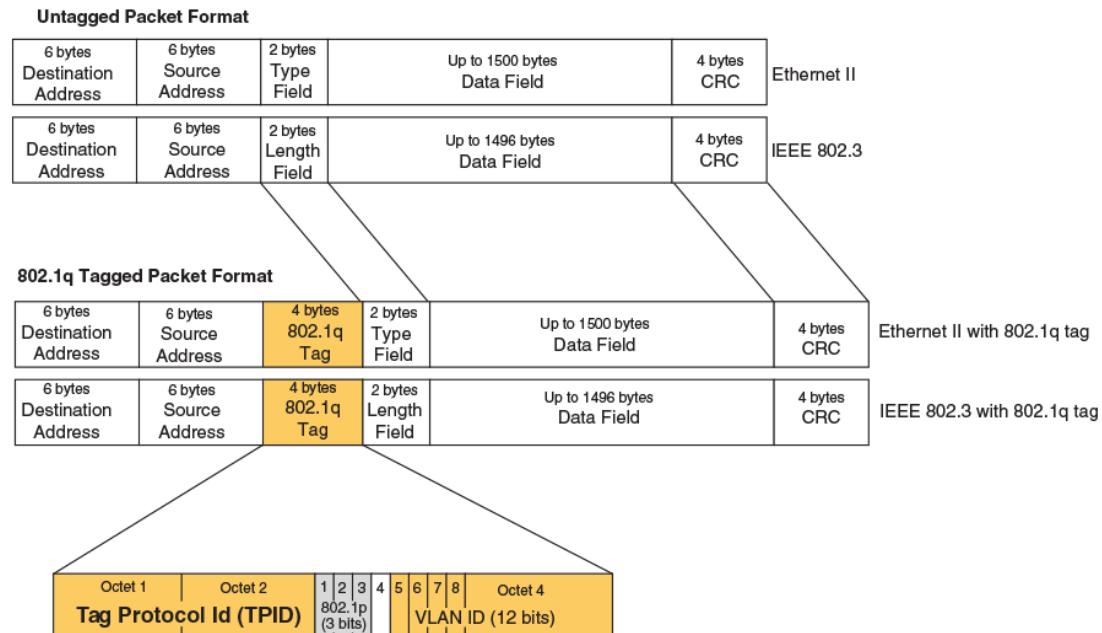
802.1Q tagging

802.1Q tagging is an IEEE standard that allows a networking device to add information to a Layer 2 packet in order to identify the VLAN membership of the packet. Brocade devices tag a packet by adding a four-byte tag to the packet. The tag contains the tag value, which identifies the data as a tag, and also contains the VLAN ID of the VLAN from which the packet is sent.

- The default tag value is 8100 (hexadecimal). This value comes from the 802.1Q specification. You can change this tag value on a global basis on Brocade devices if needed to be compatible with other vendors' equipment.
- The VLAN ID is determined by the VLAN on which the packet is being forwarded.

The following figure shows the format of packets with and without the 802.1Q tag. The tag format is vendor-specific. To use the tag for VLANs configured across multiple devices, make sure all the devices support the same tag format.

FIGURE 74 Packet containing a Brocade 802.1Q VLAN tag

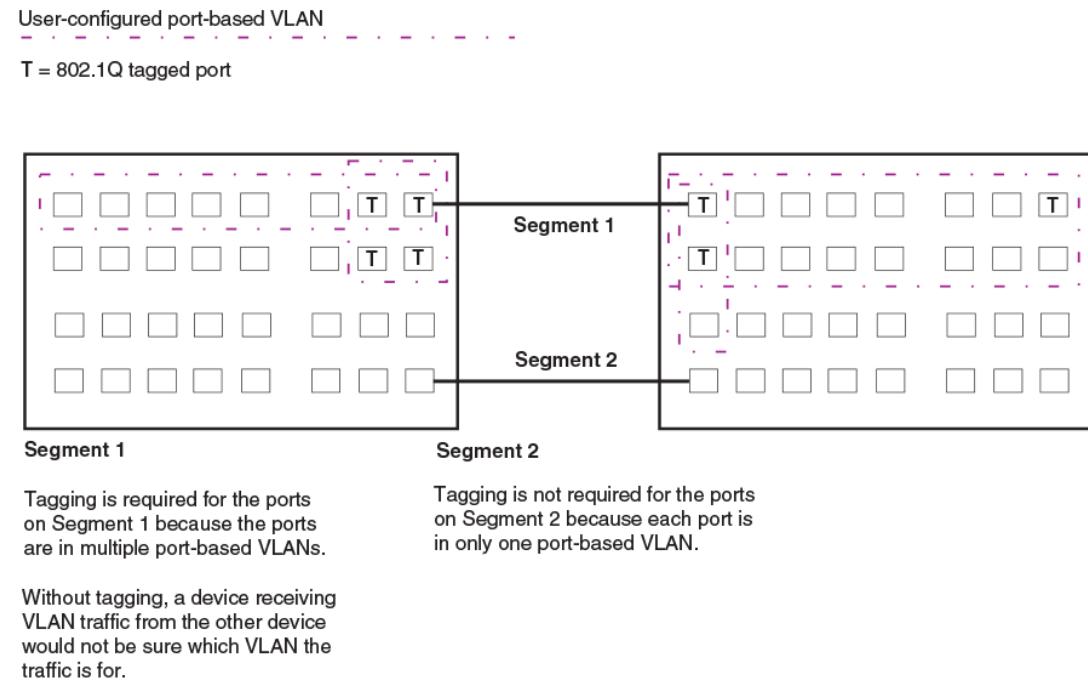


If you configure a VLAN that spans multiple devices, you need to use tagging only if a port connecting one of the devices to the other is a member of more than one port-based VLAN. If a port connecting one device to the other is a member of only a single port-based VLAN, tagging is not required.

If you use tagging on multiple devices, each device must be configured for tagging and must use the same tag value. In addition, the implementation of tagging must be compatible on the devices. The tagging on all Brocade devices is compatible with other Brocade devices.

The following figure shows an example of two devices that have the same Layer 2 port-based VLANs configured across them. Notice that only one of the VLANs requires tagging.

FIGURE 75 VLANs configured across multiple devices



Support for 802.1ad (Q-in-Q) tagging

Brocade devices provide finer granularity for configuring 802.1Q tagging, enabling you to configure 802.1Q tag-types on a group of ports, thereby enabling the creation of two identical 802.1Q tags (802.1ad tagging) on a single device. This enhancement improves SAV interoperability between Brocade devices and other vendors' devices that support the 802.1Q tag-types, but are not very flexible with the tag-types they accept.

- Q-in-Q is not supported on ICX 6430-C and ICX 6450-C devices.
- FastIron X Series devices support one value for **tag-type**, which is defined at the global level.
- Brocade FCX Series and ICX 6650 devices support one value for **tag-type**, which is defined at the global level, and one value for **tag-profile**, which is defined at the global and interface level of the CLI.

802.1 ad tagging for Brocade FCX Series and ICX 6650 devices

The following enhancements allow the Brocade FCX Series and ICX 6650 devices, including those in an IronStack, to use Q-in-Q and SAV, by allowing the changing of a tag profile for ports:

- In addition to the default tag type 0x8100, you can now configure one additional global tag profile with a number from 0xffff.
- Tag profiles on a single port, or a group of ports can be configured to point to the global tag profile.

For example applications and configuration details, refer to [802.1ad tagging configuration](#) on page 395.

To configure a global tag profile, enter the following command in the configuration mode.

```
device(config)# tag-profile 9500
```

Syntax: [no] tag-profile tag-no

tag-no - the number of the tag, can be 0x8100 (default), or 0xffff

To direct individual ports or on a range of ports to this tag profile, enter commands similar to the following.

```
device(config)# interface ethernet 1/1/1
device(config-if-e1000-1/1/1)# tag-profile enable
device(config-mif-1/1/1,1/2/1)# tag-profile enable
```

Spanning Tree Protocol (STP)

The default state of STP depends on the device type:

- STP is disabled by default on Brocade Layer 3 Switches.
- STP is enabled by default on Brocade Layer 2 Switches.

Also by default, each port-based VLAN has a separate instance of STP. Thus, when STP is globally enabled, each port-based VLAN on the device runs a separate spanning tree.

You can enable or disable STP on the following levels:

- **Globally** - Affects all ports on the device.

NOTE

If you configure a port-based VLAN on the device, the VLAN has the same STP state as the default STP state on the device. Thus, on Layer 2 Switches, new VLANs have STP enabled by default. On Layer 3 Switches, new VLANs have STP disabled by default. You can enable or disable STP in each VLAN separately. In addition, you can enable or disable STP on individual ports.

- **Port-based VLAN** - Affects all ports within the specified port-based VLAN.

STP is a Layer 2 protocol. Thus, you cannot enable or disable STP for individual protocol VLANs or for IP subnet, IPX network, or AppleTalk cable VLANs. The STP state of a port-based VLAN containing these other types of VLANs determines the STP state for all the Layer 2 broadcasts within the port-based VLAN. This is true even though Layer 3 protocol broadcasts are sent on Layer 2 within the VLAN.

It is possible that STP will block one or more ports in a protocol VLAN that uses a virtual routing interface to route to other VLANs. For IP protocol and IP subnet VLANs, even though some of the physical ports of the virtual routing interface are blocked, the virtual routing interface can still route so long as at least one port in the virtual routing interface protocol VLAN is not blocked by STP.

If you enable Single STP (SSTP) on the device, the ports in all VLANs on which STP is enabled become members of a single spanning tree. The ports in VLANs on which STP is disabled are excluded from the single spanning tree.

For more information, refer to [Spanning Tree Protocol](#) on page 241.

Virtual routing interfaces

A virtual routing interface is a logical routing interface that Brocade Layer 3 Switches use to route Layer 3 protocol traffic between protocol VLANs.

Brocade devices send Layer 3 traffic at Layer 2 within a protocol VLAN. However, Layer 3 traffic from one protocol VLAN to another must be routed.

If you want the device to be able to send Layer 3 traffic from one protocol VLAN to another, you must configure a virtual routing interface on each protocol VLAN, then configure routing parameters on the virtual routing interfaces. For example, to enable a Layer 3 Switch to route IP traffic from one IP subnet VLAN to another, you must configure a virtual routing interface on each IP subnet VLAN, then configure the appropriate IP routing parameters on each of the virtual routing interfaces.

The following figure shows an example of Layer 3 protocol VLANs that use virtual routing interfaces for routing.

FIGURE 76 Use virtual routing interfaces for routing between Layer 3 protocol VLANs

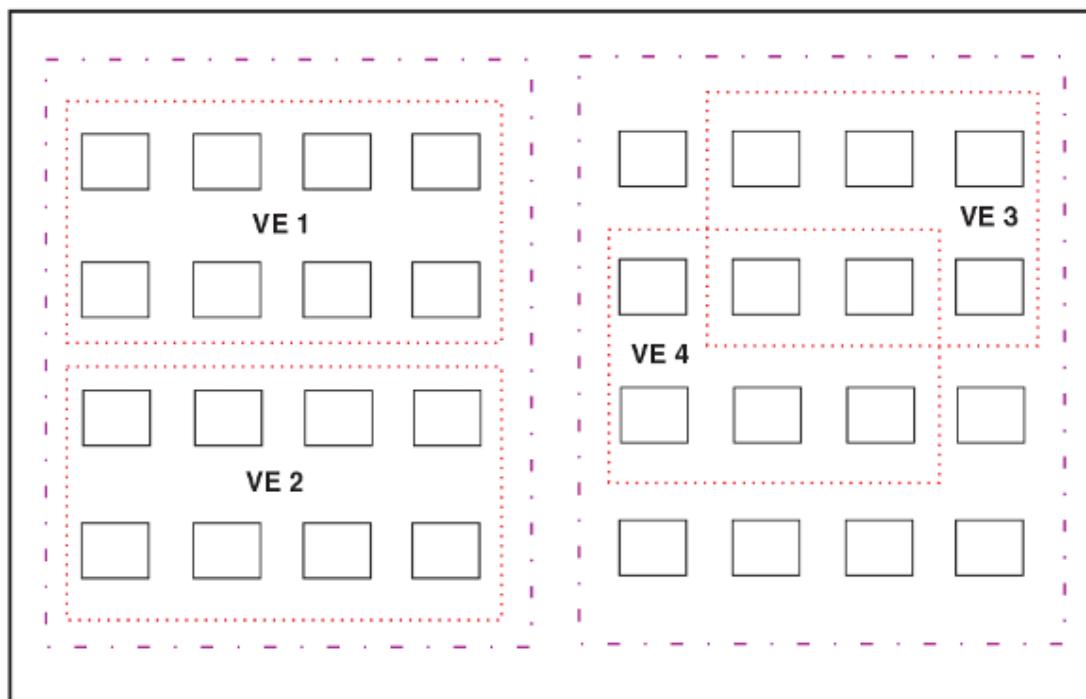
User-configured port-based VLAN



User-configured protocol VLAN, IP sub-net VLAN,
IPX network VLAN, or AppleTalk cable VLAN



VE = virtual interface
("VE" stands for "Virtual Ethernet")



Layer 2 and Layer 3 traffic within a VLAN
is bridged at Layer 2.

Layer 3 traffic between protocol VLANs
is routed using virtual interfaces (VE).
To route to one another, each protocol
VLAN must have a virtual interface.

VLAN and virtual routing interface groups

Brocade FastIron devices support the configuration of VLAN groups. To simplify configuration, you can configure VLAN groups and virtual routing interface groups. When you create a VLAN group, the VLAN parameters you configure for the group apply to all the VLANs within the group. Additionally, you can

easily associate the same IP subnet interface with all the VLANs in a group by configuring a virtual routing interface group with the same ID as the VLAN group.

For configuration information, refer to [VLAN groups and virtual routing interface group](#) on page 383.

Dynamic, static, and excluded port membership

When you add ports to a protocol VLAN, IP subnet VLAN, IPX network VLAN, or AppleTalk cable VLAN, you can add them dynamically or statically:

- Dynamic ports
- Static ports

You also can explicitly exclude ports.

Dynamic ports

Dynamic ports are added to a VLAN when you create the VLAN. However, if a dynamically added port does not receive any traffic for the VLAN protocol within ten minutes, the port is removed from the VLAN. However, the port remains a candidate for port membership. Thus, if the port receives traffic for the VLAN protocol, the device adds the port back to the VLAN.

After the port is added back to the VLAN, the port can remain an active member of the VLAN up to 20 minutes without receiving traffic for the VLAN protocol. If the port ages out, it remains a candidate for VLAN membership and is added back to the VLAN when the VLAN receives protocol traffic. At this point, the port can remain in the VLAN up to 20 minutes without receiving traffic for the VLAN protocol, and so on.

Unless you explicitly add a port statically or exclude a port, the port is a dynamic port and thus can be an active member of the VLAN, depending on the traffic it receives.

NOTE

You cannot configure dynamic ports in an AppleTalk cable VLAN. The ports in an AppleTalk cable VLAN must be static. However, ports in an AppleTalk protocol VLAN can be dynamic or static.

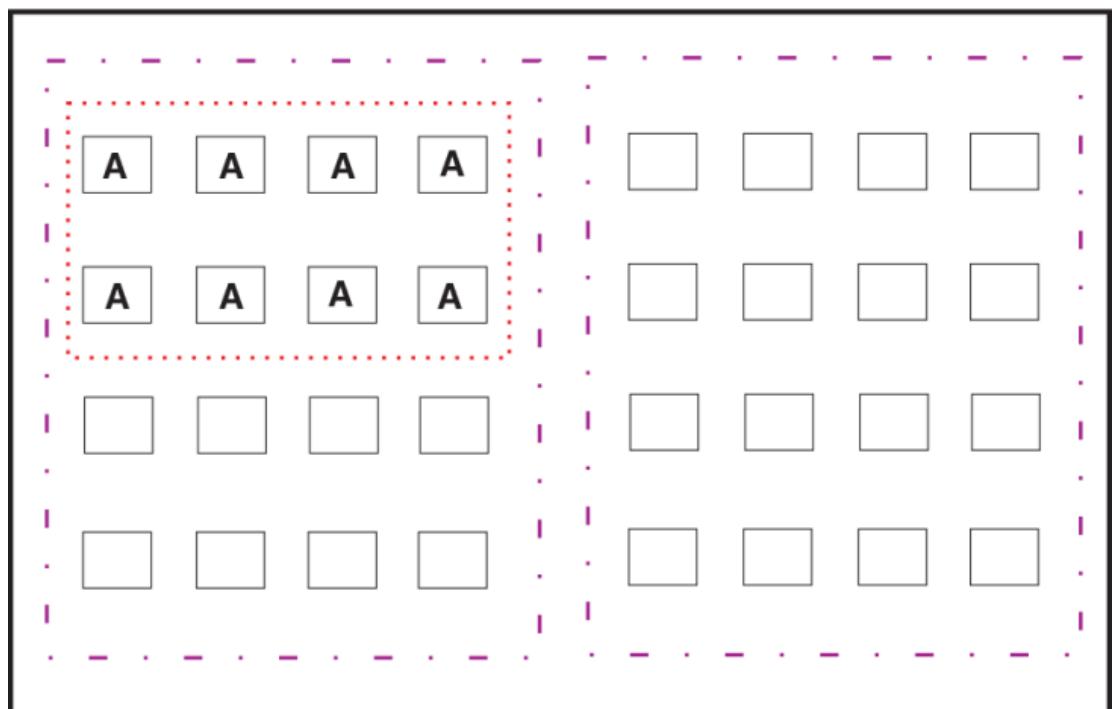
The following figure shows an example of a VLAN with dynamic ports. Dynamic ports not only join and leave the VLAN according to traffic, but also allow some broadcast packets of the specific protocol to "leak" through the VLAN. Refer to [Broadcast leaks](#) on page 359.

FIGURE 77 VLAN with dynamic ports--all ports are active when you create the VLAN

A = active port

C = candidate port

When you add ports dynamically,
all the ports are added when you add
the VLAN.

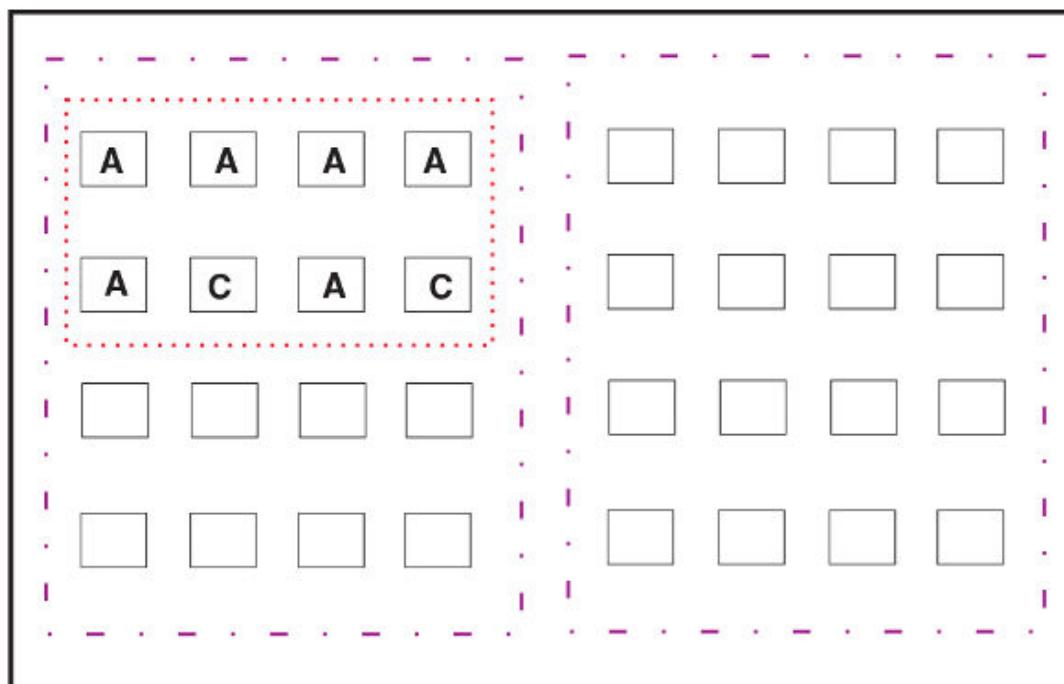


SUBNET Ports in a new protocol VLAN that do not receive traffic for the VLAN protocol age out after 10 minutes and become candidate ports. The above figure shows what happens if a candidate port receives traffic for the VLAN protocol.

FIGURE 78 VLAN with dynamic ports--candidate ports become active again if they receive protocol traffic

Ports that time out remain candidates for membership in the VLAN and become active again if they receive traffic for the VLAN's protocol, IP sub-net, IPX network, or AppleTalk cable range.

When a candidate port rejoins a VLAN, the timeout for that port becomes 20 minutes. Thus, the port remains an active member of the VLAN even if it does not receive traffic for 20 minutes. After that, the port becomes a candidate port again.



Static ports

Static ports are permanent members of the protocol VLAN. The ports remain active members of the VLAN regardless of whether the ports receive traffic for the VLAN protocol. You must explicitly identify the port as a static port when you add it to the VLAN. Otherwise, the port is dynamic and is subject to aging out.

Excluded ports

If you want to prevent a port in a port-based VLAN from ever becoming a member of a protocol, IP subnet, IPX network, or AppleTalk cable VLAN configured in the port-based VLAN, you can explicitly exclude the port. You exclude the port when you configure the protocol, IP subnet, IPX network, or AppleTalk cable VLAN.

Excluded ports do not leak broadcast packets. Refer to [Broadcast leaks](#) on page 359.

Broadcast leaks

A dynamic port becomes a member of a Layer 3 protocol VLAN when traffic from the VLAN's protocol is received on the port. After this point, the port remains an active member of the protocol VLAN, unless the port does not receive traffic from the VLAN's protocol for 20 minutes. If the port does not receive traffic for the VLAN's protocol for 20 minutes, the port ages out and is no longer an active member of the VLAN.

To enable a host that has been silent for awhile to send and receive packets, the dynamic ports that are currently members of the Layer 3 protocol VLAN "leak" Layer 3 broadcast packets to the ports that have aged out. When a host connected to one of the aged out ports responds to a leaked broadcast, the port is added to the protocol VLAN again.

To "leak" Layer 3 broadcast traffic, an active port sends 1/8th of the Layer 3 broadcast traffic to the inactive (aged out) ports.

Static ports do not age out and do not leak broadcast packets.

Super aggregated VLANs

Brocade FastIron devices support Super Aggregated VLANs. You can aggregate multiple VLANs within another VLAN. This feature allows you to construct Layer 2 paths and channels. This feature is particularly useful for Virtual Private Network (VPN) applications in which you need to provide a private, dedicated Ethernet connection for an individual client to transparently reach its subnet across multiple networks.

For an application example and configuration information, refer to [Super-aggregated VLAN configuration](#) on page 388.

Trunk group ports and VLAN membership

A trunk group is a set of physical ports that are configured to act as a single physical interface. Each trunk group port configuration is based on the configuration of the lead port, which is the lowest numbered port in the group.

If you add a trunk group lead port to a VLAN, all of the ports in the trunk group become members of that VLAN.

Summary of VLAN configuration rules

A hierarchy of VLANs exists between the Layer 2 and Layer 3 protocol-based VLANs:

- Port-based VLANs are at the lowest level of the hierarchy.
- Layer 3 protocol-based VLANs, IP, IPv6, IPX, AppleTalk, Decnet, and NetBIOS are at the middle level of the hierarchy.
- IP subnet, IPX network, and AppleTalk cable VLANs are at the top of the hierarchy.

NOTE

You cannot have a protocol-based VLAN and a subnet or network VLAN of the same protocol type in the same port-based VLAN. For example, you can have an IPX protocol VLAN and IP subnet VLAN in the same port-based VLAN, but you cannot have an IP protocol VLAN and an IP subnet VLAN in the same port-based VLAN, nor can you have an IPX protocol VLAN and an IPX network VLAN in the same port-based VLAN.

As a Brocade device receives packets, the VLAN classification starts from the highest level VLAN first. Therefore, if an interface is configured as a member of both a port-based VLAN and an IP protocol VLAN, IP packets coming into the interface are classified as members of the IP protocol VLAN because that VLAN is higher in the VLAN hierarchy.

Multiple VLAN membership rules

- A port can belong to multiple, unique, overlapping Layer 3 protocol-based VLANs without VLAN tagging.
- A port can belong to multiple, overlapping Layer 2 port-based VLANs only if the port is a tagged port. Packets sent out of a tagged port use an 802.1Q-tagged frame.
- When both port and protocol-based VLANs are configured on a given device, all protocol VLANs must be strictly contained within a port-based VLAN. A protocol VLAN cannot include ports from multiple port-based VLANs. This rule is required to ensure that port-based VLANs remain loop-free Layer 2 broadcast domains.
- IP protocol VLANs and IP subnet VLANs cannot operate concurrently on the system or within the same port-based VLAN.
- IPX protocol VLANs and IPX network VLANs cannot operate concurrently on the system or within the same port-based VLAN.
- If you first configure IP and IPX protocol VLANs before deciding to partition the network by IP subnet and IPX network VLANs, then you need to delete those VLANs before creating the IP subnet and IPX network VLANs.
- One of each type of protocol VLAN is configurable within each port-based VLAN on the Layer 2 Switch.
- Multiple IP subnet and IPX network VLANs are configurable within each port-based VLAN on the Layer 2 Switch.
- Removing a configured port-based VLAN from a Brocade Communication Systems, Inc Layer 2 Switch or Layer 3 Switch automatically removes any protocol-based VLAN, IP subnet VLAN, AppleTalk cable VLAN, or IPX network VLAN, or any Virtual Ethernet router interfaces defined within the Port-based VLAN.

Routing between VLANs

Brocade Layer 3 Switches can locally route IP, IPX, and Appletalk between VLANs defined within a single router. All other routable protocols or protocol VLANs (for example, DecNet) must be routed by another external router capable of routing the protocol.

Virtual routing interfaces (Layer 2 Switches only)

You need to configure virtual routing interfaces if an IP, IPX, or Appletalk protocol VLAN, IP subnet VLAN, AppleTalk cable VLAN, or IPX network VLAN needs to route protocols to another port-based

VLAN on the same router. A virtual routing interface can be associated with the ports in only a single port-based VLAN. Virtual router interfaces must be defined at the highest level of the VLAN hierarchy.

If you do not need to further partition the port-based VLAN by defining separate Layer 3 VLANs, you can define a single virtual routing interface at the port-based VLAN level and enable IP, IPX, and Appletalk routing on a single virtual routing interface.

Some configurations may require simultaneous switching and routing of the same single protocol across different sets of ports on the same router. When IP, IPX, or Appletalk routing is enabled on a Brocade Layer 3 Switch, you can route these protocols on specific interfaces while bridging them on other interfaces. In this scenario, you can create two separate backbones for the same protocol, one bridged and one routed.

To bridge IP, IPX, or Appletalk at the same time these protocols are being routed, you need to configure an IP protocol, IP subnet, IPX protocol, IPX network, or Appletalk protocol VLAN and not assign a virtual routing interface to the VLAN. Packets for these protocols are bridged or switched at Layer 2 across ports on the router that are included in the Layer 3 VLAN. If these VLANs are built within port-based VLANs, they can be tagged across a single set of backbone fibers to create separate Layer 2 switched and Layer 3 routed backbones for the same protocol on a single physical backbone.

Routing between VLANs using virtual routing interfaces (Layer 3 Switches only)

Brocade calls the ability to route between VLANs with virtual routing interfaces Integrated Switch Routing (ISR) . There are some important concepts to understand before designing an ISR backbone.

Virtual router interfaces can be defined on port-based, IP protocol, IP subnet, IPX protocol, IPX network, AppleTalk protocol, and AppleTalk cable VLANs.

To create any type of VLAN on a Brocade Layer 3 Switch, Layer 2 forwarding must be enabled. When Layer 2 forwarding is enabled, the Layer 3 Switch becomes a Switch on all ports for all non-routable protocols.

If the router interfaces for IP, IPX, or AppleTalk are configured on physical ports, then routing occurs independent of the Spanning Tree Protocol (STP). However, if the router interfaces are defined for any type VLAN, they are virtual routing interfaces and are subject to the rules of STP.

If your backbone consists of virtual routing interfaces all within the same STP domain, it is a bridged backbone, not a routed one. This means that the set of backbone interfaces that are blocked by STP will be blocked for routed protocols as well. The routed protocols will be able to cross these paths only when the STP state of the link is FORWARDING. This problem is easily avoided by proper network design.

When designing an ISR network, pay attention to your use of virtual routing interfaces and the spanning-tree domain. If Layer 2 switching of your routed protocols (IP, IPX, AppleTalk) is not required across the backbone, then the use of virtual routing interfaces can be limited to edge switch ports within each router. Full backbone routing can be achieved by configuring routing on each physical interface that connects to the backbone. Routing is independent of STP when configured on a physical interface.

If your ISR design requires that you switch IP, IPX, or Appletalk at Layer 2 while simultaneously routing the same protocols over a single backbone, then create multiple port-based VLANs and use VLAN tagging on the backbone links to separate your Layer 2 switched and Layer 3 routed networks.

There is a separate STP domain for each port-based VLAN. Routing occurs independently across port-based VLANs or STP domains. You can define each end of each backbone link as a separate tagged port-based VLAN. Routing will occur independently across the port-based VLANs. Because each port-based VLAN STP domain is a single point-to-point backbone connection, you are guaranteed to never have an STP loop. STP will never block the virtual router interfaces within the tagged port-based VLAN, and you will have a fully routed backbone.

Dynamic port assignment (Layer 2 Switches and Layer 3 Switches)

All Switch ports are dynamically assigned to any Layer 3 VLAN on Brocade Layer 2 Switches and any non-routable VLAN on Brocade Layer 3 Switches. To maintain explicit control of the VLAN, you can explicitly exclude ports when configuring any Layer 3 VLAN on a Brocade Layer 2 Switch or any non-routable Layer 3 VLAN on a Brocade Layer 3 Switch.

If you do not want the ports to have dynamic membership, you can add them statically. This eliminates the need to explicitly exclude the ports that you do not want to participate in a particular Layer 3 VLAN.

Assigning a different VLAN ID to the default VLAN

When you enable port-based VLANs, all ports in the system are added to the default VLAN. By default, the default VLAN ID is "VLAN 1". The default VLAN is not configurable. If you want to use the VLAN ID "VLAN 1" as a configurable VLAN, you can assign a different VLAN ID to the default VLAN.

To reassign the default VLAN to a different VLAN ID, enter the following command.

```
device(config)# default-vlan-id 4095
```

Syntax: [no] default-vlan-id *vlan-id*

You must specify a valid VLAN ID that is not already in use. For example, if you have already defined VLAN 10, do not try to use "10" as the new VLAN ID for the default VLAN. Valid VLAN IDs are numbers from 1 - 4095.

NOTE

This command does not change the properties of the default VLAN. Changing the name allows you to use the VLAN ID "1" as a configurable VLAN.

Assigning different VLAN IDs to reserved VLANs 4091 and 4092

If you want to use VLANs 4091 and 4092 as configurable VLANs, you can assign them to different VLAN IDs.

For example, to reassign reserved VLAN 4091 to VLAN 10, enter the following commands.

```
device(config)# reserved-vlan-map vlan 4091 new-vlan 10
Reload required. Please write memory and then reload or power cycle.
device(config)# write mem
device(config)# exit
device# reload
```

NOTE

You must save the configuration (write mem) and reload the software to place the change into effect.

The above configuration changes the VLAN ID of 4091 to 10. After saving the configuration and reloading the software, you can configure VLAN 4091 as you would any other VLAN.

Syntax: [no] reserved-vlan-map *vlan* 4091 | 4092 *new-vlan* *vlan-id*

For *vlan-id*, enter a valid VLAN ID that is not already in use. For example, if you have already defined VLAN 20, do not try to use "20" as the new VLAN ID. Valid VLAN IDs are numbers from 1 - 4090, 4093, and 4095. VLAN ID 4094 is reserved for use by the Single Spanning Tree feature.

Viewing reassigned VLAN IDs for reserved VLANs 4091 and 4092

To view the assigned VLAN IDs for reserved VLANs 4091 and 4092, use the **show reserved-vlan-map** command. The reassigned VLAN IDs also display in the output of the **show running-config** and **show config** commands.

The following shows example output for the **show reserved-vlan-map** command.

```
device # show reserved-vlan-map
Reserved Purpose          Default      Re-assign      Current
  CPU VLAN                4091           10            10
  All Ports VLAN           4092           33            33
```

Syntax: show reserved-vlan-map

The following table defines the fields in the output of the **show reserved-vlan-map** command.

TABLE 54 Output of the **show reserved-vlan-map** command

Field	Description
Reserved Purpose	Describes for what the VLAN is reserved. Note that the description is for Brocade internal VLAN management.
Default	The default VLAN ID of the reserved VLAN.
Re-assign	The VLAN ID to which the reserved VLAN was reassigned. ¹
Current	The current VLAN ID for the reserved VLAN. ¹

1. If you reassign a reserved VLAN without saving the configuration and reloading the software, the reassigned VLAN ID will display in the Re-assign column. However, the previously configured or default VLAN ID will display in the Current column until the configuration is saved and the device reloaded.

Assigning trunk group ports

When a "lead" trunk group port is assigned to a VLAN, all other members of the trunk group are automatically added to that VLAN. A lead port is the first port of a trunk group port range; for example, "1" in 1 - 4 or "5" in 5 - 8.

Enable spanning tree on a VLAN

The spanning tree bridge and port parameters are configurable using one CLI command set at the Global Configuration Level of each Port-based VLAN. Suppose you want to enable the IEEE 802.1D STP across VLAN 3. To do so, use the following method.

NOTE

When port-based VLANs are not operating on the system, STP is set on a system-wide level at the global CONFIG level of the CLI.

1. Access the global CONFIG level of the CLI on FSX-A by entering the following commands.

```
device-A> enable
No password has been assigned yet...
```

```
device-A# configure terminal
device-A(config) #
```

2. Access the level of the CLI for configuring port-based VLAN 3 by entering the following command.

```
device-A(config)# 
device-A(config)# vlan 3
device-A(config-vlan-3) #
```

3. From VLAN 3 configuration level of the CLI, enter the following command to enable STP on all tagged and untagged ports associated with VLAN 3.

```
device-B(config-vlan-3) #
device-B(config-vlan-3) # spanning-tree
device-B(config-vlan-3) #
```

4. Enter the following commands to exit the VLAN CONFIG mode and save the configuration to the system-config file on flash memory.

```
device-B(config-vlan-3) #
device-B(config-vlan-3) # end
device-B# write memory
device-B#
```

5. Repeat steps 1 - 4 on FSX-B.

NOTE

You do not need to configure values for the STP parameters. All parameters have default values as noted below. Additionally, all values will be globally applied to all ports on the system or on the port-based VLAN for which they are defined.

To configure a specific path-cost or priority value for a given port, enter those values using the key words in the brackets [] shown in the syntax summary below. If you do not want to specify values for any given port, this portion of the command is not required.

Syntax: **vlan *vlan-id* by port**

Syntax: **[no] spanning-tree**

Syntax: **spanning-tree [ethernet [slotnum/] portnum path-cost value priority value] forward-delay value hello-time value maximum-age time priority value**

Bridge STP parameters (applied to all ports within a VLAN):

- Forward Delay - the period of time a bridge will wait (the listen and learn period) before forwarding data packets. Possible values: 4 - 30 seconds. Default is 15.
- Maximum Age - the interval a bridge will wait for receipt of a hello packet before initiating a topology change. Possible values: 6 - 40 seconds. Default is 20.
- Hello Time - the interval of time between each configuration BPDU sent by the root bridge. Possible values: 1 - 10 seconds. Default is 2.
- Priority - a parameter used to identify the root bridge in a network. The bridge with the lowest value has the highest priority and is the root. Possible values: 1 - 65,535. Default is 32,678.

Port parameters (applied to a specified port within a VLAN):

- Path Cost - a parameter used to assign a higher or lower path cost to a port. Possible values: 1 - 65535. Default is (1000/Port Speed) for Half-Duplex ports and is (1000/Port Speed)/2 for Full-Duplex ports.
- Priority - value determines when a port will be rerouted in relation to other ports. Possible values: 0 - 255. Default is 128.

Configuring IP subnet, IPX network and protocol-based VLANs

Protocol-based VLANs provide the ability to define separate broadcast domains for several unique Layer 3 protocols within a single Layer 2 broadcast domain. Some applications for this feature might include security between departments with unique protocol requirements. This feature enables you to limit the amount of broadcast traffic end-stations, servers, and routers need to accept.

IP subnet, IPX network, and protocol-based VLAN configuration example

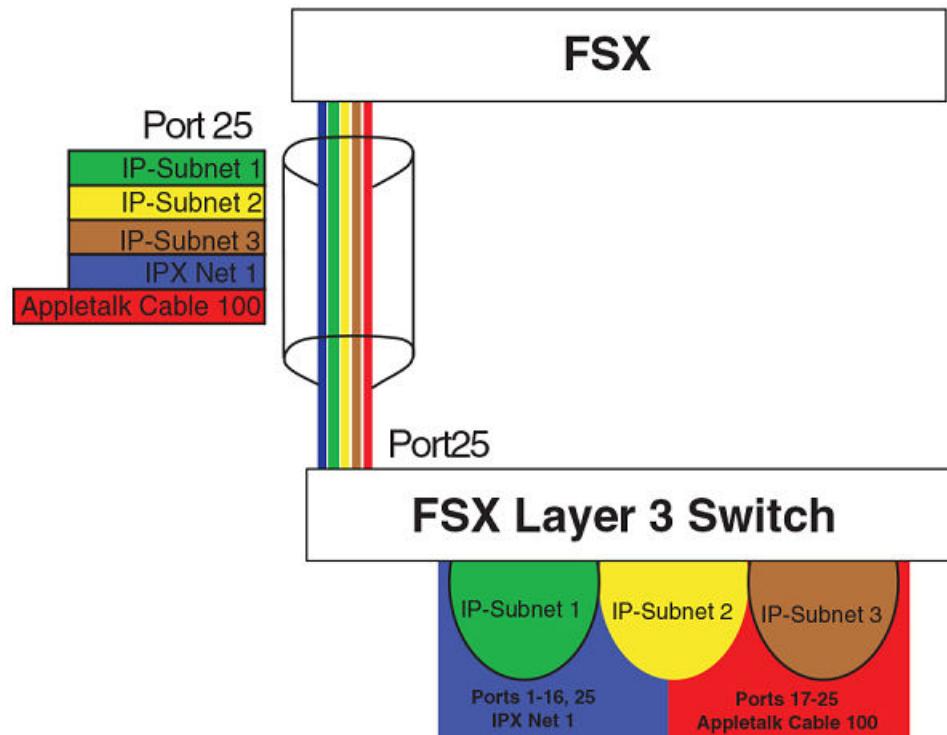
Suppose you want to create five separate Layer 3 broadcast domains within a single Layer 2 STP broadcast domain:

- Three broadcast domains, one for each of three separate IP subnets
- One for IPX Network 1
- One for the Appletalk protocol

Also suppose you want a single router interface to be present within all of these separate broadcast domains, without using IEEE 802.1Q VLAN tagging or any proprietary form of VLAN tagging.

The following figure shows this configuration.

FIGURE 79 Protocol-based (Layer 3) VLANs



To configure the VLANs shown in the above figure, use the following procedure.

1. To permanently assign ports 1 - 8 and port 25 to IP subnet VLAN 10.1.1.0, enter the following commands.

```
device(config-vlan-2)# ip-subnet 10.1.1.0/24 name Green  
device(config-vlan-ip-subnet)# no dynamic  
device(config-vlan-ip-subnet)# static ethernet 1 to 8 ethernet 25
```

2. To permanently assign ports 9 - 16 and port 25 to IP subnet VLAN 10.1.2.0, enter the following commands.

```
device(config-vlan-3)# ip-subnet 10.1.2.0/24 name Yellow  
device(config-vlan-ip-subnet)# no dynamic  
device(config-vlan-ip-subnet)# static ethernet 9 to 16 ethernet 25
```

3. To permanently assign ports 17 - 25 to IP subnet VLAN 10.1.3.0, enter the following commands.

```
device(config-vlan-4)# ip-subnet 10.1.3.0/24 name Brown  
device(config-vlan-ip-subnet)# no dynamic  
device(config-vlan-ip-subnet)# static ethernet 17 to 25
```

4. To permanently assign ports 1 - 12 and port 25 to IPX network 1 VLAN, enter the following commands.

```
device(config-ip-subnet)# ipx-network 1 ethernet_802.3 name Blue  
device(config-ipx-network)# no dynamic  
device(config-ipx-network)# static ethernet 1 to 12 ethernet 25  
device(config-ipx-network) #
```

5. To permanently assign ports 12 - 25 to Appletalk VLAN, enter the following commands.

```
device(config-ipx-proto)# atalk-proto name Red  
device(config-atalk-proto)# no dynamic  
device(config-atalk-proto)# static ethernet 13 to 25  
device(config-atalk-proto)# end  
device# write memory  
device#
```

Syntax: `ip-subnet ip-addr ip-mask [name string]`

Syntax: `ipx-network ipx-network-number frame-encapsulation-type netbios-allow | netbios-disallow [name string]`

Syntax: `ip-proto | ipx-proto | atalk-proto | decnet-proto | netbios-proto | other-proto static | exclude | dynamic ethernet [slotnum/]portnum [to [slotnum/]portnum] [name string]`

IP subnet, IPX network, and protocol-based VLANs within port-based VLANs

If you plan to use port-based VLANs in conjunction with protocol-based VLANs, you must create the port-based VLANs first. Once you create a port-based VLAN, then you can assign Layer 3 protocol VLANs within the boundaries of the port-based VLAN. Generally, you create port-based VLANs to allow multiple separate STP domains.

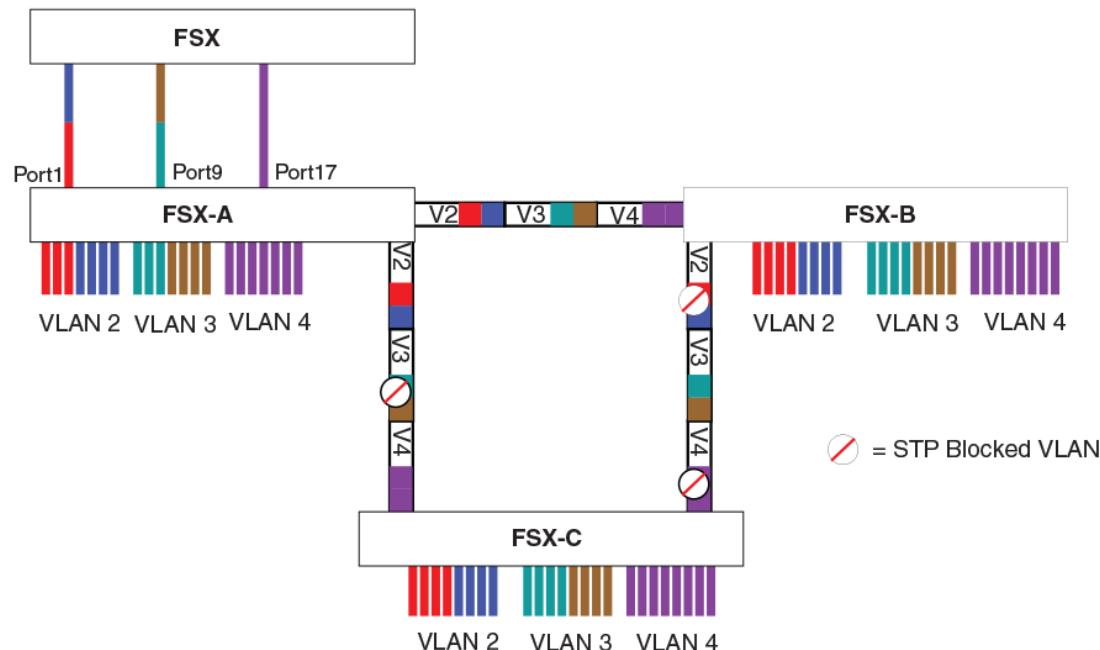
Suppose you need to provide three separate STP domains across an enterprise campus backbone. The first STP domain (VLAN 2) requires a set of ports at each Layer 2 Switch location to be statically mapped to IP only. No other protocols can enter the switches on this set of ports.

A second set of ports within STP domain VLAN 2 will be restricted to only IPX traffic. The IP and IPX protocol VLANs will overlap on Port 1 of FSX-A to support both protocols on the same router interface. The IP subnets and IPX network that span the two protocol VLANs will be determined by the FastIron router configuration. The IP and IPX Protocol VLANs ensure that only the ports included in the each Layer 3 protocol VLAN will see traffic from the FastIron router.

The second STP domain (VLAN 3) requires that half the ports in the domain are dedicated to IP subnet 10.1.1.0/24 and the other ports are dedicated to IPX network 1. Similar to VLAN 2, Port 9 from VLAN 3 will be used to carry this IP subnet and IPX network to the FastIron router. No other protocols will be allowed to enter the network on VLAN 3. Also, no IP packets with a source address on subnet 10.1.1.0/24 or IPX packets with a source address on network 1 will be allowed to enter the switches on VLAN 3.

There is no need to segment Layer 3 broadcast domains within the STP broadcast domain (VLAN 4). The FastIron router will dictate the IP subnets and IPX network that are on VLAN 4. There are no Layer 3 protocol restrictions on VLAN 4; however, the FastIron router is configured to only forward IP and IPX between STP domains.

FIGURE 80 More protocol-based VLANs



To configure the Layer 3 VLANs on the FSX Layer 2 Switches in the above figure, use the following procedure.

Configuring Layer 3 VLANs on FSX-A

Enter the following commands to configure FSX-A.

1. Create port-based VLAN 2 and assign the untagged and tagged ports that will participate in this VLAN.

```
device-A >en
device-A# config t
device-A(config)# vlan 2 name IP_IPX_Protocol
device-A(config-vlan-2)# untagged e1-to 8
device-A(config-vlan-2)# tagged e25 to 26
```

2. Enable STP and set the priority to force FSX-A to be the root bridge for VLAN 2.

```
device-A(config-vlan-2)# spanning-tree
device-A(config-vlan-2)# spanning-tree priority 500
device-A(config-vlan-2)#

```

3. Create the IP and IPX protocol-based VLANs and statically assign the ports within VLAN 2 that will be associated with each protocol-based VLAN.

```
device-A(config-vlan-2)# ip-proto name Red
device-A(config-vlan-ip-proto)# no dynamic
device-A(config-vlan-ip-proto)# static e1 to 4 e25 to 26
device-A(config-vlan-ip-proto)# exclude e5 to 8
device-A(config-vlan-ip-proto)# ipx-proto name Blue
device-A(config-vlan-ipx-proto)# no dynamic
device-A(config-vlan-ipx-proto)# static e1 e5 to 8 e25 to 26
device-A(config-vlan-ipx-proto)# exclude e2 to 4
```

4. To prevent machines with non-IP protocols from getting into the IP portion of VLAN 2, create another Layer 3 protocol VLAN to exclude all other protocols from the ports that contains the IP-protocol VLAN. To do so, enter the following commands.

```
device-A(config-vlan-ipx-proto)# other-proto name Block_other_proto
device-A(config-vlan-other-proto)# no dynamic
device-A(config-vlan-other-proto)# exclude e1 to 8
device-A(config-vlan-other-proto)#

```

5. Create port-based VLAN 3. Note that FSX-B will be the root for this STP domain, so you do not need to adjust the STP priority.

```
device-A(config-vlan-other-proto)# vlan 3 name IP-Sub_IPX-Net_Vlans
device-A(config-vlan-3)# untagged e9 to 16
device-A(config-vlan-3)# tagged e25 to 26
device-A(config-vlan-3)# spanning-tree
device-A(config-vlan-3)#

```

6. Create IP subnet VLAN 10.1.1.0/24, IPX network 1, and other-protocol VLANs

```
device-A(config-vlan-3)# ip-subnet 10.1.1.0/24 name Green
device-A(config-vlan-ip-subnet)# no dynamic
device-A(config-vlan-ip-subnet)# static e9 to 12 e25 to 26
device-A(config-vlan-ip-subnet)# exclude e13 to 16
device-A(config-vlan-ip-subnet)# ipx-net 1 ethernet_802.3 name Brown
device-A(config-vlan-ipx-network)# no dynamic
device-A(config-vlan-ipx-network)# static e9 e13 to 16 e25 to 26
device-A(config-vlan-ipx-network)# exclude e10 to 12
device-A(config-vlan-ipx-network)# other-proto name Block_other_proto
device-A(config-vlan-other-proto)# no dynamic
device-A(config-vlan-other-proto)# exclude e9 to 16
device-A(config-vlan-other-proto)#

```

7. Configure the last port-based VLAN 4. You need to set the STP priority for this VLAN because FSX-A will be the root bridge for this VLAN. Because you do not need to partition this STP domain into multiple Layer 3 broadcast domains, this is the only configuration required for VLAN 4.

```
device-A(config-vlan-other-proto)# vlan 4 name Purple_ALL-Protocols
device-A(config-vlan-4)# untagged e17 to 24
device-A(config-vlan-4)# tagged e25 to 26
device-A(config-vlan-4)# spanning-tree
device-A(config-vlan-4)# spanning-tree priority 500
device-A(config-vlan-4)#

```

Configuring Layer 3 VLANs on FSX-B

Enter the following commands to configure FSX-B.

```
device# config t
device(config)# host device-B
device-B(config)#vlan 2 name IP_IPX_Protocol
device-B(config-vlan-2)# untagged e1 to 8
device-B(config-vlan-2)# tagged e25 to 26
device-B(config-vlan-2)# spanning-tree
device-B(config-vlan-2)# ip-proto name Red
device-B(config-vlan-ip-proto)# no dynamic
device-B(config-vlan-ip-proto)# static e1 to 4 e25 to 26
device-B(config-vlan-ip-proto)# exclude e5 to 8
device-B(config-vlan-ip-proto)# ipx-proto name Blue
device-B(config-vlan-ipx-proto)# no dynamic
device-B(config-vlan-ipx-proto)# static e5 to 8 e25 to 26
```

```

device-B(config-vlan-ipx-proto)# exclude e1 to 4
device-B(config-vlan-other-proto)# vlan 3 name IP-Sub_IPX-Net_VLANs
device-B(config-vlan-3)# untagged e9 to 16
device-B(config-vlan-3)# tagged e25 to 26
device-B(config-vlan-3)# spanning-tree
device-B(config-vlan-3)# spanning-tree priority 500
device-B(config-vlan-3)# ip-sub 10.1.1.0/24 name Green
device-B(config-vlan-ip-subnet)# no dynamic
device-B(config-vlan-ip-subnet)# static e9 to 12 e25 to 26
device-B(config-vlan-ip-subnet)# exclude e13 to 16
device-B(config-vlan-ip-subnet)# ipx-net 1 ethernet_802.3 name Brown
device-B(config-vlan-ipx-network)# no dynamic
device-B(config-vlan-ipx-network)# static e13 to 16 e25 to 26
device-B(config-vlan-ipx-network)# exclude e9 to 12
device-B(config-vlan-ipx-network)# vlan 4 name Purple_ALL-Protocols
device-B(config-vlan-4)# untagged e17 to 24
device-B(config-vlan-4)# tagged e25 to 26
device-B(config-vlan-4)# spanning-tree

```

Configuring Layer 3 VLANs on FSX-C

Enter the following commands to configure FSX-C.

```

device# config t
device(config)# host device-C
device-C(config)# vlan 2 name IP_IPX_Protocol
device-C(config-vlan-2)# untagged e1 to 8
device-C(config-vlan-2)# tagged e25 to 26
device-C(config-vlan-2)# spanning-tree
device-C(config-vlan-2)# ip-proto name Red
device-C(config-vlan-ip-proto)# no dynamic
device-C(config-vlan-ip-proto)# static e1 to 4 e25 to 26
device-C(config-vlan-ip-proto)# exclude e5 to 8
device-C(config-vlan-ip-proto)# ipx-proto name Blue
device-C(config-vlan-ipx-proto)# no dynamic
device-C(config-vlan-ipx-proto)# static e5 to 8 e25 to 26
device-C(config-vlan-ipx-proto)# exclude e1 to 4
device-C(config-vlan-other-proto)# vlan 3 name IP-Sub_IPX-Net_VLANs
device-C(config-vlan-3)# untagged e9 to 16
device-C(config-vlan-3)# tagged e25 to 26
device-C(config-vlan-3)# spanning-tree
device-C(config-vlan-3)# ip-sub 10.1.1.0/24 name Green
device-C(config-vlan-ip-subnet)# no dynamic
device-C(config-vlan-ip-subnet)# static e9 to 12 e25 to 26
device-C(config-vlan-ip-subnet)# exclude e13 to 16
device-C(config-vlan-ip-subnet)# ipx-net 1 ethernet_802.3 name Brown
device-C(config-vlan-ipx-network)# no dynamic
device-C(config-vlan-ipx-network)# static e13 to 16 e25 to 26
device-C(config-vlan-ipx-network)# exclude e9 to 12
device-C(config-vlan-ipx-network)# vlan 4 name Purple_ALL-Protocols
device-C(config-vlan-4)# untagged e17 to 24
device-C(config-vlan-4)# tagged e25 to 26
device-C(config-vlan-4)# spanning-tree

```

IPv6 protocol VLAN configuration

You can configure a protocol-based VLAN as a broadcast domain for IPv6 traffic. When the Layer 3 Switch receives an IPv6 multicast packet (a packet with 06 in the version field and 0xFF as the beginning of the destination address), the Layer 3 Switch forwards the packet to all other ports.

NOTE

The Layer 3 Switch forwards all IPv6 multicast packets to all ports except the port that received the packet, and does not distinguish among subnet directed multicasts.

You can add the VLAN ports as static ports or dynamic ports. A static port is always an active member of the VLAN. Dynamic ports within any protocol VLAN age out after 10 minutes if no member protocol traffic is received on a port within the VLAN. The aged out port, however, remains as a candidate dynamic port for that VLAN. The port becomes active in the VLAN again if member protocol traffic is received on that port.

Once a port is re-activated, the aging out period for the port is reset to 20 minutes. Each time a member protocol packet is received by a candidate dynamic port (aged out port) the port becomes active again and the aging out period is reset for 20 minutes.

NOTE

You can disable VLAN membership aging of dynamically added ports. Refer to [Disabling membership aging of dynamic VLAN ports](#) on page 376).

To configure an IPv6 VLAN, enter commands such as the following.

```
device(config)# vlan 2
device(config-vlan-2)# untagged ethernet 1/1 to 1/8
device(config-vlan-2)# ipv6-proto name V6
device(config-ipv6-subnet)# static ethernet 1/1 to 1/6
device(config-ipv6-subnet)# dynamic
```

The first two commands configure a port-based VLAN and add ports 1/1 - 1/8 to the VLAN. The remaining commands configure an IPv6 VLAN within the port-based VLAN. The **static** command adds ports 1/1 - 1/6 as static ports, which do not age out. The **dynamic** command adds the remaining ports, 1/7 - 1/8, as dynamic ports. These ports are subject to aging as described above.

Syntax: [no] **ipv6-proto** [name *string*]

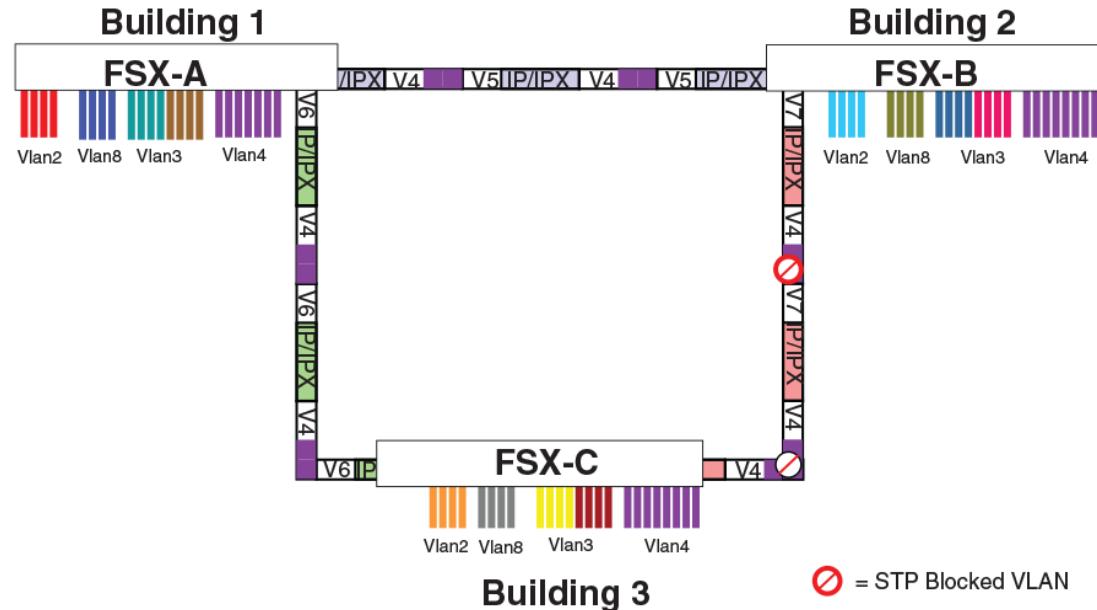
Routing between VLANs using virtual routing interfaces (Layer 3 Switches only)

Brocade Layer 3 Switches offer the ability to create a virtual routing interface within a Layer 2 STP port-based VLAN or within each Layer 3 protocol, IP subnet, or IPX network VLAN. This combination of multiple Layer 2 or Layer 3 broadcast domains, or both, and virtual routing interfaces are the basis for Brocade Communication Systems, Inc' very powerful Integrated Switch Routing (ISR) technology. ISR is very flexible and can solve many networking problems. The following example is meant to provide ideas by demonstrating some of the concepts of ISR.

Suppose you want to move routing out to each of three buildings in a network. Remember that the only protocols present on VLAN 2 and VLAN 3 are IP and IPX. Therefore, you can eliminate tagged ports 25 and 26 from both VLAN 2 and VLAN 3 and create new tagged port-based VLANs to support separate IP subnets and IPX networks for each backbone link.

You also need to create unique IP subnets and IPX networks within VLAN 2 and VLAN 3 at each building. This will create a fully routed IP and IPX backbone for VLAN 2 and VLAN 3. However, VLAN 4 has no protocol restrictions across the backbone. In fact there are requirements for NetBIOS and DecNet to be bridged among the three building locations. The IP subnet and IPX network that exists within VLAN 4 must remain a flat Layer 2 switched STP domain. You enable routing for IP and IPX on a virtual routing interface only on FSX-A. This will provide the flat IP and IPX segment with connectivity to the rest of the network. Within VLAN 4 IP and IPX will follow the STP topology. All other IP subnets and IPX networks will be fully routed and have use of all paths at all times during normal operation.

The following figure shows the configuration described above.

FIGURE 81 Routing between protocol-based VLANs

To configure the Layer 3 VLANs and virtual routing interfaces on the FSX Layer 3 Switch in the above figure, use the following procedure.

Configuring Layer 3 VLANs and virtual routing interfaces on the FSX-A

Enter the following commands to configure FSX-A. The following commands enable OSPF or RIP routing.

```
device>en
No password has been assigned yet...
device# configure terminal
device(config)# hostname FSX-A
device-A(config)# router ospf
device-A(config-ospf-router)# area 0.0.0.0 normal
Please save configuration to flash and reboot.
device-A(config-ospf-router)#

```

The following commands create the port-based VLAN 2. In the previous example, an external FSX defined the router interfaces for VLAN 2. With ISR, routing for VLAN 2 is done locally within each FSX. Therefore, there are two ways you can solve this problem. One way is to create a unique IP subnet and IPX network VLAN, each with its own virtual routing interface and unique IP or IPX address within VLAN 2 on each FSX. In this example, this is the configuration used for VLAN 3. The second way is to split VLAN 2 into two separate port-based VLANs and create a virtual router interface within each port-based VLAN. Later in this example, this second option is used to create a port-based VLAN 8 to show that there are multiple ways to accomplish the same task with ISR.

You also need to create the Other-Protocol VLAN within port-based VLAN 2 and 8 to prevent unwanted protocols from being Layer 2 switched within port-based VLAN 2 or 8. Note that the only port-based VLAN that requires STP in this example is VLAN 4. You will need to configure the rest of the network to prevent the need to run STP.

```
device-A(config-ospf-router)# vlan 2 name IP-Subnet_10.1.2.0/24
device-A(config-vlan-2)# untagged ethernet 1 to 4
device-A(config-vlan-2)# no spanning-tree
device-A(config-vlan-2)# router-interface ve1
device-A(config-vlan-2)# other-proto name block_other_protocols

```

```
device-A(config-vlan-other-proto)# no dynamic
device-A(config-vlan-other-proto)# exclude ethernet 1 to 4
```

Once you have defined the port-based VLAN and created the virtual routing interface, you need to configure the virtual routing interface just as you would configure a physical interface.

```
device-A(config-vlan-other-proto)# interface ve1
device-A(config-vif-1)# ip address 10.1.2.1/24
device-A(config-vif-1)# ip ospf area 0.0.0.0
```

Do the same thing for VLAN 8.

```
device-A(config-vif-1)# vlan 8 name IPX_Network2
device-A(config-vlan-8)# untagged ethernet 5 to 8
device-A(config-vlan-8)# no spanning-tree
device-A(config-vlan-8)# router-interface ve 2
device-A(config-vlan-8)# ipx-network 2 ethernet_802.3
device-A(config-vlan-8)# other-proto name block-other-protocols
device-A(config-vlan-other-proto)# no dynamic
device-A(config-vlan-other-proto)# exclude ethernet 5 to 8
device-A(config-vlan-other-proto)# interface ve2
device-A(config-vif-1)# ip address 10.1.2.2/24
device-A(config-vif-1)# ip ospf area 0.0.0.0
```

The next thing you need to do is create VLAN 3. This is very similar to the previous example with the addition of virtual routing interfaces to the IP subnet and IPX network VLANs. Also there is no need to exclude ports from the IP subnet and IPX network VLANs on the router.

```
device-A(config-vif-2)# vlan 3 name IP_Sub_&_IPX_Net_VLAN
device-A(config-vlan-3)# untagged ethernet 9 to 16
device-A(config-vlan-3)# no spanning-tree
device-A(config-vlan-3)# ip-subnet 10.1.1.0/24
device-A(config-vlan-ip-subnet)# static ethernet 9 to 12
device-A(config-vlan-ip-subnet)# router-interface ve3
device-A(config-vlan-ip-subnet)# ipx-network 1 ethernet_802.3
device-A(config-vlan-ipx-network)# static ethernet 13 to 16
device-A(config-vlan-ipx-network)# router-interface ve4
device-A(config-vlan-ipx-network)# other-proto name block-other-protocols
device-A(config-vlan-other-proto)# exclude ethernet 9 to 16
device-A(config-vlan-other-proto)# no dynamic
device-A(config-vlan-other-proto)# interface ve 3
device-A(config-vif-3)# ip addr 10.1.1.1/24
device-A(config-vif-3)# ip ospf area 0.0.0.0
```

Now configure VLAN 4. Remember this is a flat segment that, in the previous example, obtained its IP default gateway and IPX router services from an external FSX. In this example, FSX-A will provide the routing services for VLAN 4. You also want to configure the STP priority for VLAN 4 to make FSX-A the root bridge for this VLAN.

```
device-A(config-vif-4)# vlan 4 name Bridged_ALL_Proocols
device-A(config-vlan-4)# untagged ethernet 17 to 24
device-A(config-vlan-4)# tagged ethernet 25 to 26
device-A(config-vlan-4)# spanning-tree
device-A(config-vlan-4)# spanning-tree priority 500
device-A(config-vlan-4)# ipx network 3 ethernet_802.3
device-A(config-vlan-ipx-network)# exit
device-A(config-vlan-4)# router-interface ve5
device-A(config-vlan-4)# interface ve5
device-A(config-vif-5)# ip address 10.1.3.1/24
device-A(config-vif-5)# ip ospf area 0.0.0.0
```

It is time to configure a separate port-based VLAN for each of the routed backbone ports (Ethernet 25 and 26). If you do not create a separate tagged port-based VLAN for each point-to-point backbone link, you need to include tagged interfaces for Ethernet 25 and 26 within VLANs 2, 3, and 8. This type of configuration makes the entire backbone a single STP domain for each VLAN 2, 3, and 8. This is the configuration used in the example in [Configuring IP subnet, IPX network and protocol-based VLANs](#) on page 365. In this scenario, the virtual routing interfaces within port-based VLANs 2, 3, and 8

will be accessible using only one path through the network. The path that is blocked by STP is not available to the routing protocols until it is in the STP FORWARDING state.

```
device-A(config-vif-5)# vlan 5 name Rtr_BB_to_Bldg.2
device-A(config-vlan-5)# tagged ethernet 25
device-A(config-vlan-5)# no spanning-tree
device-A(config-vlan-5)# router-interface ve6
device-A(config-vlan-5)# vlan 6 name Rtr_BB_to_Bldg.3
device-A(config-vlan-6)# tagged ethernet 26
device-A(config-vlan-6)# no spanning-tree
device-A(config-vlan-6)# ipx-network 4 ethernet_802.3
device-A(config-vlan-6)# exit
device-A(config-vlan-6)# router-interface ve7
device-A(config-vlan-6)# interface ve6
device-A(config-vif-6)# ip addr 10.1.4.1/24
device-A(config-vif-6)# ip ospf area 0.0.0.0
device-A(config-vif-6)# interface ve7
device-A(config-vif-7)# ip addr 10.1.5.1/24
device-A(config-vif-7)# ip ospf area 0.0.0.0
device-A(config-vif-7)#

```

This completes the configuration for FSX-A. The configuration for FSX-B and C is very similar except for a few issues which are as follows:

- IP subnets and IPX networks configured on FSX-B and FSX-C must be unique across the entire network, except for the backbone port-based VLANs 5, 6, and 7 where the subnet is the same but the IP address must change.
- There is no need to change the default priority of STP within VLAN 4.
- There is no need to include a virtual router interface within VLAN 4.
- The backbone VLAN between FSX-B and FSX-C must be the same at both ends and requires a new VLAN ID. The VLAN ID for this port-based VLAN is VLAN 7.

Configuring Layer 3 VLANs and virtual routing interfaces for FSX-B

Enter the following commands to configure FSX-B.

```
device> en
No password has been assigned yet...
device# config t
device(config)# hostname FSX-B
device-B(config)# router ospf
device-B(config-ospf-router)# area 0.0.0.0 normal
device-B(config-ospf-router)# vlan 2 name IP-Subnet_10.1.6.0/24
device-B(config-vlan-2)# untagged ethernet 1 to 4
device-B(config-vlan-2)# no spanning-tree
device-B(config-vlan-2)# router-interface ve1
device-B(config-vlan-2)# other-proto name block-other-protocols
device-B(config-vlan-other-proto)# no dynamic
device-B(config-vlan-other-proto)# exclude ethernet 1 to 4
device-B(config-vlan-other-proto)# interface ve1
device-B(config-vif-1)# ip addr 10.1.6.1/24
device-B(config-vif-1)# ip ospf area 0.0.0.0
device-B(config-vif-1)# vlan 8 name IPX_Network6
device-B(config-vlan-8)# untagged ethernet 5 to 8
device-B(config-vlan-8)# no span
device-B(config-vlan-8)# router-interface ve2
device-B(config-vlan-8)# ipx-net 6 ethernet_802.3
device-B(config-vlan-ipx-network)# exit
device-B(config-vlan-8)# other-proto name block-other-protocols
device-B(config-vlan-other-proto)# no dynamic
device-B(config-vlan-other-proto)# exclude ethernet 5 to 8
device-B(config-vlan-other-proto)# interface ve2
device-B(config-vif-2)# vlan 3 name IP_Sub & IPX_Net_VLAN
device-B(config-vlan-3)# untagged ethernet 9 to 16
device-B(config-vlan-3)# no spanning-tree
device-B(config-vlan-3)# ip-subnet 10.1.7.0/24
device-B(config-vlan-ip-subnet)# static ethernet 9 to 12
device-B(config-vlan-ip-subnet)# router-interface ve3
device-B(config-vlan-ip-subnet)# ipx-network 7 ethernet_802.3
device-B(config-vlan-ipx-network)# static ethernet 13 to 16

```

```

device-B(config-vlan-ipx-network)# router-interface ve4
device-B(config-vlan-ipx-network)# other-proto name block-other-protocols
device-B(config-vlan-other-proto)# exclude ethernet 9 to 16
device-B(config-vlan-other-proto)# no dynamic
device-B(config-vlan-other-proto)# interface ve 3
device-B(config-vif-3)# ip addr 10.1.7.1/24
device-B(config-vif-3)# ip ospf area 0.0.0.0
device-B(config-vif-3)# interface ve4
device-B(config-vif-4)# vlan 4 name Bridged_ALL_Protocols
device-B(config-vlan-4)# untagged ethernet 17 to 24
device-B(config-vlan-4)# tagged ethernet 25 to 26
device-B(config-vlan-4)# spanning-tree
device-B(config-vlan-4)# ipx-network 4 ethernet_802.3
device-B(config-vlan-ipx-network)# exit
device-B(config-vlan-4)# vlan 5 name Rtr_BB_to_Bldg.1
device-B(config-vlan-5)# tagged ethernet 25
device-B(config-vlan-5)# no spanning-tree
device-B(config-vlan-5)# router-interface ve5
device-B(config-vlan-5)# vlan 7 name Rtr_BB_to_Bldg.3
device-B(config-vlan-7)# tagged ethernet 26
device-B(config-vlan-7)# no spanning-tree
device-B(config-vlan-7)# router-interface ve6
device-B(config-vlan-7)# interface ve5
device-B(config-vif-5)# ip addr 10.1.4.2/24
device-B(config-vif-5)# ip ospf area 0.0.0.0
device-B(config-vif-5)# interface ve6
device-B(config-vif-6)# ip addr 10.1.8.1/24
device-B(config-vif-6)# ip ospf area 0.0.0.0
device-B(config-vif-6)#

```

Configuring Layer 3 VLANs and virtual routing interfaces for FSX-C

Enter the following commands to configure FSX-C.

```

device> en
No password has been assigned yet...
device# config t
device(config)# hostname FSX -C
device-C(config)# router ospf
device-C(config-ospf-router)# area 0.0.0.0 normal
device-C(config-ospf-router)# vlan 2 name IP-Subnet_10.1.9.0/24
device-C(config-vlan-2)# untagged ethernet 1 to 4
device-C(config-vlan-2)# no spanning-tree
device-C(config-vlan-2)# router-interface ve1
device-C(config-vlan-2)# other-proto name block-other-protocols
device-C(config-vlan-other-proto)# no dynamic
device-C(config-vlan-other-proto)# exclude ethernet 1 to 4
device-C(config-vlan-other-proto)# interface ve1
device-C(config-vif-1)# ip addr 10.1.9.1/24
device-C(config-vif-1)# ip ospf area 0.0.0.0
device-C(config-vif-1)# vlan 8 name IPX_Network9
device-C(config-vlan-8)# untagged ethernet 5 to 8
device-C(config-vlan-8)# no span
device-C(config-vlan-8)# ipx-net 9 ethernet_802.3
device-C(config-vlan-ipx-network)# exit
device-C(config-vlan-8)# router-interface ve2
device-C(config-vlan-8)# other-proto name block-other-protocols
device-C(config-vlan-other-proto)# no dynamic
device-C(config-vlan-other-proto)# exclude ethernet 5 to 8
device-C(config-vlan-other-proto)# interface ve2
device-C(config-vif-1)# ip addr 10.1.9.2/24
device-C(config-vif-1)# ip ospf area 0.0.0.0
device-C(config-vif-2)# vlan 3 name IP_Sub_&_IPX_Net_VLAN
device-C(config-vlan-3)# untagged ethernet 9 to 16
device-C(config-vlan-3)# no spanning-tree
device-C(config-vlan-3)# ip-subnet 10.1.10.0/24
device-C(config-vlan-ip-subnet)# static ethernet 9 to 12
device-C(config-vlan-ip-subnet)# router-interface ve3
device-C(config-vlan-ip-subnet)# ipx-network 10 ethernet_802.3
device-C(config-vlan-ipx-network)# static ethernet 13 to 16
device-C(config-vlan-ipx-network)# router-interface ve4
device-C(config-vlan-ipx-network)# other-proto name block-other-protocols
device-C(config-vlan-other-proto)# exclude ethernet 9 to 16
device-C(config-vlan-other-proto)# no dynamic

```

```

device-C(config-vlan-other-proto)# interface ve 3
device-C(config-vif-3)# ip addr 10.1.10.1/24
device-C(config-vif-3)# ip ospf area 0.0.0.0
device-C(config-vif-3)# interface ve4
device-C(config-vif-4)# vlan 4 name Bridged ALL Protocols
device-C(config-vlan-4)# untagged ethernet 17 to 24
device-C(config-vlan-4)# tagged ethernet 25 to 26
device-C(config-vlan-4)# spanning-tree
device-C(config-vlan-4)# vlan 7 name Rtr_BB_to_Bldg.2
device-C(config-vlan-7)# tagged ethernet 25
device-C(config-vlan-7)# no spanning-tree
device-C(config-vlan-7)# ipx-network 8 ethernet_802.3
device-C(config-vlan-ip-subnet)# exit
device-C(config-vlan-7)# router-interface ve5
device-C(config-vlan-7)# vlan 6 name Rtr_BB_to_Bldg.1
device-C(config-vlan-6)# tagged ethernet 26
device-C(config-vlan-6)# no spanning-tree
device-C(config-vlan-6)# router-interface ve6
device-C(config-vlan-6)# ipx network 5 ethernet_802.3
device-C(config-vlan-6)# interface ve5
device-C(config-vif-5)# ip addr 10.1.8.2/24
device-C(config-vif-5)# ip ospf area 0.0.0.0
device-C(config-vif-5)# interface ve6
device-C(config-vif-6)# ip addr 10.1.5.2/24
device-C(config-vif-6)# ip ospf area 0.0.0.0
device-C(config-vif-6)#

```

Configuring protocol VLANs with dynamic ports

The configuration examples for protocol VLANs in the sections above show how to configure the VLANs using static ports. You also can configure the following types of protocol VLANs with dynamic ports:

- AppleTalk protocol
- IP protocol
- IPX protocol
- IP subnet
- IPX network

NOTE

The software does not support dynamically adding ports to AppleTalk cable VLANs. Conceptually, an AppleTalk cable VLAN consists of a single network cable, connected to a single port. Therefore, dynamic addition and removal of ports is not applicable.

NOTE

You cannot route to or from protocol VLANs with dynamically added ports.

Aging of dynamic ports

When you add the ports to the VLAN, the software automatically adds them all to the VLAN. However, dynamically added ports age out. If the age time for a dynamic port expires, the software removes the port from the VLAN. If that port receives traffic for the IP subnet or IPX network, the software adds the port to the VLAN again and starts the aging timer over. Each time the port receives traffic for the VLAN's IP subnet or IPX network, the aging timer starts over.

NOTE

You can disable VLAN membership aging of dynamically added ports. Refer to [Disabling membership aging of dynamic VLAN ports](#) on page 376).

Dynamic ports within any protocol VLAN age out after 10 minutes, if no member protocol traffic is received on a port within the VLAN. The aged out port, however, remains as a candidate dynamic port for that VLAN. The port becomes active in the VLAN again if member protocol traffic is received on that port.

Once a port is re-activated, the aging out period for the port is reset to 20 minutes. Each time a member protocol packet is received by a candidate dynamic port (aged out port) the port becomes active again and the aging out period is reset for 20 minutes.

Disabling membership aging of dynamic VLAN ports

You can disable VLAN membership aging of ports that are dynamically assigned to protocol or subnet-based VLANs. This feature resolves the connectivity issue that may occur in certain configurations when protocol or subnet VLANs are configured with dynamic port membership.

NOTE

This issue does not occur with statically assigned VLAN memberships. Thus, enable this feature only if your configuration includes dynamically assigned VLAN memberships for protocol or subnet VLANs.

To enable this feature, enter commands such as the following.

```
device(config)# vlan 10 by port
device(config-vlan-10)# interface ethernet 1/1 to 1/5
device(config-vlan-10)# ip-proto name IP_Prot_VLAN
device(config-vlan-ip-proto)# no-dynamic-aging
device(config-vlan-ip-proto)# write memory
```

These commands create an IP protocol VLAN and disable the VLAN membership aging of ports that are dynamically assigned to the protocol VLAN.

Syntax: [no] no-dynamic-aging

Enter the **no** form of the command to disable this feature after it has been enabled.

By default, VLAN membership of dynamically assigned ports will age out after a period of time if no packets belonging to that protocol or subnet VLAN are received by the CPU.

The output of the **show running-config** command indicates if the no-dynamic-aging feature is enabled for a specific protocol or subnet VLAN.

Configuration guidelines for membership aging of dynamic VLAN ports

- You cannot dynamically add a port to a protocol VLAN if the port has any routing configuration parameters. For example, the port cannot have a virtual routing interface, IP subnet address, IPX network address, or AppleTalk network address configured on it.
- Once you dynamically add a port to a protocol VLAN, you cannot configure routing parameters on the port.
- Dynamic VLAN ports are not required or supported on AppleTalk cable VLANs.
- When protocol VLANs with dynamic ports are configured, the output of the **show running-config** command in the Router image will show the "dynamic" keyword. In the Switch image, the keyword is not shown in the output of the **show running-config** command.

NOTE

In the Switch image, all the ports are dynamic ports by-default, so the **dynamic** command does not appear in the **show running-config** command output. If you configure the **no dynamic** command, it will appear in the output of the **show running-config** command. Similarly in Router image, no ports are dynamic by-default, so the **no dynamic** command does not appear in the output of the **show running-config** command. If you configure the **dynamic** command, it will appear in the output of the **show running-config** command.

Configuring an IP, IPX, or AppleTalk Protocol VLAN with Dynamic Ports

To configure an IP, IPX, or AppleTalk protocol VLAN with dynamic ports, use the following method.

To configure port-based VLAN 10, then configure an IP protocol VLAN within the port-based VLAN with dynamic ports, enter the following commands such as the following.

```
device(config)# vlan 10 by port
device(config-vlan=10)# untagged ethernet 1/1 to 1/6
added untagged port eth1/1 to 1/6 to port-vlan 30.
device(config-vlan=10)# ip-proto name IP_Prot_VLAN
device(config-vlan-ip-proto)# dynamic
device(config)# write memory
```

Syntax: **vlan** *vlan-id* **name** *string* [**by port**]

Syntax: **untagged ethernet** [*slotnum!*/*portnum* **to** [*slotnum!*/*portnum*]]

or

Syntax: **untagged ethernet** [*slotnum!*/*portnum*] **ethernet** [*slotnum!*/*portnum*]

NOTE

Use the first **untagged** command for adding a range of ports. Use the second command for adding separate ports (not in a range).

Syntax: **ip-proto** [**name** *string*]

Syntax: **ipx-proto** [**name** *string*]

Syntax: **appletalk-cable-vlan** *num* [**name** *string*]

Syntax: **dynamic**

The procedure is similar for IPX and AppleTalk protocol VLANs. Enter **ipx-proto** or **atalk-proto** instead of **ip-proto**.

Configuring an IP subnet VLAN with dynamic ports

To configure port-based VLAN 10, then configure an IP subnet VLAN within the port-based VLAN with dynamic ports, enter commands such as the following.

```
device(config)# vlan 10 name IP_VLAN by port
device(config-vlan=10)# untagged ethernet 1/1 to 1/6
added untagged port eth1/1 to 1/6 to port-vlan 10.
device(config-vlan=10)# ip-subnet 10.1.1.0/24 name Mktg-LAN
device(config-vlan-ip-subnet)# dynamic
device(config)# write memory
```

These commands create a port-based VLAN on chassis ports 1/1 - 1/6 named "Mktg-LAN", configure an IP subnet VLAN within the port-based VLAN, and then add ports from the port-based VLAN dynamically.

Syntax: **vlan** *vlan-id* name *string* [**by port**]

Syntax: **untagged ethernet** [*slotnum*]*portnum* to [*slotnum*]*portnum*

or

Syntax: **untagged ethernet** [*slotnum*]*portnum* **ethernet** [*slotnum*]*portnum*

NOTE

Use the first **untagged** command for adding a range of ports. Use the second command for adding separate ports (not in a range).

Syntax: **ip-subnet** *ip-addr* *ip-mask*[**name** *string*]

or

Syntax: **ip-subnet** *ip-addr/mask-bits* [**name** *string*]

Syntax: **dynamic**

Configuring an IPX network VLAN with dynamic ports

To configure port-based VLAN 20, then configure an IPX network VLAN within the port-based VLAN with dynamic ports, enter commands such as the following.

```
device(config)# vlan 20 name IPX_VLAN by port
device(config-vlan-10)# untagged ethernet 2/1 to 2/6
added untagged port ethe 2/1 to 2/6 to port-vlan 20.
device(config-vlan-10)# ipx-network abcd ethernet_ii name Eng-LAN
device(config-vlan-ipx-network)# dynamic
device(config)# write memory
```

These commands create a port-based VLAN on chassis ports 2/1 - 2/6 named "Eng-LAN", configure an IPX network VLAN within the port-based VLAN, and then add ports from the port-based VLAN dynamically.

Syntax: **vlan** *vlan-id* name *string* [**by port**]

Syntax: **untagged ethernet** [*slotnum*]*portnum* to [*slotnum*]*portnum*

or

Syntax: **untagged ethernet** [*slotnum*]*portnum* **ethernet** [*slotnum*]*portnum*

NOTE

Use the first **untagged** command for adding a range of ports. Use the second command for adding separate ports (not in a range).

Syntax: **ipx-network** *network-addr* **ethernet_ii|ethernet_802.2 | ethernet_802.3 | ethernet_snap [**name** *string*]**

Syntax: **dynamic**

Configuring uplink ports within a port-based VLAN

You can configure a subset of the ports in a port-based VLAN as uplink ports. When you configure uplink ports in a port-based VLAN, the device sends all broadcast and unknown-unicast traffic from a port in the VLAN to the uplink ports, but not to other ports within the VLAN. Thus, the uplink ports provide tighter broadcast control within the VLAN.

This uplink port feature behaves the same as the private VLAN (PVLAN) feature, but with the ability to support tagged ports. This feature also supports two PVLAN modes: the Primary ports (uplink ports) and Isolated ports (host ports).

For example, if two ports within a port-based VLAN are Gbps ports attached to the network and the other ports are 10/100 ports attached to clients, you can configure the two ports attached to the network as uplink ports. In this configuration, broadcast and unknown-unicast traffic in the VLAN does not go to all ports. The traffic goes only to the uplink ports. The clients on the network do not receive broadcast and unknown-unicast traffic from other ports, including other clients.

Configuration considerations for uplink ports within a port-based VLAN

- When this feature is enabled, flooded traffic (unknown unicast, unregistered multicast, and broadcast traffic) is software forwarded. ACL-based rate limit is non-functional with uplink switch configuration.
- This feature should not be enabled with protocol VLANs or PVLANS in the same VLAN.
- CPU spike is expected when this feature is enabled.

Configuration syntax for uplink ports within a port-based VLAN

To configure a port-based VLAN containing uplink ports, enter commands such as the following.

```
Device(config)# vlan 10 by port
Device(config-vlan-10)# untagged ethernet 1/1 to 1/24
Device(config-vlan-10)# untagged ethernet 2/1 to 2/2
Device(config-vlan-10)# uplink-switch ethernet 2/1 to 2/2
```

Syntax: [no] **uplink-switch ethernet slotnum/portnum [to slotnum/portnum | ethernet slotnum/portnum]**

In this example, 24 ports on a 10/100 module and two Gbps ports on a Gbps module are added to port-based VLAN 10. The two Gbps ports are then configured as uplink ports.

NOTE

This feature is not supported on the Brocade ICX 7750, Brocade ICX 7450, or Brocade ICX 7250.

IP subnet address on multiple port-based VLAN configuration

For a Brocade device to route between port-based VLANs, you must add a virtual routing interface to each VLAN. Generally, you also configure a unique IP subnet address on each virtual routing interface. For example, if you have three port-based VLANs, you add a virtual routing interface to each VLAN, then add a separate IP subnet address to each virtual routing interface. The IP address on each of the virtual routing interfaces must be in a separate subnet. The Brocade device routes Layer 3 traffic between the subnets using the subnet addresses.

NOTE

This feature applies only to Layer 3 Switches.

NOTE

Before using the method described in this section, refer to [VLAN groups and virtual routing interface group](#) on page 383. You might be able to achieve the results you want using the methods in that section instead.

The following figure shows an example of this type of configuration.

FIGURE 82 Multiple port-based VLANs with separate protocol addresses

VLAN 2

VLAN 3

VLAN 4

FSX Switch

VLAN 2
VE 1
-IP 10.0.0.1/24

VLAN 3
VE 2
-IP 10.0.1.1/24

VLAN 4
VE 3
-IP 10.0.2.1/24

As shown in this example, each VLAN has a separate IP subnet address. If you need to conserve IP subnet addresses, you can configure multiple VLANs with the same IP subnet address, as shown in the following figure.

FIGURE 83 Multiple port-based VLANs with the same protocol address

VLAN 2

.....
VLAN 3-----
VLAN 4

Each VLAN still requires a separate virtual routing interface. However, all three VLANs now use the same IP subnet address.

In addition to conserving IP subnet addresses, this feature allows containment of Layer 2 broadcasts to segments within an IP subnet. For ISP environments where the same IP subnet is allocated to different customers, placing each customer in a separate VLAN allows all customers to share the IP subnet address, while at the same time isolating them from one another Layer 2 broadcasts.

NOTE

You can provide redundancy to an IP subnet address that contains multiple VLANs using a pair of Brocade Layer 3 Switches configured for Brocade VRRP (Virtual Router Redundancy Protocol).

The Brocade device performs proxy Address Resolution Protocol (ARP) for hosts that want to send IP traffic to hosts in other VLANs that are sharing the same IP subnet address. If the source and destination hosts are in the same VLAN, the Brocade device does not need to use ARP:

- If a host attached to one VLAN sends an ARP message for the MAC address of a host in one of the other VLANs using the same IP subnet address, the Brocade device performs a proxy ARP on behalf of the other host. The Brocade device then replies to the ARP by sending the virtual routing interface MAC address. The Brocade device uses the same MAC address for all virtual routing interfaces. When the host that sent the ARP then sends a unicast packet addressed to the virtual

routing interface MAC address, the device switches the packet on Layer 3 to the destination host on the VLAN.

NOTE

If the Brocade device ARP table does not contain the requested host, the Brocade device forwards the ARP request on Layer 2 to the same VLAN as the one that received the ARP request. Then the device sends an ARP for the destination to the other VLANs that are using the same IP subnet address.

- If the destination is in the same VLAN as the source, the Brocade device does not need to perform a proxy ARP.

To configure multiple VLANs to use the same IP subnet address:

- Configure each VLAN, including adding tagged or untagged ports.
- Configure a separate virtual routing interface for each VLAN, but do not add an IP subnet address to more than one of the virtual routing interfaces.
- Configure the virtual routing interfaces that do not have the IP subnet address to "follow" the virtual routing interface that does have the address.

To configure the VLANs shown in [Figure 83](#), you could enter the following commands.

```
device(config)# vlan 1 by port
device(config-vlan-1)# untagged ethernet 1/1
device(config-vlan-1)# tagged ethernet 1/8
device(config-vlan-1)# router-interface ve 1
```

Syntax: router-interface ve number

The commands above configure port-based VLAN 1. The VLAN has one untagged port (1/1) and a tagged port (1/8). In this example, all three VLANs contain port 1/8 so the port must be tagged to allow the port to be in multiple VLANs. You can configure VLANs to share a Layer 3 protocol interface regardless of tagging. A combination of tagged and untagged ports is shown in this example to demonstrate that sharing the interface does not change other VLAN features.

Notice that each VLAN still requires a unique virtual routing interface.

The following commands configure port-based VLANs 2 and 3.

```
device(config-vlan-1)# vlan 2 by port
device(config-vlan-2)# untagged ethernet 1/2
device(config-vlan-2)# tagged ethernet 1/8
device(config-vlan-2)# router-interface ve 2
device(config-vlan-2)# vlan 3 by port
device(config-vlan-3)# untagged ethernet 1/5 to 1/6
device(config-vlan-3)# tagged ethernet 1/8
device(config-vlan-3)# router-interface ve 3
```

The following commands configure an IP subnet address on virtual routing interface 1.

```
device(config-vlan-3)# interface ve 1
device(config-vif-1)# ip address 10.0.0.1/24
```

The following commands configure virtual routing interfaces 2 and 3 to "follow" the IP subnet address configured on virtual routing interface 1.

```
device(config-vif-1)# interface ve 2
device(config-vif-2)# ip follow ve 1
device(config-vif-2)# interface ve 3
device(config-vif-3)# ip follow ve 1
```

NOTE

Because virtual routing interfaces 2 and 3 do not have their own IP subnet addresses but instead are "following" virtual routing interface a IP address, you still can configure an IPX or AppleTalk interface on virtual routing interfaces 2 and 3.

VLAN groups and virtual routing interface group

To simplify configuration when you have many VLANs with the same configuration, you can configure VLAN groups and virtual routing interface groups.

NOTE

VLAN groups are supported on Layer 3 Switches and Layer 2 Switches. Virtual routing interface groups are supported only on Layer 3 Switches.

When you create a VLAN group, the VLAN parameters you configure for the group apply to all the VLANs within the group. Additionally, you can easily associate the same IP subnet interface with all the VLANs in a group by configuring a virtual routing interface group with the same ID as the VLAN group.

- The VLAN group feature allows you to create multiple port-based VLANs with identical port members. Because the member ports are shared by all the VLANs within the group, you must add the ports as tagged ports. This feature not only simplifies VLAN configuration but also allows you to have a large number of identically configured VLANs in a startup-config file on the device flash memory module. Normally, a startup-config file with a large number of VLANs might not fit on the flash memory module. By grouping the identically configured VLANs, you can conserve space in the startup-config file so that it fits on the flash memory module.
- The virtual routing interface group feature is useful when you want to configure the same IP subnet address on all the port-based VLANs within a VLAN group. You can configure a virtual routing interface group only after you configure a VLAN group with the same ID. The virtual routing interface group automatically applies to the VLANs in the VLAN group that has the same ID and cannot be applied to other VLAN groups or to individual VLANs.

You can create up to 32 VLAN groups and 32 virtual routing interface groups. A virtual routing interface group always applies only to the VLANs in the VLAN group with the same ID.

NOTE

Depending on the size of the VLAN ID range you want to use for the VLAN group, you might need to allocate additional memory for VLANs. On Layer 3 Switches, if you allocate additional memory for VLANs, you also need to allocate the same amount of memory for virtual routing interfaces. This is true regardless of whether you use the virtual routing interface groups. To allocate additional memory, refer to [Allocating memory for more VLANs or virtual routing interfaces](#) on page 387.

NOTE

Configuring a VLAN group

To configure a VLAN group, enter commands such as the following.

```
device(config)# vlan-group 1 vlan 2 to 257
device(config-vlan-group-1)# tagged 1/1 to 1/2
```

The first command in this example begins configuration for VLAN group 1, and assigns VLANs 2 through 257 to the group. The second command adds ports 1/1 and 1/2 as tagged ports. Because all the VLANs in the group share the ports, you must add the ports as tagged ports.

Syntax: **vlan-group num vlan vlan-id to vlan-id**

Syntax: **tagged ethernet [slotnum/]portnum [to [slotnum/]portnum] ethernet [slotnum/]portnum]**

The **vlan-groupnum** parameter specifies the VLAN group ID and can be from 1 - 32. The **vlan**/**vlan-id** to **vlan-id** parameters specify a contiguous range (a range with no gaps) of individual VLAN IDs. Specify the low VLAN ID first and the high VLAN ID second. The command adds all of the specified VLANs to the VLAN group.

You can add up to 256 VLANs with the command at one time. To add more than 256 VLANs, enter separate commands. For example, to configure VLAN group 1 and add 512 VLANs to the group, enter the following commands.

```
device(config)# vlan-group 1 vlan 2 to 257
device(config-vlan-group-1)# add-vlan 258 to 513
```

NOTE

The device memory must be configured to contain at least the number of VLANs you specify for the higher end of the range. For example, if you specify 2048 as the VLAN ID at the high end of the range, you first must increase the memory allocation for VLANs to 2048 or higher. Additionally, on Layer 3 Switches, if you allocate additional memory for VLANs, you also need to allocate the same amount of memory for virtual routing interfaces, before you configure the VLAN groups. This is true regardless of whether you use the virtual routing interface groups. The memory allocation is required because the VLAN groups and virtual routing interface groups have a one-to-one mapping. Refer to [Allocating memory for more VLANs or virtual routing interfaces](#) on page 387.

If a VLAN within the range you specify is already configured, or if the range contains more than 256 VLANs, the CLI does not add the group but instead displays an error message.

```
device(config)#vlan-group 1 vlan 2 to 1000
VLAN group 1 is too big. Only 256 vlans are allowed at a time
```

In this case, create the group by specifying a valid contiguous range. Then add more VLANs to the group after the CLI changes to the configuration level for the group. See the following example.

```
device(config)#vlan-group 2 vlan 1000 to 1250
device(config-vlan-group-2)#add-vlan 1251 to 1500
device(config-vlan-group-2)#add-vlan 1501 to 1750
device(config-vlan-group-2)#add-vlan 1751 to 2000
```

You can add or remove individual VLANs or VLAN ranges from the VLAN group at configuration level. For example, if you want to add VLANs 1001 and 1002 to VLAN group 1 and remove VLANs 900 through 1000, enter the following commands.

```
device(config-vlan-group-1)# add-vlan 1001 to 1002
device(config-vlan-group-1)# remove-vlan 900 to 1000
```

Syntax: **add-vlan** *vlan-id* [**to** *vlan-id*]

Syntax: **remove-vlan** *vlan-id* [**to** *vlan-id*]

The *vlan-id* to *vlan-id* parameters specify a contiguous range (a range with no gaps) of individual VLAN IDs. Specify the low VLAN ID first and the high VLAN ID second. You can add or remove up to 256 VLANs at a time. To add or remove more than 256 VLANs, do so using separate commands. For example, to remove 512 VLANs from VLAN group 1, enter the following commands.

```
device(config-vlan-group-1)# remove-vlan 400 to 654
device(config-vlan-group-1)# remove-vlan 655 to 910
```

Displaying information about VLAN groups

To display VLAN group configuration information, use the **show vlan-group** command.

```
device# show vlan-group
vlan-group 1 vlan 2 to 20
  tagged ethe 1/1 to 1/2
!
vlan-group 2 vlan 21 to 40
  tagged ethe 1/1 to 1/2
!
```

Syntax: **show vlan-group** [*group-id*]

The *group-id* specifies a VLAN group. If you do not use this parameter, the configuration information for all the configured VLAN groups is displayed.

Configuring a virtual routing interface group

A virtual routing interface group allows you to associate the same IP subnet interface with multiple port-based VLANs. For example, if you associate a virtual routing interface group with a VLAN group, all the VLANs in the group have the IP interface of the virtual routing interface group.

Configuration notes and feature limitations for virtual routing interface group

- When you configure a virtual routing interface group, all members of the group have the same IP subnet address. This feature is useful in collocation environments where the device has many IP addresses and you want to conserve the IP address space.
- The **group-router-interface** command creates router interfaces for each VLAN in the VLAN group by using the VLAN IDs of each of the VLANs as the corresponding virtual interface number. Therefore, if a VLAN group contains VLAN IDs greater than the maximum virtual interface number allowed, the **group-router-interface** command will be rejected.

CLI syntax for virtual routing interface group

To configure a virtual routing interface group, enter commands such as the following.

```
device(config)# vlan-group 1
```

```
device(config-vlan-group-1) # group-router-interface  
device(config-vlan-group-1) # exit  
device(config)# interface group-ve 1  
device(config-vif-group-1) # ip address 10.10.10.1/24
```

These commands enable VLAN group 1 to have a group virtual routing interface, then configure virtual routing interface group 1. The software always associates a virtual routing interface group only with the VLAN group that has the same ID. In this example, the VLAN group ID is 1, so the corresponding virtual routing interface group also must have ID 1.

Syntax: **group-router-interface**

Syntax: **interface group-ve num**

Syntax: [no] **ip address ip-addr ip-mask [secondary]**

or

Syntax: [no] **ip address ip-addr/mask-bits [secondary]**

The **router-interface-group** command enables a VLAN group to use a virtual routing interface group. Enter this command at the configuration level for the VLAN group. This command configures the VLAN group to use the virtual routing interface group that has the same ID as the VLAN group. You can enter this command when you configure the VLAN group for the first time or later, after you have added tagged ports to the VLAN and so on.

The *num* parameter in the **interface group-venum** command specifies the ID of the VLAN group with which you want to associate this virtual routing interface group. The VLAN group must already be configured and enabled to use a virtual routing interface group. The software automatically associates the virtual routing interface group with the VLAN group that has the same ID. You can associate a virtual routing interface group only with the VLAN group that has the same ID.

NOTE

IPv6 is not supported with **group-ve**.

NOTE

FCX devices do not support ACLs with **group-ve**.

NOTE

FastIron devices support **group-ve** with OSPF, VRRP v2 and VRRP-E v2 protocols only.

The syntax and usage for the **ip address** command is the same as when you use the command at the interface level to add an IP interface.

Displaying the VLAN group and virtual routing interface group information

To verify configuration of VLAN groups and virtual routing interface groups, display the running-config file. If you have saved the configuration to the startup-config file, you also can verify the configuration by displaying the startup-config file. The following example shows the running-config information for the VLAN group and virtual routing interface group configured in the previous examples. The information appears in the same way in the startup-config file.

```
device# show running-config  
lines not related to the VLAN group omitted...  
vlan-group 1 vlan 2 to 20
```

```

add-vlan 1001 to 1002
tagged ethe 1/1 to 1/2
router-interface-group
lines not related to the virtual routing interface group omitted...
interface group-ve 1
  ip address 10.10.10.1 255.255.255.0

```

NOTE

If you have enabled display of subnet masks in CIDR notation, the IP address information is shown as follows: 10.10.10.1/24.

Allocating memory for more VLANs or virtual routing interfaces

Brocade Layer 2 and Layer 3 Switches support up to 4095 VLANs. In addition, Layer 3 switches support up to 512 virtual routing interfaces.

The number of VLANs and virtual routing interfaces supported on your product depends on the device and, for Chassis devices, the amount of DRAM on the management module. The following table lists the default and configurable maximum numbers of VLANs and virtual routing interfaces for Layer 2 and Layer 3 Switches. Unless otherwise noted, the values apply to both types of switches.

TABLE 55 VLAN and virtual routing interface support

VLANs	Virtual routing interfaces		
Default maximum	Configurable maximum	Default maximum	Configurable maximum
64	4094	255	512

NOTE

If many of your VLANs will have an identical configuration, you might want to configure VLAN groups and virtual routing interface groups after you increase the system capacity for VLANs and virtual routing interfaces. Refer to [VLAN groups and virtual routing interface group](#) on page 383.

Increasing the number of VLANs you can configure

NOTE

Although you can specify up to 4095 VLANs, you can configure only 4094 VLANs. VLAN ID 4094 is reserved for use by the Single Spanning Tree feature.

To increase the maximum number of VLANs you can configure, enter commands such as the following at the global CONFIG level of the CLI.

```

device(config)# system-max vlan 2048
device(config)# write memory
device(config)# end
device# reload

```

Syntax: **system-max vlan num**

The num parameter indicates the maximum number of VLANs. The range of valid values depends on the device you are configuring. Refer to [Table 55](#) on page 387.

Increasing the number of virtual routing interfaces you can configure

To increase the maximum number of virtual routing interfaces you can configure, enter commands such as the following at the global CONFIG level of the CLI.

```
device(config)# system-max virtual-interface 512  
device(config)# write memory  
device(config)# end  
device# reload
```

Syntax: system-max virtual-interface num

The num parameter indicates the maximum number of virtual routing interfaces. The range of valid values depends on the device you are configuring. Refer to [Table 55](#) on page 387.

Super-aggregated VLAN configuration

You can aggregate multiple VLANs within another VLAN. This feature allows you to construct Layer 2 paths and channels. This feature is particularly useful for Virtual Private Network (VPN) applications in which you need to provide a private, dedicated Ethernet connection for an individual client to transparently reach its subnet across multiple networks.

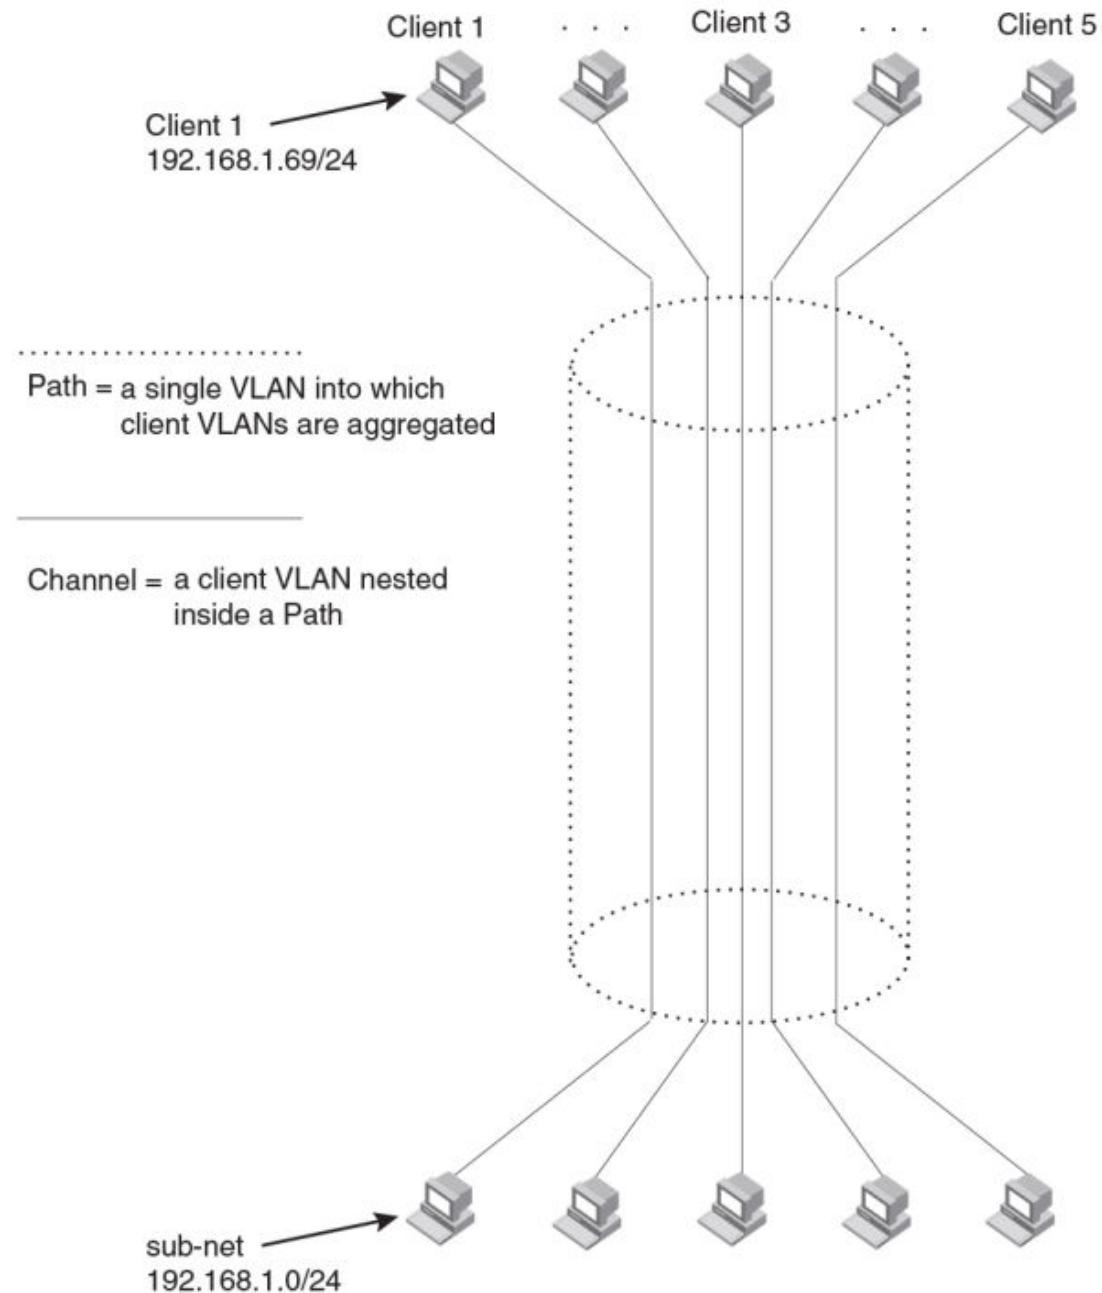
Conceptually, the paths and channels are similar to Asynchronous Transfer Mode (ATM) paths and channels. A path contains multiple channels, each of which is a dedicated circuit between two end points. The two devices at the end points of the channel appear to each other to be directly attached. The network that connects them is transparent to the two devices.

You can aggregate up to 4094 VLANs within another VLAN. This provides a total VLAN capacity on one Brocade device of 16,760,836 channels (4094 * 4094).

The devices connected through the channel are not visible to devices in other channels. Therefore, each client has a private link to the other side of the channel.

The feature allows point-to-point and point-to-multipoint connections.

The following figure shows a conceptual picture of the service that aggregated VLANs provide. Aggregated VLANs provide a path for multiple client channels. The channels do not receive traffic from other channels. Thus, each channel is a private link.

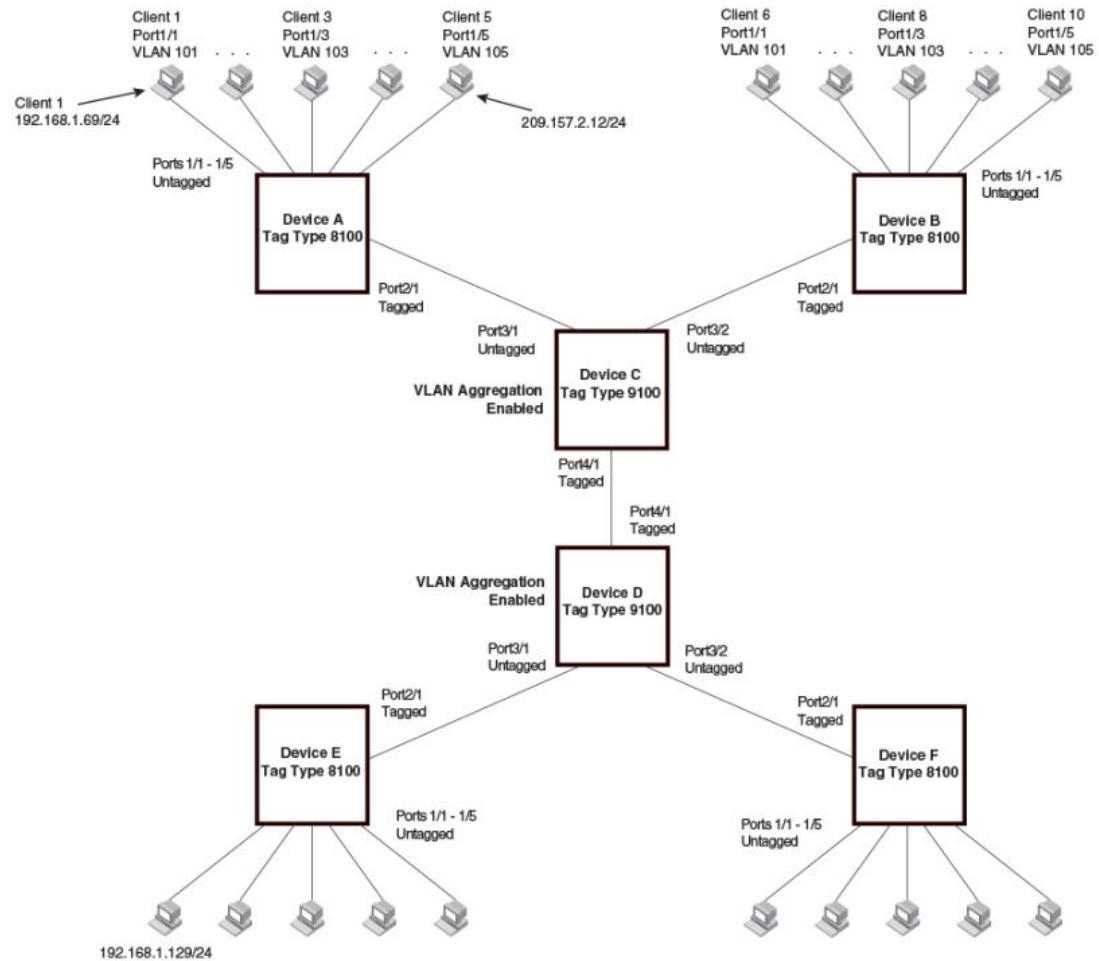
FIGURE 84 Conceptual model of the super aggregated VLAN application

Each client connected to the edge device is in its own port-based VLAN, which is like an ATM channel. All the clients' VLANs are aggregated by the edge device into a single VLAN for connection to the core. The single VLAN that aggregates the clients' VLANs is like an ATM path.

The device that aggregates the VLANs forwards the aggregated VLAN traffic through the core. The core can consist of multiple devices that forward the aggregated VLAN traffic. The edge device at the other end of the core separates the aggregated VLANs into the individual client VLANs before forwarding the traffic. The edge devices forward the individual client traffic to the clients. For the clients' perspective, the channel is a direct point-to-point link.

The following figure shows an example application that uses aggregated VLANs. This configuration includes the client connections shown in [Figure 84](#).

FIGURE 85 Example of a super aggregated VLAN application



In this example, a collocation service provides private channels for multiple clients. Although the same devices are used for all the clients, the VLANs ensure that each client receives its own Layer 2 broadcast domain, separate from the broadcast domains of other clients. For example, client 1 cannot ping client 5.

The clients at each end of a channel appear to each other to be directly connected and thus can be on the same subnet and use network services that require connection to the same subnet. In this example, client 1 is in subnet 192.168.1.0/24 and so is the device at the other end of client 1 channel.

Because each VLAN configured on the core devices is an aggregate of multiple client VLANs, the aggregated VLANs greatly increase the number of clients a core device can accommodate.

This example shows a single link between the core devices. However, you can use a trunk group to add link-level redundancy.

Configuration notes for aggregated VLANs

- This feature is not supported on the 48-port 10/100/1000 Mbps (RJ45) Ethernet POE interface module (SX-FI48GPP).
- Super Aggregated VLANs and VSRP are not supported together on the same device.
- Super Aggregated VLANs and Q-in-Q are supported using the tag-profile command.

Configuring aggregated VLANs

To configure aggregated VLANs, perform the following tasks:

- On each edge device, configure a separate port-based VLAN for each client connected to the edge device. In each client VLAN:
 - Add the port connected to the client as an untagged port.
 - Add the port connected to the core device (the device that will aggregate the VLANs) as a tagged port. This port must be tagged because all the client VLANs share the port as an uplink to the core device.
- On each core device:
 - Enable VLAN aggregation. This support allows the core device to add an additional tag to each Ethernet frame that contains a VLAN packet from the edge device. The additional tag identifies the aggregate VLAN (the path). However, the additional tag can cause the frame to be longer than the maximum supported frame size. The larger frame support allows Ethernet frames up to 1530 bytes long.
 - To allow frames larger than 1522, you must enable jumbo frames. To globally enable jumbo support , enter commands such as the following.

```
device(config)# jumbo
device(config)# write memory
device(config)# end
device# reload
```

NOTE

Enable the VLAN aggregation option only on the core devices.

- Configure a VLAN tag type (tag ID) that is different than the tag type used on the edge devices. If you use the default tag type (8100) on the edge devices, set the tag type on the core devices to another value, such as 9100. The tag type must be the same on all the core devices. The edge devices also must have the same tag type but the type must be different from the tag type on the core devices.

NOTE

You can enable the Spanning Tree Protocol (STP) on the edge devices or the core devices, but not both. If you enable STP on the edge devices and the core devices, STP will prevent client traffic from travelling through the core to the other side.

Configuring aggregated VLANs on an edge device

To configure the aggregated VLANs on device A in [Figure 85](#) on page 390, enter the following commands.

```
device(config)# vlan 101 by port
device(config-vlan-101)# tagged ethernet 2/1
device(config-vlan-101)# untagged ethernet 1/1
device(config-vlan-101)# exit
device(config)# vlan 102 by port
```

```
device(config-vlan-102)# tagged ethernet 2/1
device(config-vlan-102)# untagged ethernet 1/2
device(config-vlan-102)# exit
device(config)# vlan 103 by port
device(config-vlan-103)# tagged ethernet 2/1
device(config-vlan-103)# untagged ethernet 1/3
device(config-vlan-103)# exit
device(config)# vlan 104 by port
device(config-vlan-104)# tagged ethernet 2/1
device(config-vlan-104)# untagged ethernet 1/4
device(config-vlan-104)# exit
device(config)# vlan 105 by port
device(config-vlan-105)# tagged ethernet 2/1
device(config-vlan-105)# untagged ethernet 1/5
device(config-vlan-105)# exit
device(config)# write memory
```

Syntax: [no] **vlan** *vlan-id* [**by port**]

Syntax: [no] **tagged ethernet** [*slotnum/]portnum* [**to** [*slotnum/]portnum] | **ethernet** [*slotnum/]portnum]**

Syntax: [no] **untagged ethernet** [*slotnum/]portnum[**to** [*slotnum/]portnum] | **ethernet**[*slotnum/]portnum]***

Use the **tagged** command to add the port that the device uses for the uplink to the core device. Use the **untagged** command to add the ports connected to the individual clients.

Configuring aggregated VLANs on a core device

To configure the aggregated VLANs on device C in [Figure 85](#) on page 390, enter the following commands.

```
device(config)# tag-type 9100
device(config)# aggregated-vlan
device(config)# vlan 101 by port
device(config-vlan-101)# tagged ethernet 4/1
device(config-vlan-101)# untagged ethernet 3/1
device(config-vlan-101)# exit
device(config)# vlan 102 by port
device(config-vlan-102)# tagged ethernet 4/1
device(config-vlan-102)# untagged ethernet 3/2
device(config-vlan-102)# exit
device(config)# write memory
```

Syntax: [no] **tag-type** *num*

Syntax: [no] **aggregated-vlan**

The *num* parameter specifies the tag type can be a hexadecimal value from 0 - ffff. The default is 8100.

NOTE

tag-type is not applicable to ICX 7xxx devices. However, the same functionality can be achieved using the tag-profile CLI.

Verifying the aggregated VLAN configuration

You can verify the VLAN, VLAN aggregation option, and tag configuration by viewing the running-config. To display the running-config, enter the **show running-config** command from any CLI prompt. After you save the configuration changes to the startup-config, you also can display the settings in that file by entering the **show configuration** command from any CLI prompt.

Complete CLI examples for aggregated VLANs

The following sections show all the Aggregated VLAN configuration commands on the devices in [Figure 85](#) on page 390.

NOTE

In these examples, the configurations of the edge devices (A, B, E, and F) are identical. The configurations of the core devices (C and D) also are identical. The aggregated VLAN configurations of the edge and core devices on one side must be symmetrical (in fact, a mirror image) to the configurations of the devices on the other side. For simplicity, the example in [Figure 85](#) on page 390 is symmetrical in terms of the port numbers. This allows the configurations for both sides of the link to be the same. If your configuration does not use symmetrically arranged port numbers, the configurations should not be identical but must use the correct port numbers.

Commands for configuring aggregated VLANs on device A

```
deviceA(config)# vlan 101 by port
deviceA(config-vlan-101)# tagged ethernet 2/1
deviceA(config-vlan-101)# untagged ethernet 1/1
deviceA(config-vlan-101)# exit
deviceA(config)# vlan 102 by port
deviceA(config-vlan-102)# tagged ethernet 2/1
deviceA(config-vlan-102)# untagged ethernet 1/2
deviceA(config-vlan-102)# exit
deviceA(config)# vlan 103 by port
deviceA(config-vlan-103)# tagged ethernet 2/1
deviceA(config-vlan-103)# untagged ethernet 1/3
deviceA(config-vlan-103)# exit
deviceA(config)# vlan 104 by port
deviceA(config-vlan-104)# tagged ethernet 2/1
deviceA(config-vlan-104)# untagged ethernet 1/4
deviceA(config-vlan-104)# exit
deviceA(config)# vlan 105 by port
deviceA(config-vlan-105)# tagged ethernet 2/1
deviceA(config-vlan-105)# untagged ethernet 1/5
deviceA(config-vlan-105)# exit
vA(config)# write memory
```

Commands for configuring aggregated VLANs on device B

The commands for configuring device B are identical to the commands for configuring device A. Notice that you can use the same channel VLAN numbers on each device. The devices that aggregate the VLANs into a path can distinguish between the identically named channel VLANs based on the ID of the path VLAN.

```
deviceB(config)# vlan 101 by port
deviceB(config-vlan-101)# tagged ethernet 2/1
deviceB(config-vlan-101)# untagged ethernet 1/1
deviceB(config-vlan-101)# exit
deviceB(config)# vlan 102 by port
deviceB(config-vlan-102)# tagged ethernet 2/1
deviceB(config-vlan-102)# untagged ethernet 1/2
deviceB(config-vlan-102)# exit
deviceB(config)# vlan 103 by port
deviceB(config-vlan-103)# tagged ethernet 2/1
deviceB(config-vlan-103)# untagged ethernet 1/3
deviceB(config-vlan-103)# exit
deviceB(config)# vlan 104 by port
deviceB(config-vlan-104)# tagged ethernet 2/1
deviceB(config-vlan-104)# untagged ethernet 1/4
deviceB(config-vlan-104)# exit
deviceB(config)# vlan 105 by port
deviceB(config-vlan-105)# tagged ethernet 2/1
```

```
deviceB(config-vlan-105)# untagged ethernet 1/5
deviceB(config-vlan-105)# exit
deviceB(config)# write memory
```

Commands for configuring aggregated VLANs on device C

Because device C is aggregating channel VLANs from devices A and B into a single path, you need to change the tag type and enable VLAN aggregation.

```
deviceC(config)# tag-type 9100
deviceC(config)# aggregated-vlan
deviceC(config)# vlan 101 by port
deviceC(config-vlan-101)# tagged ethernet 4/1
deviceC(config-vlan-101)# untagged ethernet 3/1
deviceC(config-vlan-101)# exit
deviceC(config)# vlan 102 by port
deviceC(config-vlan-102)# tagged ethernet 4/1
deviceC(config-vlan-102)# untagged ethernet 3/2
deviceC(config-vlan-102)# exit
deviceC(config)# write memory
```

Commands for configuring aggregated VLANs on device D

Device D is at the other end of path and separates the channels back into individual VLANs. The tag type must be the same as tag type configured on the other core device (Device C). In addition, VLAN aggregation also must be enabled.

```
deviceD(config)# tag-type 9100
deviceD(config)# aggregated-vlan
deviceD(config)# vlan 101 by port
deviceD(config-vlan-101)# tagged ethernet 4/1
deviceD(config-vlan-101)# untagged ethernet 3/1
deviceD(config-vlan-101)# exit
deviceD(config)# vlan 102 by port
deviceD(config-vlan-102)# tagged ethernet 4/1
deviceD(config-vlan-102)# untagged ethernet 3/2
deviceD(config-vlan-102)# exit
deviceD(config)# write memory
```

Commands for configuring aggregated VLANs on device E

Because the configuration in [Figure 85](#) on page 390 is symmetrical, the commands for configuring device E are identical to the commands for configuring device A.

```
deviceE(config)# vlan 101 by port
deviceE(config-vlan-101)# tagged ethernet 2/1
deviceE(config-vlan-101)# untagged ethernet 1/1
deviceE(config-vlan-101)# exit
deviceE(config)# vlan 102 by port
deviceE(config-vlan-102)# tagged ethernet 2/1
deviceE(config-vlan-102)# untagged ethernet 1/2
deviceE(config-vlan-102)# exit
deviceE(config)# vlan 103 by port
deviceE(config-vlan-103)# tagged ethernet 2/1
deviceE(config-vlan-103)# untagged ethernet 1/3
deviceE(config-vlan-103)# exit
deviceE(config)# vlan 104 by port
deviceE(config-vlan-104)# tagged ethernet 2/1
deviceE(config-vlan-104)# untagged ethernet 1/4
deviceE(config-vlan-104)# exit
deviceE(config)# vlan 105 by port
deviceE(config-vlan-105)# tagged ethernet 2/1
deviceE(config-vlan-105)# untagged ethernet 1/5
deviceE(config-vlan-105)# exit
deviceE(config)# write memory
```

Commands for configuring aggregated VLANs on device F

The commands for configuring device F are identical to the commands for configuring device E. In this example, Because the port numbers on each side of the configuration in [Figure 85](#) on page 390 are symmetrical, the configuration of device F is also identical to the configuration of device A and device B.

```
deviceF(config)# vlan 101 by port
deviceF(config-vlan-101)# tagged ethernet 2/1
deviceF(config-vlan-101)# untagged ethernet 1/1
deviceF(config-vlan-101)# exit
deviceF(config)# vlan 102 by port
deviceF(config-vlan-102)# tagged ethernet 2/1
deviceF(config-vlan-102)# untagged ethernet 1/2
deviceF(config-vlan-102)# exit
deviceF(config)# vlan 103 by port
deviceF(config-vlan-103)# tagged ethernet 2/1
deviceF(config-vlan-103)# untagged ethernet 1/3
deviceF(config-vlan-103)# exit
deviceF(config)# vlan 104 by port
deviceF(config-vlan-104)# tagged ethernet 2/1
deviceF(config-vlan-104)# untagged ethernet 1/4
deviceF(config-vlan-104)# exit
deviceF(config)# vlan 105 by port
deviceF(config-vlan-105)# tagged ethernet 2/1
deviceF(config-vlan-105)# untagged ethernet 1/5
deviceF(config-vlan-105)# exit
deviceF(config)# write memory
```

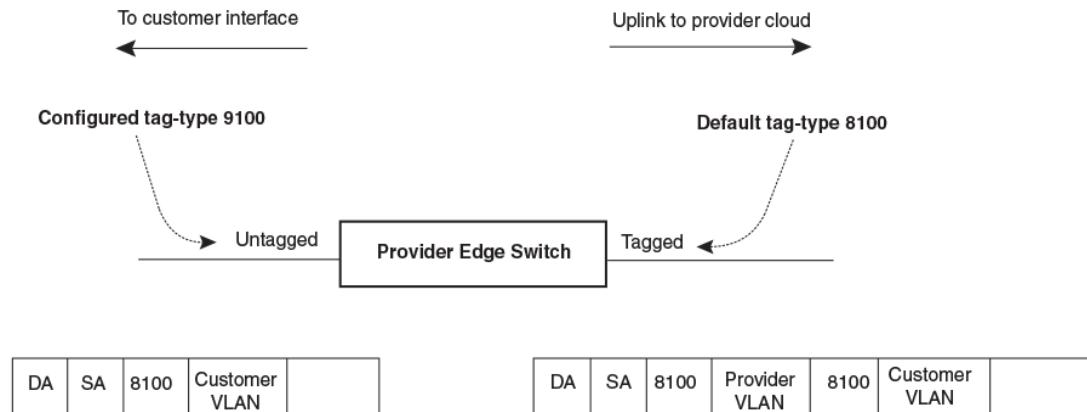
802.1ad tagging configuration

802.1ad tagging provides finer granularity for configuring 802.1Q tagging, enabling you to configure 802.1Q tag-types on a group of ports. This feature allows you to create two identical 802.1Q tags (802.1ad tagging) on a single device. This enhancement improves SAV interoperability between Brocade devices and other vendors' devices that support the 802.1Q tag-types, but are not very flexible with the tag-types they accept.

NOTE

Brocade devices treat a double-tagged Ethernet frame as a Layer 2 only frame. The packets are not inspected for Layer 3 and Layer 4 information, and operations are not performed on the packet utilizing Layer 3 or Layer 4 information.

The following figure shows an example application with 802.1ad tagging.

FIGURE 86 802.1ad configuration example

In the above figure, the untagged ports (to customer interfaces) accept frames that have any 802.1Q tag other than the configured tag-type 9100. These packets are considered untagged on this incoming port and are re-tagged when they are sent out of the uplink towards the provider. The 802.1Q tag-type on the uplink port is 8100, so the Brocade device will switch the frames to the uplink device with an additional 8100 tag, thereby supporting devices that only support this method of VLAN tagging.

Configuration rules for 802.1ad tagging

- On devices that support port regions, if you configure a port with an 802.1Q tag-type, the Brocade device automatically applies the 802.1Q tag-type to all ports within the same port region. Likewise, if you remove the 802.1Q tag-type from a port, the Brocade device automatically removes the 802.1Q tag-type from all ports within the same port region.
- Because the uplink (to the provider cloud) and the edge link (to the customer port) must have different 802.1Q tags, make sure the uplink and edge link are in different port regions. Refer to [About port regions](#) on page 17 for a list of valid port regions. Many FastIron devices have only a single port region. The above statement means that q-in-q is not supported on such devices, which is not true. There are two ways to achieve q-in-q in FastIron:
 - Using the tag-type option in the CLI: When you enable tag-type on one port, it is applied to complete port-region (in most cases port-region can be seen as a single device). To achieve q-in-q using tag-type, ingress port and egress port have to be in different port-regions. So FI units, which have only one port regions per device, cannot achieve q-in-q using tag-type CLI.

NOTE

This option is not applicable to the ICX 7xxx series devices. Only the tag-profile option is supported.

- Using the tag-profile option in the CLI: When you set tag-profile in global config, a second port region index value gets added. You need to enable tag-profile on per port basis using the "tag-profile enable" CLI. After enabling tag-profile on the port, the port's ingress and egress values point to different port region index values. Using this method, you can achieve q-in-q, even if ingress and egress ports are in a single device. The tag-profile provides more functionality compared to tag-type and should be preferred to enable q-in-q.
- The FCX, ICX 6430, ICX 6450, ICX 6610 and ICX 7xxx devices support tag-profile. For details on ICX 6610 port regions refer to Layer 2 Switching Configuration Guide. In ICX 6610, the port regions are exactly mapped to the devices.
- Q-in-Q and SAV are not supported on the following :

- 48-port 10/100/1000 Mbps (RJ45) Ethernet POE interface module (SX-FI48GPP)
- 2-port 10-Gigabit Ethernet interface module (SX-FI-2XG)
- 8-port 10-Gigabit Ethernet interface module (SX-FI-8XG)
- 24-port Gigabit Ethernet fiber interface module (SX-FI-24HF)
- 24-port Gigabit Ethernet copper interface module with PoE+ (SX-FI-24GPP)
- ICX6430C devices.
- FastIron X Series devices support one configured tag-type per device along with the default tag-type of 8100. For example, if you configure an 802.1Q tag of 9100 on ports 1 - 12, then later configure an 802.1Q tag of 5100 on port 15, the device automatically applies the 5100 tag to all ports in the same port region as port 15, and also changes the 802.1Q tag-type on ports 1 - 12 to 5100.
- 802.1ad tagging and VSRP are not supported together on the same device.
- In addition to **tag-type** Brocade FCX and ICX devices support **tag-profile**. For more information, refer to [Configuring 802.1ad tag profiles](#) on page 398 .

Enabling 802.1ad tagging

To enable 802.1ad tagging, configure an 802.1Q tag on the untagged edge links (the customer ports) to any value other than the 802.1Q tag for incoming traffic. For example, in [Figure 87](#) on page 398, the 802.1Q tag on the untagged edge links (ports 11 and 12) is 9100, whereas, the 802.1Q tag for incoming traffic is 8100.

To configure 802.1 ad tagging as shown in [Figure 87](#) on page 398, enter commands such as the following on the untagged edge links of devices C and D.

```
device(config)# tag-type 9100 ethernet 11 to 12
device(config)# aggregated-vlan
```

Note that because ports 11 and 12 belong to the port region 1 - 12, the 802.1Q tag actually applies to ports 1 - 12.

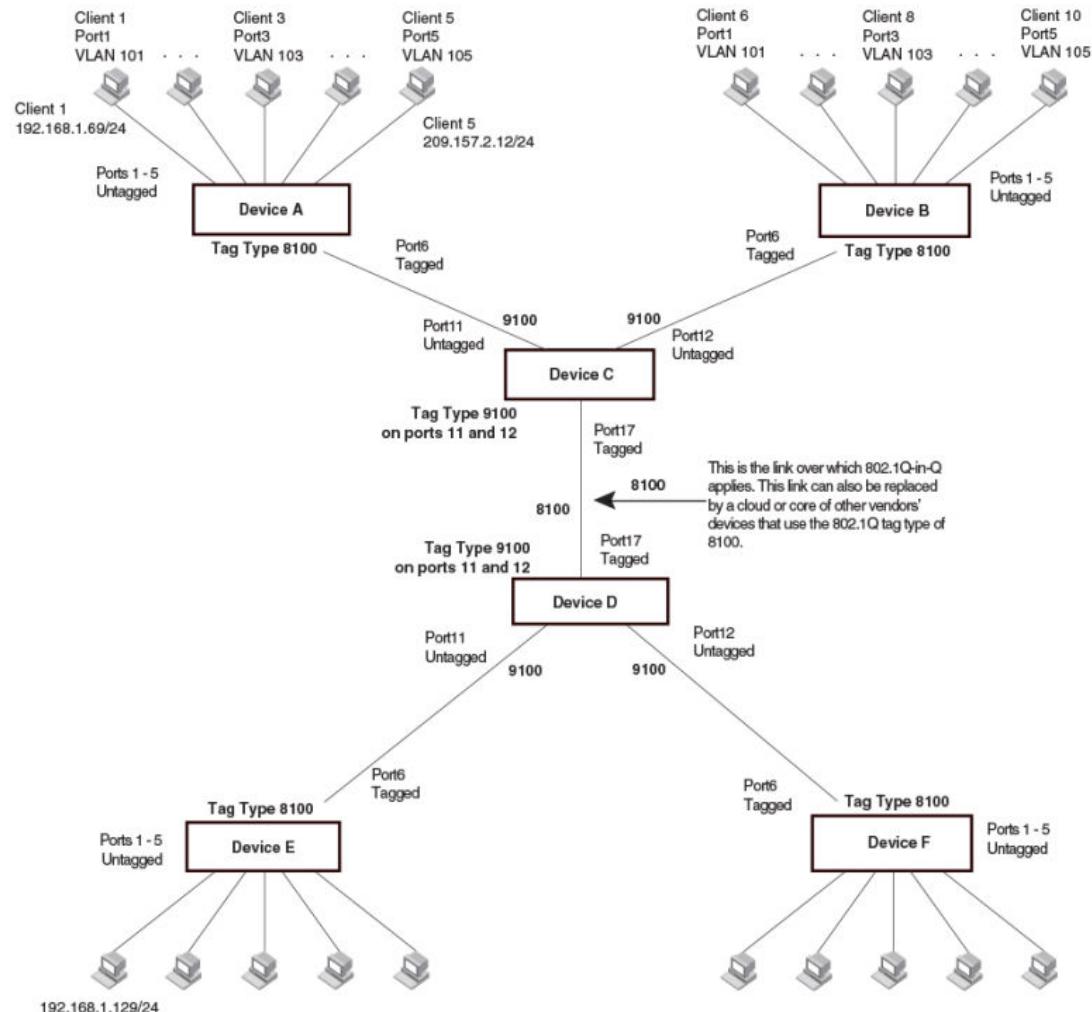
Syntax: [no] **tag-type** *num* [**ethernet** *port* [**to** *port*]]

The **ethernet** *port* to *port* parameter specifies the ports that will use the defined 802.1Q tag. This parameter operates with the following rules:

- If you specify a single port number, the 802.1Q tag applies to all ports within the port region. For example, if you enter the command **tag-type 9100 ethernet 1**, the Brocade device automatically applies the 802.1Q tag to ports 1 - 12 because all of these ports are in the same port region. You can use the **show running-config** command to view how the command has been applied.
- If you do not specify a port or range of ports, the 802.1Q tag applies to all Ethernet ports on the device.

Example 802.1ad configuration

The following figure shows an example 802.1ad configuration.

FIGURE 87 Example 802.1ad configuration

Configuring 802.1ad tag profiles

NOTE

802.1ad tag profiles are not supported on FastIron X Series devices.

The 802.1ad tagging feature supports a **tag-profile** command that allows you to add a tag profile with a value of 0 to 0xffff in addition to the default tag-type 0x8100. This enhancement also allows you to add a tag profile for a single port, or to direct a group of ports to a globally-configured tag profile.

Configuration notes for 802.1ad tagging

- One global tag profile with a number between 0 and 0xffff can be configured on stackable devices.
- On individual ports, if **tag-profile** is enabled, it points to the global tag profile.
- **Tag-profile** can also be enabled for provisional ports.

- **Tag-type** and **tag-profile** cannot be configured at the same time. You will see the message "un-configure the tag-type to set the tag-profile." If **tag-type** is already configured, you will need to unconfigure it and then add the **tag-profile**.
- Do not use the **tag-type** command in conjunction with the **tag-profile** command. If a **tag-type** has already been configured and you try to use the **tag-profile** command, you will see an error message telling you to remove the tag-type before you add the tag-profile.
- For devices operating in an IronStack topology, when a tag-type for a port is changed, the tag-type for all of the ports on a stack unit also changes. Because of this limitation, SAV and Q-in-Q cannot be used at the same time on stacking devices.

CLI Syntax for 802.1ad tagging

To add a global tag-profile enter the following command.

```
device(config) # tag-profile 9500
```

This command adds a profile in addition to the default profile of 0x8100.

Syntax: [no] tag-profile tag-no

where tag-no can be 0x8100 (the default) or 0xffff.

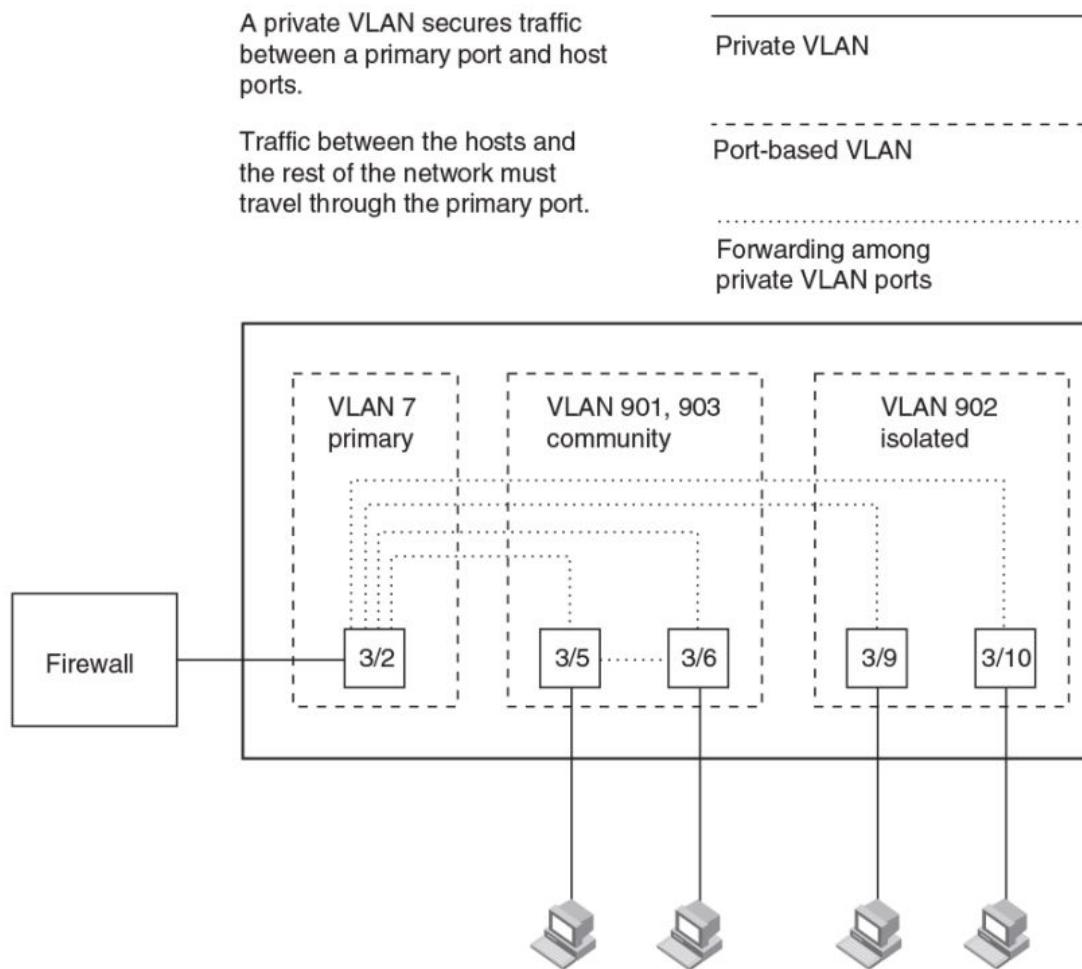
To enable the new profile on individual ports, enter commands similar to the following.

```
device(config)# interface ethernet 1/1/1
device(config-if-e1000-1/1/1)# tag-profile enable
device(config-mif-1/1/1,1/2/1)# tag-profile enable
```

Syntax: [no] tag-profile enable

Private VLAN configuration

A private VLAN (PVLAN) is a VLAN that has the properties of standard Layer 2 port-based VLANs but also provides additional control over flooding packets on a VLAN. The following table shows an example of an application using a PVLAN.

FIGURE 88 PVLAN used to secure communication between a workstation and servers

This example uses a PVLAN to secure traffic between hosts and the rest of the network through a firewall. Five ports in this example are members of a PVLAN. The first port (port 3/2) is attached to a firewall. The next four ports (ports 3/5, 3/6, 3/9, and 3/10) are attached to hosts that rely on the firewall to secure traffic between the hosts and the rest of the network. In this example, two of the hosts (on ports 3/5 and 3/6) are in a community PVLAN, and thus can communicate with one another as well as through the firewall. The other two hosts (on ports 3/9 and 3/10), are in an isolated VLAN and thus can communicate only through the firewall. The two hosts are secured from communicating with one another even though they are in the same VLAN.

By default, unknown-unicast, unregistered multicast, and broadcast are flooded in PVLAN.

By default, on all the FastIron platforms, the device will forward broadcast, unregistered multicast, and unknown unicast packets from outside sources into the PVLAN.

By default, in FastIron FSX platforms , the device will not forward broadcast, unregistered multicast, and unknown unicast packets from outside sources into the PVLAN. If needed, you can override this behavior for broadcast packets, unknown-unicast packets, or both.

You can configure a combination of the following types of PVLANS:

- Primary - The primary PVLAN ports are "promiscuous". They can communicate with all the isolated PVLAN ports and community PVLAN ports in the isolated and community VLANs that are mapped to the promiscuous port.
- Isolated - Broadcasts and unknown unicasts received on isolated ports are sent only to the promiscuous ports and switch - switch ports. They are not flooded to other ports in the isolated VLAN.

NOTE

On all devices, however, private VLANs will act as a normal VLAN and will flood unknown destinations, broadcast and multicast traffic to all ports in the VLAN if the primary VLAN does not have the PVLAN mapping that defines the uplink port for the isolated VLAN.

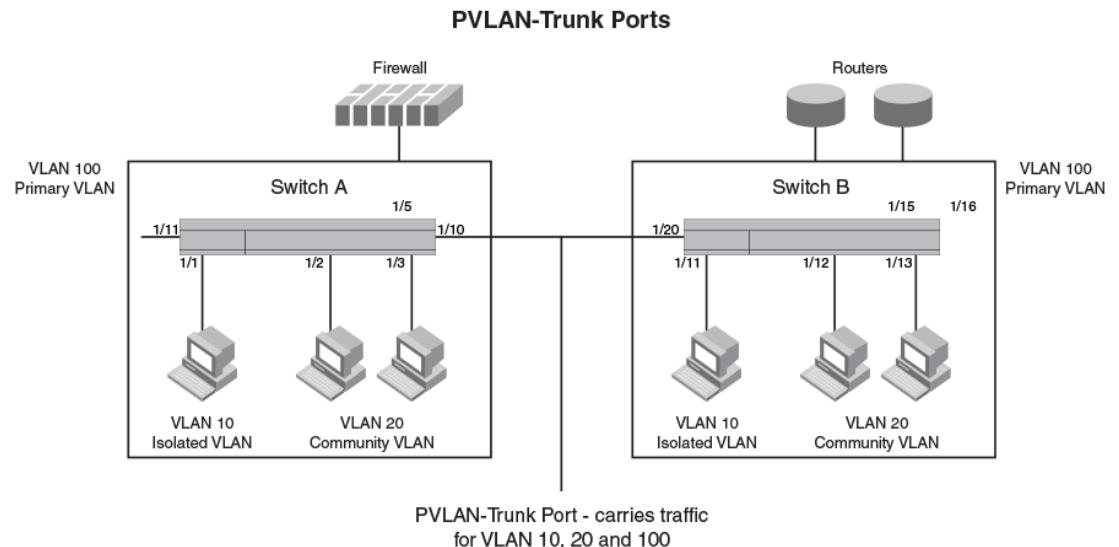
- Community - Broadcasts and unknown unicasts received on community ports are sent to the primary port and also are flooded to the other ports in the community VLAN.

Each PVLAN must have a primary VLAN. The primary VLAN is the interface between the secured ports and the rest of the network. The PVLAN can have any combination of community and isolated VLANs.

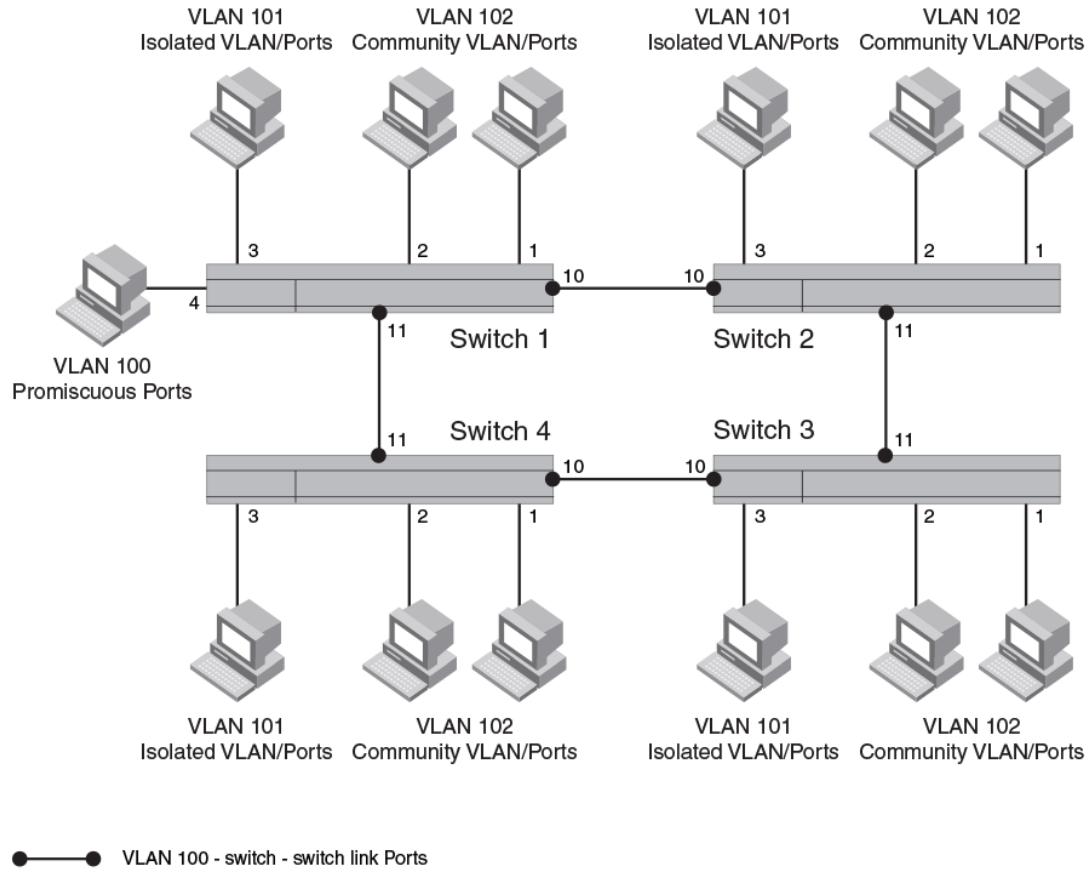
As with regular VLANs, PVLANs can span multiple switches. The PVLAN is treated like any other VLAN by the PVLAN-trunk ports. [Figure 89](#) shows an example of a PVLAN network across switches:

- Broadcast, unknown unicast or unregistered multicast traffic from the primary VLAN port is forwarded to all ports in isolated and community VLANs in both the switches.
- Broadcast, unknown unicast or unregistered multicast traffic from an isolated port in switch A is not forwarded to an isolated port in switch A. It will not be forwarded to an isolated port in switch B across the PVLAN-trunk port.
- Broadcast, unknown unicast or unregistered multicast traffic from a community port in switch A will be forwarded to a community port in switch B through the PVLAN-trunk port. It is forwarded to the promiscuous ports and PVLAN-trunk ports of the primary VLAN.

FIGURE 89 PVLAN across switches



The following figure shows an example PVLAN network with tagged switch-switch link ports.

FIGURE 90 Example PVLAN network with tagged ports

The following table lists the differences between PVLANS and standard VLANs.

TABLE 56 Comparison of PVLANS and standard port-based VLANs

Forwarding behavior	Private VLANs	Standard VLANs
All ports within a VLAN constitute a common layer broadcast domain	No	Yes
Broadcasts and unknown unicasts are forwarded to all the VLAN ports by default	No (isolated VLAN) Yes (community VLAN) Yes (Primary)	Yes
Known unicasts	No (isolated VLAN) Yes (community VLAN) Yes (Primary)	Yes

Configuration notes for PVLANS and standard VLANs

- PVLANS are supported on untagged ports on all FastIron platforms. PVLANS are also supported on tagged ports on devices other than FSX, ICX 6430, ICX 6430-C12 ICX 7450, ICX 7250 and ICX 7750.
- Normally, in any port-based VLAN, the Brocade device floods unknown unicast, unregistered multicast, and broadcast packets in hardware, although selective packets, such as IGMP, may be sent only to the CPU for analysis, based on the IGMP snooping configuration. When protocol or subnet VLANs are enabled, or if PVLAN mappings are enabled, the Brocade device will flood unknown unicast, unregistered multicast, and broadcast packets in software. The flooding of broadcast or unknown unicast from the community or isolated VLANs to other secondary VLANs will be governed by the PVLAN forwarding rules. The switching is done in hardware and thus the CPU does not enforce packet restrictions. The hardware forwarding behavior is supported on the FCX, ICX 6650, ICX 6610, ICX 6450, ICX 6430, ICX 6650, ICX 6430-C12, ICX 6450-C12, ICX 7250, ICX 7450 and ICX 7750.
- There is currently no support for IGMP snooping within PVLANS. In order for clients in PVLANS to receive multicast traffic, IGMP snooping must be disabled so that all multicast packets are treated as unregistered packets and are flooded in software to all the ports.
- The FastIron forwards all known unicast traffic in hardware. This differs from the way the BigIron implements PVLANS, in that the BigIron uses the CPU to forward packets on the primary VLAN "promiscuous" port. In addition, on the BigIron, support for the hardware forwarding sometimes results in multiple MAC address entries for the same MAC address in the device MAC address table. On the FastIron, multiple MAC entries do not appear in the MAC address table because the FastIron transparently manages multiple MAC entries in hardware.
- To configure a PVLAN, configure each of the component VLANs (isolated, community, and public) as a separate port-based VLAN:
 - Use standard VLAN configuration commands to create the VLAN and add ports.
 - Identify the PVLAN type (isolated, community, or public)
 - For the primary VLAN, map the other PVLANs to the ports in the primary VLAN
- A primary VLAN can have multiple ports. All these ports are active, but the ports that will be used depends on the PVLAN mappings. Also, secondary VLANs (isolated and community VLANs) can be mapped to more than one primary VLAN port.
- You can configure PVLANs and dual-mode VLAN ports on the same device. However, the dual-mode VLAN ports cannot be members of PVLANs.
- VLAN identifiers configured as part of a PVLAN (primary, isolated, or community) should be consistent across the switched network. The same VLAN identifiers cannot be configured as a normal VLAN or a part of any other PVLAN.
- Promiscuous and switch-switch link ports are member ports of the primary VLAN only. All switch-switch link ports are tagged ports.
- Member ports of isolated and community VLANs cannot be member ports of any other VLAN.
- VLAN classification is performed for all ports in primary and community VLANs based on the PVLAN ID (PVID) only (no VLAN classification by port, protocol, ACL and so on, if any). For isolated VLAN ports there is not classification at all (not even on PVLAN ID).
- PVST, when needed in PVLANs, should be enabled on all (primary and secondary) private VLANs.
- On ICX 7750, 7450 and 7250 devices, VLAN classification will be performed for isolated VLAN ports.
- LAG ports are not allowed in any PVLAN.

TABLE 57 PVLAN support matrix

Platform	Forwarding Type	Tagged Port	Untagged Port	ISL Port	Multiple Promiscuous Port
ICX-6650	Hardware	Yes	Yes	Yes	Yes
ICX-6610	Hardware	Yes	Yes	Yes	Yes

TABLE 57 PVLAN support matrix (Continued)

Platform	Forwarding Type	Tagged Port	Untagged Port	ISL Port	Multiple Promiscuous Port
ICX-6450	Hardware	Yes	Yes	Yes	Yes
ICX-6650 C12	Hardware	Yes	Yes	Yes	Yes
ICX-6430 C12	Hardware	No	Yes	No	Yes
FCX	Hardware	Yes	Yes	Yes	Yes
SXL	Software	No	Yes	No	Yes

Configuring the primary VLAN

To configure a primary VLAN, enter commands such as the following.

```
device(config)# vlan 7
device(config-vlan-7)# untagged ethernet 3/2
device(config-vlan-7)# pvlan type primary
device(config-vlan-7)# pvlan mapping 901 ethernet 3/2
```

These commands create port-based VLAN 7, add port 3/2 as an untagged port, identify the VLAN as the primary VLAN in a PVLAN, and map the other secondary VLANs to the ports in this VLAN.

To map the secondary VLANs to the primary VLAN and to configure the tagged switch link port, enter commands such as the following.

```
device(config)# vlan 100
device(config-vlan-100)# tagged ethernet 1/1/10 to 1/1/11
device(config-vlan-100)# untagged ethernet 1/1/4
device(config-vlan-100)# pvlan type primary
device(config-vlan-100)# pvlan mapping 101 ethernet 1/1/4
device(config-vlan-100)# pvlan mapping 102 ethernet 1/1/4
device(config-vlan-100)# pvlan pvlan-trunk 101 ethernet 1/1/10 to 1/1/11
```

These commands create port-based VLAN 100, add port 1/1/10 to 1/1/11 as a tagged port, identify the VLAN as the primary VLAN in a PVLAN, map the other secondary VLANs to the ports in this VLAN, and configure the tagged switch link port.

Syntax: **untagged ethernet [stack-unit[slotnum/]portnum [to [stack-unit[slotnum/]portnum | ethernet [stack-unit[slotnum/]portnum]**

or

Syntax: **tagged ethernet [stack-unit[slotnum/]portnum [to [stack-unit[slotnum/]portnum | ethernet[stack-unit[slotnum/]portnum]**

Syntax: **[no] pvlan type community | isolated | primary**

Syntax: **[no] pvlan mapping vlan-id ethernet [stack-unit[slotnum/]portnum**

Syntax: **[no] pvlan pvlan-trunk vlan-id ethernet [stack-unit[slotnum/]portnum [to [stack-unit[slotnum/]portnum]**

The **untagged** or **tagged** command adds the ports to the VLAN.

The **pvlan type** command specifies that this port-based VLAN is a PVLAN. Specify **primary** as the type.

The **pvlan mapping** command identifies the other PVLANS for which this VLAN is the primary. The command also specifies the primary VLAN ports to which you are mapping the other secondary VLANs. The mapping command is not allowed on the secondary VLANs. The parameters of the **pvlan mapping** command are as follows:

- The *vlan-id* parameter specifies another PVLAN. The other PVLAN you want to specify must already be configured.
- The *etherent portnum* parameter specifies the primary VLAN port to which you are mapping all the ports in the other PVLAN (the one specified by *vlan-id*).

The **pvlan pvlan-trunk** command identifies the switch-switch link for the PVLAN. There can be more than one switch-switch link for a single community VLAN.

NOTE

The **pvlan pvlan-trunk** command is not allowed on the secondary VLANs.

Configuring an isolated or community PVLAN

You can use the **pvlan type** command to configure the PVLAN as an isolated or community PVLAN. The following are some configuration considerations to be noted for configuring isolated and community PVLANS.

Isolated VLANs

- A port being added to the isolated VLAN can be either a tagged port or an untagged port.
- An isolated port (member of an isolated VLAN) communicates with the promiscuous port, if a promiscuous port is configured. If a switch-switch port is configured, the isolated port communicates with the switch-switch port also.
- An isolated VLAN must be associated with the primary VLAN for traffic from the isolated port to be switched. An isolated VLAN is associated with only one primary VLAN and to the same primary VLAN in the entire switched network.
- A primary VLAN is associated with only one isolated VLAN. An isolated VLAN can only be mapped to a promiscuous port and a switch-switch link port that belong to the same primary VLAN.

Community VLANs

- A port being added to the community VLAN can be either a tagged port or an untagged port.
- A member port of a community VLAN classifies a frame based on VLAN ID only.
- A community VLAN is associated with only one primary VLAN and to the same primary VLAN in the entire switched network. A primary VLAN is associated with multiple community VLANs.
- A community VLAN must be associated with the primary VLAN for traffic from the community port to be switched.

To configure a community PVLAN, enter commands such as the following.

```
device(config)# vlan 901
device(config-vlan-901)# untagged ethernet 3/5 to 3/6
device(config-vlan-901)# pvlan type community
```

These commands create port-based VLAN 901, add ports 3/5 and 3/6 to the VLAN as untagged ports, then specify that the VLAN is a community PVLAN.

Syntax: untagged ethernet [slotnum/]portnum [to [slotnum/]portnum | ethernet [slotnum/]portnum]

or

Syntax: tagged ethernet [slotnum/]portnum [to [slotnum/]portnum | ethernet [slotnum/]portnum]

Syntax: [no] pvlan type community | isolated | primary

The **untagged ethernet** or **tagged ethernet** command adds the ports to the VLAN.

The **pvlan type** command specifies that this port-based VLAN is a PVLAN and can be of the following types:

- **community** - Broadcasts and unknown unicasts received on community ports are sent to the primary port and also are flooded to the other ports in the community VLAN.
- **isolated** - Broadcasts and unknown unicasts received on isolated ports are sent only to the primary port. They are not flooded to other ports in the isolated VLAN.
- **primary** - The primary PVLAN ports are "promiscuous". They can communicate with all the isolated PVLAN ports and community PVLAN ports in the isolated and community VLANs that are mapped to the promiscuous port.

Changing from one PVLAN type to another (for example, from primary to community or vice versa) is allowed but the mapping will be removed.

Enabling broadcast or unknown unicast traffic to the PVLAN on FSX devices

To enhance PVLAN security, the primary PVLAN does not forward broadcast or unknown unicast packets to its community and isolated VLANs. For example, if port 3/2 in [Figure 88](#) on page 400 receives a broadcast packet from the firewall, the port does not forward the packet to the other PVLAN ports (3/5, 3/6, 3/9, and 3/10).

This forwarding restriction does not apply to traffic from the secondary PVLAN. The primary port does forward broadcast and unknown unicast packets that are received from the isolated and community VLANs. For example, if the host on port 3/9 sends an unknown unicast packet, port 3/2 forwards the packet to the firewall.

If you want to remove the forwarding restriction of the primary VLAN, you can enable the primary port to forward broadcast or unknown unicast traffic, if desired, using the following CLI method. You can enable or disable forwarding of broadcast or unknown unicast packets separately.

NOTE

On Layer 2 switches and Layer 3 switches, you also can use MAC address filters to control the traffic forwarded into and out of the PVLAN. In addition, if you are using a Layer 2 switch, you also can use ACLs.

Command syntax

To configure the ports in the primary VLAN to forward broadcast or unknown unicast and multicast traffic received from sources outside the PVLAN, enter the following commands at the global CONFIG level of the CLI.

```
Brocade(config)# pvlan-preference broadcast flood  
Brocade(config)# pvlan-preference unknown-unicast flood
```

These commands enable forwarding of broadcast and unknown-unicast packets to ports within the PVLAN. To again disable forwarding, enter a command such as the following.

```
Brocade(config)# no pvlan-preference broadcast flood
```

This command disables forwarding of broadcast packets within the PVLAN.

Syntax: [no] pvlan-preference broadcast | unknown-unicast flood

NOTE

The pvlan-preference broadcast and pvlan-preference unknown-unicast commands are not supported on the FCX and ICX platforms. These are supported on all the other FastIron platforms.

CLI example for a general PVLAN network

To configure the PVLANS shown in [Figure 88](#) on page 400, enter the following commands.

```
device(config)# vlan 901
device(config-vlan-901)# untagged ethernet 3/5 to 3/6
device(config-vlan-901)# pvlan type community
device(config-vlan-901)# exit
device(config)# vlan 902
device(config-vlan-902)# untagged ethernet 3/9 to 3/10
device(config-vlan-902)# pvlan type isolated
device(config-vlan-902)# exit
device(config)# vlan 903
device(config-vlan-903)# untagged ethernet 3/7 to 3/8
device(config-vlan-903)# pvlan type community
device(config-vlan-903)# exit
device(config)# vlan 7
device(config-vlan-7)# untagged ethernet 3/2
device(config-vlan-7)# pvlan type primary
device(config-vlan-7)# pvlan mapping 901 ethernet 3/2
device(config-vlan-7)# pvlan mapping 902 ethernet 3/2
device(config-vlan-7)# pvlan mapping 903 ethernet 3/2
```

CLI example for a PVLAN network with switch-switch link ports

To configure the PVLANS with tagged switch-switch link ports as shown in [Figure 90](#) on page 402, enter the following commands.

FCX Switch 1

```
device(config)# vlan 101 by port
device(config-vlan-101)# untagged ethernet 1/1/3
device(config-vlan-101)# pvlan type isolated
device(config)# vlan 102 by port
device(config-vlan-102)# untagged ethernet 1/1/1 to 1/1/2
device(config-vlan-102)# pvlan type community
device(config)# vlan 100 by port
device(config-vlan-100)# tagged ethernet 1/1/10 to 1/1/11
device(config-vlan-100)# untagged ethernet 1/1/4
device(config-vlan-100)# pvlan type primary
device(config-vlan-100)# pvlan mapping 101 ethernet 1/1/4
device(config-vlan-100)# pvlan mapping 102 ethernet 1/1/4
device(config-vlan-100)# pvlan pvlan-trunk 101 ethernet 1/1/10 to 1/1/11
device(config-vlan-100)# pvlan pvlan-trunk 102 ethernet 1/1/10 to 1/1/11
```

FCX Switch 2

```
device(config)# vlan 101 by port
device(config-vlan-101)# untagged ethernet 1/1/3
device(config-vlan-101)# pvlan type isolated
device(config)# vlan 102 by port
device(config-vlan-102)# untagged ethernet 1/1/1 to 1/1/2
device(config-vlan-102)# pvlan type community
device(config)# vlan 100 by port
device(config-vlan-100)# tagged ethernet 1/1/10 to 1/1/11
device(config-vlan-100)# pvlan type primary
device(config-vlan-100)# pvlan pvlan-trunk 102 ethernet 1/1/10 to 1/1/11
device(config-vlan-100)# pvlan pvlan-trunk 101 ethernet 1/1/10 to 1/1/11
```

FCX Switch 3

```
device(config)# vlan 101 by port
device(config-vlan-101)# untagged ethernet 1/1/3
device(config-vlan-101)# pvlan type isolated
device(config)# vlan 102 by port
device(config-vlan-102)# untagged ethernet 1/1/1 to 1/1/2
device(config-vlan-102)# pvlan type community
device(config)# vlan 100 by port
device(config-vlan-100)# tagged ethernet 1/1/10 to 1/1/11
device(config-vlan-100)# pvlan type primary
device(config-vlan-100)# pvlan pvlan-trunk 102 ethernet 1/1/10 to 1/1/11
device(config-vlan-100)# pvlan pvlan-trunk 101 ethernet 1/1/10 to 1/1/11
```

FCX Switch 4

```
device(config)# vlan 101 by port
device(config-vlan-101)# untagged ethernet 1/1/3
device(config-vlan-101)# pvlan type isolated
device(config)# vlan 102 by port
device(config-vlan-102)# untagged ethernet 1/1/1 to 1/1/2
device(config-vlan-102)# pvlan type community
device(config)# vlan 100 by port
device(config-vlan-100)# tagged ethernet 1/1/10 to 1/1/11
device(config-vlan-100)# pvlan type primary
device(config-vlan-100)# pvlan pvlan-trunk 102 ethernet 1/1/10 to 1/1/11
device(config-vlan-100)# pvlan pvlan-trunk 101 ethernet 1/1/10 to 1/1/11
```

Multiple promiscuous ports support in private VLANs

Promiscuous ports are member ports of a primary VLAN. Prior versions of the FastIron software supported only a single promiscuous port to be mapped to a secondary VLAN. However, now each secondary VLAN can associate with one or more promiscuous ports.

An isolated VLAN with member ports, when mapped to one or more promiscuous ports of the primary VLAN achieves the same forwarding behavior as that of the uplink ports in a port based VLAN (uplink-switch). The broadcast and unknown unicast traffic from a host (isolated) port is flooded to the uplink (promiscuous) ports only. Due to the hardware forwarding functionality of the private VLAN (in the supported stackable SKUs) this method of achieving the uplink port behavior achieves a better throughput than the conventional method of achieving uplink port.

Mapping secondary VLAN to primary VLAN by multiple promiscuous ports

1. Add the tagged ethernet ports to a VLAN.
2. Configure the VLAN as isolated.
3. Create VLAN 101.
4. Configure the VLAN 101 as primary VLAN in a PVLAN.
5. Map the isolated VLAN to the primary VLAN with uplink-ports as promiscuous ports. All broadcast and unknown-unicast traffic from isolated VLAN ports will be sent to only promiscuous ports.

In the following example configuration, the isolated VLAN 100 has multiple promiscuous ports 1/1/3 and 1/1/4.

```
device(config-vlan-100)# tagged ethernet 1/1/1 to 1/1/2
device(config-vlan-100)# pvlan type isolated
device(config-vlan-101)# tagged ethernet 1/1/3 to 1/1/5
device(config-vlan-101)# pvlan type primary
device(config-vlan-101)# pvlan map 100 ethernet 1/1/3
device(config-vlan-101)# pvlan map 100 ethernet 1/1/4
```

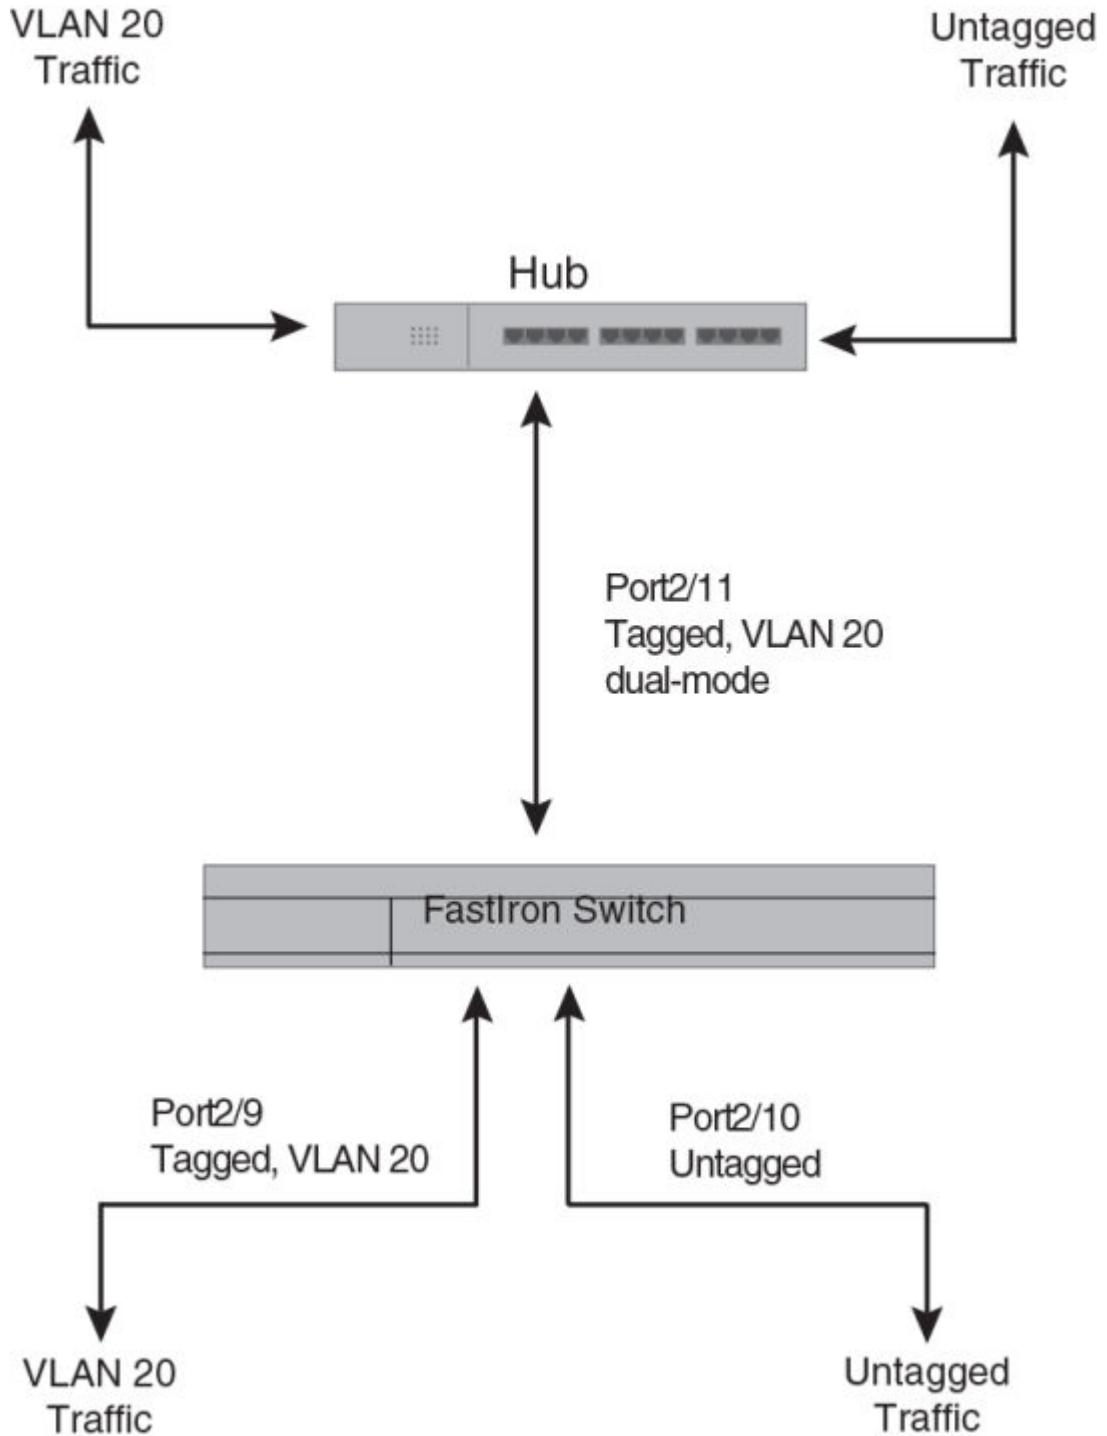
NOTE

Multiple promiscuous ports support for PVLAN is available on FCX, ICX 6610, ICX 6430, ICX 6450, ICX 6450-C12, ICX 6430-C12, ICX 6650, ICX 7250, ICX 7450 and ICX 7750 devices. The Brocade ICX 6430, ICX6430-C12 and SXL do not support tagged ports as part of PVLANS.

Dual-mode VLAN ports

Configuring a tagged port as a dual-mode port allows it to accept and transmit both tagged traffic and untagged traffic at the same time. A dual-mode port accepts and transmits frames belonging to VLANs configured for the port, as well as frames belonging to the default VLAN (that is, untagged traffic).

For example, in the following figure, port 2/11 is a dual-mode port belonging to VLAN 20. Traffic for VLAN 20, as well as traffic for the default VLAN, flows from a hub to this port. The dual-mode feature allows traffic for VLAN 20 and untagged traffic to go through the port at the same time.

FIGURE 91 Dual-mode VLAN port example

To enable the dual-mode feature on port 2/11 in the above figure, enter the following commands.

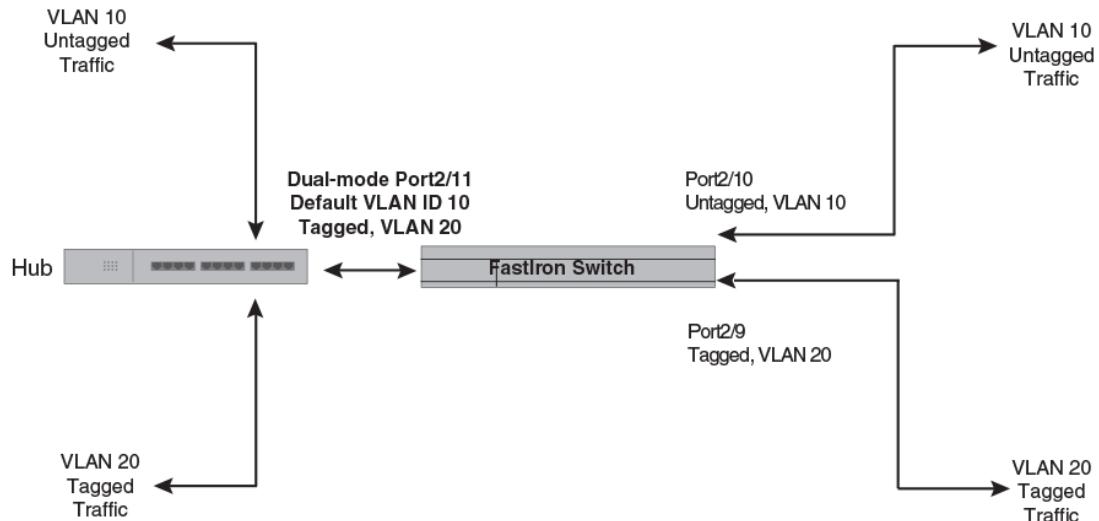
```
device(config)# vlan 20
device(config-vlan-20)# tagged ethernet 2/11
device(config-vlan-20)# tagged ethernet 2/9
device(config-vlan-20)# interface ethernet 2/11
```

```
device(config-if-e1000-2/11) # dual-mode
device(config-if-e1000-2/11) # exit
```

Syntax: [no] dual-mode

You can configure a dual-mode port to transmit traffic for a specified VLAN (other than the DEFAULT-VLAN) as untagged, while transmitting traffic for other VLANs as tagged. The following figure illustrates this enhancement.

FIGURE 92 Specifying a default VLAN ID for a dual-mode port



In the above figure, tagged port 2/11 is a dual-mode port belonging to VLANs 10 and 20. The default VLAN assigned to this dual-mode port is 10. This means that the port transmits tagged traffic on VLAN 20 (and all other VLANs to which the port belongs) and transmits untagged traffic on VLAN 10.

The dual-mode feature allows tagged traffic for VLAN 20 and untagged traffic for VLAN 10 to go through port 2/11 at the same time. A dual-mode port transmits only untagged traffic on its default VLAN (that is, either VLAN 1, or a user-specified VLAN ID), and only tagged traffic on all other VLANs.

The following commands configure VLANs 10 and 20 in [Figure 92](#). Tagged port 2/11 is added to VLANs 10 and 20, then designated a dual-mode port whose specified default VLAN is 10. In this configuration, port 2/11 transmits only untagged traffic on VLAN 10 and only tagged traffic on VLAN 20.

```
device(config)# vlan 10 by port
device(config-vlan-10)# untagged ethernet 2/10
device(config-vlan-10)# tagged ethernet 2/11
device(config-vlan-10)# exit
device(config)# vlan 20 by port
device(config-vlan-20)# tagged ethernet 2/9
device(config-vlan-20)# tagged ethernet 2/11
device(config-vlan-20)# exit
device(config)# interface ethernet 2/11
device(config-if-e1000-2/11) # dual-mode 10
device(config-if-e1000-2/11) # exit
```

Syntax: [no] dual-mode [*vlan-id*]

You can configure multiple ports using commands such as the following.

```
Brocade# interface ethernet 1/1/6 to 1/1/9
Brocade (config-mif-1/1/6-1/1/9) # dual-mode
```

NOTE

An error message is displayed while attempting to configure an existing dual-mode on a port range.

Example:

```
Port 1/1/6 has already been configured as dual mode on VLAN 20
Port 1/1/7 has already been configured as dual mode on VLAN 20
Port 1/1/8 has already been configured as dual mode on VLAN 20
Port 1/1/9 has already been configured as dual mode on VLAN 1
```

Notes:

- If you do not specify a *vlan-id* in the **dual mode** command, the port default VLAN is set to 1. The port transmits untagged traffic on the DEFAULT-VLAN.
- The dual-mode feature is disabled by default. Only tagged ports can be configured as dual-mode ports.
- In trunk group, either all of the ports must be dual-mode, or none of them can be.

The **show vlan** command displays a separate row for dual-mode ports on each VLAN.

```
device# show vlan
Total PORT-VLAN entries: 3
Maximum PORT-VLAN entries: 16
legend: [S=Slot]
PORT-VLAN 1, Name DEFAULT-VLAN, Priority level0, Spanning tree Off
    Un tagged Ports: (S1) 1 2 3 4 5 6 7 8
    Un tagged Ports: (S2) 1 2 3 4 5 6 7 8 12 13 14 15 16 17 18 19
    Un tagged Ports: (S2) 20 21 22 23 24
        Tagged Ports: None
        Uplink Ports: None
        DualMode Ports: None
PORT-VLAN 10, Name [None], Priority level0, Spanning tree Off
    Un tagged Ports: (S2) 10
        Tagged Ports: None
        Uplink Ports: None
        DualMode Ports: (S2) 11
PORT-VLAN 20, Name [None], Priority level0, Spanning tree Off
    Un tagged Ports: None
        Tagged Ports: (S2) 9
        Uplink Ports: None
        DualMode Ports: (S2) 11
```

Displaying VLAN information

After you configure the VLANs, you can verify the configuration using the **show** commands described in this section.

NOTE

If a VLAN name begins with "GVRP_VLAN_", the VLAN was created by the GARP VLAN Registration Protocol (GVRP). If a VLAN name begins with "STATIC_VLAN_", the VLAN was created by GVRP and then was converted into a statically configured VLAN.

Displaying VLANs in alphanumeric order

By default, VLANs are displayed in alphanumeric order, as shown in the following example.

```
device# show run
...
vlan 2 by port
...
vlan 10 by port
...
vlan 100 by port
...
```

Displaying system-wide VLAN information

Use the **show vlans** command to display VLAN information for all the VLANs configured on the device.

The following example shows the display for the IP subnet and IPX network VLANs configured in the examples in [Configuring an IP subnet VLAN with dynamic ports](#) on page 377 and [Configuring an IPX network VLAN with dynamic ports](#) on page 378.

```
device# show vlans
Total PORT-VLAN entries: 2
Maximum PORT-VLAN entries: 8
legend: [S=Slot]
PORT-VLAN 1, Name DEFAULT-VLAN, Priority level0, Spanning tree Off
  Untagged Ports: (S2) 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16
  Untagged Ports: (S2) 17 18 19 20 21 22 23 24
  Untagged Ports: (S4) 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16
  Untagged Ports: (S4) 17 18 19 20 21 22 23 24
    Tagged Ports: None
PORT-VLAN 10, Name IP VLAN, Priority level0, Spanning tree Off
  Untagged Ports: (S1) 1 2 3 4 5 6
    Tagged Ports: None
IP-subnet VLAN 10.1.1.0 255.255.255.0, Dynamic port enabled
  Name: Mktg-LAN
  Static ports: None
  Exclude ports: None
  Dynamic ports: (S1) 1 2 3 4 5 6
PORT-VLAN 20, Name IPX VLAN, Priority level0, Spanning tree Off
  Untagged Ports: (S2) 1 2 3 4 5 6
    Tagged Ports: None
IPX-network VLAN 0000ABCD, frame type ethernet_ii, Dynamic port enabled
  Name: Eng-LAN
  Static ports: None
  Exclude ports: None
  Dynamic ports: (S2) 1 2 3 4 5 6
```

In the **show vlans** output, ports that are tagged but are not dual-mode ports are listed as tagged ports. In the following example display output, ports 7 and 8 are dual-mode ports in port-based VLAN 4. Ports 7 and 8 also belong to port-based VLAN 3, but they are tagged ports only in VLAN 3 and are not configured as dual-mode ports.

```
device# show vlan 4
Total PORT-VLAN entries: 5
Maximum PORT-VLAN entries: 3210
PORT-VLAN 4, Name [None], Priority level0, Spanning tree Off
  Untagged Ports: None
    Tagged Ports: 6 9 10 11
    Uplink Ports: None
    DualMode Ports: 7 8
ESX624FE+2XG Router# show vlan 3
Total PORT-VLAN entries: 5
Maximum PORT-VLAN entries: 3210
PORT-VLAN 3, Name [None], Priority level0, Spanning tree Off
  Untagged Ports: None
    Tagged Ports: 6 7 8 9 10
    Uplink Ports: None
    DualMode Ports: None
```

Syntax: show vlans [vlan-id | ethernet [slotnum/]portnum]

The *vlan-id* parameter specifies a VLAN for which you want to display the configuration information.

The *slotnum* parameter is required on chassis devices.

The *portnum* parameter specifies a port. If you use this parameter, the command lists all the VLAN memberships for the port.

Displaying global VLAN information

The **show vlan brief** command displays the following information:

- The system-max VLAN values (maximum, default, and current)
- The default VLAN ID number
- The total number of VLANs configured on the device
- The VLAN ID numbers of the VLANs configured on the device

The following shows example output.

```
device# show vlan brief
System-max vlan Params: Max(4095) Default(64) Current(3210)
Default vlan Id :1
Total Number of Vlan Configured :5
VLANs Configured :1 to 4 10
```

Syntax: show vlan brief

Displaying VLAN information for specific ports

Use one of the following methods to display VLAN information for specific ports.

To display VLAN information for all the VLANs of which port 7/1 is a member, enter the following command.

```
device# show vlans ethernet 7/1
Total PORT-VLAN entries: 3
Maximum PORT-VLAN entries: 8
legend: [S=Slot]
PORT-VLAN 100, Name [None], Priority level0, Spanning tree Off
  Untagged Ports: (S7) 1 2 3 4
  Tagged Ports: None
```

Syntax: show vlans [vlan-id | ethernet [slotnum/]portnum]

The *vlan-id* parameter specifies a VLAN for which you want to display the configuration information.

The *slotnum* parameter is required on chassis devices.

The *portnum* parameter specifies a port. If you use this parameter, the command lists all the VLAN memberships for the port.

Displaying a port VLAN membership

To display VLAN membership for a specific port on the device, enter a command such as the following.

```
device# show vlan brief ethernet 7
Port 7 is a member of 3 VLANs
VLANs 3 to 4 10
```

Syntax: show vlan brief ethernet [slotnum/]portnum

The *slotnum* parameter is required on chassis devices.

Displaying a port dual-mode VLAN membership

The output of the **show interfaces** command lists dual-mode configuration and corresponding VLAN numbers. The following shows an example output.

```
device# show interfaces ethernet 7
GigabitEthernet7 is down, line protocol is down
Port down for 2 days 1 hour 40 minutes 5 seconds
Hardware is GigabitEthernet, address is 0000.00a8.4706 (bia 0000.00a8.4706)
Configured speed auto, actual unknown, configured duplex fdx, actual unknown
Configured mdi mode AUTO, actual unknown
Member of 3 L2 VLANs, port is dual mode in Vlan 4
, port state is BLOCKING
```

Syntax: show interfaces ethernet [slotnum/]portnum [to [slotnum/]portnum [ethernet [slotnum/]portnum...]]

The *slotnum* parameter is required on chassis devices.

NOTE

The port up/down time is required only for physical ports and not for loopback/ ve/ tunnel ports.

Displaying port default VLAN IDs (PVIDs)

The output of the **show interfaces brief** command lists the port default VLAN IDs (PVIDs) for each port. PVIDs are displayed as follows:

- For untagged ports, the PVID is the VLAN ID number.
- For dual-mode ports, the PVID is the dual-mode VLAN ID number.
- For tagged ports without dual-mode, the PVID is always Not Applicable (NA).

```
device# show interfaces brief
Port Link      State   Dupl Speed Trunk Tag Pvid
  Pri MAC          Name
    1 Up        Forward Full 1G   None  No   1
      0 0000.00a8.4700 a12345678901
    2 Up        Forward Full 1G   None  Yes  1
      0 0000.00a8.4701
    3 Up        Forward Full 1G   None  Yes NA
      0 0000.00a8.4702
    4 Up        Forward Full 1G   None  Yes NA
      0 0000.00a8.4703
    5 Up        Forward Full 1G   None  No   2
      0 0000.00a8.4704
    6 Down      None     None None  None Yes NA
      0 0000.00a8.4705
    7 Down      None     None None  None Yes  4
      0 0000.00a8.4706
    8 Down      None     None None  None Yes  4
      0 0000.00a8.4707
    9 Down      None     None None  None Yes NA
      0 0000.00a8.4708
   10 Down      None     None None  None Yes NA
      0 0000.00a8.4709
```

Syntax: show interfaces brief [ethernet [slotnum/]portnum [to [slotnum/]portnum [ethernet [slotnum/]portnum...]]]

The *slotnum* parameter is required on chassis devices.

Displaying PVLAN information

To display the PVLAN configuration with respect to the primary VLAN and its associated secondary VLANs and to display the member ports, promiscous ports, and the switch-switch link ports of a PVLAN, enter a command such as the following.

```
device# show pvlan
PVLAN: primary VLAN 100
  Port 1/1/4 1/1/10 1/1/11
Community VLAN 102
  Port 1/1/1 1/1/2 1/1/10 1/1/11
  Promiscuous Port: 1/1/4
  Inter switch link Port: 1/1/10 1/1/11
  BpduGuard enabled Port: 1/1/1 1/1/2
Isolate VLAN 101
  Port 1/1/3 1/1/10 1/1/11
  Promiscuous Port: 1/1/4
  Inter switch link Port: 1/1/10 1/1/11
  BpduGuard enabled Port: 1/1/1 1/1/2
```

Syntax: **show pvlan *vid***

The *vid* variable specifies the VLAN ID of the PVLAN. If the VLAN ID is not specified, the command displays the default VLAN ID.

NOTE

The **show pvlan** command is not supported on software forwarding platforms.
